

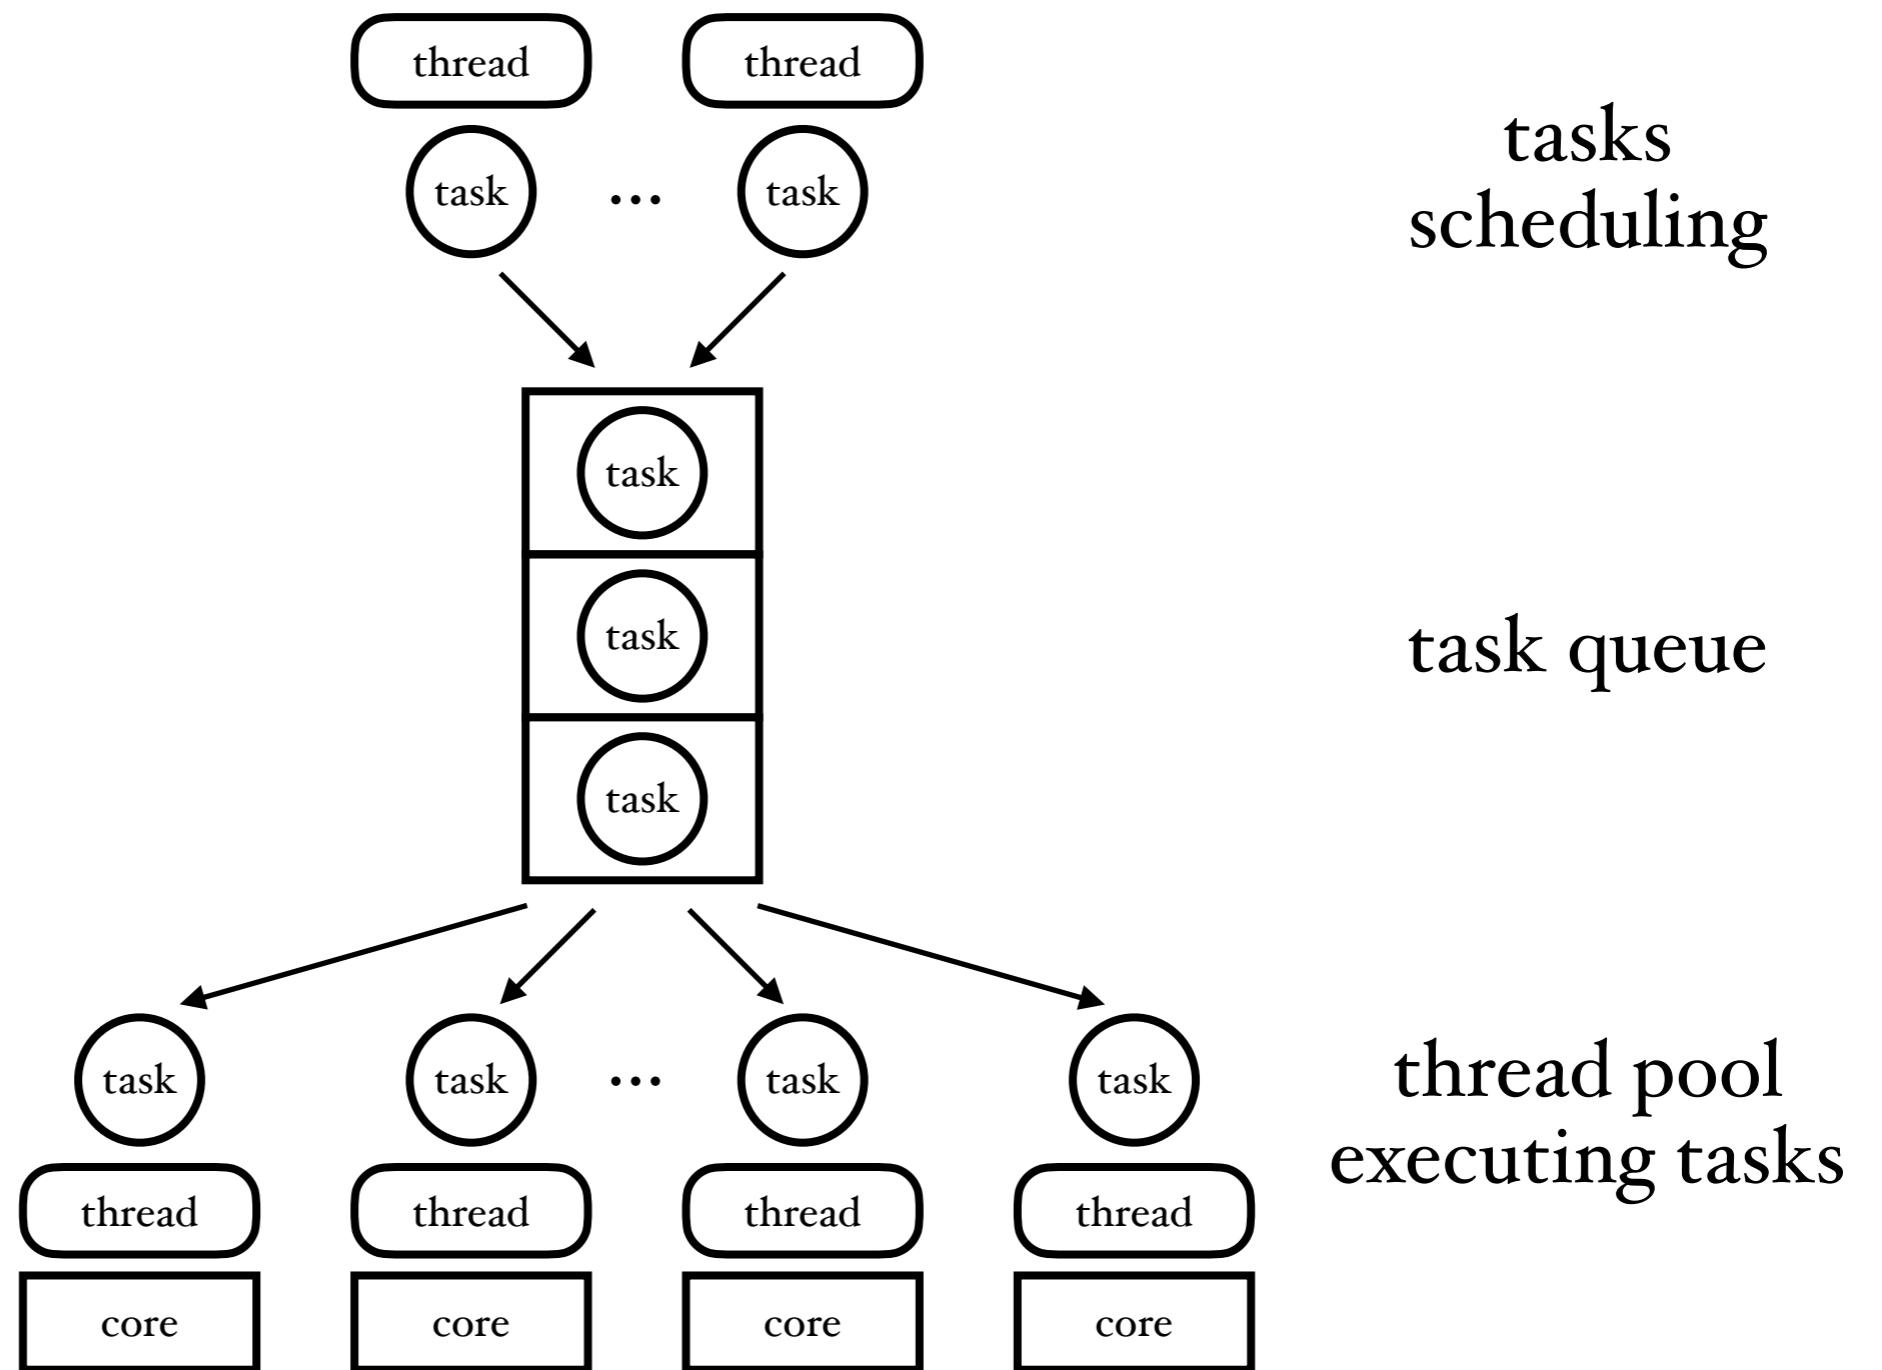
Towards better C++ compilation model

Mateusz Zych

6 July 2020

Async Task Library

Diagram



Async Task Library

Task System - Interface

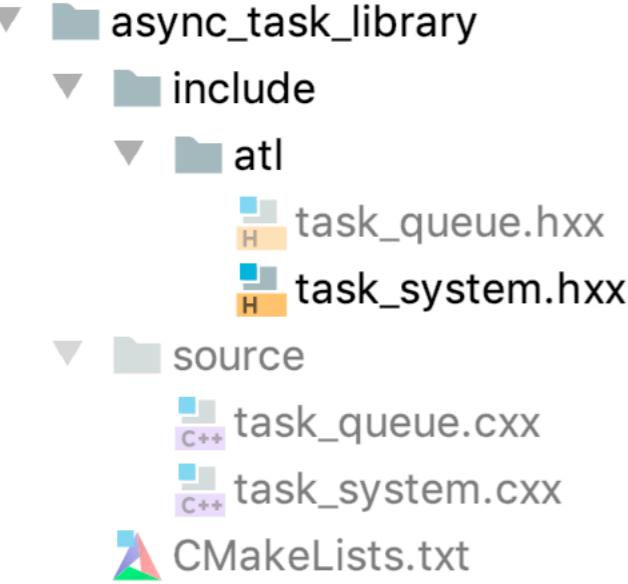
```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

#include ...

namespace atl
{
    class task_system
    {
        ...
    public:
        task_system ();
        ~task_system ();

        template <typename callable>
        auto async (callable&& task) -> void { ... }
    };
}

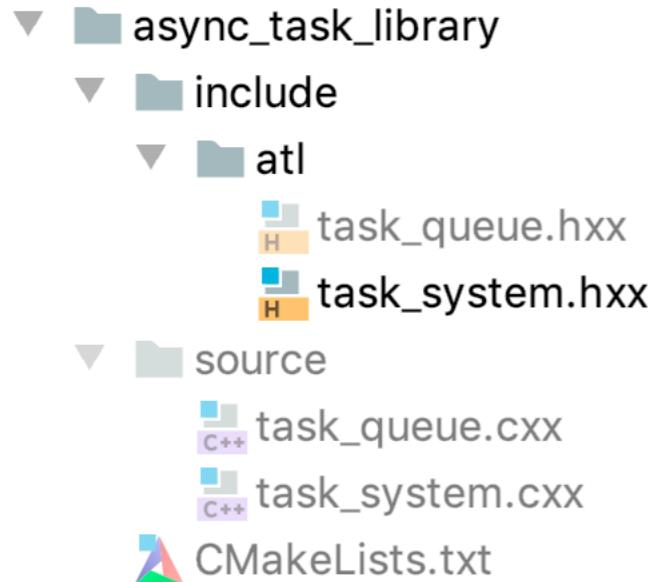
#endif
```



The image shows a project file tree for an Async Task Library. The root folder contains a CMakeLists.txt file. Inside the root are two main directories: 'async_task_library' and 'source'. The 'async_task_library' directory contains an 'include' folder which further contains an 'atl' folder with 'task_queue.hxx' and 'task_system.hxx' files. The 'source' folder contains 'task_queue.cxx' and 'task_system.cxx' files.

Async Task Library

Task System - Interface



```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

#include ...

namespace atl
{
    class task_system
    {
        ...
    public:
        task_system ();
        ~task_system ();

        template <typename callable>
        auto async (callable&& task) -> void { ... }

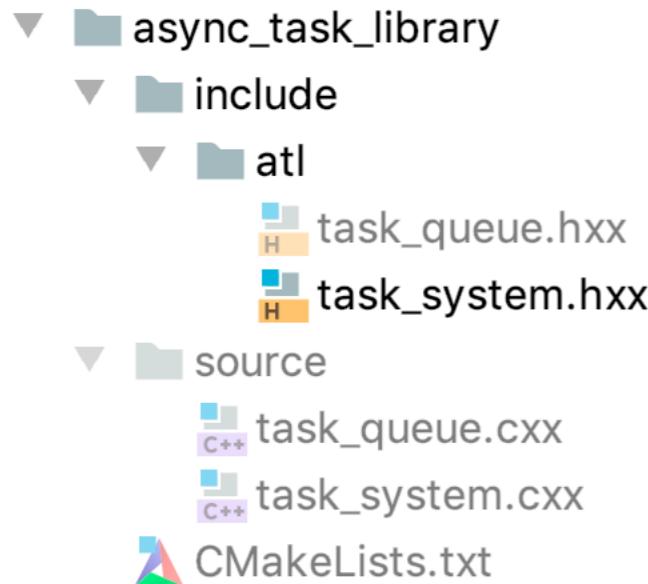
    };
}

#endif
```

task_system can be
constructed and destructed

Async Task Library

Task System - Interface



```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

#include ...

namespace atl
{
    class task_system
    {
        ...
    public:
        task_system();
        ~task_system();

        template <typename callable>
        auto async(callable&& task) -> void { ... }

    };
}

#endif
```

task_system can be constructed and destructed

scheduling a task is done via async() member function

Async Task Library

Task System - Implementation

```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

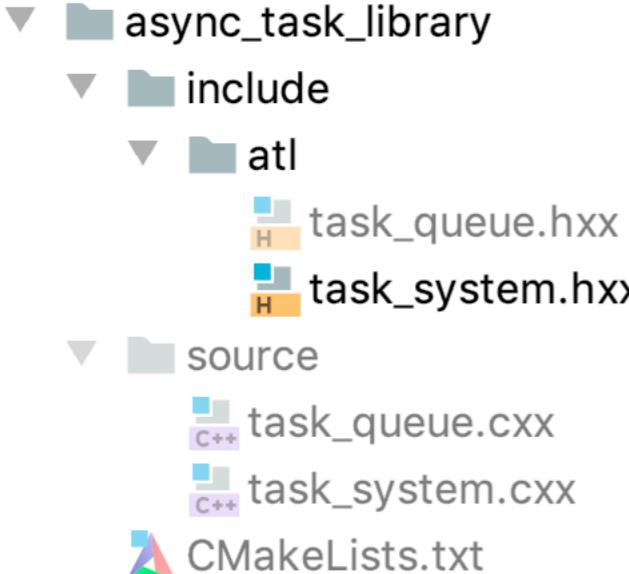
#include <atl/task_queue.hxx>

#include <vector>
#include <thread>

namespace atl
{
    class task_system
    {
        private:
            std::vector<std::thread> threads;
            task_queue queue;

            ...
    };
}

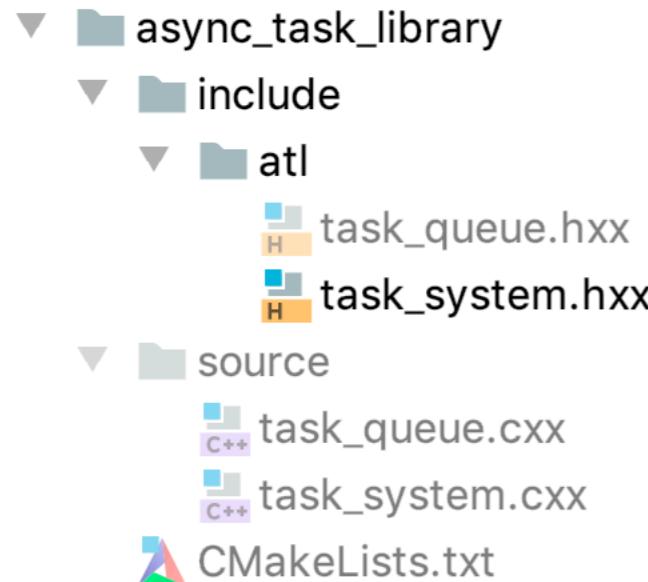
#endif
```



```
async_task_library
├── include
│   └── atl
│       ├── task_queue.hxx
│       └── task_system.hxx
└── source
    ├── task_queue.cxx
    └── task_system.cxx
CMakeLists.txt
```

Async Task Library

Task System - Implementation



```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

#include <atl/task_queue.hxx>

#include <vector>
#include <thread>

namespace atl
{
    class task_system
    {
        private:
            std::vector<std::thread> threads;
            task_queue queue;

            ...
    };
}

#endif
```

thread pool, on which
tasks will be executed

Async Task Library

Task System - Implementation

```
▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    CMakeLists.txt
```

```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

#include <atl/task_queue.hxx>

#include <vector>
#include <thread>

namespace atl
{
    class task_system
    {
    private:
        std::vector<std::thread> threads;
        task_queue queue;
    ...
};

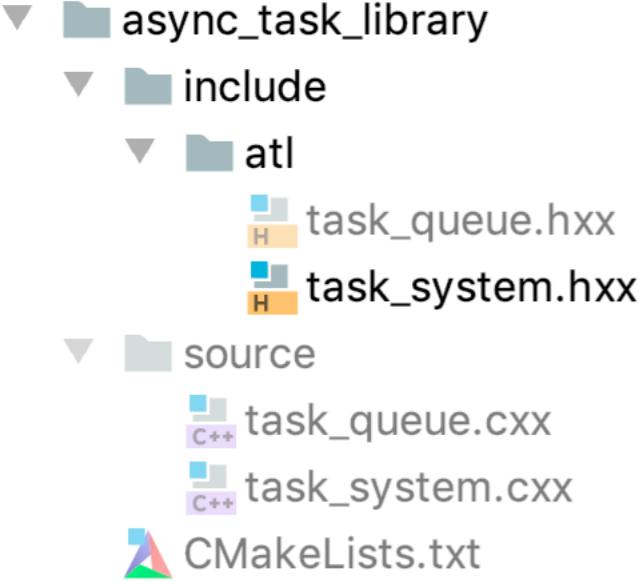
#endif
```

thread pool, on which
tasks will be executed

queue of tasks
to be executed

Async Task Library

Task System - Implementation



```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

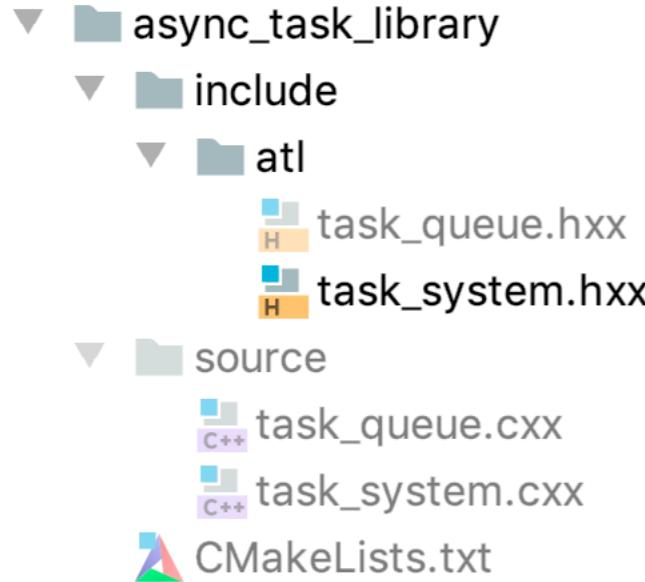
#include ...

namespace atl
{
    class task_system
    {
        ...
        template <typename callable>
        auto async (callable&& task) -> void
        {
            queue.push(std::forward<callable>(task));
        }
    };
}

#endif
```

Async Task Library

Task System - Implementation



```
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

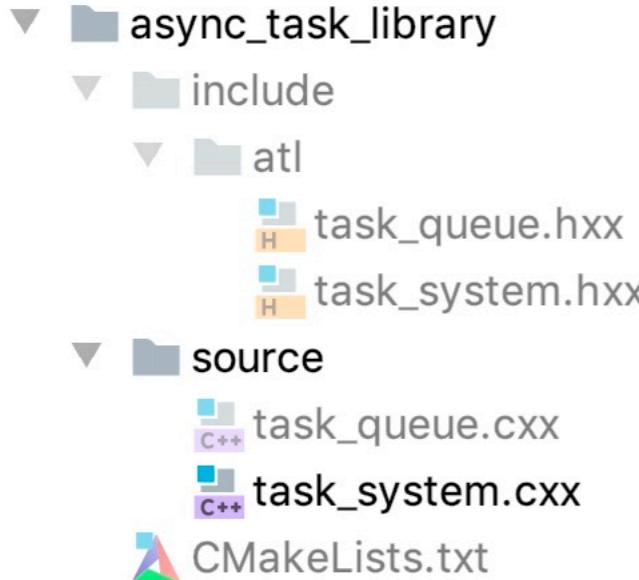
#include ...

namespace atl
{
    class task_system
    {
        ...
        template <typename callable>
        auto async (callable&& task) -> void
        {
            queue.push(std::forward<callable>(task));
        }
    };
#endif
```

A callout box points from the text "pushes it to the queue" to the line "queue.push(std::forward<callable>(task));".

Async Task Library

Task System - Implementation



```
#include <atl/task_system.hxx>

namespace atl
{
    ...

    task_system::task_system()
    {
        auto thread_count = std::thread::hardware_concurrency();

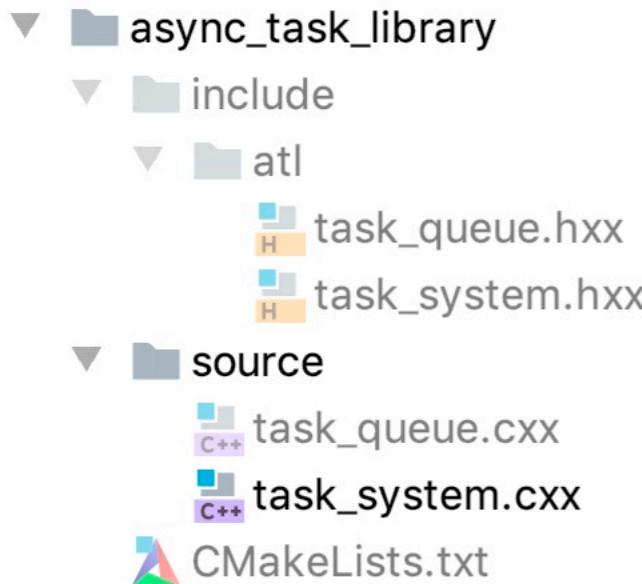
        for (auto i = 0uL; i != thread_count; ++i)
        {
            threads.emplace_back([this] () { run_thread(queue); });
        }
    }

    task_system::~task_system()
    {
        queue.finish();

        for (auto& thread : threads)
        {
            thread.join();
        }
    }
}
```

Async Task Library

Task System - Implementation



```
#include <atl/task_system.hxx>

namespace atl
{
    ...
    task_system::task_system()
    {
        auto thread_count = std::thread::hardware_concurrency();
        for (auto i = 0uL; i != thread_count; ++i)
        {
            threads.emplace_back([this] () { run_thread(queue); });
        }
    }

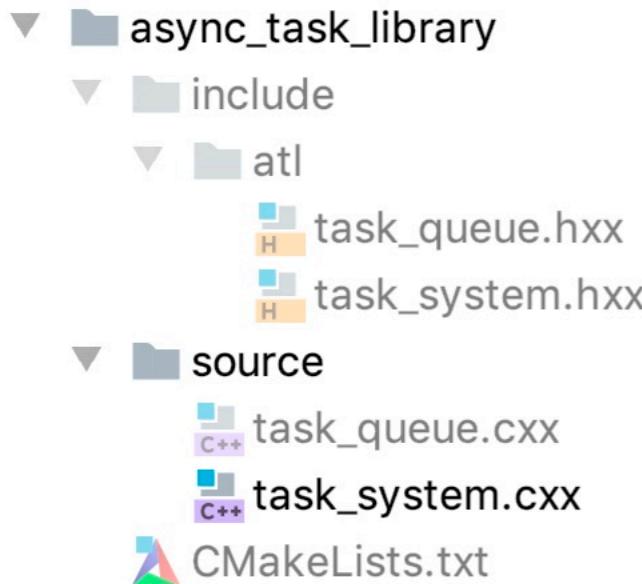
    task_system::~task_system()
    {
        queue.finish();

        for (auto& thread : threads)
        {
            thread.join();
        }
    }
}
```

spawn as many threads,
as HW is capable of running in parallel

Async Task Library

Task System - Implementation



```
#include <atl/task_system.hxx>

namespace atl
{
    ...
    task_system::task_system()
    {
        auto thread_count = std::thread::hardware_concurrency();
        for (auto i = 0uL; i != thread_count; ++i)
        {
            threads.emplace_back([this] () { run_thread(queue); });
        }
    }

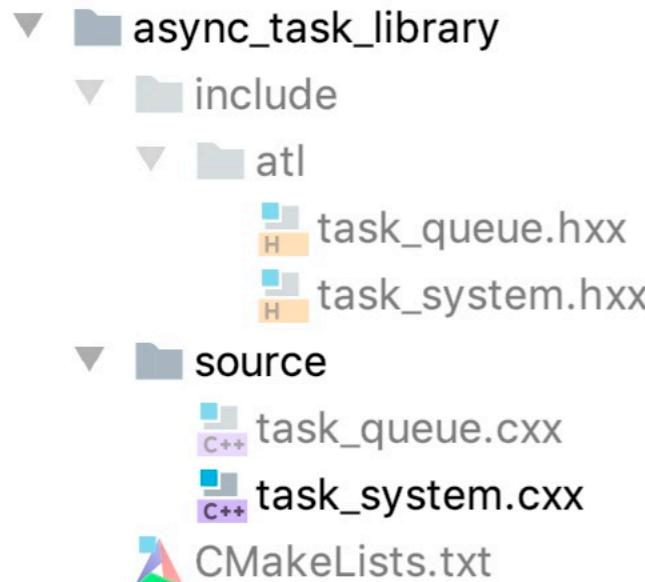
    task_system::~task_system()
    {
        queue.finish();
        for (auto& thread : threads)
        {
            thread.join();
        }
    }
}
```

spawn as many threads,
as HW is capable of running in parallel

signal to the queue,
that all tasks has been scheduled

Async Task Library

Task System - Implementation



```
#include <atl/task_system.hxx>

namespace atl
{
    ...
    task_system::task_system()
    {
        auto thread_count = std::thread::hardware_concurrency();
        for (auto i = 0uL; i != thread_count; ++i)
        {
            threads.emplace_back([this] () { run_thread(queue); });
        }
    }

    task_system::~task_system()
    {
        queue.finish();
        for (auto& thread : threads)
        {
            thread.join();
        }
    }
}
```

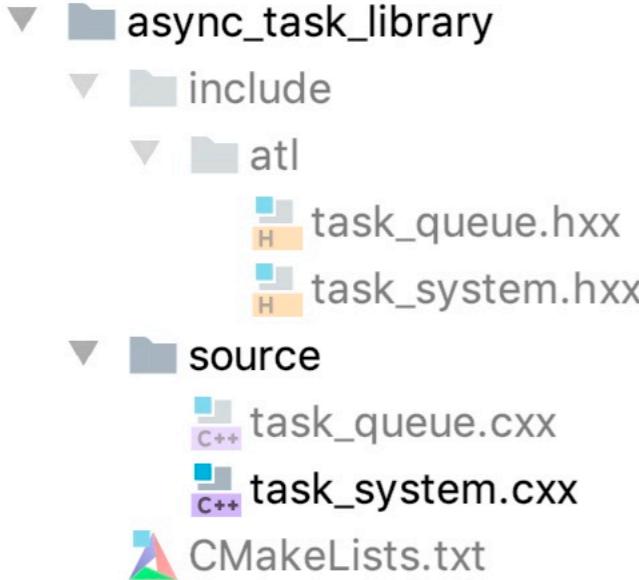
spawn as many threads,
as HW is capable of running in parallel

signal to the queue,
that all tasks has been scheduled

wait on threads,
until they finish all scheduled work

Async Task Library

Task System - Implementation



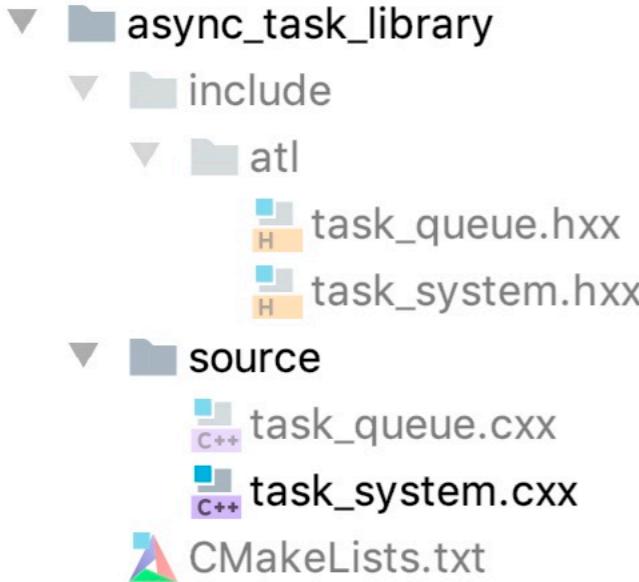
```
#include <atl/task_system.hxx>

namespace atl
{
    namespace
    {
        auto run_thread (task_queue& queue) -> void
        {
            while (true)
            {
                auto task = queue.pop();

                if (task)
                {
                    std::invoke(*task);
                }
                else
                {
                    break;
                }
            }
        }
    }
}
```

Async Task Library

Task System - Implementation



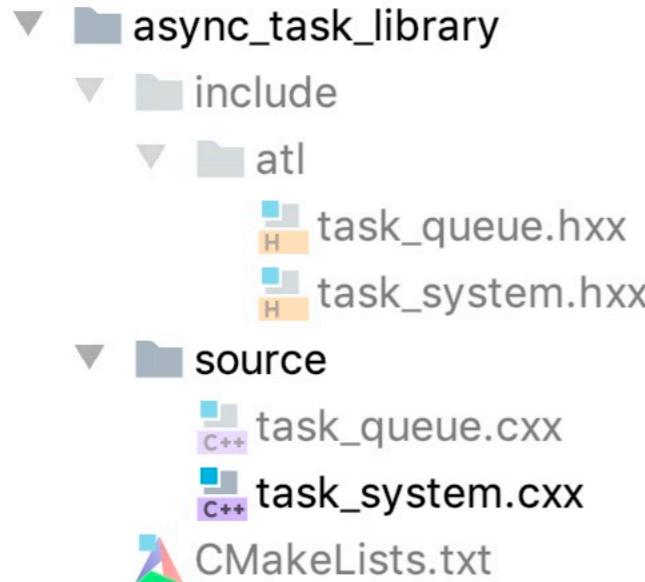
```
#include <atl/task_system.hxx>

namespace atl
{
    namespace
    {
        auto run_thread (task_queue& queue) -> void
        {
            while (true)
            {
                auto task = queue.pop();
                if (task)
                {
                    std::invoke(*task);
                }
                else
                {
                    break;
                }
            }
        }
    ...
}
```

threads pop() tasks from the queue and execute them

Async Task Library

Task System - Implementation



```
#include <atl/task_system.hxx>

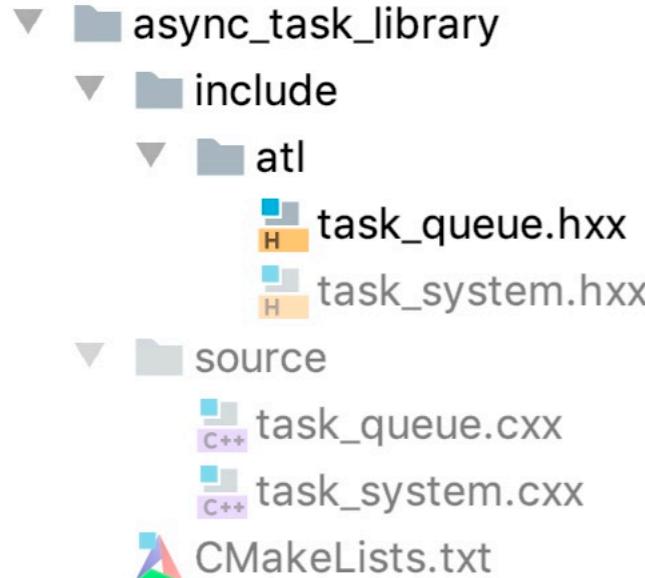
namespace atl
{
    namespace
    {
        auto run_thread (task_queue& queue) -> void
        {
            while (true)
            {
                auto task = queue.pop();
                if (task)
                {
                    std::invoke(*task);
                }
                else
                {
                    break;
                }
            }
            ...
        }
    }
}
```

threads pop() tasks from the queue and execute them

threads exit when all tasks has been executed, i.e., there are no tasks to pop()

Async Task Library

Task Queue - Interface



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <functional>
#include <optional>

namespace atl
{
    class task_queue
    {
        ...

        public:
            auto finish () -> void;

            auto pop () -> std::optional<std::function<void()>>;

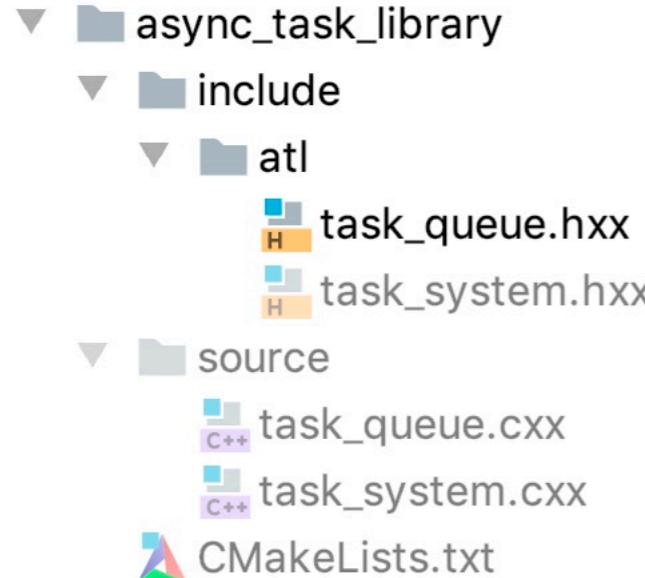
            template <typename callable>
            auto push (callable&& task) -> void { ... }

    };
}

#endif
```

Async Task Library

Task Queue - Interface



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <functional>
#include <optional>

namespace atl
{
    class task_queue
    {
        ...

        public:
            auto finish () -> void;

            auto pop () -> std::optional<std::function<void()>>;

            template <typename callable>
            auto push (callable&& task) -> void { ... }

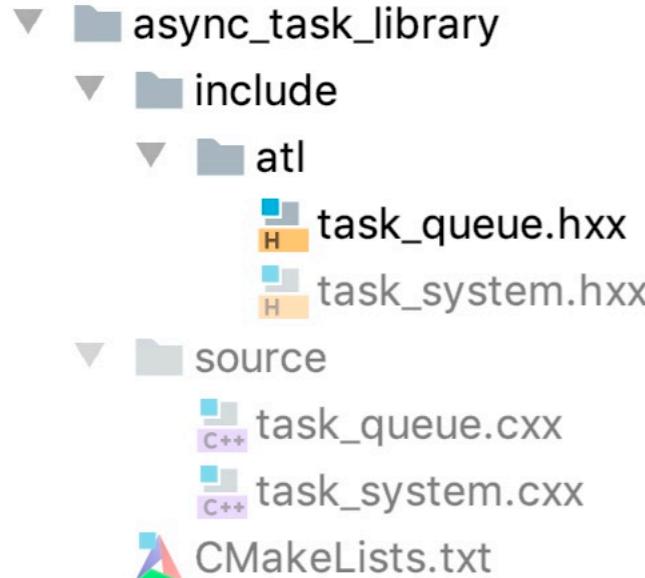
    };
}

#endif
```

A callout box points from the word 'push' in the code to the explanatory text 'push() task to the queue'.

Async Task Library

Task Queue - Interface



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <functional>
#include <optional>

namespace atl
{
    class task_queue
    {
        ...

        public:
            auto finish () -> void;

            auto pop () -> std::optional<std::function<void()>>;

            template <typename callable>
            auto push (callable&& task) -> void { ... }

    };
}

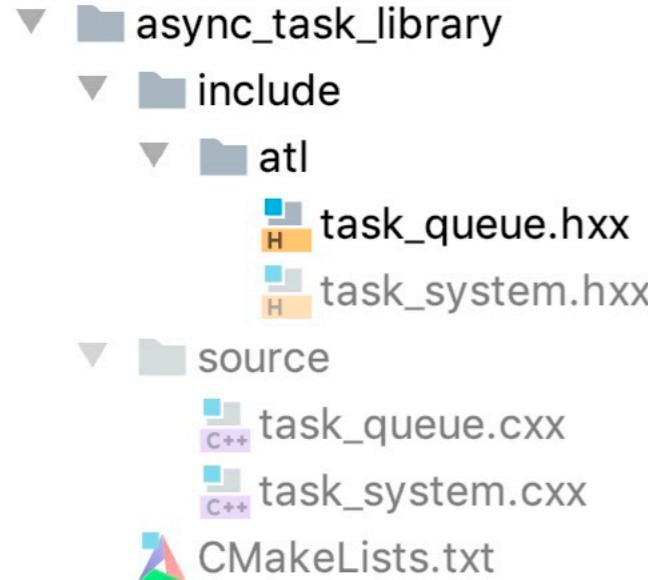
#endif
```

push() task to the queue

pop() task from the queue, if available

Async Task Library

Task Queue - Interface



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <functional>
#include <optional>

namespace atl
{
    class task_queue
    {
        ...
    public:
        auto finish () -> void;

        auto pop () -> std::optional<std::function<void()>>;
        template <typename callable>
        auto push (callable&& task) -> void { ... }
    };
}

#endif
```

signal to the queue that no more tasks will be scheduled, i.e., push()'ed to the queue

push() task to the queue

pop() task from the queue, if available

Async Task Library

Task Queue - Implementation

```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    CMakeLists.txt

#include <deque>
#include <mutex>
#include <condition_variable>
#include <functional>

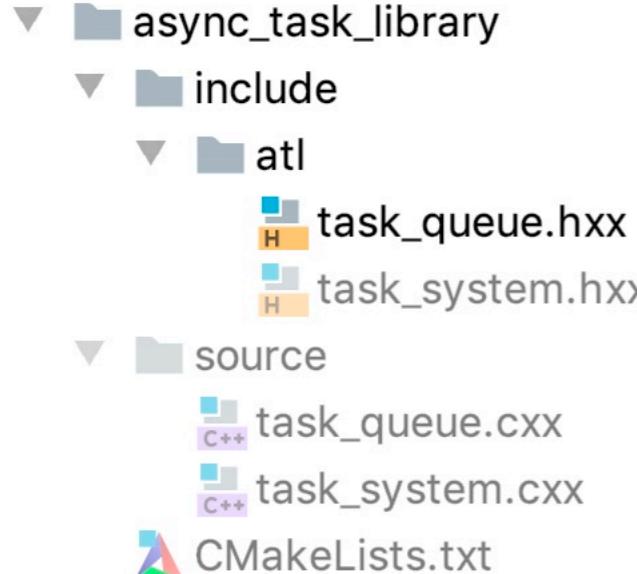
namespace atl
{
    class task_queue
    {
    private:
        std::deque<std::function<void()>> tasks;
        std::mutex mutex;
        std::condition_variable ready;
        bool done { false };

        ...
    };
}

#endif
```

Async Task Library

Task Queue - Implementation



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    └─ CMakeLists.txt
```

```
#include <deque>
#include <mutex>
#include <condition_variable>
#include <functional>

namespace atl
{
    class task_queue
    {
    private:
        std::deque<std::function<void()>> tasks;
        std::mutex mutex;
        std::condition_variable ready;
        bool done { false };

        ...
    };
}

#endif
```

A callout bubble points from the word "tasks" in the code to the variable `tasks` in the class definition, which is annotated with the text "tasks to be executed".

Async Task Library

Task Queue - Implementation

```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    CMakeLists.txt

#include <deque>
#include <mutex>
#include <condition_variable>
#include <functional>

namespace atl
{
    class task_queue
    {
    private:
        std::deque<std::function<void()>> tasks;
        std::mutex
        std::condition_variable
        bool
            ...
    };
}

#endif
```

tasks to be executed

flag done defining whether more tasks will be pushed to the queue

Async Task Library

Task Queue - Implementation

```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    CMakeLists.txt

#include <deque>
#include <mutex>
#include <condition_variable>
#include <functional>

namespace atl
{
    class task_queue
    {
        private:
            std::deque<std::function<void()>> tasks;
            std::mutex mutex;
            std::condition_variable ready;
            bool done { false };

            ...
    };
}

#endif
```

mutex protecting shared state, i.e., task queue & done flag

Async Task Library

Task Queue - Implementation

```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    CMakeLists.txt

#include <deque>
#include <mutex>
#include <condition_variable>
#include <functional>

namespace atl
{
    class task_queue
    {
    private:
        std::deque<std::function<void()>> tasks;
        std::mutex mutex;
        std::condition_variable ready;
        bool done { false };

        ...
    };
}

#endif
```

mutex protecting shared state, i.e., task queue & done flag

condition variable specifying whether a task is ready to be pop()ed from the queue

Async Task Library

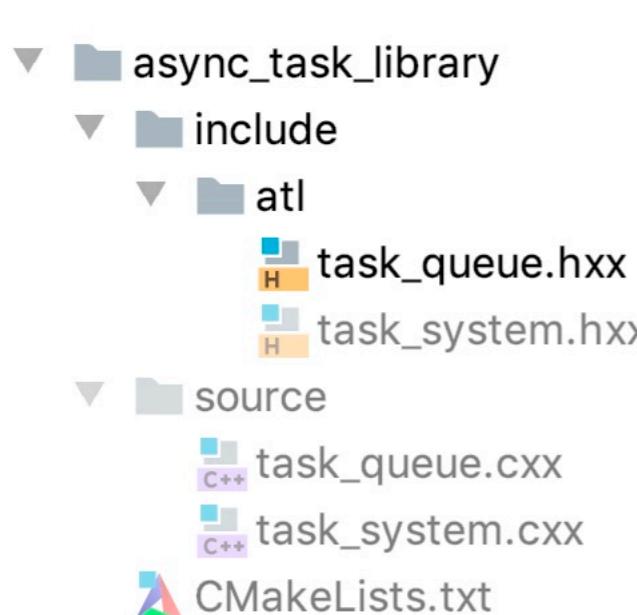
Task Queue - Implementation

```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <utility>

namespace atl
{
    class task_queue
    {
        ...
        template <typename callable>
        auto push (callable&& task) -> void
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            tasks.emplace_back(std::forward<callable>(task));
        }
        ready.notify_one();
    };
}

#endif
```

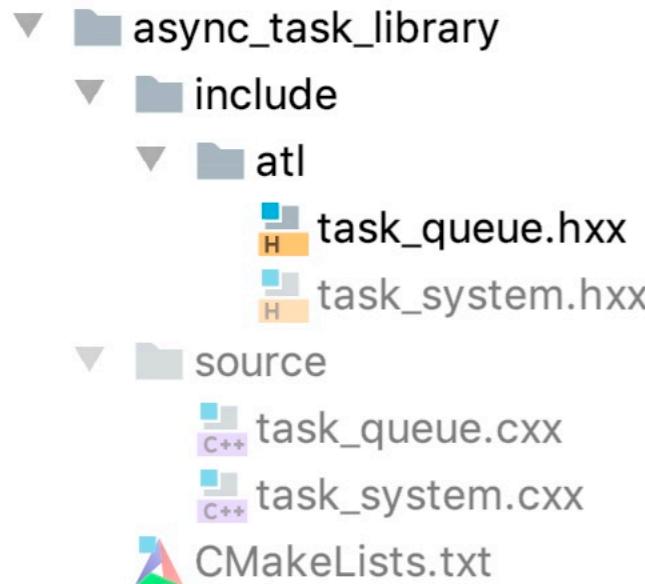


```
async_task_library
├── include
│   └── atl
│       ├── task_queue.hxx
│       └── task_system.hxx
└── source
    ├── task_queue.cxx
    └── task_system.cxx
CMakeLists.txt
```

The image shows a project file tree for the 'async_task_library'. It contains two main directories: 'include' and 'source'. The 'include' directory contains an 'atl' subdirectory with two header files: 'task_queue.hxx' and 'task_system.hxx'. The 'source' directory contains two C++ source files: 'task_queue.cxx' and 'task_system.cxx'. A 'CMakeLists.txt' file is also present at the root level.

Async Task Library

Task Queue - Implementation



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <utility>

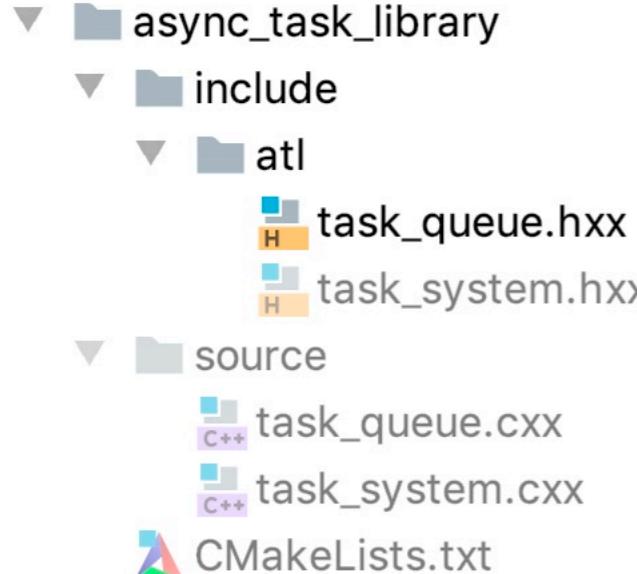
namespace atl
{
    class task_queue
    {
        ...
        template <typename callable>
        auto push (callable&& task) -> void
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            tasks.emplace_back(std::forward<callable>(task));
        }
        ready.notify_one();
    };
}

#endif
```

A callout bubble points to the line `auto lock = std::lock_guard<std::mutex> { mutex };` with the text "enter critical section".

Async Task Library

Task Queue - Implementation



```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <utility>

namespace atl
{
    class task_queue
    {
        ...
        template <typename callable>
        auto push (callable&& task) -> void
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            tasks.emplace_back(std::forward<callable>(task));
        }
        ready.notify_one();
    };
}

#endif
```

The image shows a file tree for the 'async_task_library' project. It includes an 'include' directory containing 'atl' and header files like 'task_queue.hxx' and 'task_system.hxx'. The 'source' directory contains implementation files 'task_queue.cxx' and 'task_system.cxx', along with a 'CMakeLists.txt' file.

Annotations highlight specific code sections:

- A callout box points to the 'push' function with the text "enter critical section".
- A callout box points to the 'tasks.emplace_back' line with the text "put task at the end of the queue".

Async Task Library

Task Queue - Implementation

```
#ifndef ATL_TASK_QUEUE
#define ATL_TASK_QUEUE

#include ...
#include <utility>

namespace atl
{
    class task_queue
    {
        ...
        template <typename callable>
        auto push (callable&& task) -> void
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            tasks.emplace_back(std::forward<callable>(task));
        }
        ready.notify_one();
    };
}

#endif
```

enter critical section

put task at the end of the queue

notify one thread waiting inside pop()
that a task is ready to be executed

The diagram illustrates the file structure of the Async Task Library. It shows a folder named 'async_task_library' containing 'include' and 'source' subfolders. Inside 'include', there is an 'atl' folder containing 'task_queue.hxx' and 'task_system.hxx'. Inside 'source', there are 'task_queue.cxx' and 'task_system.cxx' files, along with a 'CMakeLists.txt' file.

The code snippet shows the implementation of the 'push' method in 'task_queue.hxx'. It includes necessary headers, defines a namespace 'atl', and implements a class 'task_queue'. The 'push' method takes a callable object by rvalue reference, enters a critical section using a lock guard, adds the task to a vector ('tasks'), and then notifies one thread waiting inside the 'pop' method that a task is ready to be executed.

Async Task Library

Task Queue - Implementation

```
#include <atl/task_queue.hxx>

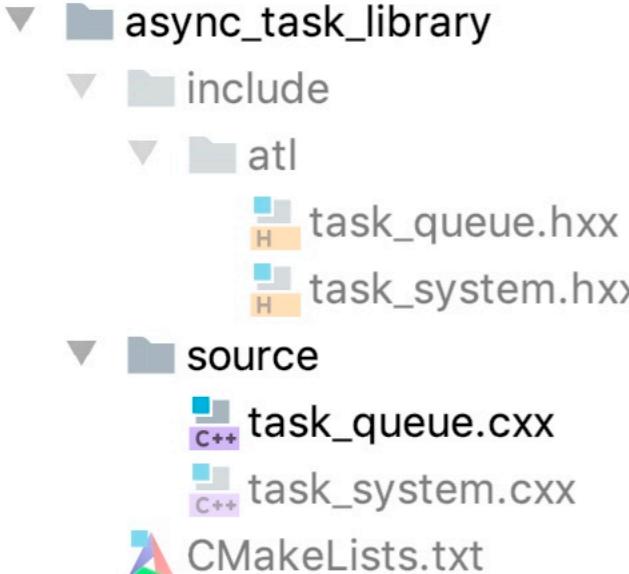
namespace atl
{
    ...

    auto task_queue::pop () -> std::optional<std::function<void()>>
    {
        auto lock = std::unique_lock<std::mutex> { mutex };

        ready.wait(lock, [&] () { return !tasks.empty() || done; });

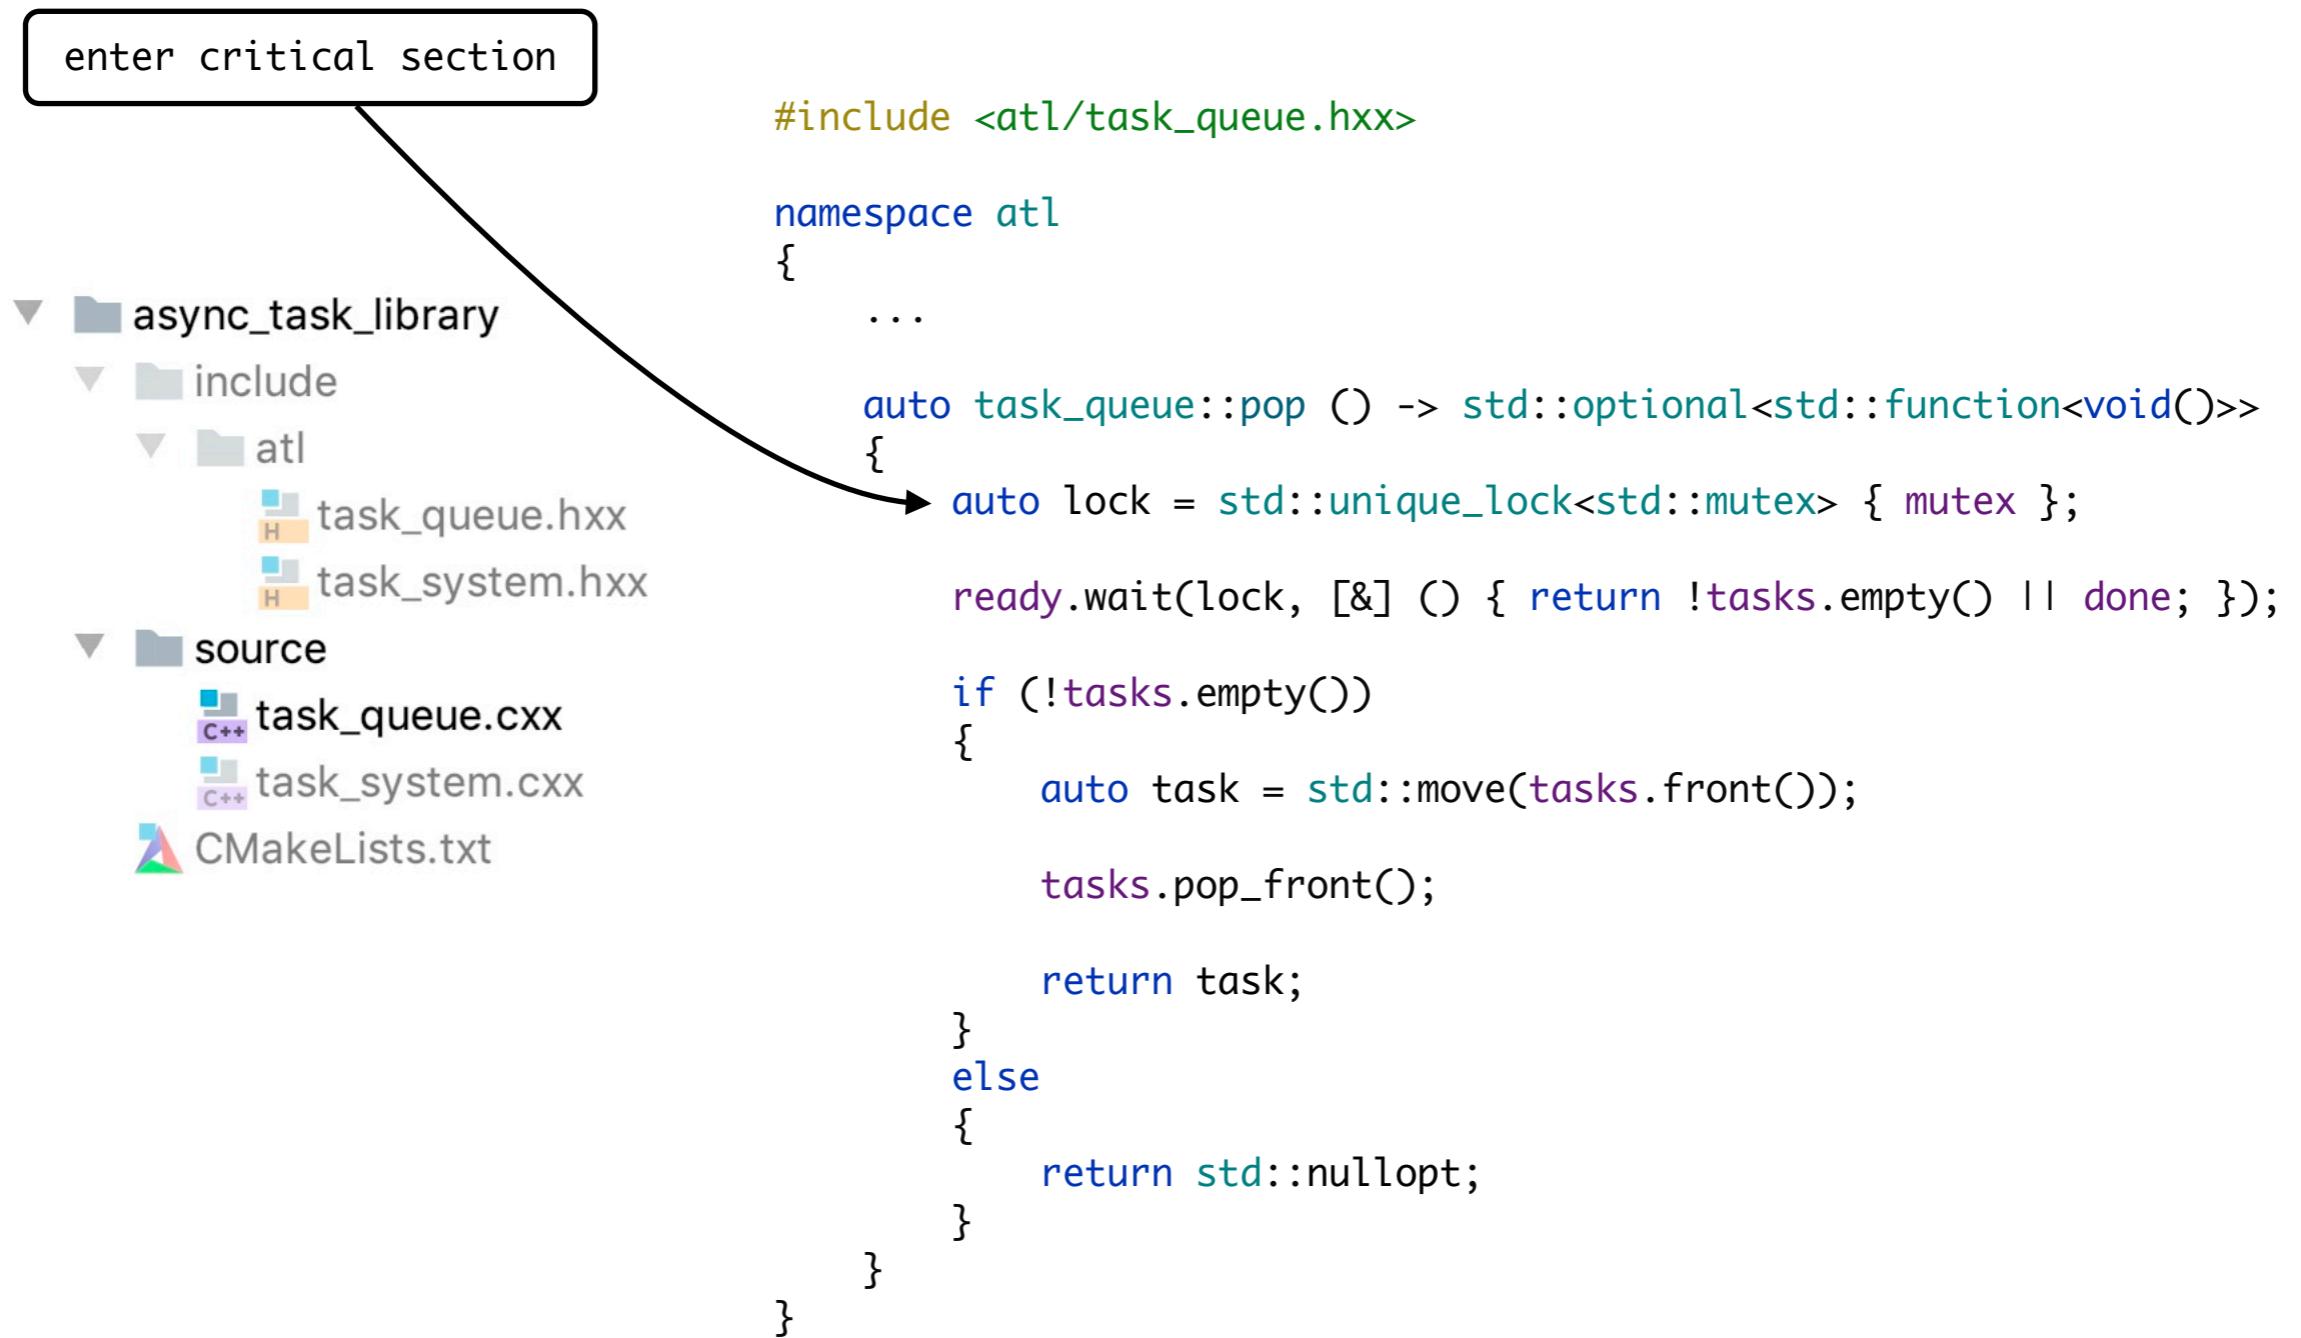
        if (!tasks.empty())
        {
            auto task = std::move(tasks.front());
            tasks.pop_front();

            return task;
        }
        else
        {
            return std::nullopt;
        }
    }
}
```



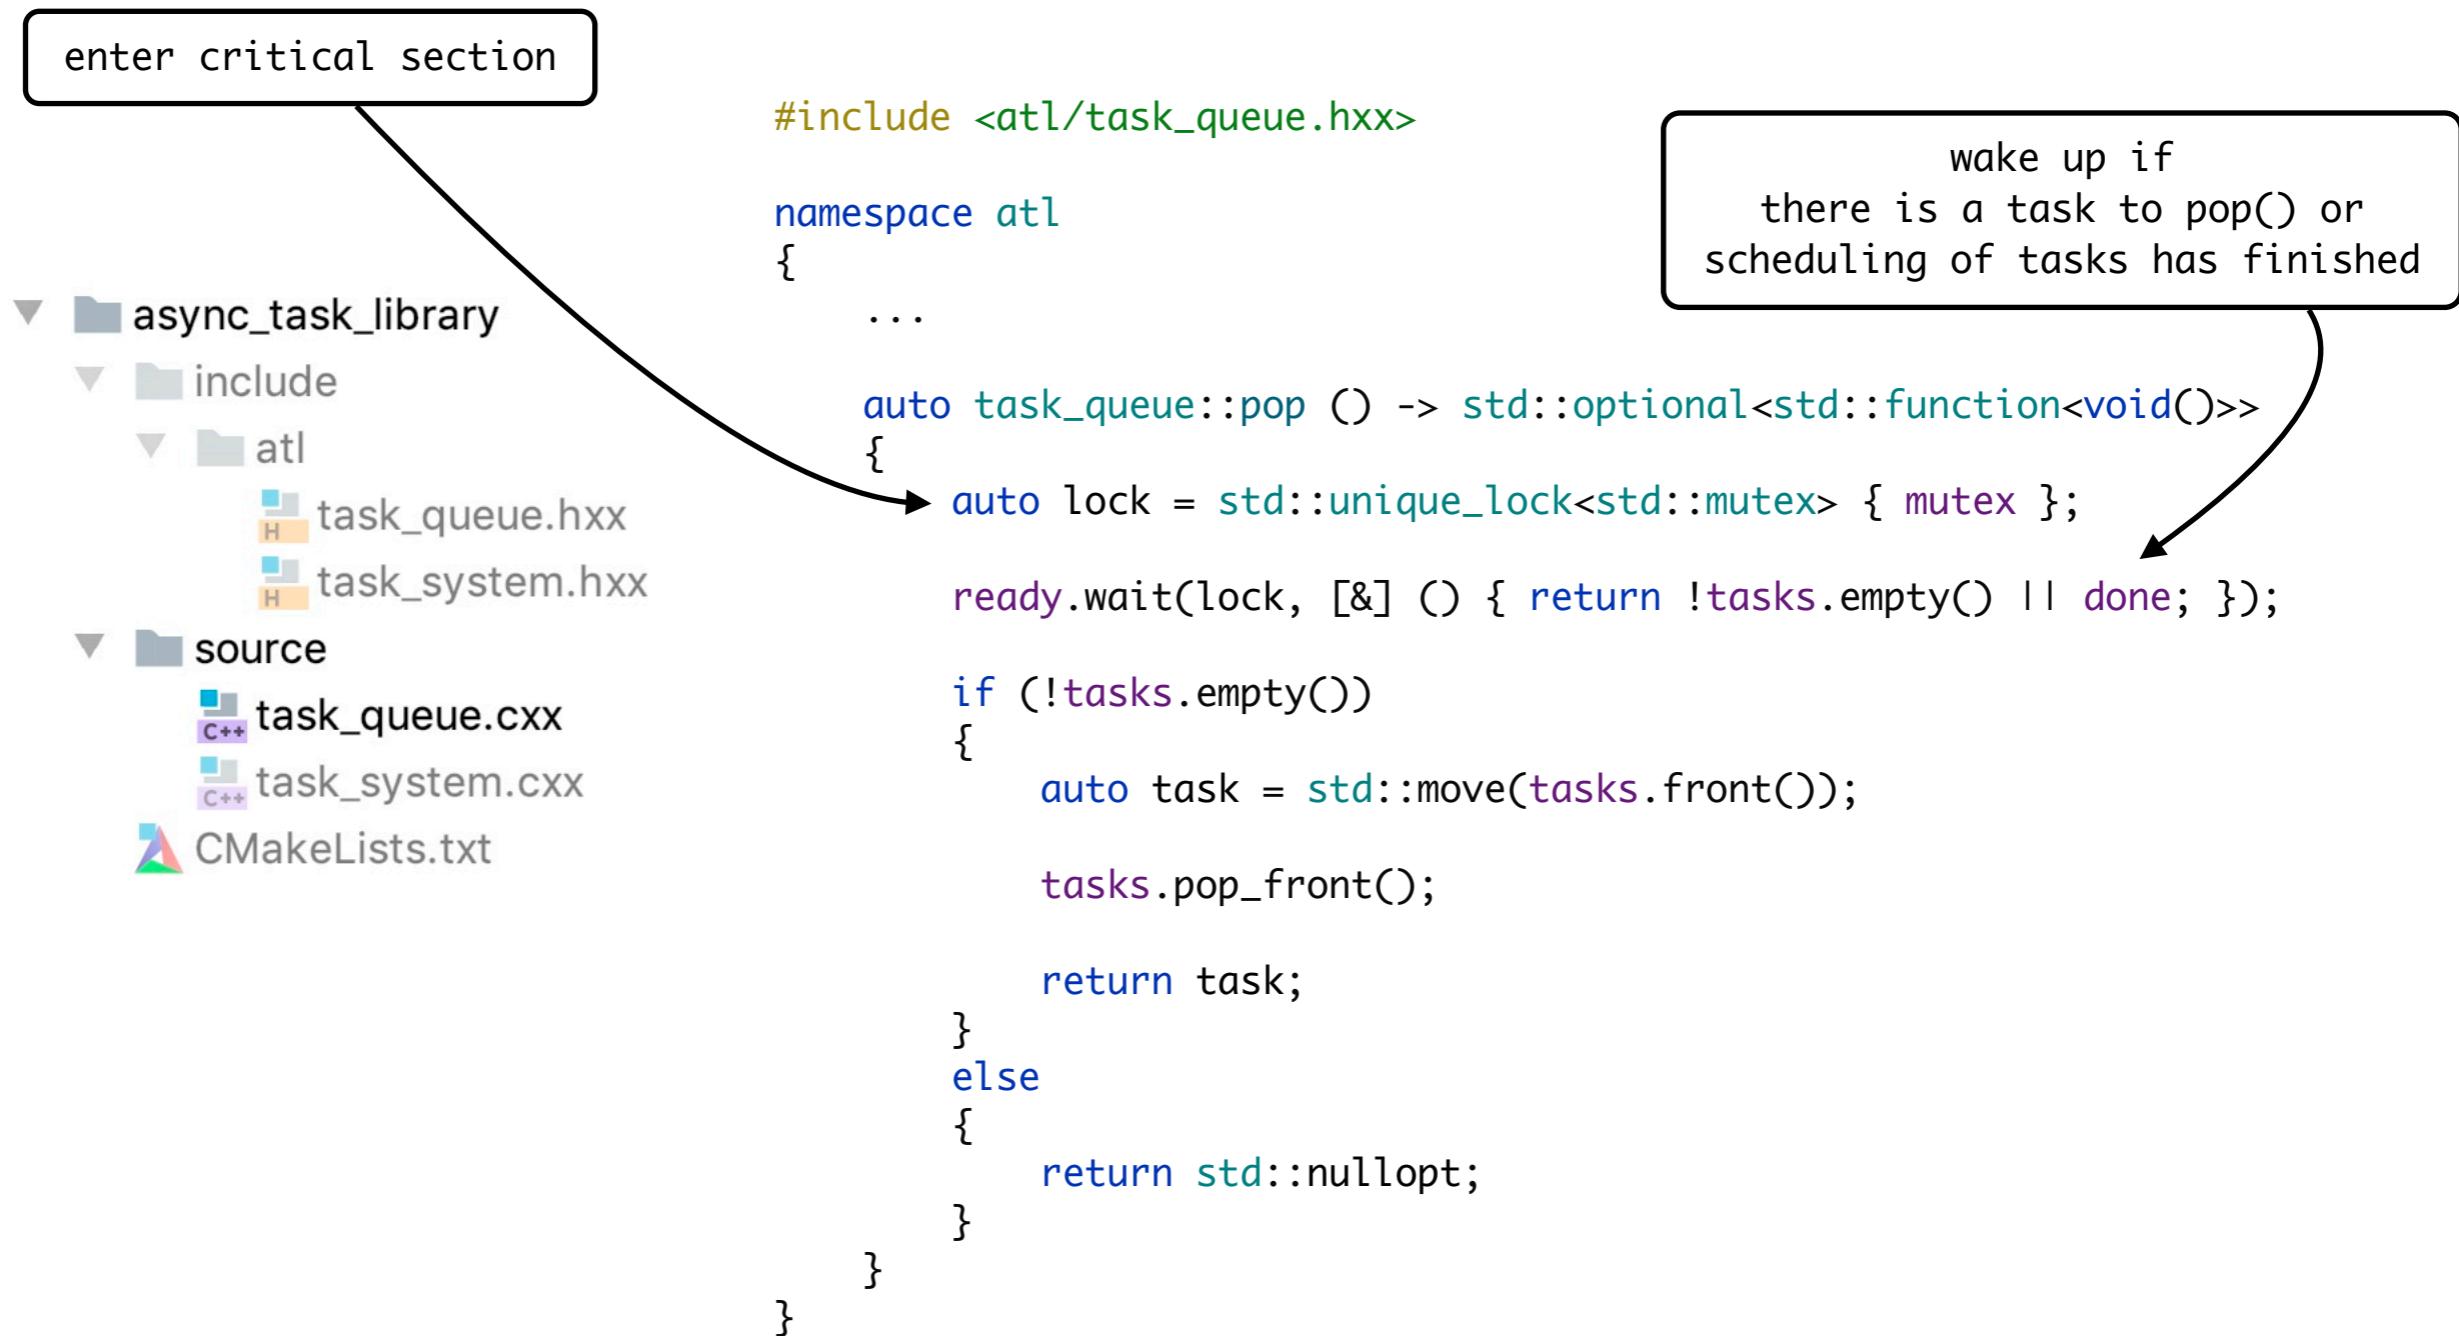
Async Task Library

Task Queue - Implementation



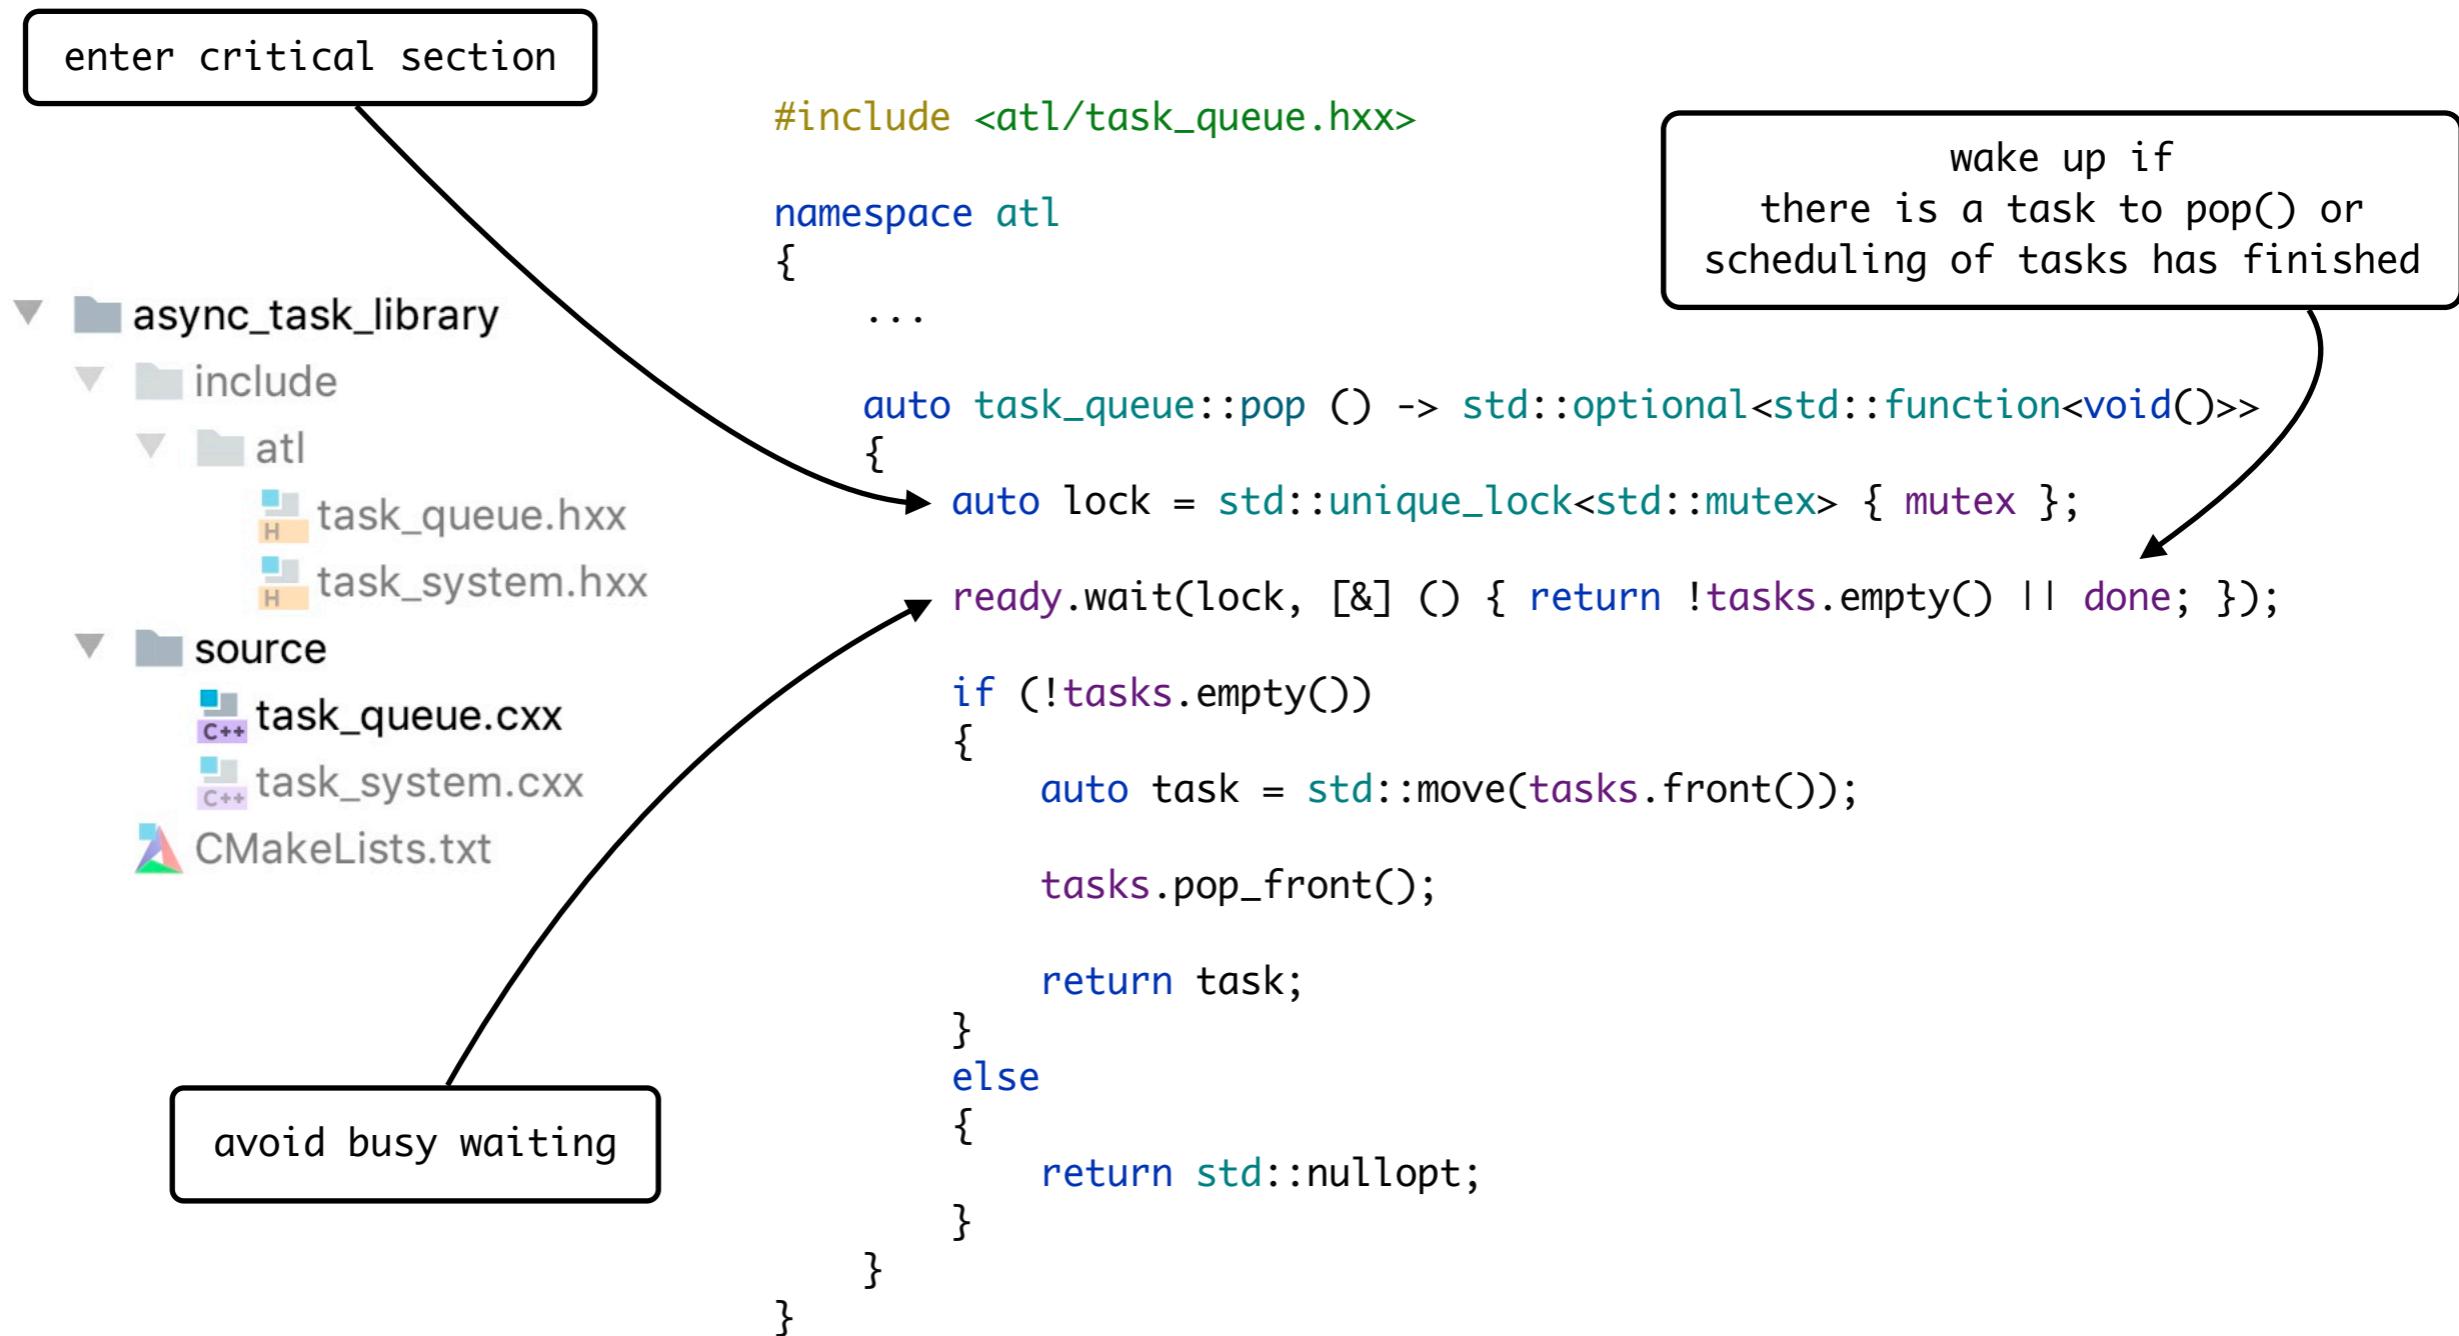
Async Task Library

Task Queue - Implementation



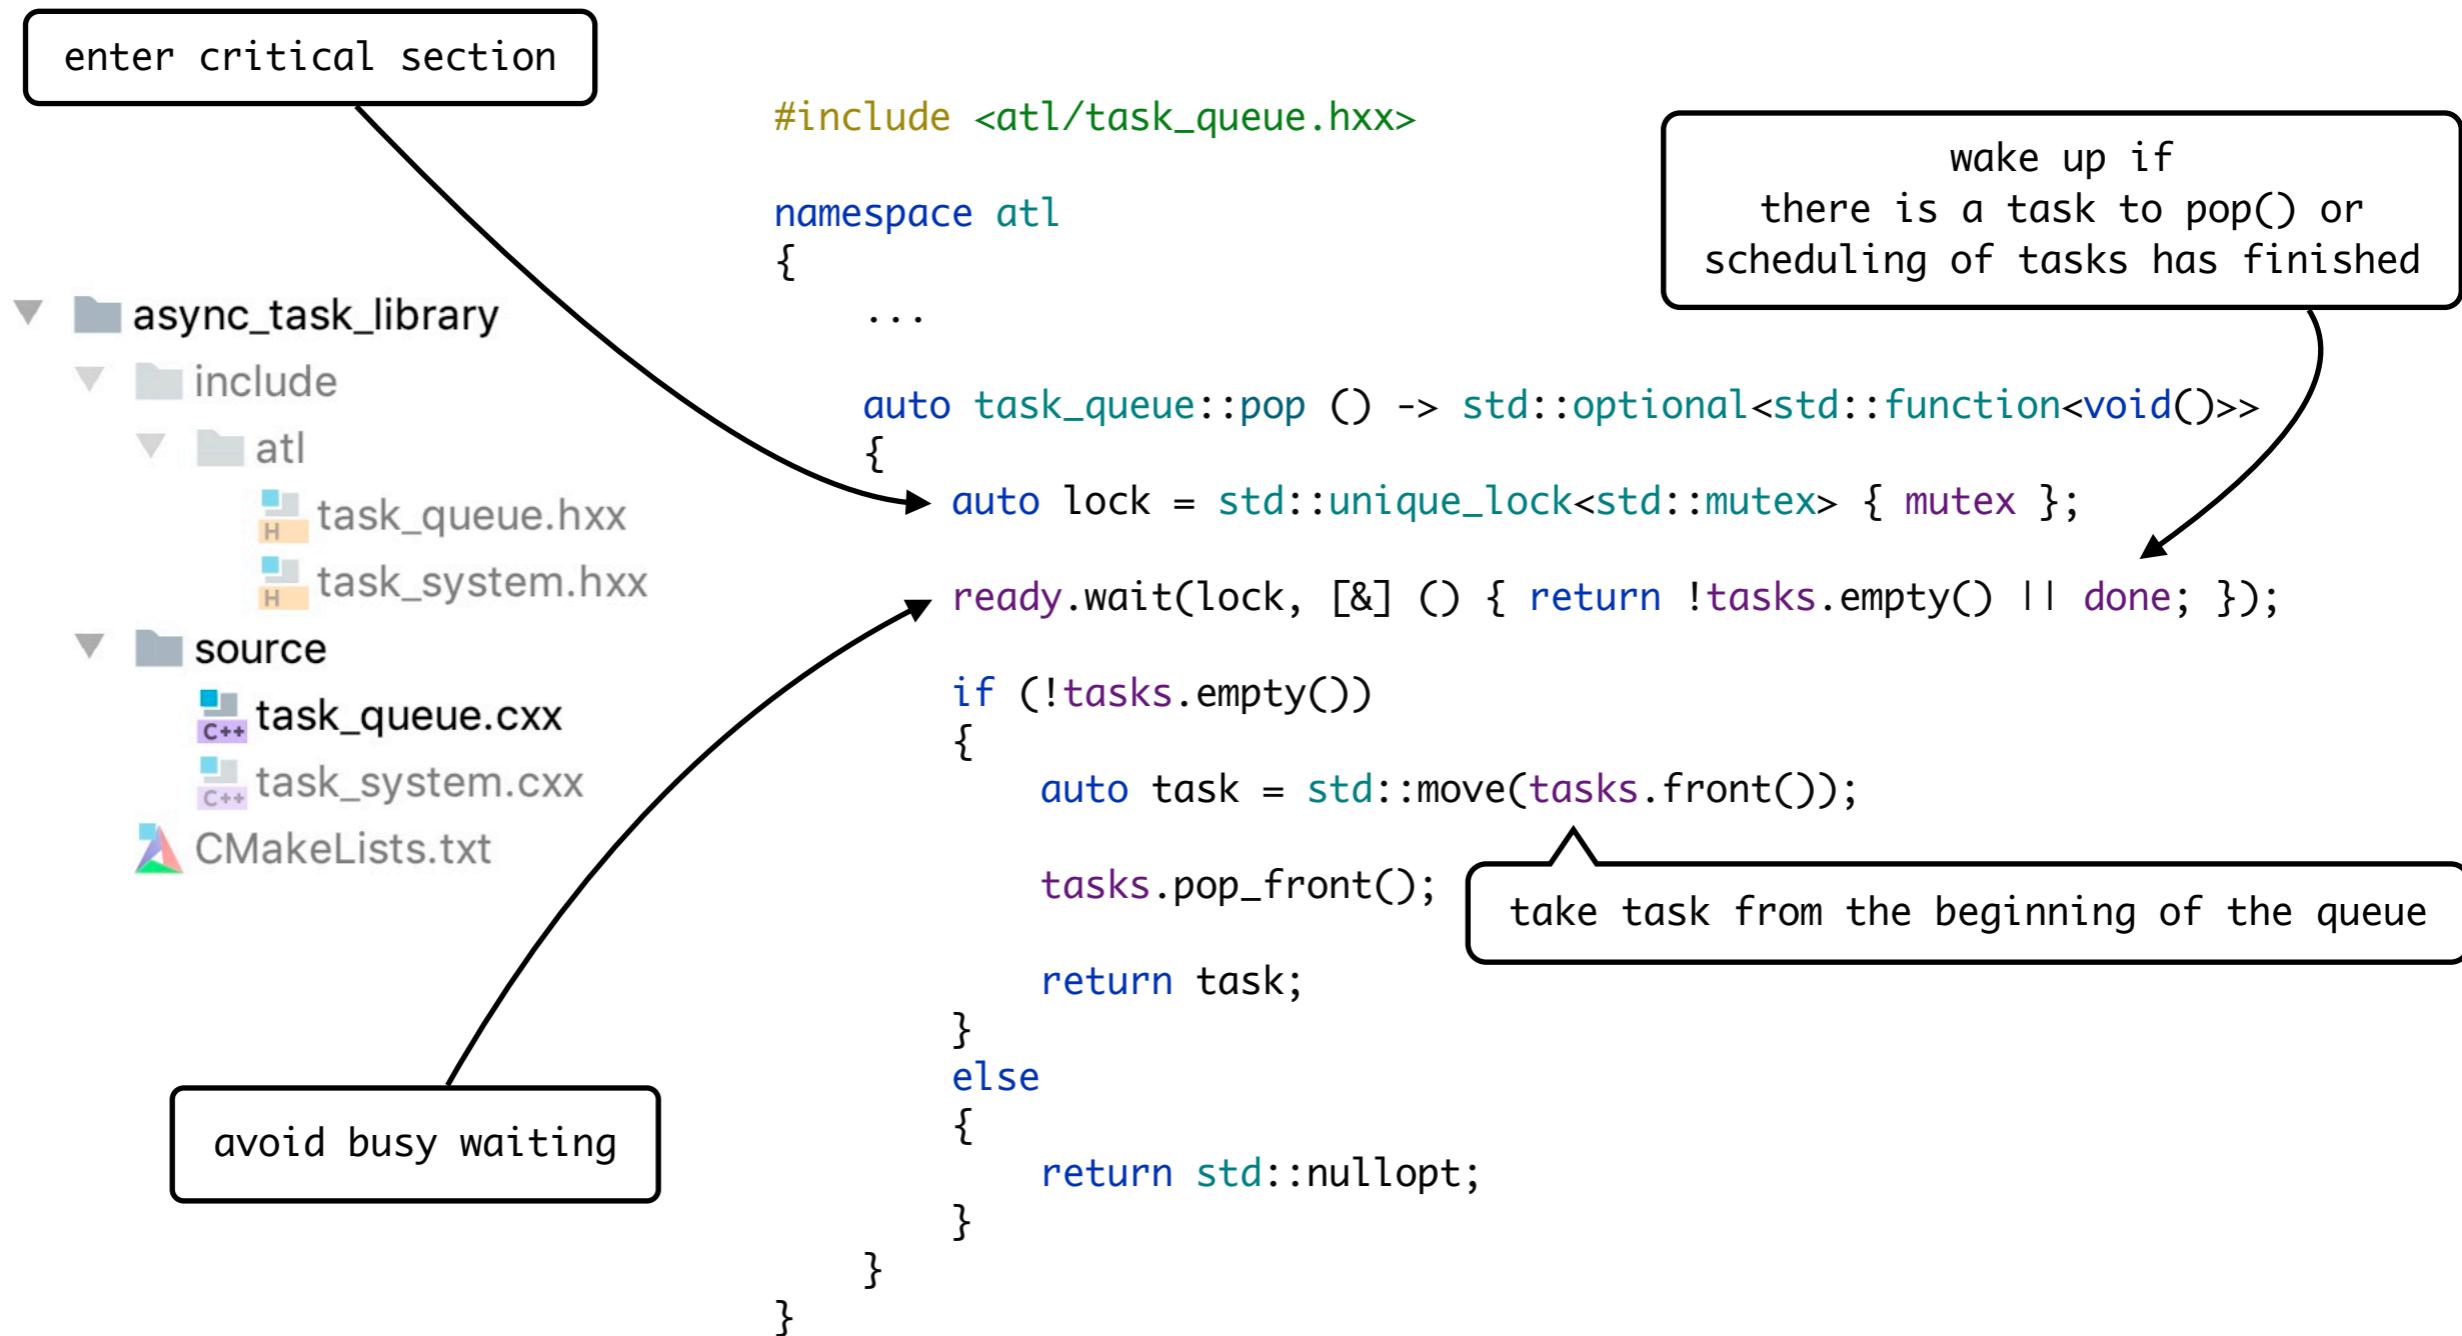
Async Task Library

Task Queue - Implementation



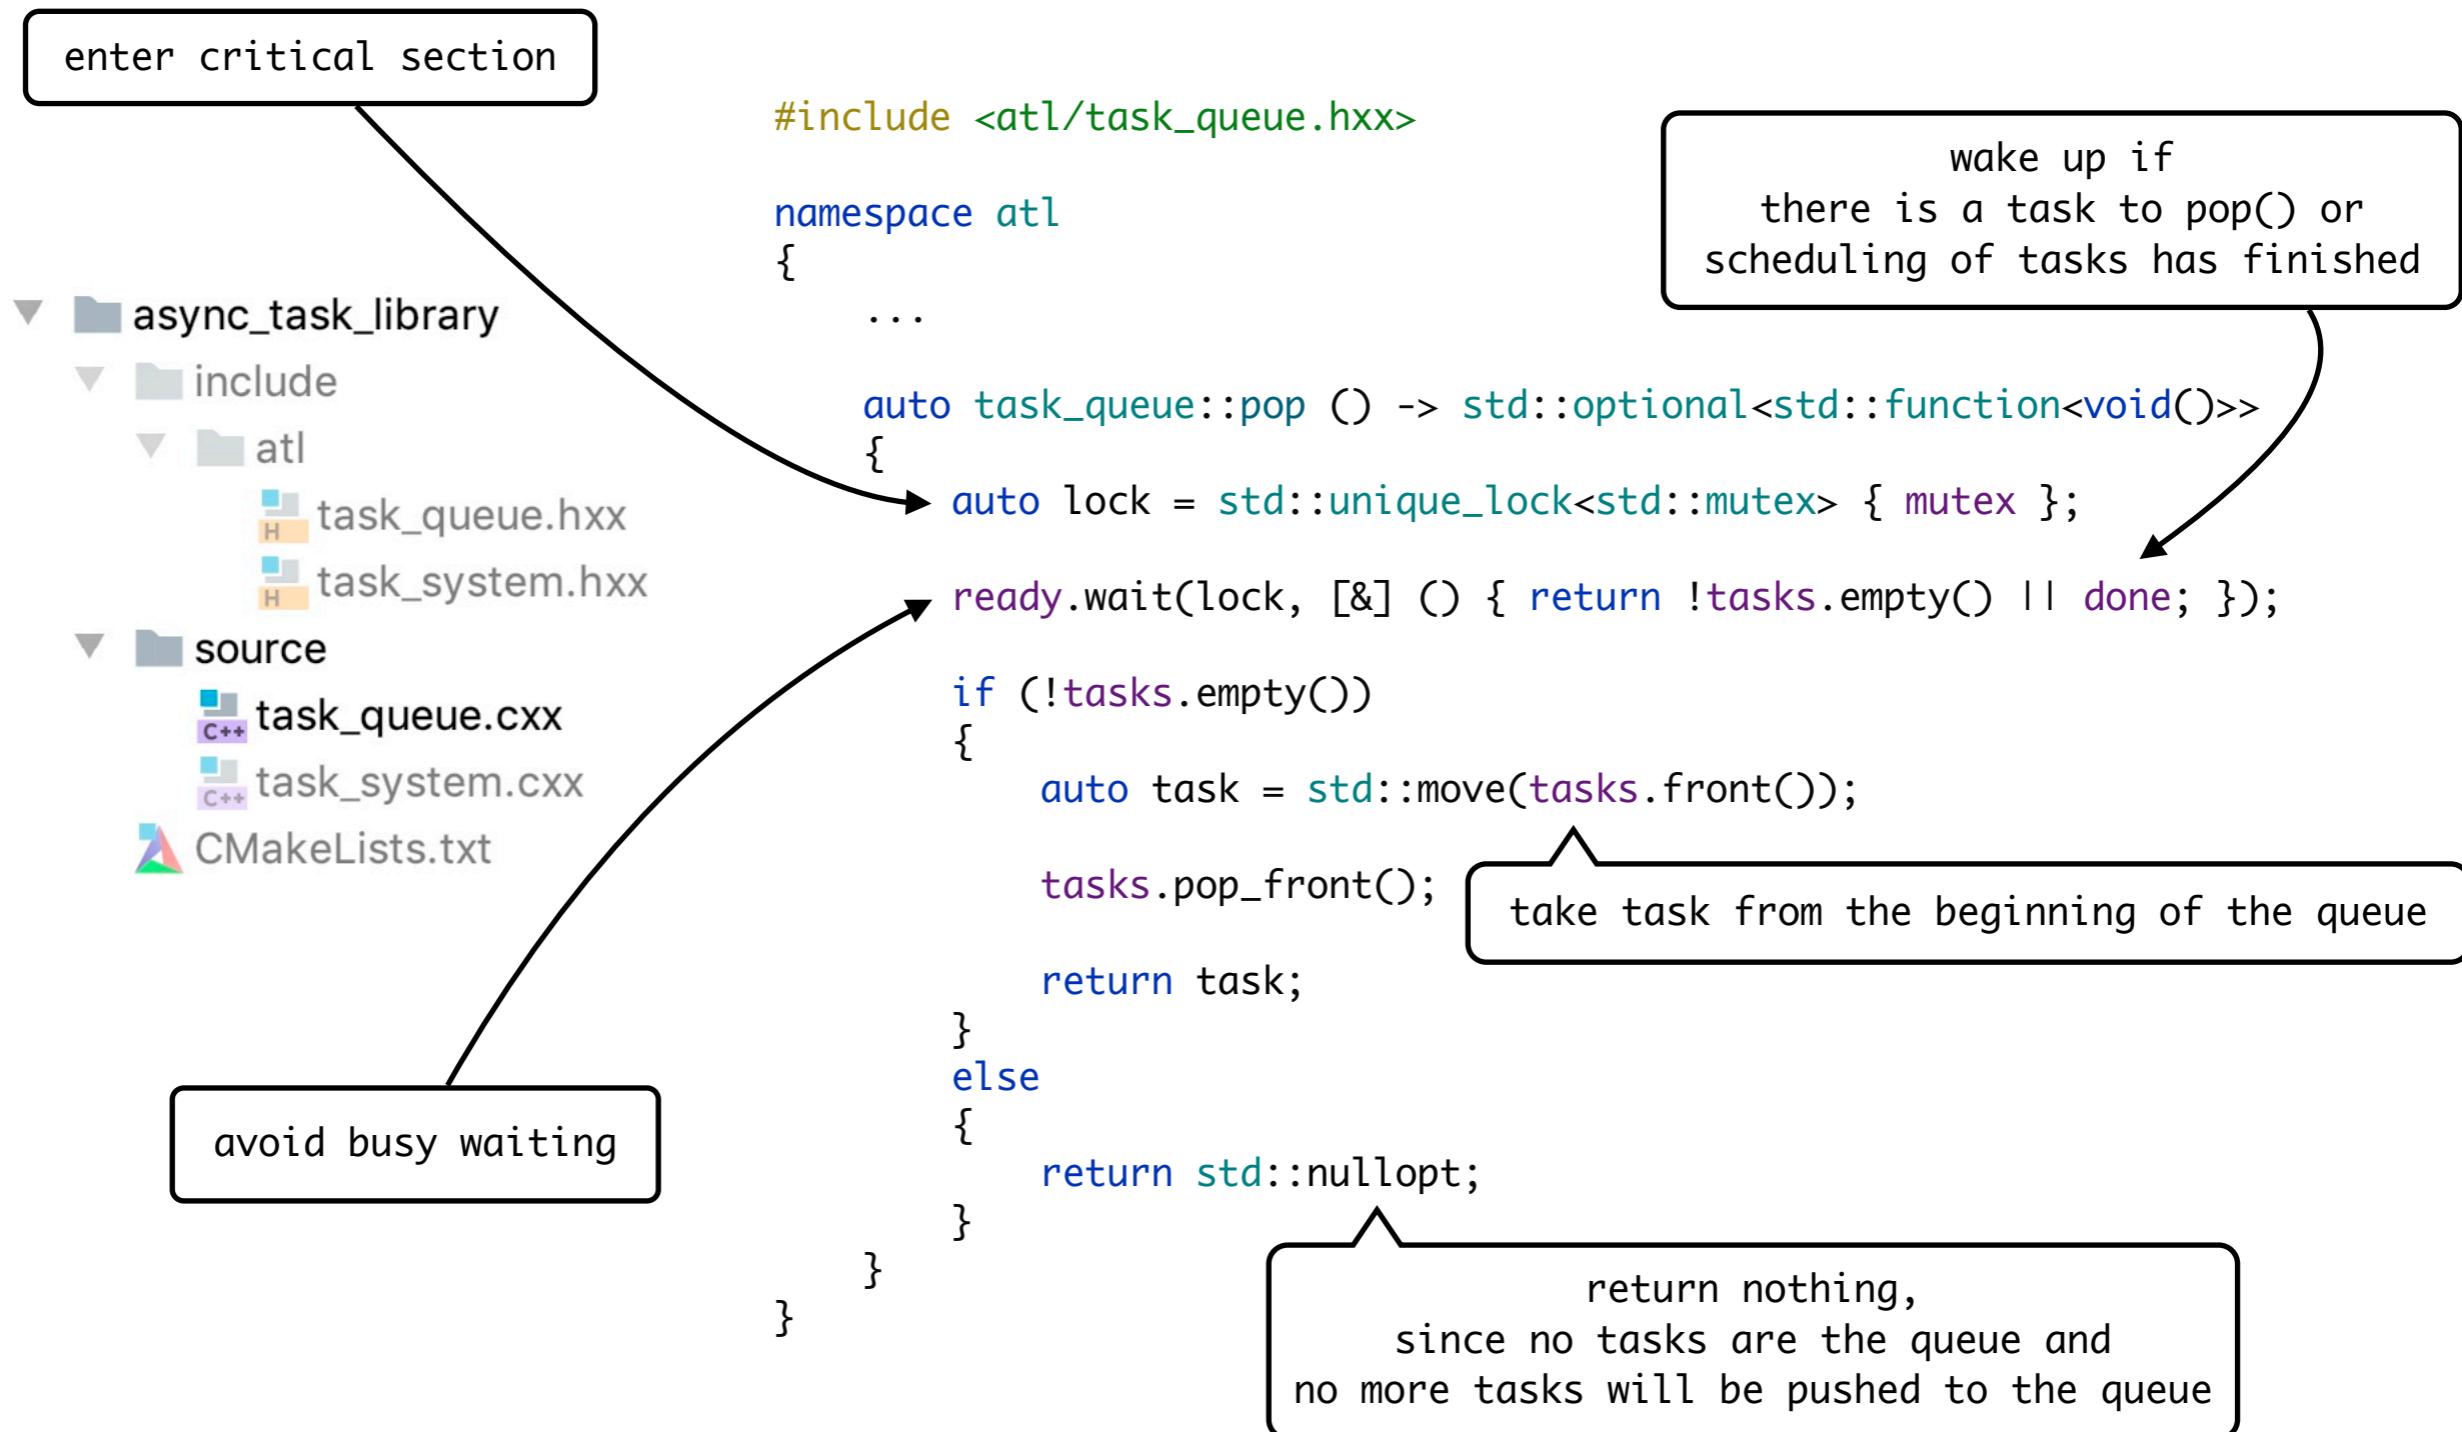
Async Task Library

Task Queue - Implementation



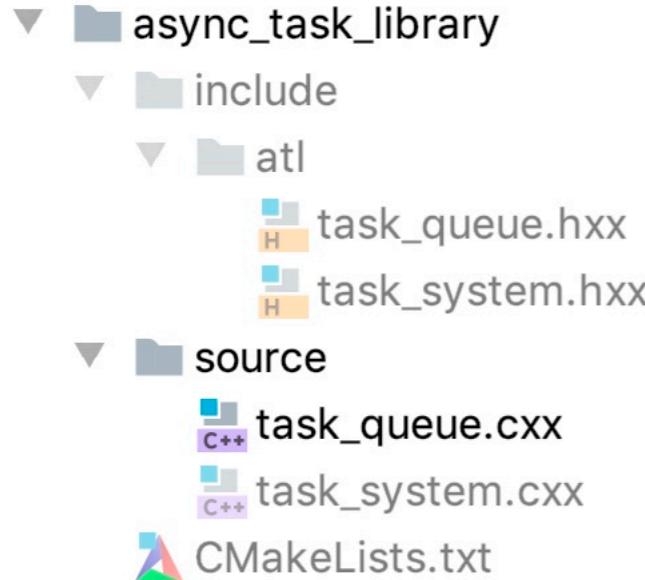
Async Task Library

Task Queue - Implementation



Async Task Library

Task Queue - Implementation

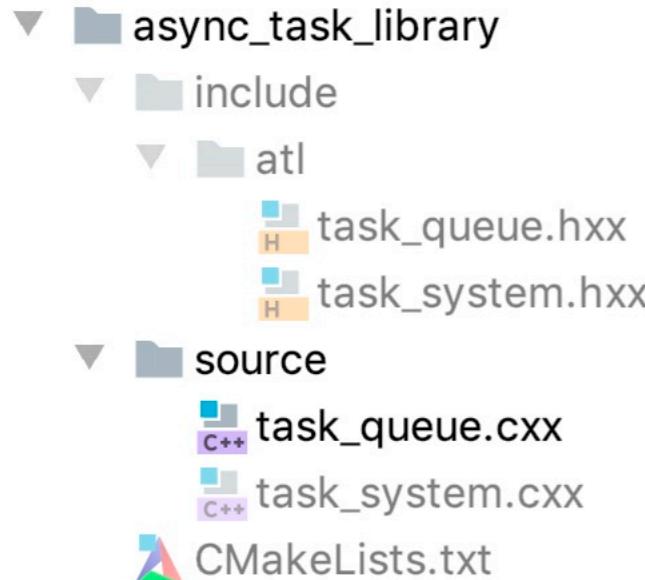


```
#include <atl/task_queue.hxx>

namespace atl
{
    auto task_queue::finish() -> void
    {
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            done = true;
        }
        ready.notify_all();
    }
    ...
}
```

Async Task Library

Task Queue - Implementation



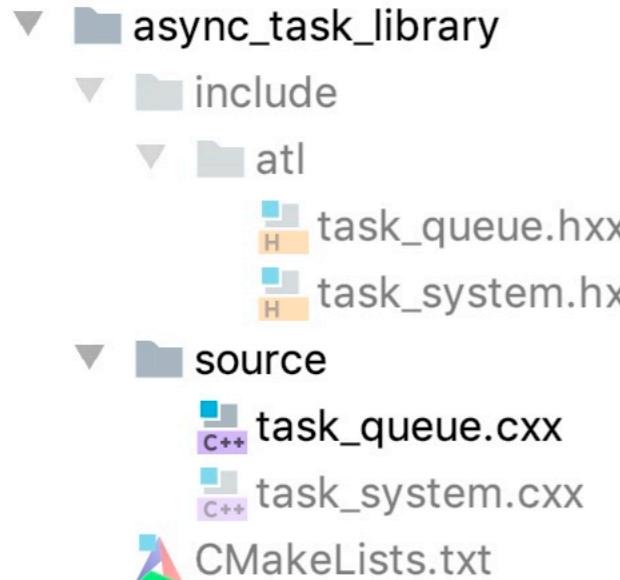
```
#include <atl/task_queue.hxx>

namespace atl
{
    auto task_queue::finish() -> void
    {
        auto lock = std::lock_guard<std::mutex> { mutex };
        done = true;
    }
    ready.notify_all();
}
...
```

enter critical section

Async Task Library

Task Queue - Implementation



```
#include <atl/task_queue.hxx>

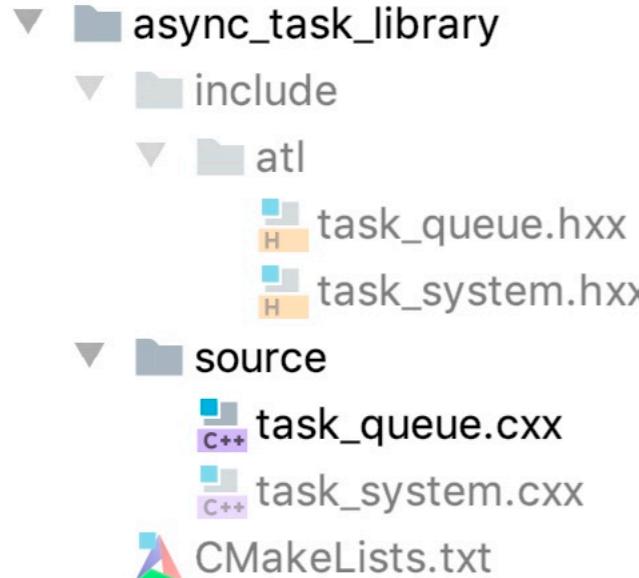
namespace atl
{
    auto task_queue::finish() -> void
    {
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            done = true; ←
        }
        ready.notify_all(); → set done flag
    }
    ...
}
```

Diagram annotations:

- A callout box points to the line `auto lock = std::lock_guard<std::mutex> { mutex };` with the text "enter critical section".
- A callout box points to the line `done = true;` with the text "set done flag".

Async Task Library

Task Queue - Implementation



```
#include <atl/task_queue.hxx>

namespace atl
{
    auto task_queue::finish() -> void
    {
        {
            auto lock = std::lock_guard<std::mutex> { mutex };
            done = true; // set done flag
        }
        ready.notify_all(); // notify all threads waiting inside pop(),
                            // that no more tasks will be push()ed to the queue
    }
}
```

Async Task Library

CMake build system

```
cmake_minimum_required (VERSION 3.15)
project(AsyncTaskLibrary VERSION 1.0.0 LANGUAGES CXX)

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ task_queue.cxx
        └─ task_system.cxx
    └─ CMakeLists.txt

find_package(Threads REQUIRED)

add_library           (async-task-lib STATIC)
target_compile_features (async-task-lib PUBLIC cxx_std_17)
target_include_directories(async-task-lib PUBLIC include)
target_link_libraries   (async-task-lib PUBLIC ${CMAKE_THREAD_LIBS_INIT})
target_sources          (async-task-lib PRIVATE
                           include/atl/task_queue.hxx
                           include/atl/task_system.hxx
                           source/task_queue.cxx
                           source/task_system.cxx)
```

Async Task Library

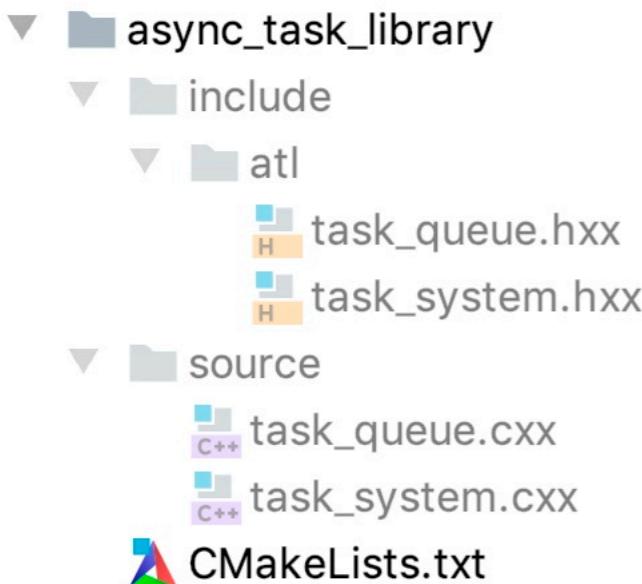
CMake build system

```
cmake_minimum_required(VERSION 3.15)
project(AsyncTaskLibrary VERSION 1.0.0 LANGUAGES CXX)

find_package(Threads REQUIRED)

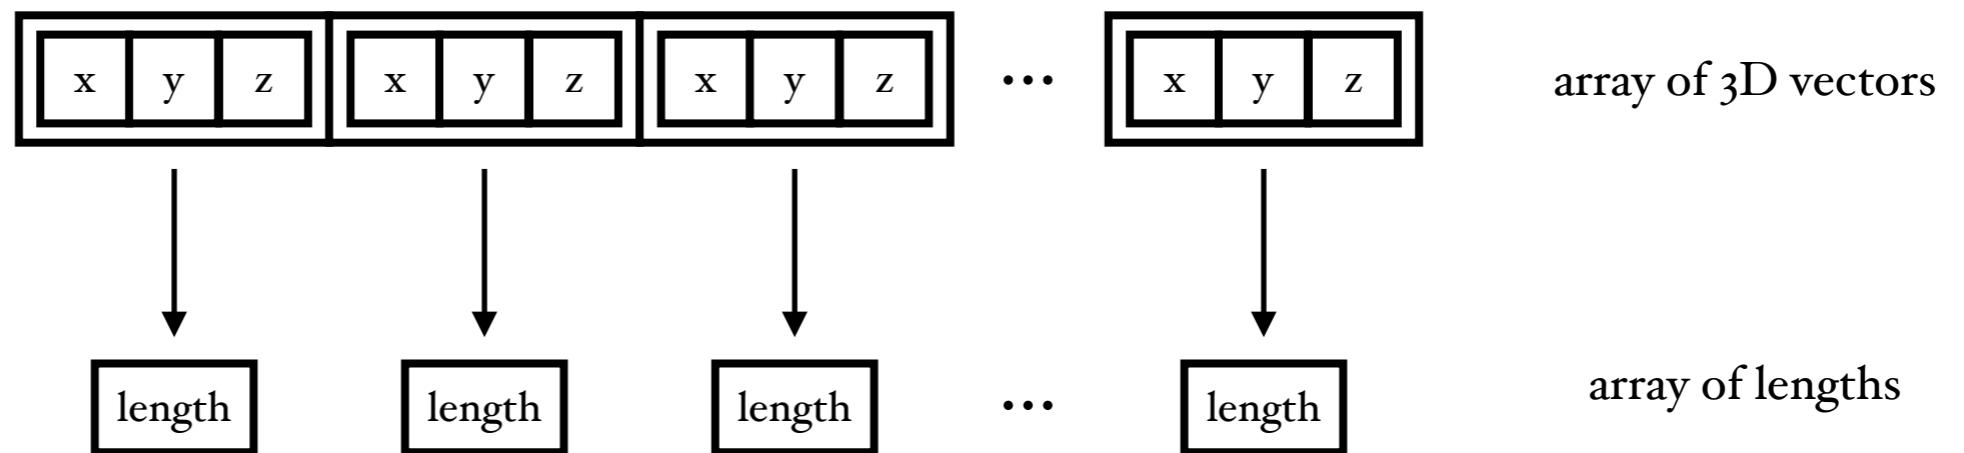
add_library(async-task-lib STATIC)
target_compile_features(async-task-lib PUBLIC cxx_std_17)
target_include_directories(async-task-lib PUBLIC include)
target_link_libraries(async-task-lib PUBLIC ${CMAKE_THREAD_LIBS_INIT})
target_sources(async-task-lib PRIVATE
    include/at1/task_queue.hxx
    include/at1/task_system.hxx
    source/task_queue.cxx
    source/task_system.cxx)
```

allow clients to #include header files
from the include directory



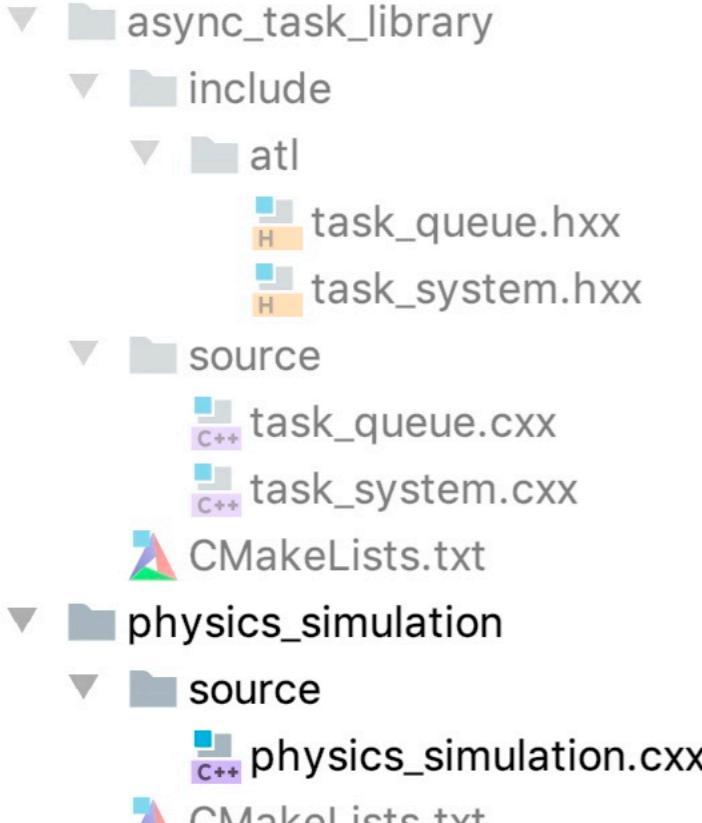
Physics Simulation

Diagram



Physics Simulation

Data Flow



```
#include ...

namespace phy_sim
{
    struct vec3 { float x, y, z; };

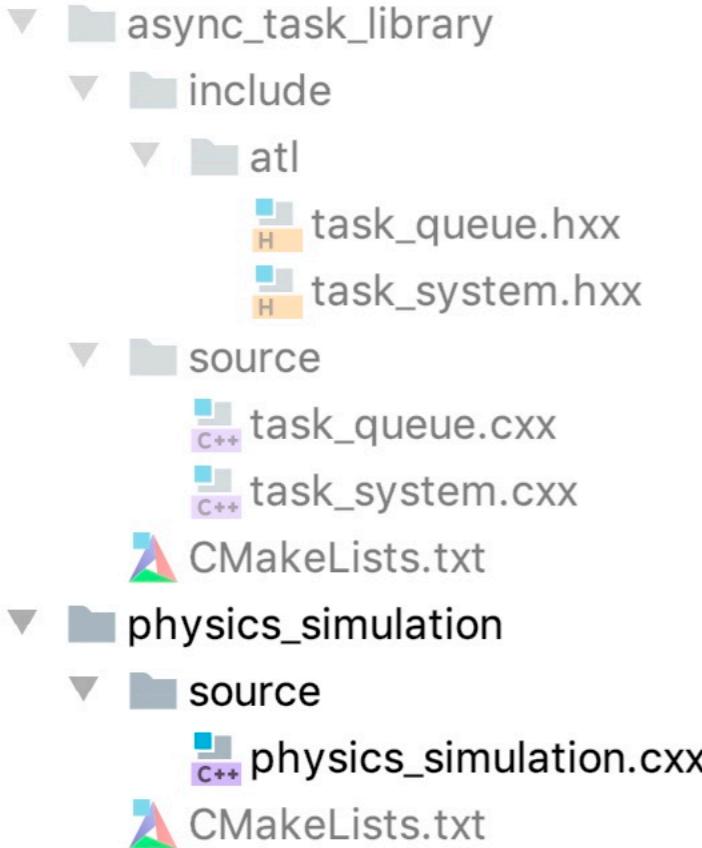
    auto read_input () -> array<vec3>
    { ... }
    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    { ... }
    auto print_results (const array<float>& lengths) -> void
    { ... }

    auto main () -> int
    {
        const auto vectors = phy_sim::read_input();
        const auto lengths = phy_sim::compute_lengths(vectors);
        phy_sim::print_results(lengths);

        return 0;
    }
}
```

Physics Simulation

Data Flow



```
#include ...

namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto read_input () -> array<vec3>
    { ... }
    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    { ... }
    auto print_results (const array<float>& lengths) -> void
    { ... }

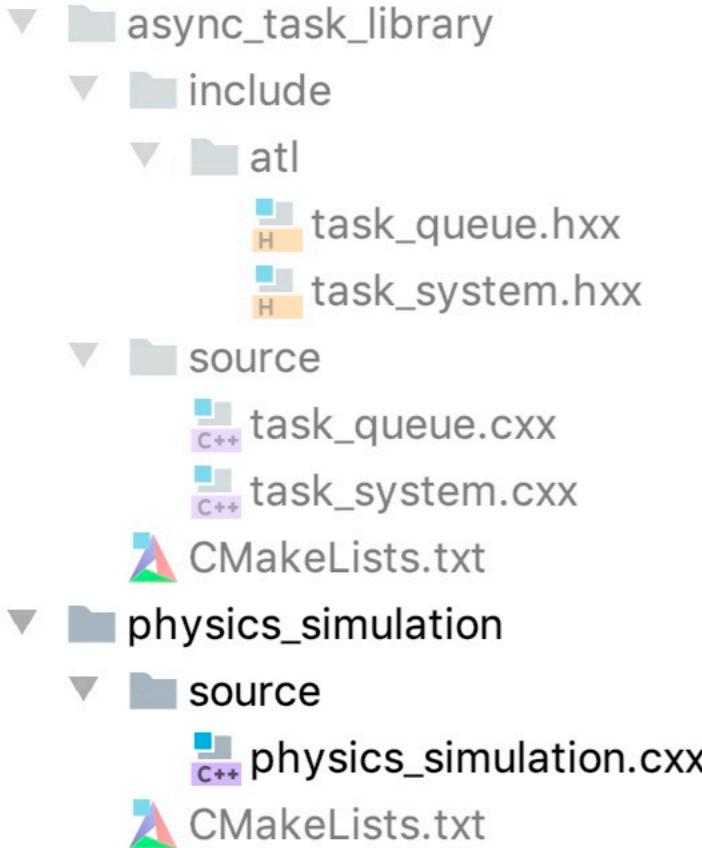
    auto main () -> int
    {
        const auto vectors = phy_sim::read_input();
        const auto lengths = phy_sim::compute_lengths(vectors);
        phy_sim::print_results(lengths);

        return 0;
    }
}
```

A callout box labeled "3D vector" points to the `vec3` struct definition.

Physics Simulation

Data Flow



```
#include ...

namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto read_input () -> array<vec3>
    {
        ...
    }

    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    {
        ...
    }

    auto print_results (const array<float>& lengths) -> void
    {
        ...
    }
}

auto main () -> int
{
    const auto vectors = phy_sim::read_input();
    const auto lengths = phy_sim::compute_lengths(vectors);
    phy_sim::print_results(lengths);

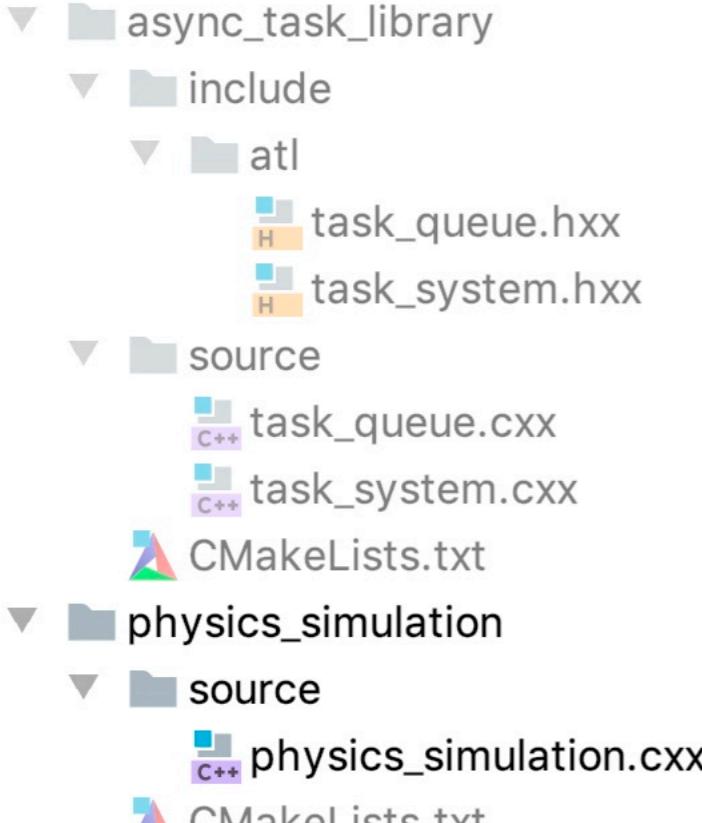
    return 0;
}
```

3D vector

obtain input array of 3D vectors

Physics Simulation

Data Flow



```
#include ...

namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto read_input () -> array<vec3>
    {
        ...
    }

    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    {
        ...
    }

    auto print_results (const array<float>& lengths) -> void
    {
        ...
    }
}

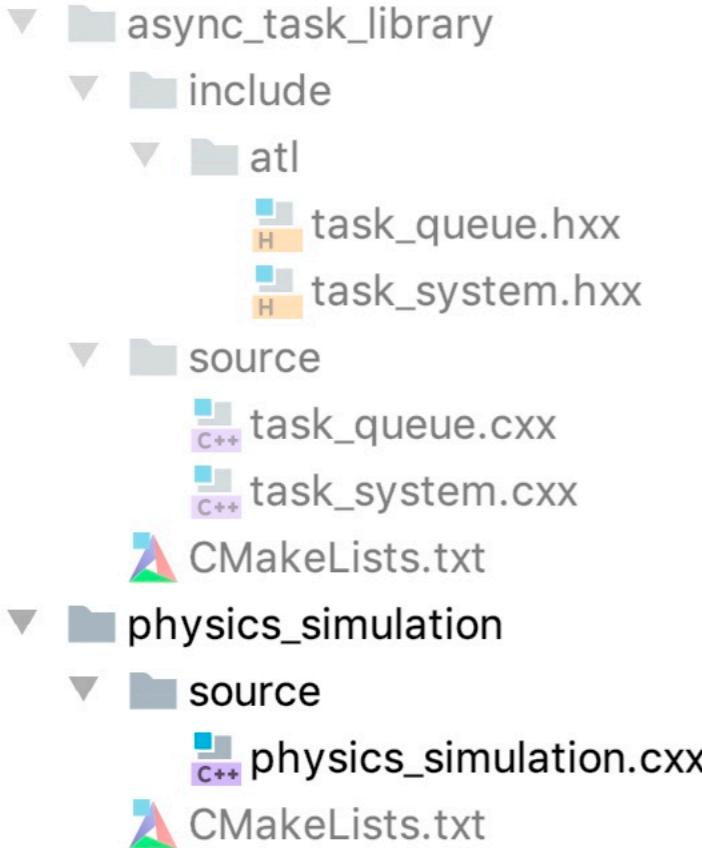
auto main () -> int
{
    const auto vectors = phy_sim::read_input();
    const auto lengths = phy_sim::compute_lengths(vectors);
    phy_sim::print_results(lengths);
    return 0;
}
```

Annotations for the code:

- A callout box points to the 'vec3' struct definition with the text "3D vector".
- A callout box points to the first line of the 'main' function with the text "obtain input array of 3D vectors".
- A callout box points to the second line of the 'main' function with the text "compute lengths of 3D vectors".

Physics Simulation

Data Flow



```
#include ...

namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto read_input () -> array<vec3>
    { ... }

    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    { ... }

    auto print_results (const array<float>& lengths) -> void
    { ... }
}

auto main () -> int
{
    const auto vectors = phy_sim::read_input();
    const auto lengths = phy_sim::compute_lengths(vectors);
    phy_sim::print_results(lengths);
    return 0;
}
```

3D vector

obtain input array of 3D vectors

compute lengths of 3D vectors

report calculated lengths

Physics Simulation

Length of a 3D Vector

```
#include <atl/task_system.hxx>
#include <cmath>

namespace phy_sim
{
    ...

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x + v.y * v.y + v.z * v.z);
    }

    ...
}

▼ async_task_library
  ▼ include
    ▼ atl
      task_queue.hxx
      task_system.hxx
  ▼ source
    task_queue.cxx
    task_system.cxx
  CMakeLists.txt

▼ physics_simulation
  ▼ source
    physics_simulation.cxx
  CMakeLists.txt
```

Physics Simulation

Length of a 3D Vector

```
▼ └── async_task_library
    └── include
        └── atl
            ├── task_queue.hxx
            └── task_system.hxx
    └── source
        ├── task_queue.cxx
        └── task_system.cxx
└── CMakeLists.txt

▼ └── physics_simulation
    └── source
        ├── physics_simulation.cxx
    └── CMakeLists.txt
```

```
#include <atl/task_system.hxx>
#include <cmath>

namespace phy_sim
{
    ...

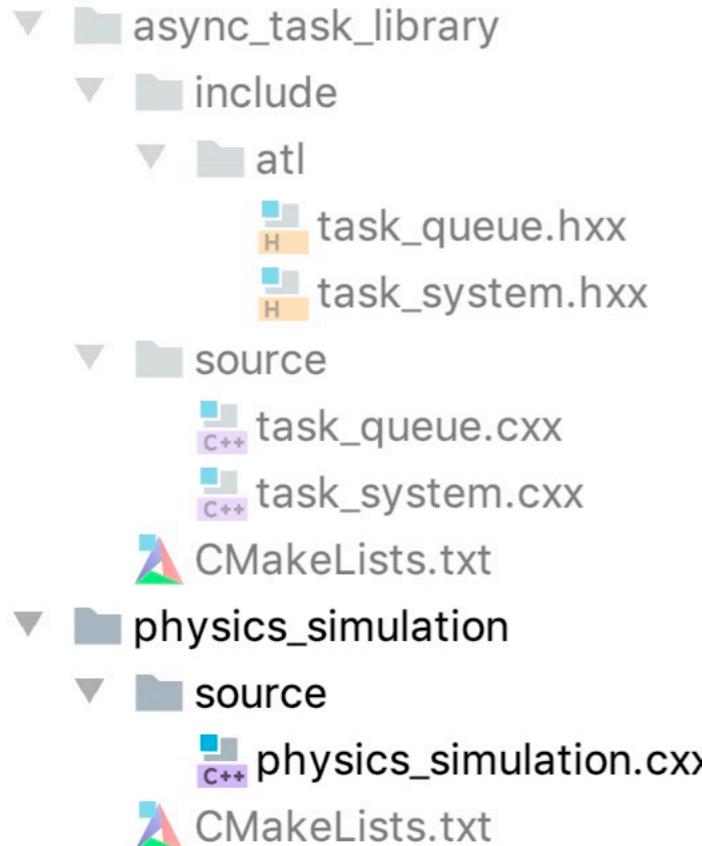
    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x + v.y * v.y + v.z * v.z);
    }
    ...
}
```

calculate length of a 3D vector
using Euclidean metric:

$$|v| = \sqrt{v_x^2 + v_y^2 + v_z^2}$$

Physics Simulation

Compute Lengths - Task System Example



```
#include <atl/task_system.hxx>
#include <cmath>

namespace phy_sim
{
    ...

    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    {
        auto lengths = array<float> { };

        auto task_system = atl::task_system { };

        for (auto index = 0; index < vectors.size(); ++index)
        {
            auto task = [index, &vectors, &lengths] ()
            {
                lengths[index] = length(vectors[index]);
            };
            task_system.async(task);
        }
        return lengths;
    }

    ...
}
```

Physics Simulation

Compute Lengths - Task System Example

The diagram illustrates a code example for a physics simulation using a task system. A callout box labeled "create array of lengths" points to the `task_system.hxx` header file within the `async_task_library` directory.

```
#include <atl/task_system.hxx>
#include <cmath>

namespace phy_sim
{
    ...

    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    {
        auto lengths = array<float> { };

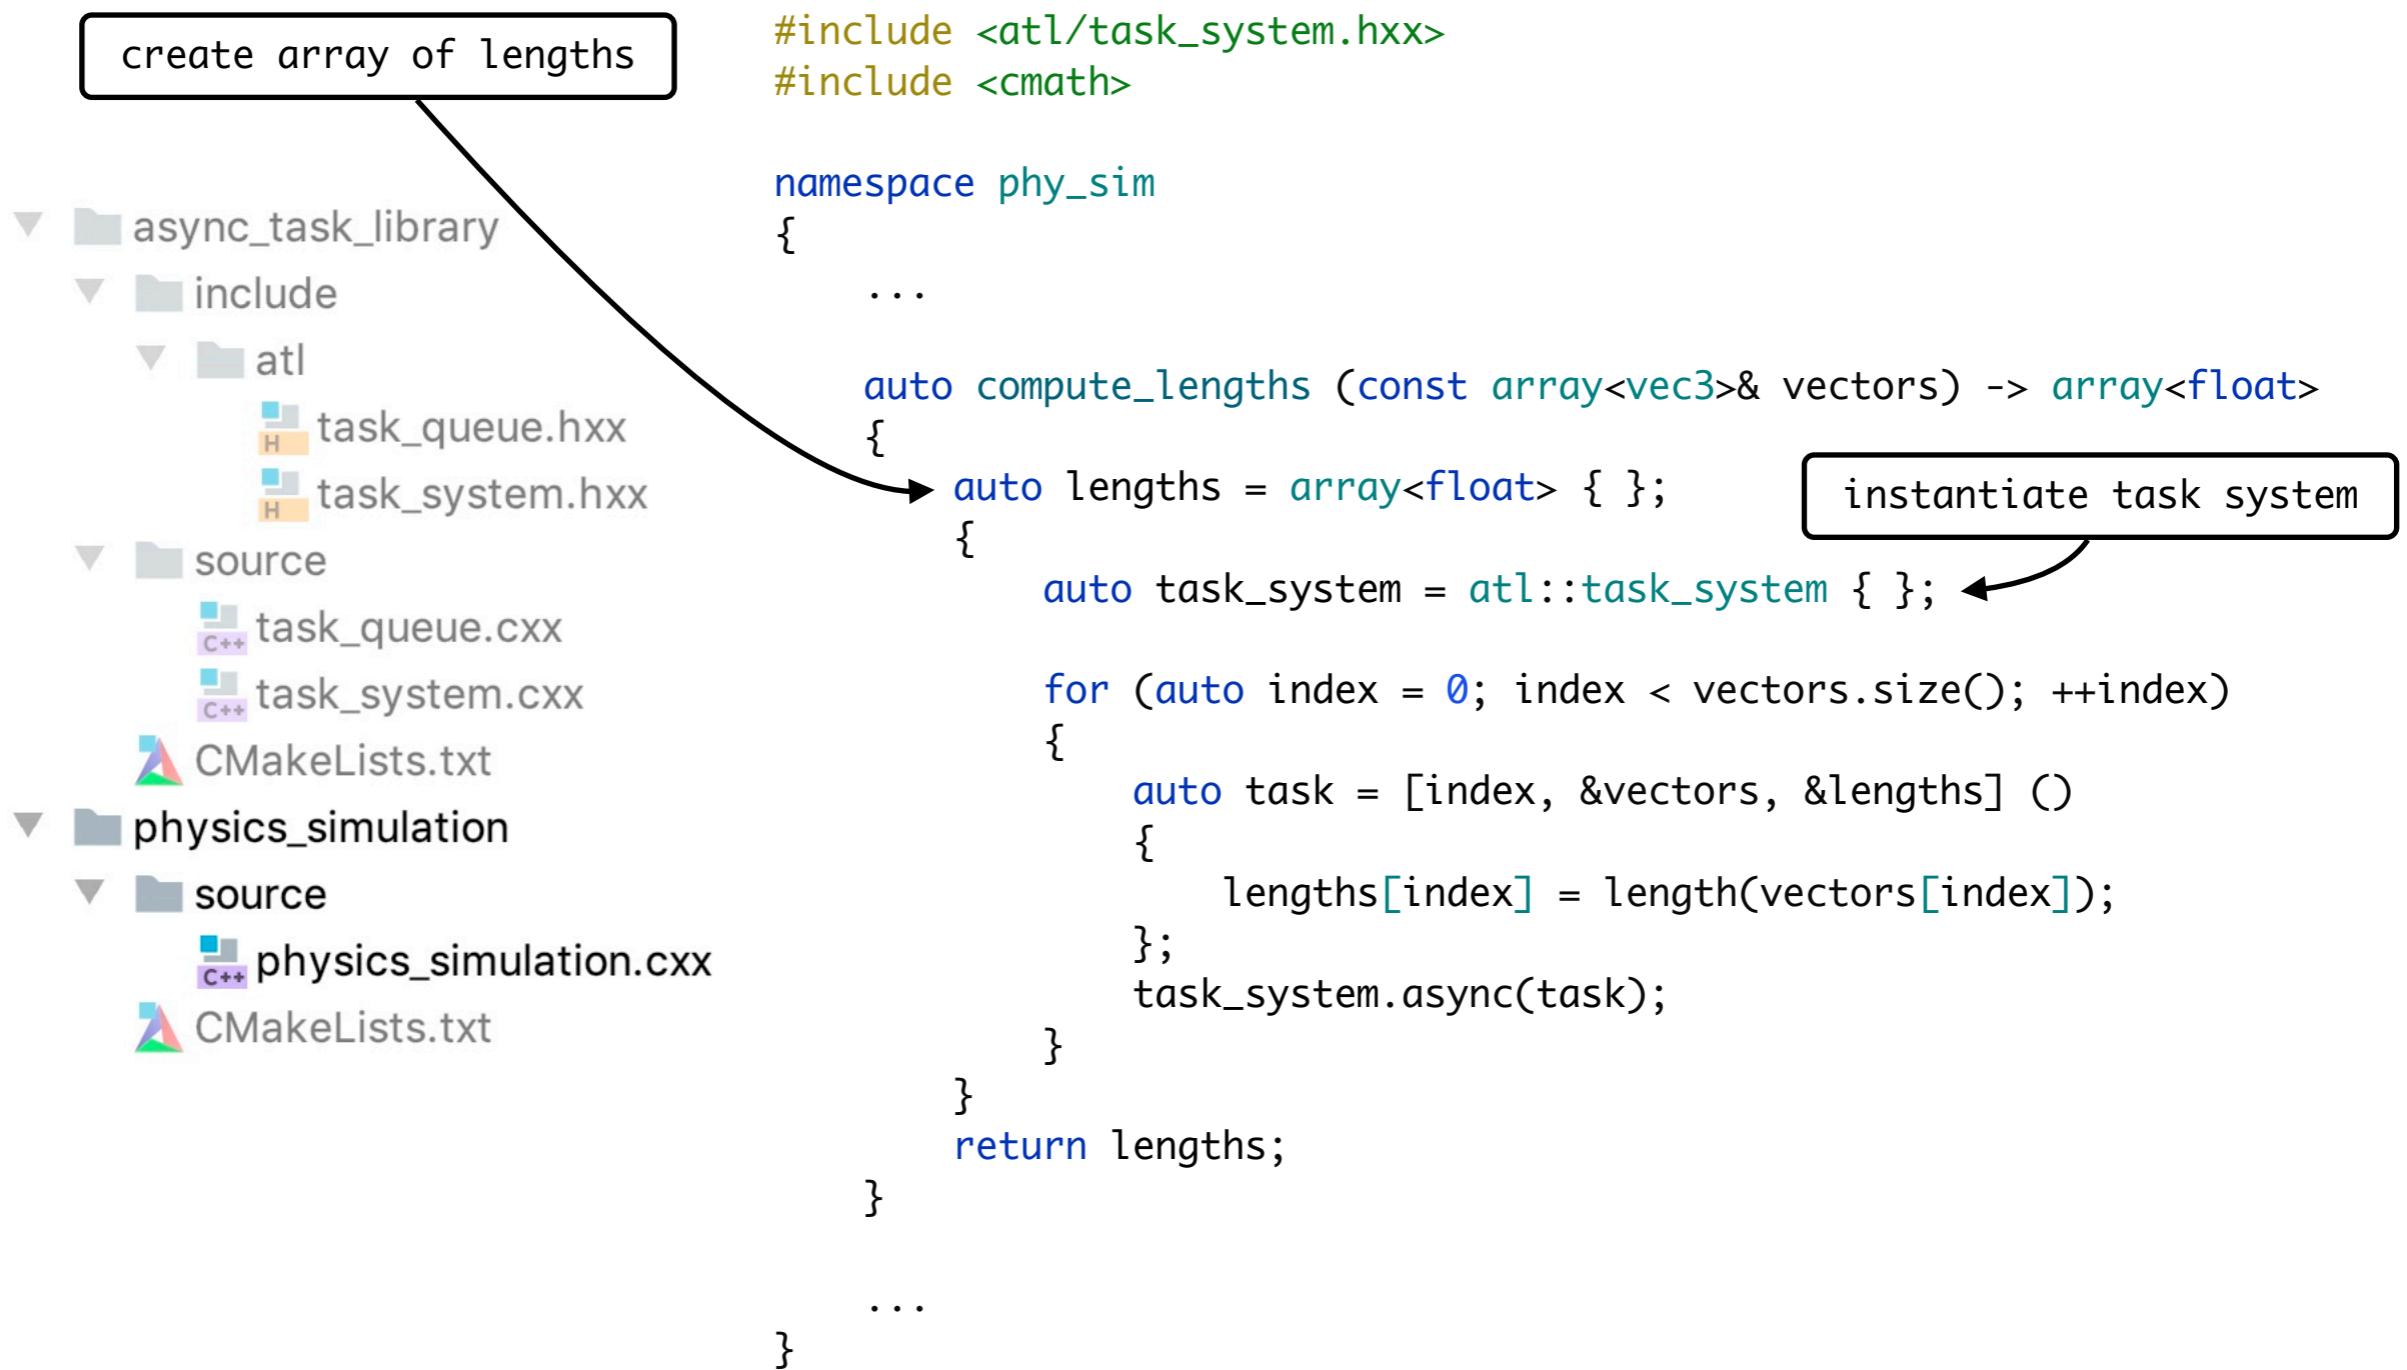
        auto task_system = atl::task_system { };

        for (auto index = 0; index < vectors.size(); ++index)
        {
            auto task = [index, &vectors, &lengths] ()
            {
                lengths[index] = length(vectors[index]);
            };
            task_system.async(task);
        }
        return lengths;
    }

    ...
}
```

Physics Simulation

Compute Lengths - Task System Example



```
#include <atl/task_system.hxx>
#include <cmath>

namespace phy_sim
{
    ...

    auto compute_lengths (const array<vec3>& vectors) -> array<float>
    {
        auto lengths = array<float> { };

        auto task_system = atl::task_system { };

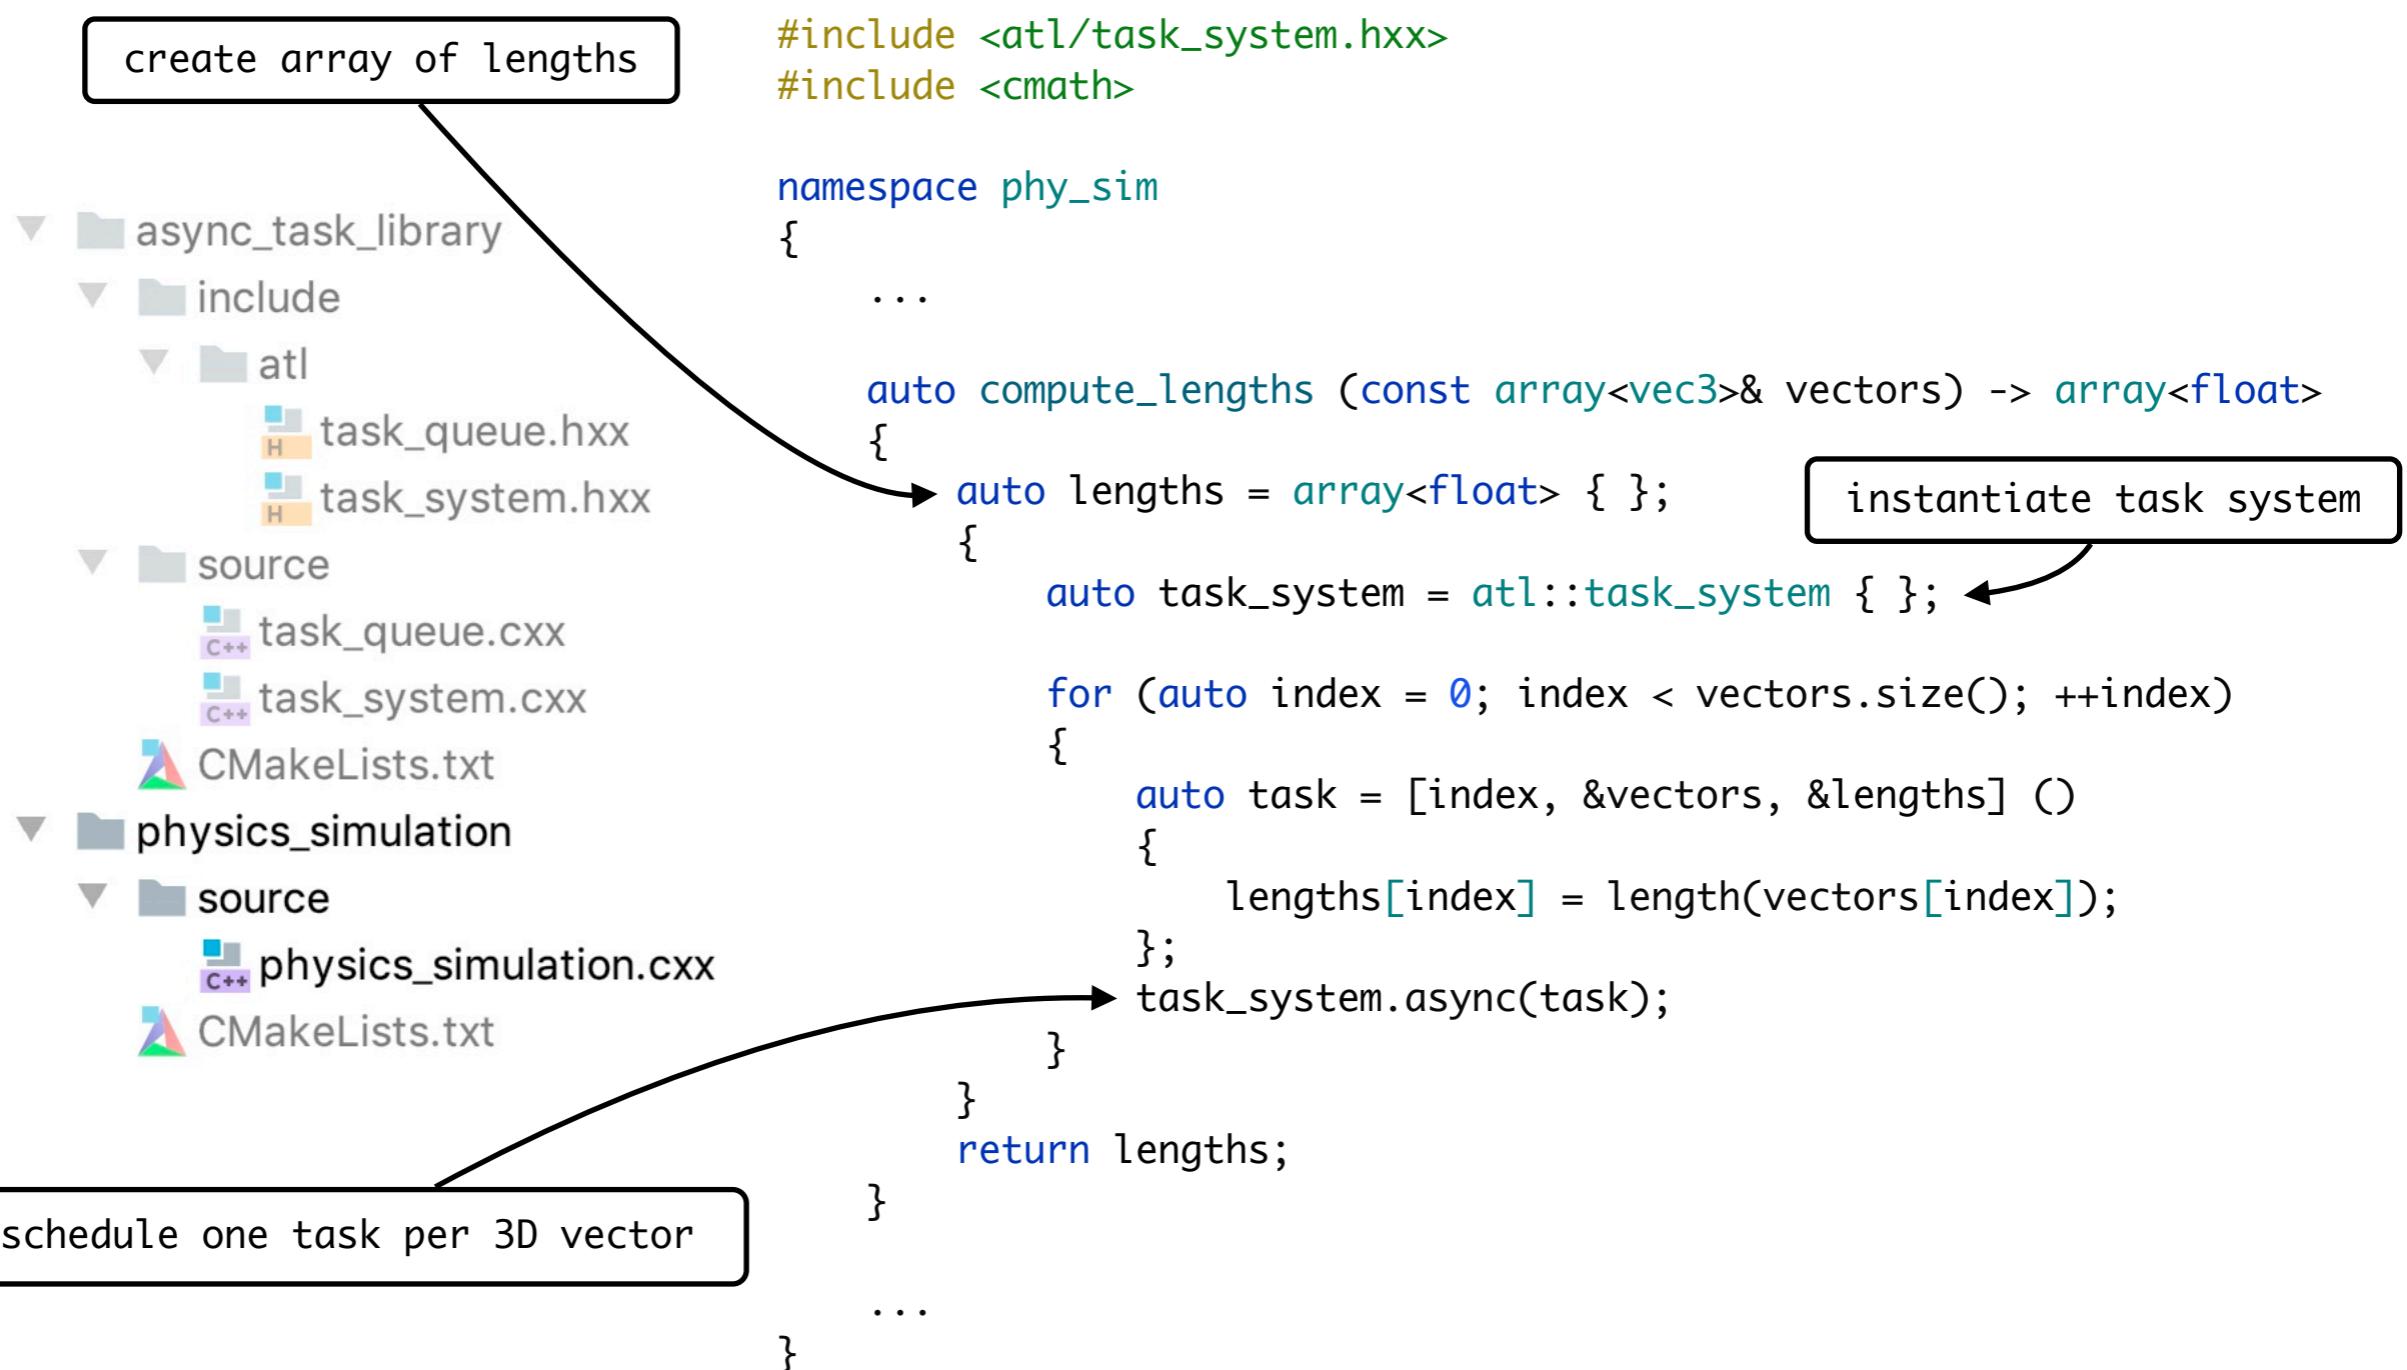
        for (auto index = 0; index < vectors.size(); ++index)
        {
            auto task = [index, &vectors, &lengths] ()
            {
                lengths[index] = length(vectors[index]);
            };
            task_system.async(task);
        }
        return lengths;
    }

    ...
}
```

instantiate task system

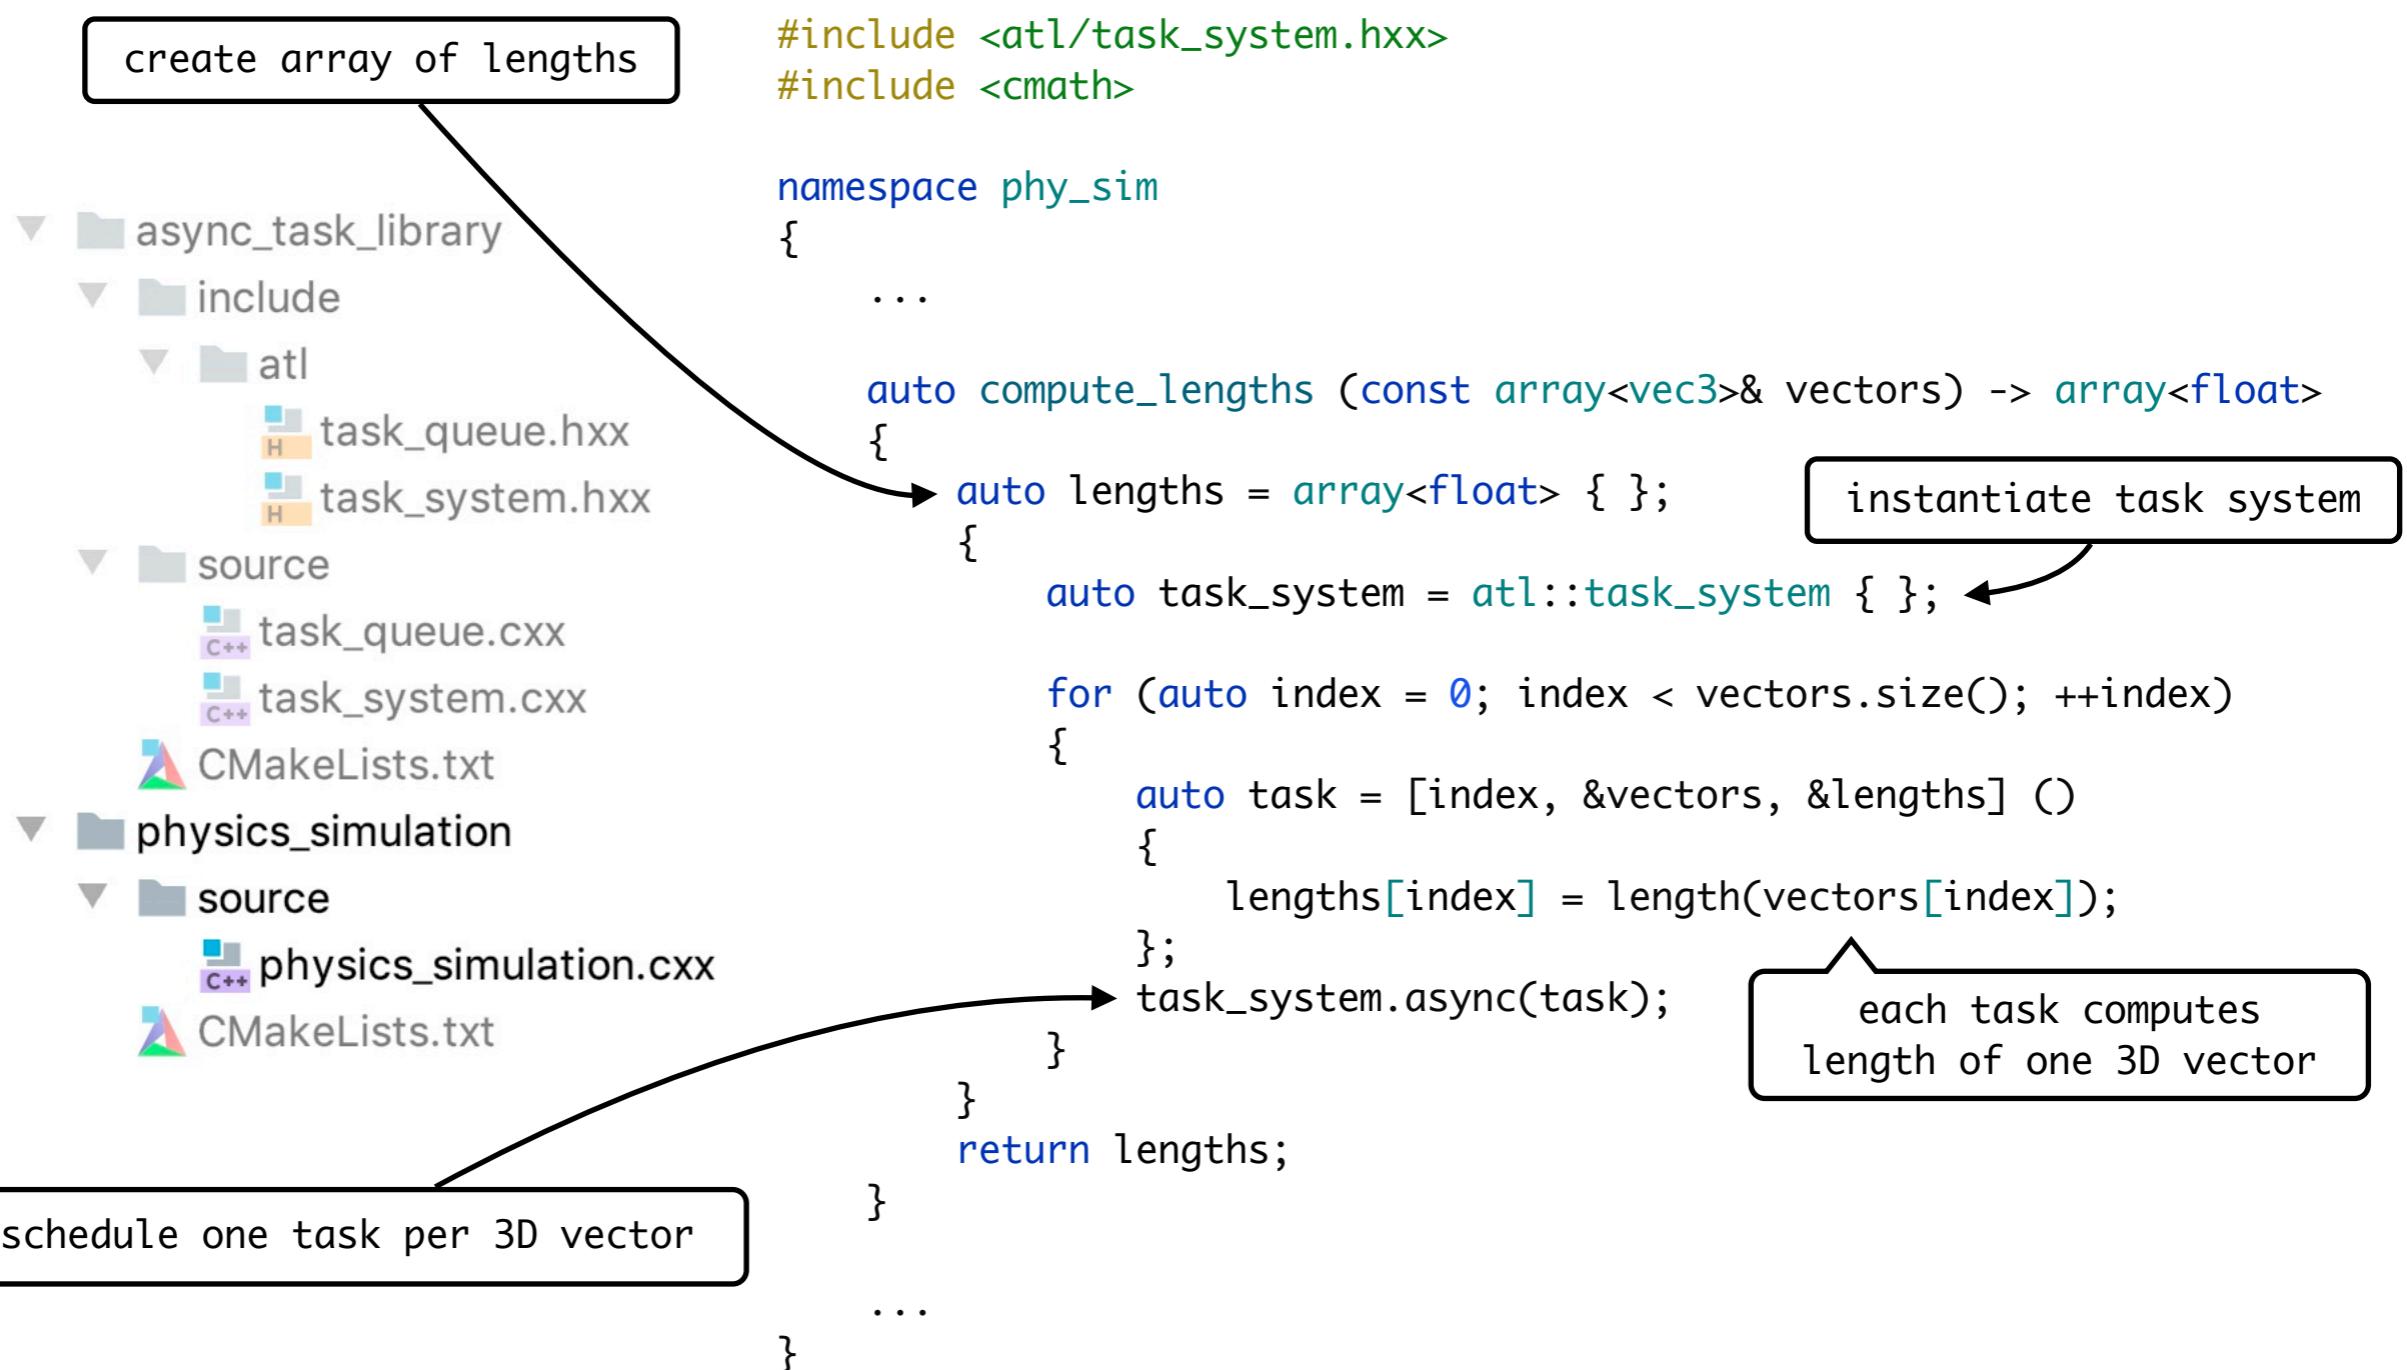
Physics Simulation

Compute Lengths - Task System Example



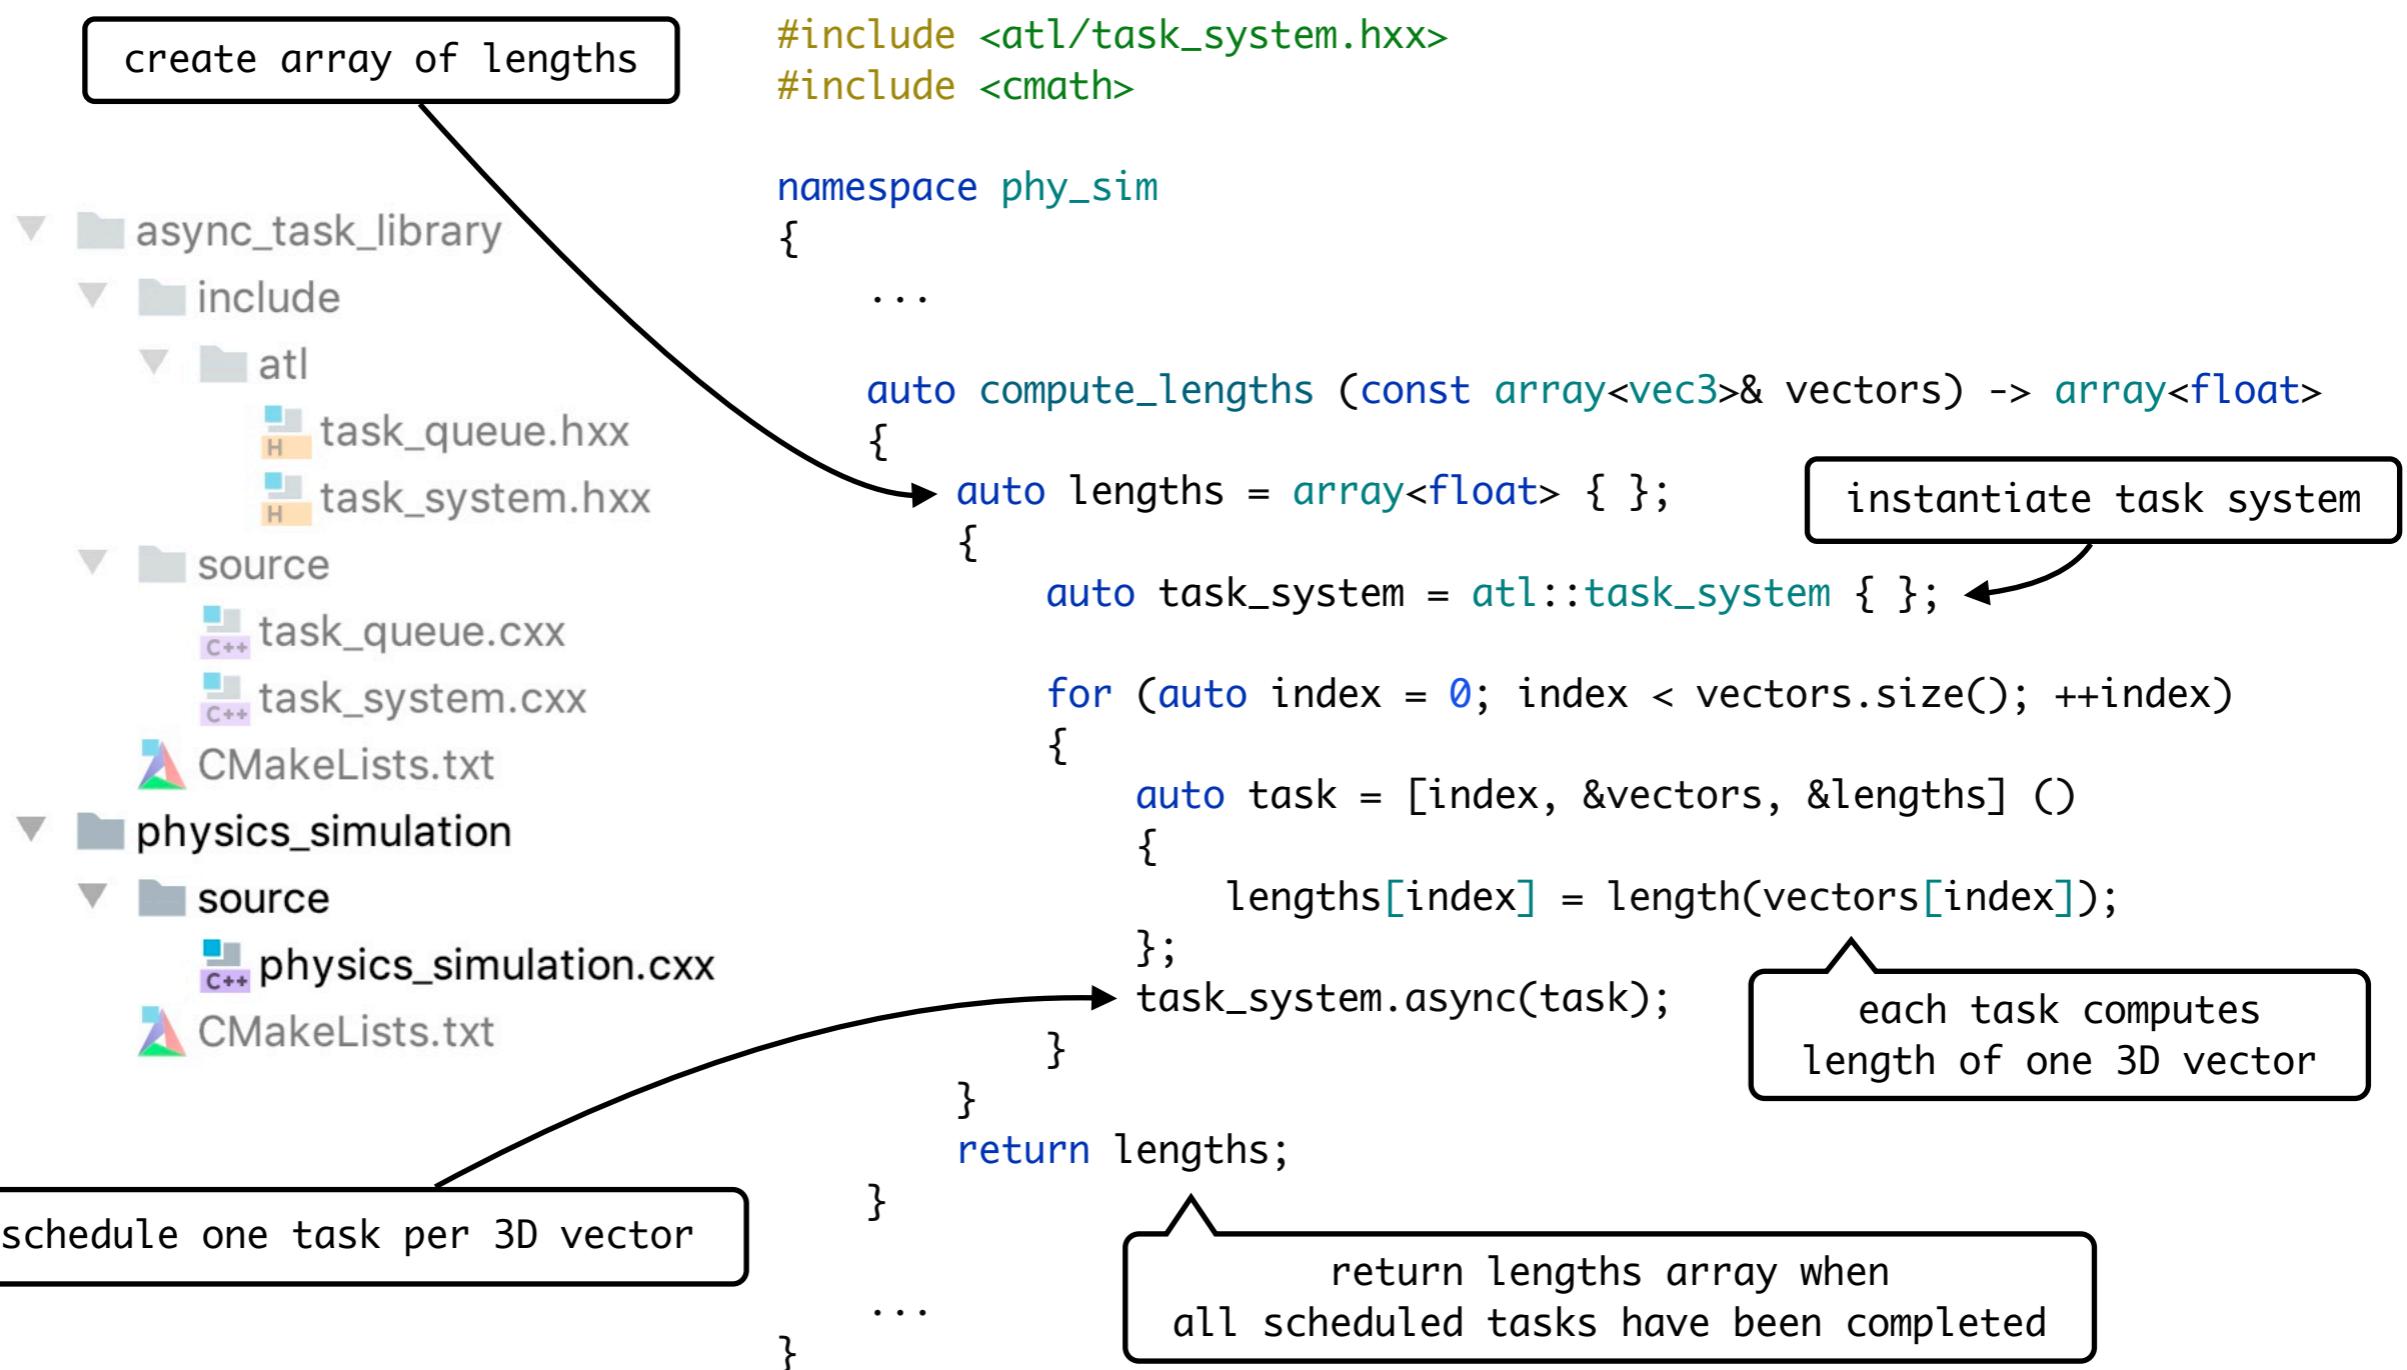
Physics Simulation

Compute Lengths - Task System Example



Physics Simulation

Compute Lengths - Task System Example



Physics Simulation

CMake build system

```
cmake_minimum_required (VERSION 3.15)

▼ └─ async_task_library
    └─ include
        └─ atl
            └─ task_queue.hxx
            └─ task_system.hxx
    └─ source
        └─ CMakeLists.txt
            └─ task_queue.cxx
            └─ task_system.cxx
            └─ CMakeLists.txt

└─ physics_simulation
    └─ source
        └─ physics_simulation.hxx
        └─ CMakeLists.txt
            └─ physics_simulation.cxx
            └─ CMakeLists.txt

add_executable (physics-simulation)
target_compile_features (physics-simulation PRIVATE cxx_std_17)
target_sources (physics-simulation PRIVATE
                source/physics_simulation.hxx)
target_link_libraries (physics-simulation PRIVATE
                        async-task-lib)
```

Physics Simulation

CMake build system

```
cmake_minimum_required (VERSION 3.15)

▼ └── async_task_library
    └── include
        └── atl
            ├── task_queue.hxx
            └── task_system.hxx
    └── source
        ├── task_queue.cxx
        └── task_system.cxx
    └── CMakeLists.txt

└── physics_simulation
    └── source
        └── physics_simulation.cxx
    └── CMakeLists.txt
```

add_executable (physics-simulation)

target_compile_features (physics-simulation PRIVATE cxx_std_17)

target_sources (physics-simulation PRIVATE source/physics_simulation.cxx)

target_link_libraries (physics-simulation PRIVATE async-task-lib)

define dependency between CMake targets
physics-simulation and async-task-lib

Physics Simulation

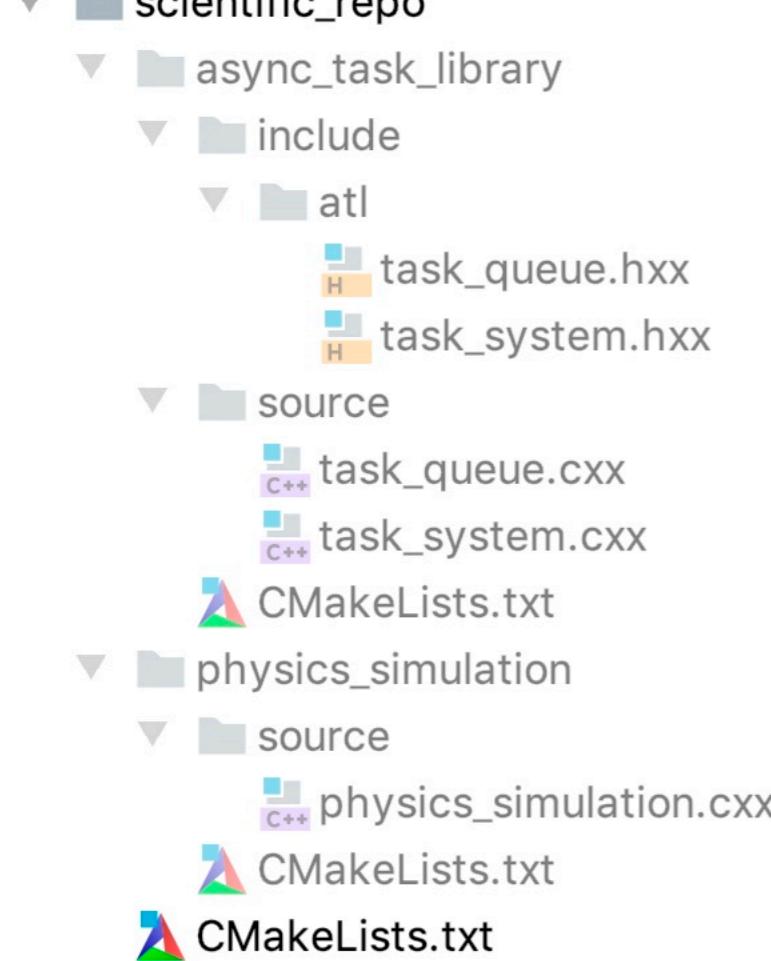
Scientific Repo

```
cmake_minimum_required (VERSION 3.15)
project(PhysicsSimulation VERSION 1.0.0 LANGUAGES CXX)

▼科学性仓库
  ▼async_task_library
    ▼include
      ▼atl
        task_queue.hxx
        task_system.hxx
    ▼source
      task_queue.cxx
      task_system.cxx
    CMakeLists.txt
  ▼physics_simulation
    ▼source
      physics_simulation.cxx
    CMakeLists.txt
CMakeLists.txt
```

Physics Simulation

Scientific Repo



```
cmake_minimum_required (VERSION 3.15)
project(PhysicsSimulation VERSION 1.0.0 LANGUAGES CXX)

add_subdirectory(async_task_library)
add_subdirectory(physics_simulation)
```

compose together CMake targets
async-task-lib and physics-simulation
into CMake project PhysicsSimulation

Question

What is wrong with the above code?

Question

What is wrong with the above code?

- Performance of the task system is very poor.
 - The task queue is under a lot of contention,
since threads scheduling and executing tasks are all trying to access the same task queue.

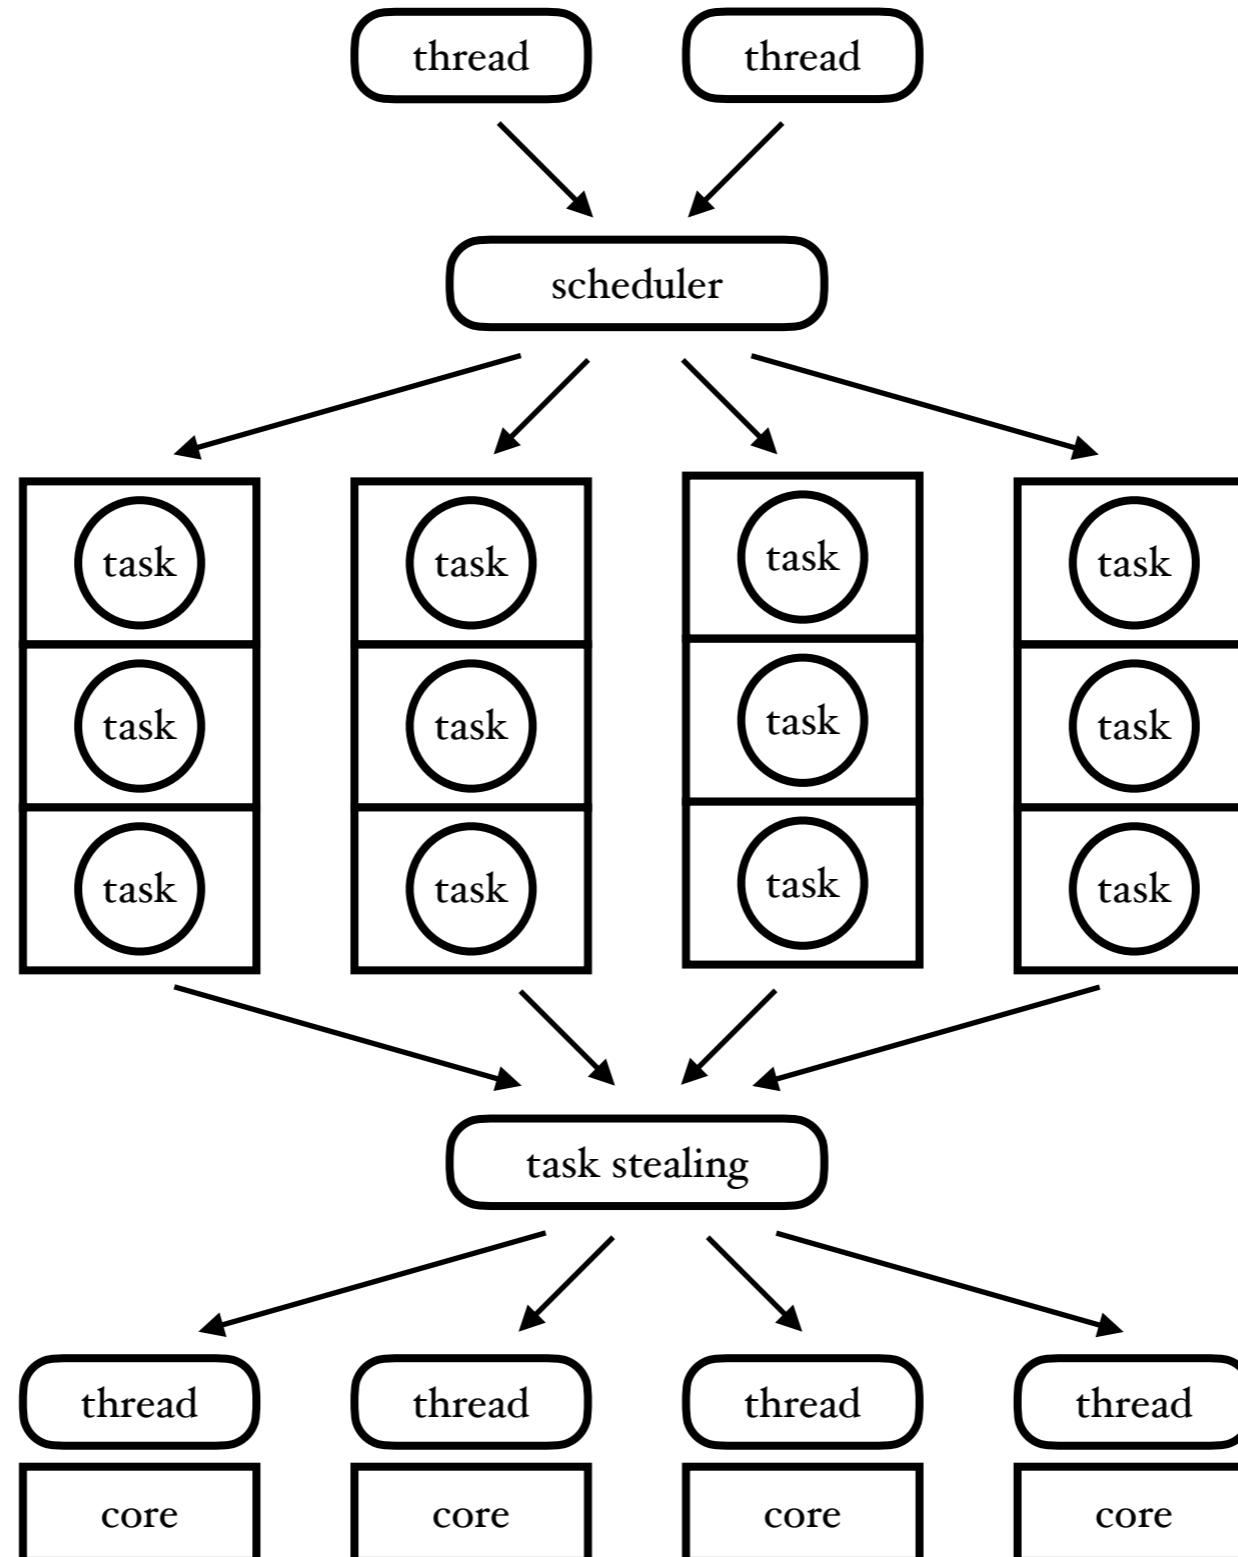
Question

What is wrong with the above code?

- Performance of the task system is very poor.
 - The task queue is under a lot of contention, since threads scheduling and executing tasks are all trying to access the same task queue.
 - In a much better task system:
 - ✓ each thread in the thread pool would have its own task queue,
 - ✓ a task scheduler would push tasks into a chosen queue using the round-robin algorithm,
 - ✓ load balancing would be done via task stealing.

Recommendation

Task System - Better Design



The scheduler `push()`s a task into the first non-busy queue.

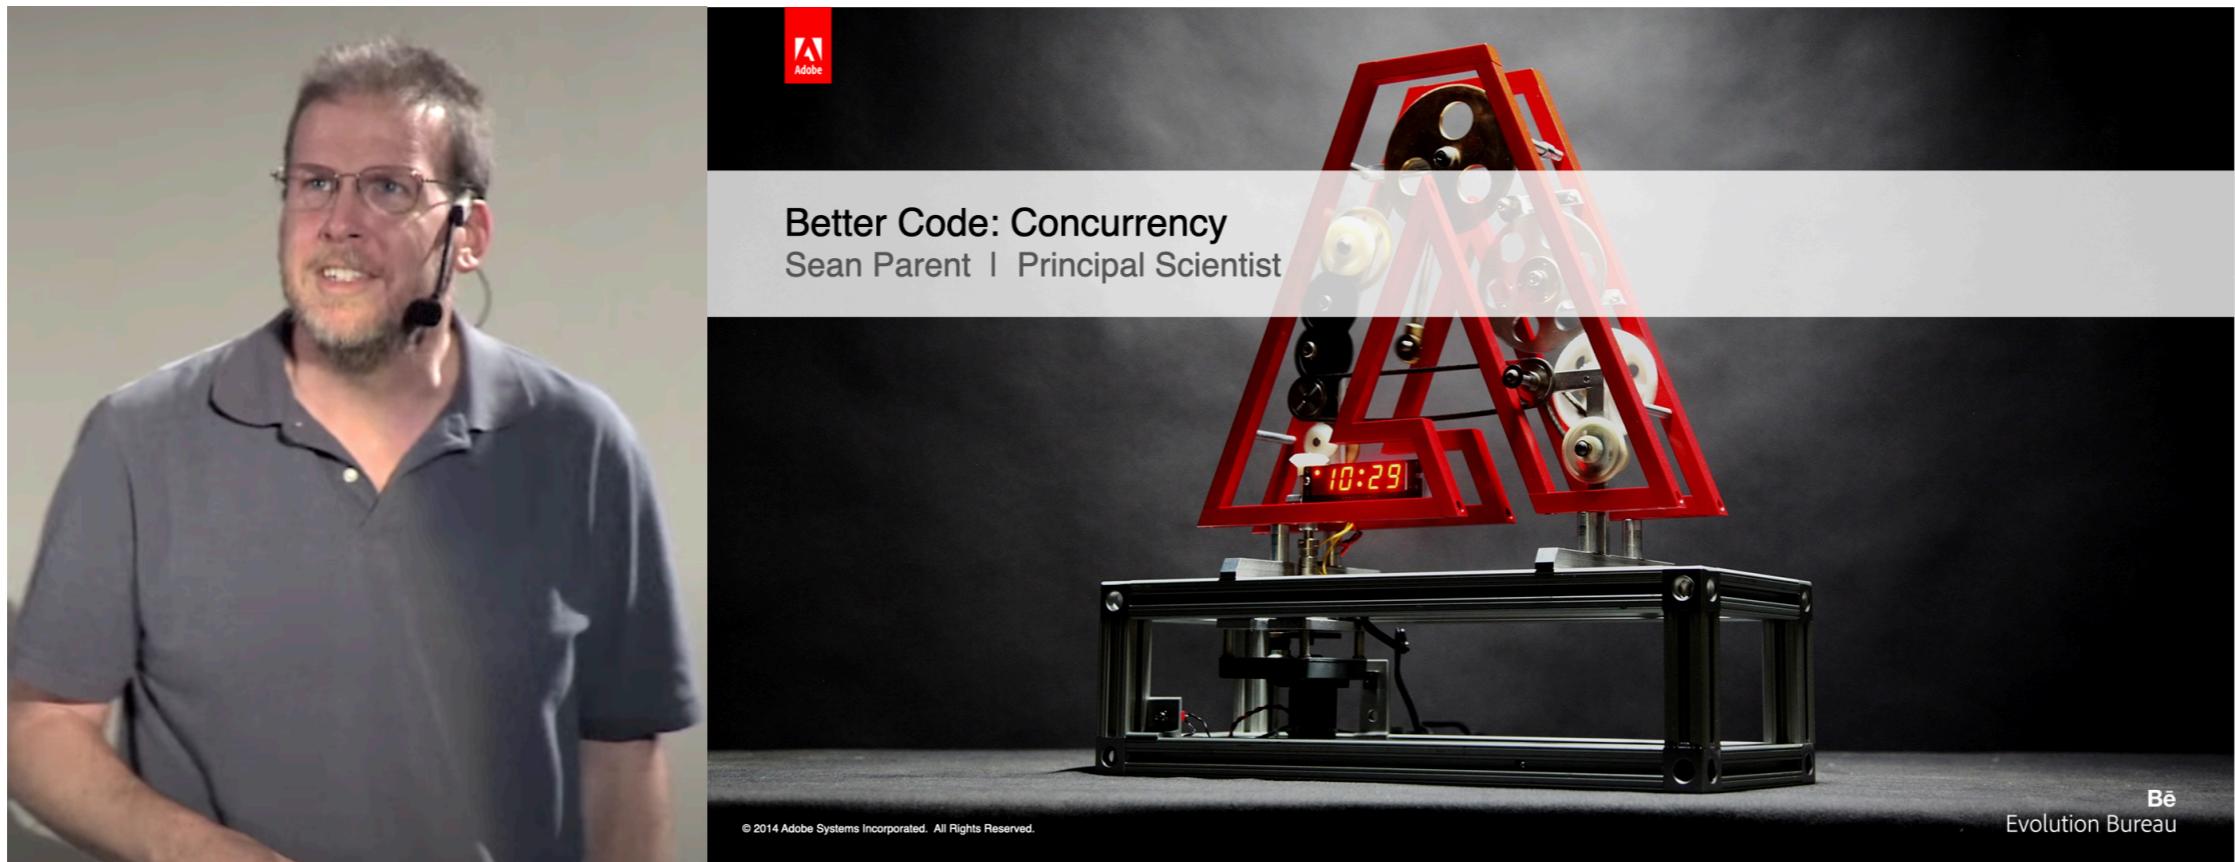
If a thread can't `pop()` a task from its own queue, because its empty or busy, then it will try to `pop()` a task from a different queue.

Recommendation

Task System - Better Design

In fact, do not listen to me!

Instead listen to the Sean Parent's talk “Better Code: Concurrency”, in which he discusses and implements properly such a task system.



Question

What else is wrong with the above code?

- Performance of the task system is very poor.

Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.

```
namespace atl
{
    class task_system
    {
        public:
            task_system();
            ~task_system();

            template <typename callable>
            auto async (callable&& task) -> void
            {
                queue.push(std::forward<callable>(task));
            }
    };
}
```

What is a task?

Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.

```
namespace atl
{
    class task_system
    {
public:
    task_system();
    ~task_system();

    template <typename callable>
    auto async (callable&& task) -> void
    {
        queue.push(std::forward<callable>(task));
    }
};
```

What is a task?

Well, it's something we can push into the task queue.

Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.

```
namespace atl
{
    class task_queue
    {
        private:
            std::deque<std::function<void()>> tasks;

        public:
            template <typename callable>
            auto push (callable&& task) -> void
            {
                {
                    auto lock = std::lock_guard<std::mutex> { mutex };
                    tasks.emplace_back(std::forward<callable>(task));
                }
                ready.notify_one();
            }
    };
}
```

OK. What can be pushed into the task queue?

Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.

```
namespace atl
{
    class task_queue
    {
        private:
            std::deque<std::function<void()>> tasks;

        public:
            template <typename callable>
            auto push (callable&& task) -> void
            {
                {
                    auto lock = std::lock_guard<std::mutex> { mutex };
                    tasks.emplace_back(std::forward<callable>(task));
                }
                ready.notify_one();
            }
    };
}
```

OK. What can be pushed into the task queue?
Well, something convertible to a `std::function<void()>`.

Recommendation

Task System - Better Design

- Performance of the task system is very poor.
- Interface of the task system is not precise.

```
namespace atl
{
    template <typename type>
    concept task = std::is_invocable_r_v<void, type>;

    class task_system
    {
        ...
    public:
        task_system();
        ~task_system();

        auto async (task auto&& t) -> void
        { ... }
    };
}
```

This thinking is all wrong!

Instead of expressing interface in terms of implementation details,
we need to define the **task concept**,
which specifies what does it mean to be a task.

Recommendation

Task System - Better Design

Again, do not listen to me!

Instead listen to the Bjarne Stroustrup's talk about interfaces of templates
“Concepts: The Future of Generic Programming”.



Concepts:
The Future of Generic Programming
(the future is here)

Bjarne Stroustrup
Morgan Stanley and Columbia University
www.stroustrup.com



Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.

Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.
- Implementation details of the task system are exposed to clients.

```
#include <atl/task_system.hxx>

namespace phy_sim
{
    auto function () -> void
    {
        auto task_system = atl::task_system {};
        auto task_queue = atl::task_queue {};
    }
}
```

The physics simulation executable
has access to the task_queue type.

```
#include <atl/task_queue.hxx>

namespace atl
{
    class task_system
    {
        ...
    private:
        task_queue queue;
    };
}
```

Question

What else is wrong with the above code?

- Performance of the task system is very poor.
- Interface of the task system is not precise.
- Implementation details of the task system are exposed to clients.

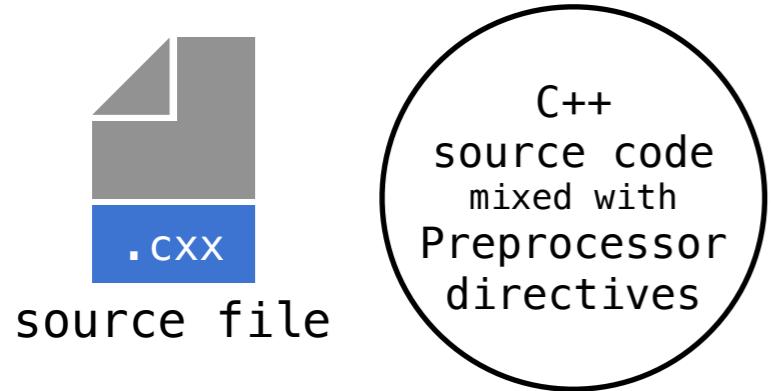
```
#include <optional>
namespace atl
{
    class task_queue
    {
        auto pop () -> std::optional<std::function<void()>>;
    };
}

#include <atl/task_system.hxx>
namespace phy_sim
{
    auto function () -> void
    {
        auto opt_int = std::optional<int> { 7 };
    }
}
```

The physics simulation executable has also access to the std::optional template.

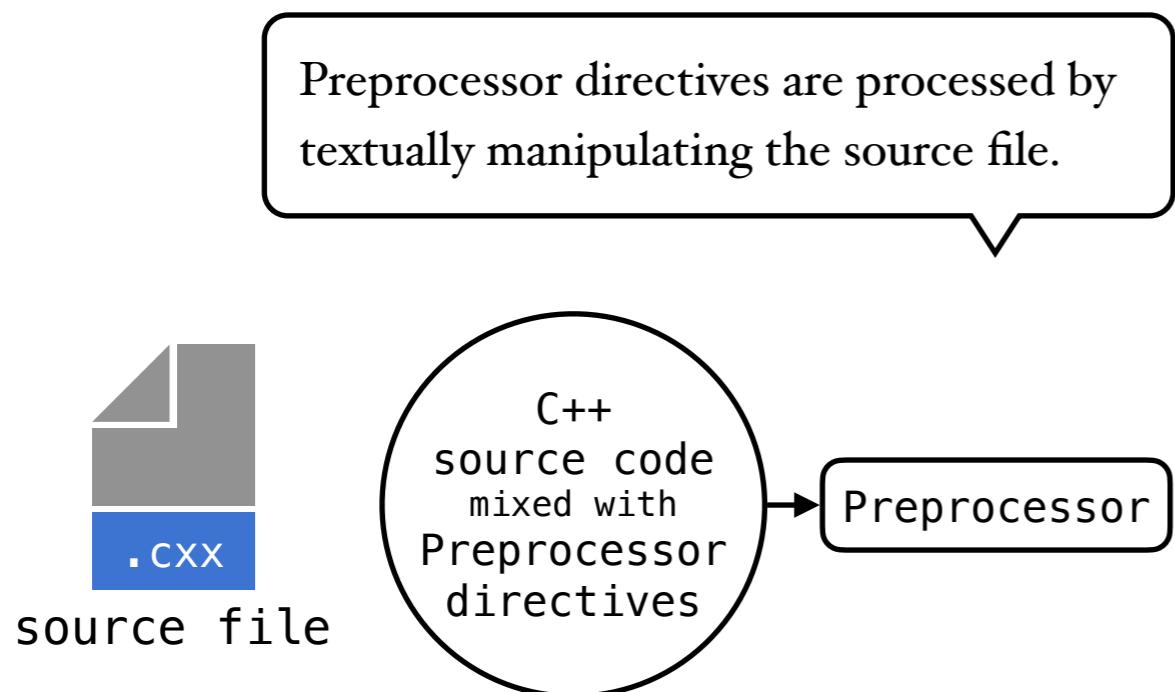
Textual Inclusion Model

Preprocessing & Compilation



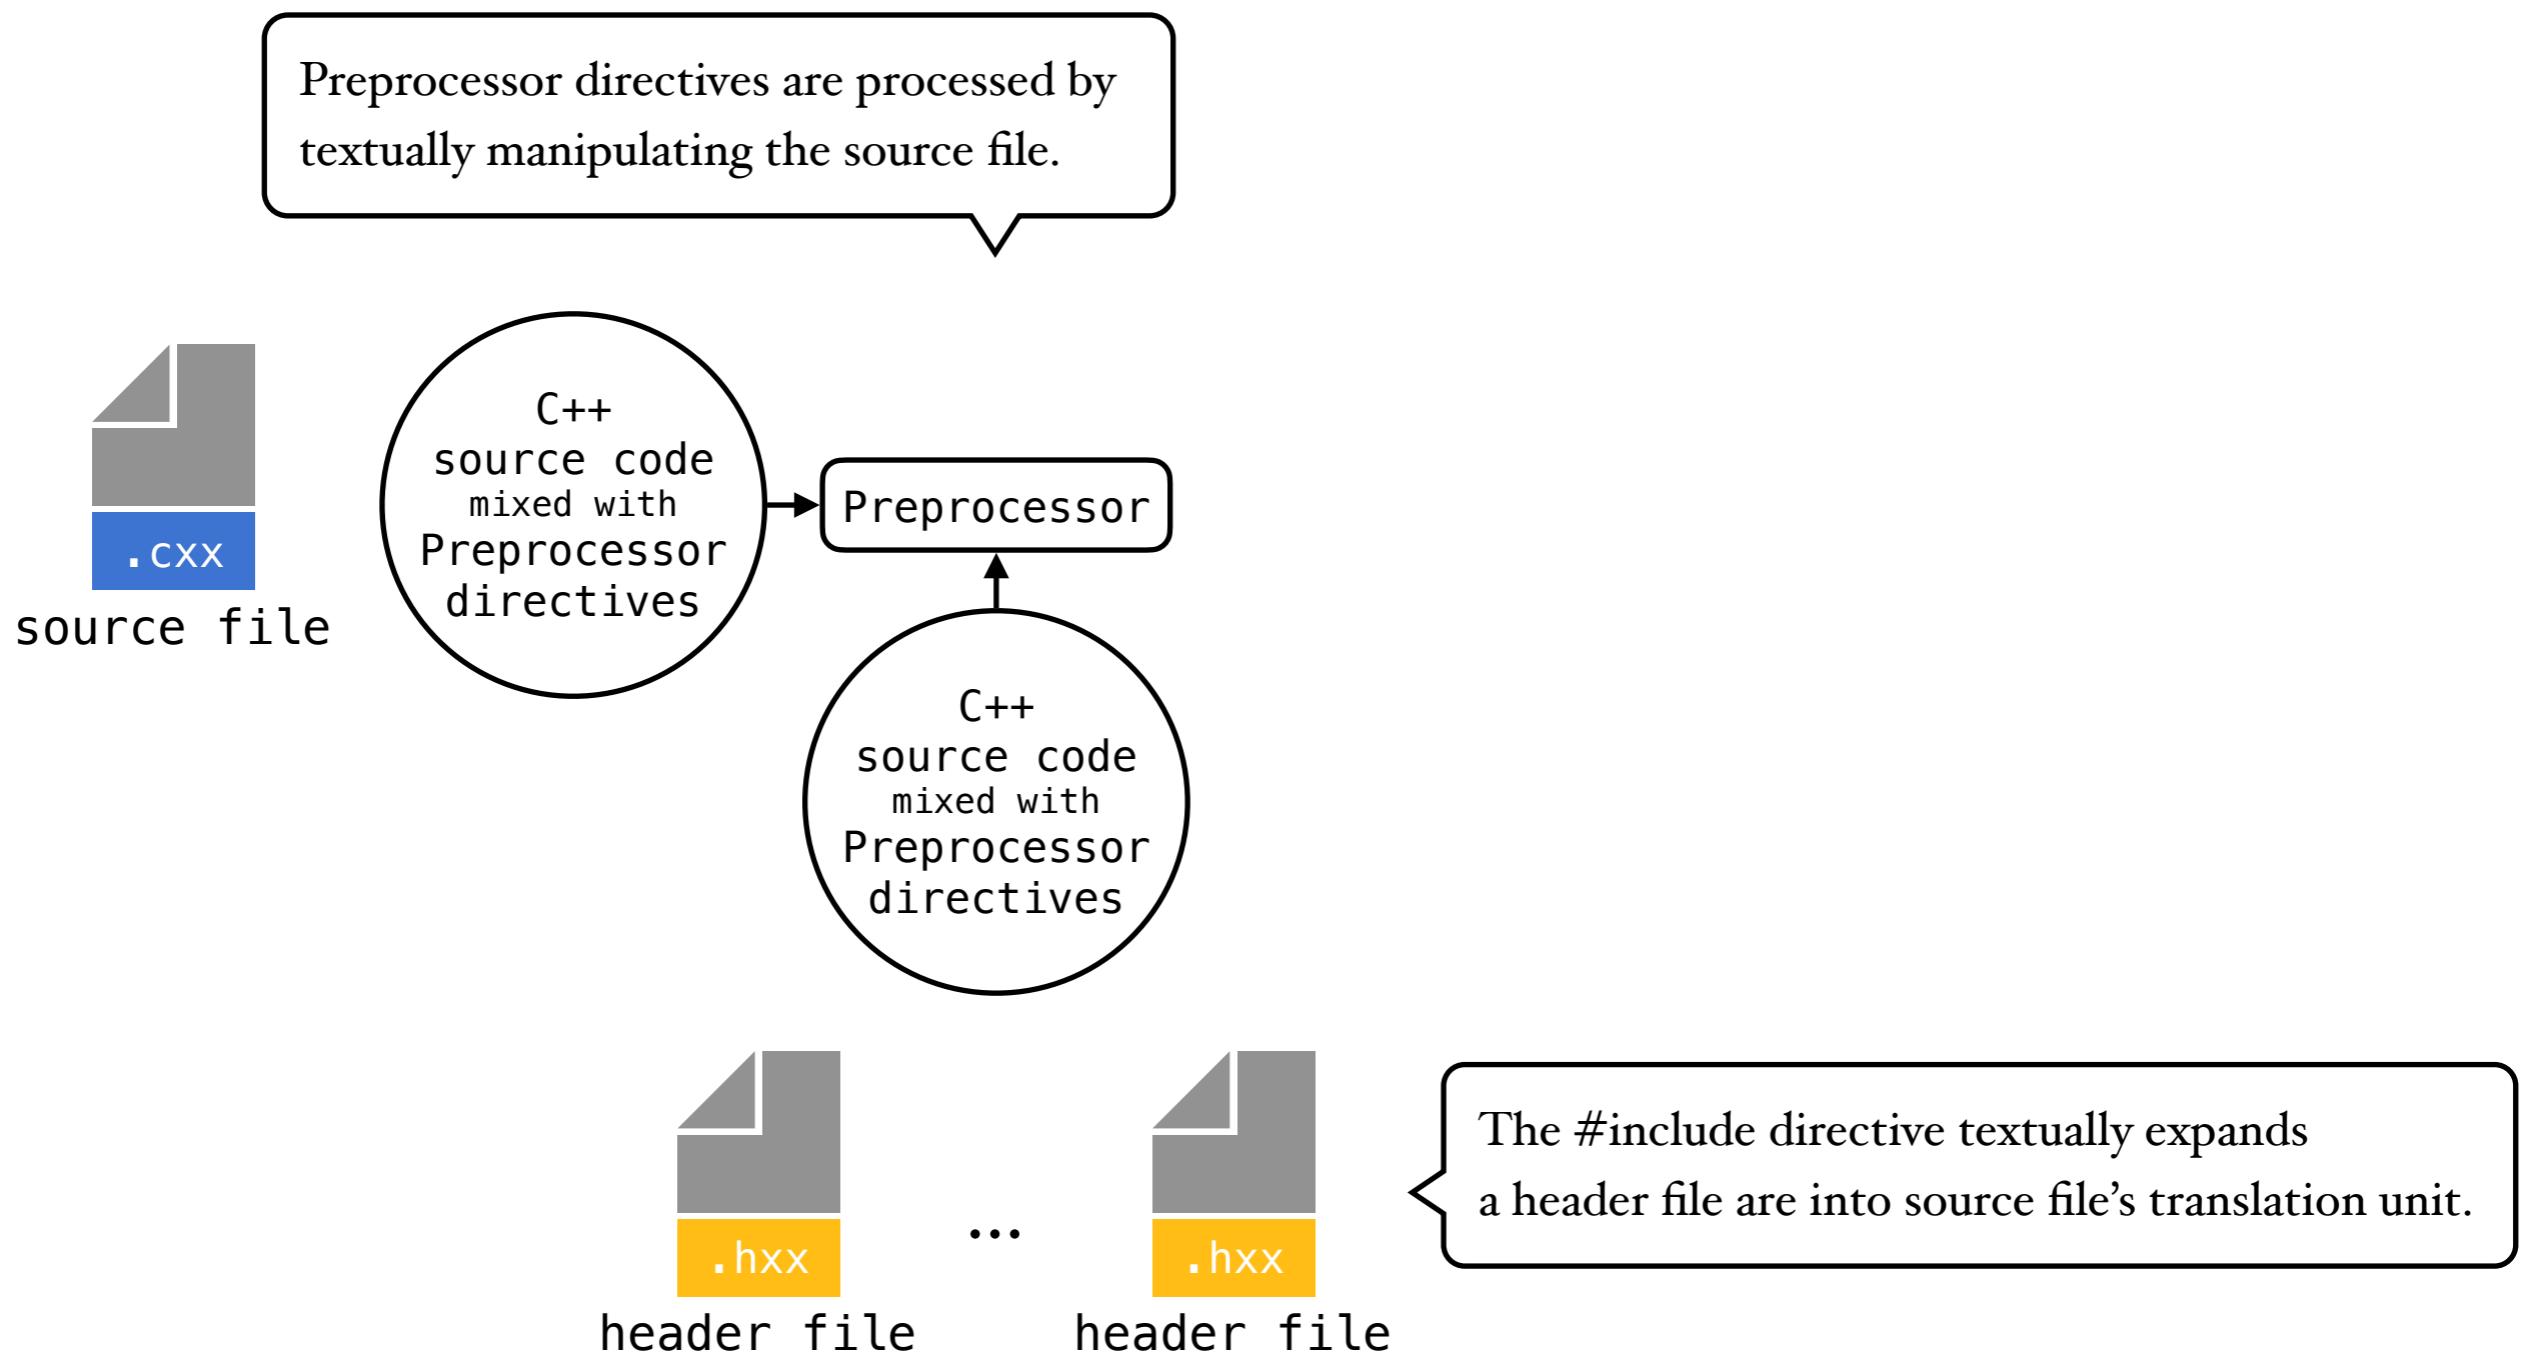
Textual Inclusion Model

Preprocessing & Compilation



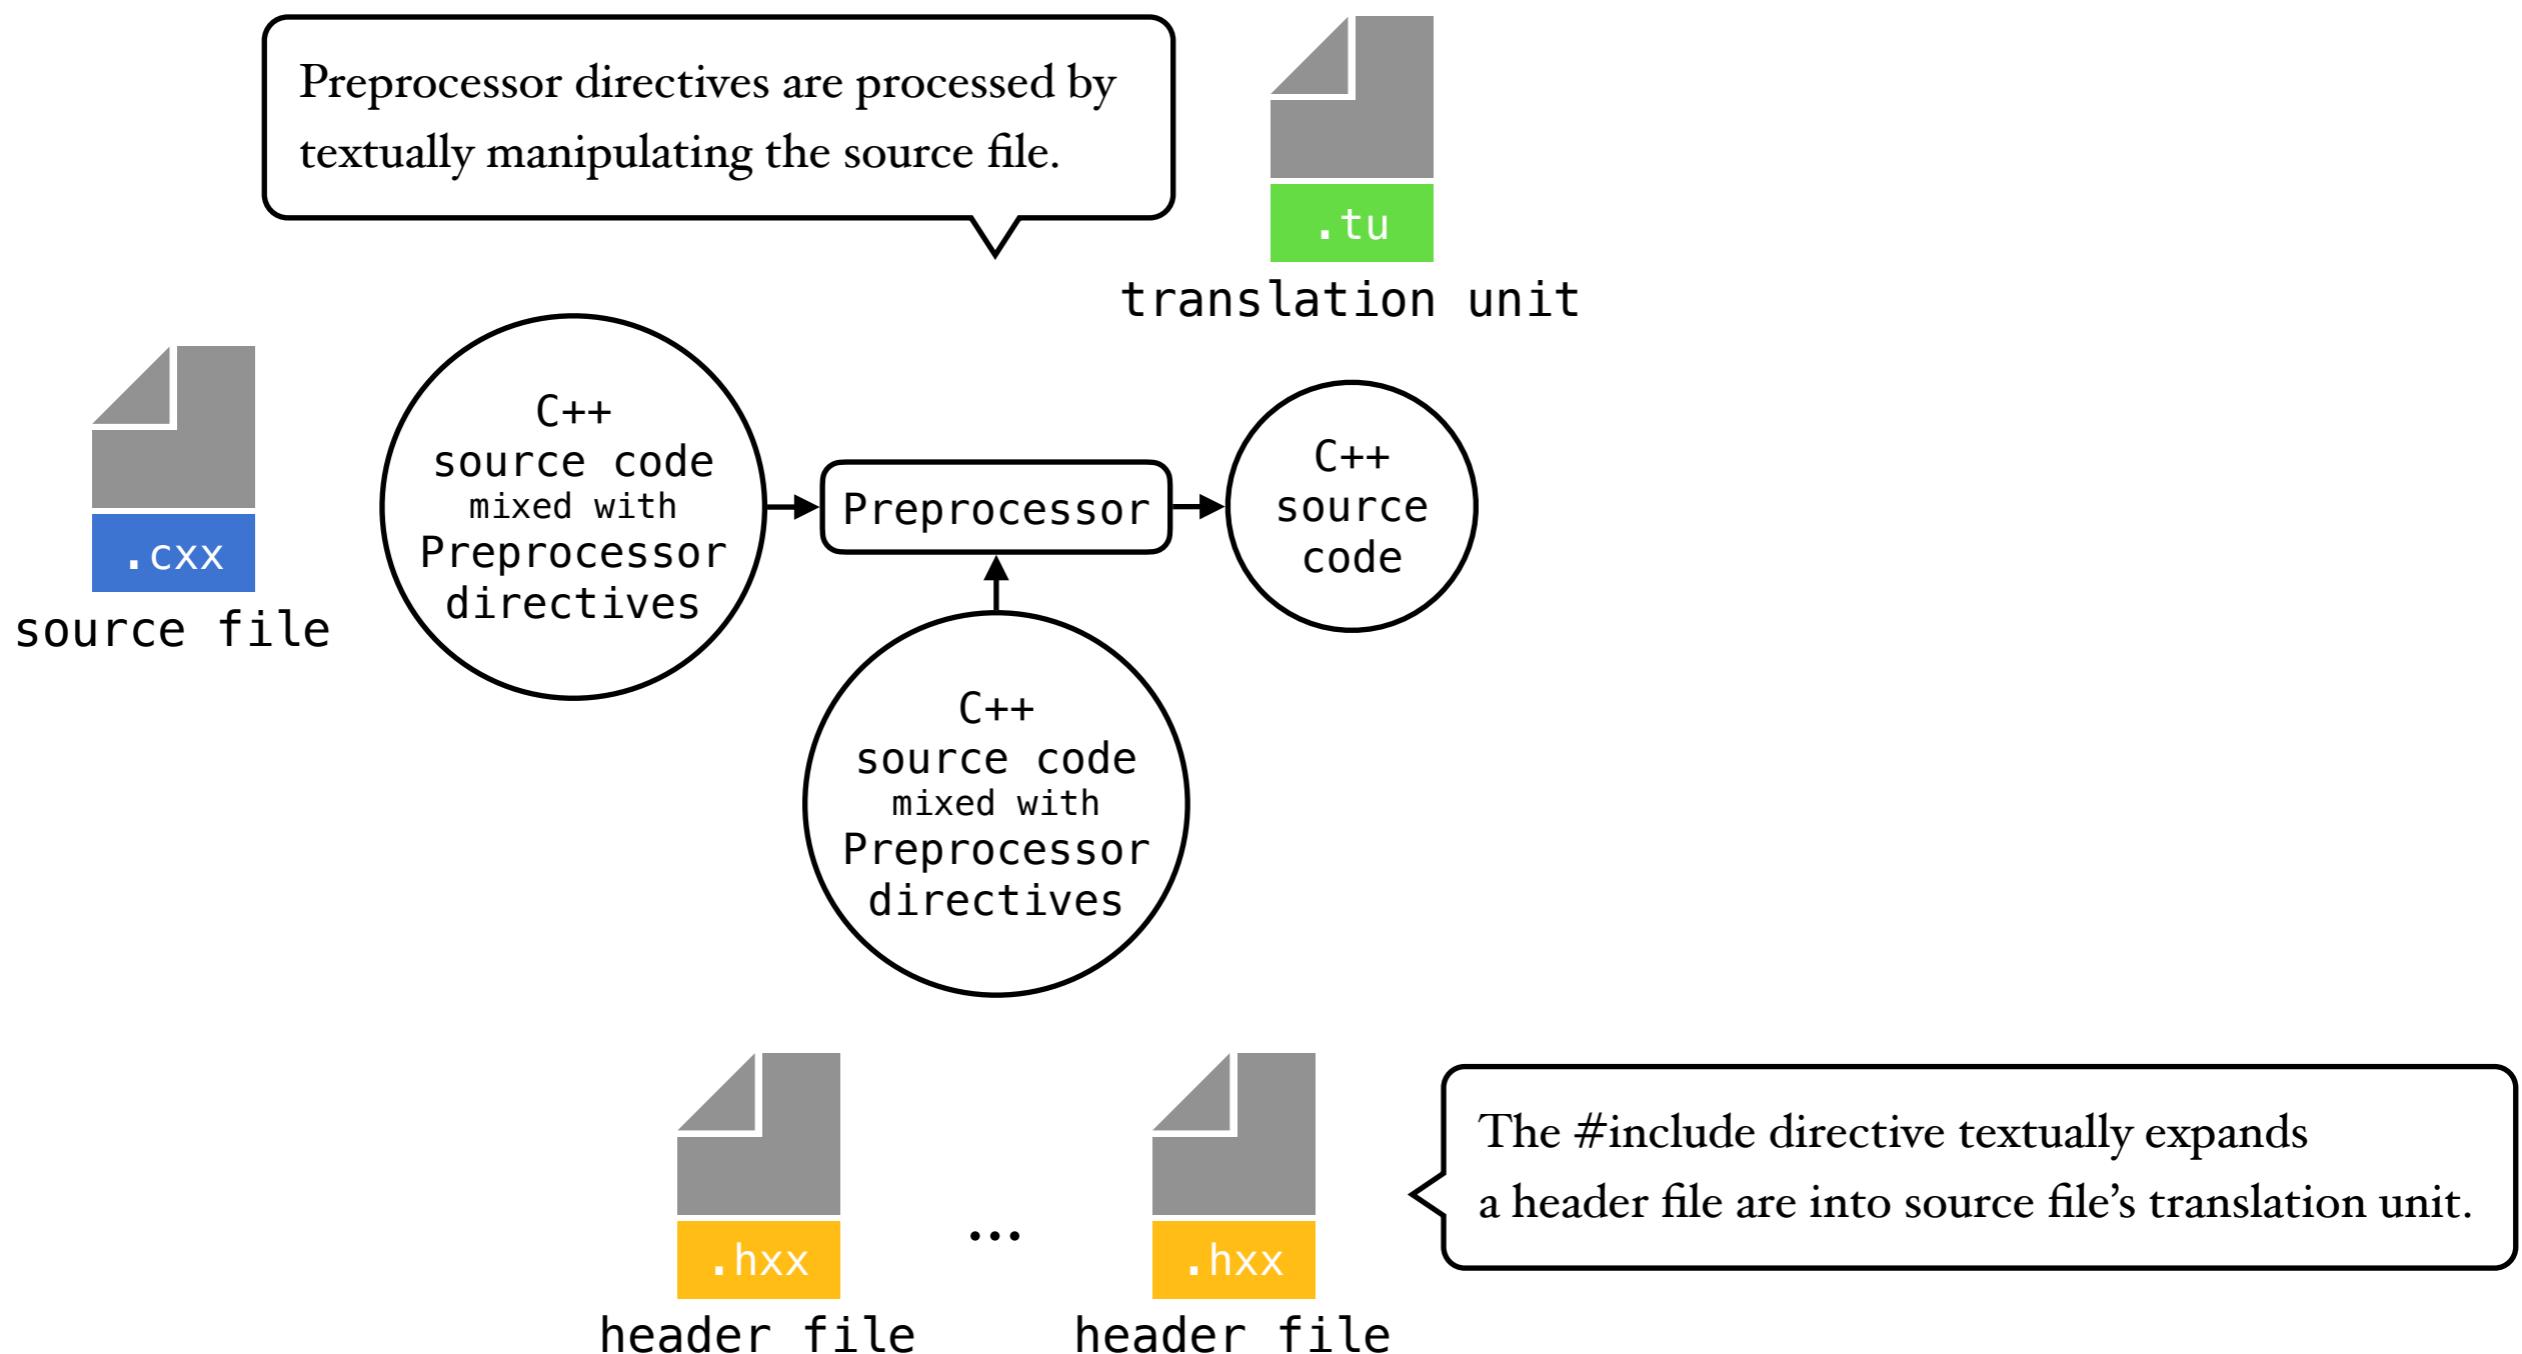
Textual Inclusion Model

Preprocessing & Compilation



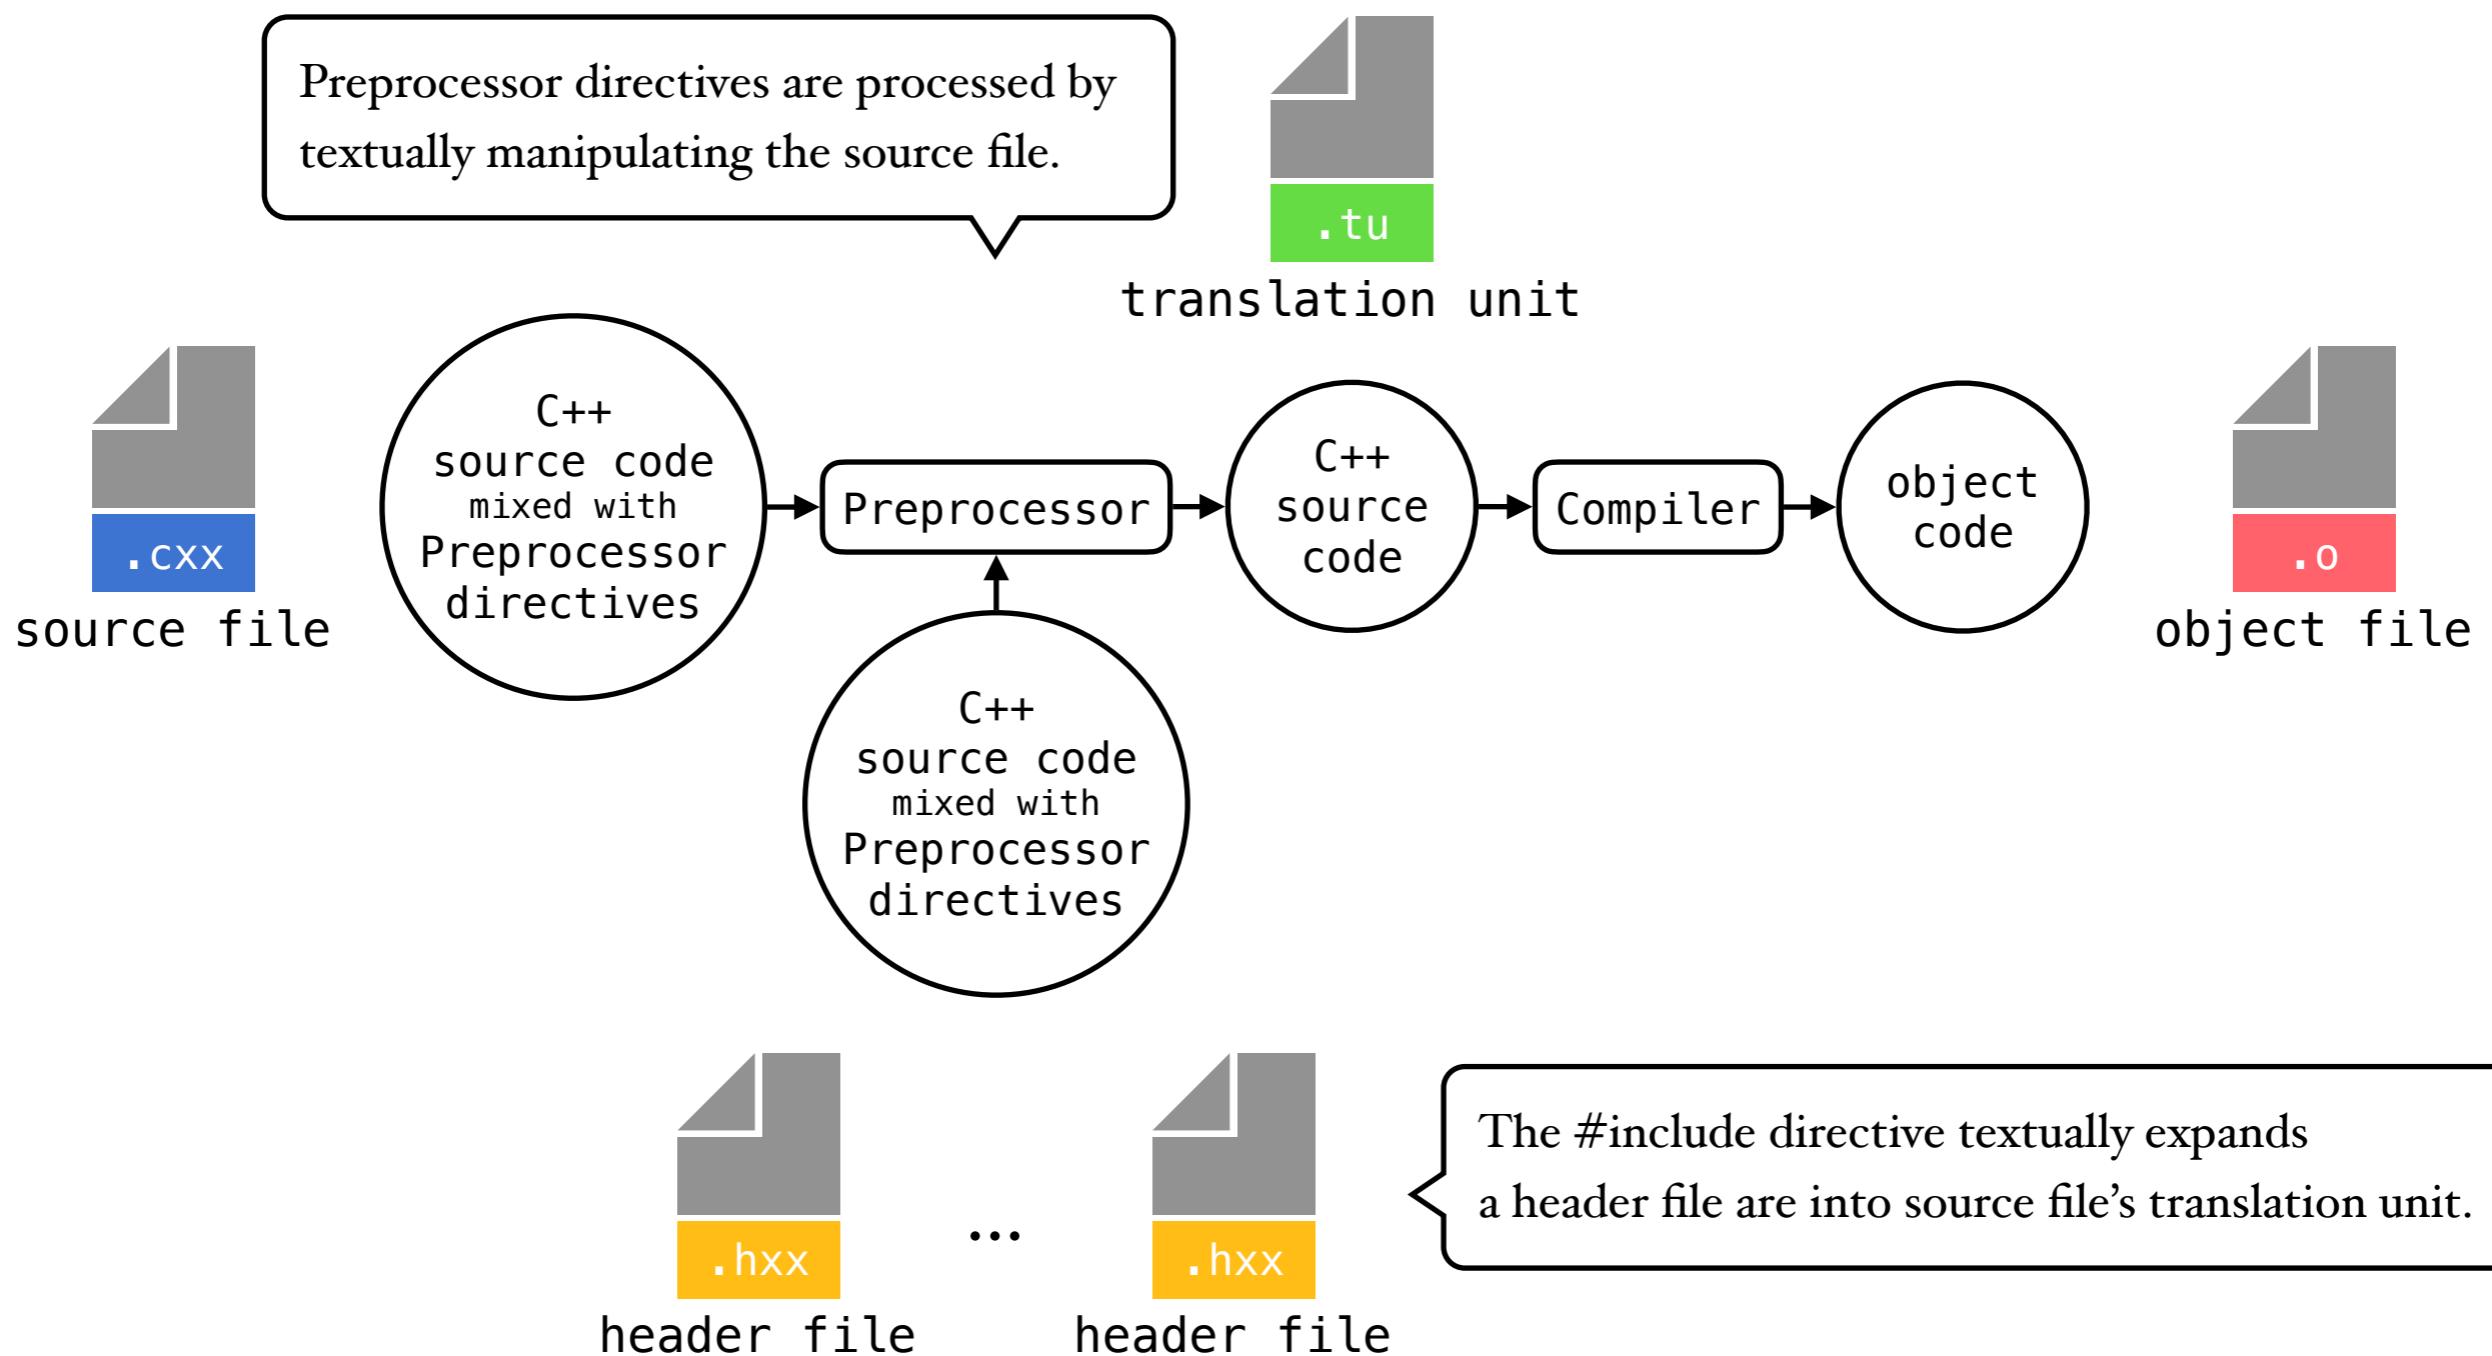
Textual Inclusion Model

Preprocessing & Compilation



Textual Inclusion Model

Preprocessing & Compilation



Textual Inclusion Model

Translation Unit

// task_system.hxx #ifndef ATL_TASK_SYSTEM #define ATL_TASK_SYSTEM #include <atl/task_queue.hxx> ... namespace atl { class task_system { ... }; } #endif	// task_queue.hxx #ifndef ATL_TASK_QUEUE #define ATL_TASK_QUEUE ... #include <optional> namespace atl { class task_queue { ... }; } #endif
// physics_simulation.cxx #include <atl/task_system.hxx> ... namespace phy_sim { ... } auto main () -> int { ... }	// optional #ifndef LIBCPP_OPTIONAL #define LIBCPP_OPTIONAL ... namespace std { ... template<class T> class optional { ... }; } #endif

Textual Inclusion Model

Translation Unit

// task_system.hxx #ifndef ATL_TASK_SYSTEM #define ATL_TASK_SYSTEM #include <atl/task_queue.hxx> ... namespace atl { class task_system { ... }; } #endif	// task_queue.hxx #ifndef ATL_TASK_QUEUE #define ATL_TASK_QUEUE ... #include <optional> namespace atl { class task_queue { ... }; } #endif	// physics_simulation.tu namespace std { ... template <class T> class optional { ... }; } namespace atl { class task_queue { ... }; } ... namespace atl { class task_system { ... }; } ... namespace phy_sim { ... } auto main () -> int { ... }
// physics_simulation.cxx #include <atl/task_system.hxx> ... namespace phy_sim { ... } auto main () -> int { ... }	// optional #ifndef LIBCPP_OPTIONAL #define LIBCPP_OPTIONAL ... namespace std { ... template<class T> class optional { ... }; } #endif	

Textual Inclusion Model

Translation Unit

```
$ clang++ -std=c++17 \
-I //scientific_repo/async_task_library/include \
-E //scientific_repo/physics_simulation/source/physics_simulation.cxx \
-o physics_simulation.tu
```



Textual Inclusion Model

Translation Unit

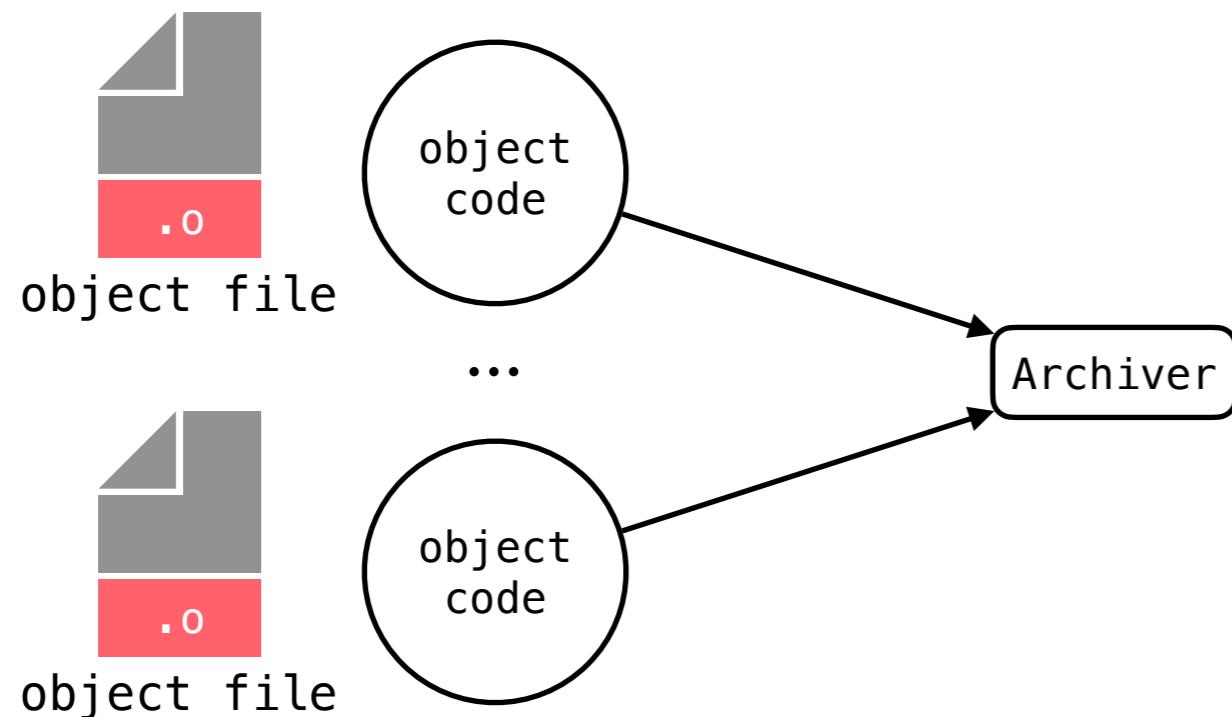
```
$ clang++ -std=c++17  
-I //scientific_repo/async_task_library/include  
-E //scientific_repo/physics_simulation/source/physics_simulation.cxx \  
-o physics_simulation.tu
```

Notice that, C++ compilers operate
on a different representation of
C++ programs than C++ programmers.

// task_system.hxx #ifndef ATL_TASK_SYSTEM #define ATL_TASK_SYSTEM #include <atl/task_queue.hxx> ... namespace atl { class task_system { ... }; } #endif	// task_queue.hxx #ifndef ATL_TASK_QUEUE #define ATL_TASK_QUEUE ... #include <optional> namespace atl { class task_queue { ... }; } #endif	→	// physics_simulation.tu namespace std { ... template <class T> class optional { ... }; } namespace atl { class task_queue { ... }; } ... namespace atl { class task_system { ... }; } ... namespace phy_sim { ... } auto main () -> int { ... }
-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	----------	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

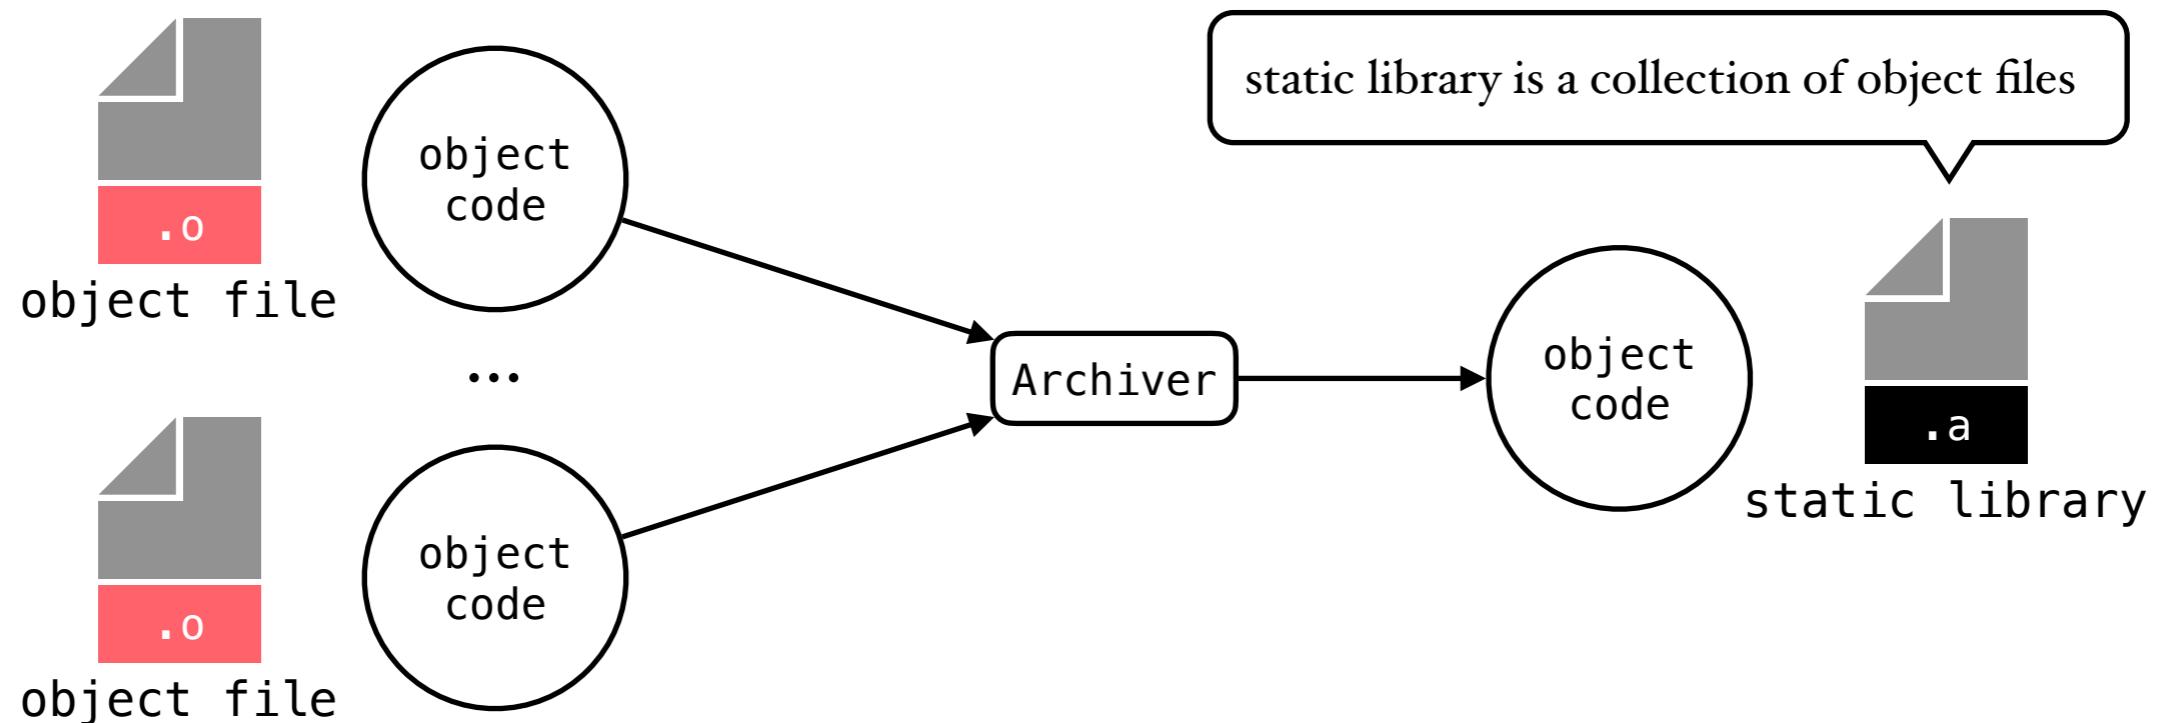
Textual Inclusion Model

Linking



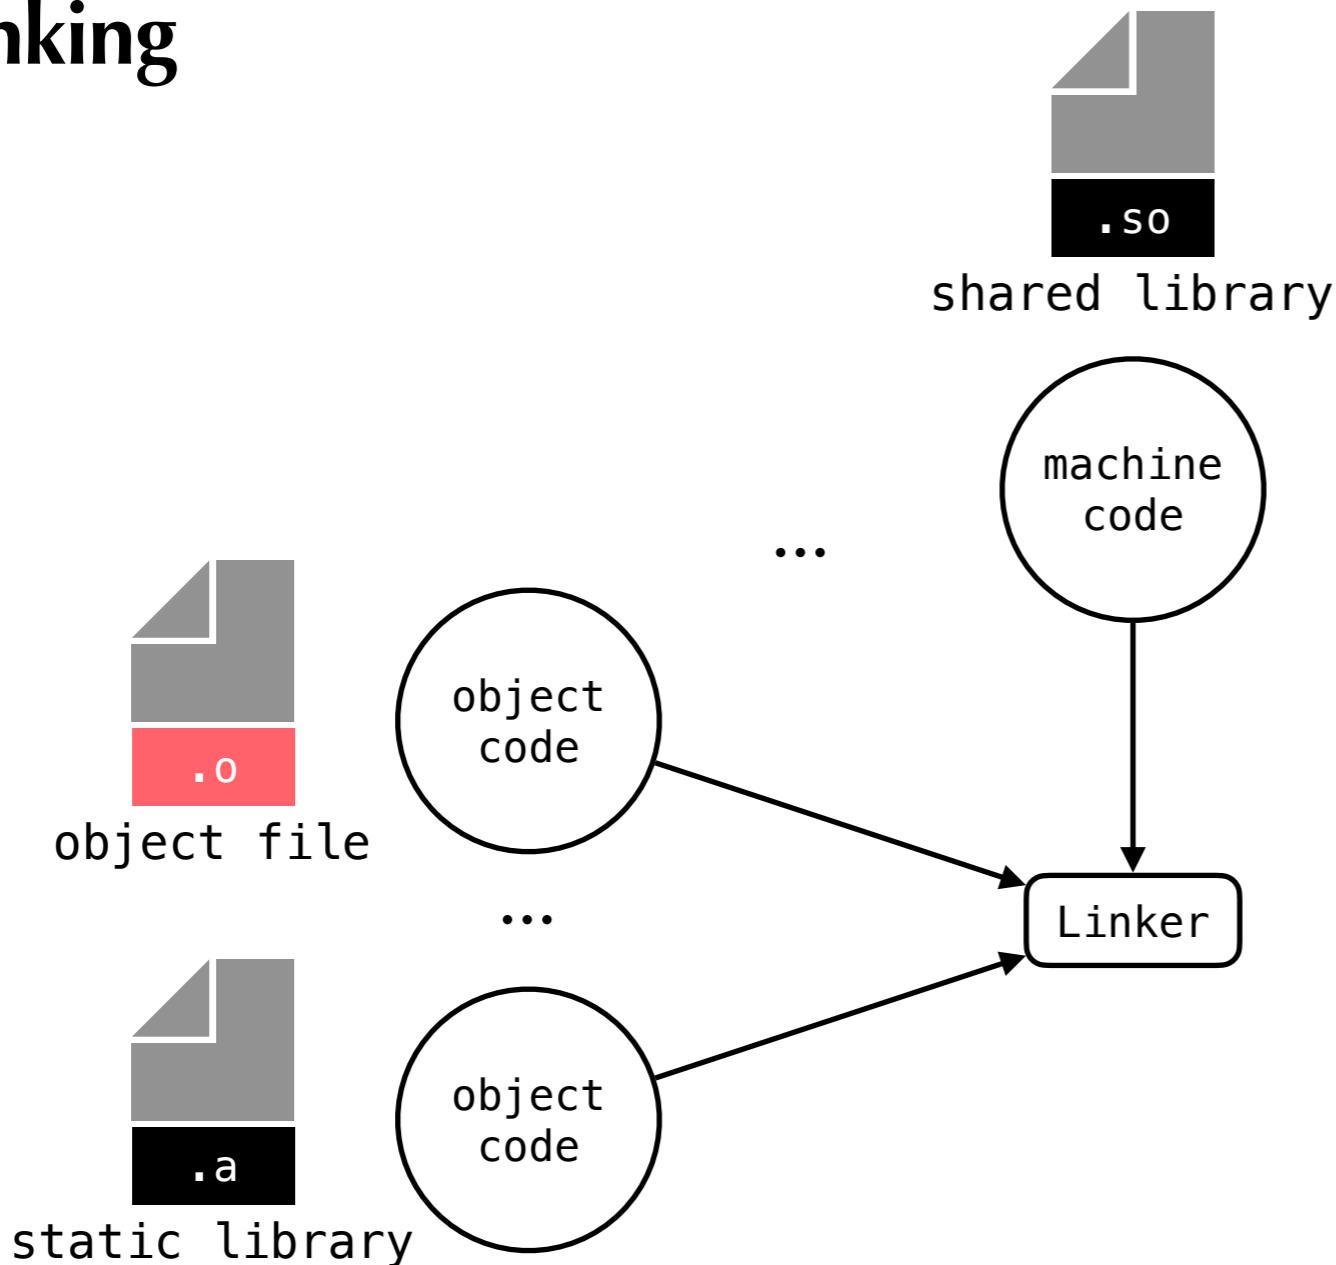
Textual Inclusion Model

Linking



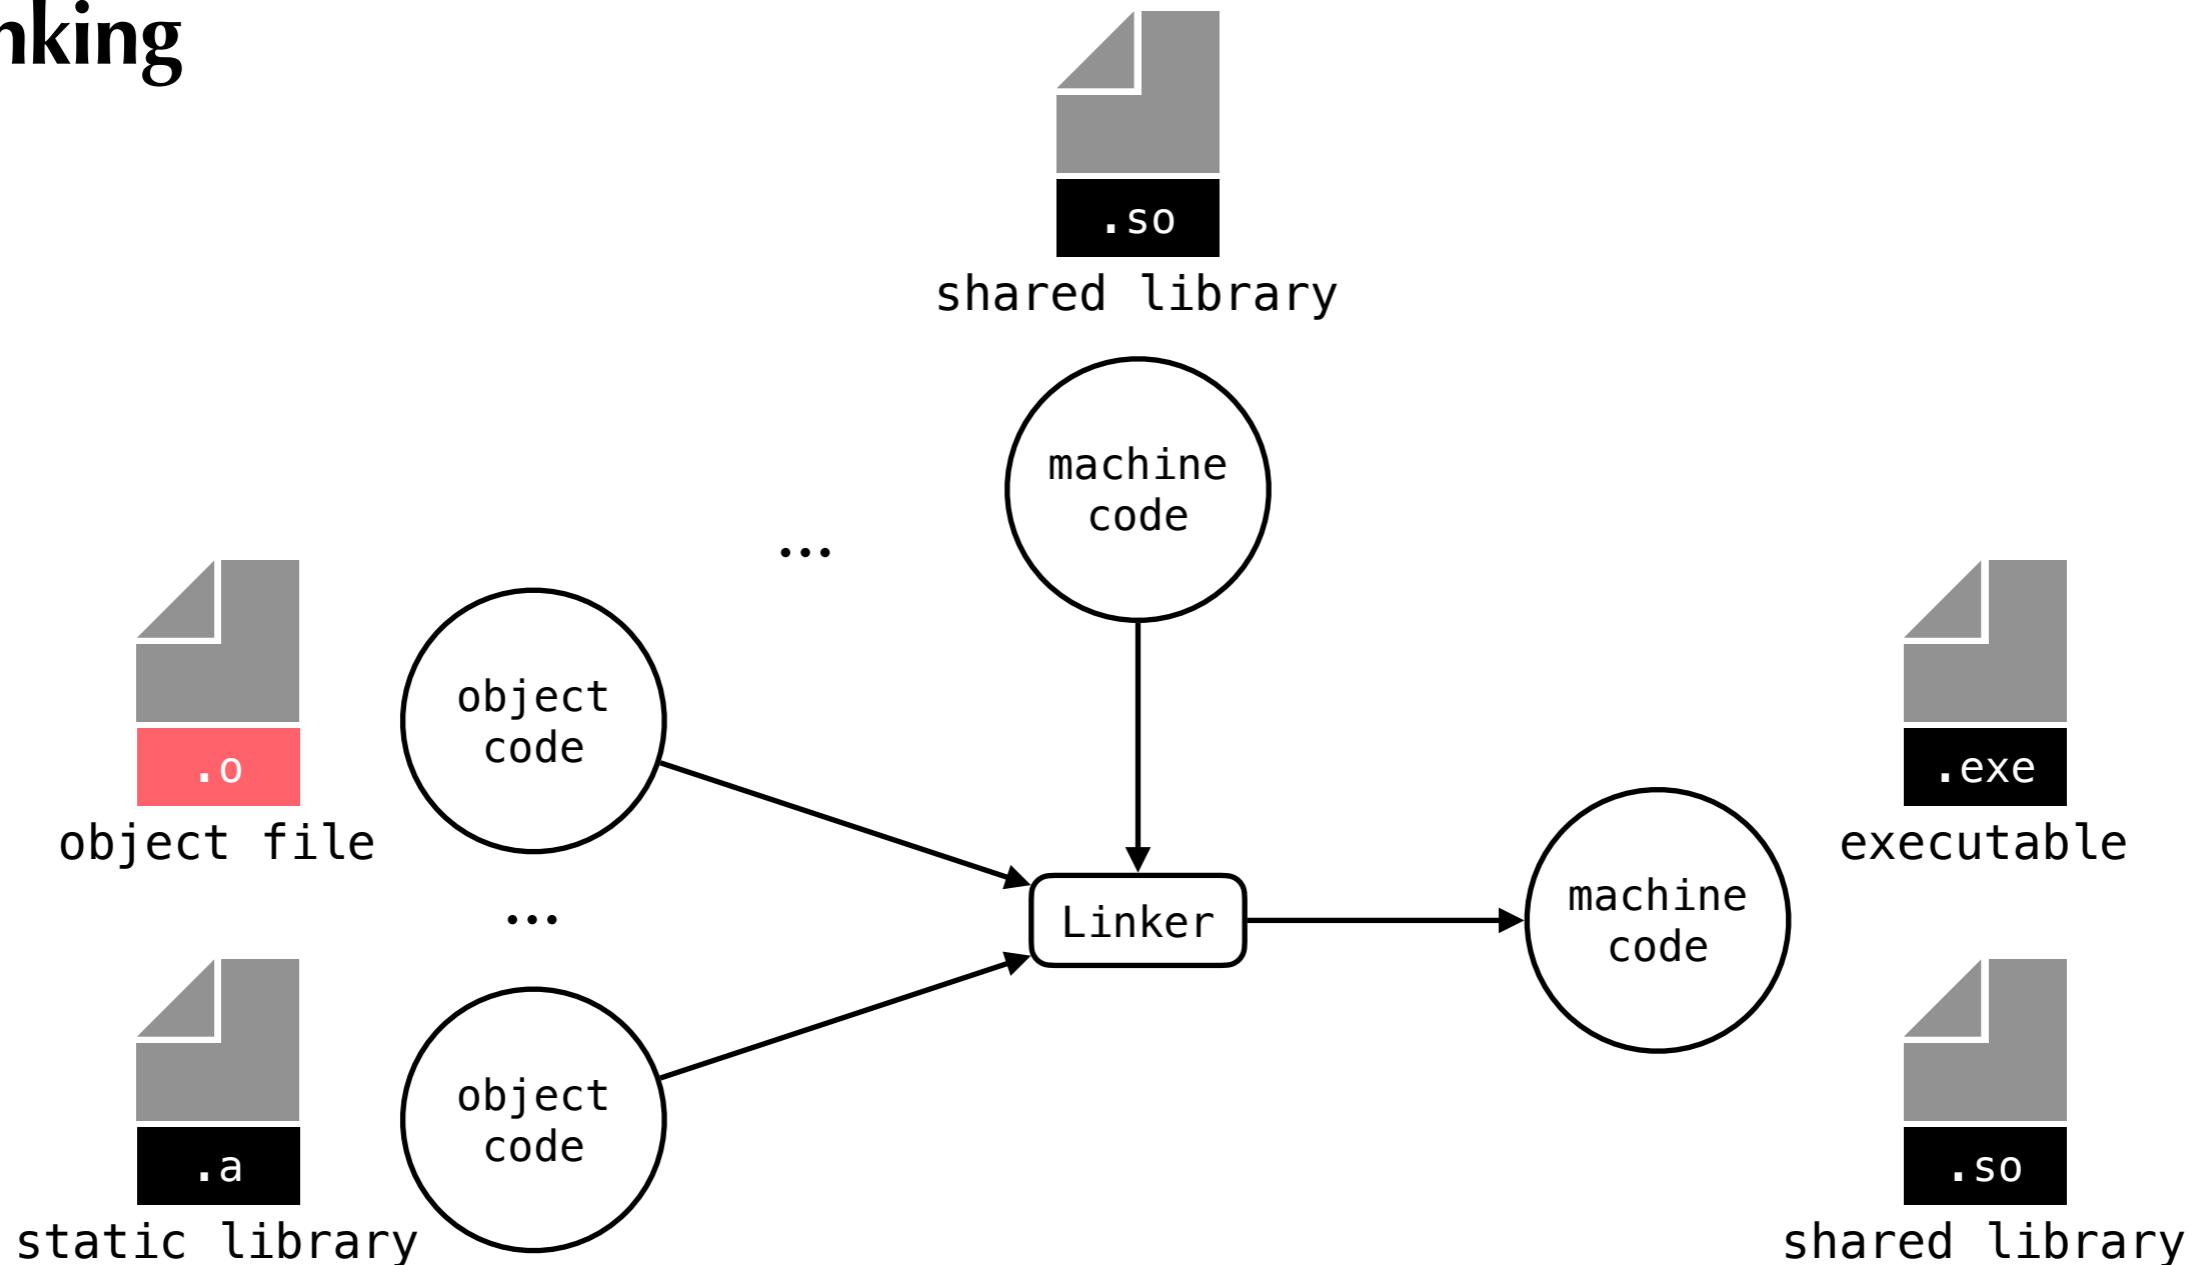
Textual Inclusion Model

Linking



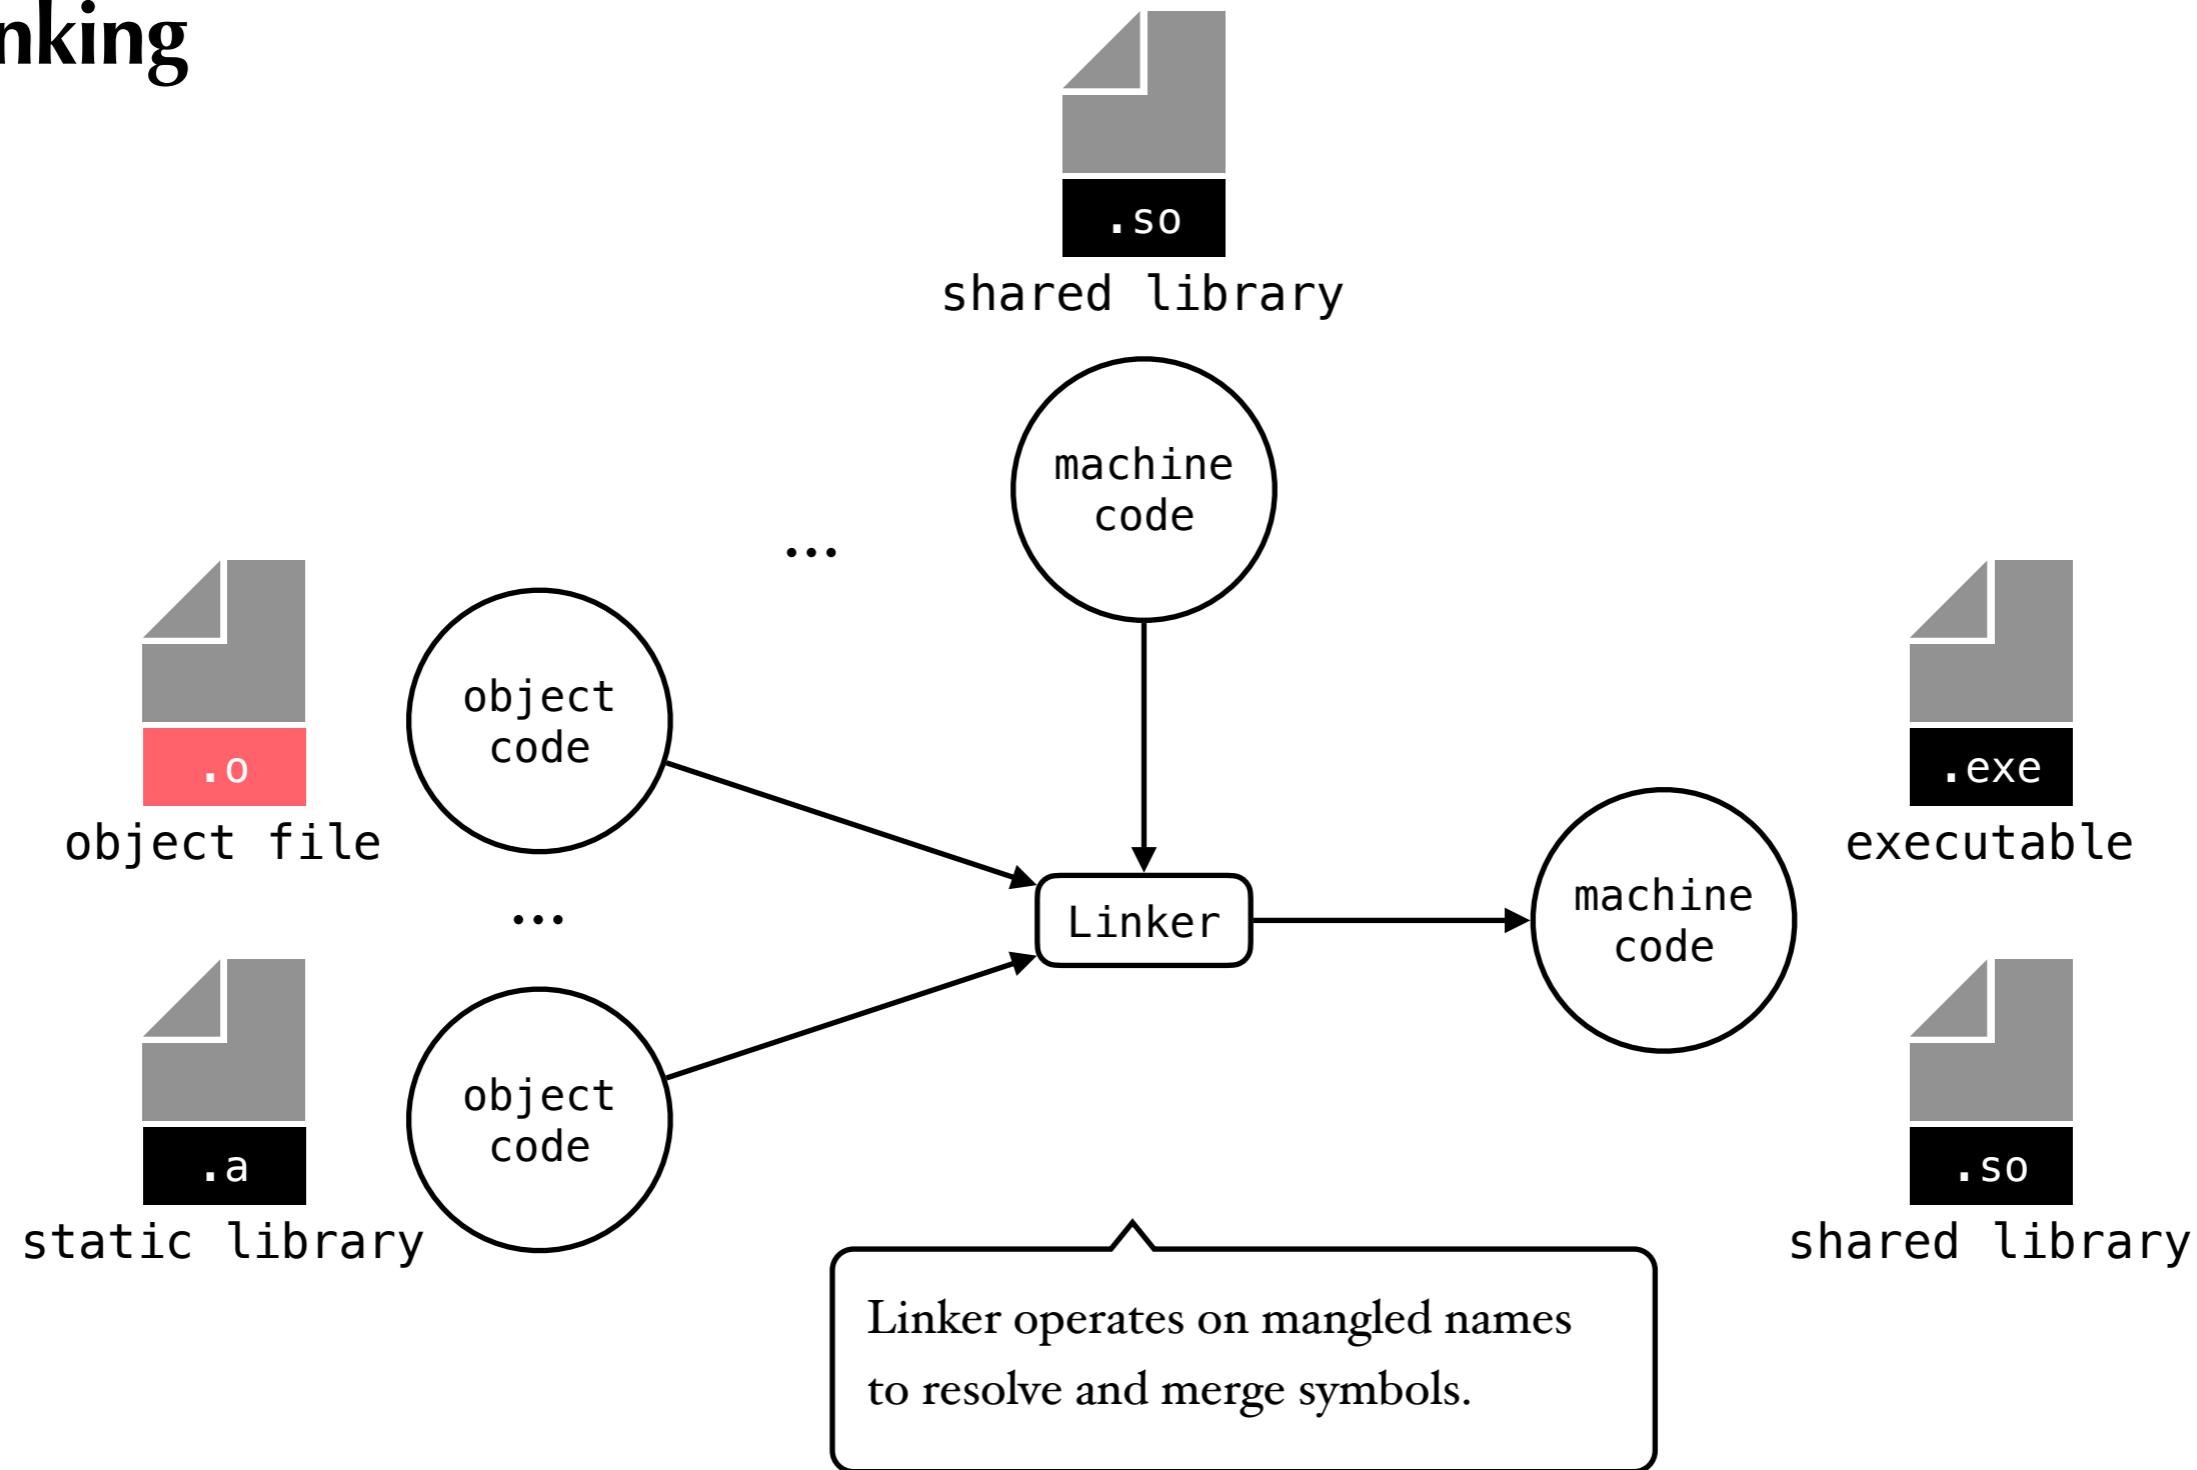
Textual Inclusion Model

Linking



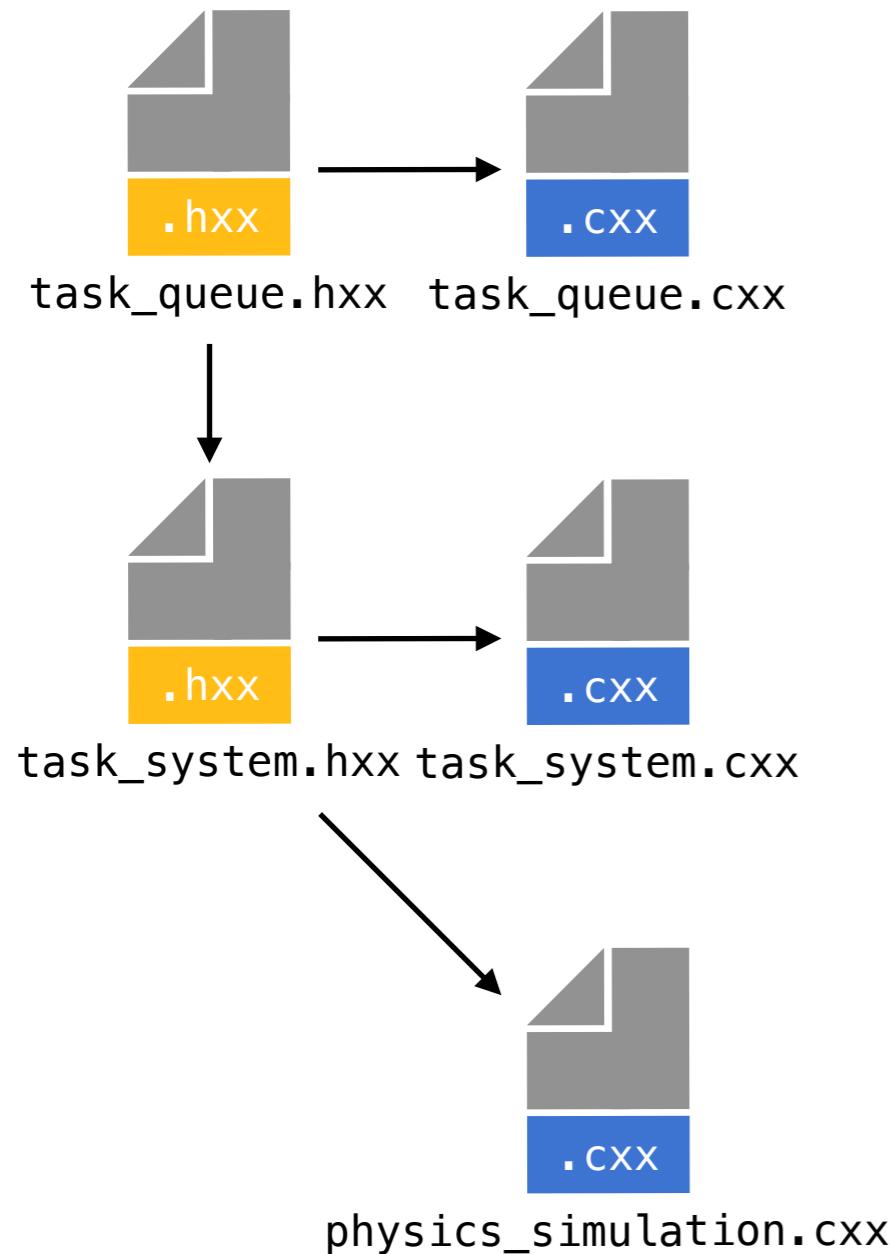
Textual Inclusion Model

Linking



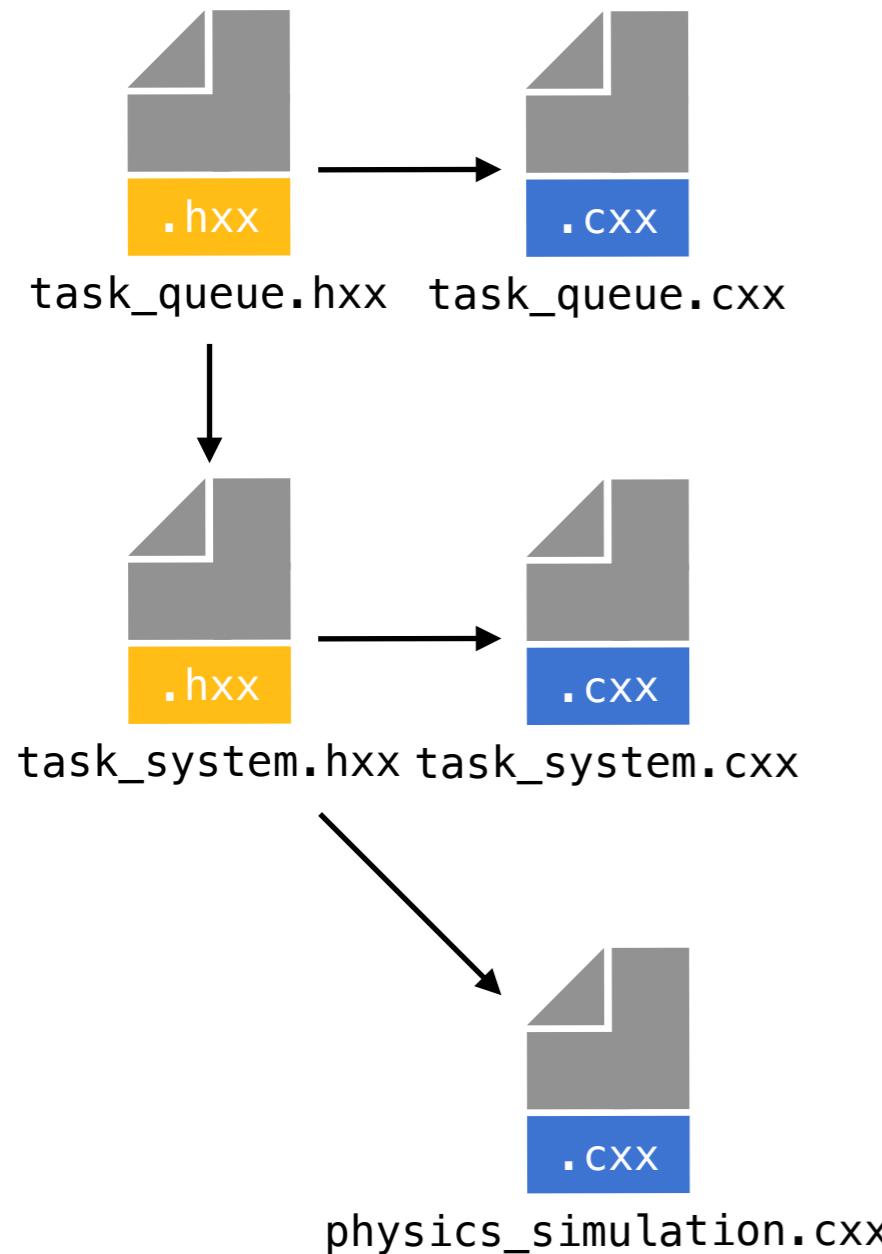
Textual Inclusion Model

Inclusion Graph



Textual Inclusion Model

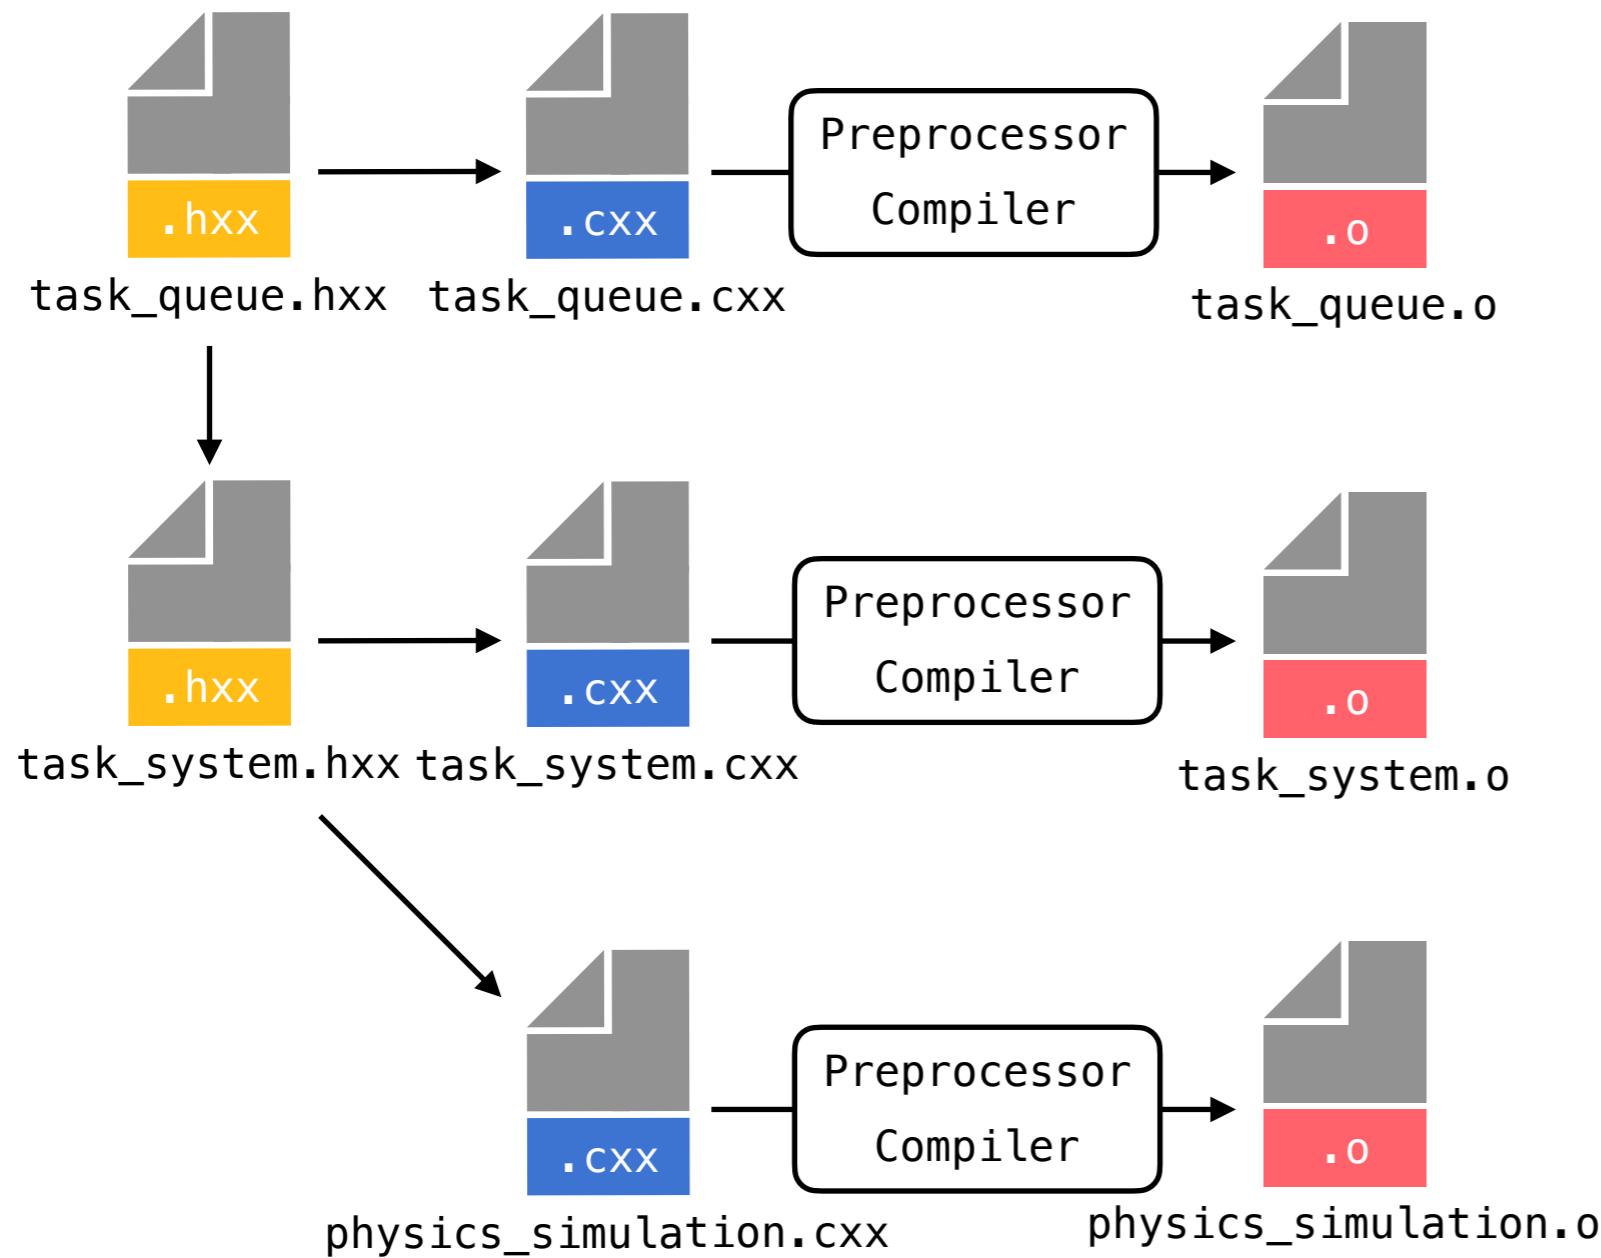
Inclusion Graph



This is very low-level view of the build process,
since it does not define dependencies between components.

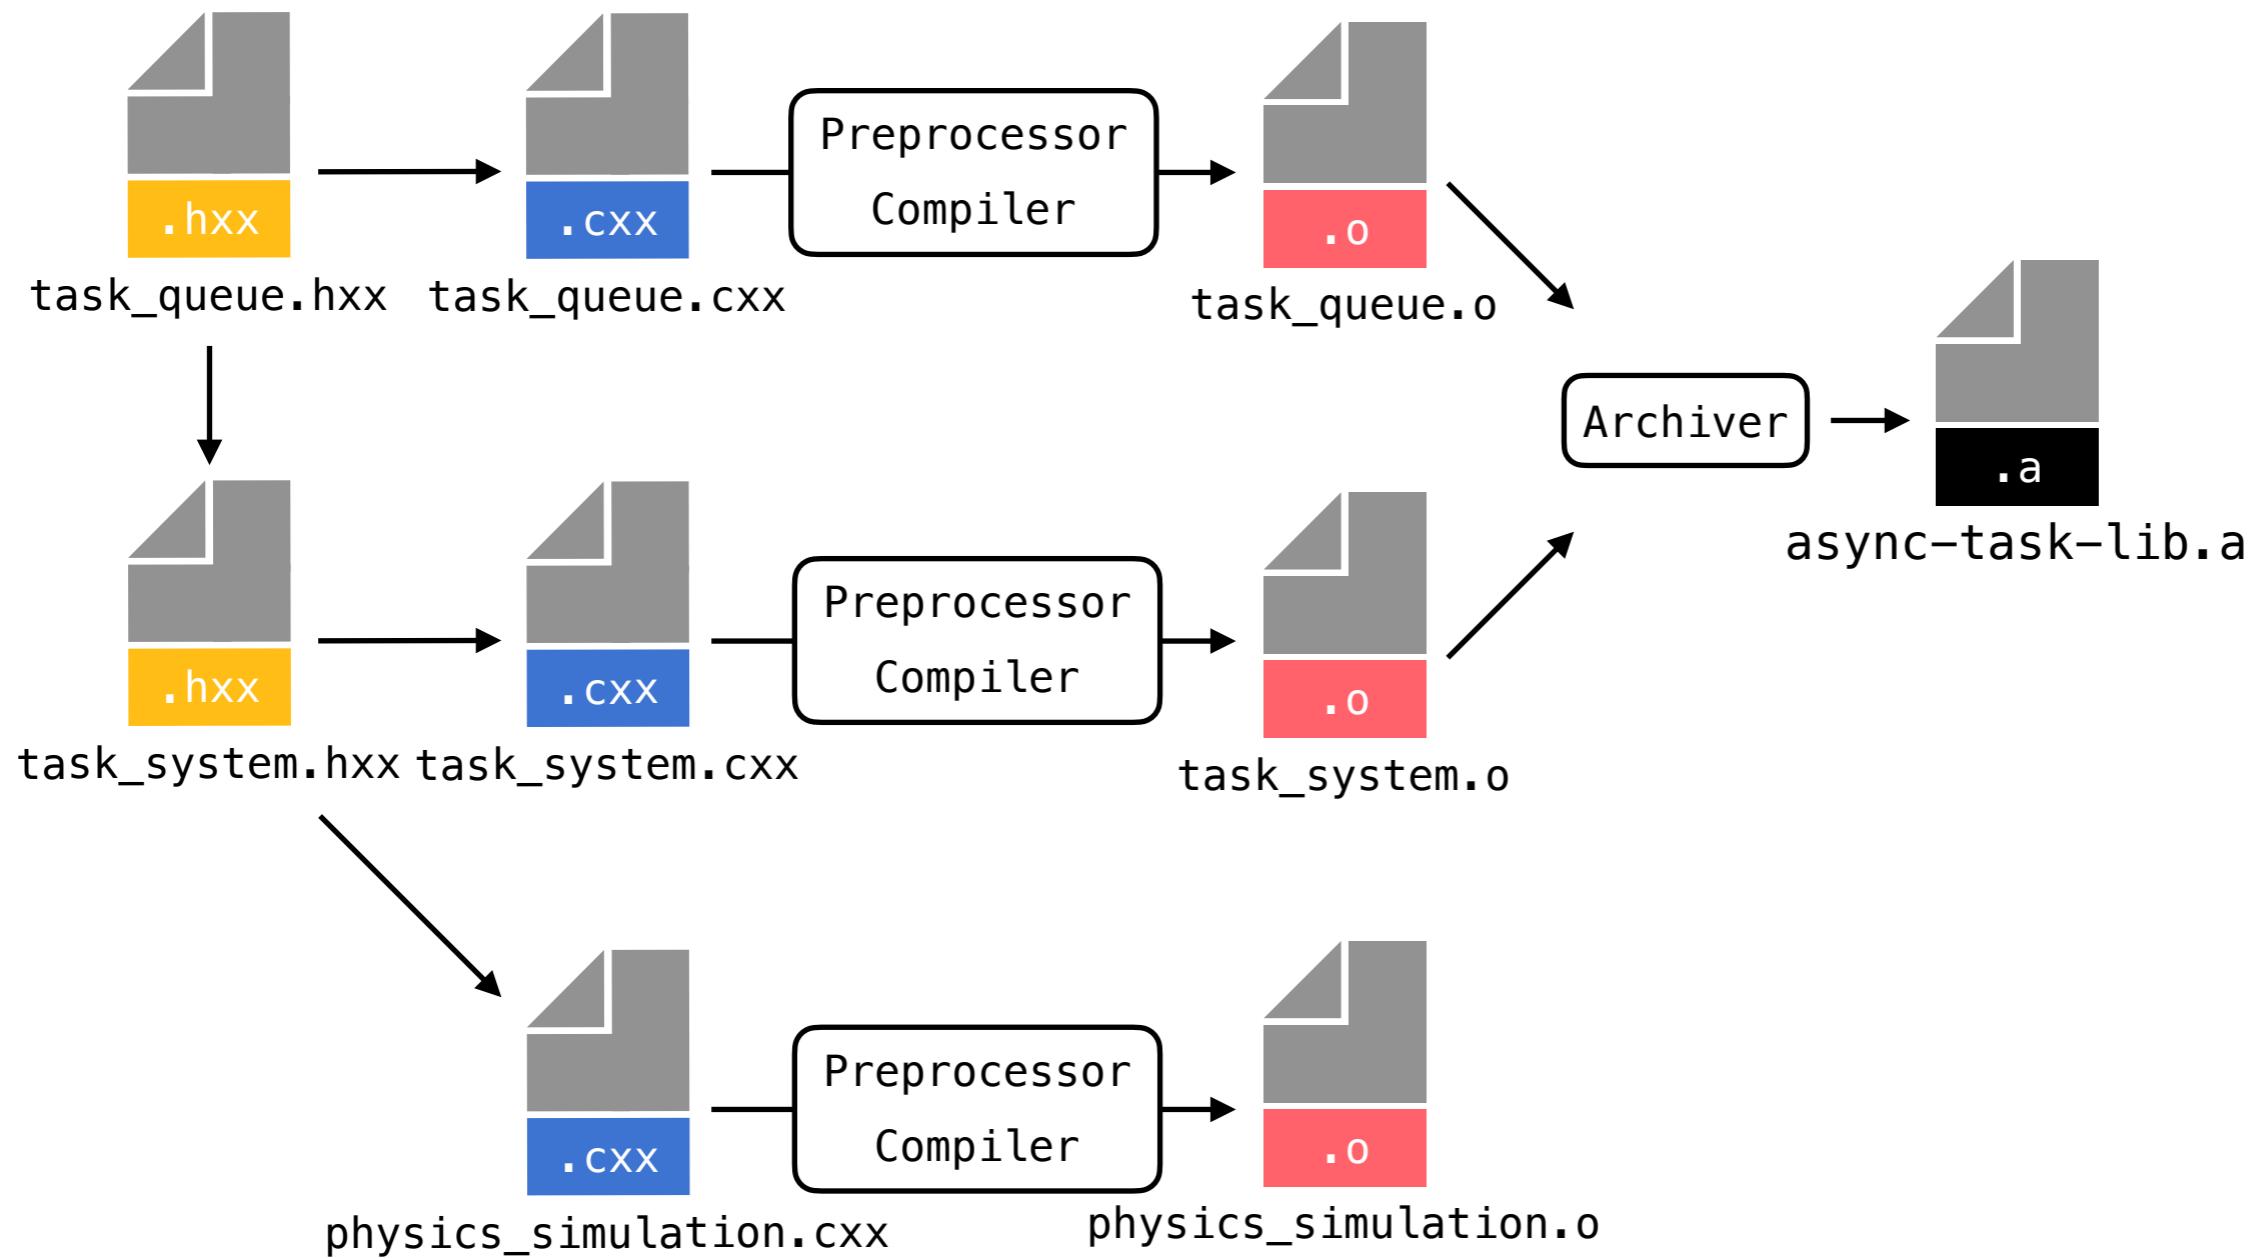
Textual Inclusion Model

Build Workflow



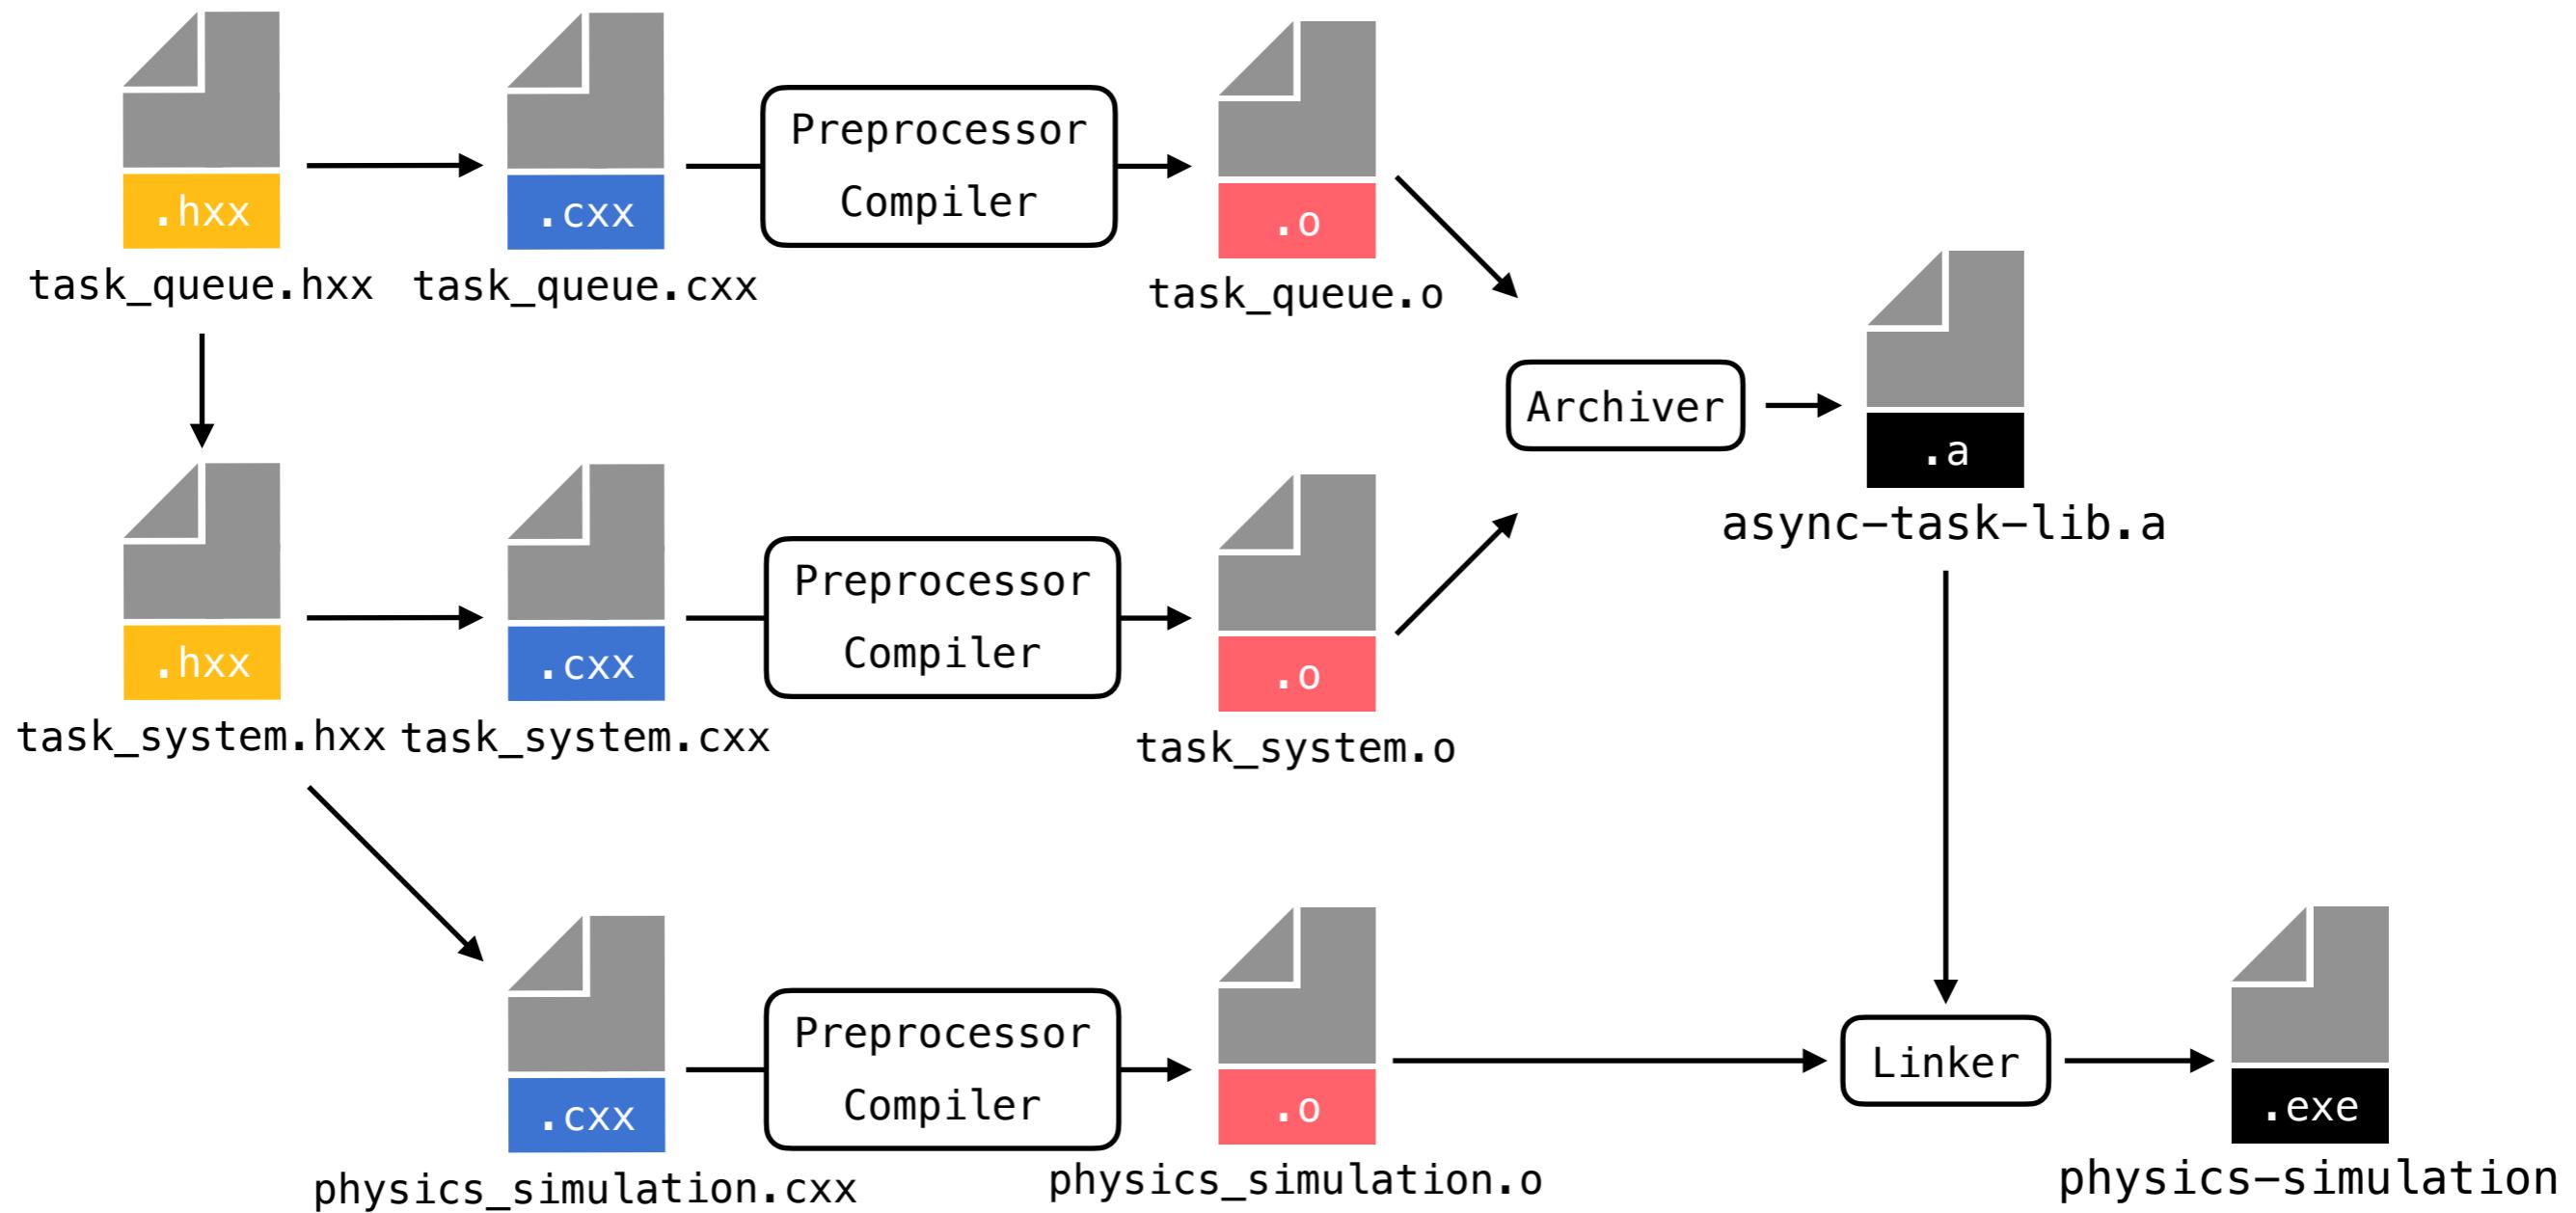
Textual Inclusion Model

Build Workflow



Textual Inclusion Model

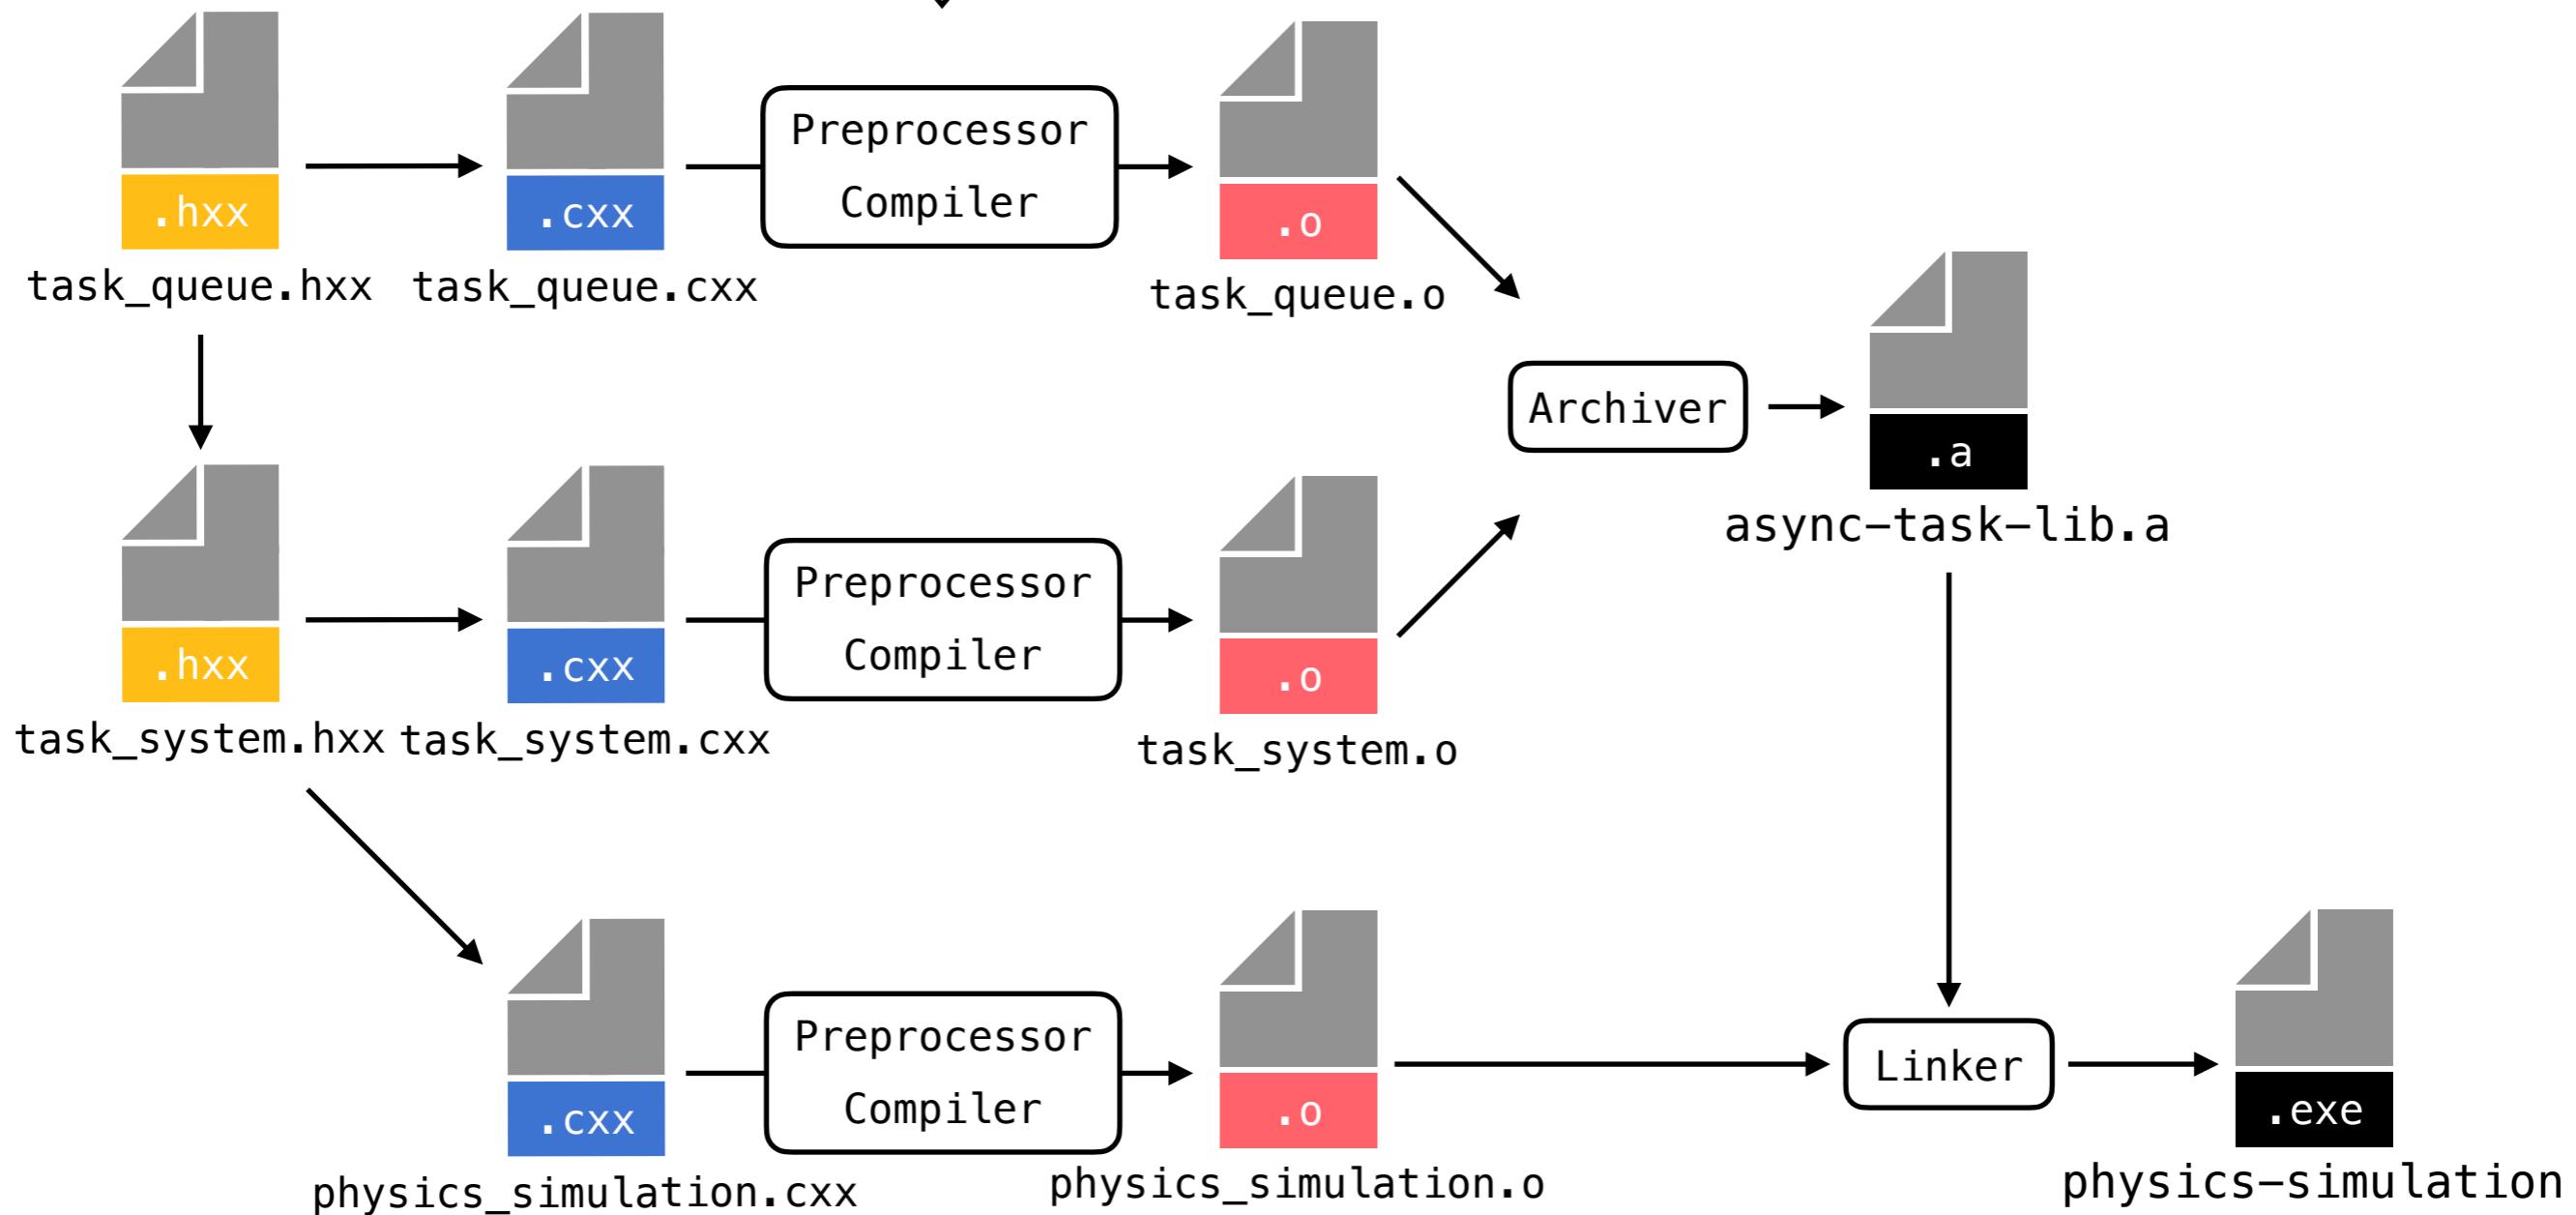
Build Workflow



Textual Inclusion Model

Build Workflow

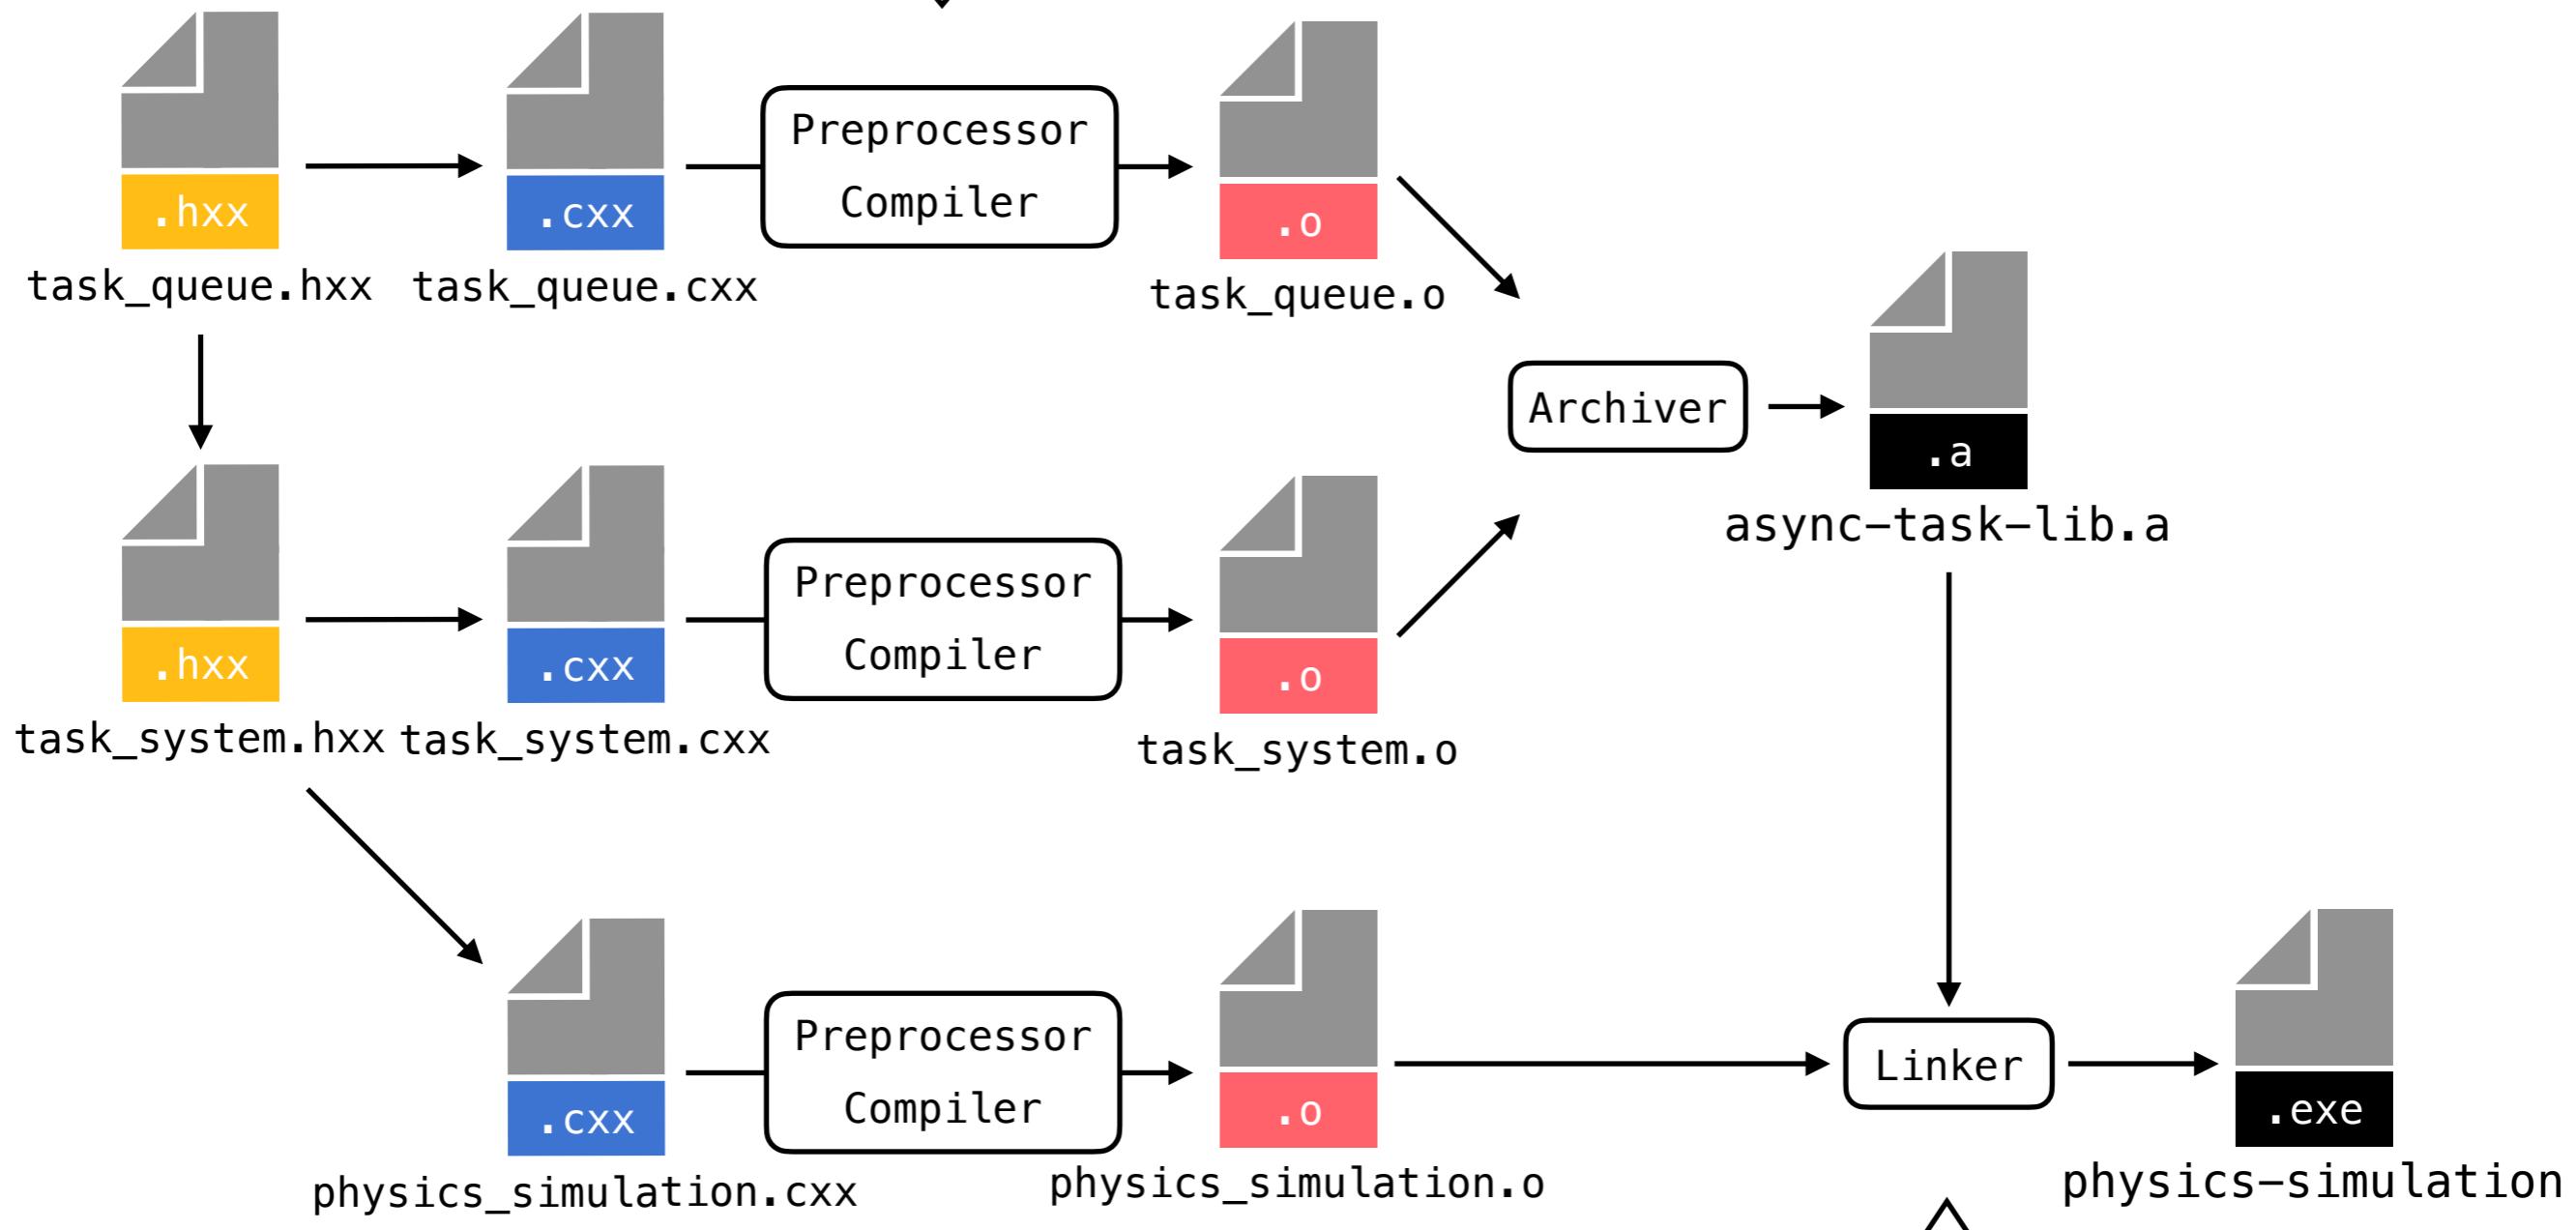
The textual inclusion of header files is very expensive, because their content is compiled multiple times, making the textual inclusion model incredibly inefficient and wasteful.



Textual Inclusion Model

Build Workflow

The textual inclusion of header files is very expensive, because their content is compiled multiple times, making the textual inclusion model incredibly inefficient and wasteful.



Linker discards duplicated symbols.

Textual Inclusion Model

Mangled Symbol Names

C++ entities	Linker symbols
<code>atl::task_queue::task_queue ()</code>	<code>__ZN3atl10task_queueC1Ev</code>
<code>atl::task_queue::~task_queue ()</code>	<code>__ZN3atl10task_queueD1Ev</code>
<code>atl::task_queue::pop ()</code>	<code>__ZN3atl10task_queue3popEv</code>
<code>atl::task_queue::finish ()</code>	<code>__ZN3atl10task_queue6finishEv</code>
<code>atl::task_system::task_system ()</code>	<code>__ZN3atl11task_systemC1Ev</code>
<code>atl::task_system::~task_system ()</code>	<code>__ZN3atl11task_systemD1Ev</code>
<code>phy_sim::length (const vec3&);</code>	<code>__ZN7phy_sim6lengthERKNS_4vec3E</code>
<code>phy_sim::compute_lengths (const std::array<phy_sim::vec3, 128ul>&);</code>	<code>__ZN7phy_sim15compute_lengthsERKNSt3__15arrayINS_4vec3ELm128EEE</code>

Textual Inclusion Model

Mangled Symbol Names

C++ entities	Linker symbols
<code>atl::task_queue::task_queue ()</code>	<code>__ZN3atl10task_queueC1Ev</code>
<code>atl::task_queue::~task_queue ()</code>	<code>__ZN3atl10task_queueD1Ev</code>
<code>atl::task_queue::pop ()</code>	<code>__ZN3atl10task_queue3popEv</code>
<code>atl::task_queue::finish ()</code>	<code>__ZN3atl10task_queue6finishEv</code>
<code>atl::task_system::task_system ()</code>	<code>__ZN3atl11task_systemC1Ev</code>
<code>atl::task_system::~task_system ()</code>	<code>__ZN3atl11task_systemD1Ev</code>
<code>phy_sim::length (const vec3&);</code>	<code>__ZN7phy_sim6lengthERKNS_4vec3E</code>
<code>phy_sim::compute_lengths (const std::array<phy_sim::vec3, 128ul>&);</code>	<code>__ZN7phy_sim15compute_lengthsERKNSt3__15arrayINS_4vec3ELm128EEE</code>

Linker symbol list is global for the whole C++ program, which implies that each C++ entity can have only one definition.

Textual Inclusion Model

One Definition Rule

```
class task_queue
{
private:
    struct task
    {
        std::function<void ()>           function;
#ifndef NDEBUG
        std::chrono::steady_clock::time_point arrival;
        std::chrono::steady_clock::time_point departure;
#endif
    };
    std::deque<task> tasks;
    ...
};
```

Textual Inclusion Model

One Definition Rule

```
class task_queue
{
private:
    struct task
    {
        std::function<void ()>
#ifndef NDEBUG
        std::chrono::steady_clock::time_point arrival;
        std::chrono::steady_clock::time_point departure;
#endif
    };
    std::deque<task> tasks;
    ...
};
```

Let's assume that the task queue in Debug mode reports how long, on average, a task waits for execution in the task queue, that is the task queue latency.

function;

Textual Inclusion Model

One Definition Rule

```
class task_queue
{
private:
    struct task
    {
        std::function<void ()>
#ifndef NDEBUG
        std::chrono::steady_clock::time_point arrival;
        std::chrono::steady_clock::time_point departure;
#endif
    };
    std::deque<task> tasks;
    ...
};
```

Let's assume that the task queue in Debug mode reports how long, on average, a task waits for execution in the task queue, that is the task queue latency.

function;

To implement this functionality, in Debug mode, each task will store its time of arrival and departure.

Textual Inclusion Model

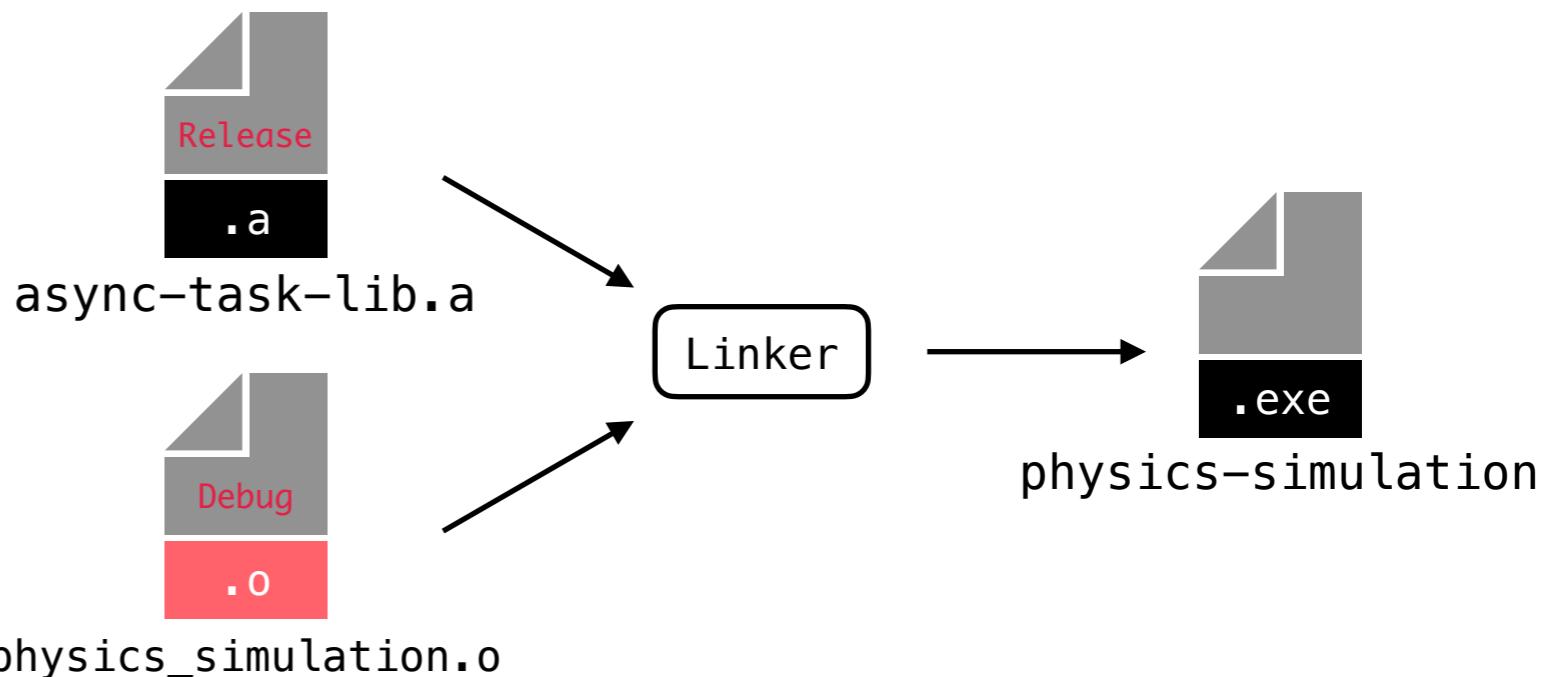
One Definition Rule

```
class task_queue
{
private:
    struct task
    {
        std::function<void ()>
#ifndef NDEBUG
        std::chrono::steady_clock::time_point arrival;
        std::chrono::steady_clock::time_point departure;
#endif
    };
    std::deque<task> tasks;
    ...
};
```

Let's assume that the task queue in Debug mode reports how long, on average, a task waits for execution in the task queue, that is the task queue latency.

function;

To implement this functionality, in Debug mode, each task will store its time of arrival and departure.



Textual Inclusion Model

One Definition Rule

```
class task_queue
{
private:
    struct task
    {
        std::function<void ()>
#ifndef NDEBUG
        std::chrono::steady_clock::time_point arrival;
        std::chrono::steady_clock::time_point departure;
#endif
    };
    std::deque<task> tasks;
    ...
};
```

Let's assume that the task queue in Debug mode reports how long, on average, a task waits for execution in the task queue, that is the task queue latency.

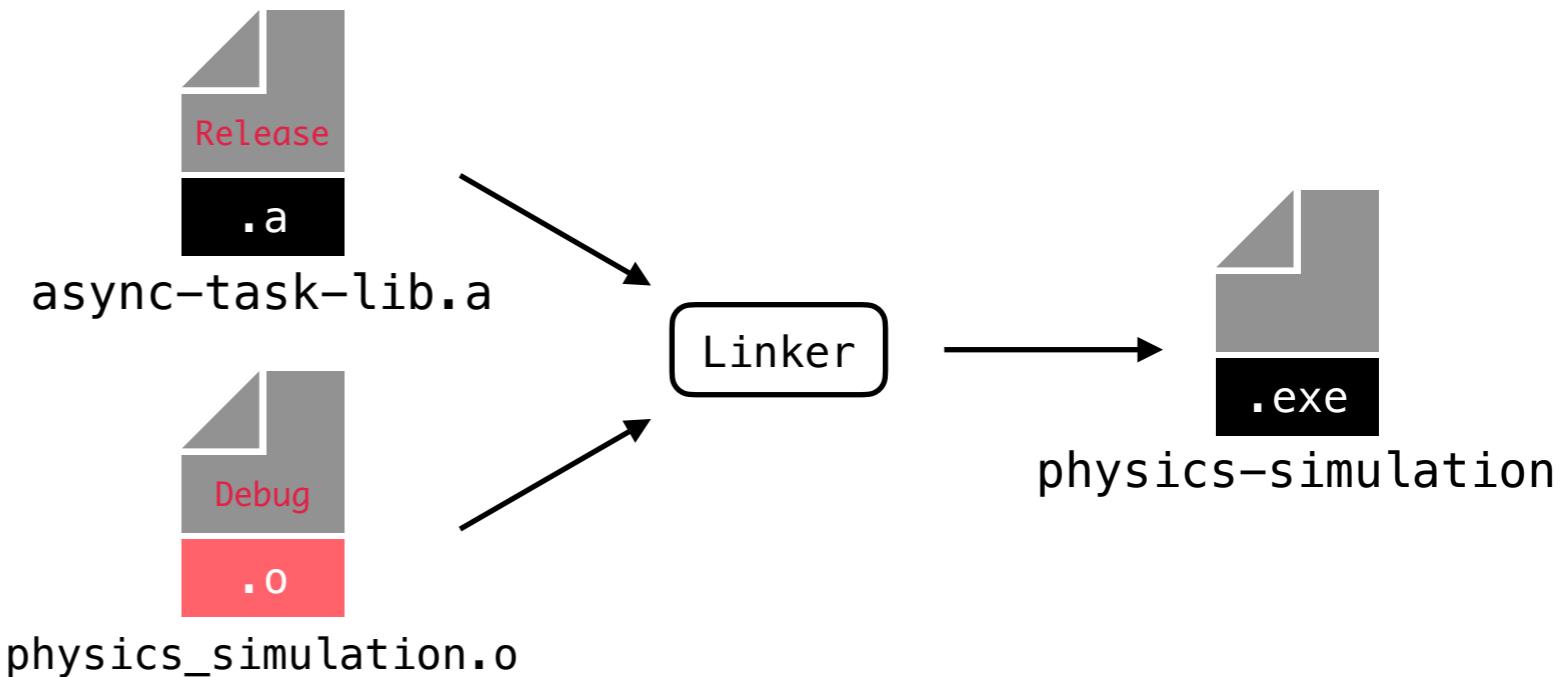
function;

To implement this functionality, in Debug mode, each task will store its time of arrival and departure.

Let's also assume that we are debugging an issue in the physics simulation.

To improve responsiveness of debugging we decided to use Release version of the Async Task Library.

The resulting program is **ill-formed**, because translation units will contain conflicting definitions of the task structure!



Textual Inclusion Model

One Definition Rule

```
class task_queue
{
private:
    struct task
    {
        std::function<void ()>
#ifndef NDEBUG
        std::chrono::steady_clock::time_point arrival;
        std::chrono::steady_clock::time_point departure;
#endif
    };
    std::deque<task> tasks;
    ...
};
```

Let's also assume that we are debugging an issue in the physics simulation.

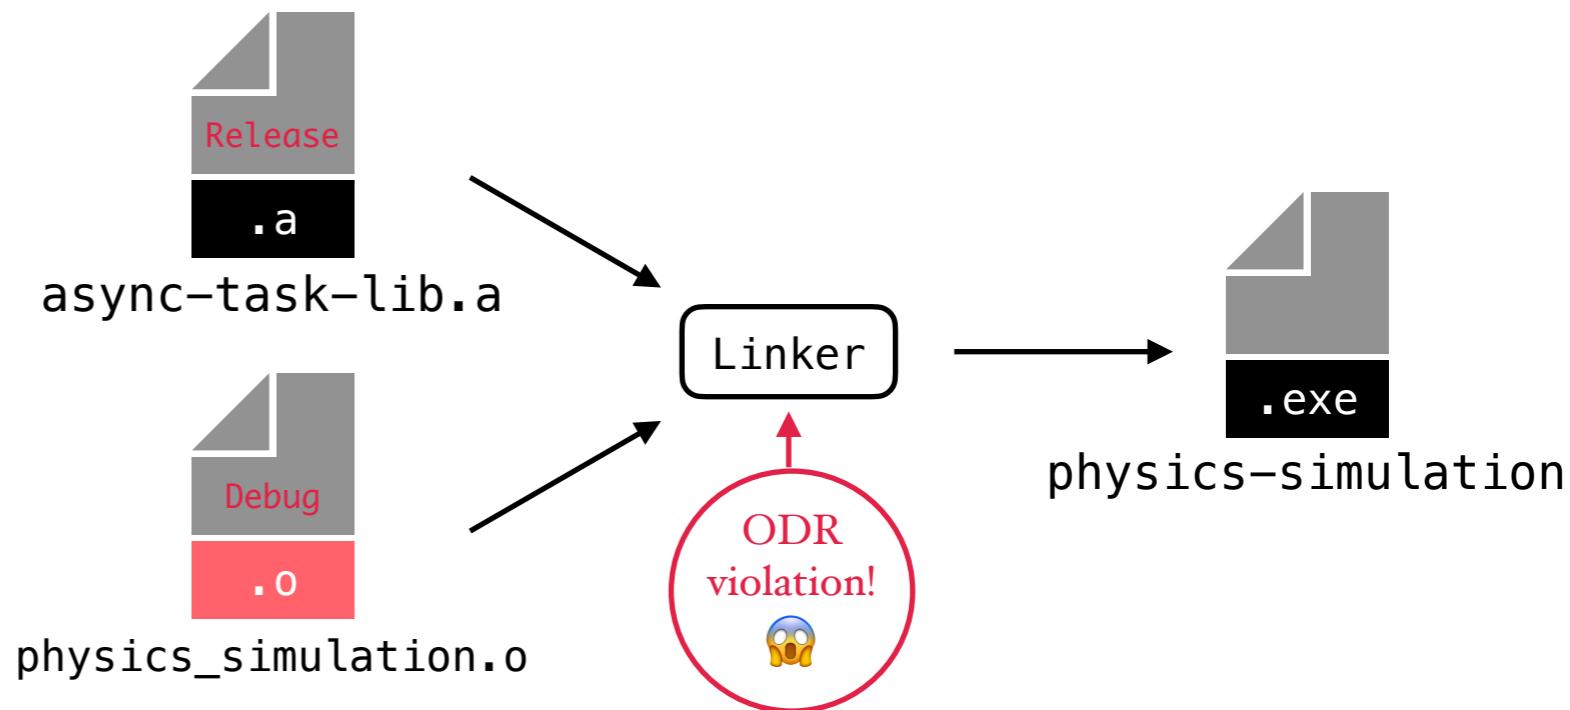
To improve responsiveness of debugging we decided to use Release version of the Async Task Library.

The resulting program is **ill-formed**, because translation units will contain conflicting definitions of the task structure!

Let's assume that the task queue in Debug mode reports how long, on average, a task waits for execution in the task queue, that is the task queue latency.

function;

To implement this functionality, in Debug mode, each task will store its time of arrival and departure.



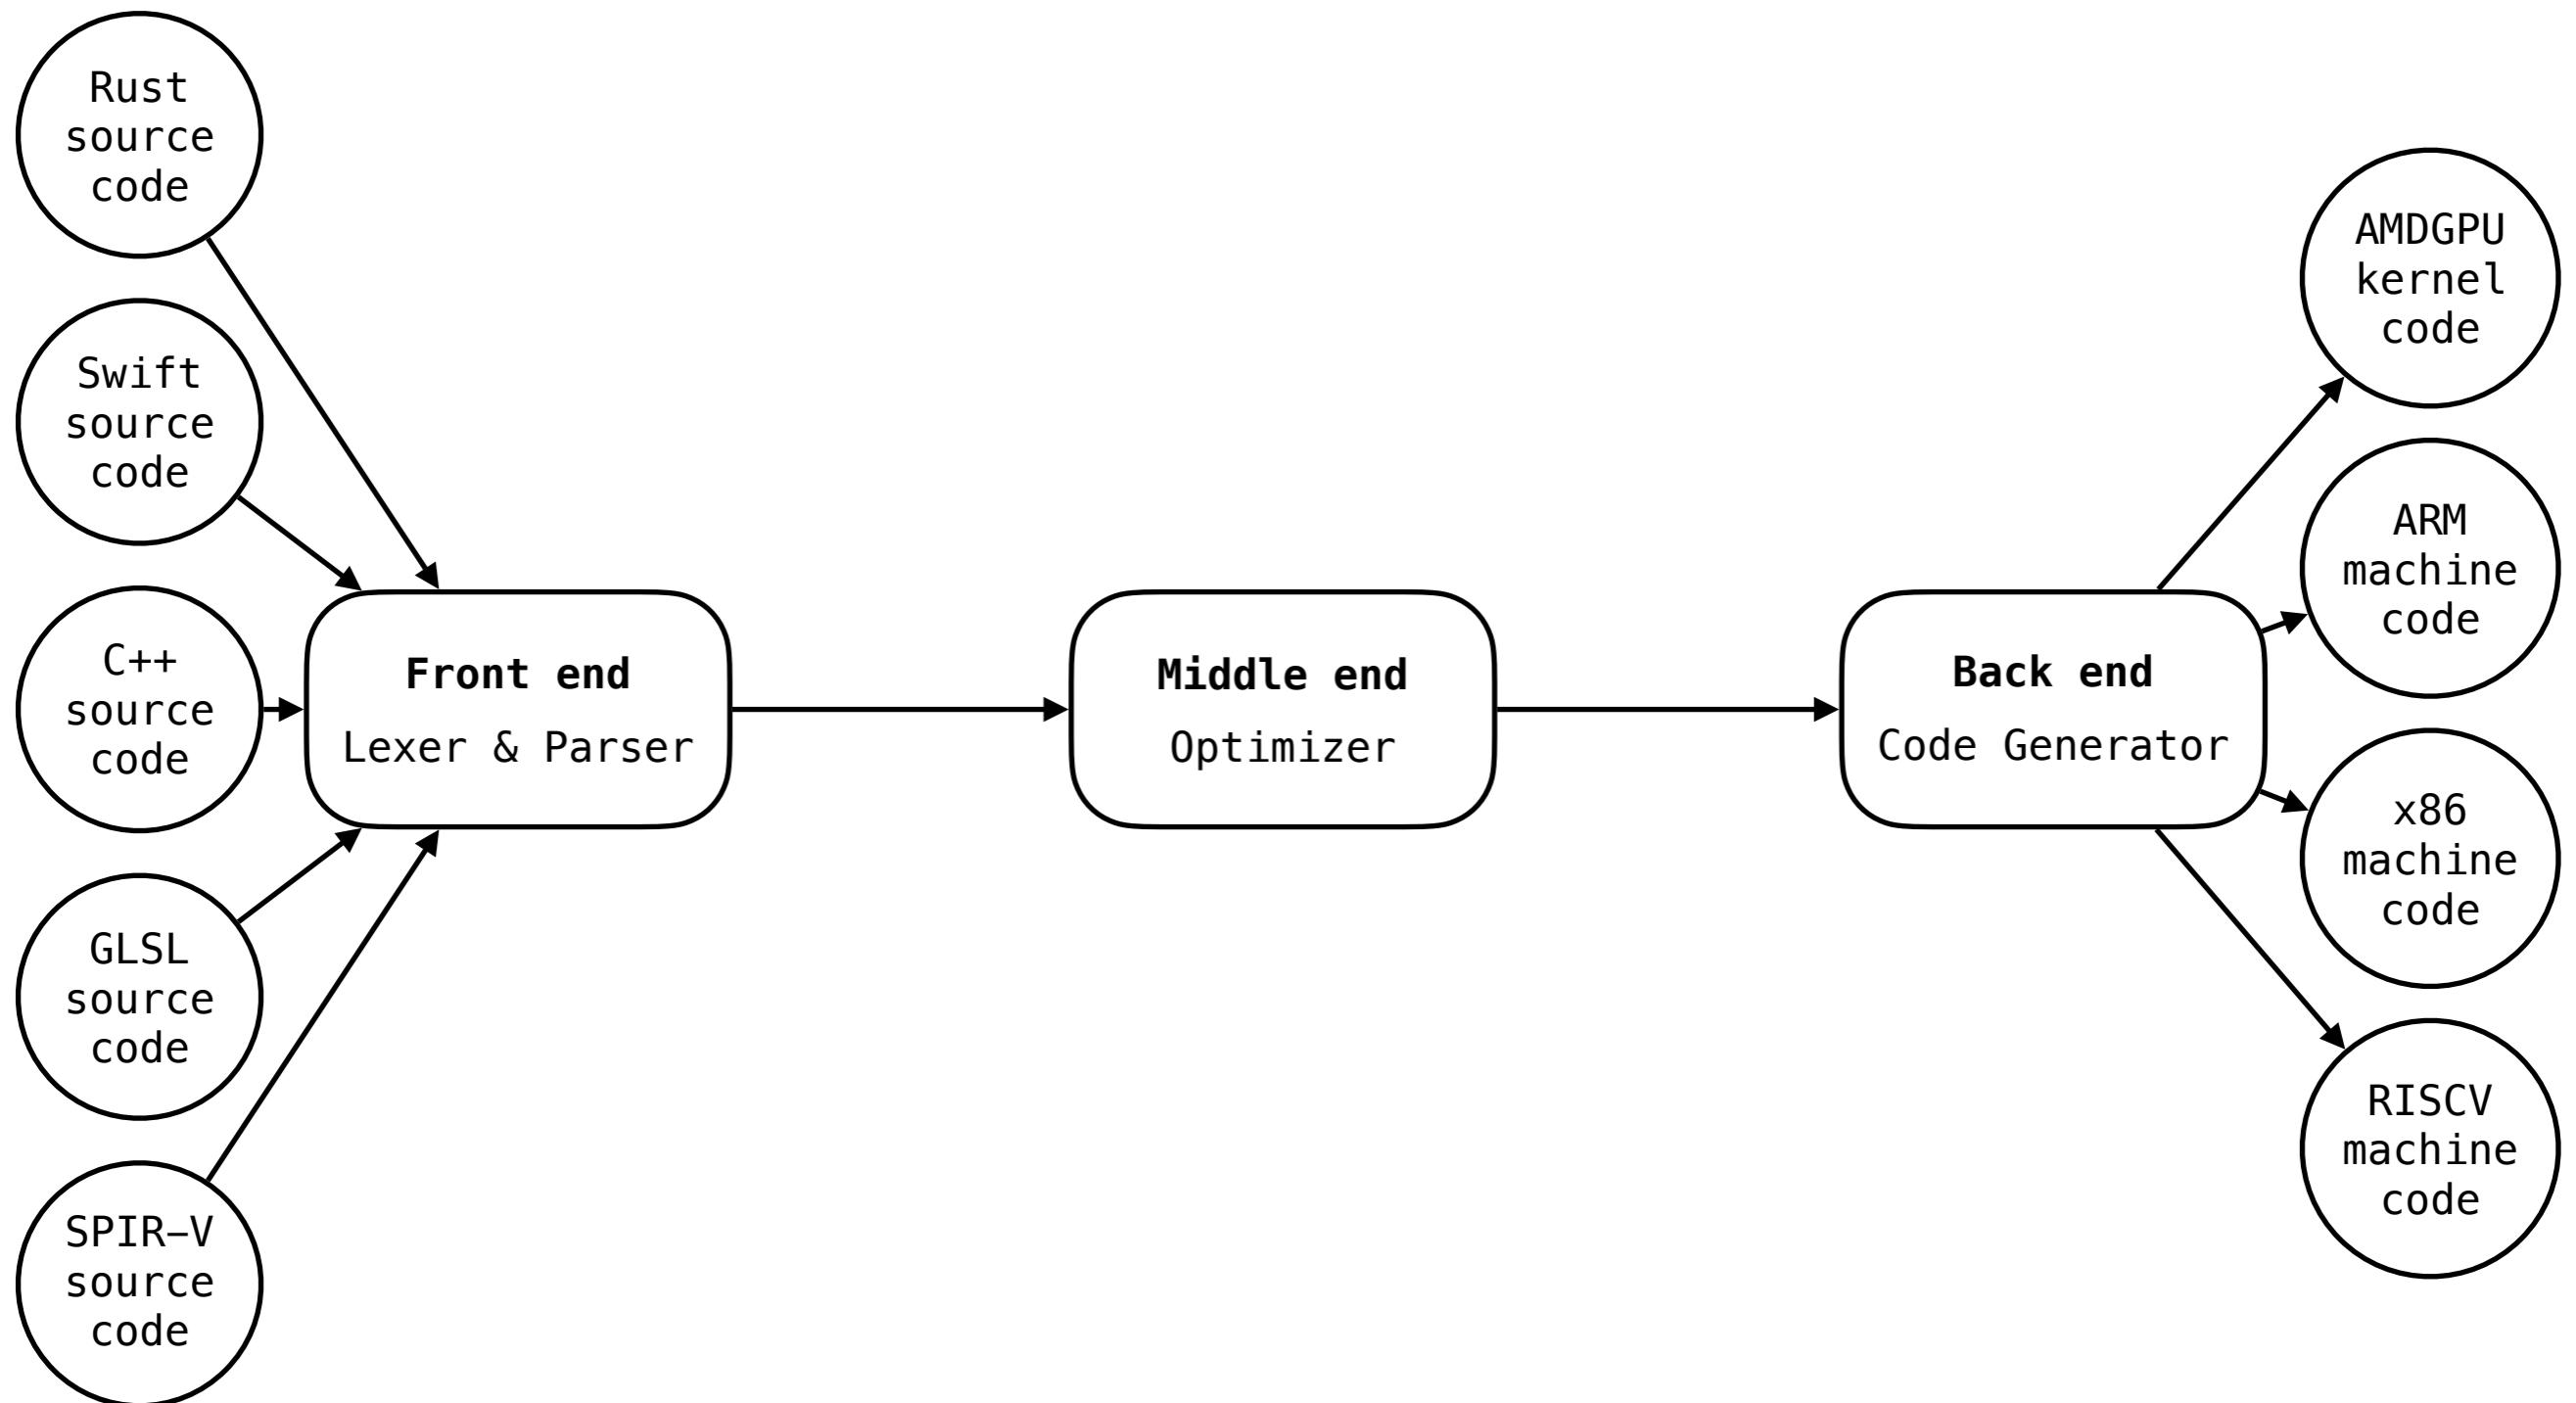
Compiler Architecture

Three stage compiler design



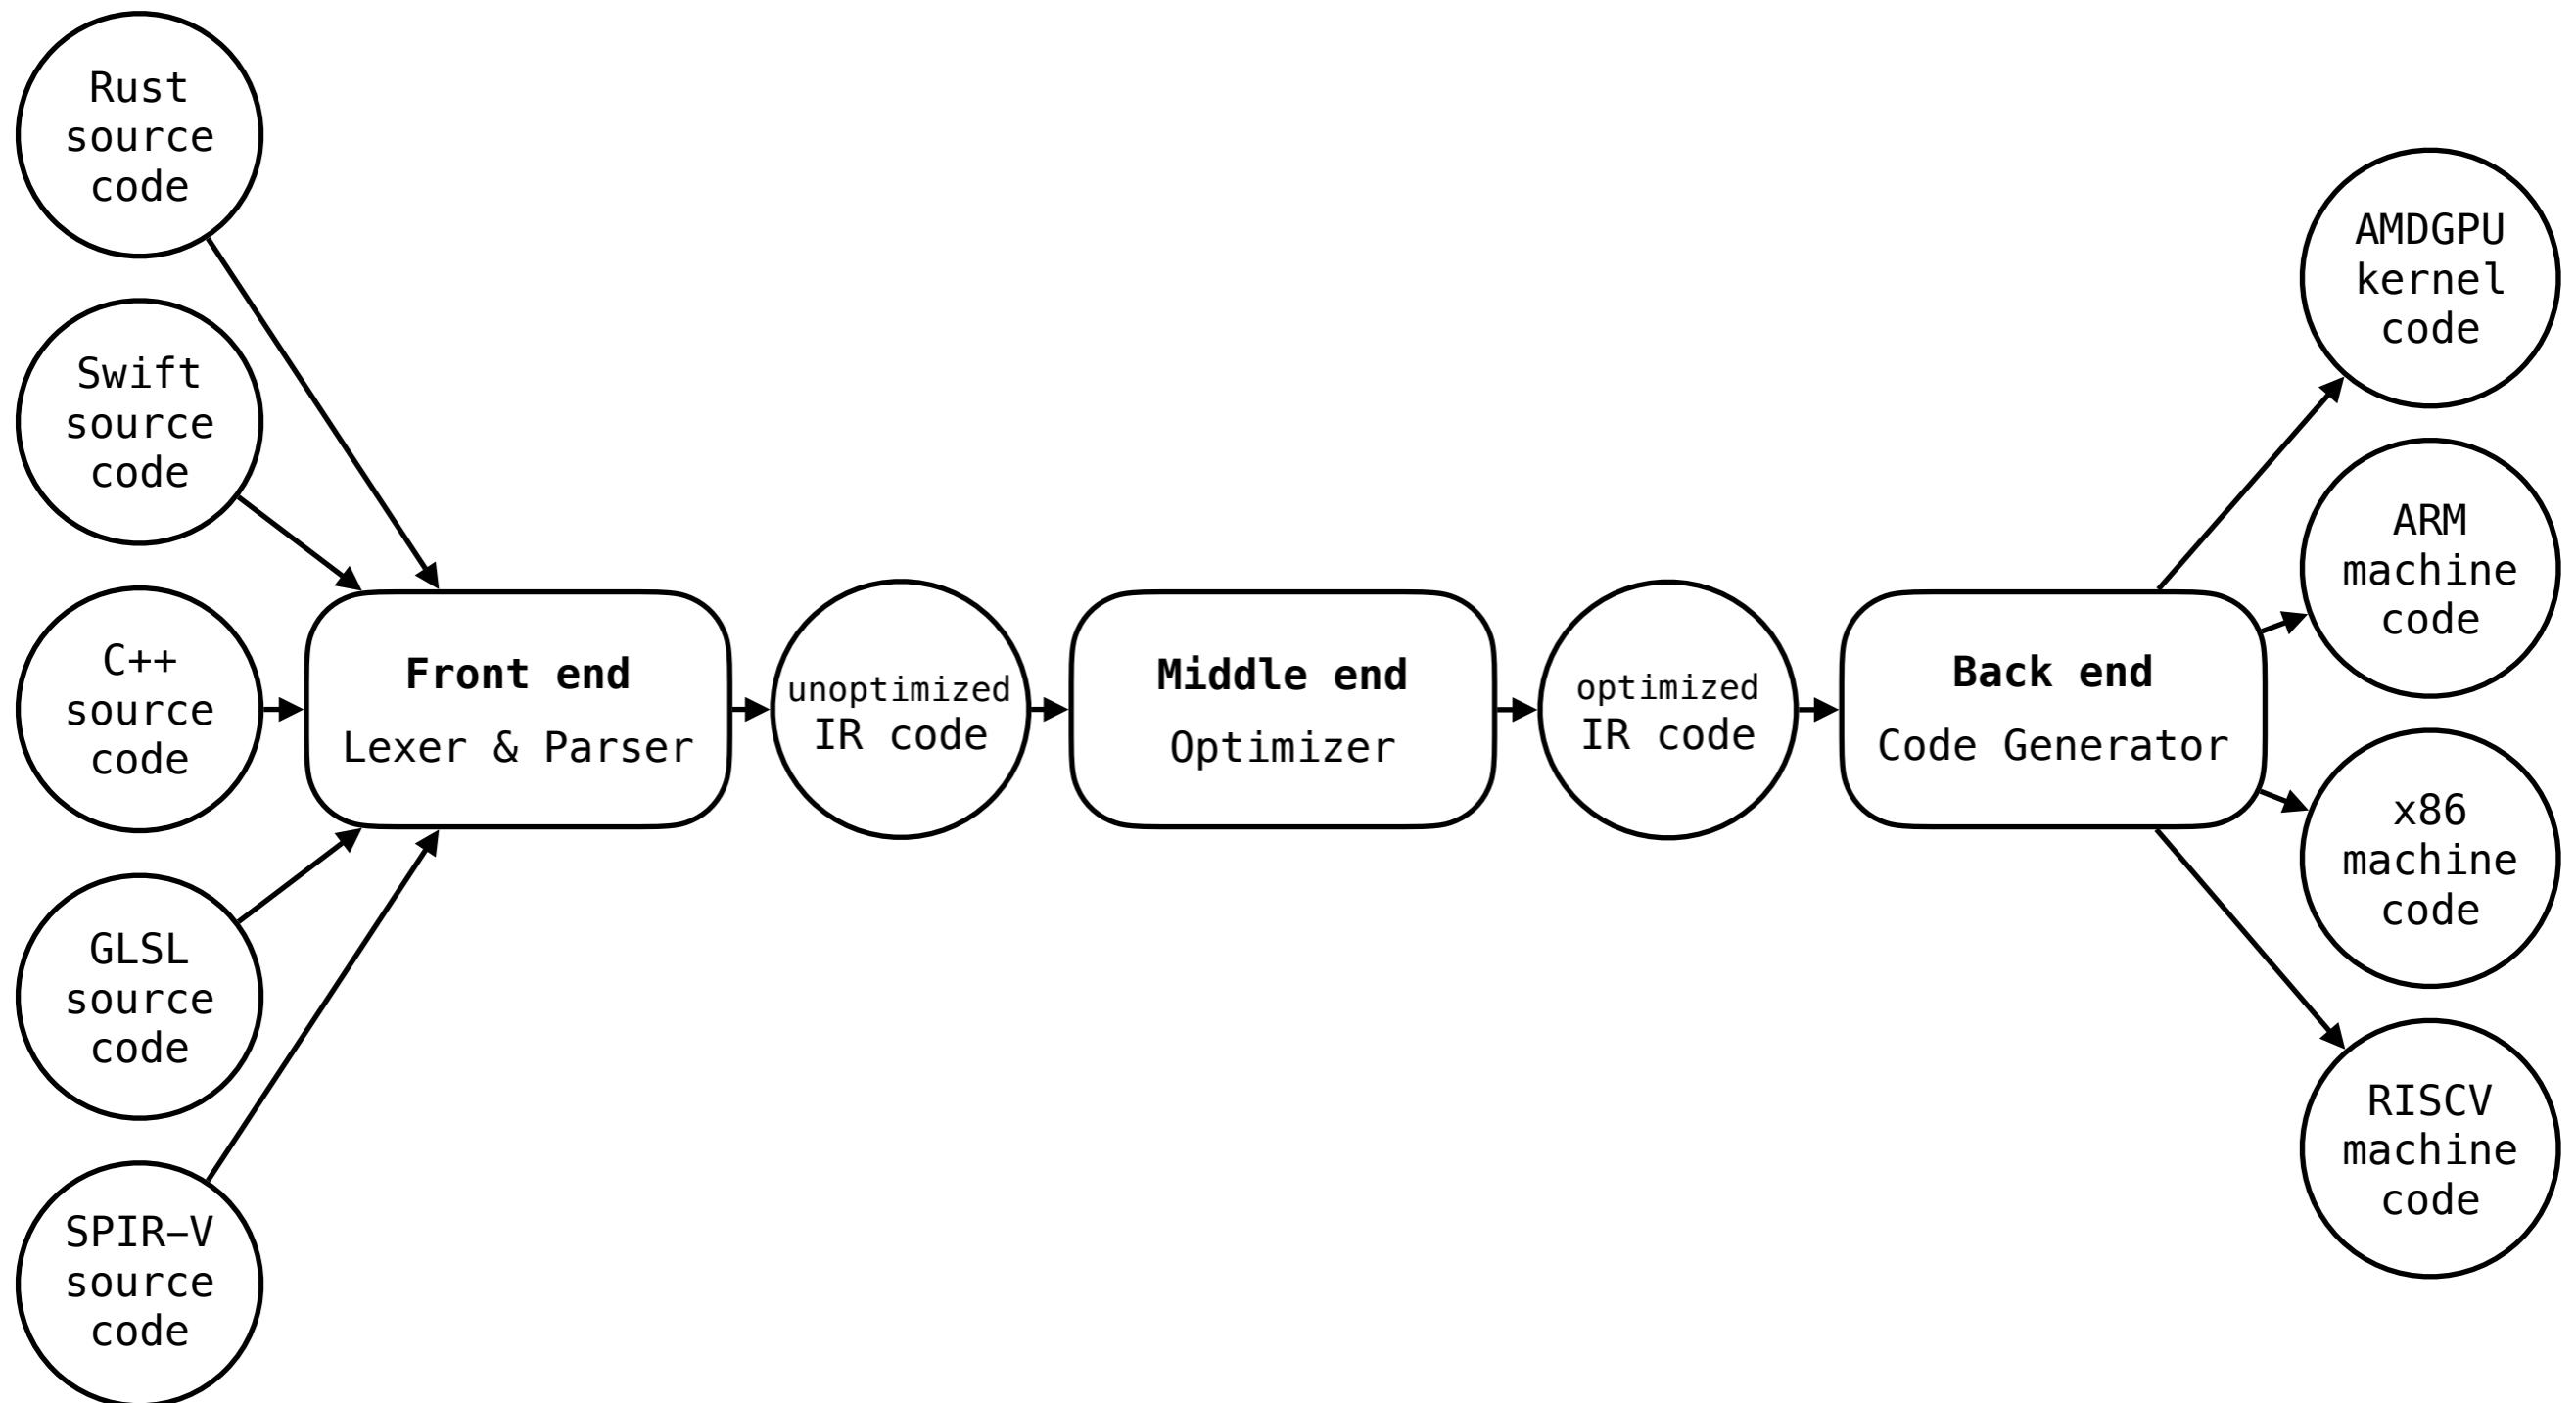
Compiler Architecture

Three stage compiler design



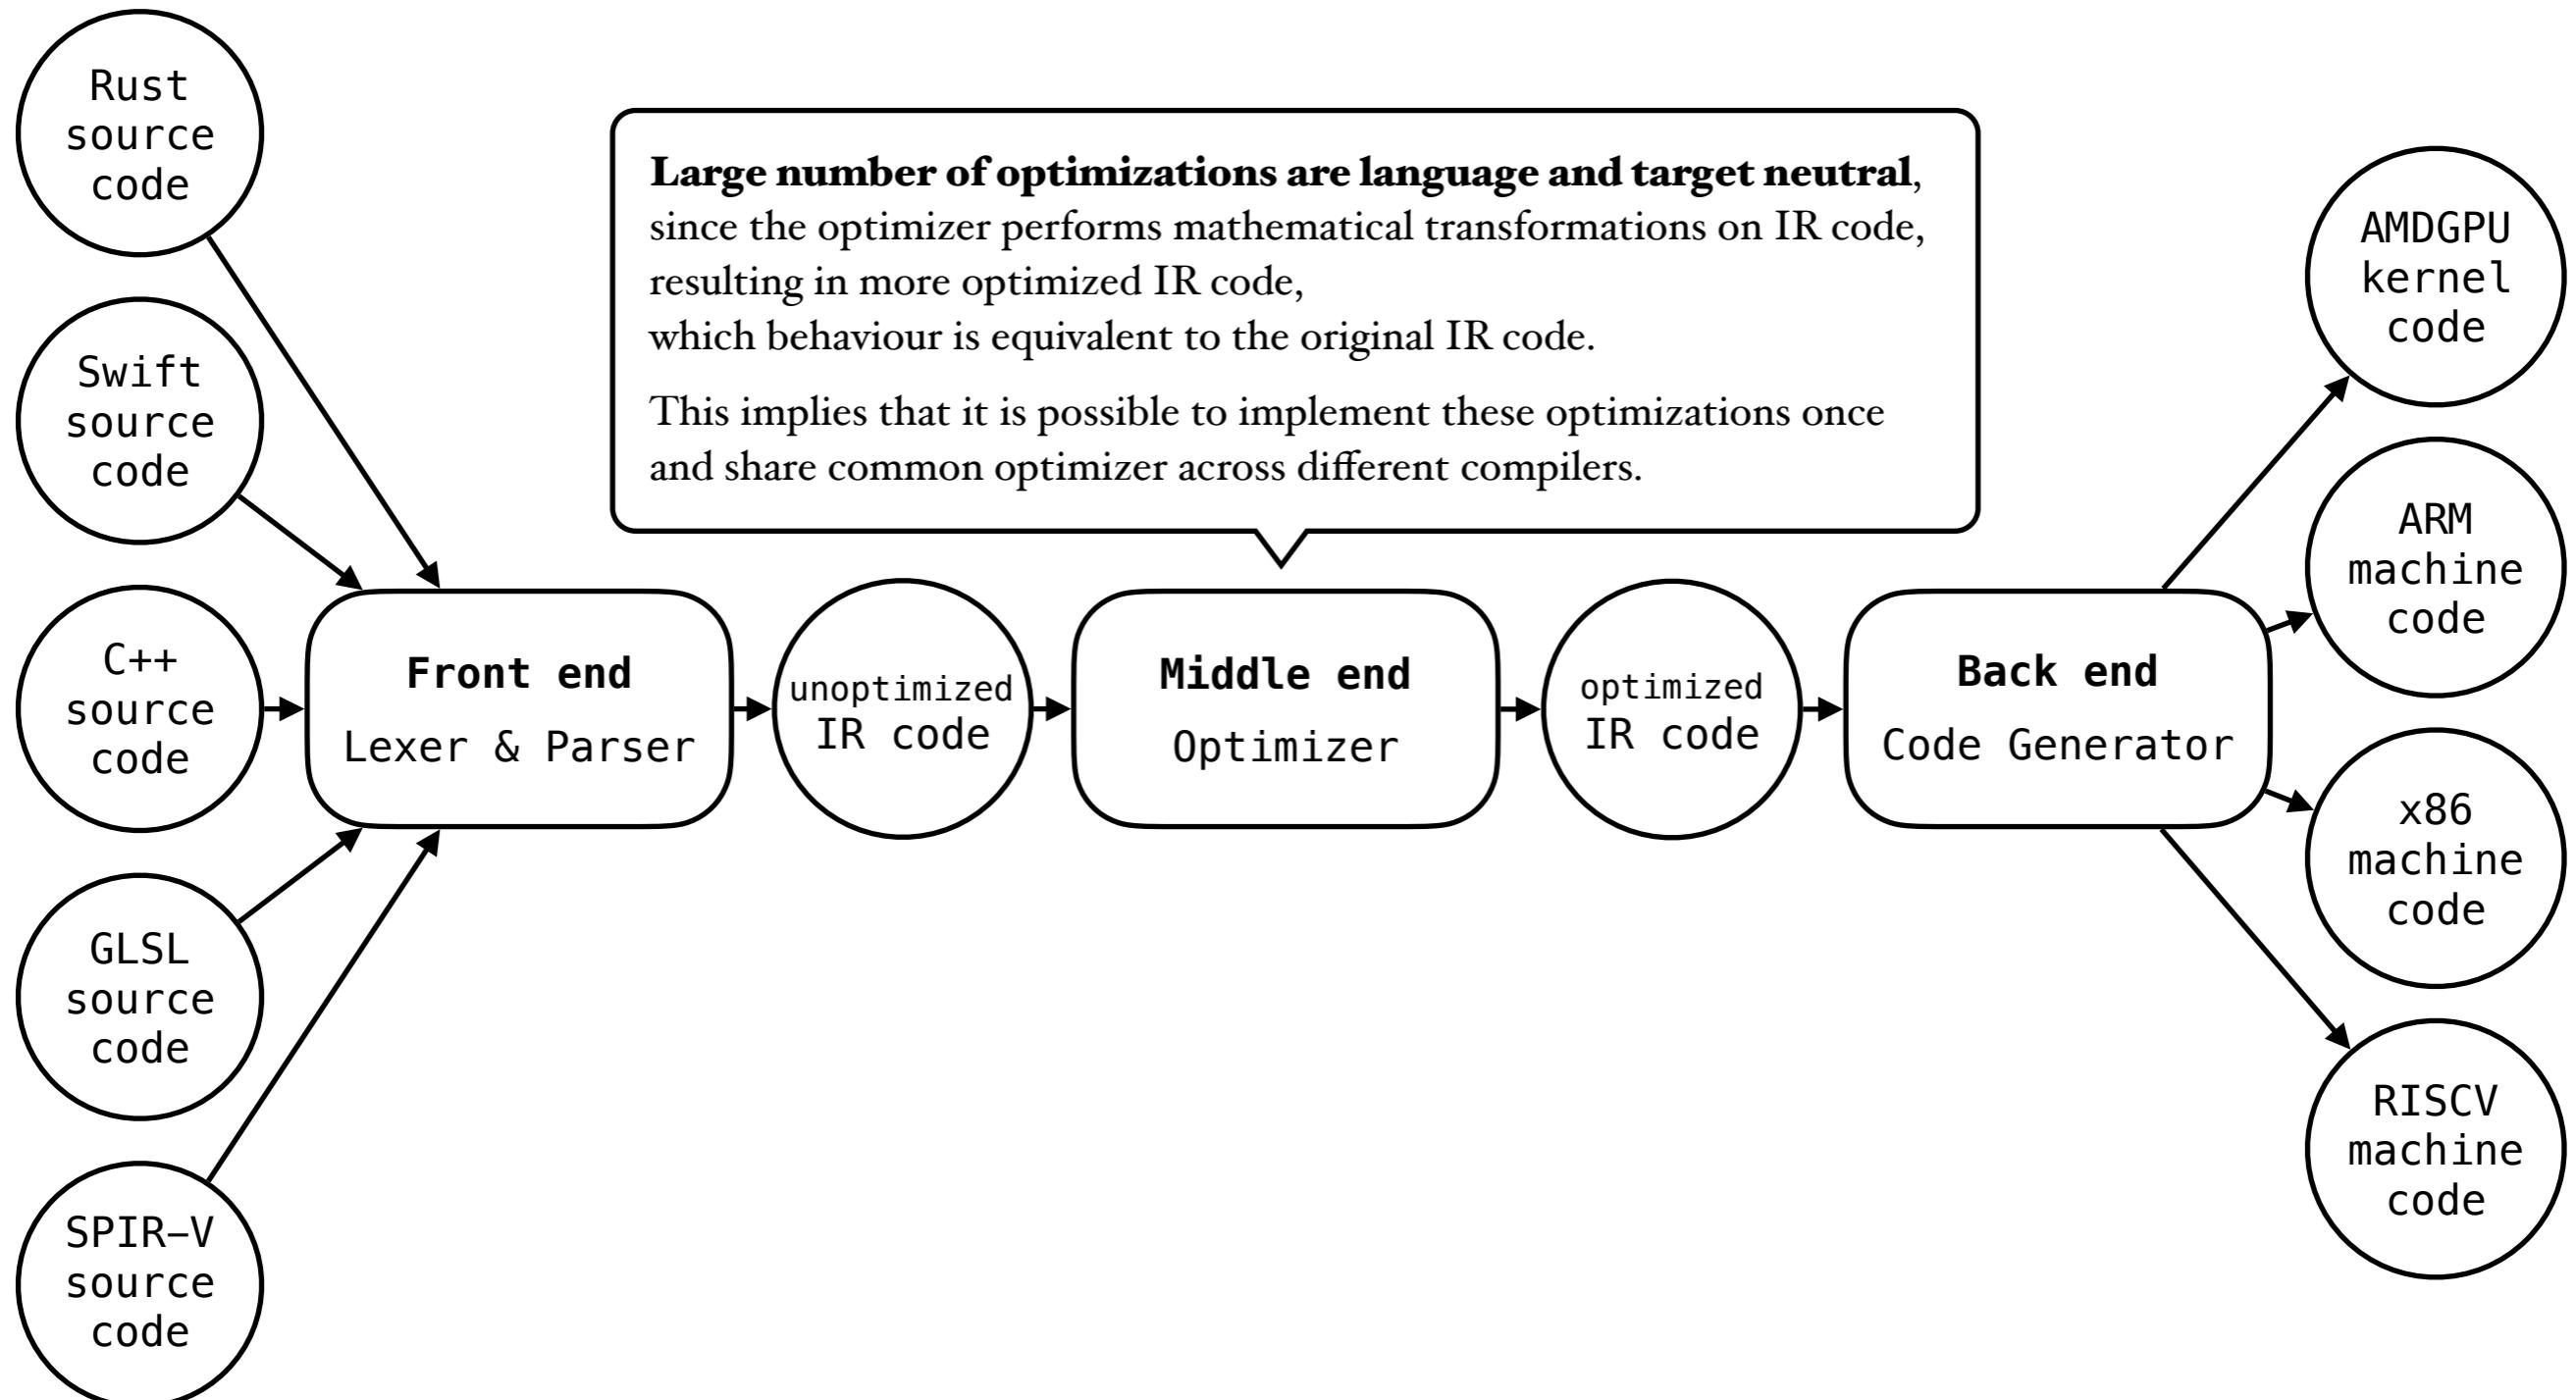
Compiler Architecture

Three stage compiler design



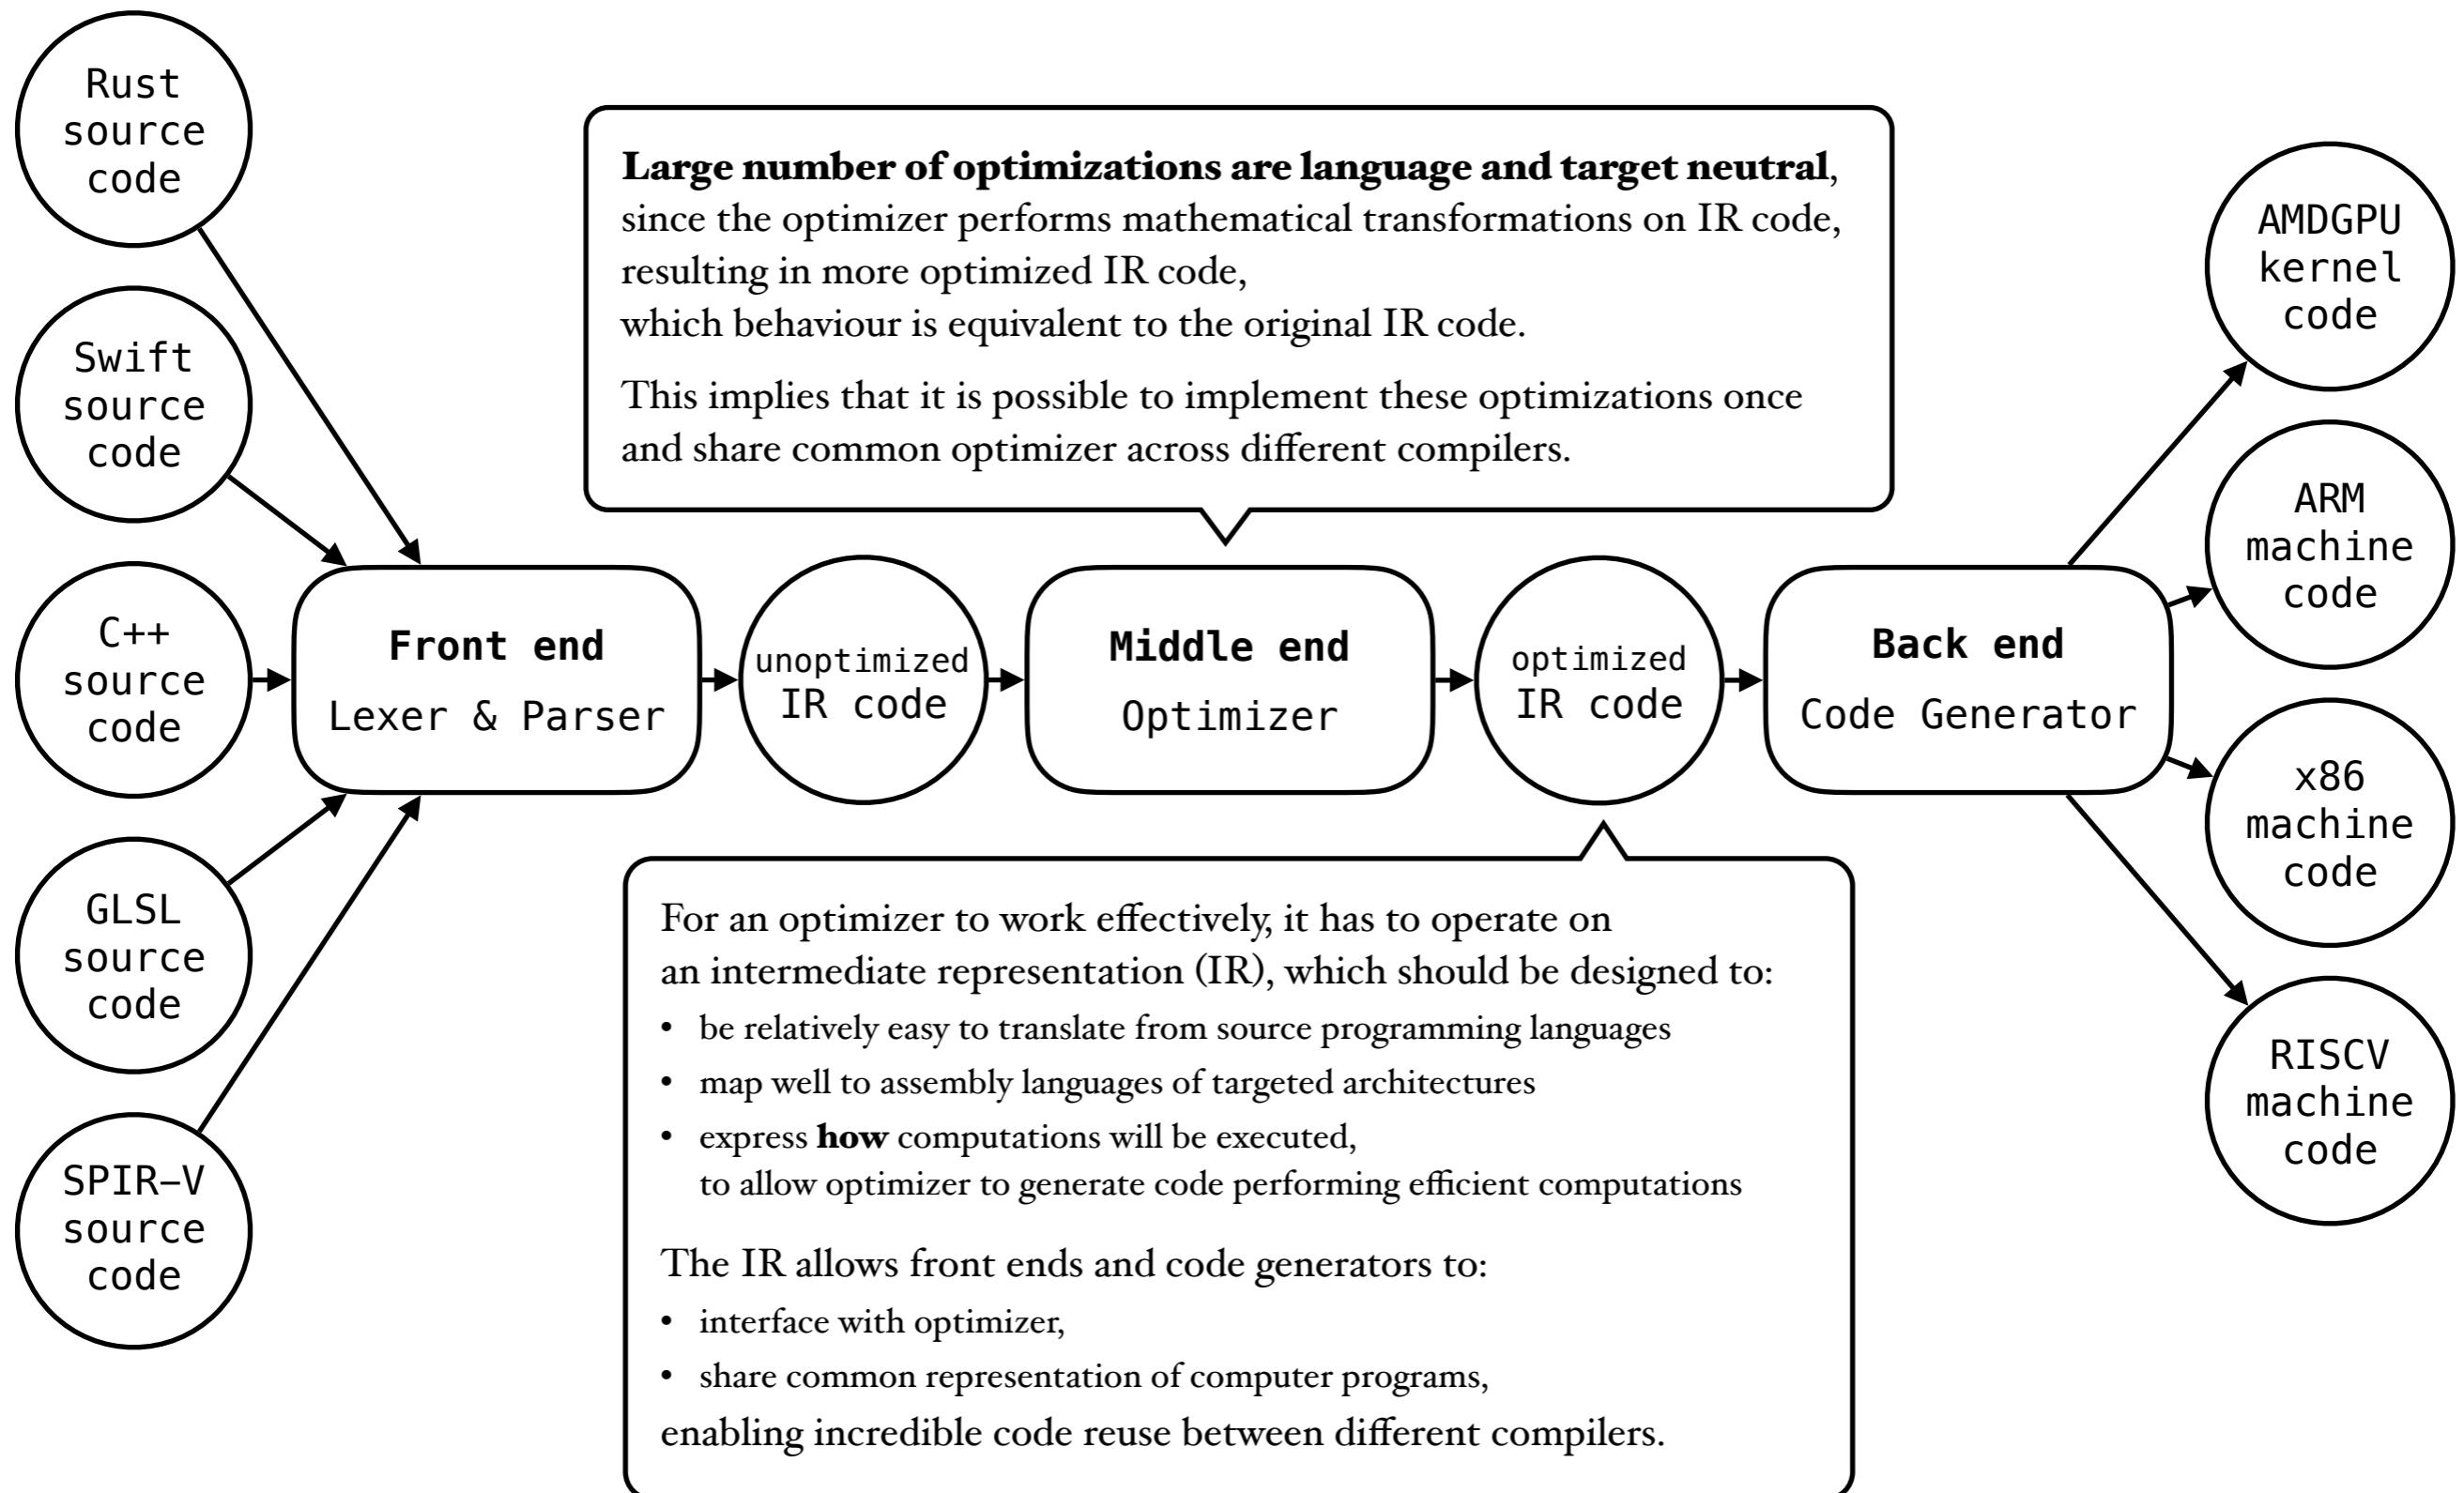
Compiler Architecture

Three stage compiler design



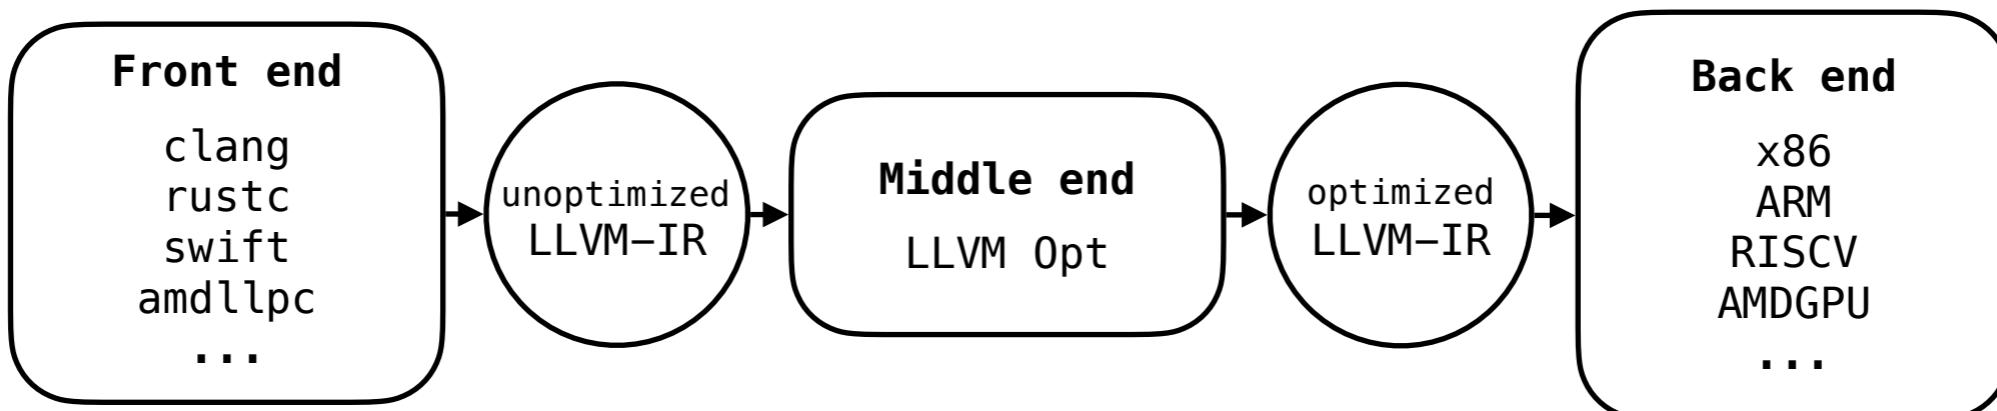
Compiler Architecture

Three stage compiler design



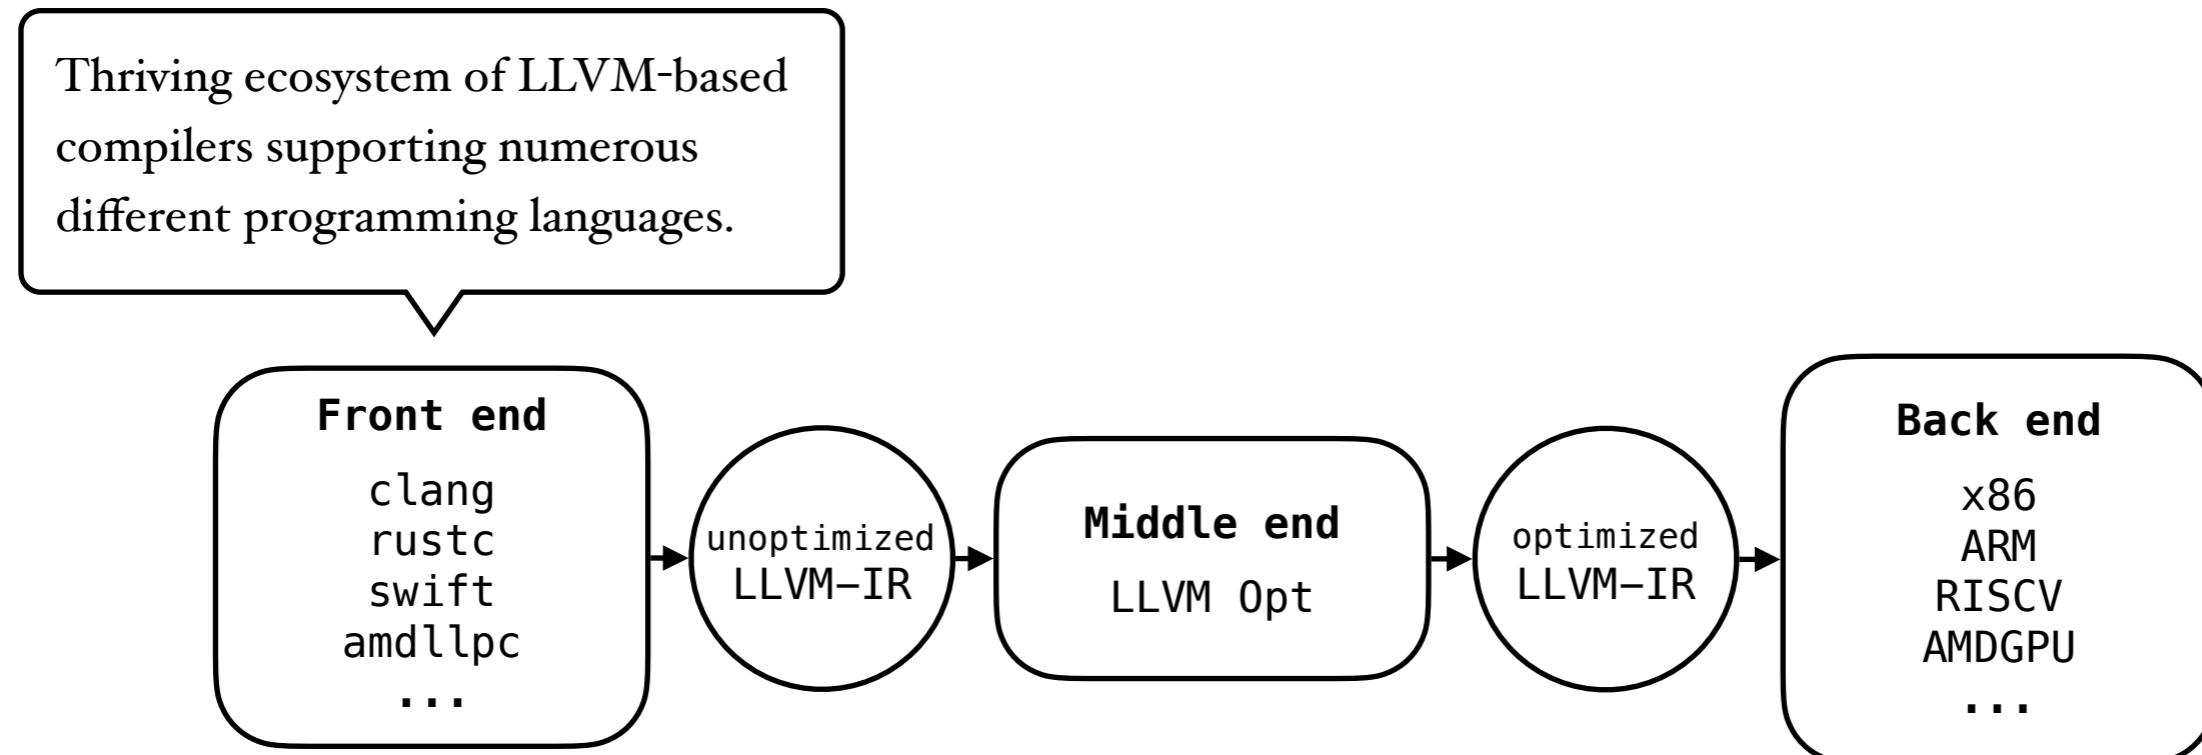
Compiler Architecture

The LLVM Compiler Infrastructure & Clang Compiler



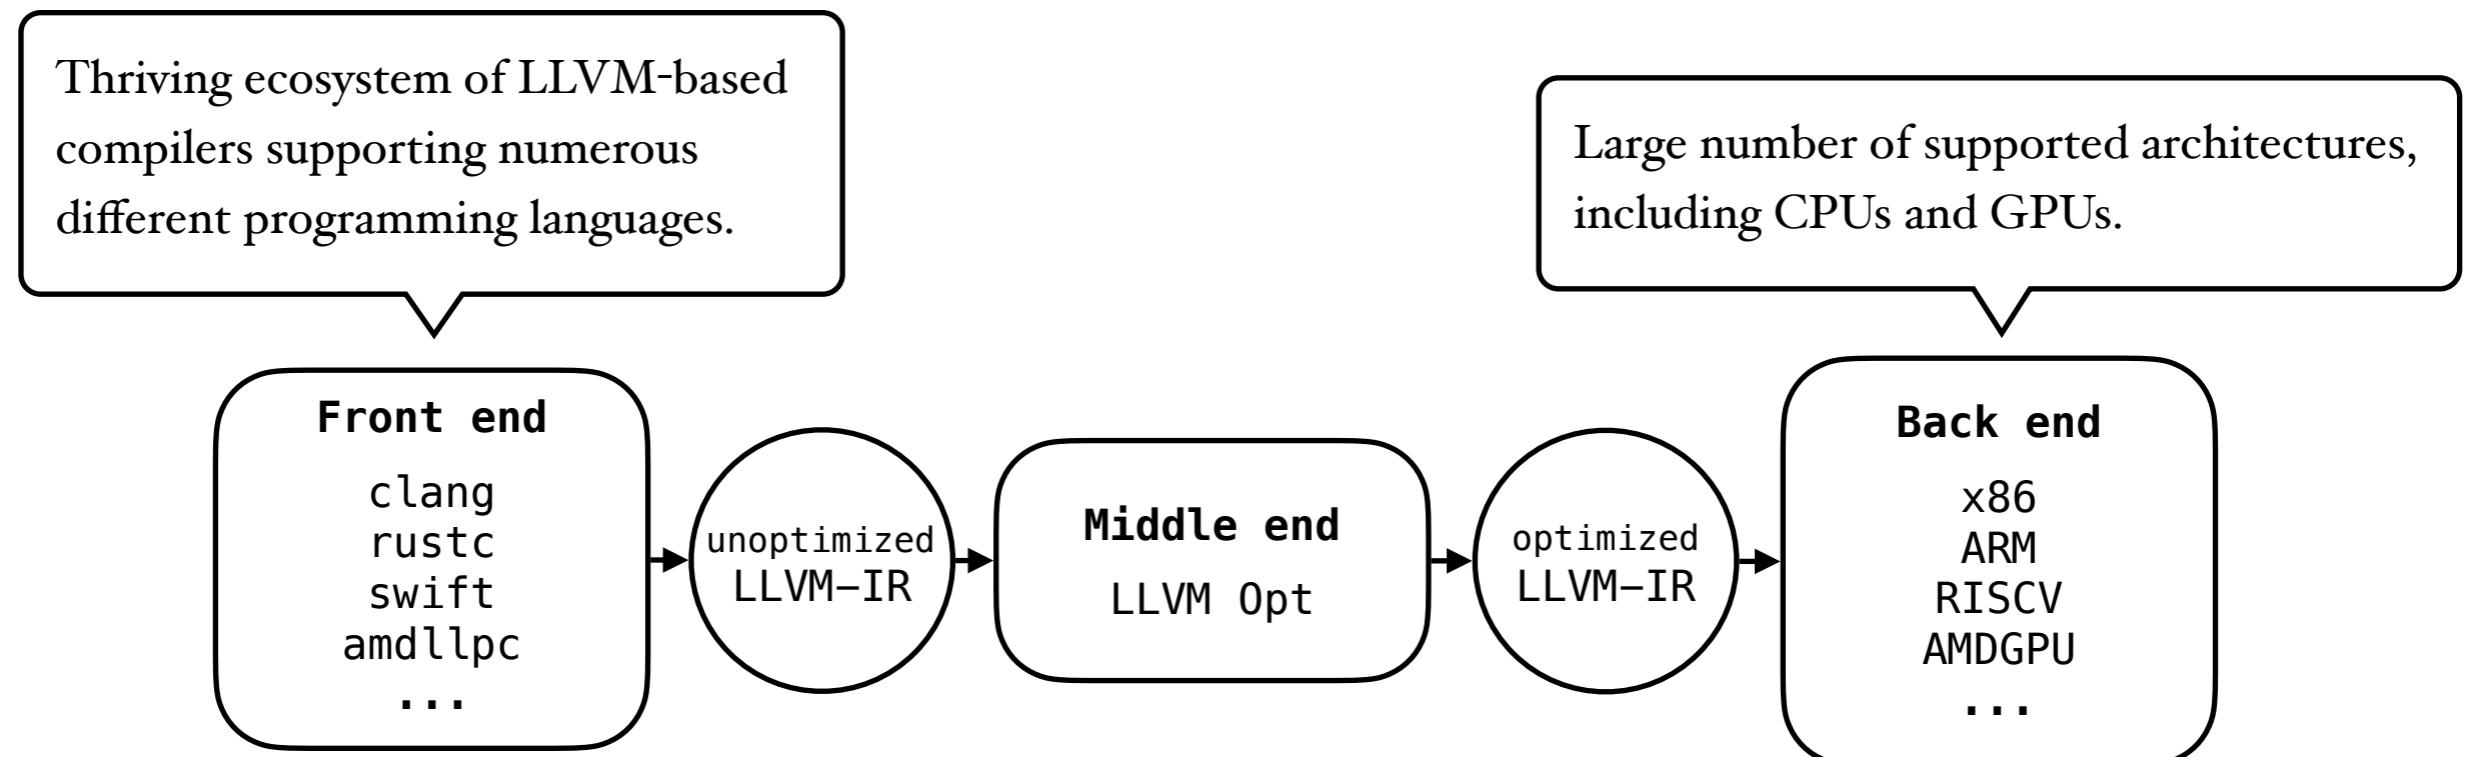
Compiler Architecture

The LLVM Compiler Infrastructure & Clang Compiler



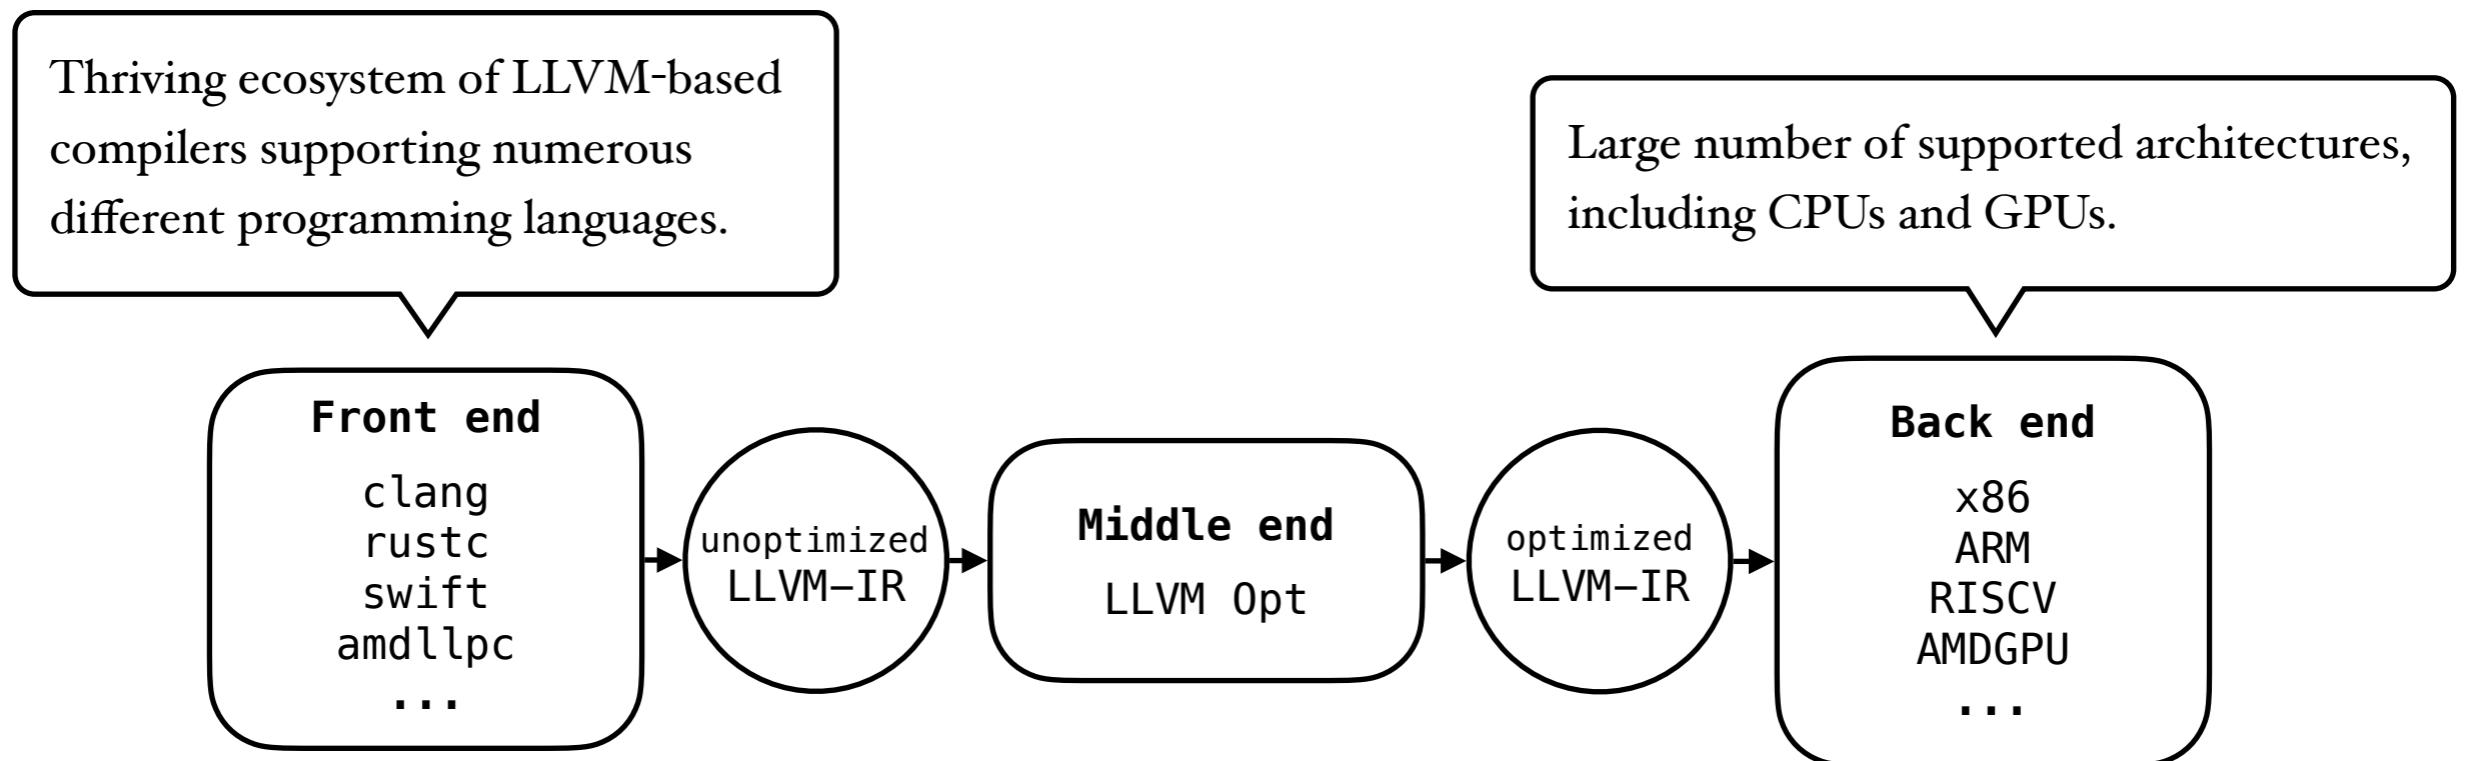
Compiler Architecture

The LLVM Compiler Infrastructure & Clang Compiler



Compiler Architecture

The LLVM Compiler Infrastructure & Clang Compiler



I highly recommend reading the following:

- [Architecture of Open Source Applications - LLVM by Chris Lattner](#)
- [My First Language Frontend with LLVM Tutorial](#)

References:

- <https://llvm.org>
- <https://github.com/llvm/llvm-project>

Compiler Architecture

LLVM-IR

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Compiler Architecture

LLVM-IR

```
$ clang++ -std=c++17 -emit-llvm \
-S //physics_simulation/source/vec3_length.cxx \
-o vec3_length.ll
```

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Compiler Architecture

LLVM-IR

```
$ clang++ -std=c++17 -emit-llvm \
-s //physics_simulation/source/vec3_length.cxx \
-o vec3_length.ll
```

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare float @sqrt (float %x)

%vec3 = type { float, float, float }

define float @_ZN7phy_sim6lengthERKNS_4vec3E (%vec3* %v)
{
entry:
    %v.x_ptr      = getelementptr %vec3, %vec3* %v, i64 0, i32 0
    %v.x          = load float, float* %v.x_ptr
    %v.x2         = fmul float %v.x, %v.x

    %v.y_ptr      = getelementptr %vec3, %vec3* %v, i64 0, i32 1
    %v.yz_ptr     = bitcast float* %v.y_ptr to <2 x float>*
    %v.yz         = load <2 x float>, <2 x float>* %v.yz_ptr
    %v.y2z2       = fmul <2 x float> %v.yz, %v.yz

    %v.y2         = extractelement <2 x float> %v.y2z2, i32 0
    %v.z2         = extractelement <2 x float> %v.y2z2, i32 1

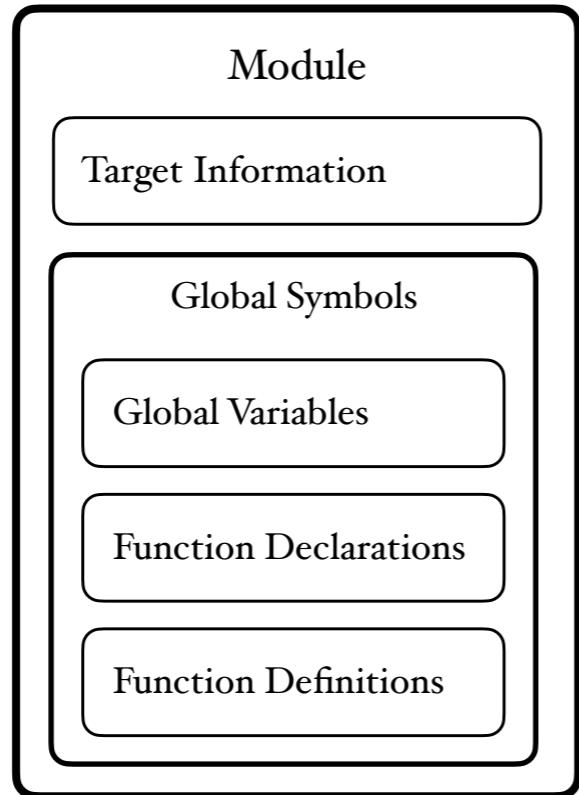
    %v.x2y2z2_sum = fadd float %v.x2      , %v.y2
    %v.x2y2z2z2_sum = fadd float %v.x2y2z2_sum, %v.z2
    %v.length     = call float @sqrt(float %v.x2y2z2z2_sum)

    ret float %v.length
}
```

Compiler Architecture

LLVM-IR

```
$ clang++ -std=c++17 -emit-llvm \\\n-s //physics_simulation/source/vec3_length.cxx \\n-o vec3_length.ll
```

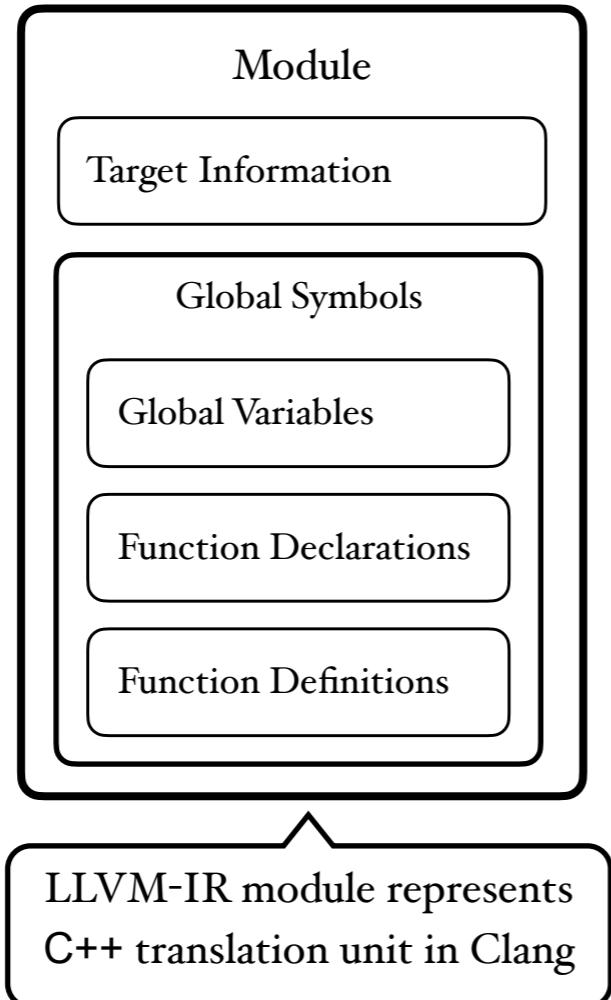


```
namespace phy_sim\n{\n    struct vec3 { float x, y, z; };\n\n    auto length (const vec3& v) -> float\n    {\n        return std::sqrt(v.x * v.x +\n                         v.y * v.y +\n                         v.z * v.z);\n    }\n}\n\ntarget datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"\ntarget triple = "x86_64-apple-macosx10.14.0"\n\ndeclare float @sqrt (float %x)\n\n%vec3 = type { float, float, float }\n\ndefine float @_ZN7phy_sim6lengthERKNS_4vec3E (%vec3* %v)\n{\nentry:\n    %v.x_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 0\n    %v.x = load float, float* %v.x_ptr\n    %v.x2 = fmul float %v.x, %v.x\n\n    %v.y_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 1\n    %v.yz_ptr = bitcast float* %v.y_ptr to <2 x float>*\n    %v.yz = load <2 x float>, <2 x float>* %v.yz_ptr\n    %v.y2z2 = fmul <2 x float> %v.yz, %v.yz\n\n    %v.y2 = extractelement <2 x float> %v.y2z2, i32 0\n    %v.z2 = extractelement <2 x float> %v.y2z2, i32 1\n\n    %v.x2y2_sum = fadd float %v.x2 , %v.y2\n    %v.x2y2z2_sum = fadd float %v.x2y2_sum, %v.z2\n    %v.length = call float @sqrt(float %v.x2y2z2_sum)\n\n    ret float %v.length\n}
```

Compiler Architecture

LLVM-IR

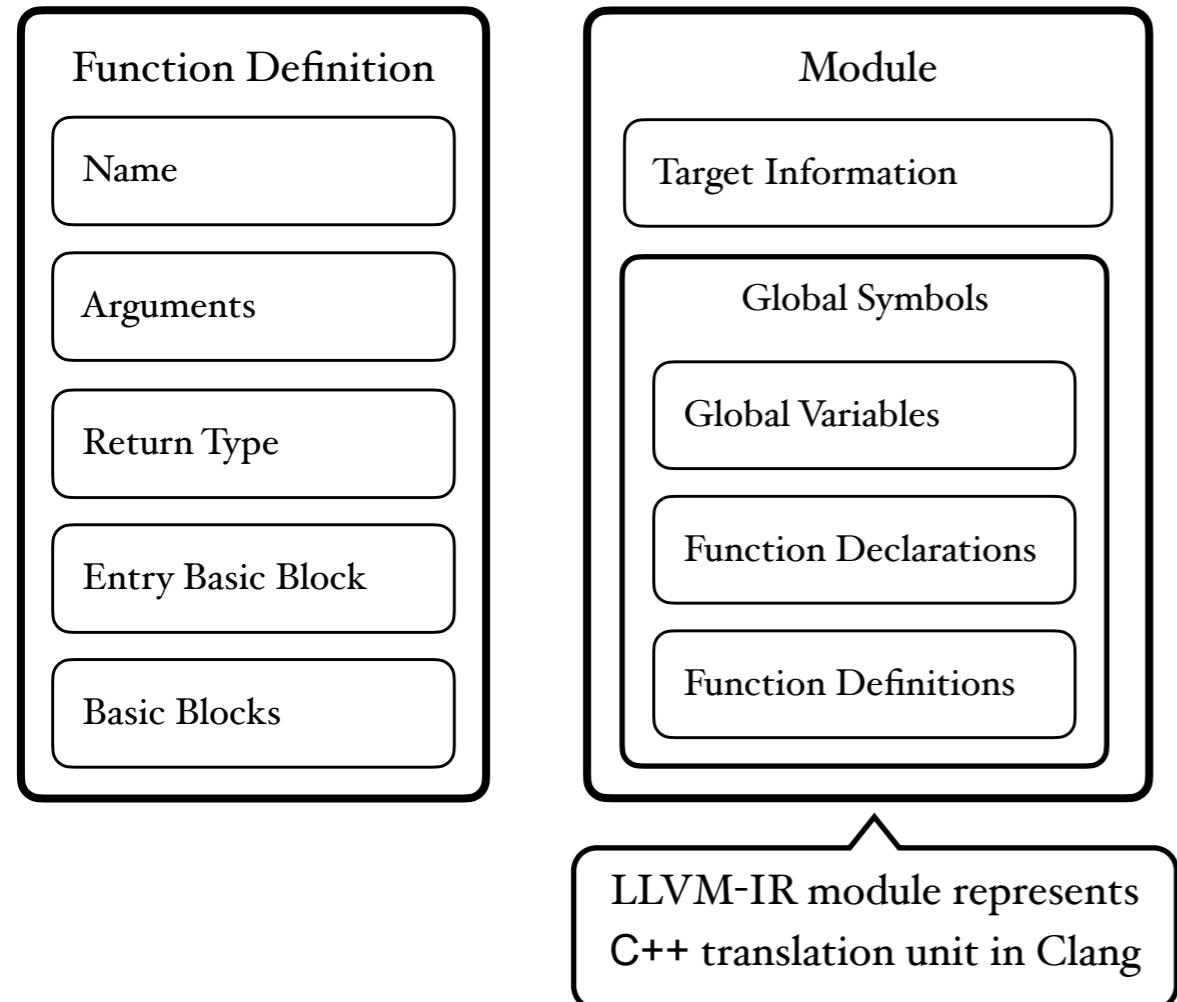
```
$ clang++ -std=c++17 -emit-llvm \\\n-s //physics_simulation/source/vec3_length.cxx \\n-o vec3_length.ll
```



```
namespace phy_sim\n{\n    struct vec3 { float x, y, z; };\n\n    auto length (const vec3& v) -> float\n    {\n        return std::sqrt(v.x * v.x +\n                         v.y * v.y +\n                         v.z * v.z);\n    }\n}\n\ntarget datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"\ntarget triple = "x86_64-apple-macosx10.14.0"\n\ndeclare float @sqrt (float %x)\n\n%vec3 = type { float, float, float }\n\ndefine float @_ZN7phy_sim6lengthERKNS_4vec3E (%vec3* %v)\n{\nentry:\n    %v.x_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 0\n    %v.x = load float, float* %v.x_ptr\n    %v.x2 = fmul float %v.x, %v.x\n\n    %v.y_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 1\n    %v.yz_ptr = bitcast float* %v.y_ptr to <2 x float>*\n    %v.yz = load <2 x float>, <2 x float>* %v.yz_ptr\n    %v.y2z2 = fmul <2 x float> %v.yz, %v.yz\n\n    %v.y2 = extractelement <2 x float> %v.y2z2, i32 0\n    %v.z2 = extractelement <2 x float> %v.y2z2, i32 1\n\n    %v.x2y2_sum = fadd float %v.x2 , %v.y2\n    %v.x2y2z2_sum = fadd float %v.x2y2_sum, %v.z2\n    %v.length = call float @sqrt(float %v.x2y2z2_sum)\n\n    ret float %v.length\n}
```

Compiler Architecture

LLVM-IR



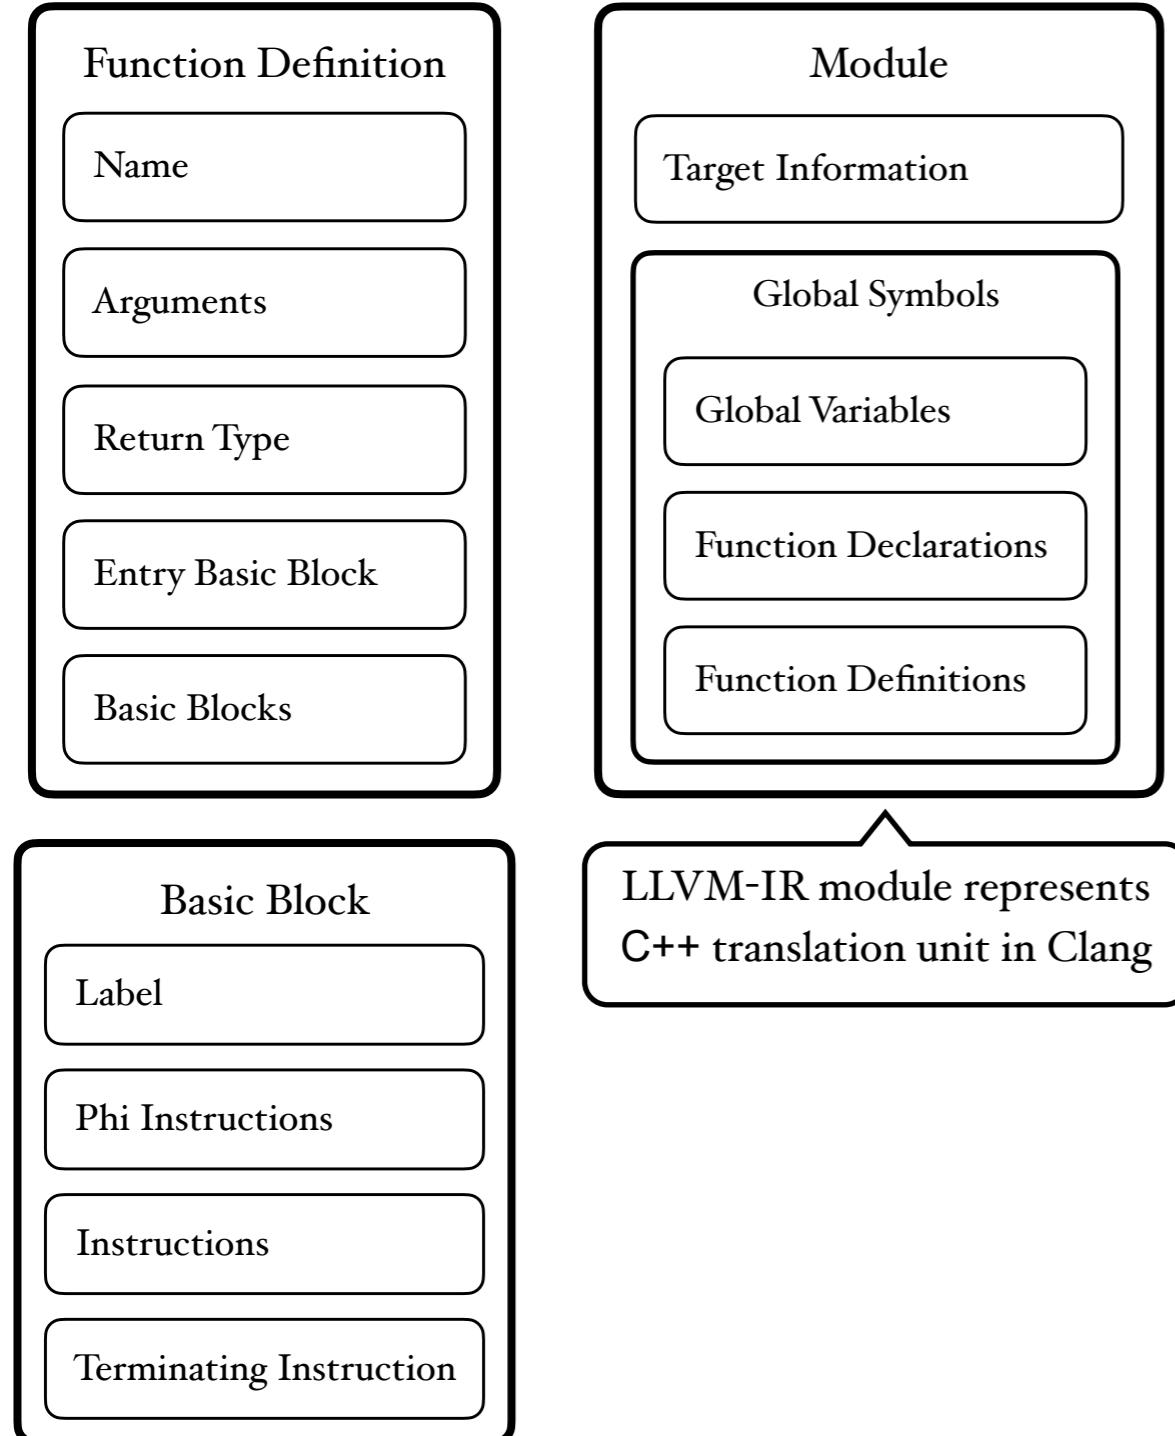
```
$ clang++ -std=c++17 -emit-llvm \\\n-s //physics_simulation/source/vec3_length.cxx \\n-o vec3_length.ll
```

```
namespace phy_sim\n{\n    struct vec3 { float x, y, z; };\n\n    auto length (const vec3& v) -> float\n    {\n        return std::sqrt(v.x * v.x +\n                         v.y * v.y +\n                         v.z * v.z);\n    }\n}\n\ntarget datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"\ntarget triple = "x86_64-apple-macosx10.14.0"\n\ndeclare float @sqrt (float %x)\n\n%vec3 = type { float, float, float }\n\ndefine float @_ZN7phy_sim6lengthERKNS_4vec3E (%vec3* %v)\n{\nentry:\n    %v.x_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 0\n    %v.x = load float, float* %v.x_ptr\n    %v.x2 = fmul float %v.x, %v.x\n\n    %v.y_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 1\n    %v.yz_ptr = bitcast float* %v.y_ptr to <2 x float>*\n    %v.yz = load <2 x float>, <2 x float>* %v.yz_ptr\n    %v.y2z2 = fmul <2 x float> %v.yz, %v.yz\n\n    %v.y2 = extractelement <2 x float> %v.y2z2, i32 0\n    %v.z2 = extractelement <2 x float> %v.y2z2, i32 1\n\n    %v.x2y2_sum = fadd float %v.x2 , %v.y2\n    %v.x2y2z2_sum = fadd float %v.x2y2_sum, %v.z2\n    %v.length = call float @sqrt(float %v.x2y2z2_sum)\n\nret float %v.length\n}
```

Compiler Architecture

LLVM-IR

```
$ clang++ -std=c++17 -emit-llvm \n
-s //physics_simulation/source/vec3_length.cxx \n
-o vec3_length.ll
```

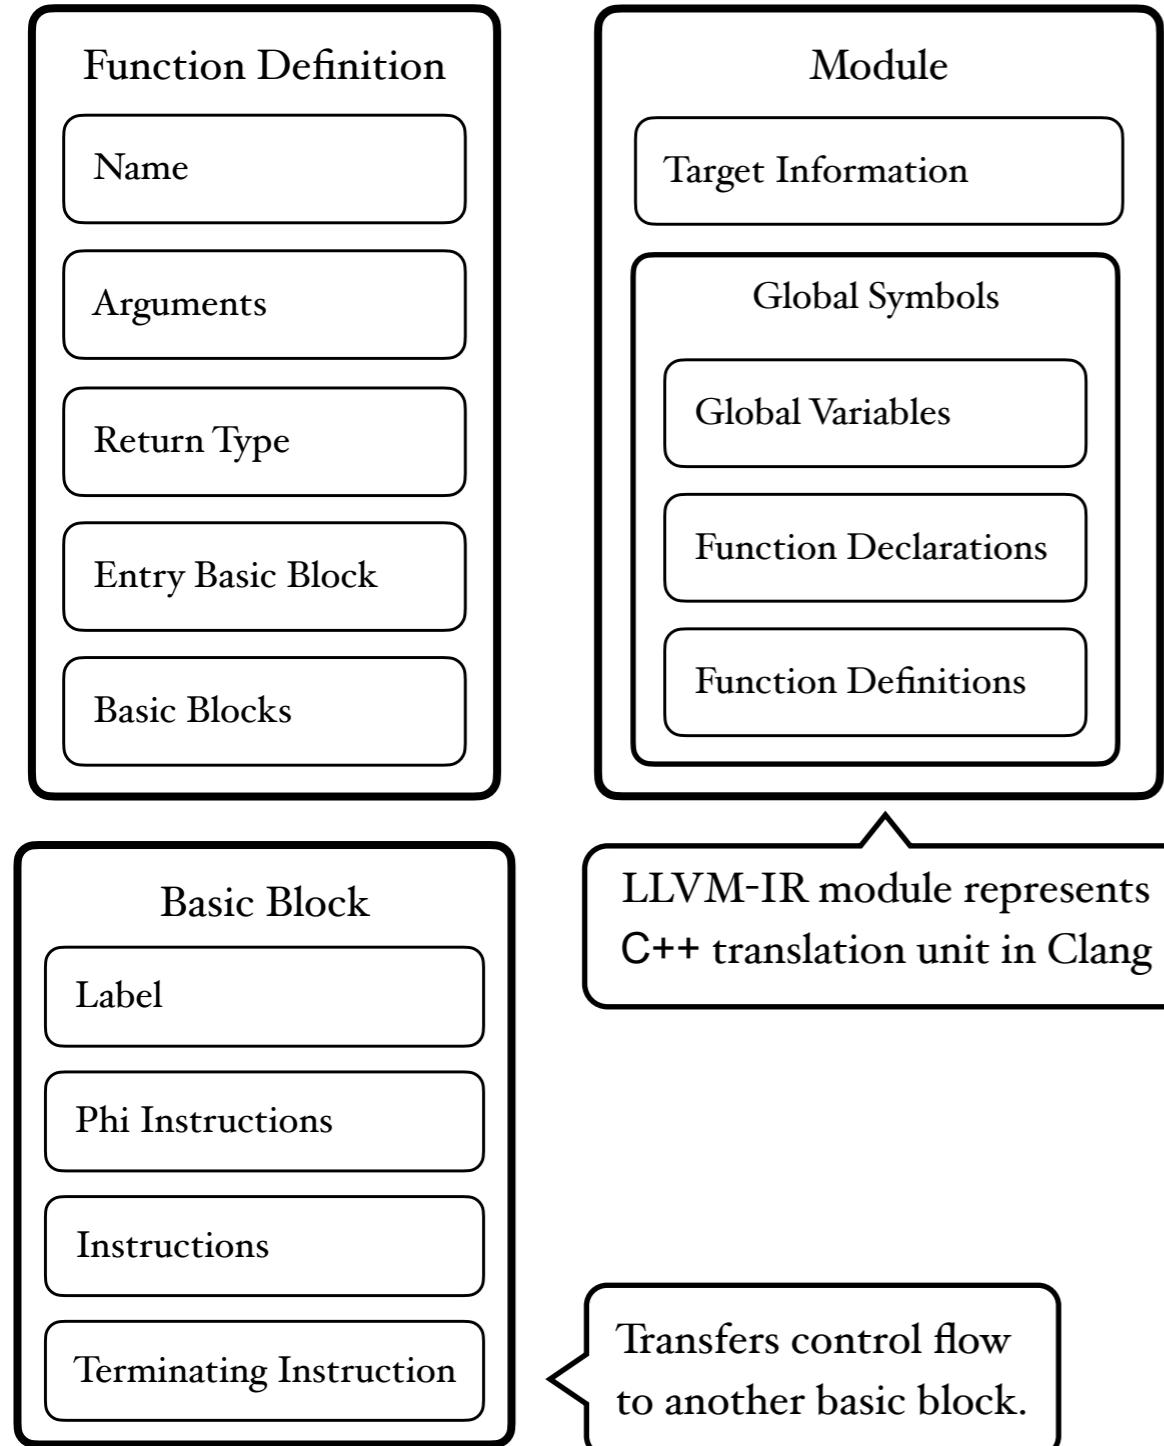


```
namespace phy_sim\n
{\n    struct vec3 { float x, y, z; };\n\n    auto length (const vec3& v) -> float\n    {\n        return std::sqrt(v.x * v.x +\n                         v.y * v.y +\n                         v.z * v.z);\n    }\n}\n\ntarget datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"\ntarget triple = "x86_64-apple-macosx10.14.0"\n\ndeclare float @sqrt (float %x)\n\n%vec3 = type { float, float, float }\n\ndefine float @_ZN7phy_sim6lengthERKNS_4vec3E (%vec3* %v)\n{\nentry:\n    %v.x_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 0\n    %v.x = load float, float* %v.x_ptr\n    %v.x2 = fmul float %v.x, %v.x\n\n    %v.y_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 1\n    %v.yz_ptr = bitcast float* %v.y_ptr to <2 x float>*\n    %v.yz = load <2 x float>, <2 x float>* %v.yz_ptr\n    %v.y2z2 = fmul <2 x float> %v.yz, %v.yz\n\n    %v.y2 = extractelement <2 x float> %v.y2z2, i32 0\n    %v.z2 = extractelement <2 x float> %v.y2z2, i32 1\n\n    %v.x2y2_sum = fadd float %v.x2 , %v.y2\n    %v.x2y2z2_sum = fadd float %v.x2y2_sum, %v.z2\n    %v.length = call float @sqrt(float %v.x2y2z2_sum)\n\nret float %v.length\n}
```

Compiler Architecture

LLVM-IR

```
$ clang++ -std=c++17 -emit-llvm \n-s //physics_simulation/source/vec3_length.cxx \n-o vec3_length.ll
```



```
namespace phy_sim\n{\n    struct vec3 { float x, y, z; };\n\n    auto length (const vec3& v) -> float\n    {\n        return std::sqrt(v.x * v.x +\n                         v.y * v.y +\n                         v.z * v.z);\n    }\n}\n\ntarget datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"\ntarget triple = "x86_64-apple-macosx10.14.0"\n\ndeclare float @sqrt (float %x)\n\n%vec3 = type { float, float, float }\n\ndefine float @_ZN7phy_sim6lengthERKNS_4vec3E (%vec3* %v)\n{\nentry:\n    %v.x_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 0\n    %v.x = load float, float* %v.x_ptr\n    %v.x2 = fmul float %v.x, %v.x\n\n    %v.y_ptr = getelementptr %vec3, %vec3* %v, i64 0, i32 1\n    %v.yz_ptr = bitcast float* %v.y_ptr to <2 x float>*\n    %v.yz = load <2 x float>, <2 x float>* %v.yz_ptr\n    %v.y2z2 = fmul <2 x float> %v.yz, %v.yz\n\n    %v.y2 = extractelement <2 x float> %v.y2z2, i32 0\n    %v.z2 = extractelement <2 x float> %v.y2z2, i32 1\n\n    %v.x2y2_sum = fadd float %v.x2 , %v.y2\n    %v.x2y2z2_sum = fadd float %v.x2y2_sum, %v.z2\n    %v.length = call float @sqrt(float %v.x2y2z2_sum)\n\nret float %v.length\n}
```

Compiler Architecture

LLVM-IR

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```

Compiler Architecture

LLVM-IR

```
; copy_file.ll

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare i32 @getchar()
declare i32 @putchar(i32)

define i32 @main()
{
entry:
%first_char      = call i32 @getchar()
%first_char_eq_eof = icmp eq i32 %first_char, -1;EOF

br i1 %first_char_eq_eof, label %exit, label %loop

loop:
%out_char        = phi i32 [ %next_char, %loop ],
                  [ %first_char, %entry ]
call i32 @putchar(i32 %out_char)

%next_char       = call i32 @getchar()
%next_char_eq_eof = icmp eq i32 %next_char, -1;EOF

br i1 %next_char_eq_eof, label %exit, label %loop

exit:
ret i32 0
}
```

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```

Compiler Architecture

LLVM-IR

LLVM-IR has strong static typing,
supporting floating point and integer types,
as well as pointers, functions, structures, arrays and vectors.

```
; copy_file.ll

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare i32 @getchar()

declare i32 @putchar (i32)

define i32 @main ()
{
entry:
%first_char      = call i32 @getchar()
%first_char_eq_eof = icmp eq i32 %first_char, -1;EOF

br i1 %first_char_eq_eof, label %exit, label %loop

loop:
%out_char        = phi i32 [ %next_char, %loop
                           [ %first_char, %entry ] ],
                  call i32 @putchar(i32 %out_char)

%next_char       = call i32 @getchar()
%next_char_eq_eof = icmp eq i32 %next_char, -1;EOF

br i1 %next_char_eq_eof, label %exit, label %loop

exit:
ret i32 0
}
```

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```

Compiler Architecture

LLVM-IR

LLVM-IR has strong static typing,
supporting floating point and integer types,
as well as pointers, functions, structures, arrays and vectors.

```
; copy_file.ll

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare i32 @getchar()
declare i32 @putchar(i32)
define i32 @main()
{
entry:
%first_char      = call i32 @getchar()
%first_char_eq_eof = icmp eq i32 %first_char, -1;EOF

br i1 %first_char_eq_eof, label %exit, label %loop

loop:
%out_char        = phi i32 [ %next_char, %loop ],
                  [ %first_char, %entry ]
call i32 @putchar(i32 %out_char)

%next_char       = call i32 @getchar()
%next_char_eq_eof = icmp eq i32 %next_char, -1;EOF

br i1 %next_char_eq_eof, label %exit, label %loop

exit:
ret i32 0
}
```

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```

LLVM-IR code has **Static Single Assignment** (SSA) form,
which means that once value is returned from an instruction
and assigned to a virtual register it must never change.

Compiler Architecture

LLVM-IR

```
; copy_file.ll

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare i32 @getchar()
declare i32 @putchar(i32)
define i32 @main()
{
entry:
%first_char      = call i32 @getchar()
%first_char_eq_eof = icmp eq i32 %first_char, -1;EOF

br i1 %first_char_eq_eof, label %exit, label %loop

loop:
%out_char        = phi i32 [ %next_char, %loop ], [
call    i32 @putchar(i32 %out_char)

%next_char       = call i32 @getchar()
%next_char_eq_eof = icmp eq i32 %next_char, -1;EOF

br i1 %next_char_eq_eof, label %exit, label %loop

exit:
ret i32 0
}
```

LLVM-IR has strong static typing, supporting floating point and integer types, as well as pointers, functions, structures, arrays and vectors.

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```

LLVM-IR code has **Static Single Assignment** (SSA) form, which means that once value is returned from an instruction and assigned to a virtual register it must never change.

Phi node selects value from one of predeceasing basic blocks, depending on control flow.

Compiler Architecture

LLVM-IR

LLVM-IR has strong static typing, supporting floating point and integer types, as well as pointers, functions, structures, arrays and vectors.

```
; copy_file.ll
```

```
target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"
```

```
declare i32 @getchar()
```

```
declare i32 @putchar(i32)
```

```
define i32 @main()
```

```
{  
entry:
```

```
  %first_char      = call i32 @getchar()  
  %first_char_eq_eof = icmp eq i32 %first_char, -1;EOF
```

```
  br i1 %first_char_eq_eof, label %exit, label %loop
```

```
loop:
```

```
  %out_char       = phi i32 [ %next_char, %loop  
                           [ %first_char, %entry ] ],  
                           call i32 @putchar(i32 %out_char)
```

```
  %next_char      = call i32 @getchar()  
  %next_char_eq_eof = icmp eq i32 %next_char, -1;EOF
```

```
  br i1 %next_char_eq_eof, label %exit, label %loop
```

```
exit:
```

```
  ret i32 0
```

```
#include <stdio.h>  
  
int main ()  
{  
    int c;  
  
    while ((c = getchar()) != EOF)  
    {  
        putchar(c);  
    }  
}
```

LLVM-IR code has **Static Single Assignment** (SSA) form, which means that once value is returned from an instruction and assigned to a virtual register it must never change.

Phi node selects value from one of predeceasing basic blocks, depending on control flow.

Each LLVM-IR function defines a **Control Flow Graph** (CFG), in which vertices are basic blocks, from which that function consists of, and edges are control flow transfers from one basic block to another.

Compiler Architecture

LLVM-IR

```
; copy_file.ll

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare i32 @getchar()
declare i32 @putchar(i32)

define i32 @main()
{
entry:
%first_char      = call i32 @getchar()
%first_char_eq_eof = icmp eq i32 %first_char, -1;EOF

br i1 %first_char_eq_eof, label %exit, label %loop

loop:
%out_char        = phi i32 [ %next_char, %loop
                           [ %first_char, %entry ] ],
                  call i32 @putchar(i32 %out_char)

%next_char       = call i32 @getchar()
%next_char_eq_eof = icmp eq i32 %next_char, -1;EOF

br i1 %next_char_eq_eof, label %exit, label %loop

exit:
ret i32 0
}
```

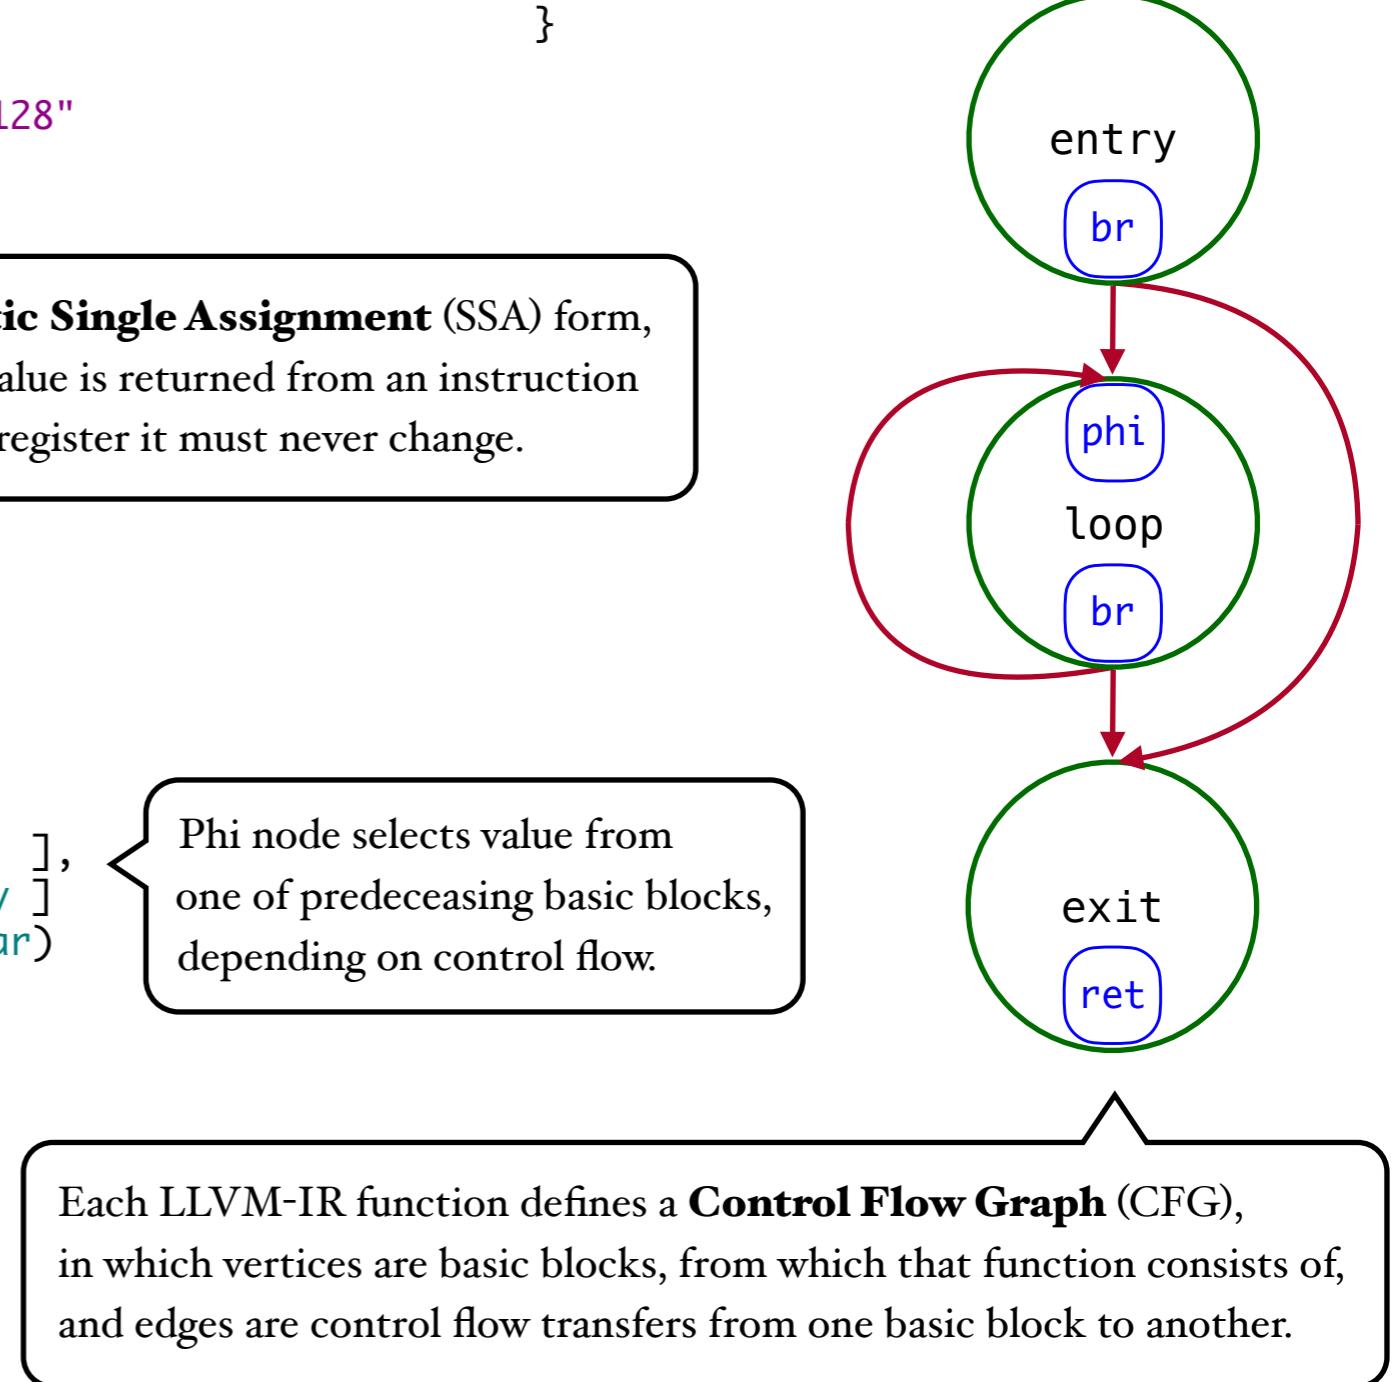
LLVM-IR has strong static typing, supporting floating point and integer types, as well as pointers, functions, structures, arrays and vectors.

LLVM-IR code has **Static Single Assignment** (SSA) form, which means that once value is returned from an instruction and assigned to a virtual register it must never change.

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```



Compiler Architecture

LLVM-IR

```
; copy_file.ll

target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
target triple    = "x86_64-apple-macosx10.14.0"

declare i32 @getchar()
declare i32 @putchar(i32)

define i32 @main()
{
entry:
%first_char      = call i32 @getchar()
%first_char_eq_eof = icmp eq i32 %first_char, -1;EOF

br i1 %first_char_eq_eof, label %exit, label %loop

loop:
%out_char        = phi i32 [ %next_char, %loop
                           [ %first_char, %entry ] ],
                   call i32 @putchar(i32 %out_char)

%next_char       = call i32 @getchar()
%next_char_eq_eof = icmp eq i32 %next_char, -1;EOF

br i1 %next_char_eq_eof, label %exit, label %loop

exit:
ret i32 0
}
```

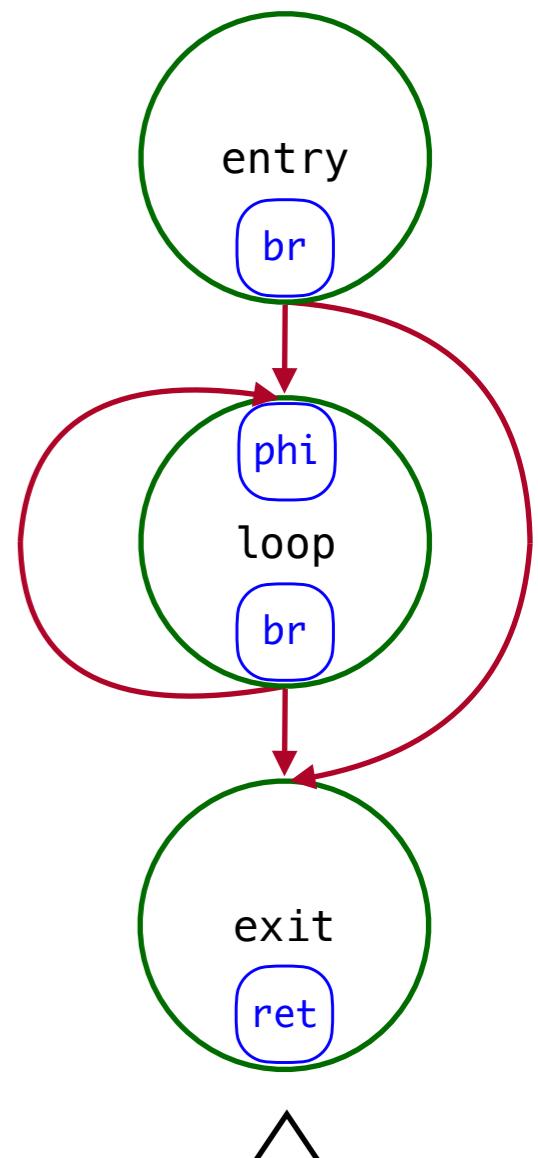
LLVM-IR has strong static typing, supporting floating point and integer types, as well as pointers, functions, structures, arrays and vectors.

LLVM-IR code has **Static Single Assignment** (SSA) form, which means that once value is returned from an instruction and assigned to a virtual register it must never change.

```
#include <stdio.h>

int main ()
{
    int c;

    while ((c = getchar()) != EOF)
    {
        putchar(c);
    }
}
```



Phi node selects value from one of predeceasing basic blocks, depending on control flow.

Each LLVM-IR function defines a **Control Flow Graph** (CFG), in which vertices are basic blocks, from which that function consists of, and edges are control flow transfers from one basic block to another.

```
$ llc copy_file.ll -o copy_file.s
$ as copy_file.s -o copy_file.o
```

Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

C++ objects representing LLVM-IR code implemented in LLVM Libraries

2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference

Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

C++ objects representing LLVM-IR code implemented in LLVM Libraries

2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference

```
// llvm.IR
llvm::Module
llvm::GlobalVariable
llvm::Function
llvm::BasicBlock
llvm::Instruction
llvm::PHINode
llvm::CallInst
llvm::AllocaInst
llvm::LoadInst
llvm::StoreInst
llvm::CmpInst
llvm::ConstantInt
llvm::GetElementPtrInst
llvm::BitCastInst
llvm::ExtractElementInst
llvm::BinaryOperator
llvm::BranchInst
llvm::ReturnInst
```

Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

C++ objects representing LLVM-IR code implemented in LLVM Libraries

2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference

```
// llvm.IR
llvm::Module
llvm::GlobalVariable
llvm::Function
llvm::BasicBlock
llvm::Instruction
llvm::PHINode
llvm::CallInst
llvm::AllocaInst
llvm::LoadInst
llvm::StoreInst
llvm::CmpInst
llvm::ConstantInt
llvm::GetElementPtrInst
llvm::BitCastInst
llvm::ExtractElementInst
llvm::BinaryOperator
llvm::BranchInst
llvm::ReturnInst
```



bitcode.bc

Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

C++ objects representing LLVM-IR code implemented in LLVM Libraries

2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference

```
// llvm.IR
llvm::Module
llvm::GlobalVariable
llvm::Function
llvm::BasicBlock
llvm::Instruction
llvm::PHINode
llvm::CallInst
llvm::AllocaInst
llvm::LoadInst
llvm::StoreInst
llvm::CmpInst
llvm::ConstantInt
llvm::GetElementPtrInst
llvm::BitCastInst
llvm::ExtractElementInst
llvm::BinaryOperator
llvm::BranchInst
llvm::ReturnInst
```



bitcode.bc



assembly.ll

Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

C++ objects representing LLVM-IR code implemented in LLVM Libraries

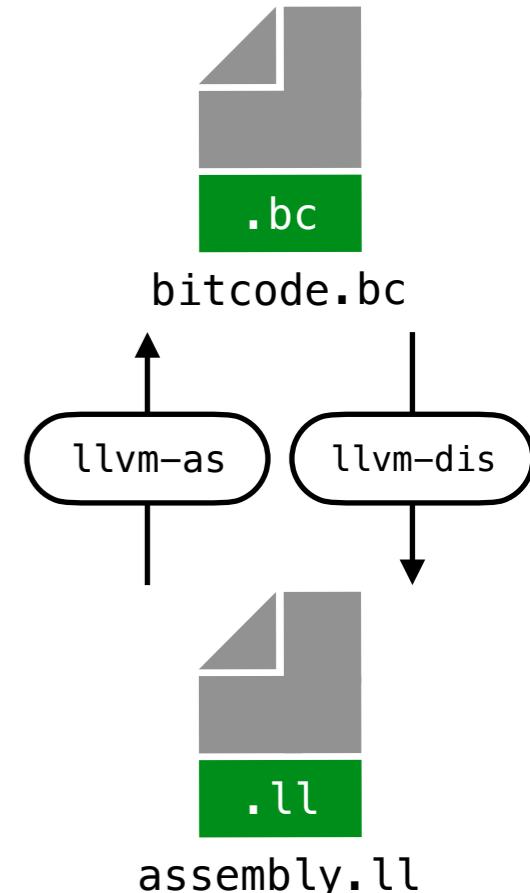
2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference

```
// llvm.IR
llvm::Module
llvm::GlobalVariable
llvm::Function
llvm::BasicBlock
llvm::Instruction
llvm::PHINode
llvm::CallInst
llvm::AllocaInst
llvm::LoadInst
llvm::StoreInst
llvm::CmpInst
llvm::ConstantInt
llvm::GetElementPtrInst
llvm::BitCastInst
llvm::ExtractElementInst
llvm::BinaryOperator
llvm::BranchInst
llvm::ReturnInst
```



Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

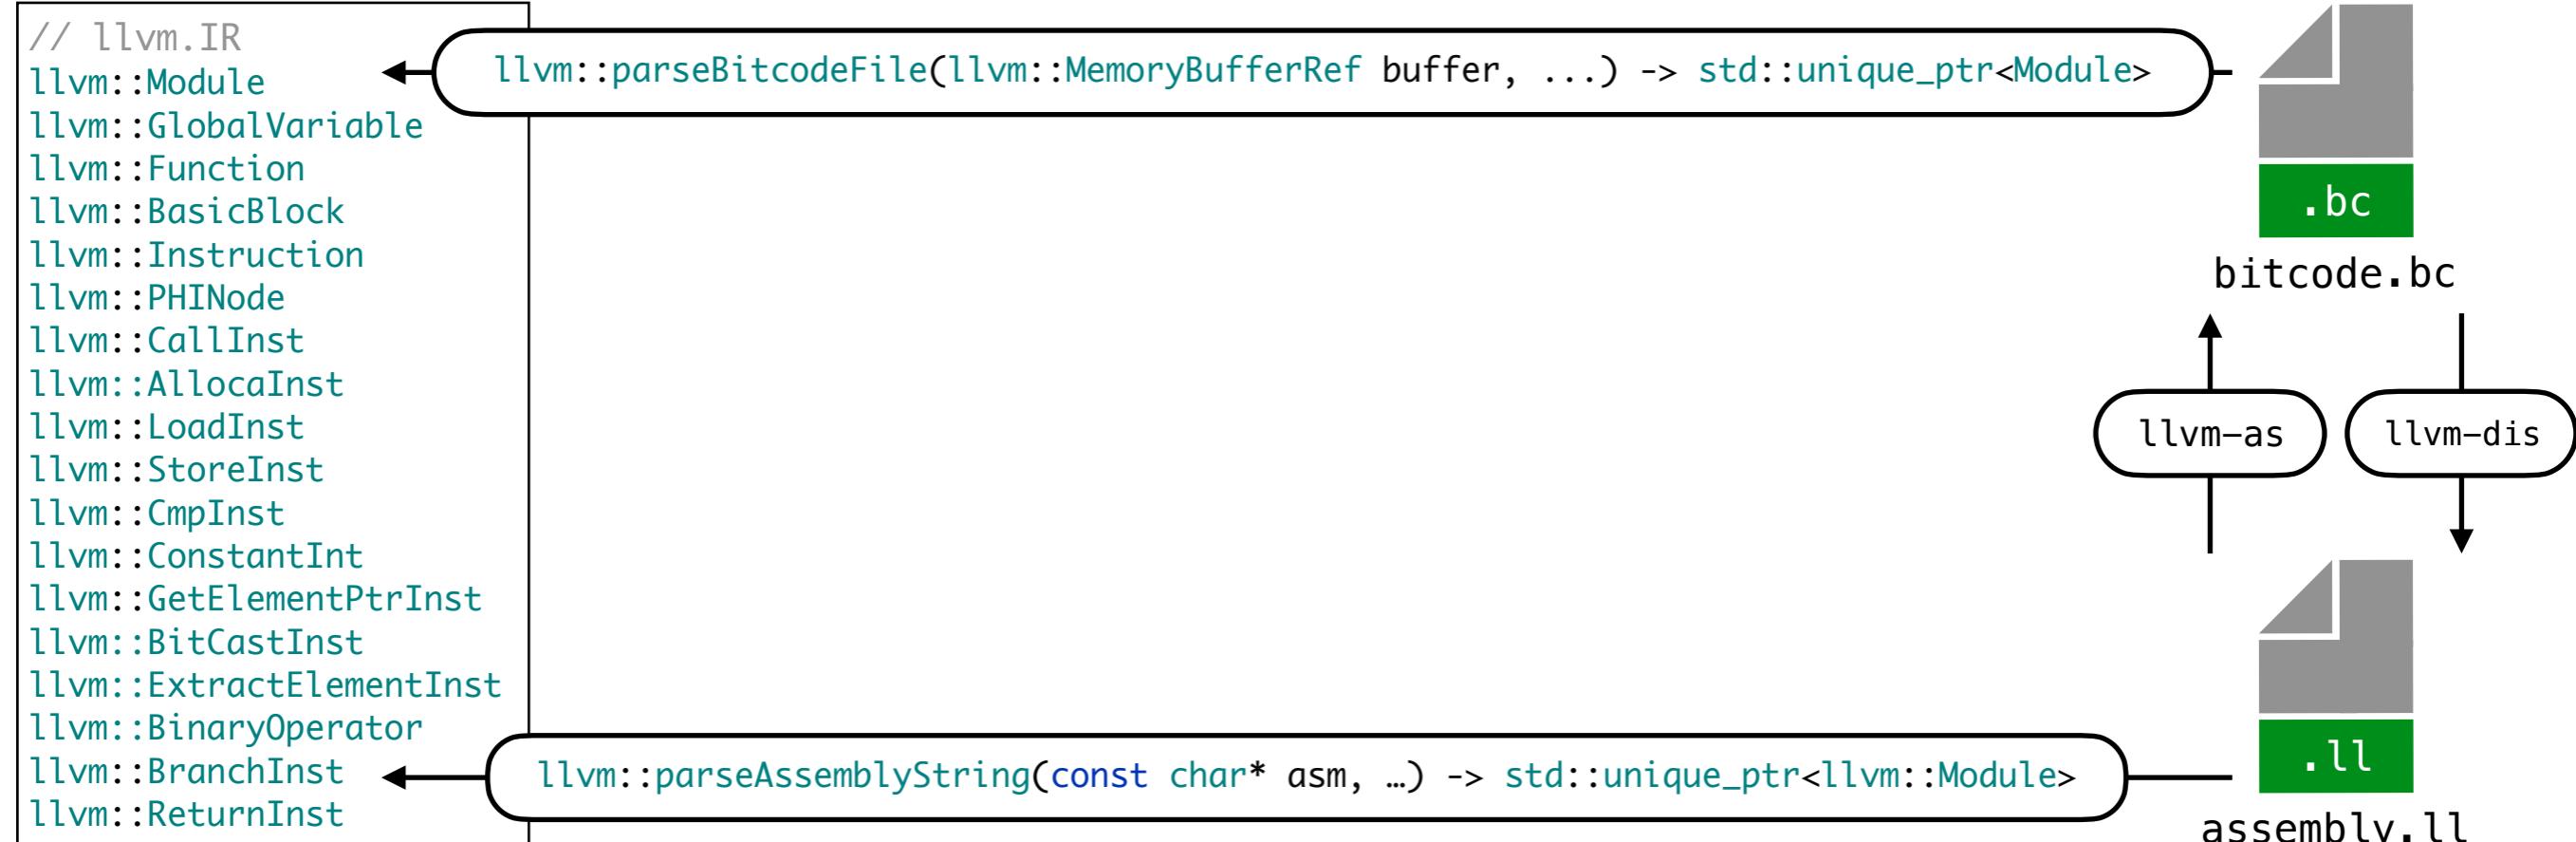
C++ objects representing LLVM-IR code implemented in LLVM Libraries

2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference



Compiler Architecture

Representation of LLVM-IR

LLVM-IR has 3 equivalent representations:

1. **in-memory representation**

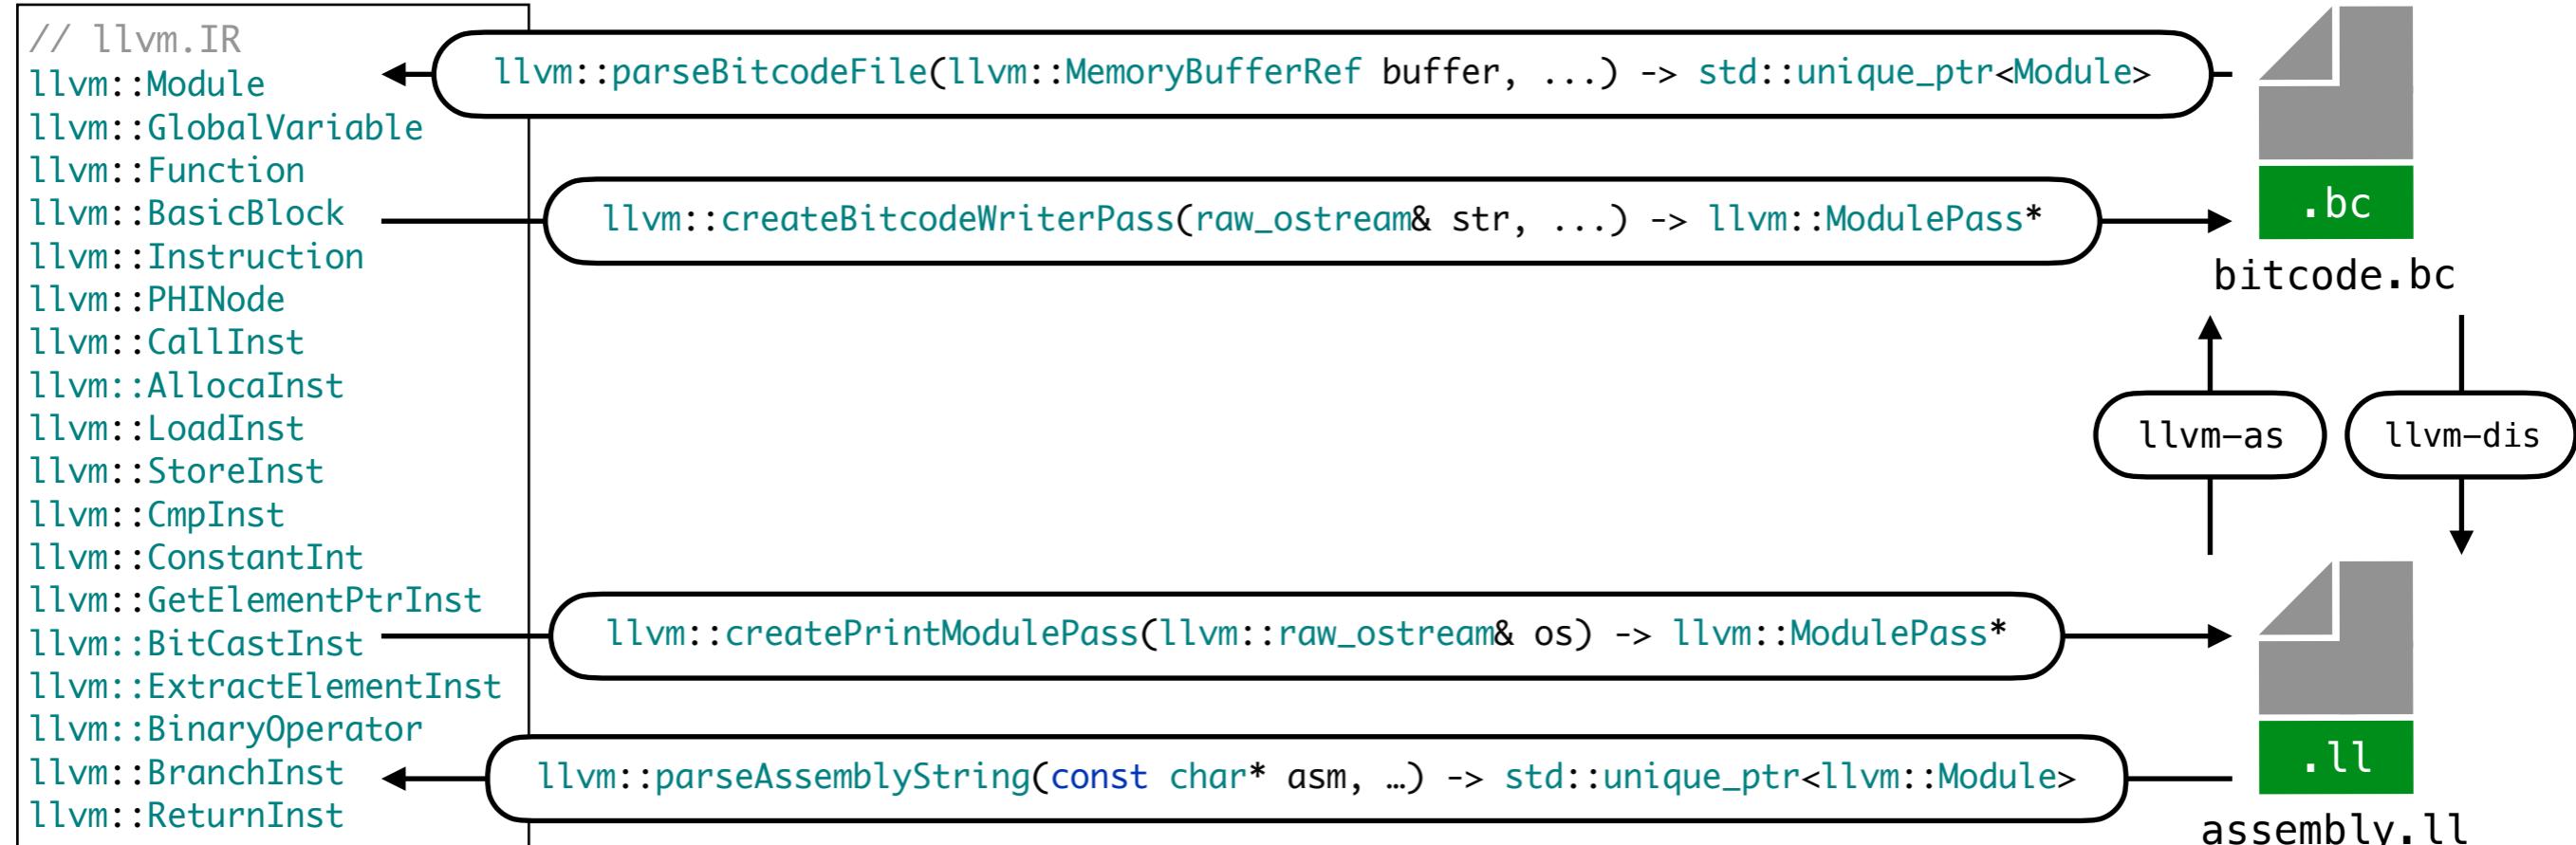
C++ objects representing LLVM-IR code implemented in LLVM Libraries

2. **on-disk representation**

stream of bits representing LLVM-IR code encoded according to LLVM Bitcode Format

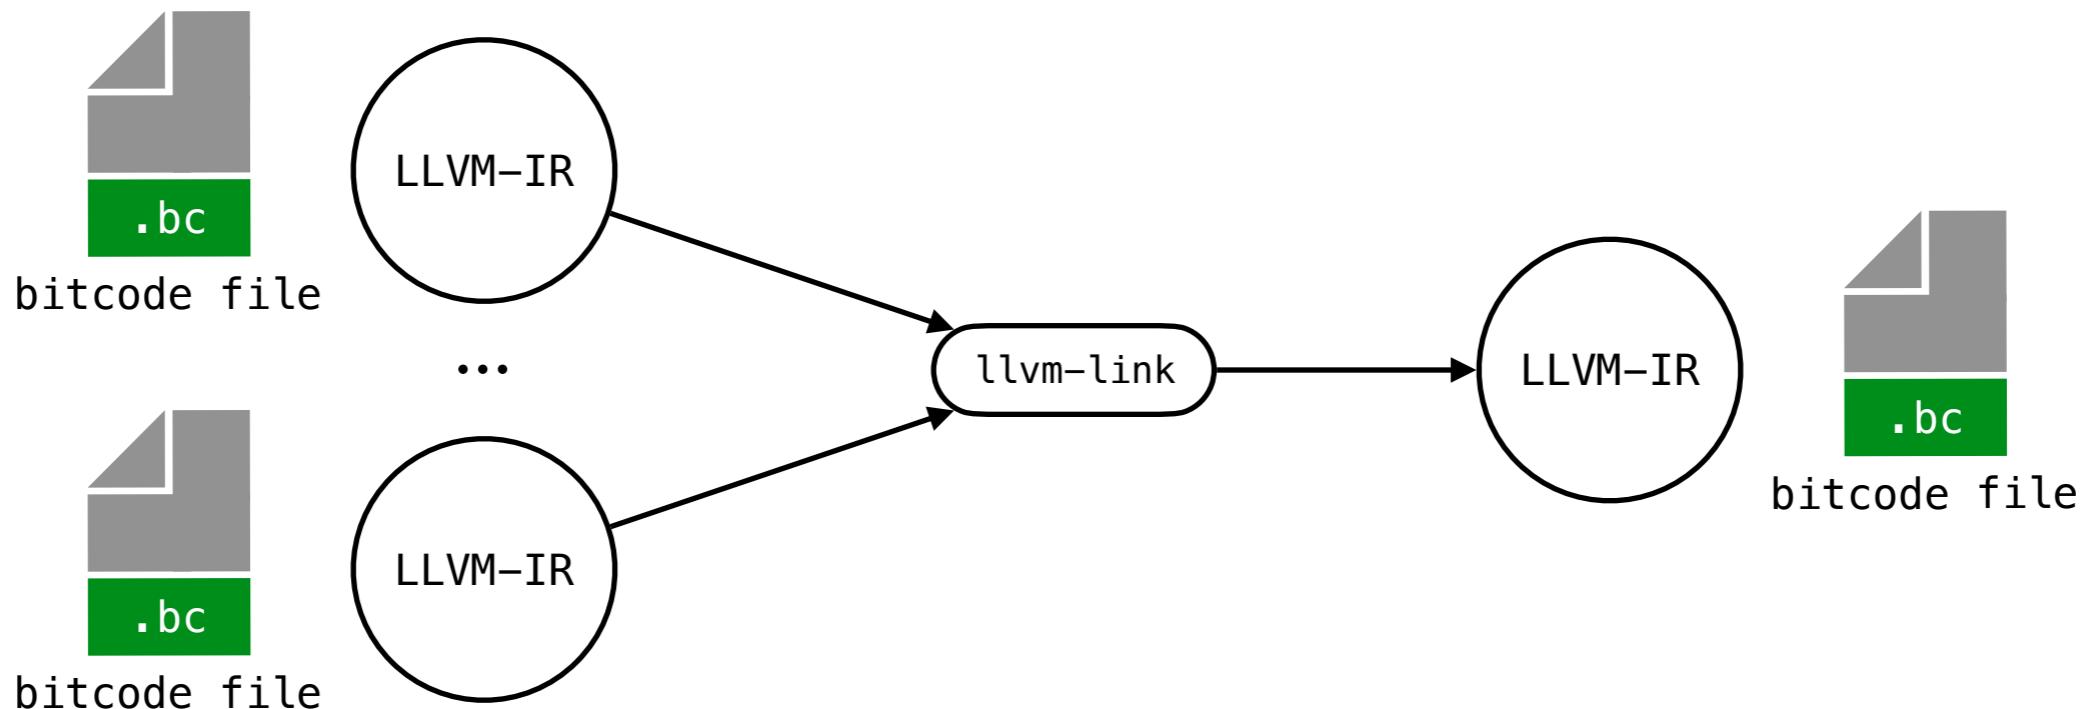
3. **human readable assembly**

stream of characters representing LLVM-IR code interpreted according to the LLVM Language Reference



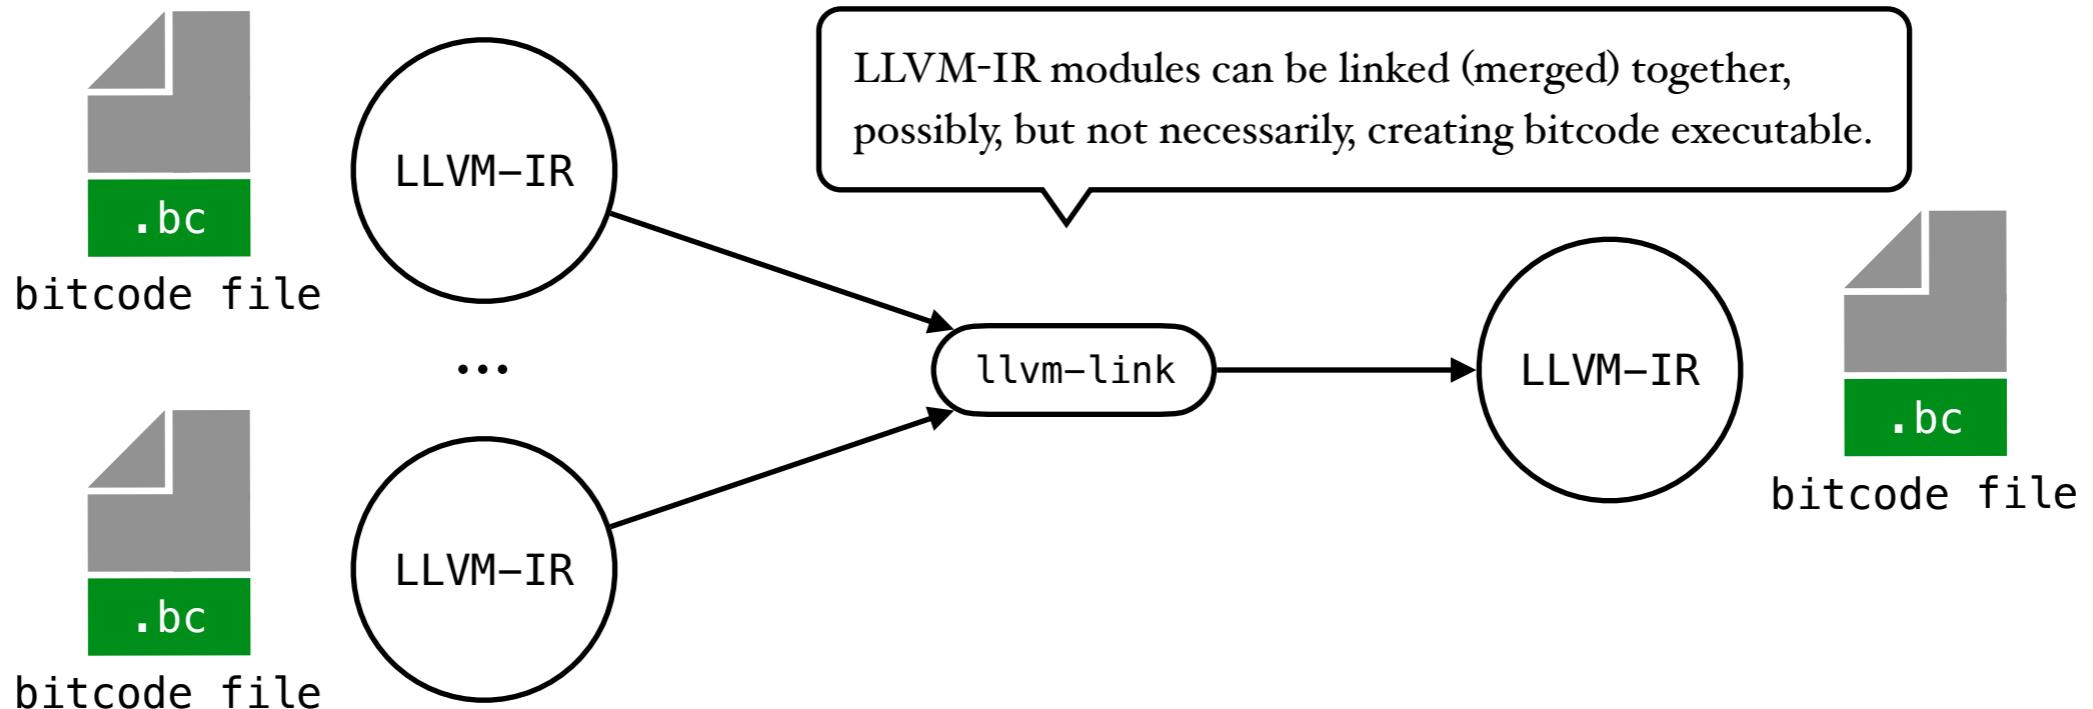
Compiler Architecture

Linking & Interpreting LLVM-IR



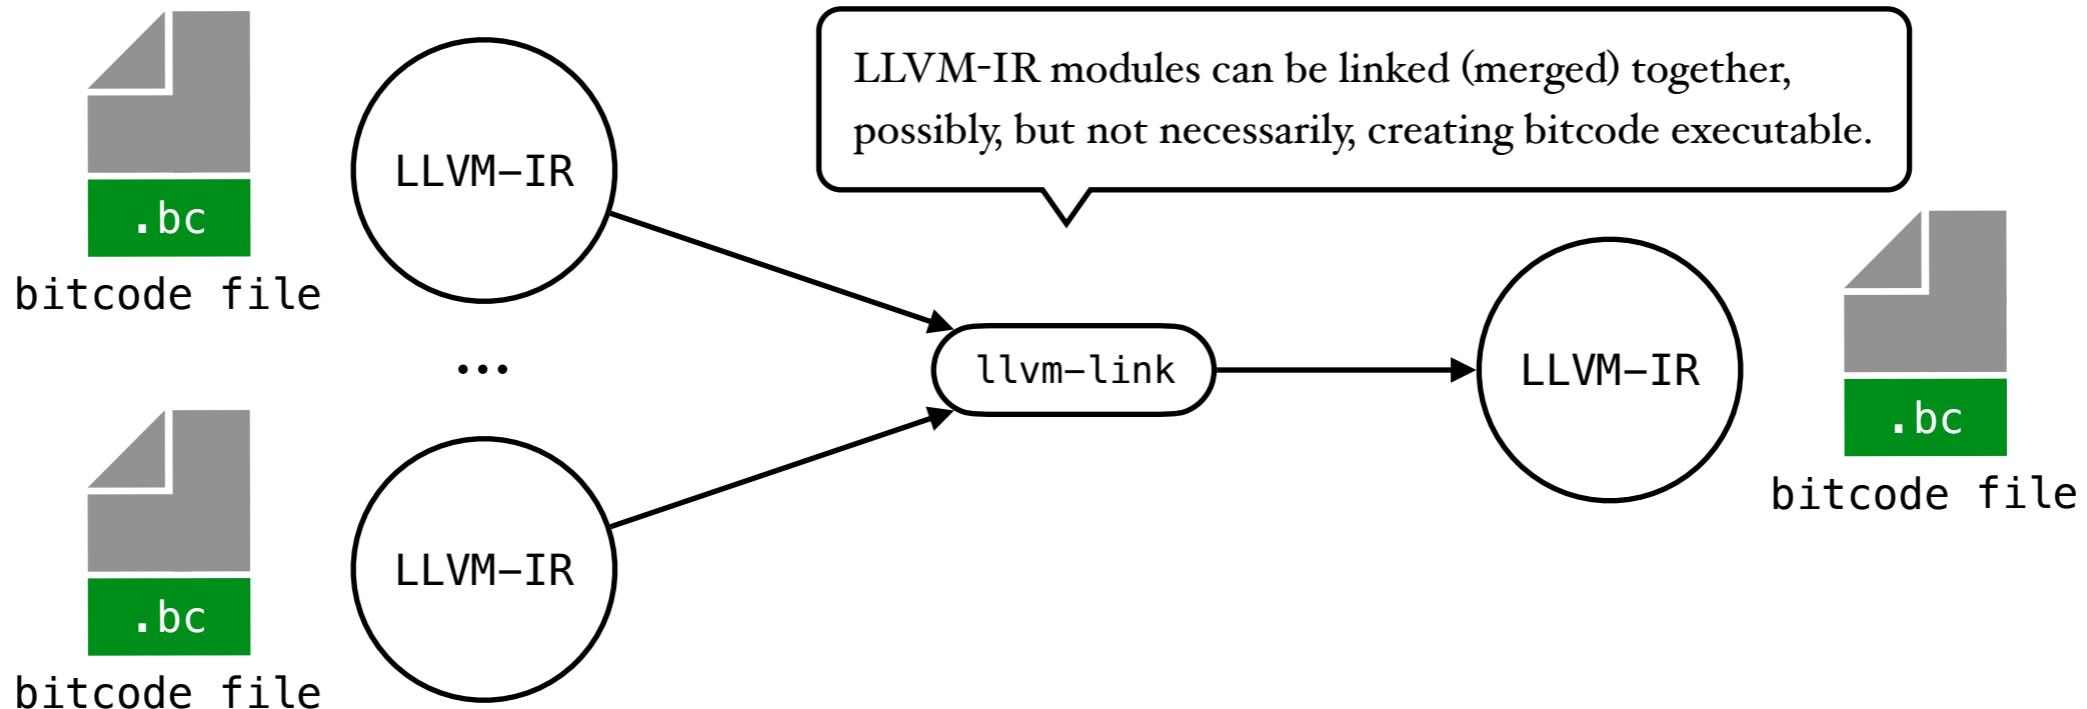
Compiler Architecture

Linking & Interpreting LLVM-IR

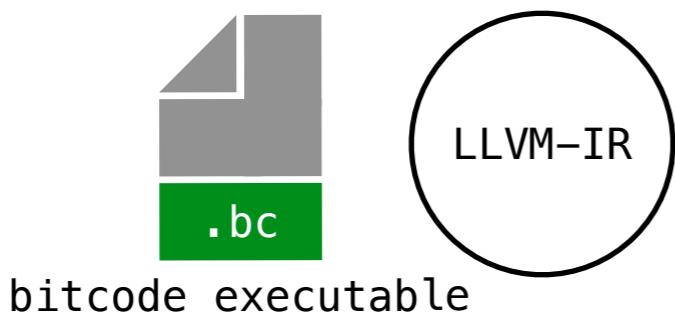


Compiler Architecture

Linking & Interpreting LLVM-IR

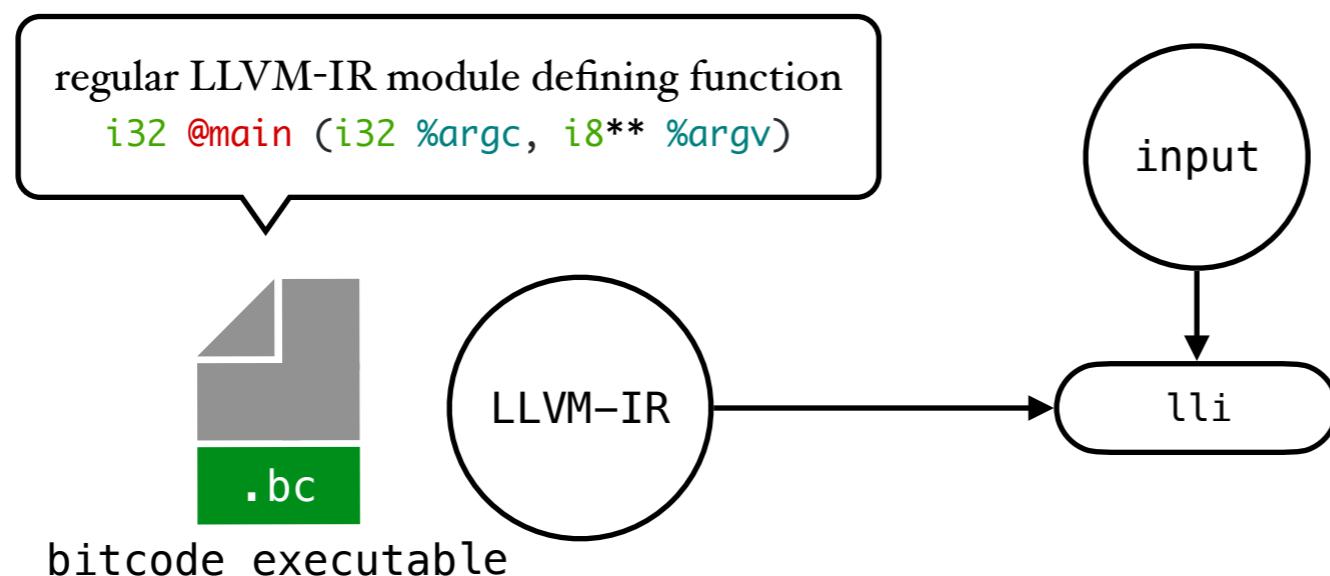
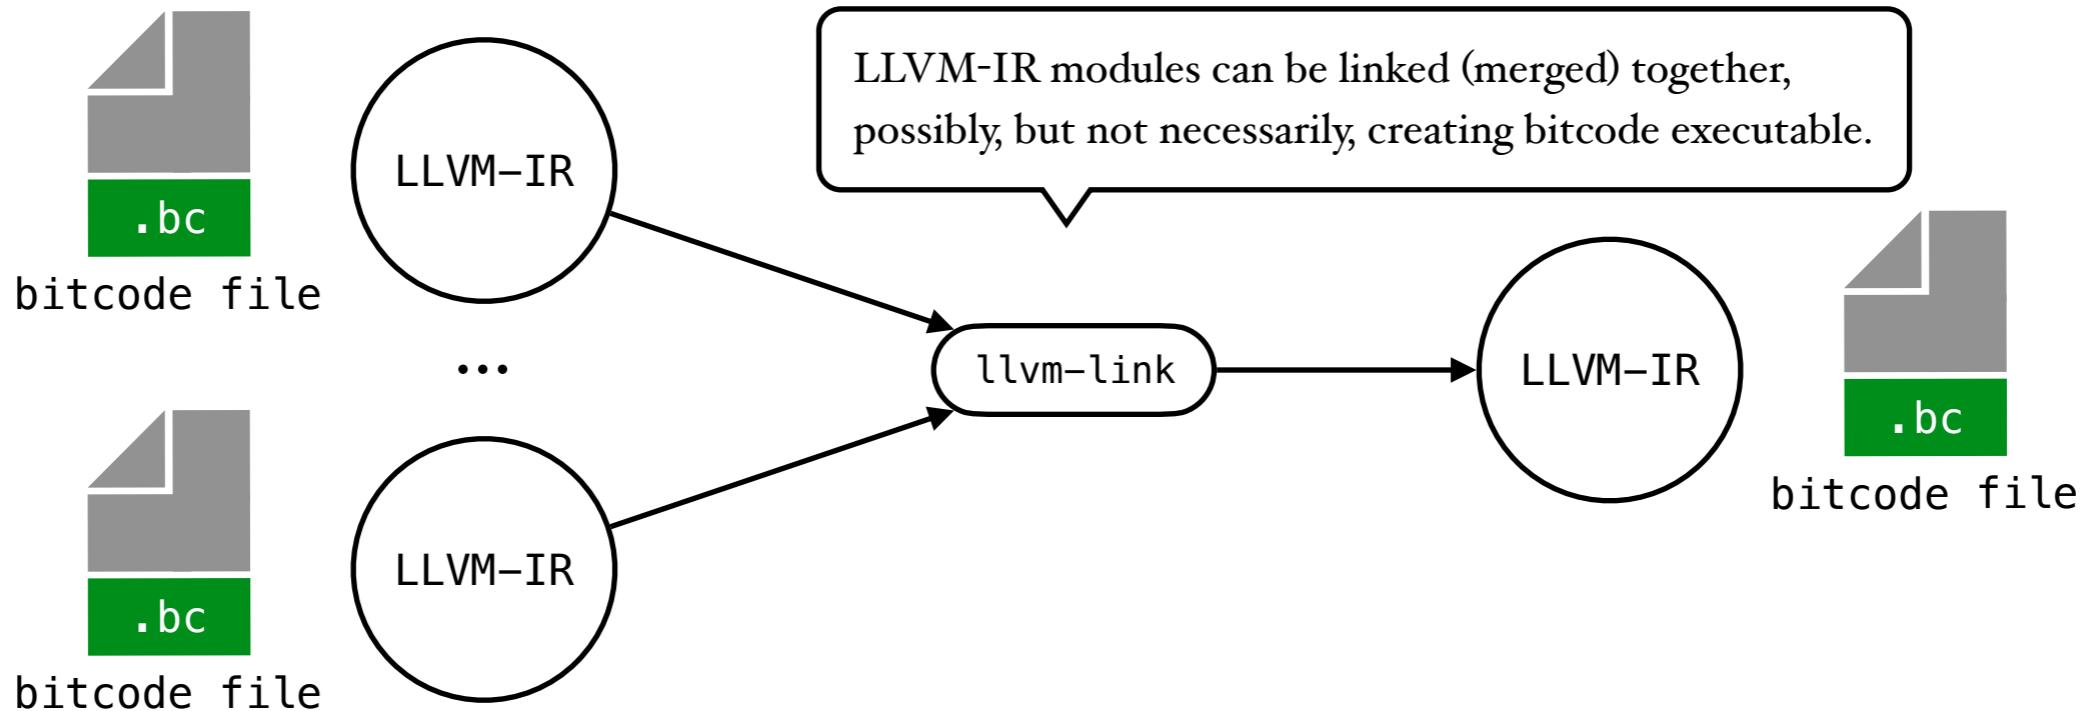


regular LLVM-IR module defining function
`i32 @main (i32 %argc, i8** %argv)`



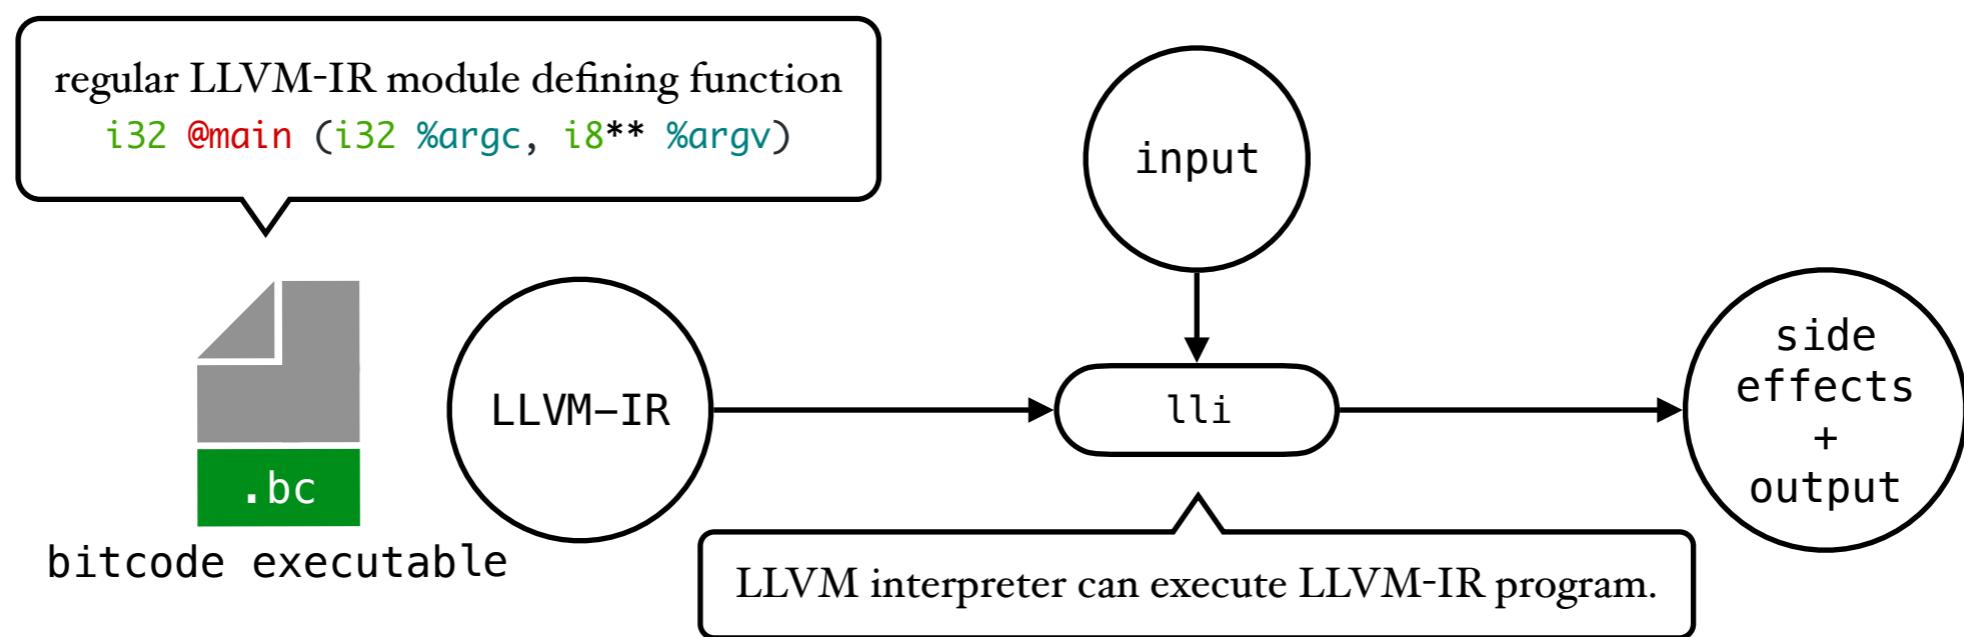
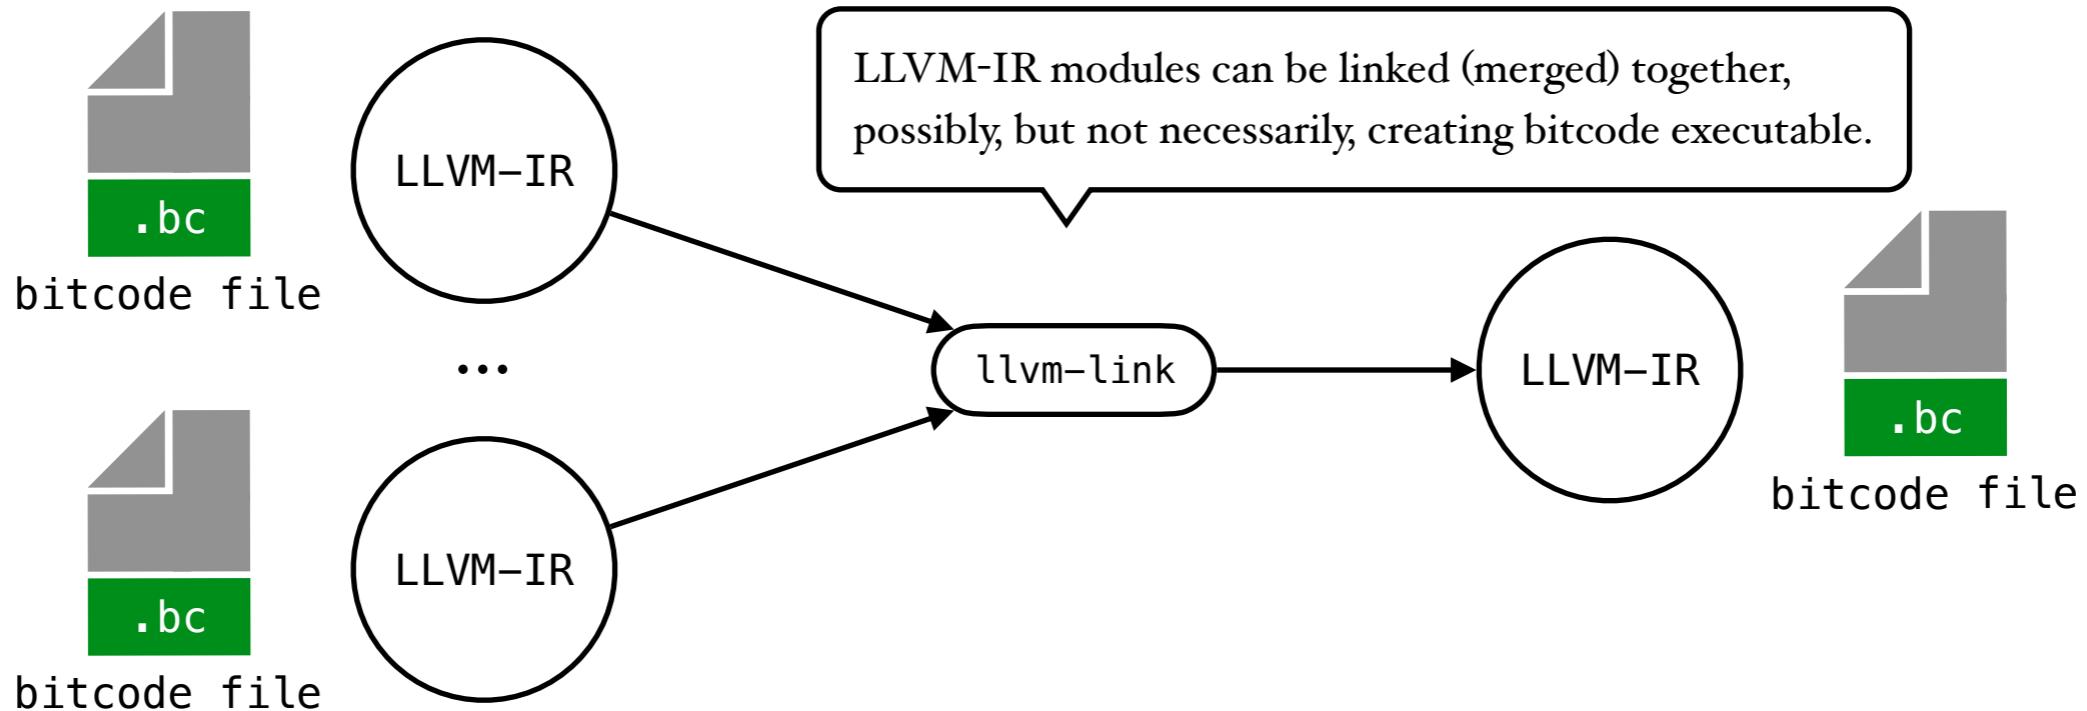
Compiler Architecture

Linking & Interpreting LLVM-IR



Compiler Architecture

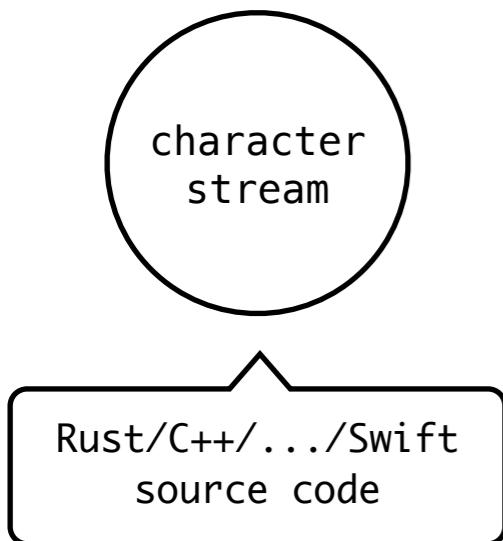
Linking & Interpreting LLVM-IR



Compiler Architecture

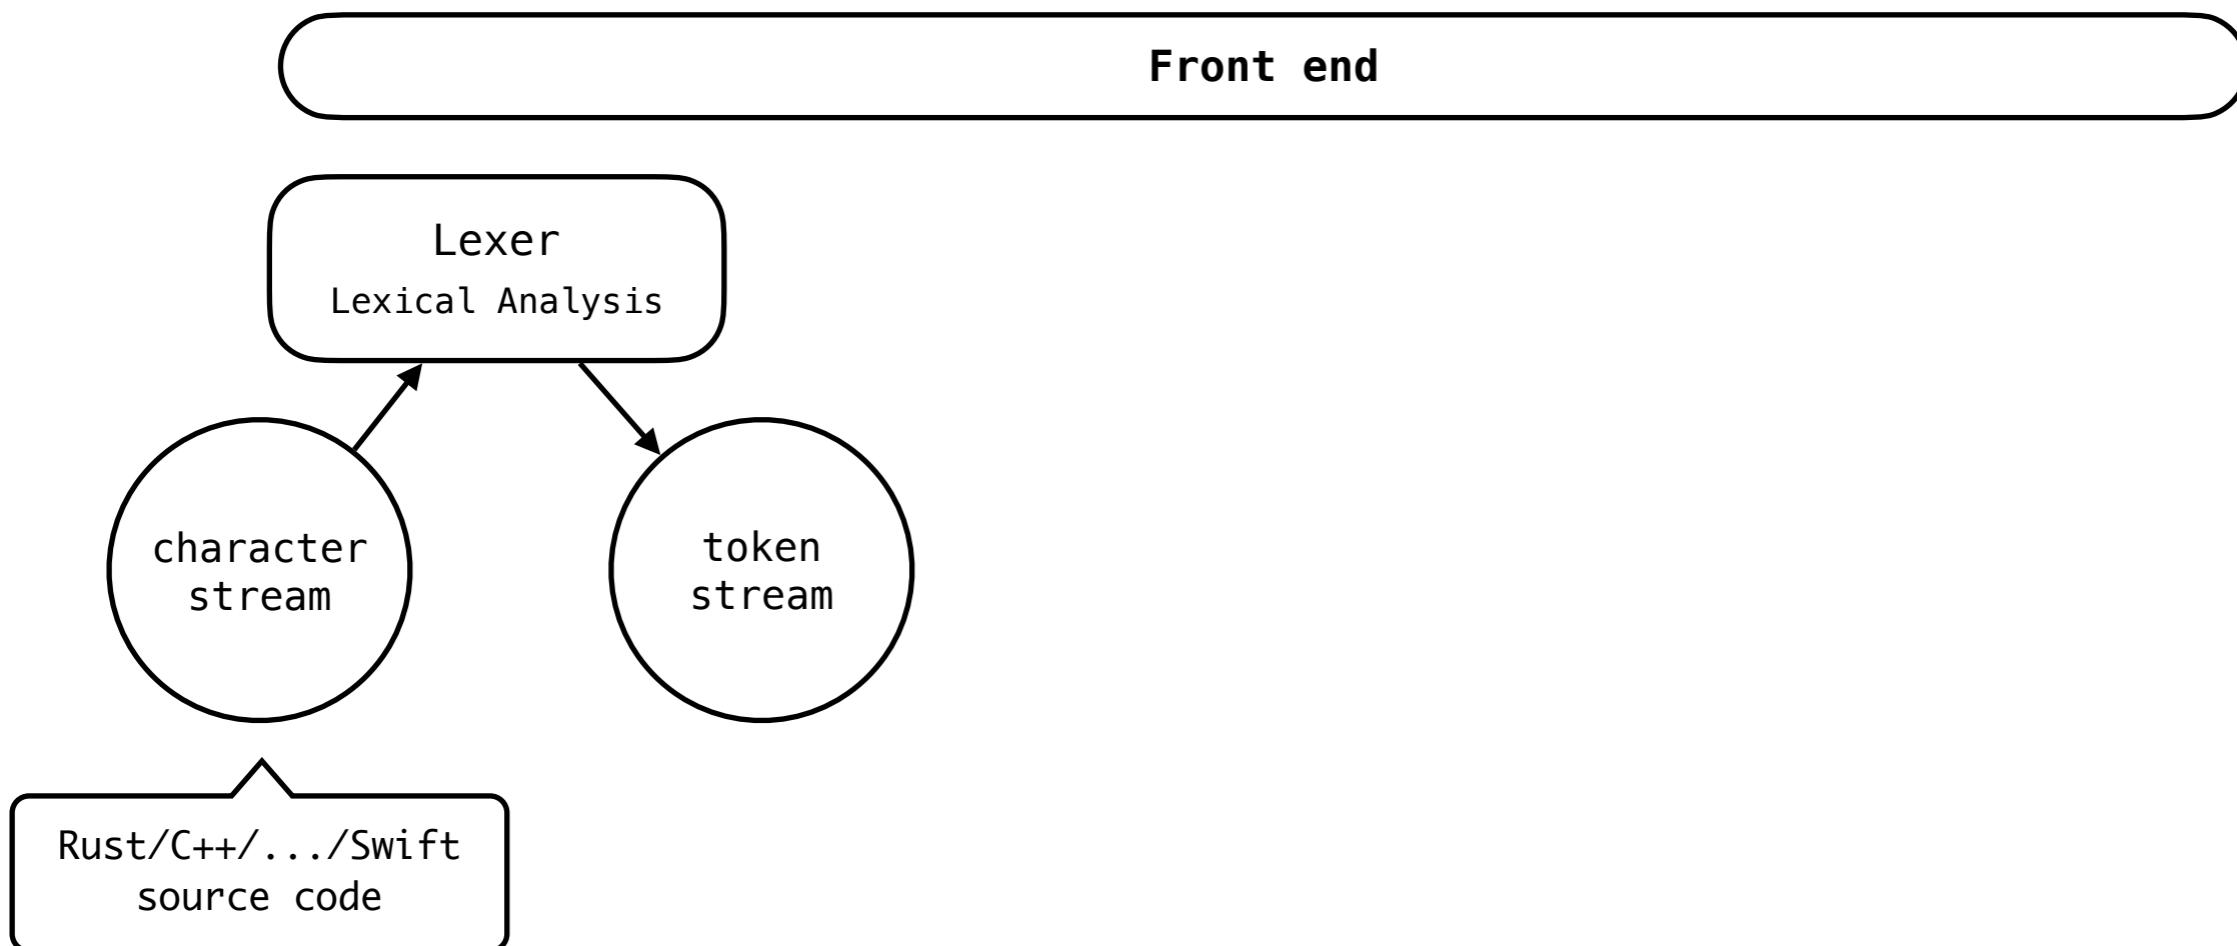
Compiler Front-end

Front end



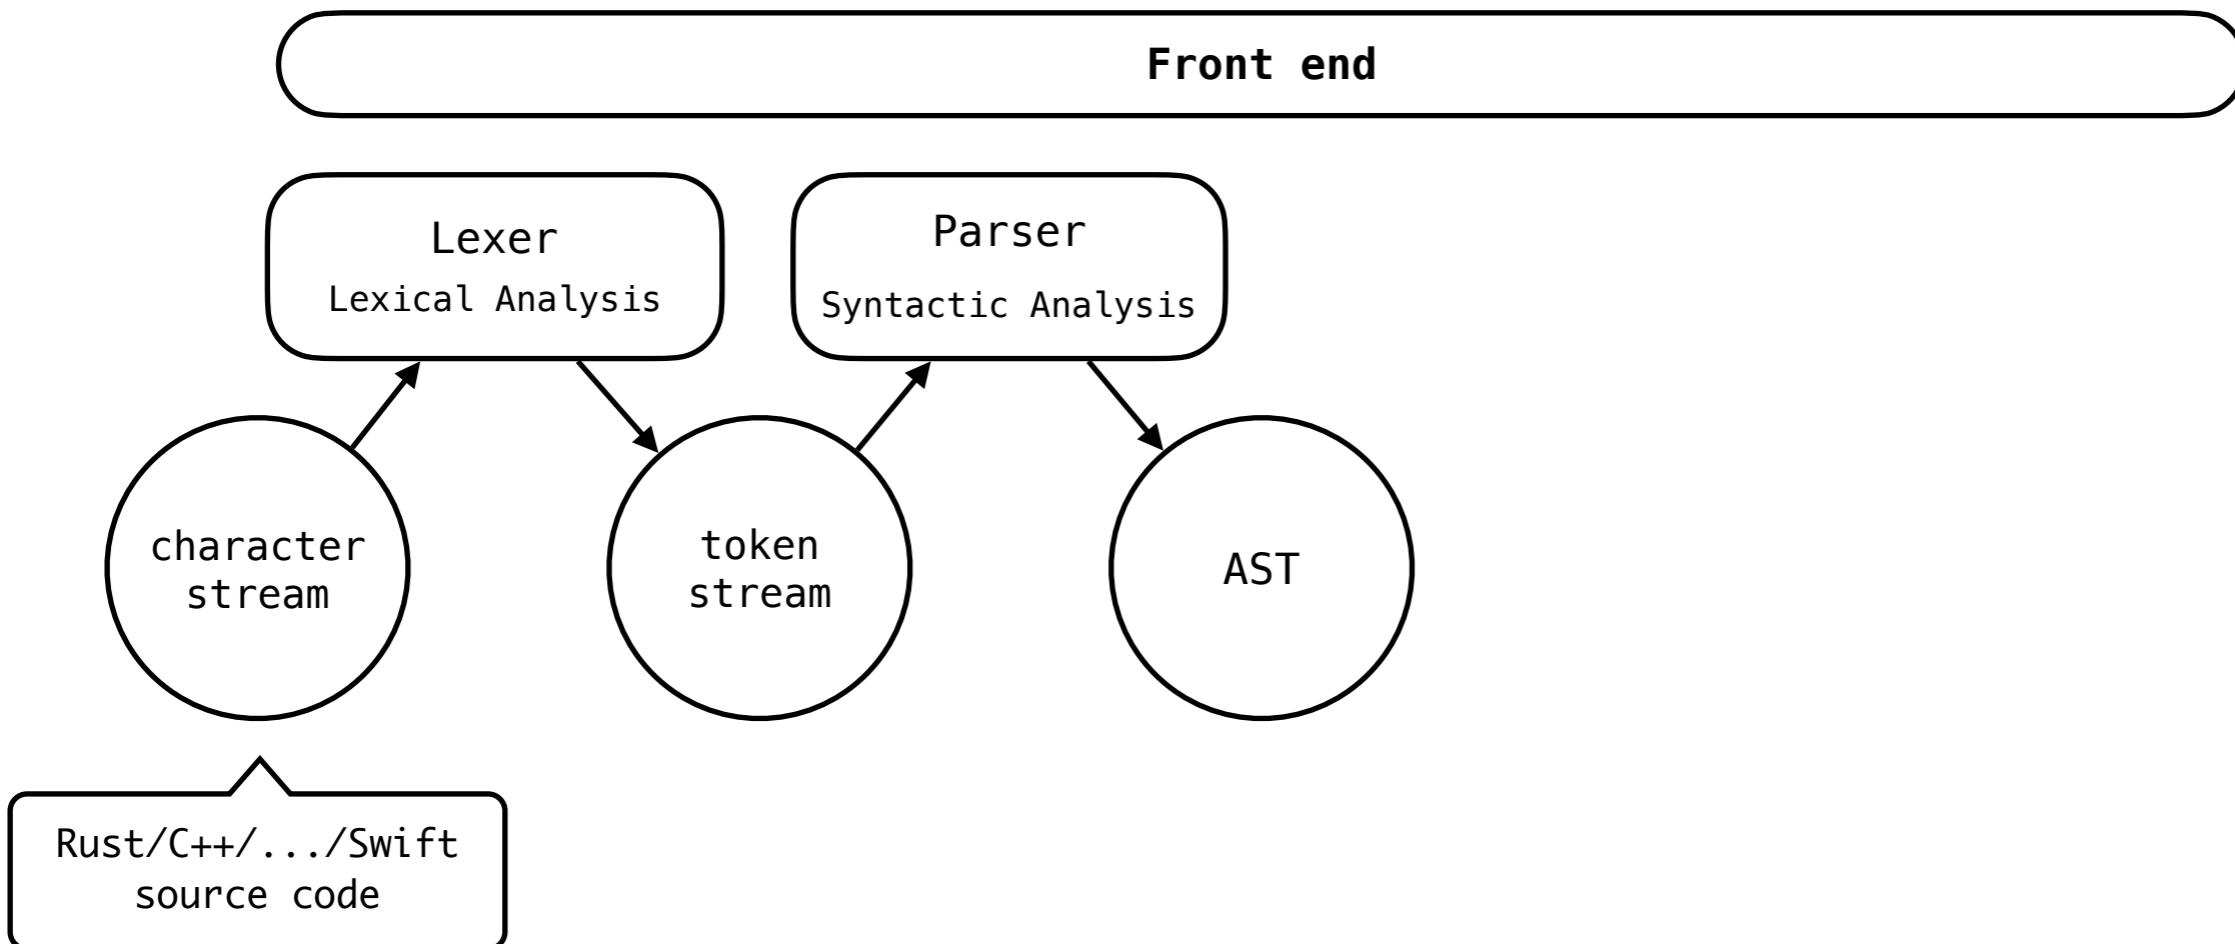
Compiler Architecture

Compiler Front-end



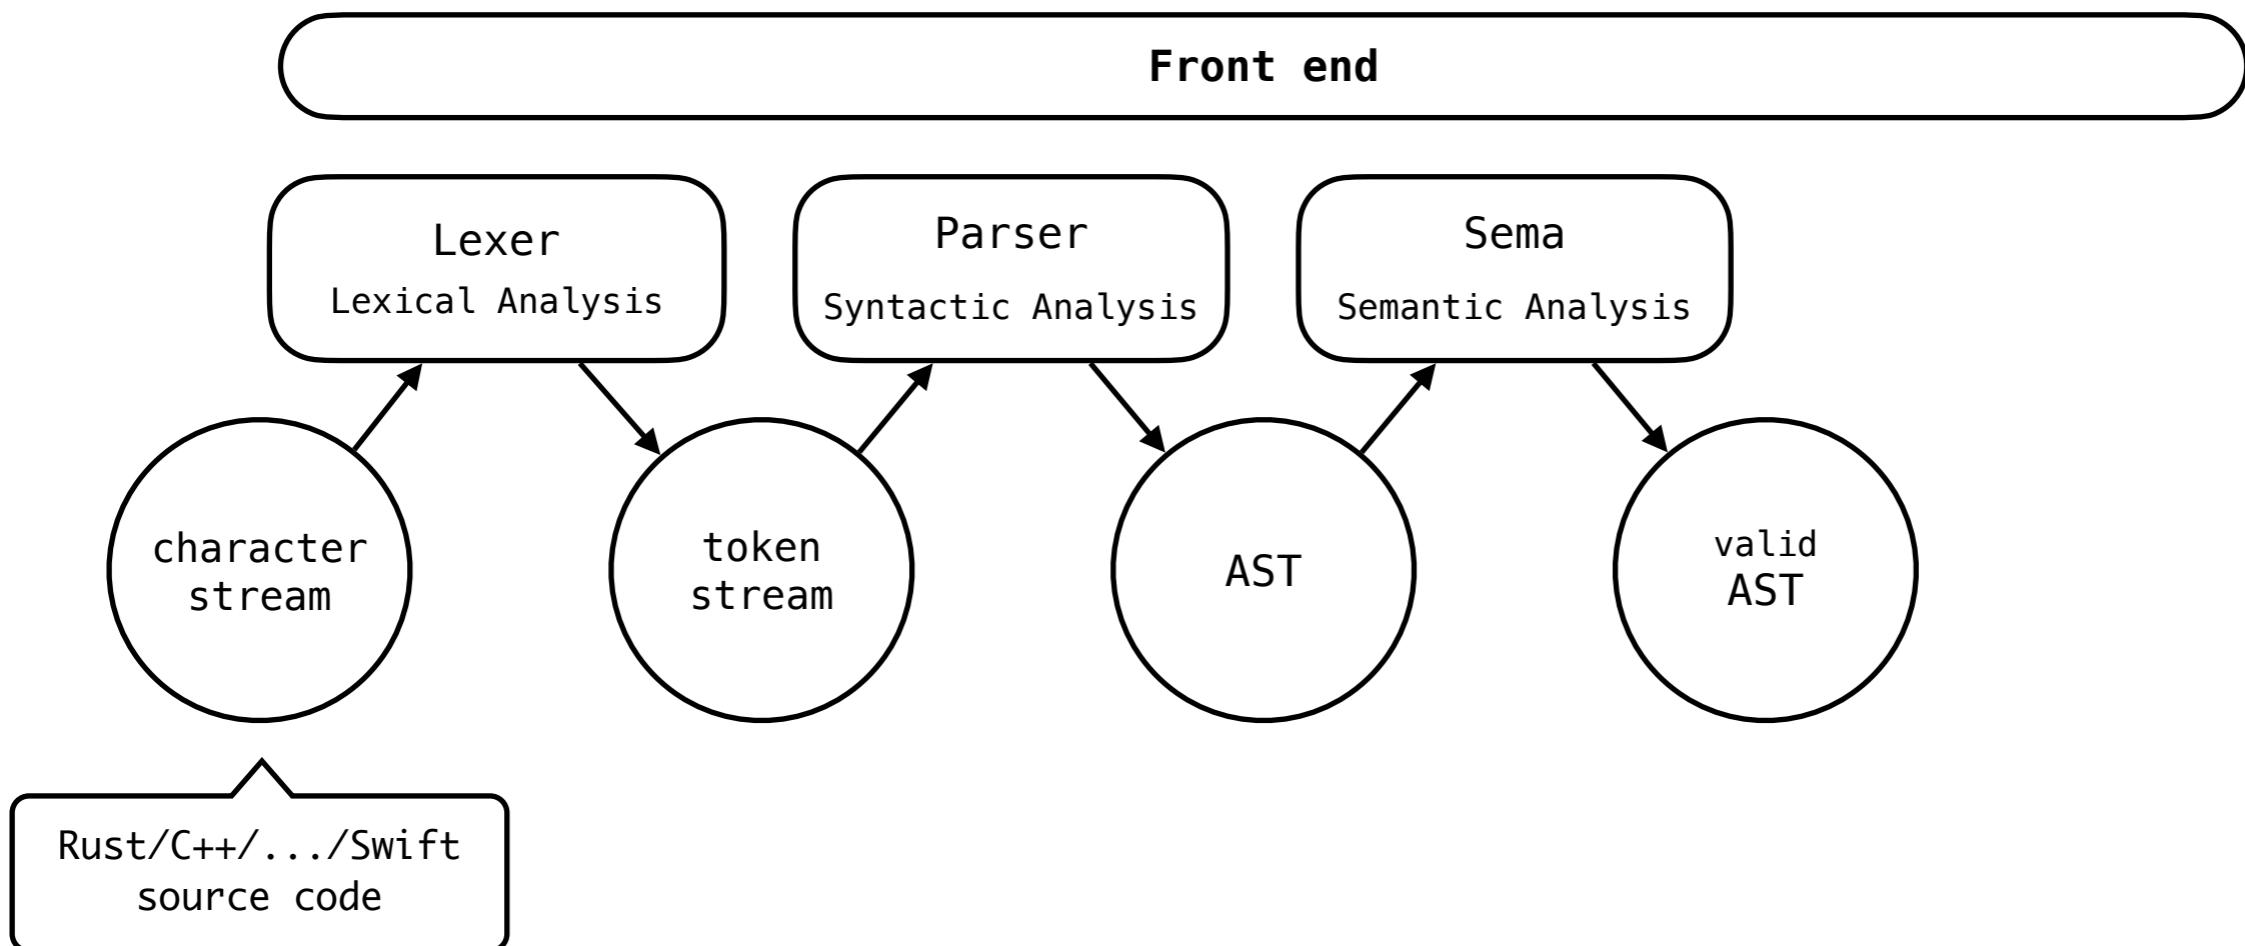
Compiler Architecture

Compiler Front-end



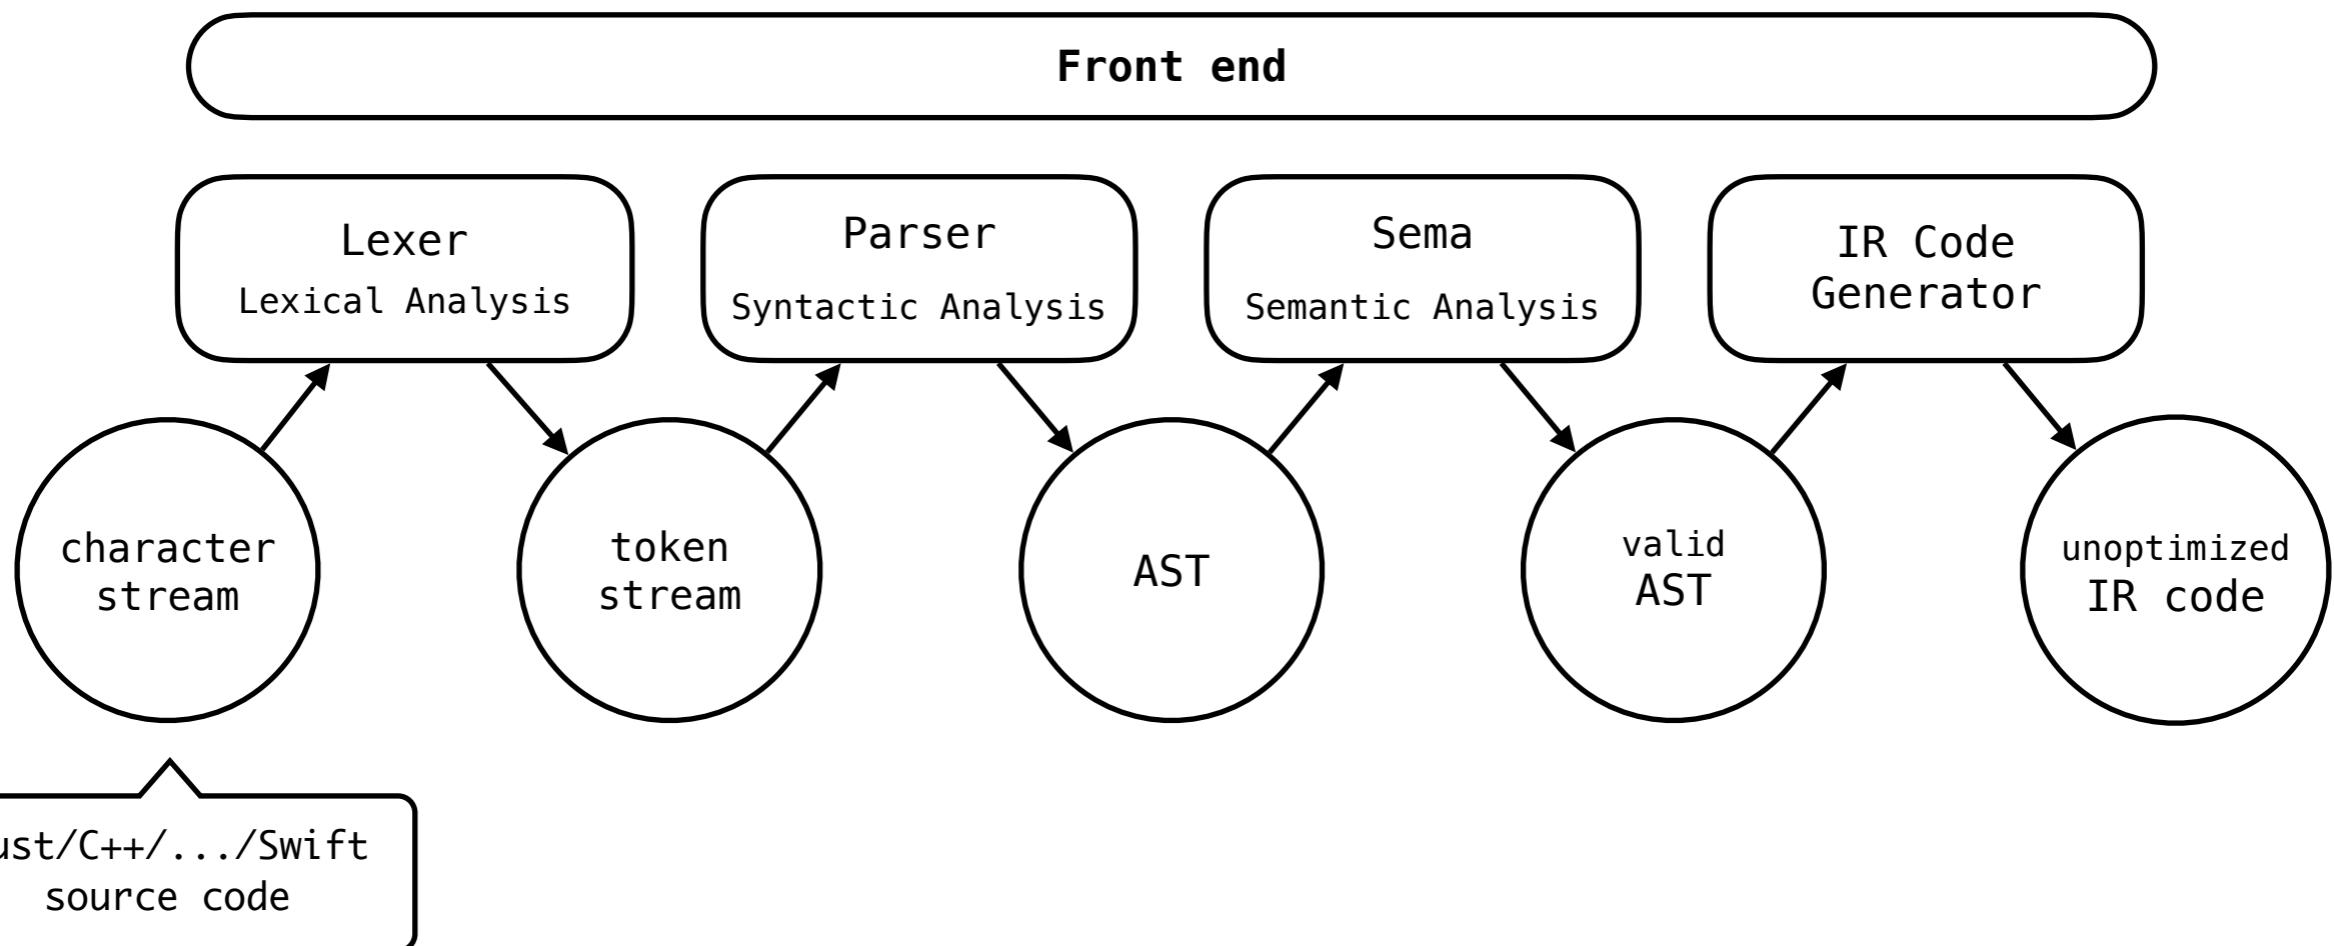
Compiler Architecture

Compiler Front-end



Compiler Architecture

Compiler Front-end



Compiler Architecture

Token

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Compiler Architecture

Token

```
$ clang++ -std=c++17 -Xclang -dump-tokens \ -c //physics_simulation/source/vec3_length.cxx
```

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Compiler Architecture

Token

```
$ clang++ -std=c++17 -Xclang -dump-tokens \
-c //physics_simulation/source/vec3_length.cxx
```

Token Kind	Token Value	Token Kind	Token Value	Token Kind	Token Value
namespace	'namespace'	amp	'&'	period	'. '
identifier	'phy_sim'	identifier	'v'	identifier	'y'
l_brace	'{'	r_paren	')'	star	'*'
struct	'struct'	arrow	'->'	identifier	'v'
identifier	'vec3'	float	'float'	period	'. '
l_brace	'{'	l_brace	'{'	identifier	'y'
float	'float'	return	'return'	plus	'+'
identifier	'x'	identifier	'std'	identifier	'v'
comma	', '	coloncolon	::'	period	'. '
identifier	'y'	identifier	'sqrt'	identifier	'z'
comma	', '	l_paren	'('	star	'*'
identifier	'z'	identifier	'v'	identifier	'v'
semi	'; '	period	'. '	period	'. '
r_brace	'}'	identifier	'x'	identifier	'z'
semi	'; '	star	'*'	r_paren	')'
auto	'auto'	identifier	'v'	semi	';'
identifier	'length'	period	'. '	r_brace	'}'
l_paren	'('	identifier	'x'	r_brace	'}'
const	'const'	plus	'+'		
identifier	'vec3'	identifier	'v'		

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Compiler Architecture

Token

```
$ clang++ -std=c++17 -Xclang -dump-tokens \
-c //physics_simulation/source/vec3_length.cxx
```

Token Kind	Token Value	Token Kind	Token Value	Token Kind	Token Value
namespace	'namespace'	amp	'&'	period	'. '
identifier	'phy_sim'	identifier	'v'	identifier	'y'
l_brace	'{'	r_paren	')'	star	'*'
struct	'struct'	arrow	'->'	identifier	'v'
identifier	'vec3'	float	'float'	period	'. '
l_brace	'{'	l_brace	'{'	identifier	'y'
float	'float'	return	'return'	plus	'+'
identifier	'x'	identifier	'std'	identifier	'v'
comma	', '	coloncolon	::'	period	'. '
identifier	'y'	identifier	'sqrt'	identifier	'z'
comma	', '	l_paren	'('	star	'*'
identifier	'z'	identifier	'v'	identifier	'v'
semi	'; '	period	'. '	period	'. '
r_brace	'}'	identifier	'x'	identifier	'z'
semi	'; '	star	'*'	r_paren	')'
auto	'auto'	identifier	'v'	semi	'; '
identifier	'length'	period	'. '	r_brace	'}'
l_paren	'('	identifier	'x'	r_brace	'}'
const	'const'	plus	'+'		
identifier	'vec3'	identifier	'v'		

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Token is a sequence of characters, which has identified semantics.

Tokens represent keywords, operators, literals or identifiers.

Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```

Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```

TranslationUnitDecl

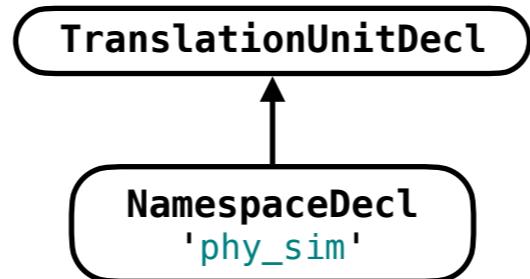
Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```



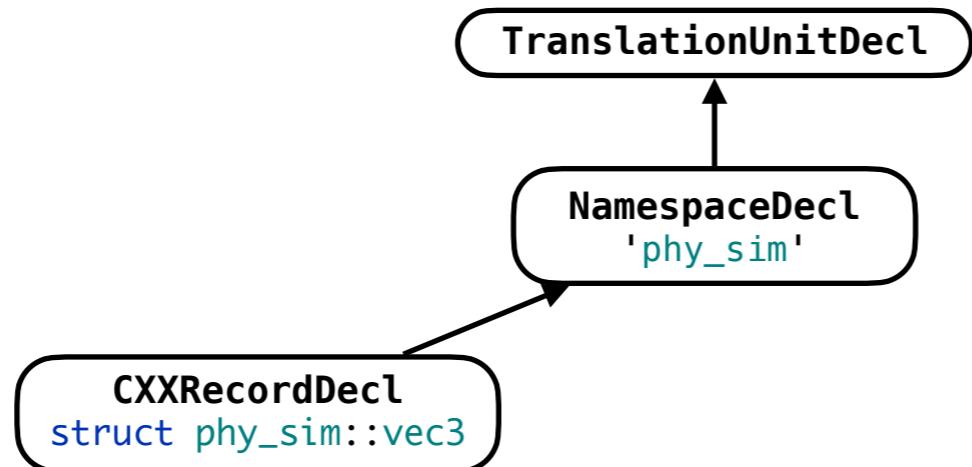
Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```



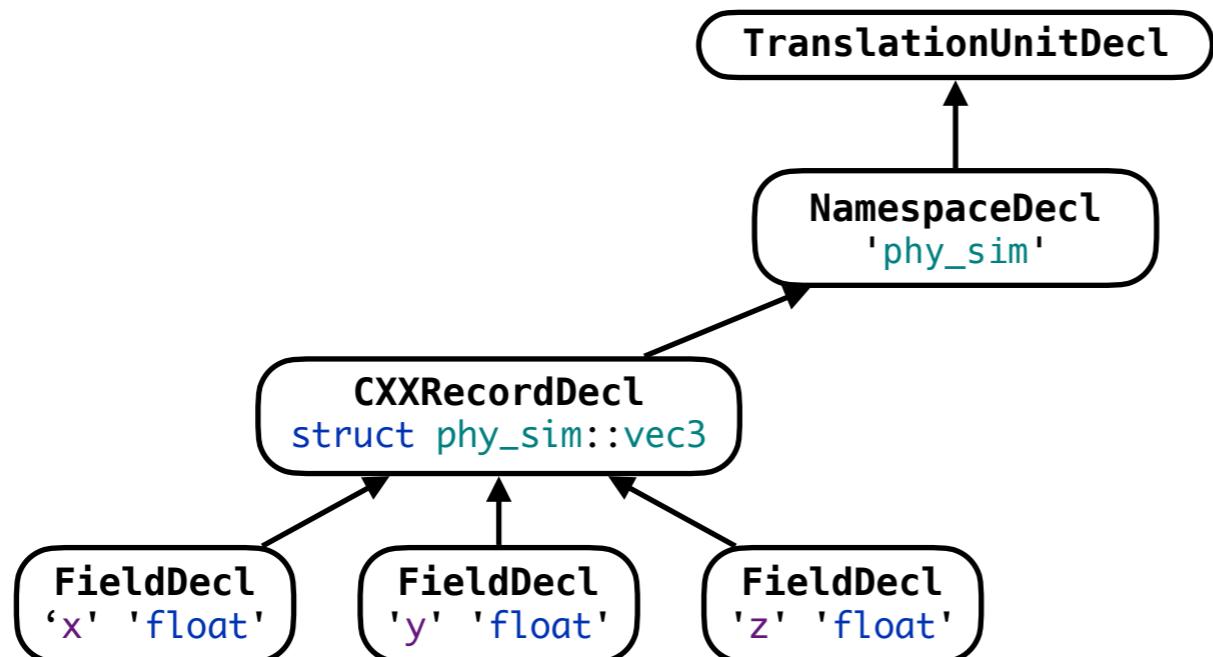
Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```



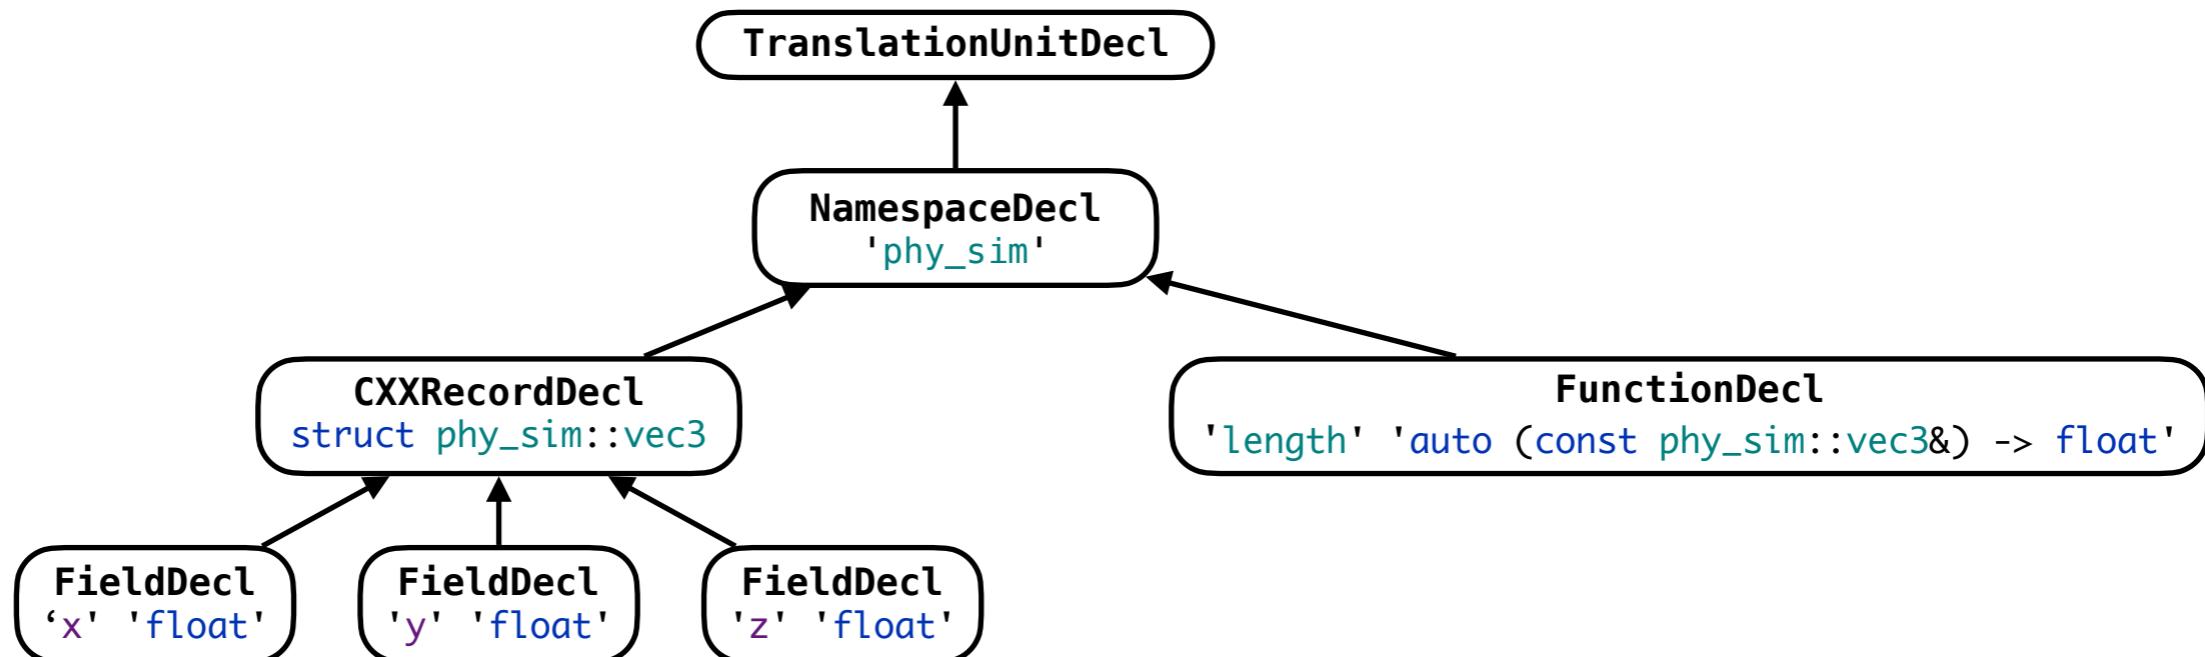
Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```



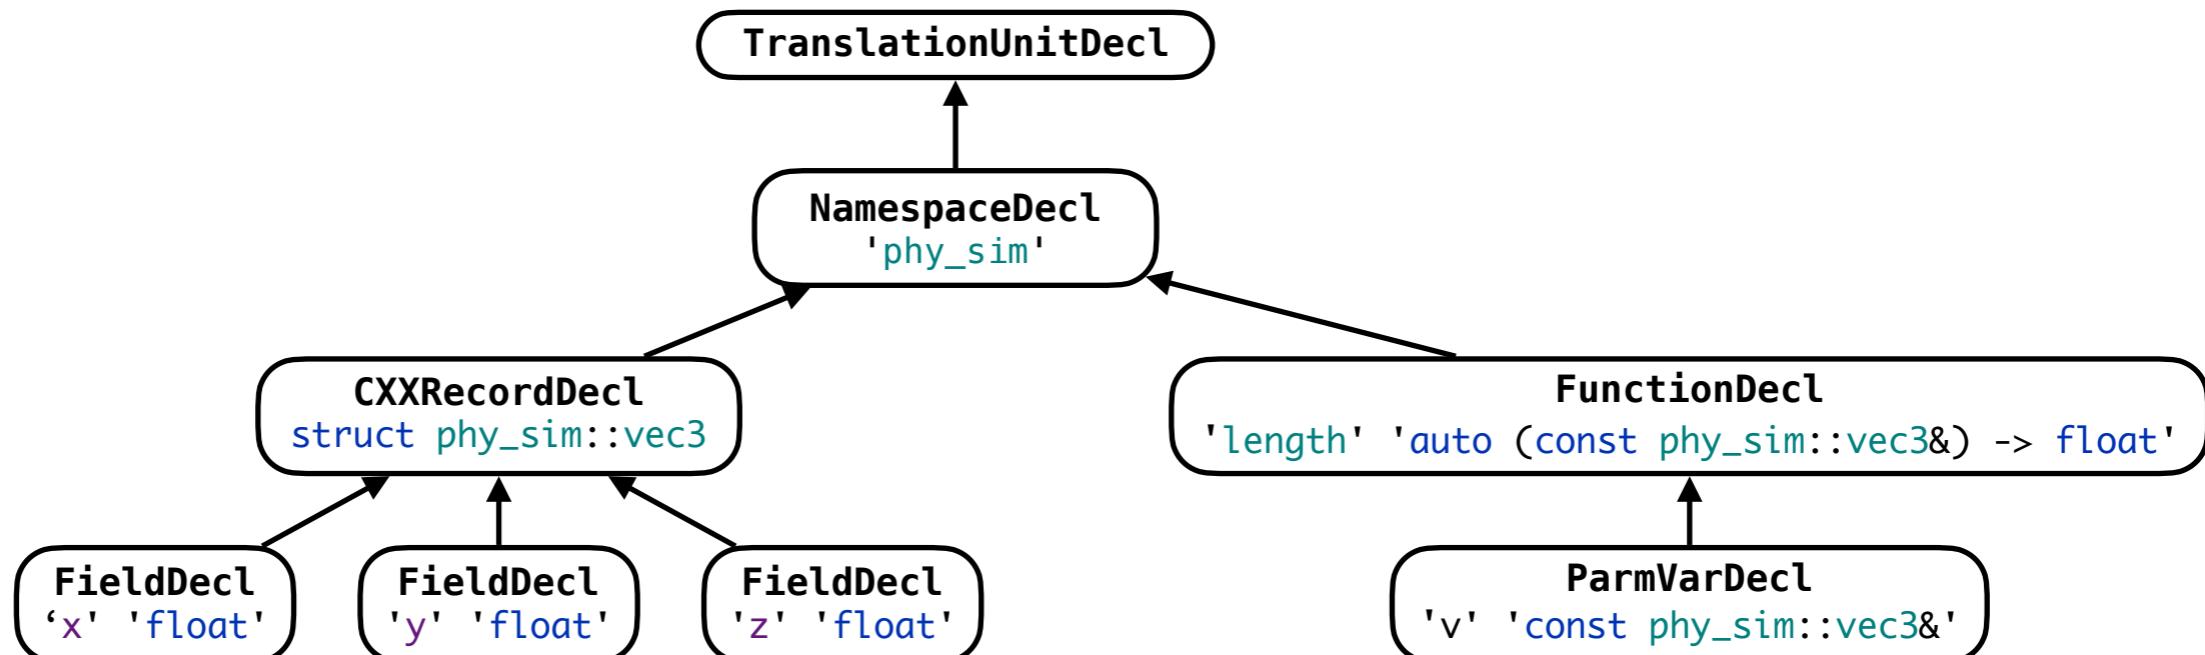
Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump      \
-c //physics_simulation/source/vec3_length.cxx
```



Compiler Architecture

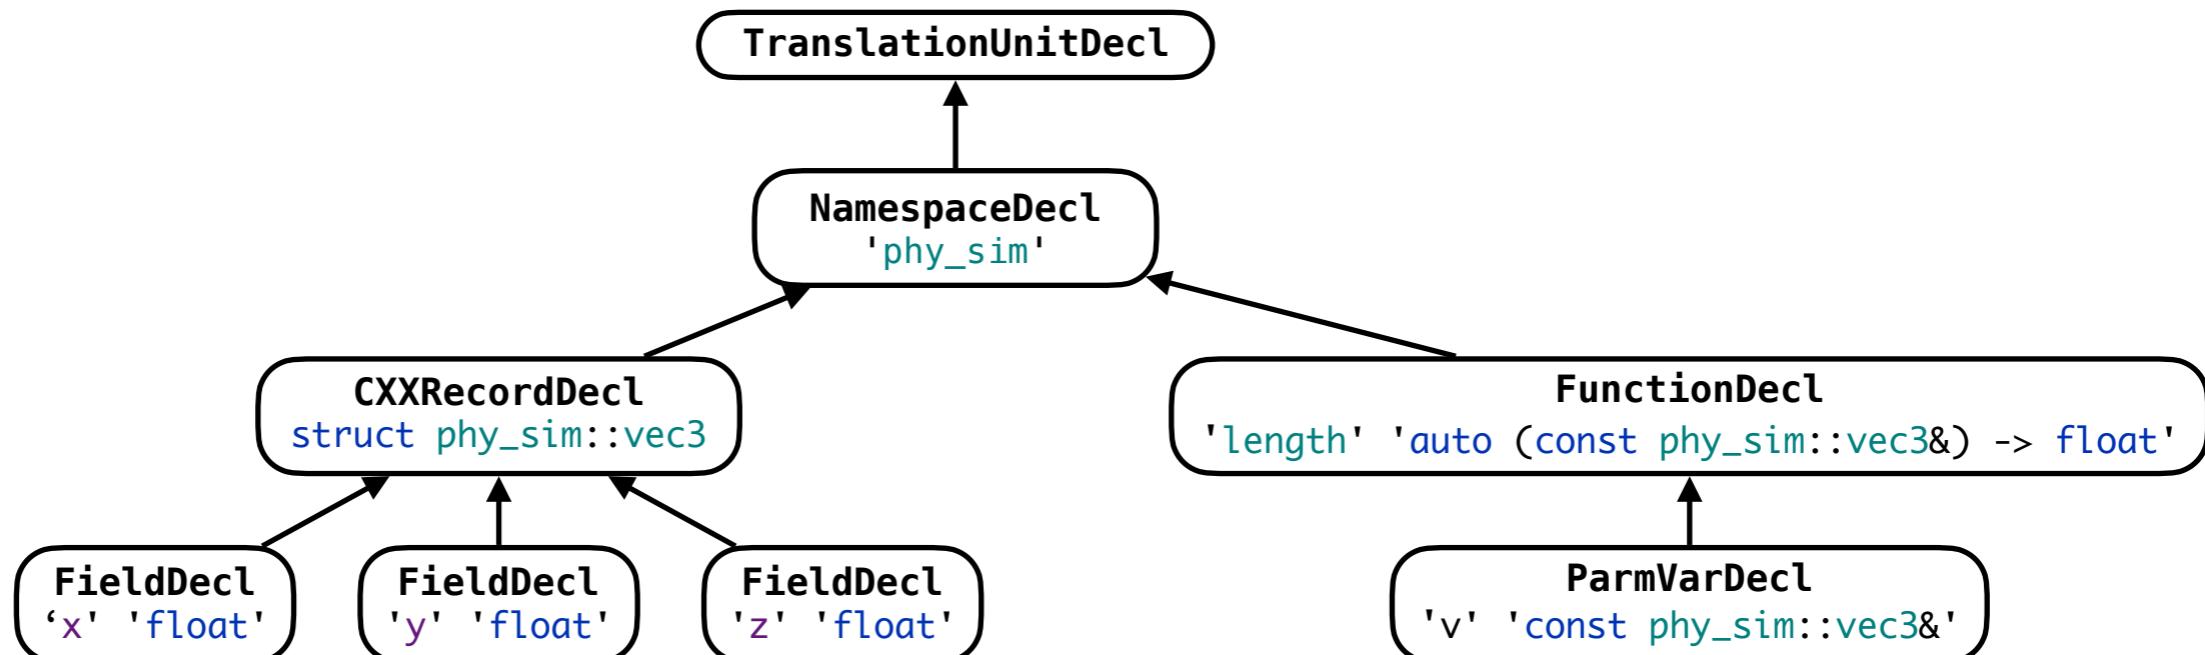
Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump \
-c //physics_simulation/source/vec3_length.cxx
```

Abstract Syntax Tree (AST) is concrete representation of C++ programs, in which their syntax & semantics is represented by a directed graph of C++ entities (vertices) and relations between them (edges).



Compiler Architecture

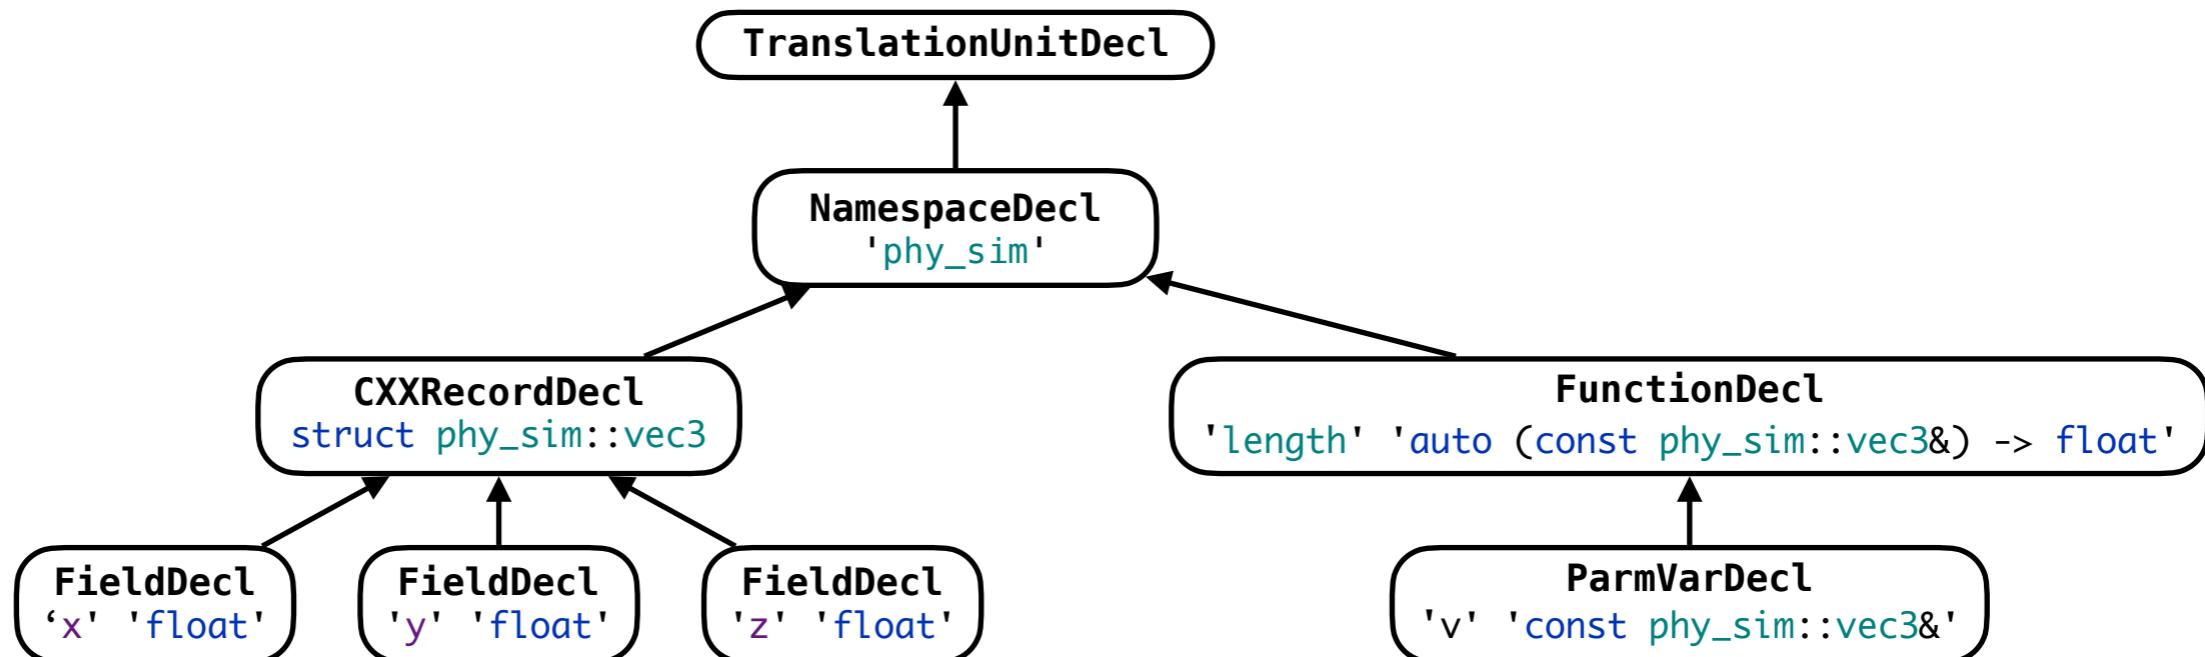
Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

```
$ clang++ -std=c++17 -Xclang -ast-dump \
-c //physics_simulation/source/vec3_length.cxx
```

Abstract Syntax Tree (AST) is concrete representation of C++ programs, in which their syntax & semantics is represented by a directed graph of C++ entities (vertices) and relations between them (edges).



AST is representation of C++ programs alternative to C++ source code, that is conversion of C++ source code to AST is reversible without losing semantics.

Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

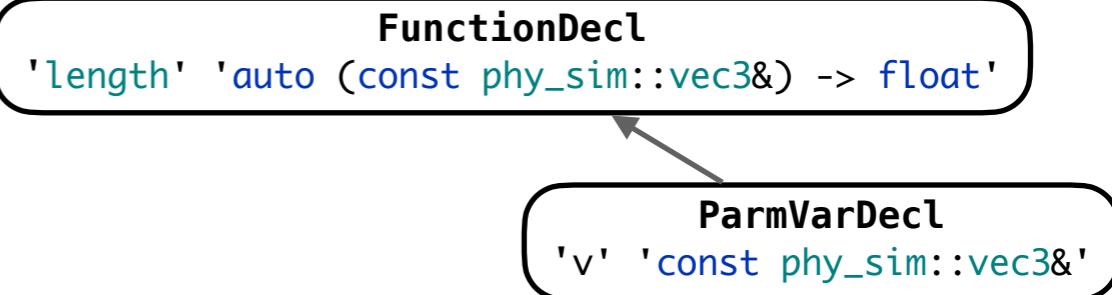
FunctionDecl
'length' 'auto (const phy_sim::vec3&) -> float'

Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

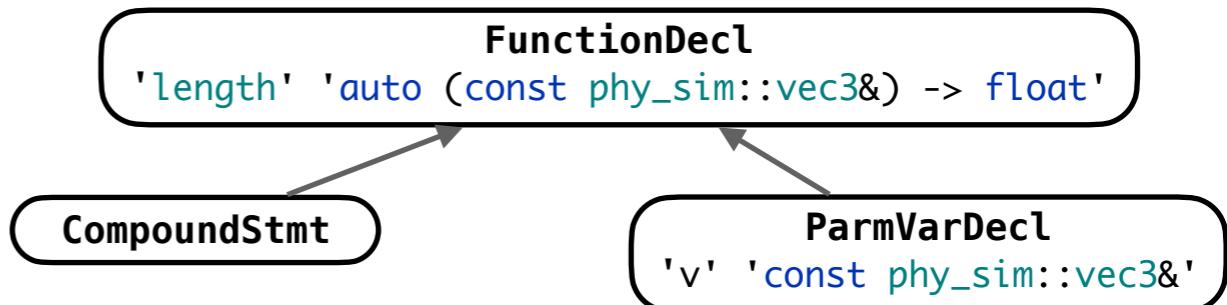


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

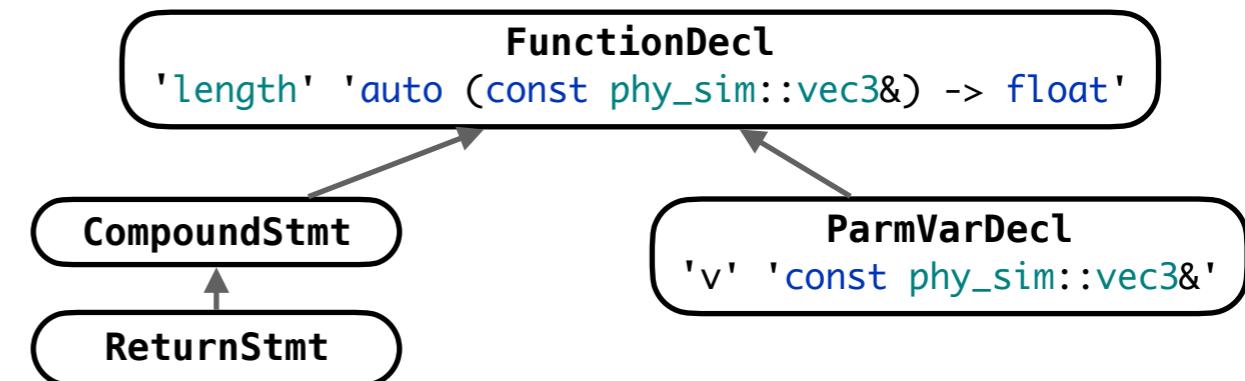


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

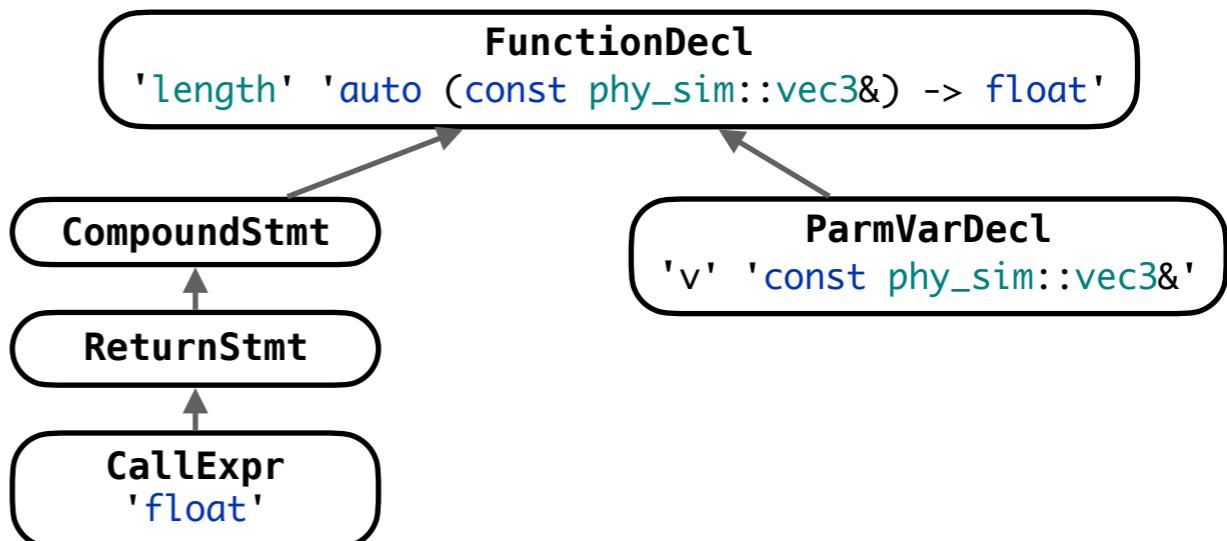


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

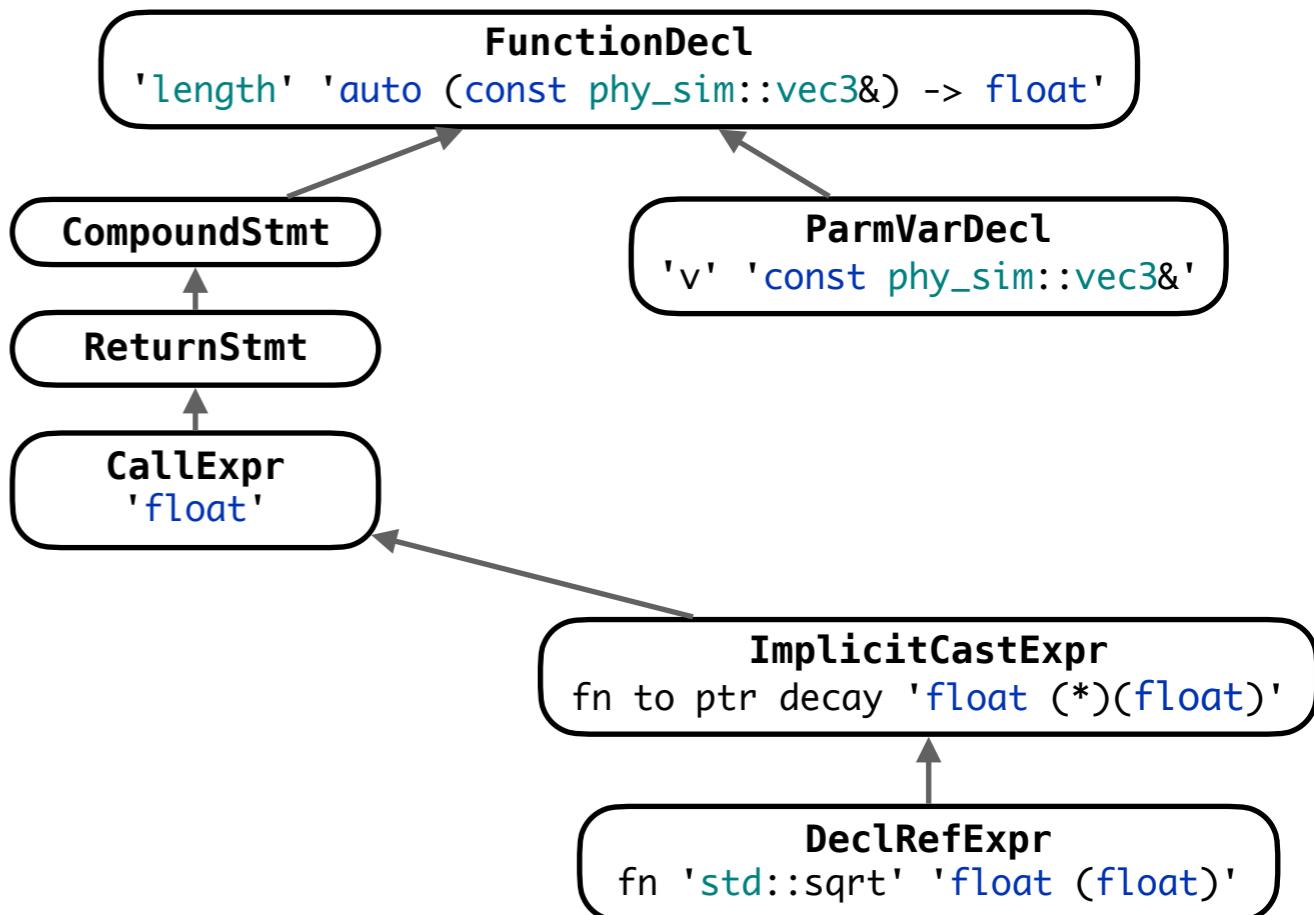


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

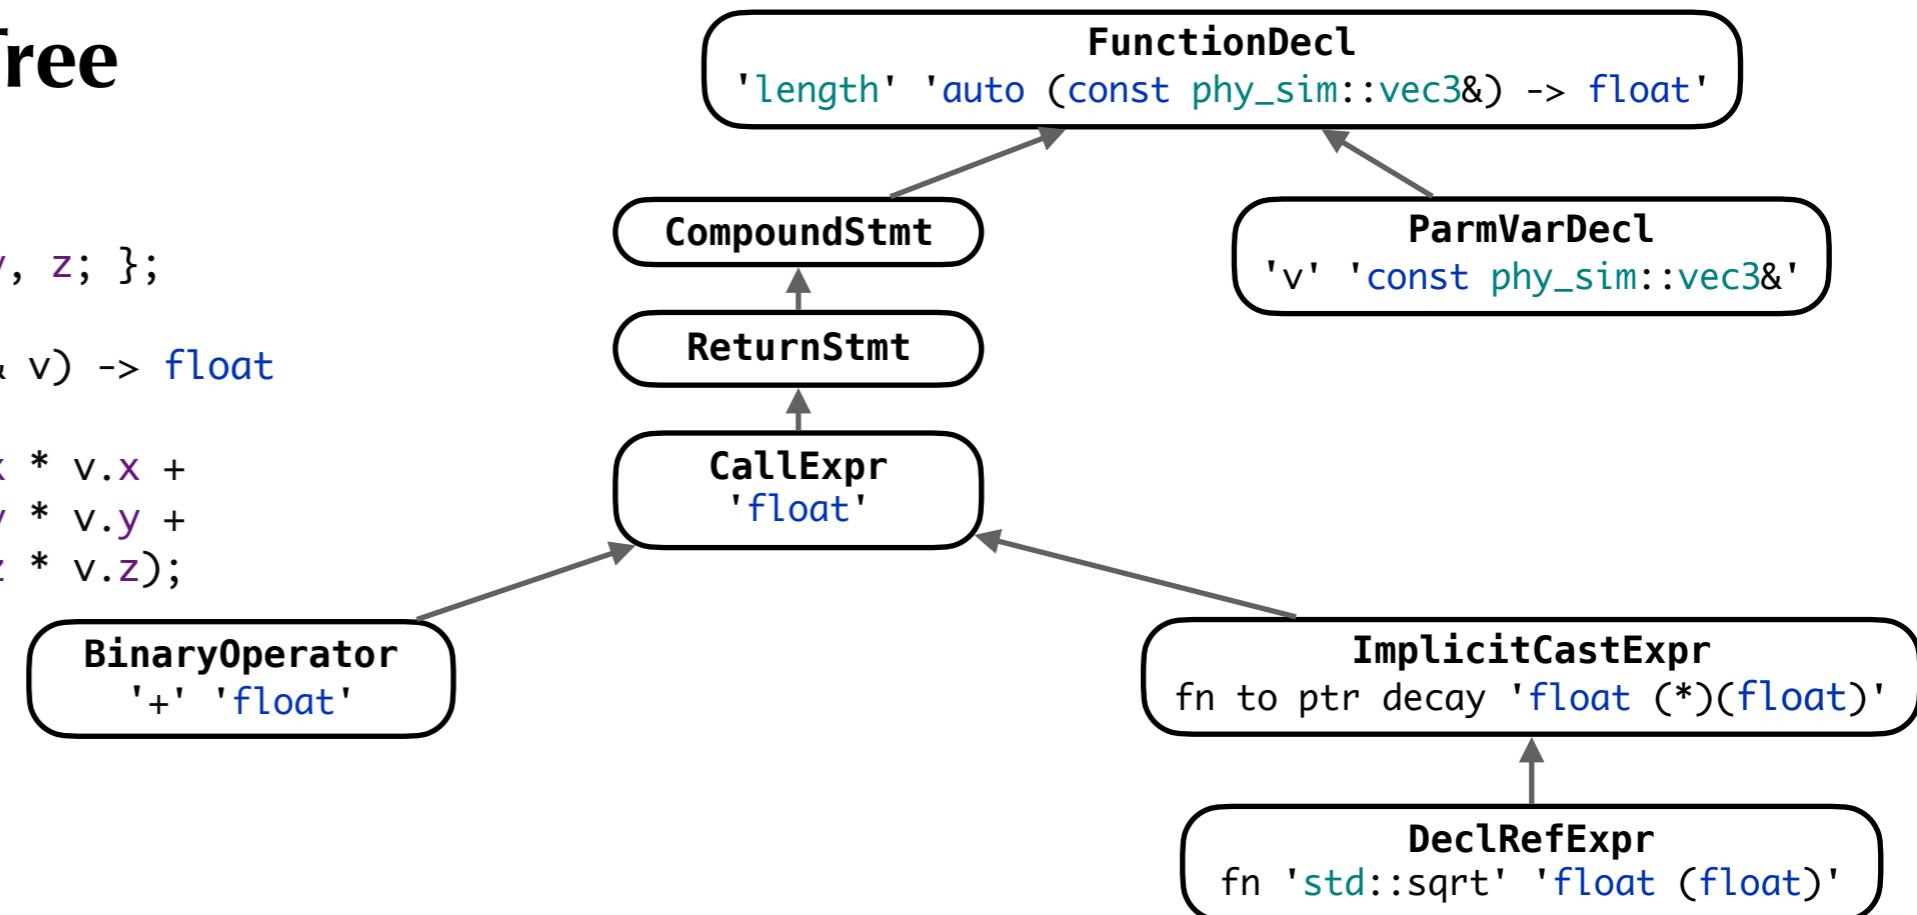


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

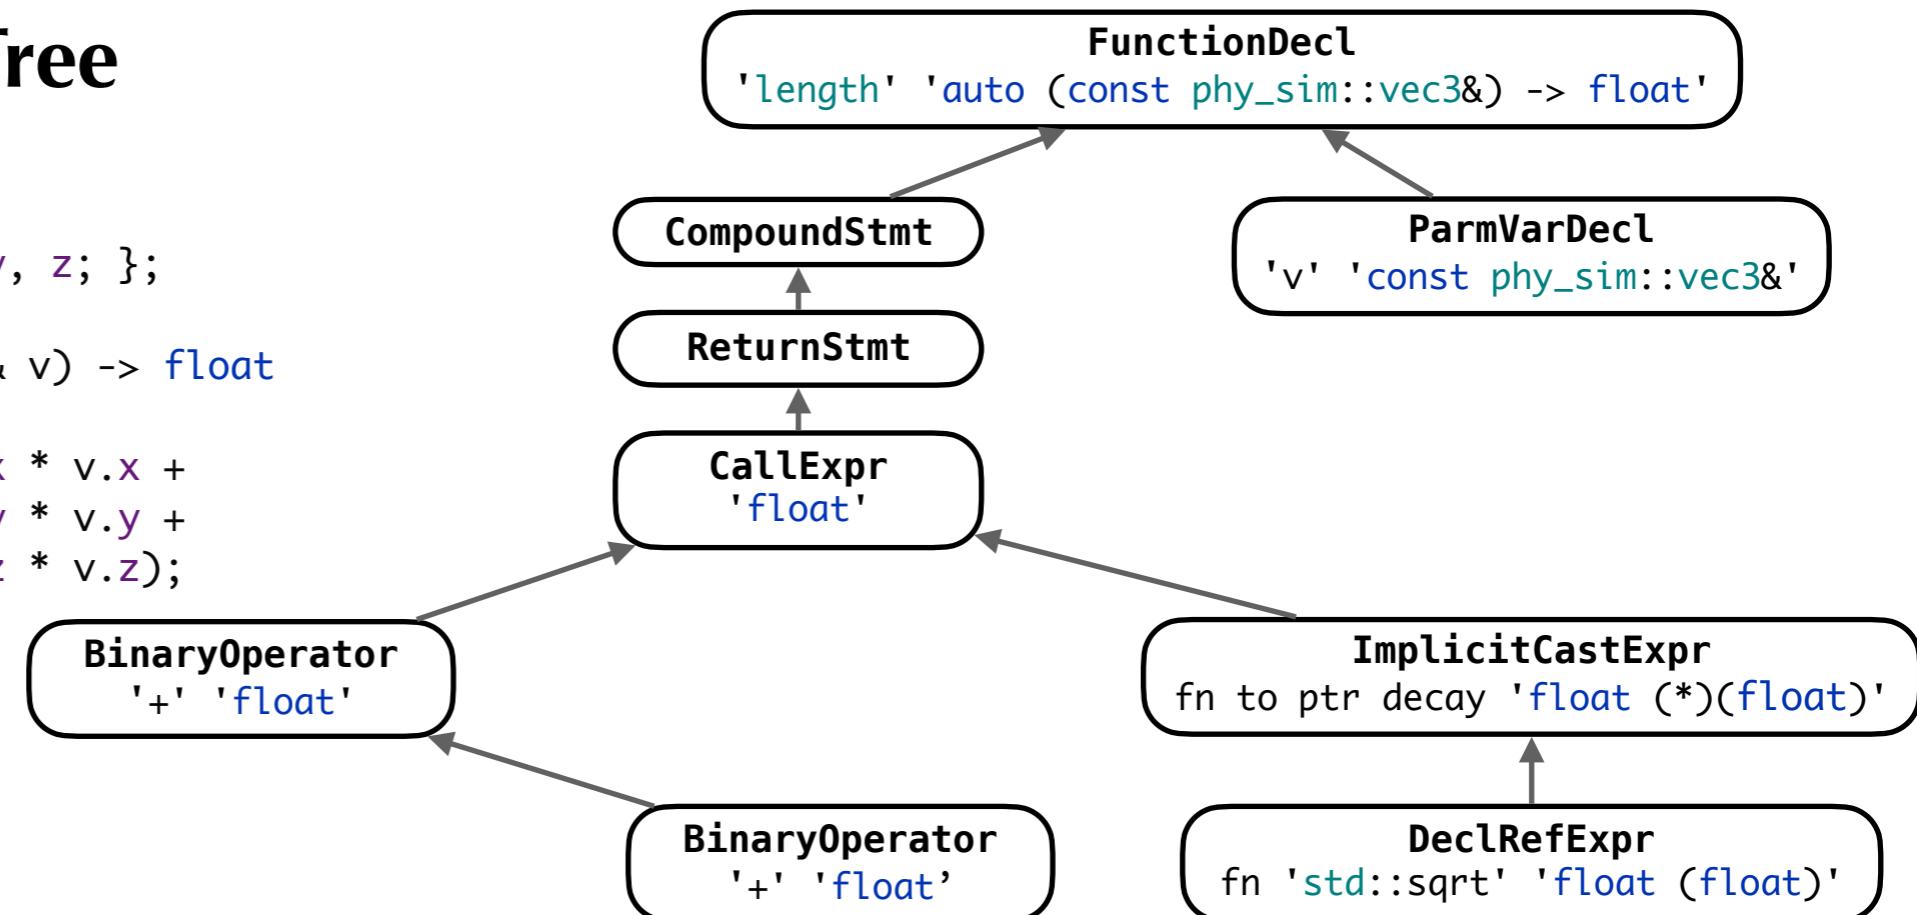


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

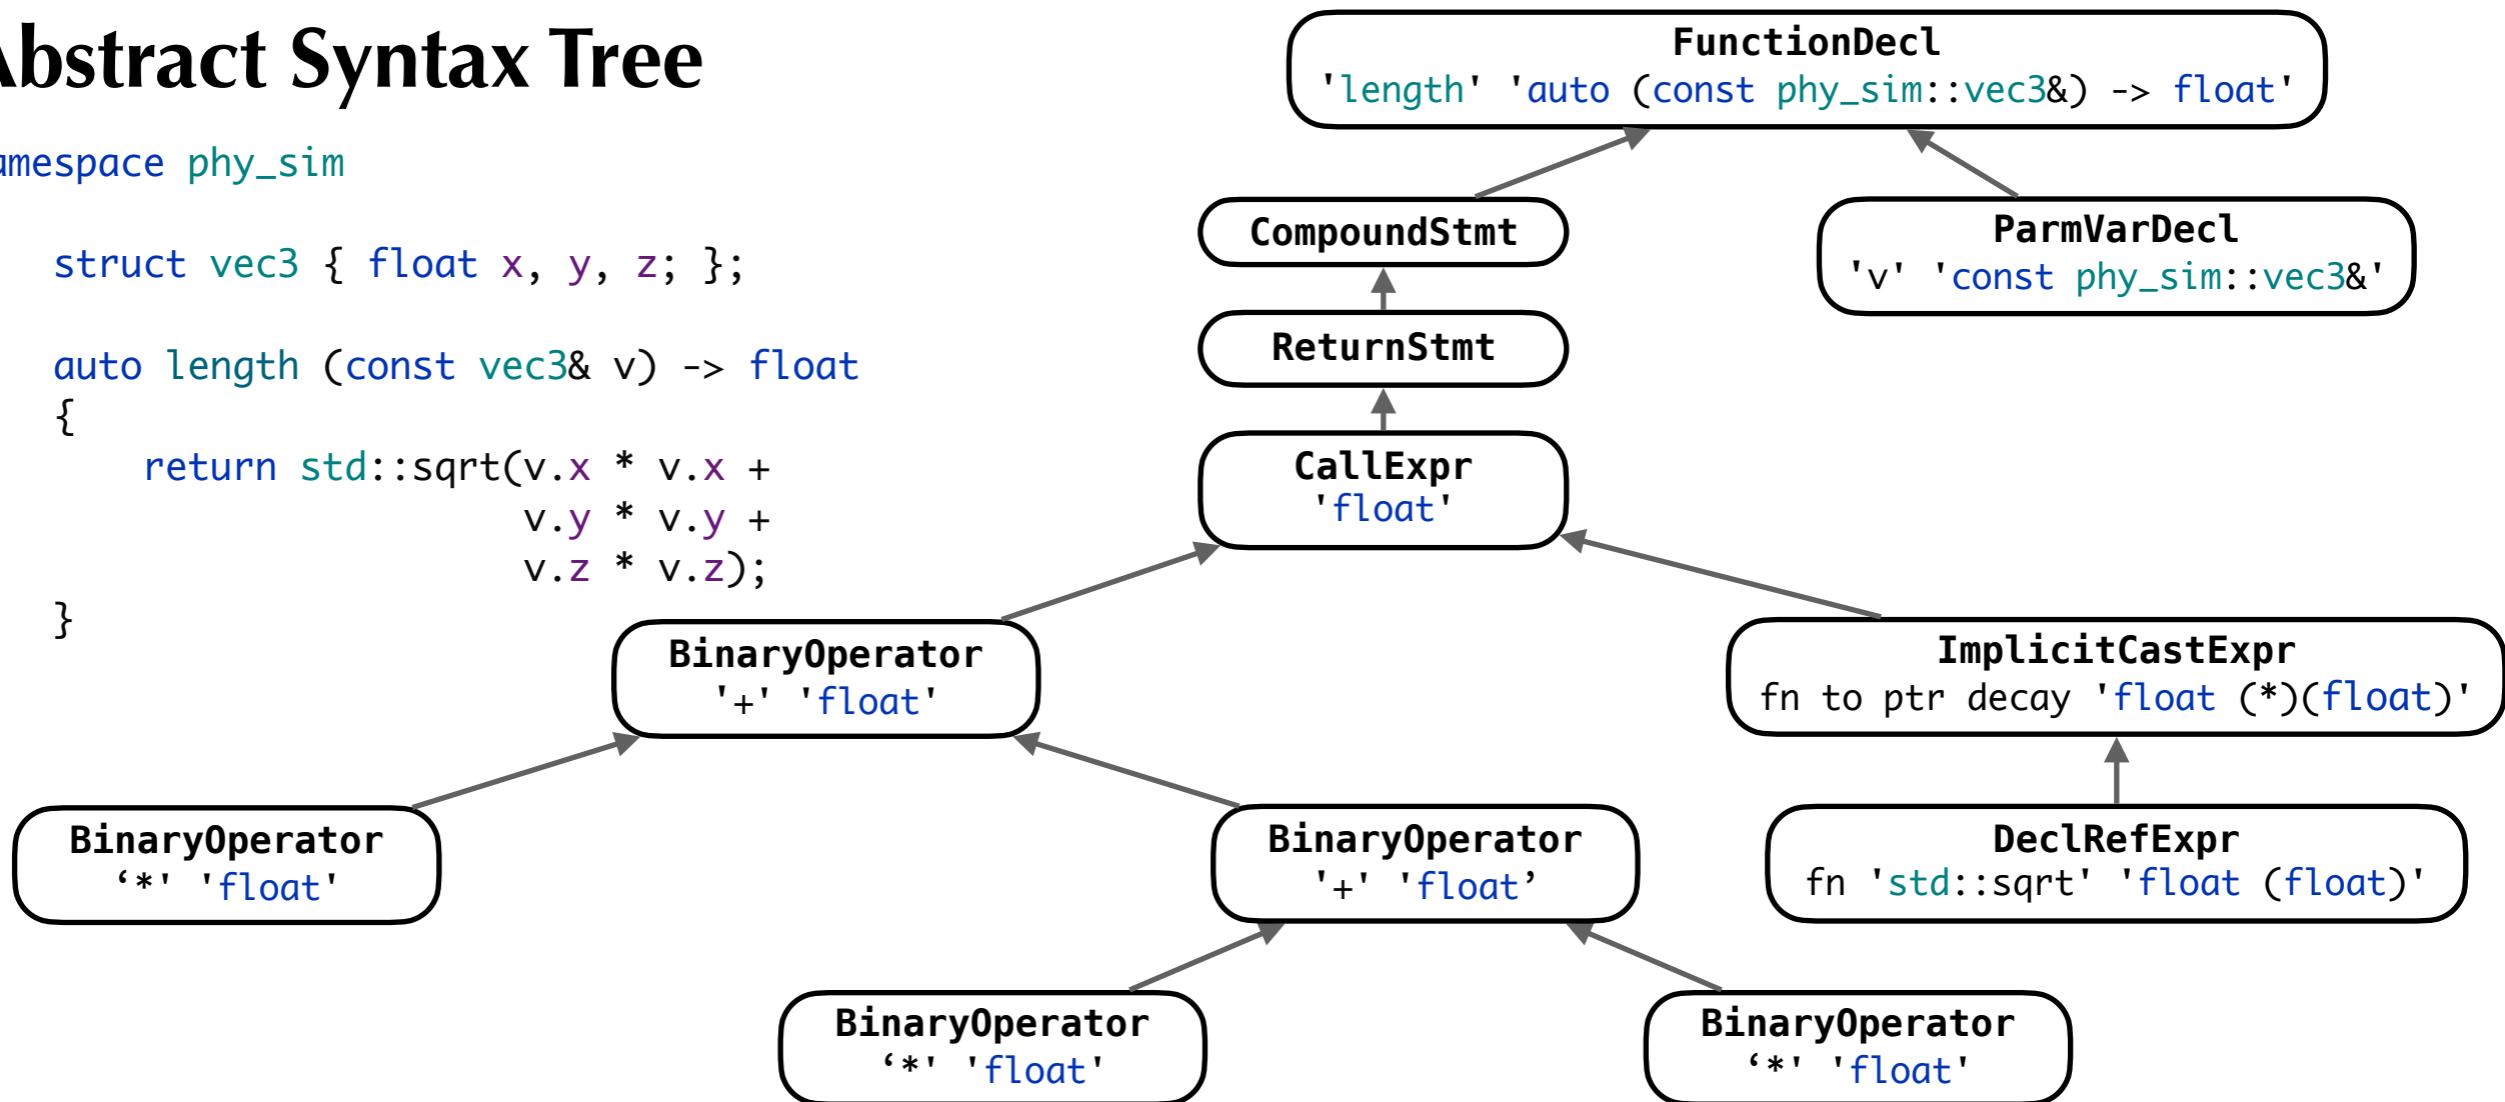


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

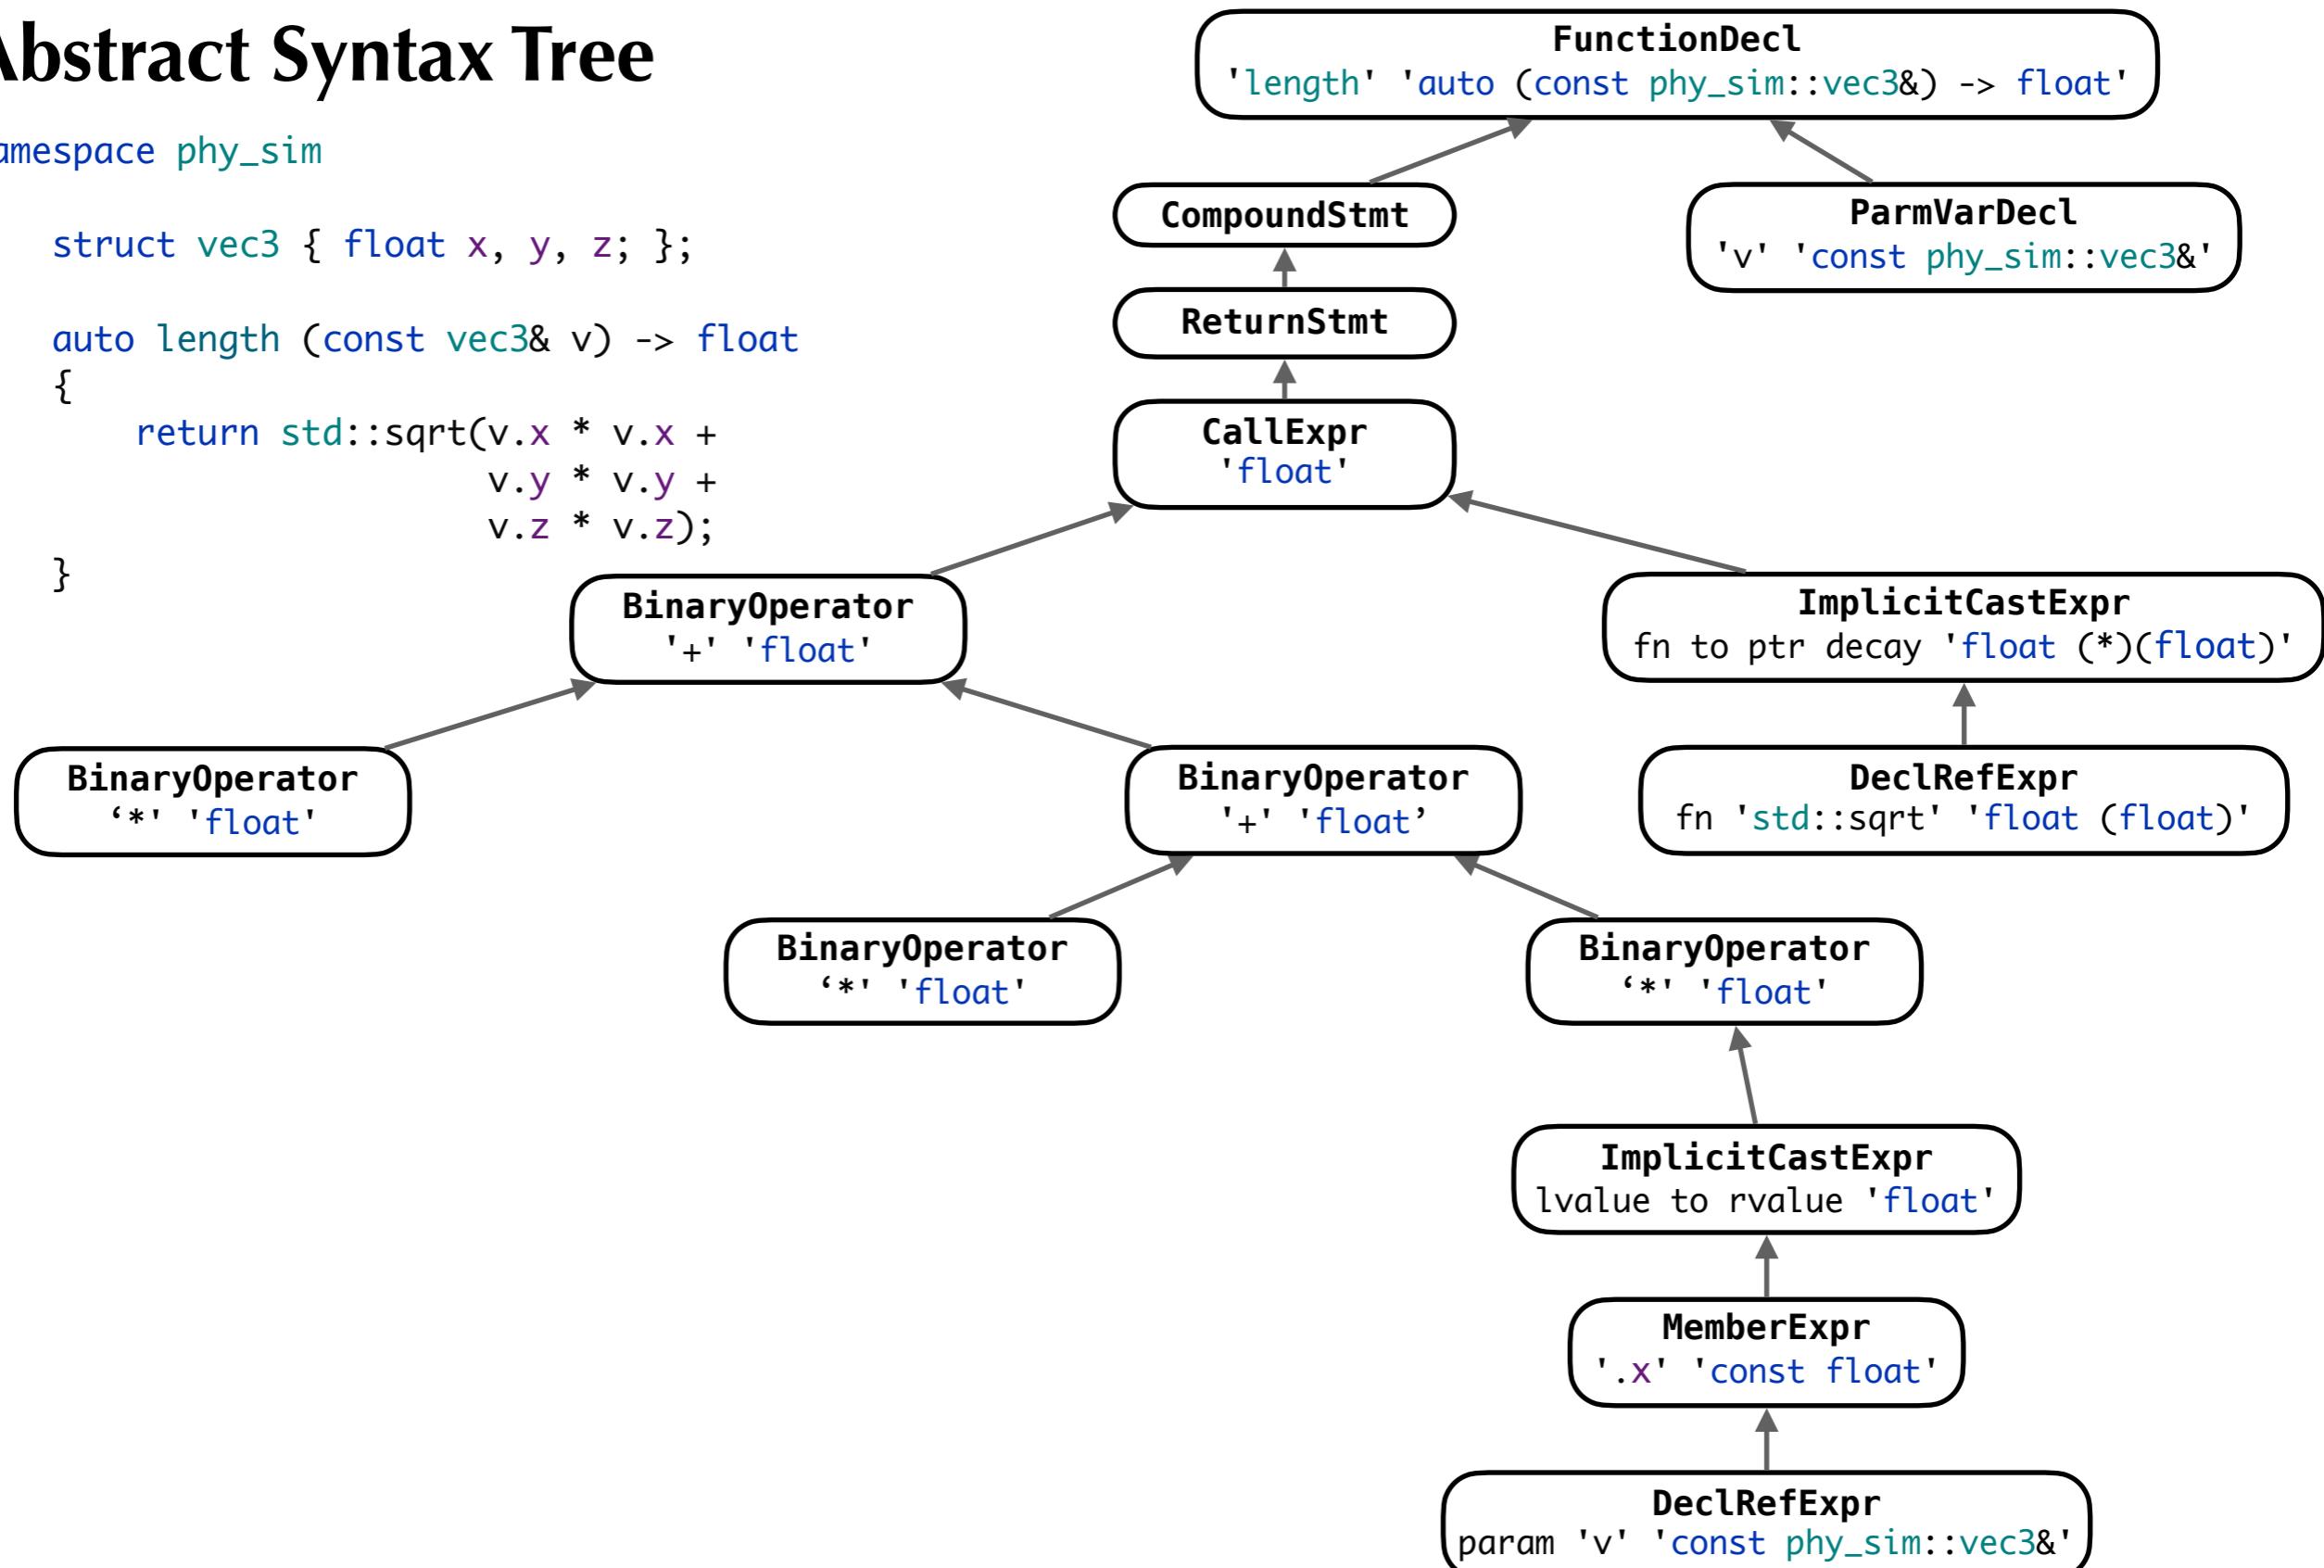


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

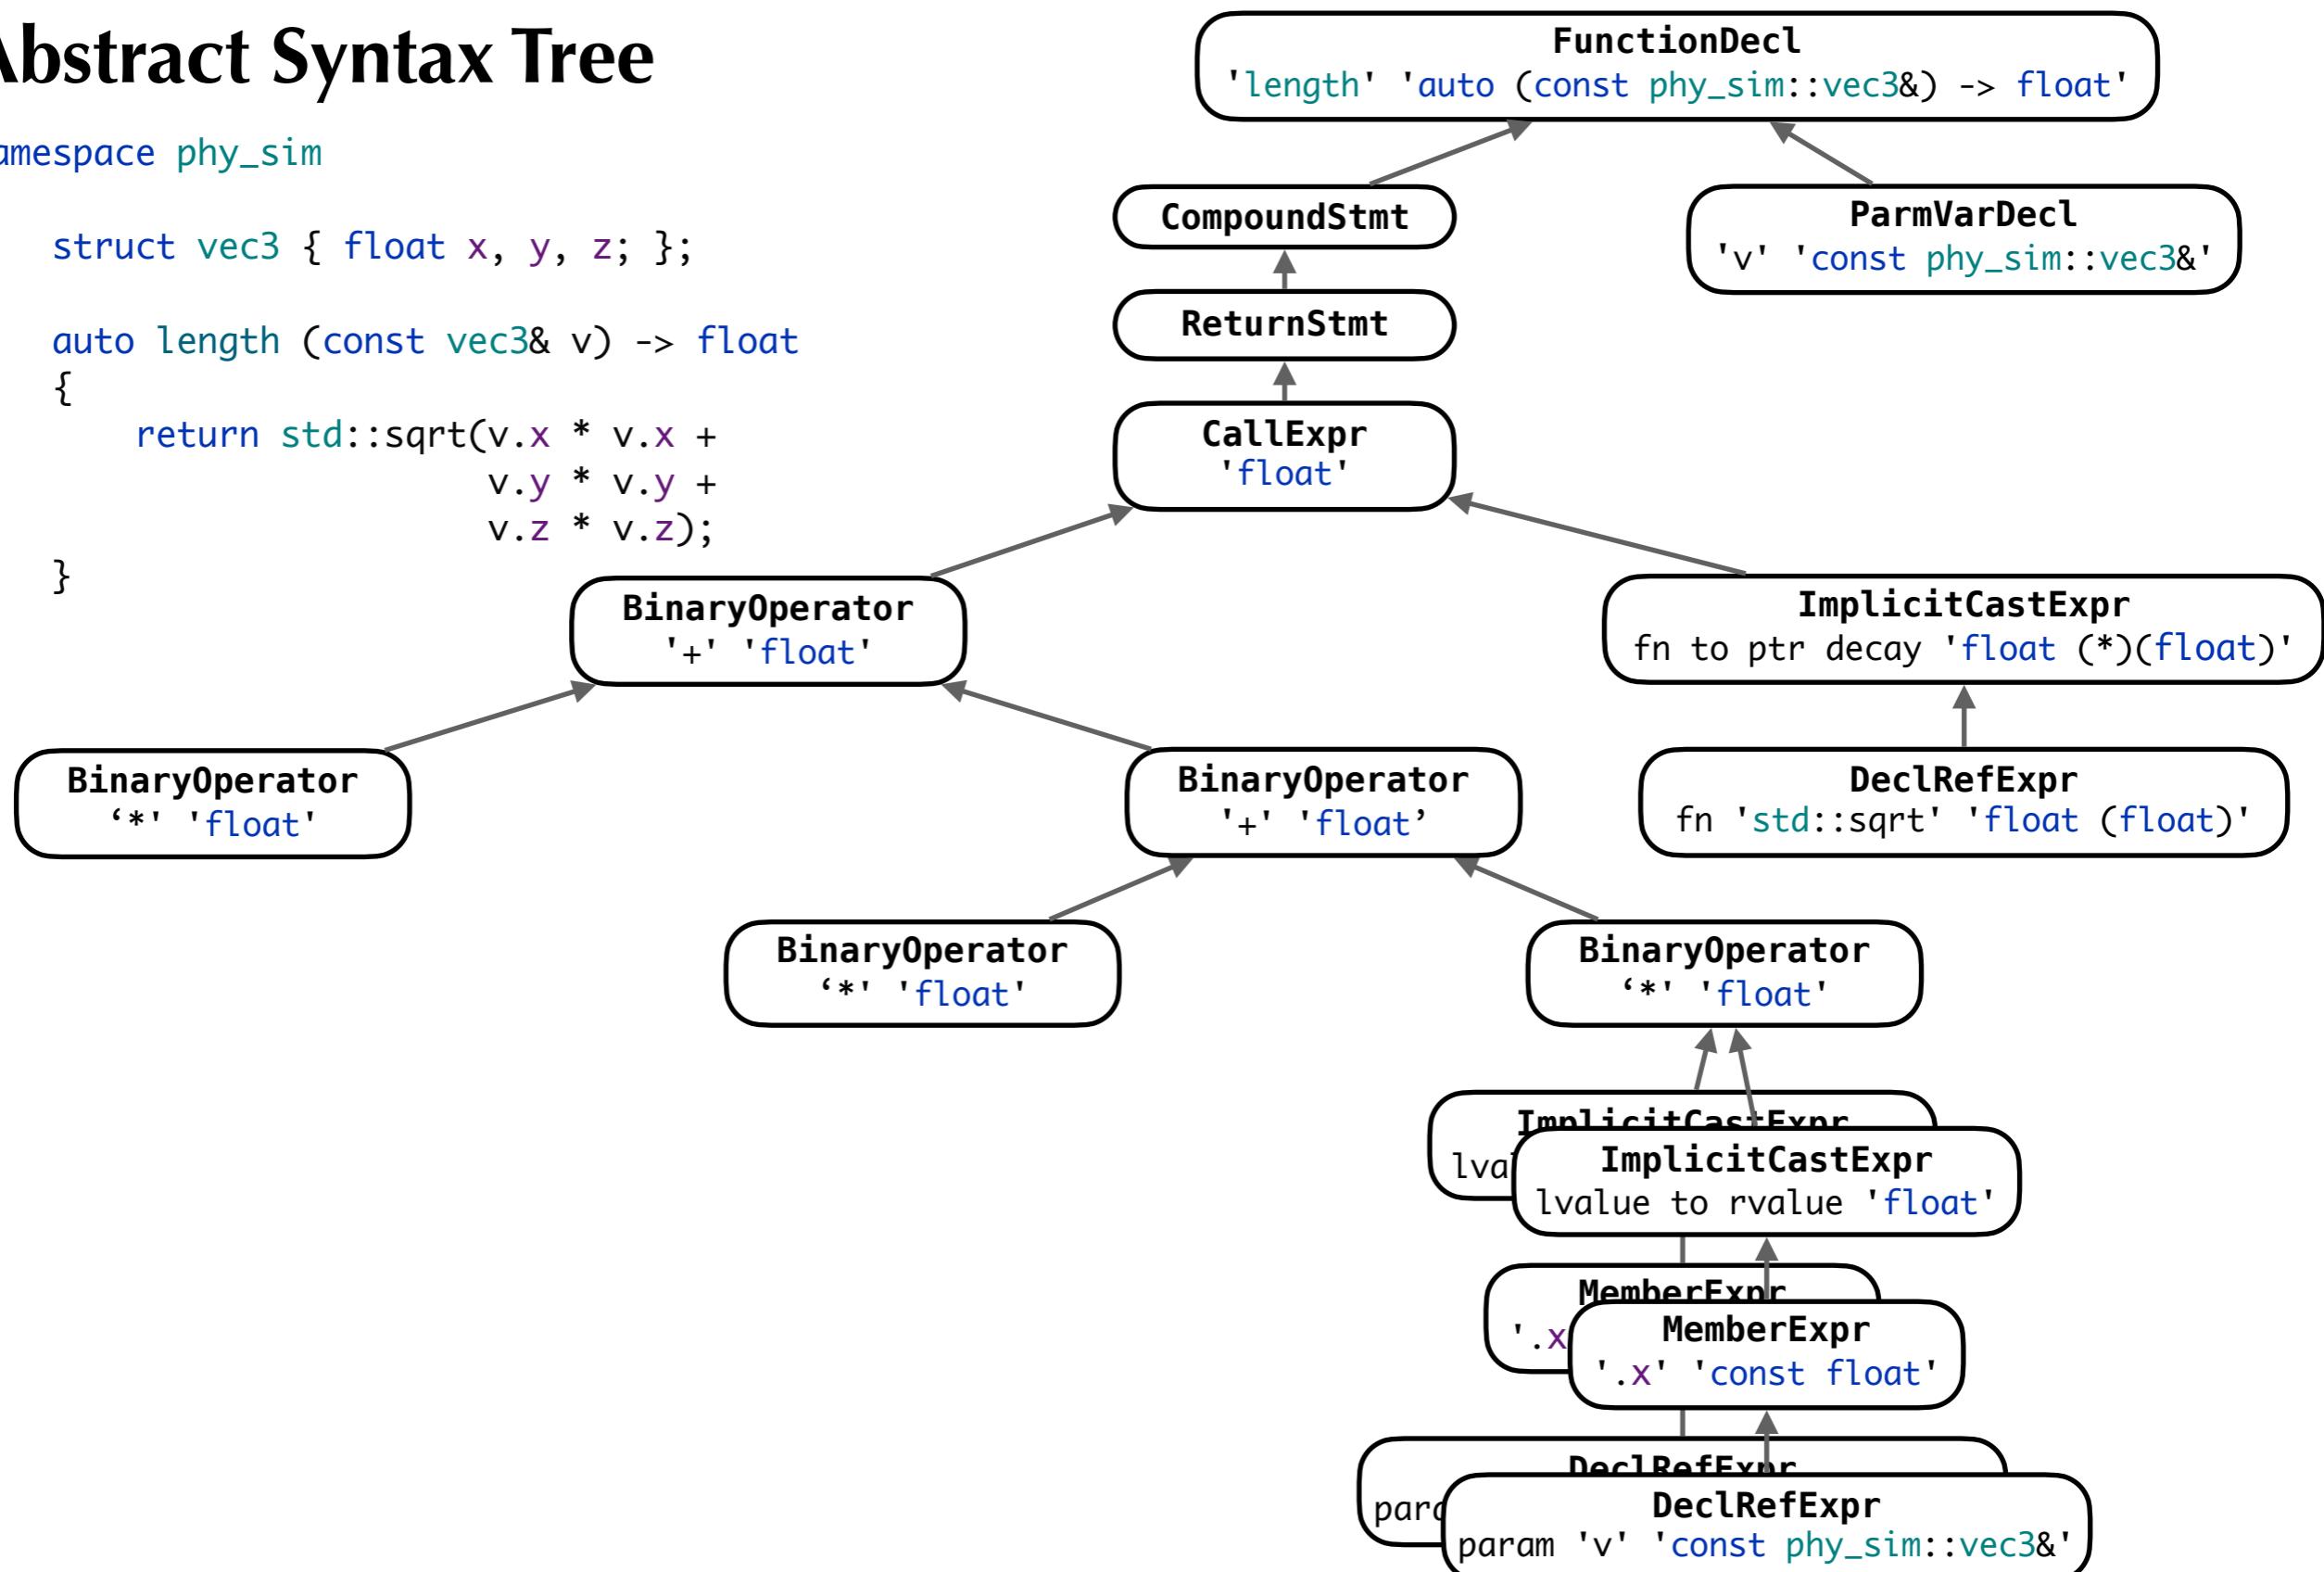


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

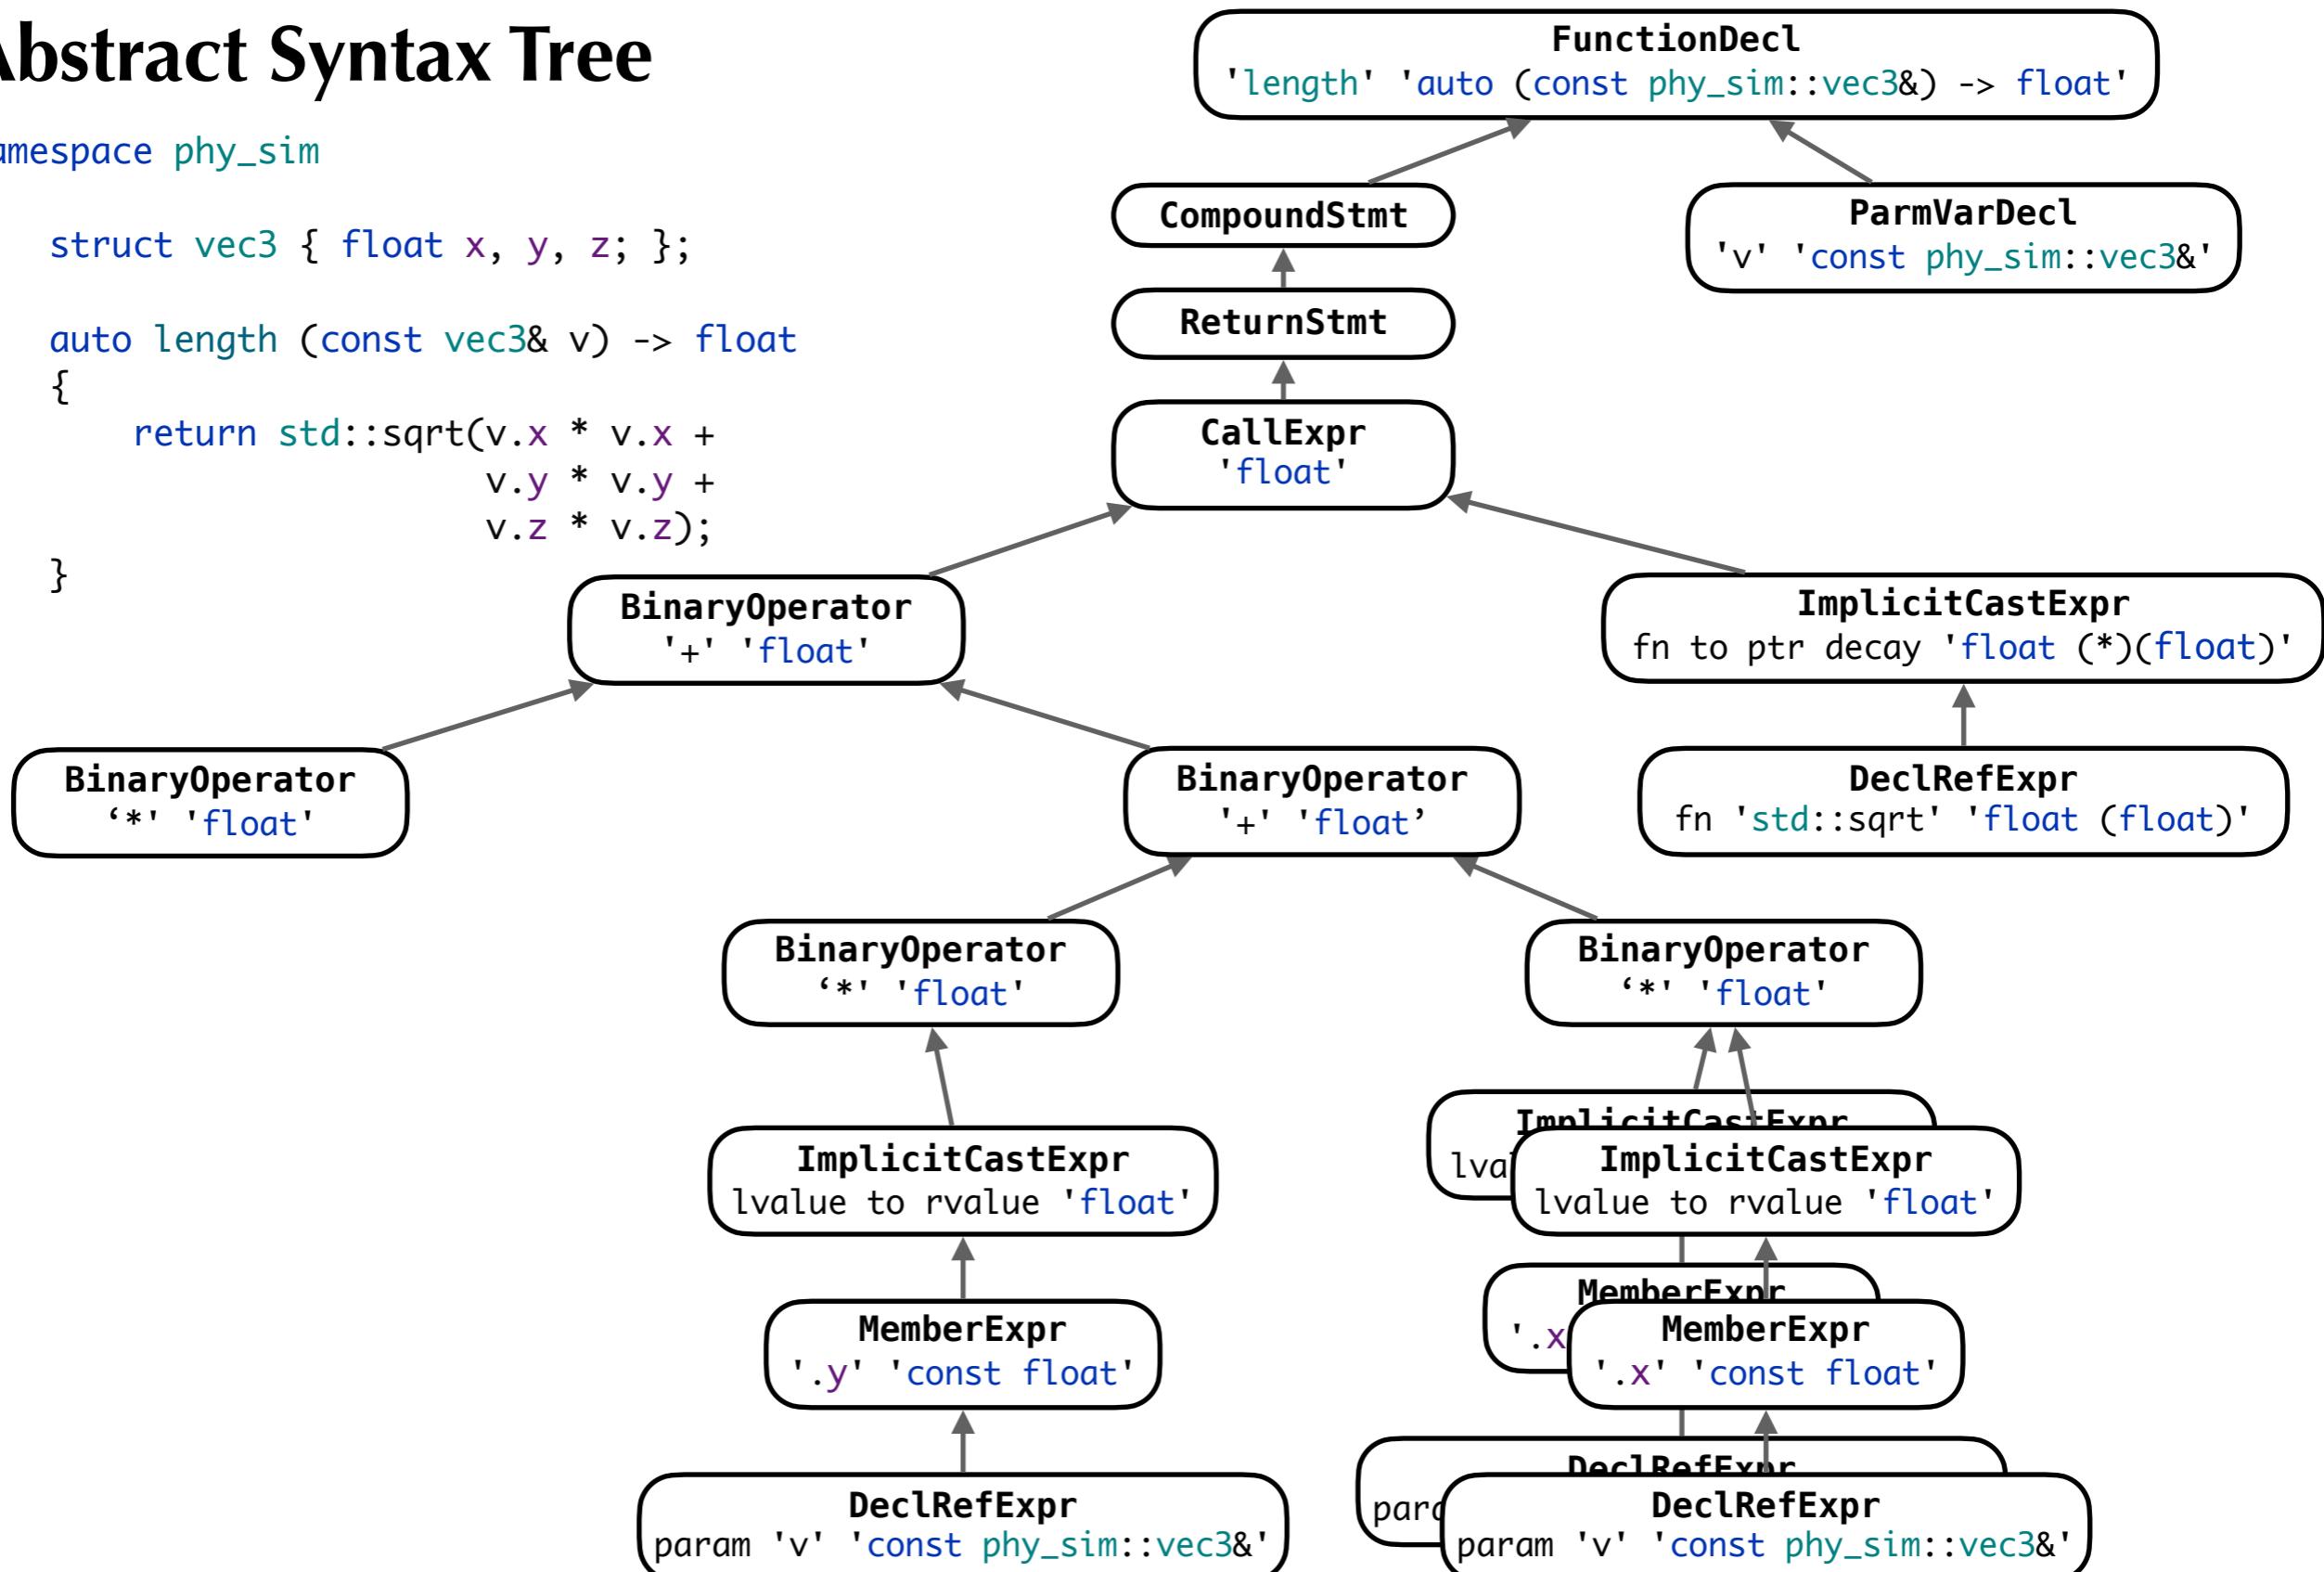


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

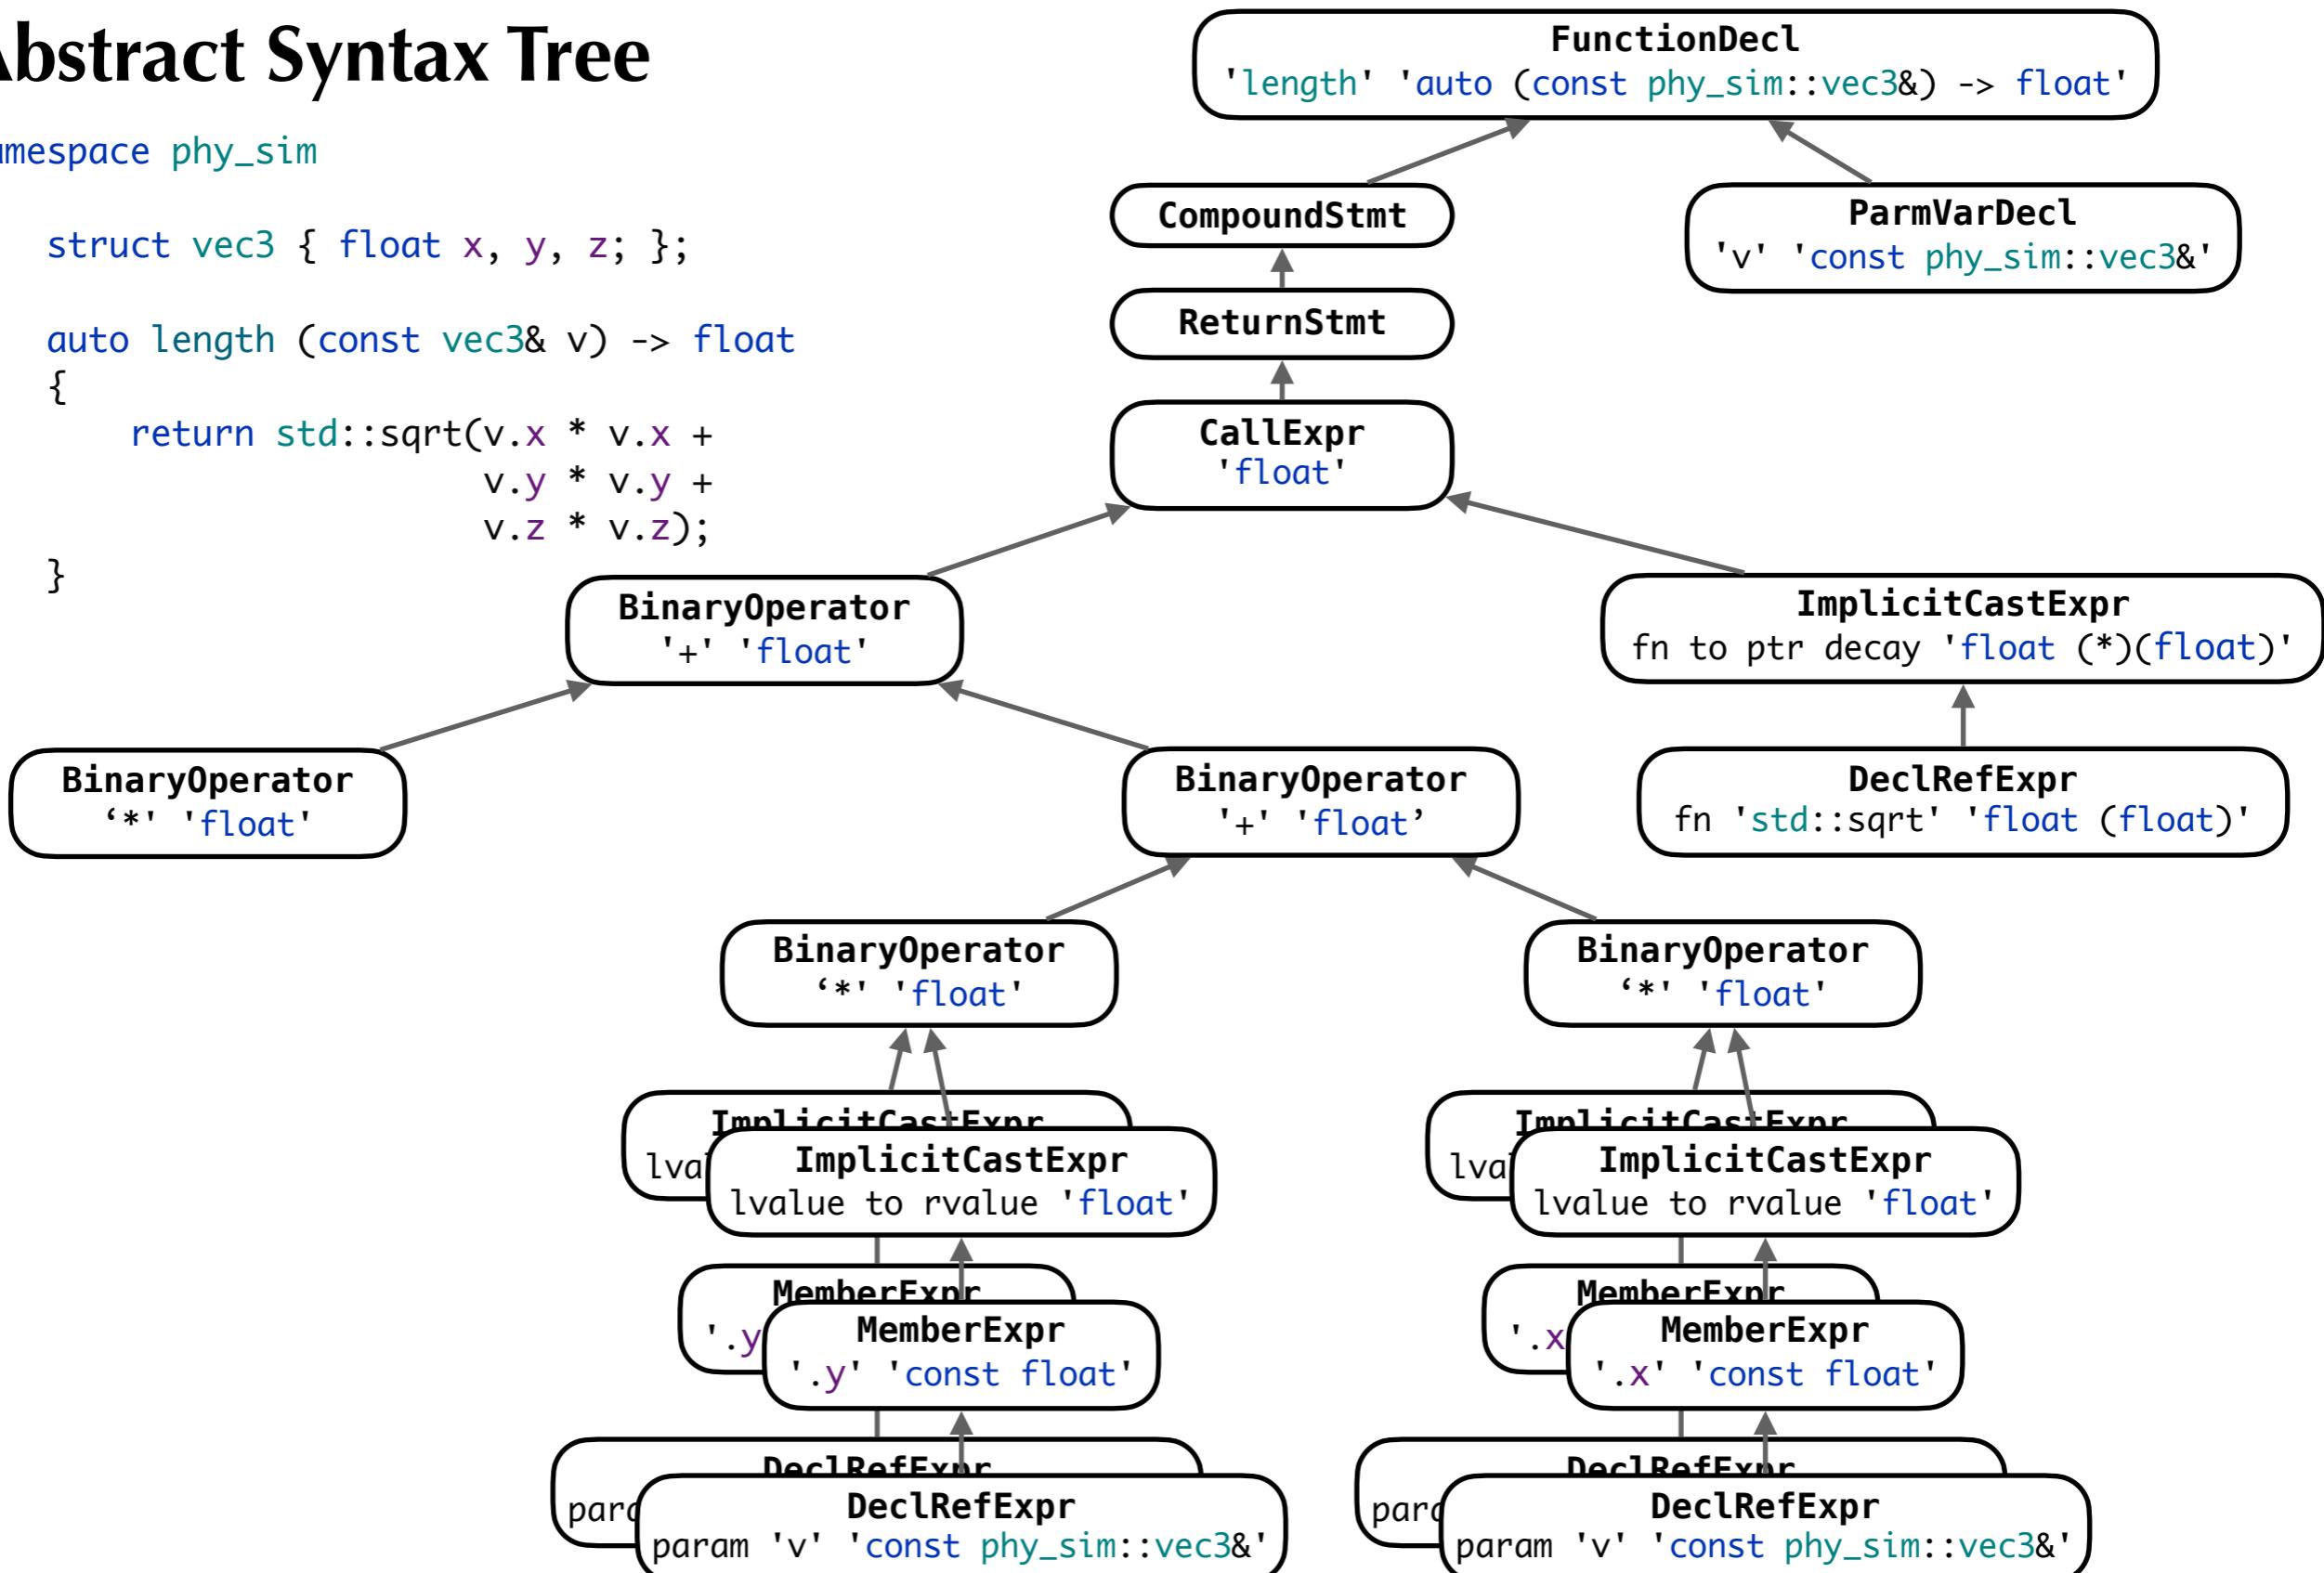


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```

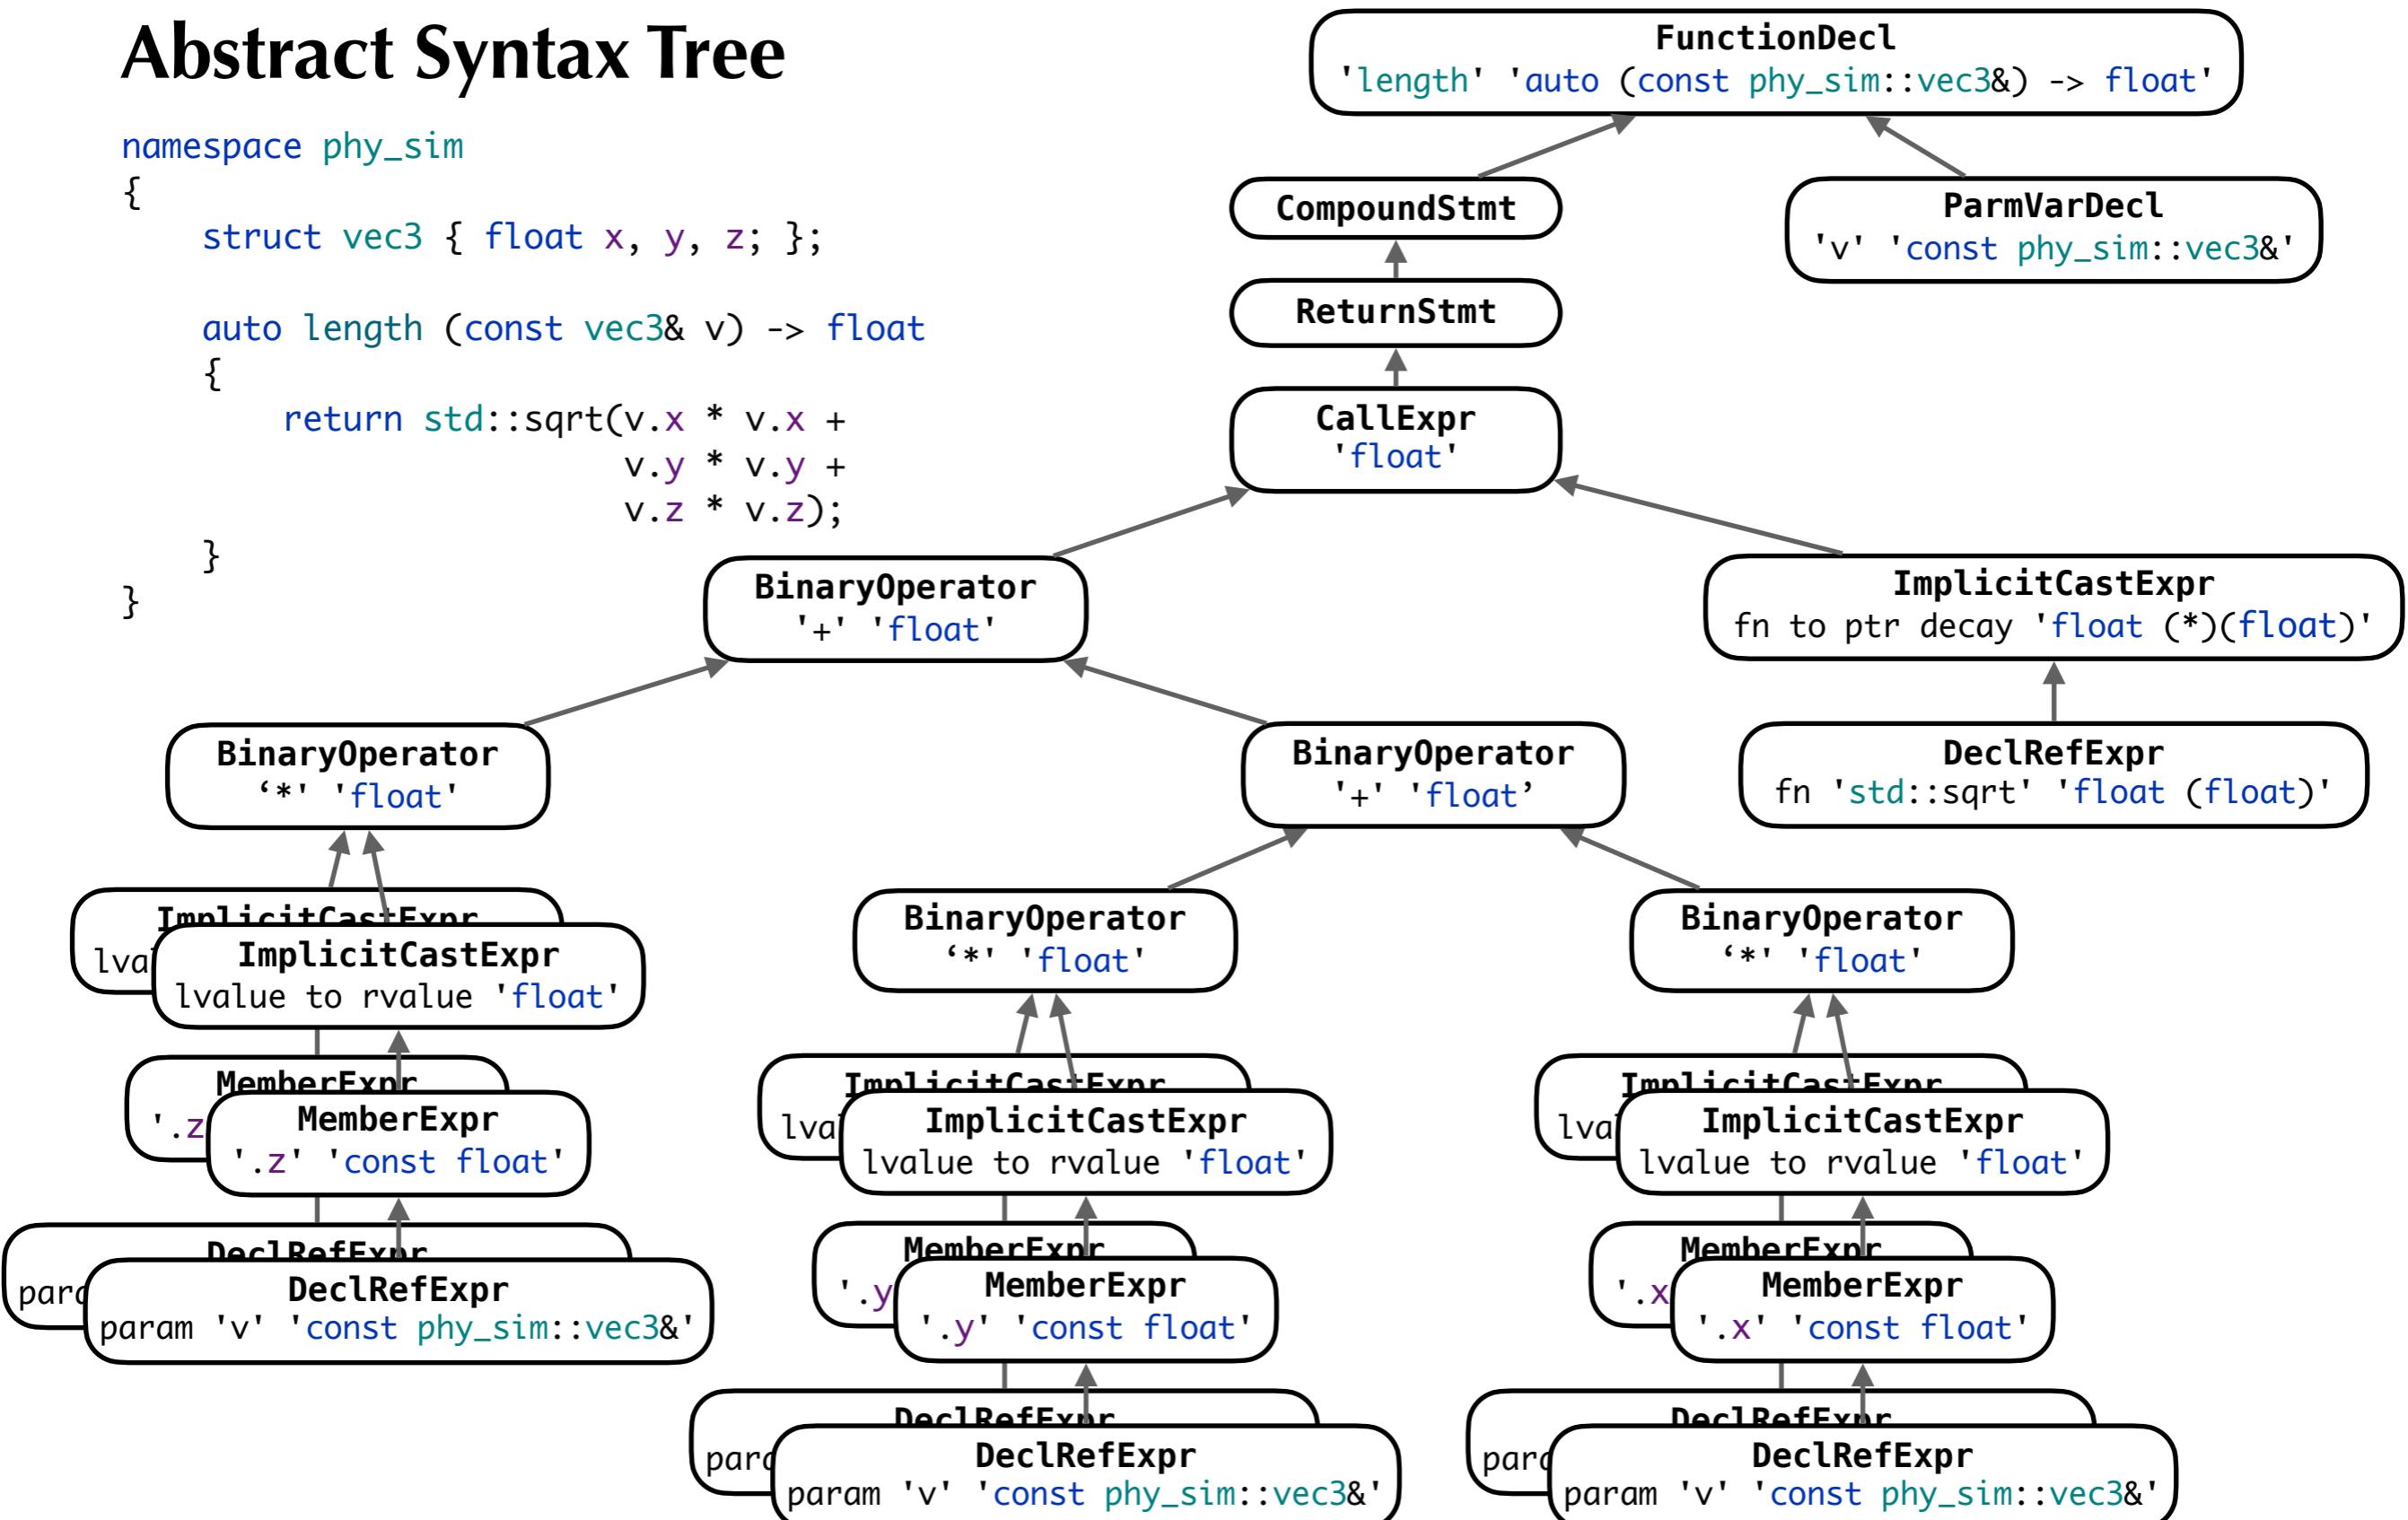


Compiler Architecture

Abstract Syntax Tree

```
namespace phy_sim
{
    struct vec3 { float x, y, z; };

    auto length (const vec3& v) -> float
    {
        return std::sqrt(v.x * v.x +
                         v.y * v.y +
                         v.z * v.z);
    }
}
```



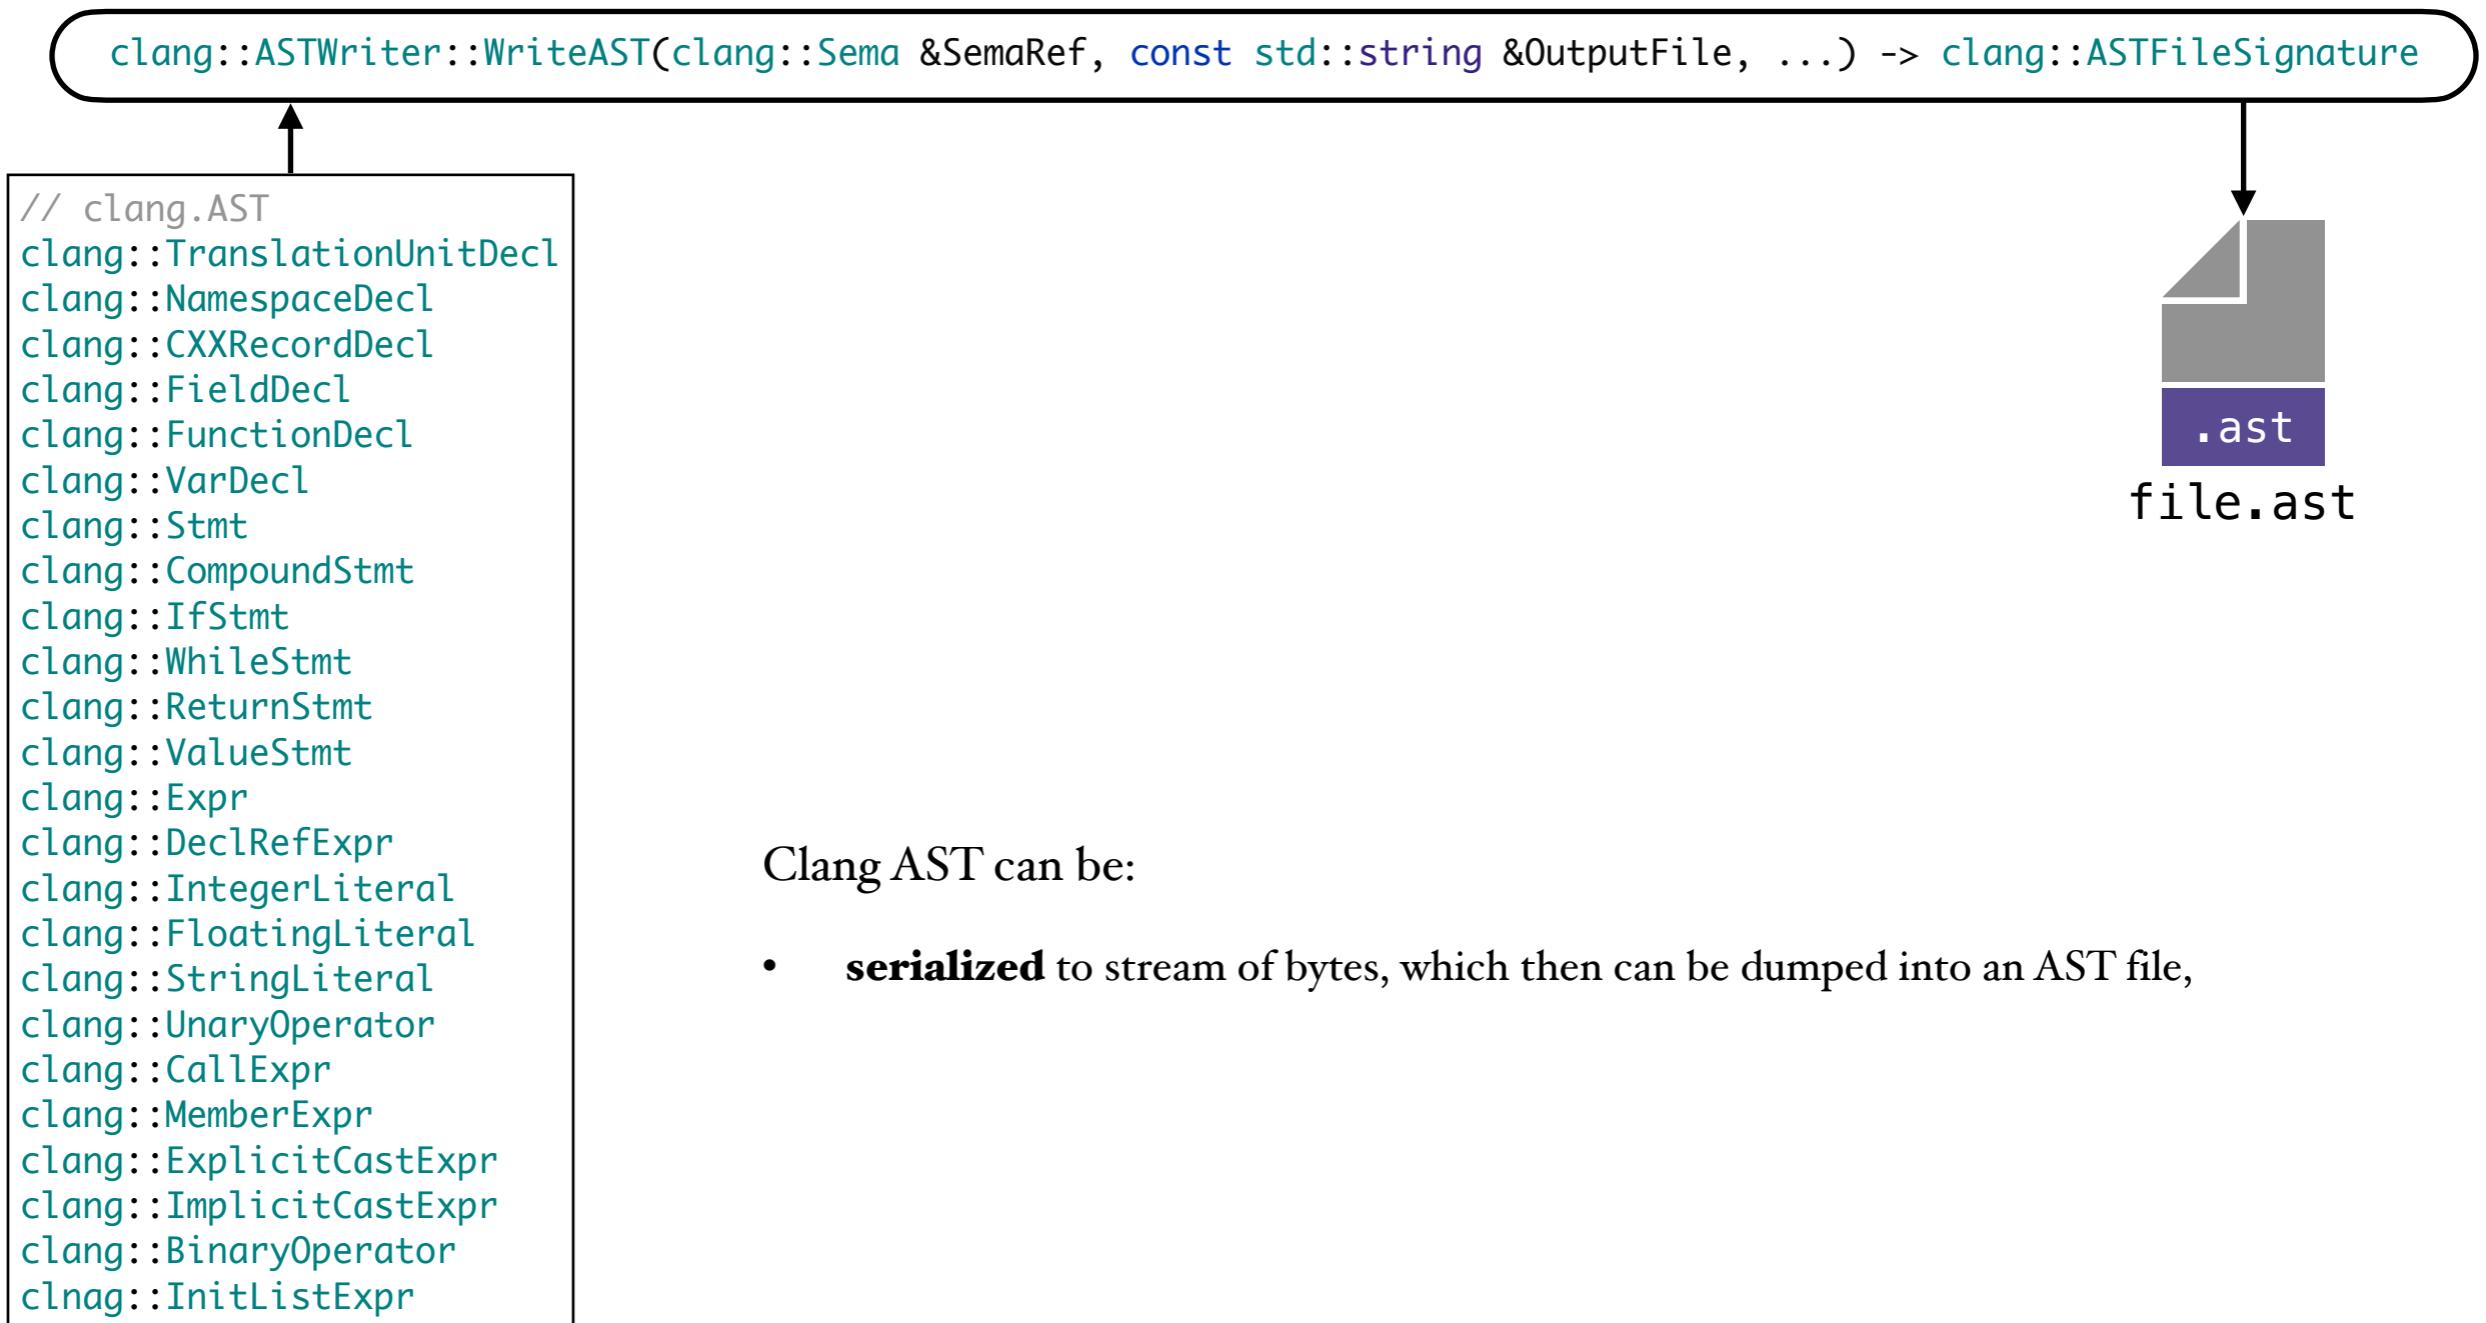
Compiler Architecture

Serialization & Deserialization of Clang AST

```
// clang.AST
clang::TranslationUnitDecl
clang::NamespaceDecl
clang::CXXRecordDecl
clang::FieldDecl
clang::FunctionDecl
clang::VarDecl
clang::Stmt
clang::CompoundStmt
clang::IfStmt
clang::WhileStmt
clang::ReturnStmt
clang::ValueStmt
clang::Expr
clang::DeclRefExpr
clang::IntegerLiteral
clang::FloatingLiteral
clang::StringLiteral
clang::UnaryOperator
clang::CallExpr
clang::MemberExpr
clang::ExplicitCastExpr
clang::ImplicitCastExpr
clang::BinaryOperator
clang::InitListExpr
```

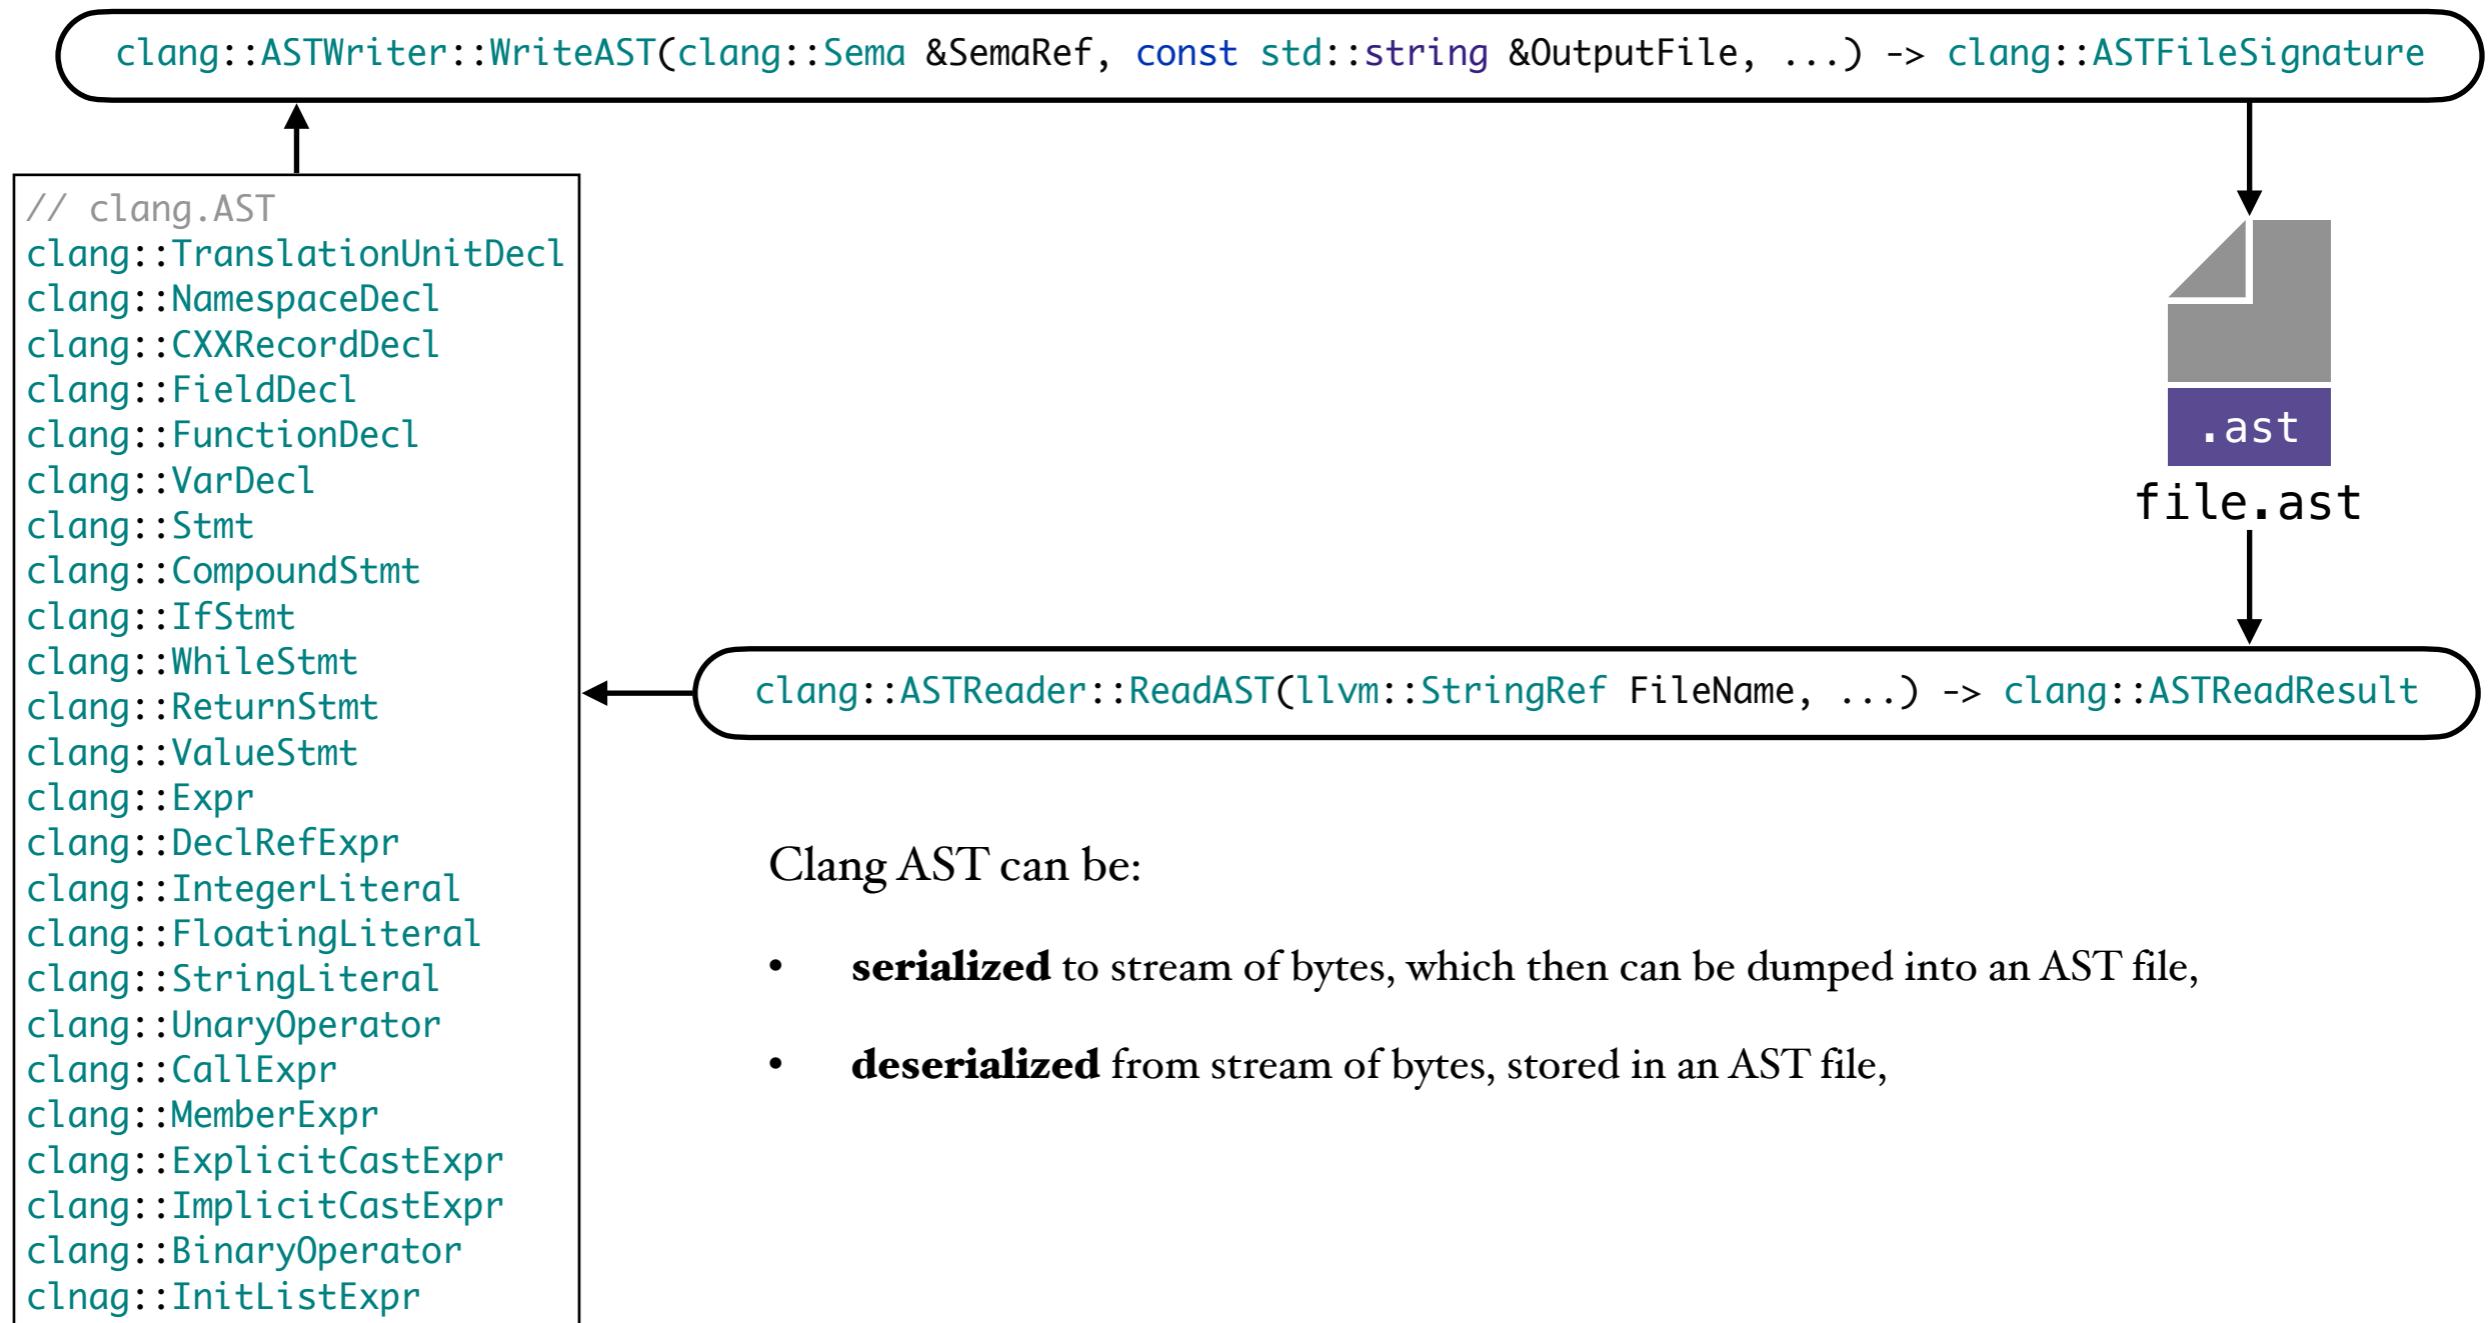
Compiler Architecture

Serialization & Deserialization of Clang AST



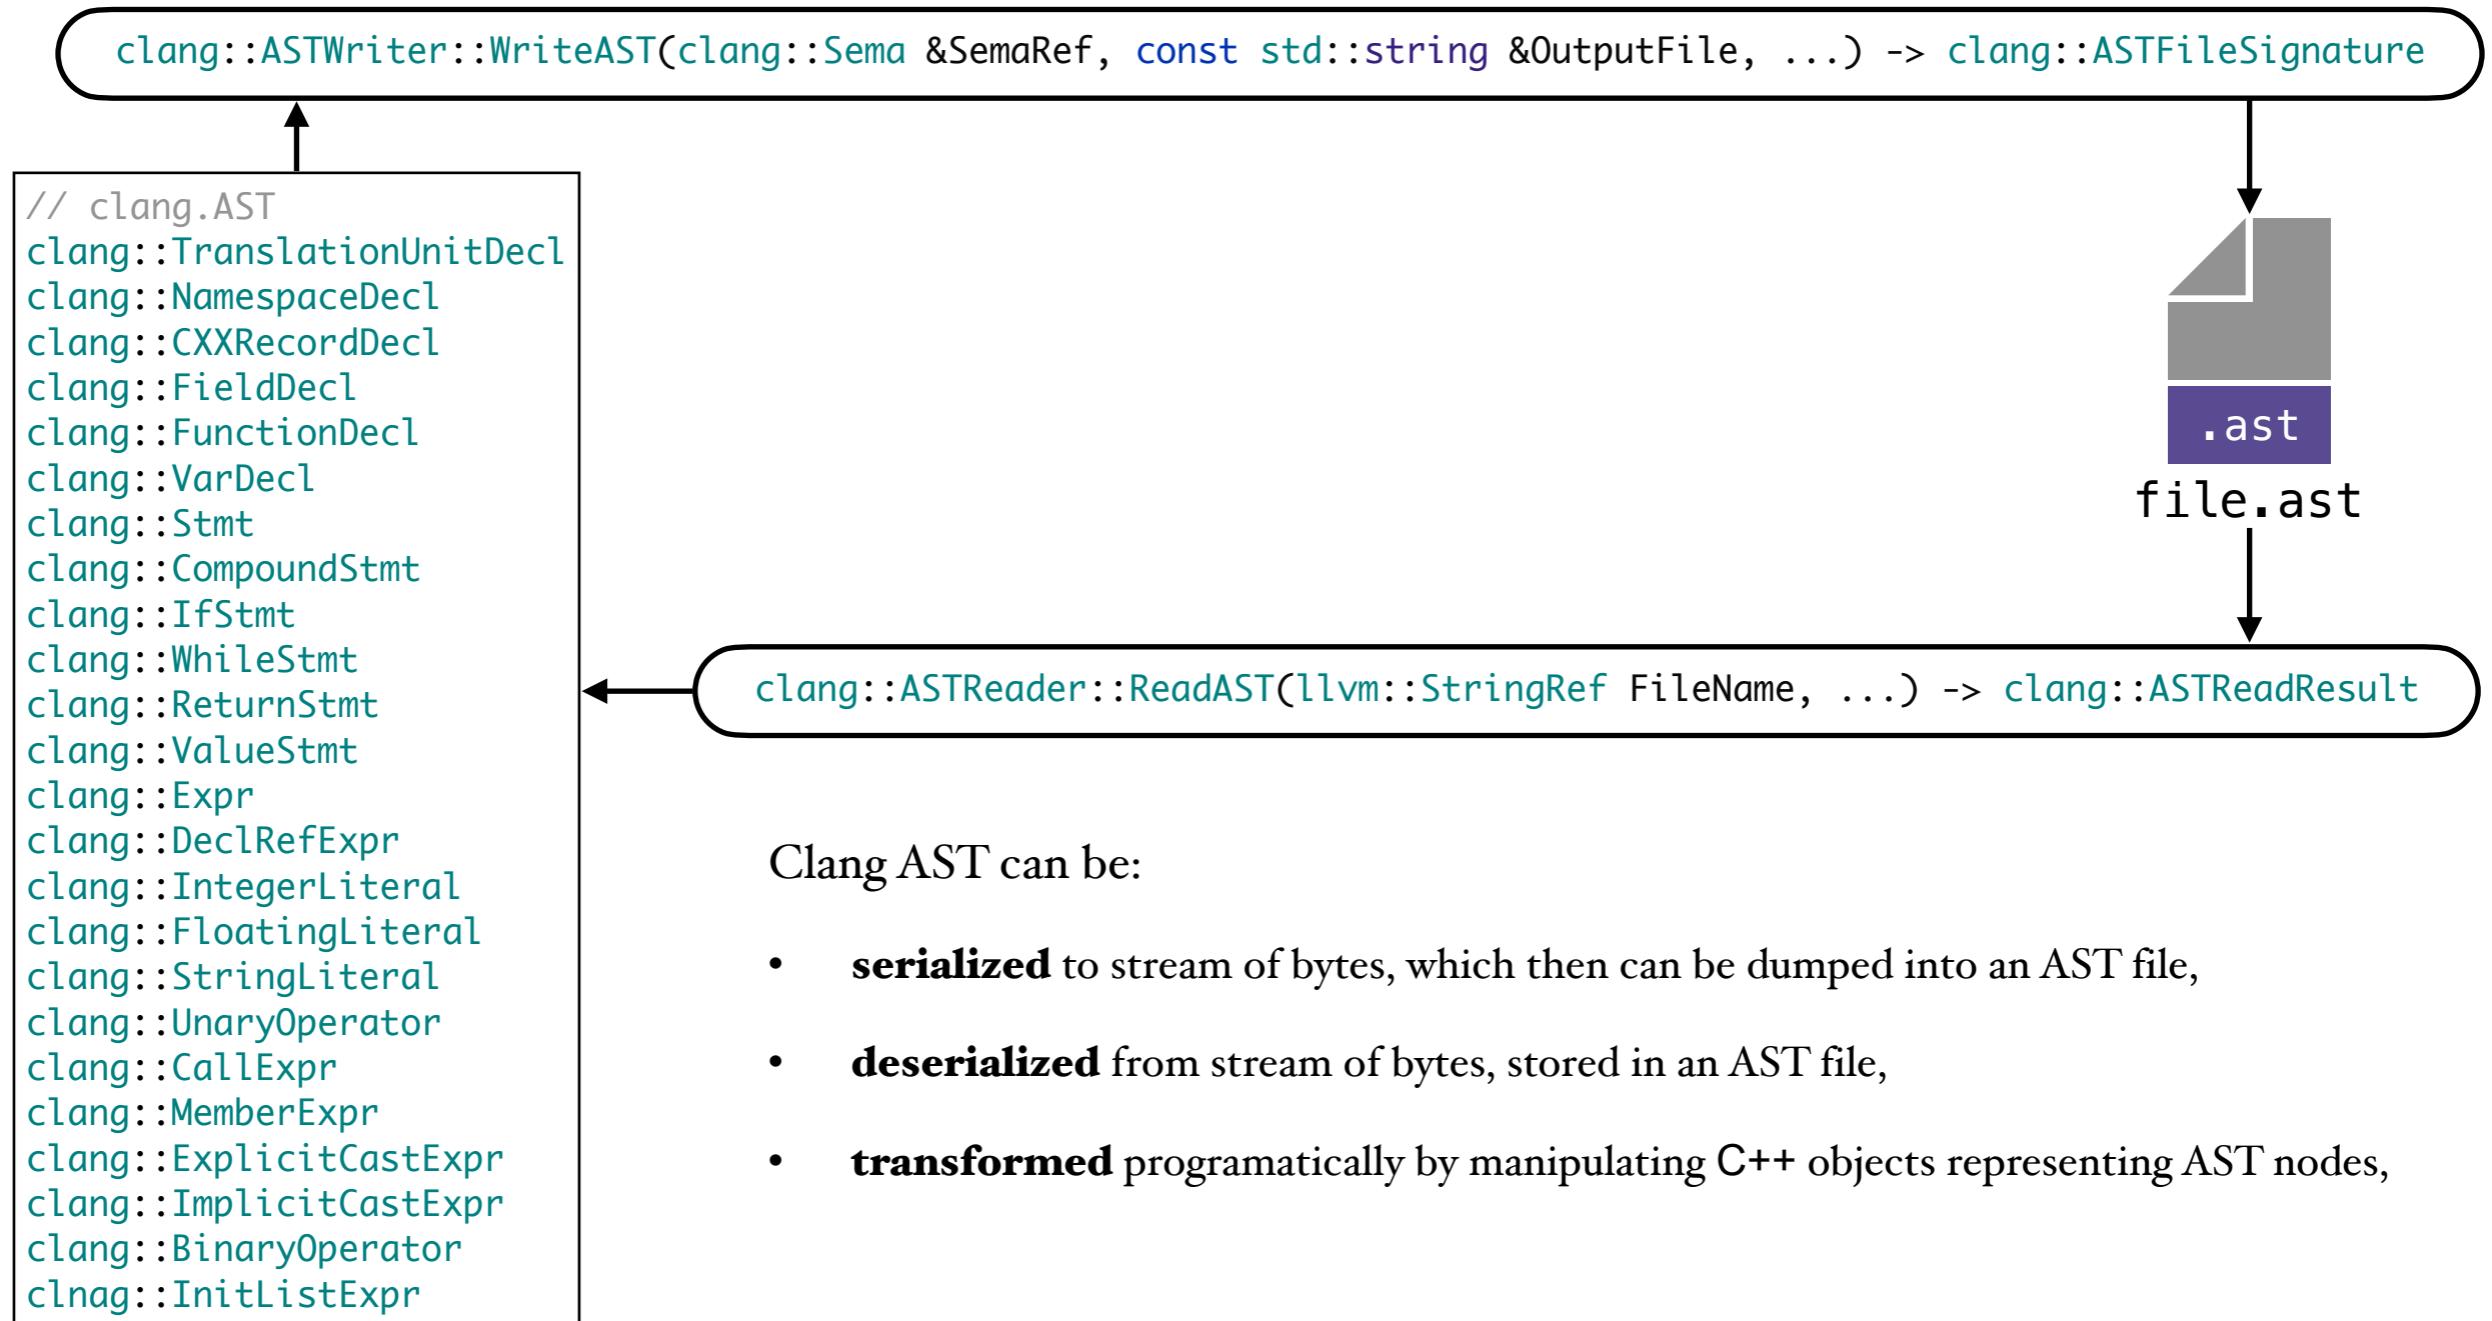
Compiler Architecture

Serialization & Deserialization of Clang AST



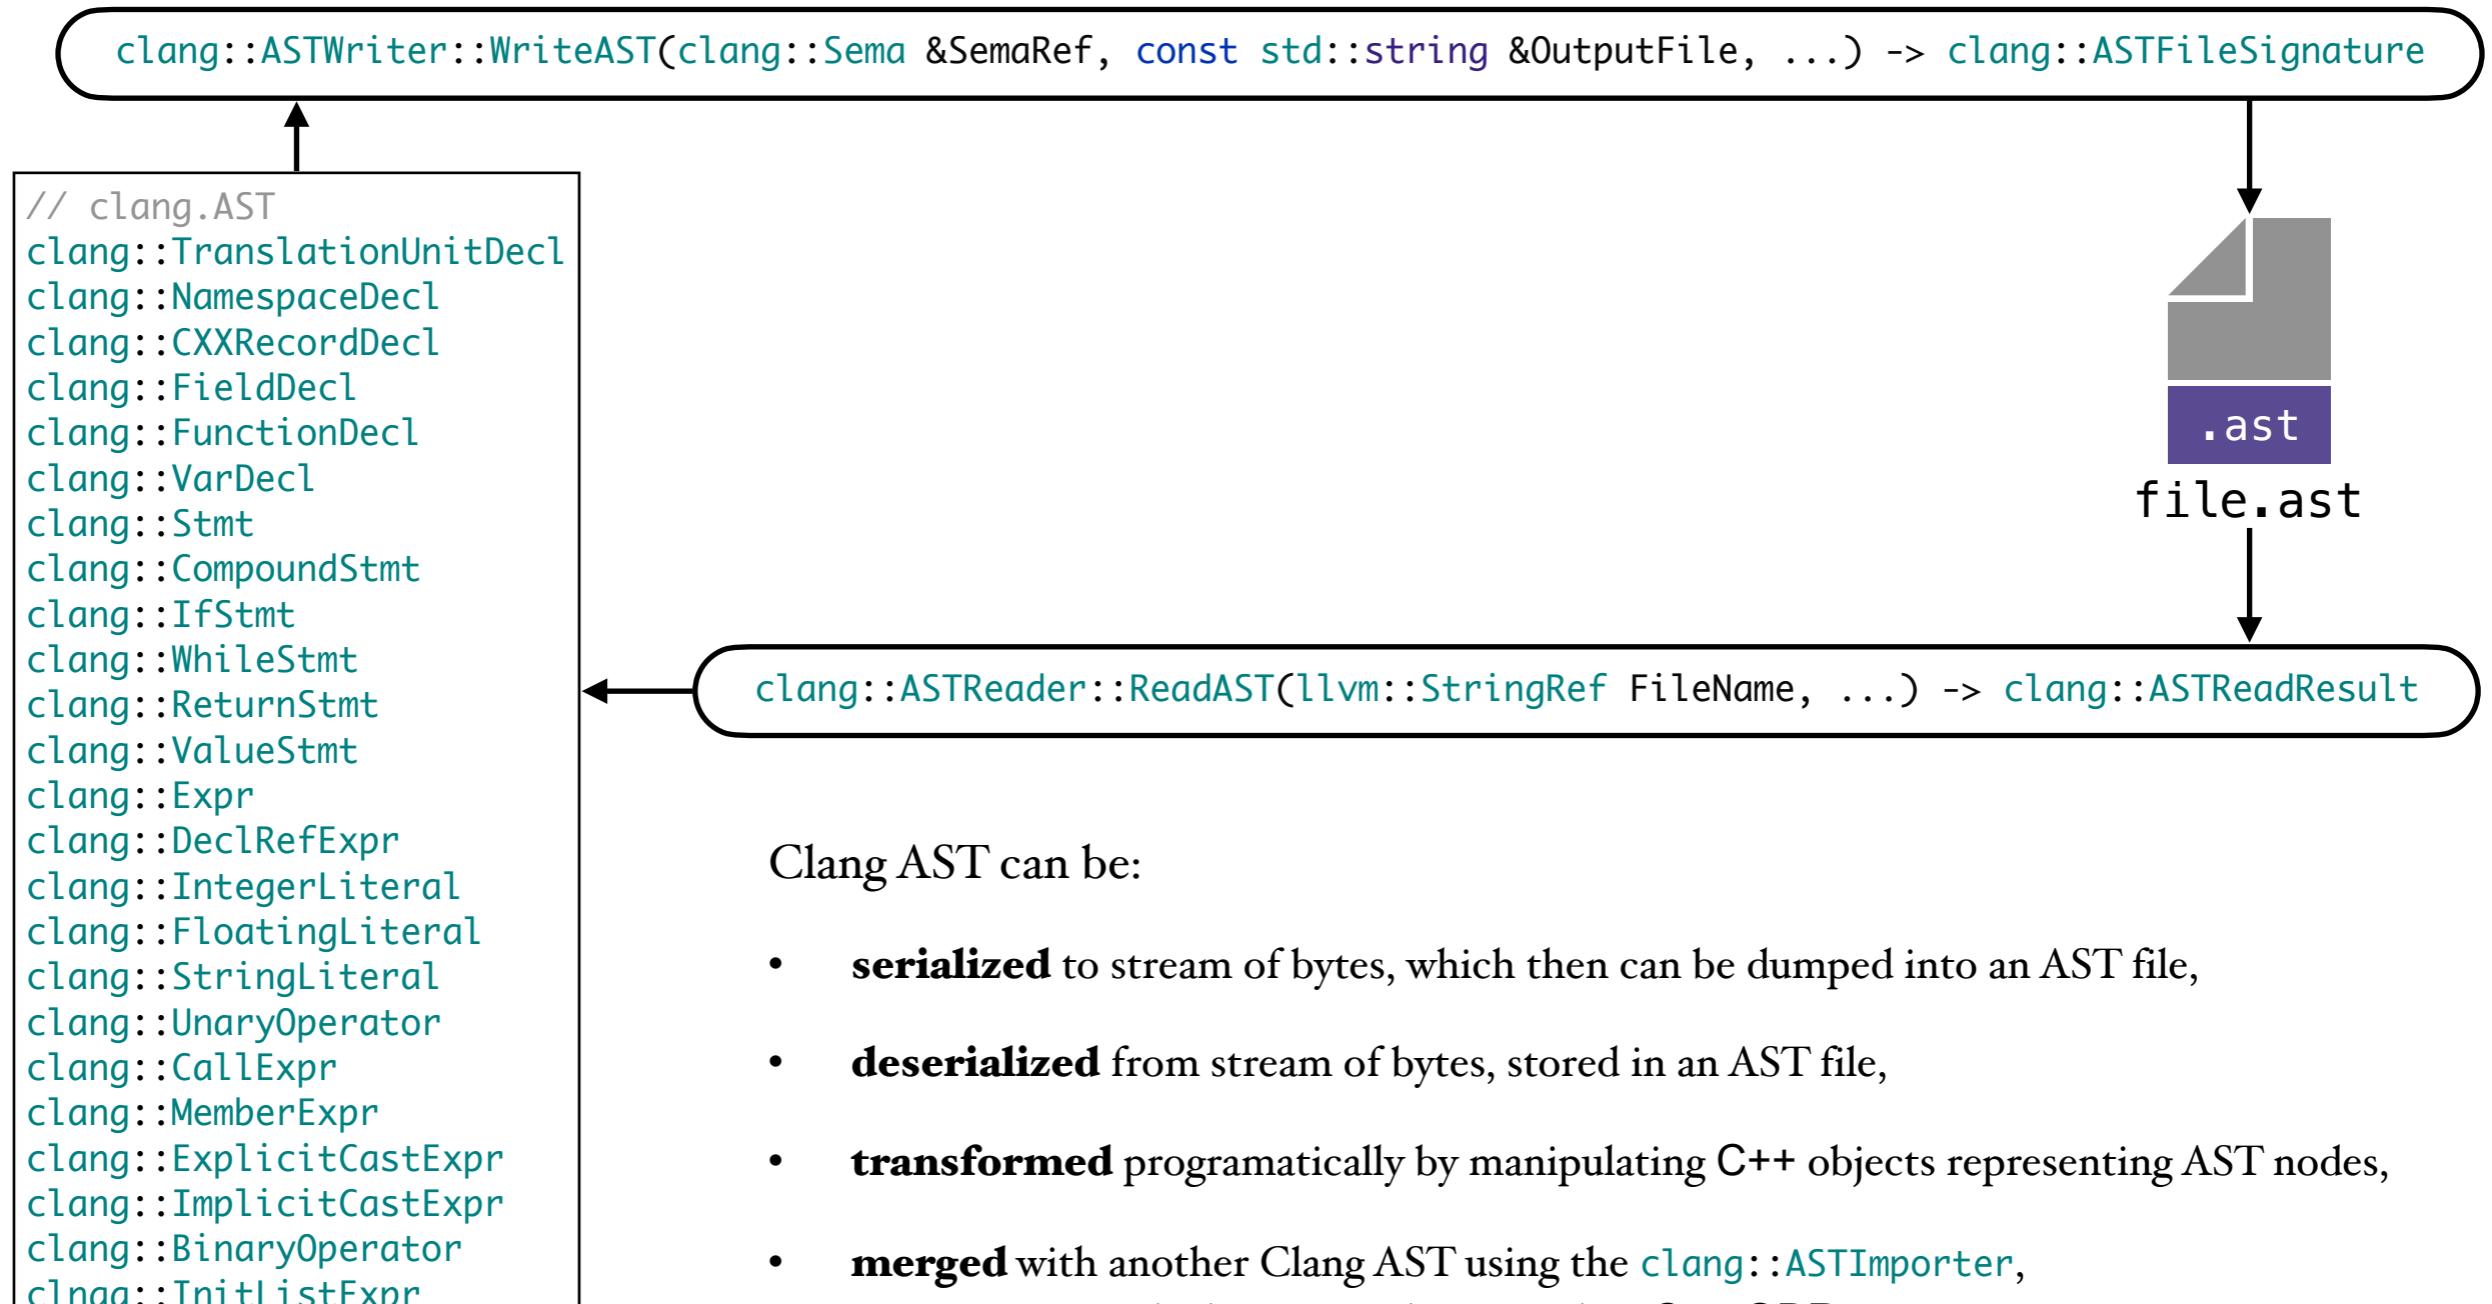
Compiler Architecture

Serialization & Deserialization of Clang AST



Compiler Architecture

Serialization & Deserialization of Clang AST



Question

How to design better compilation model for C++?

Question

How to design better compilation model for C++?

Key ideas:

- I. Introduce **C++ modules**, a new kind of C++ entities, consisting of **interface** and **implementation units**, to which already existing C++ entities could belong to, to provide true **encapsulation** of implementation details.

Question

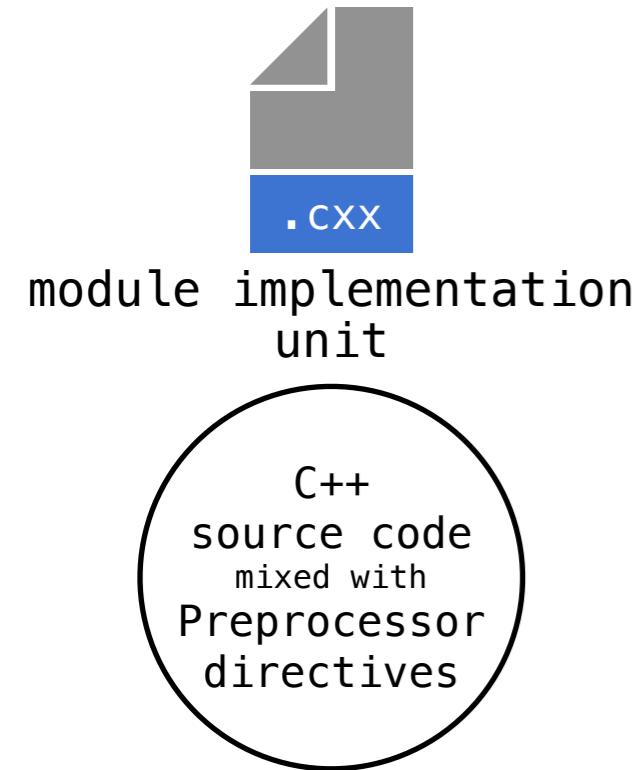
How to design better compilation model for C++?

Key ideas:

1. Introduce **C++ modules**, a new kind of C++ entities, consisting of **interface** and **implementation units**, to which already existing C++ entities could belong to, to provide true **encapsulation** of implementation details.
2. Store on a disk **binary representation** of module interface units, which will be **efficient to ingest** by a C++ compiler in subsequent compilations of dependent module units, to **avoid redundant** and **expensive processing** performed by the C++ compiler.

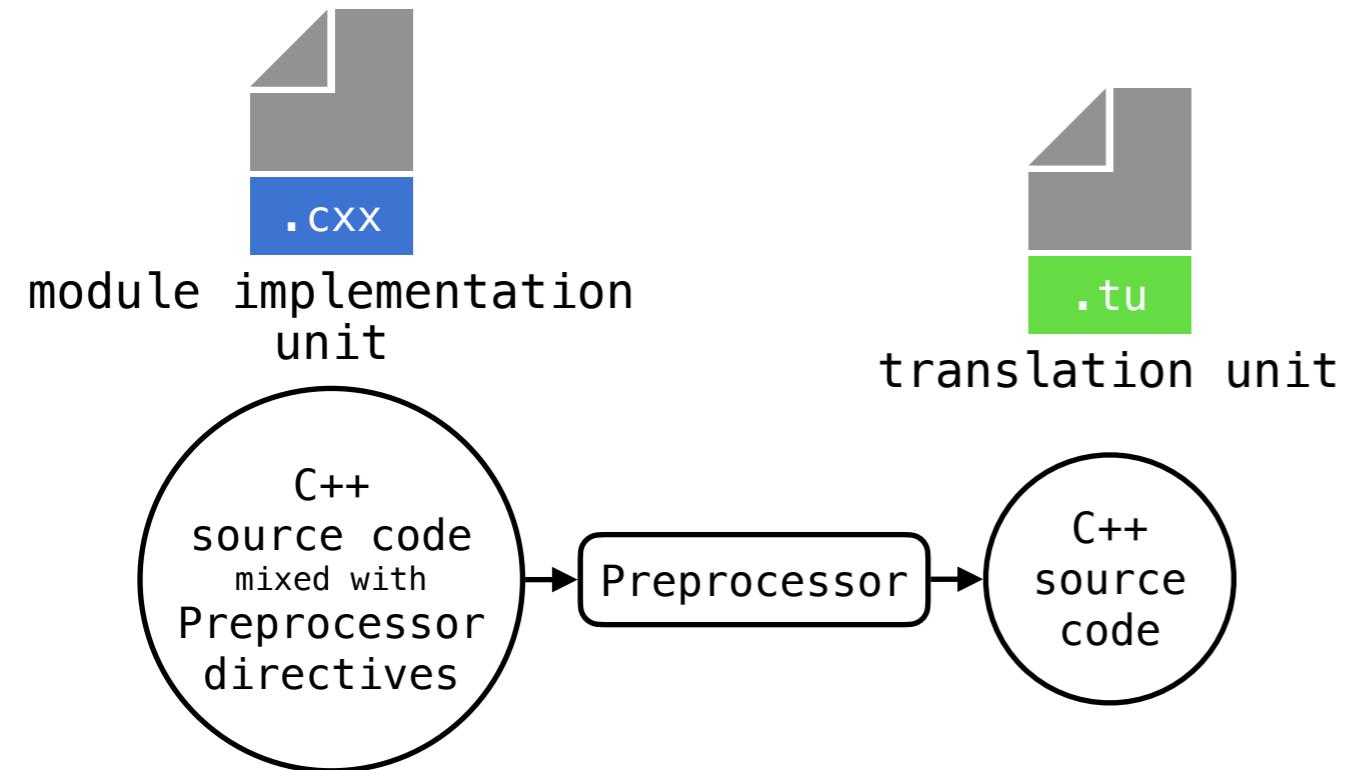
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



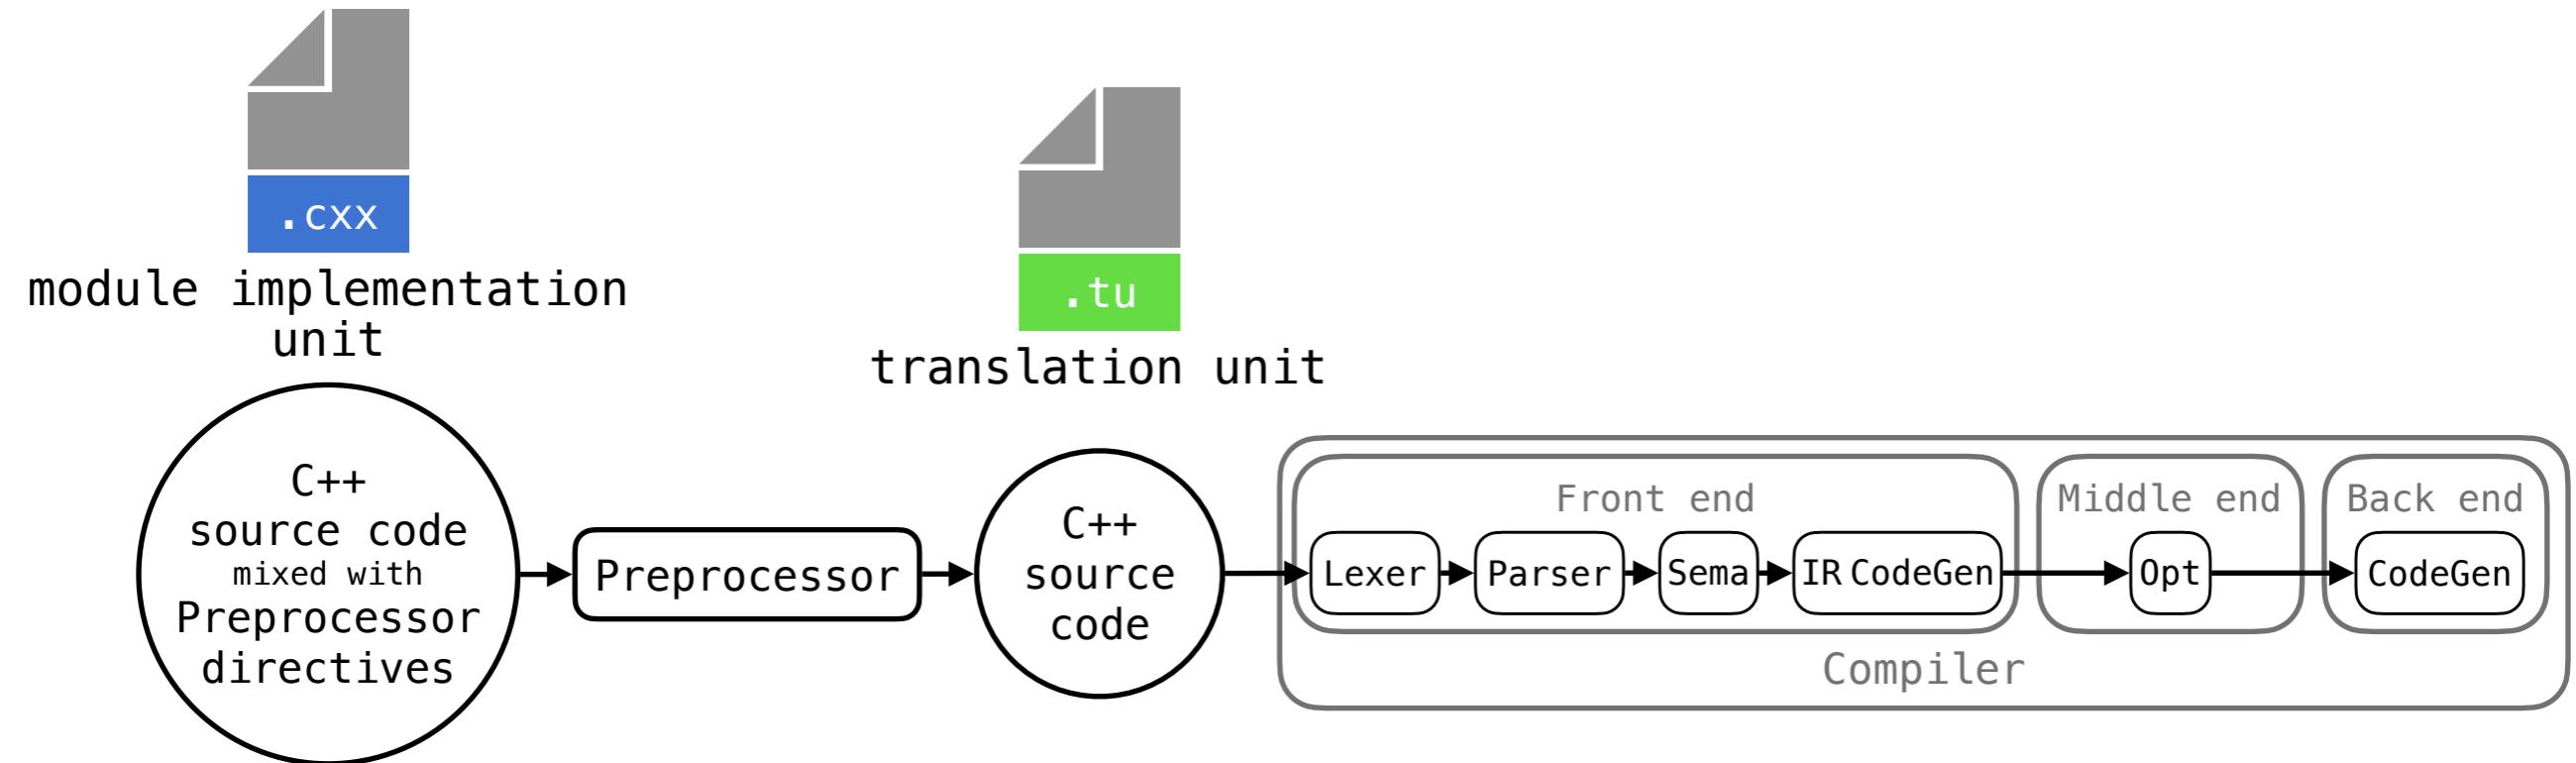
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



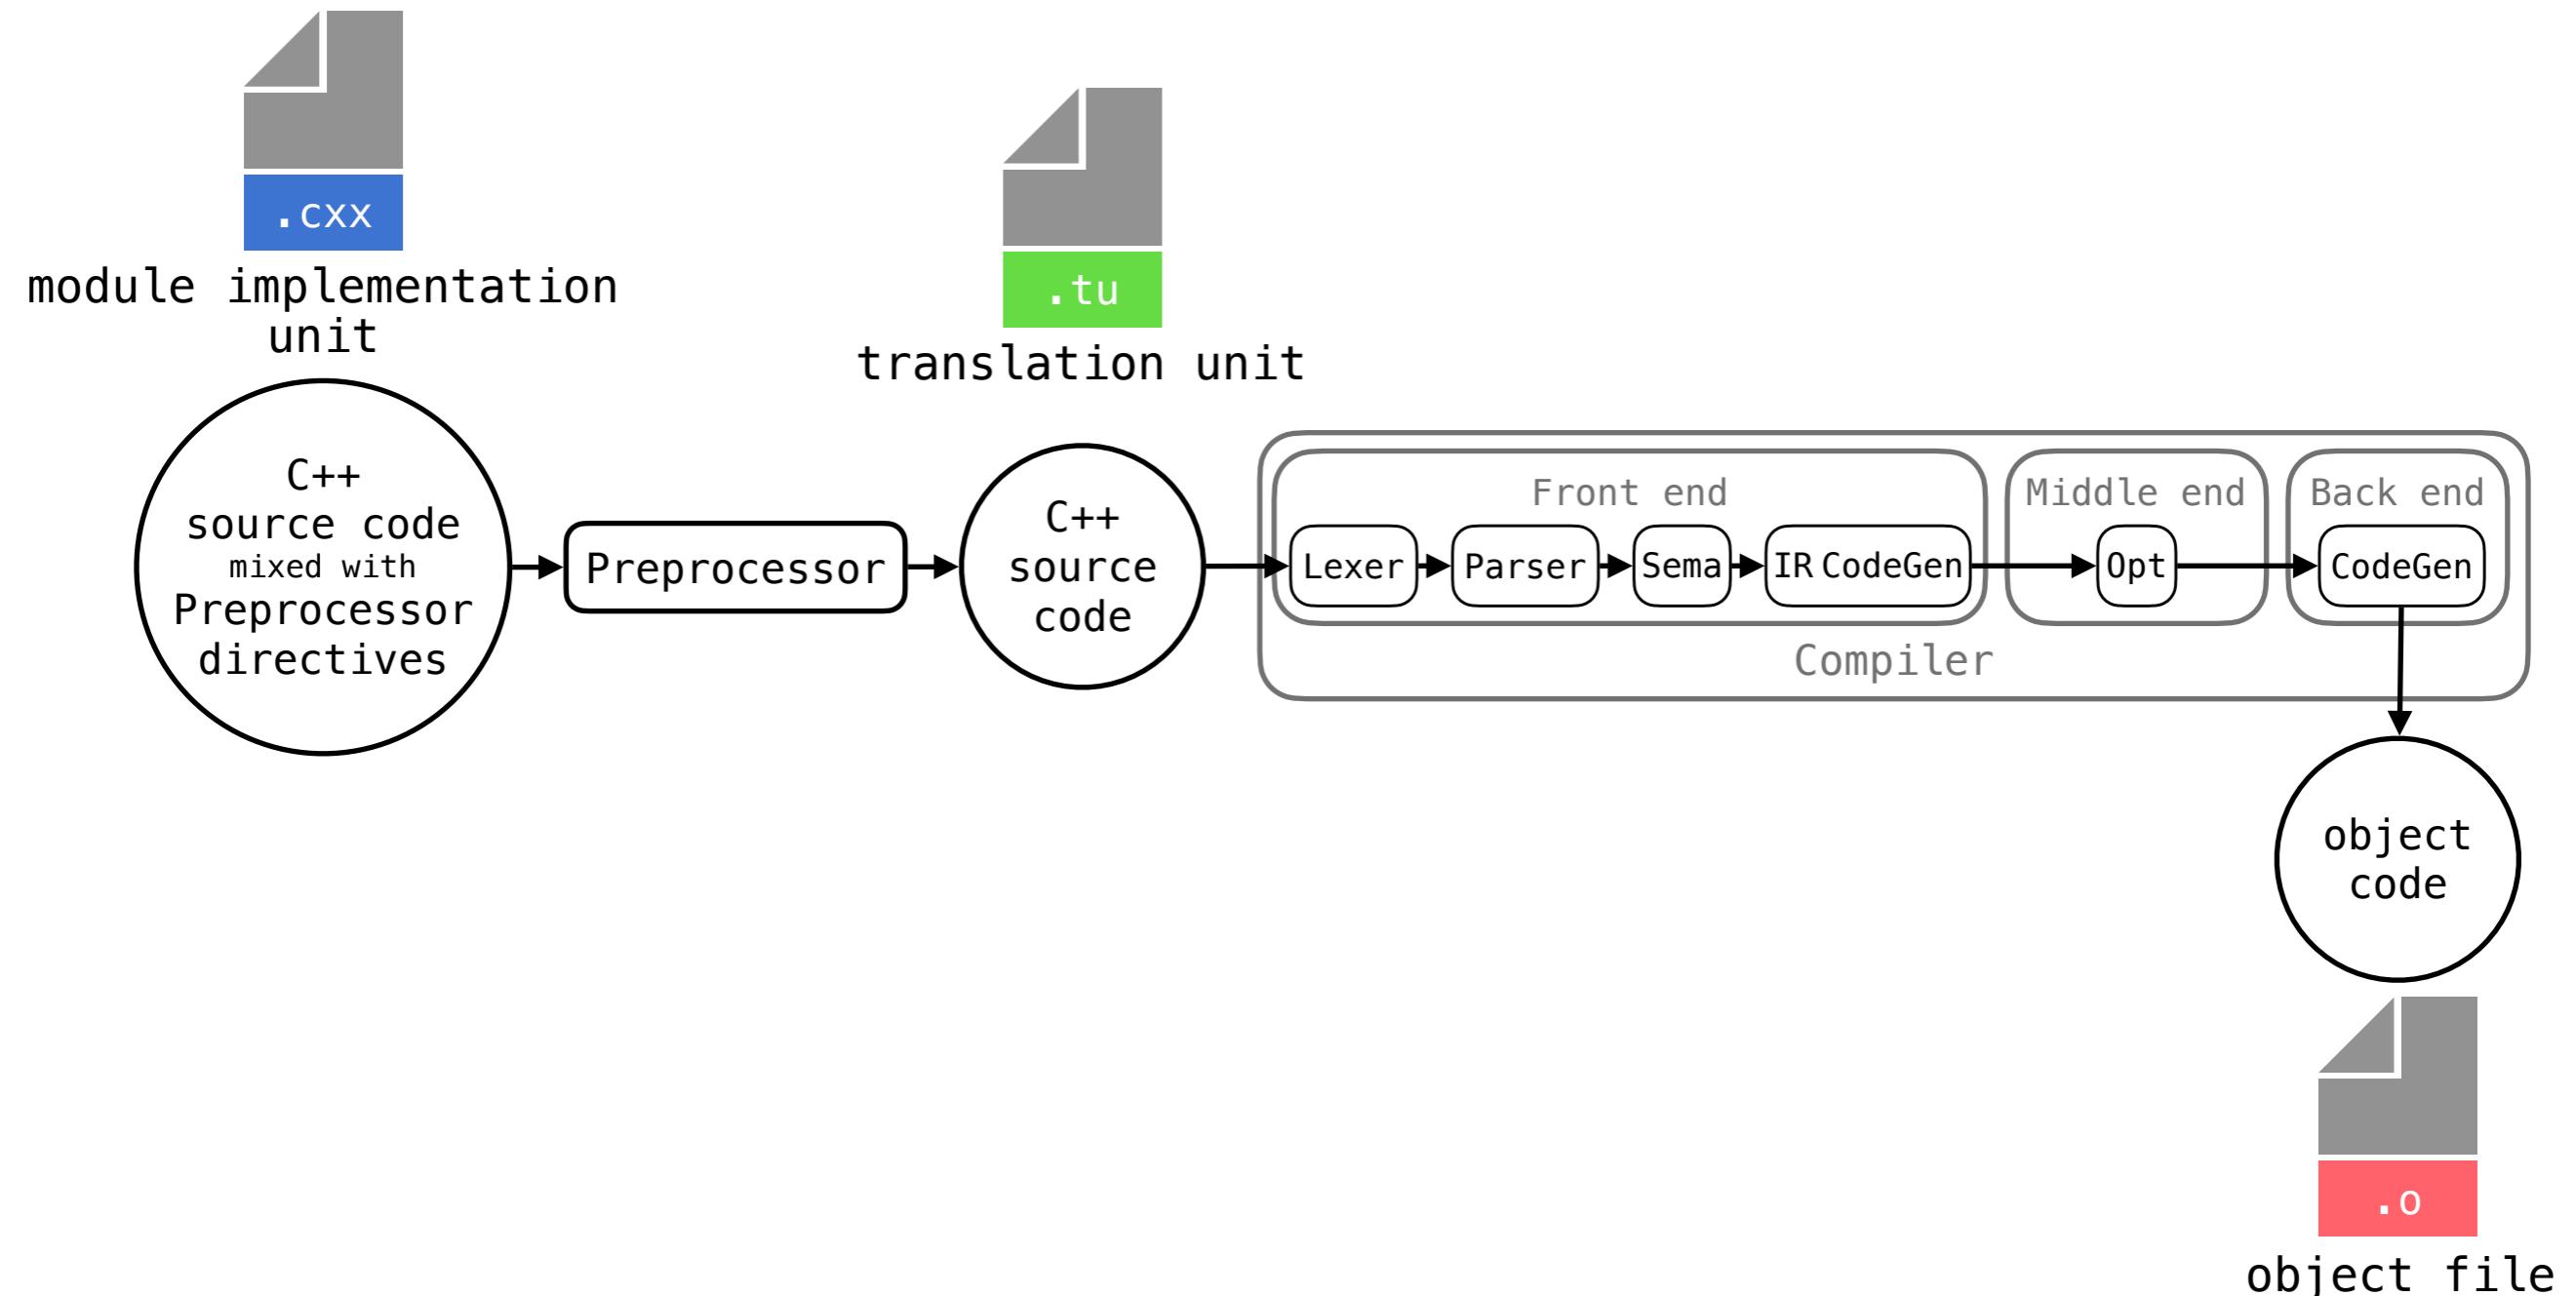
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



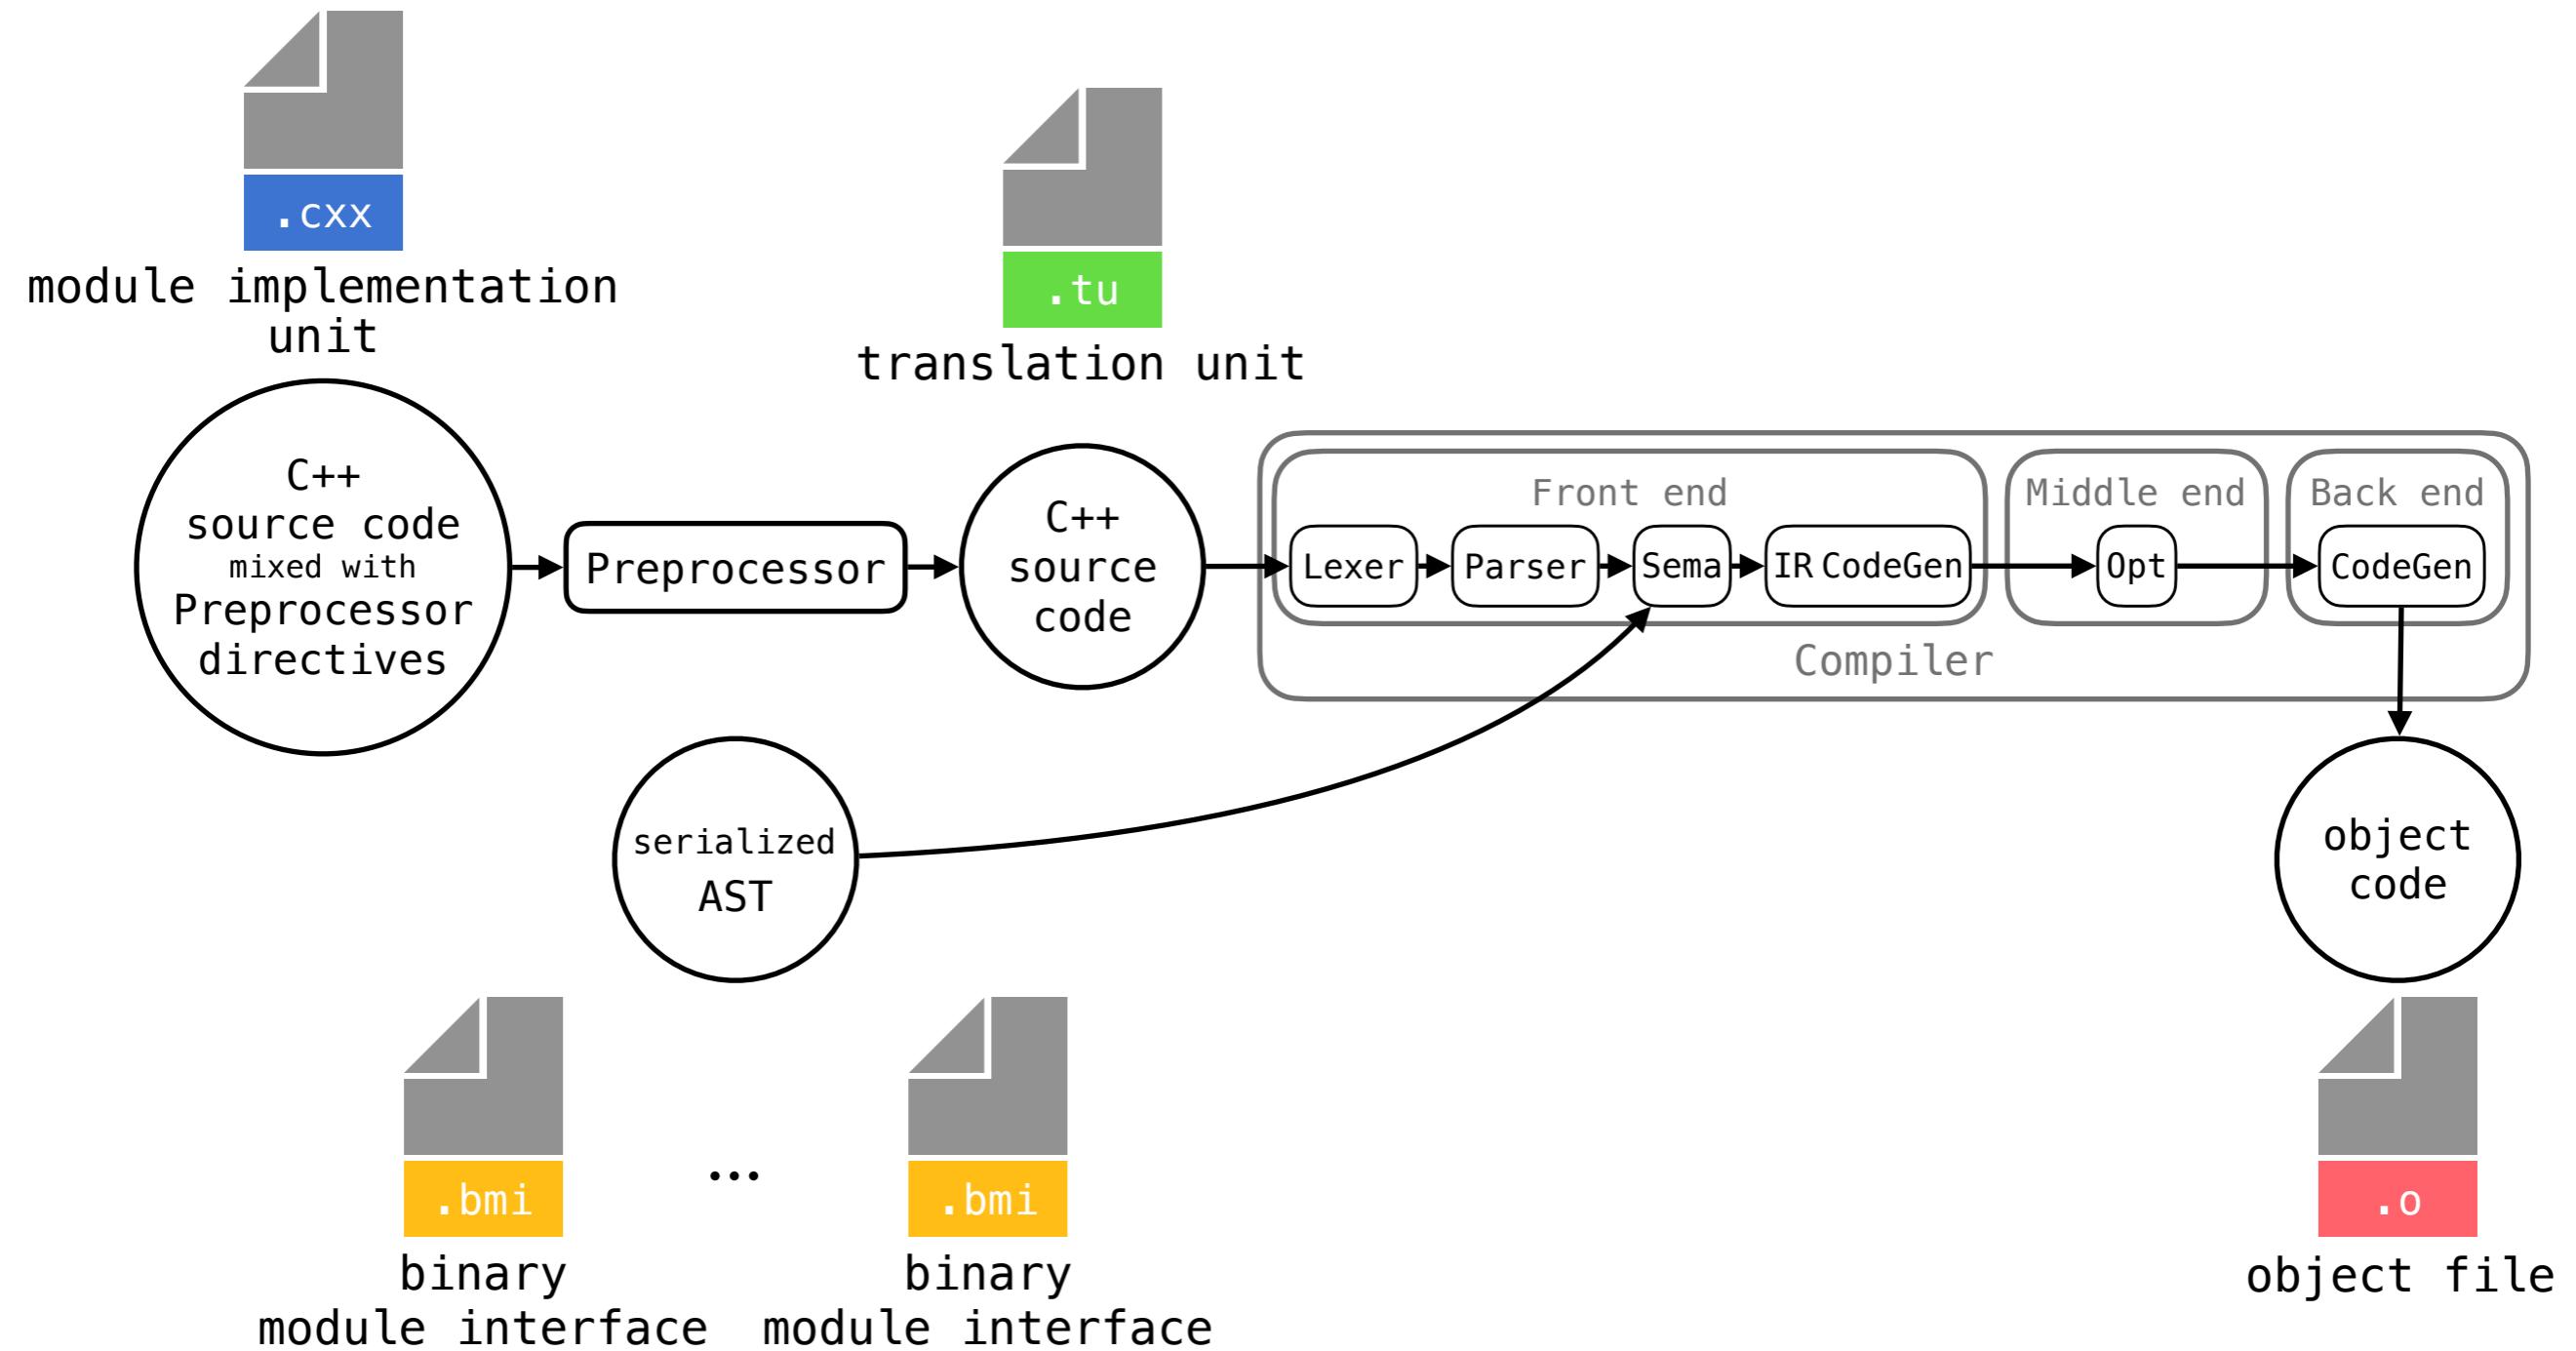
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



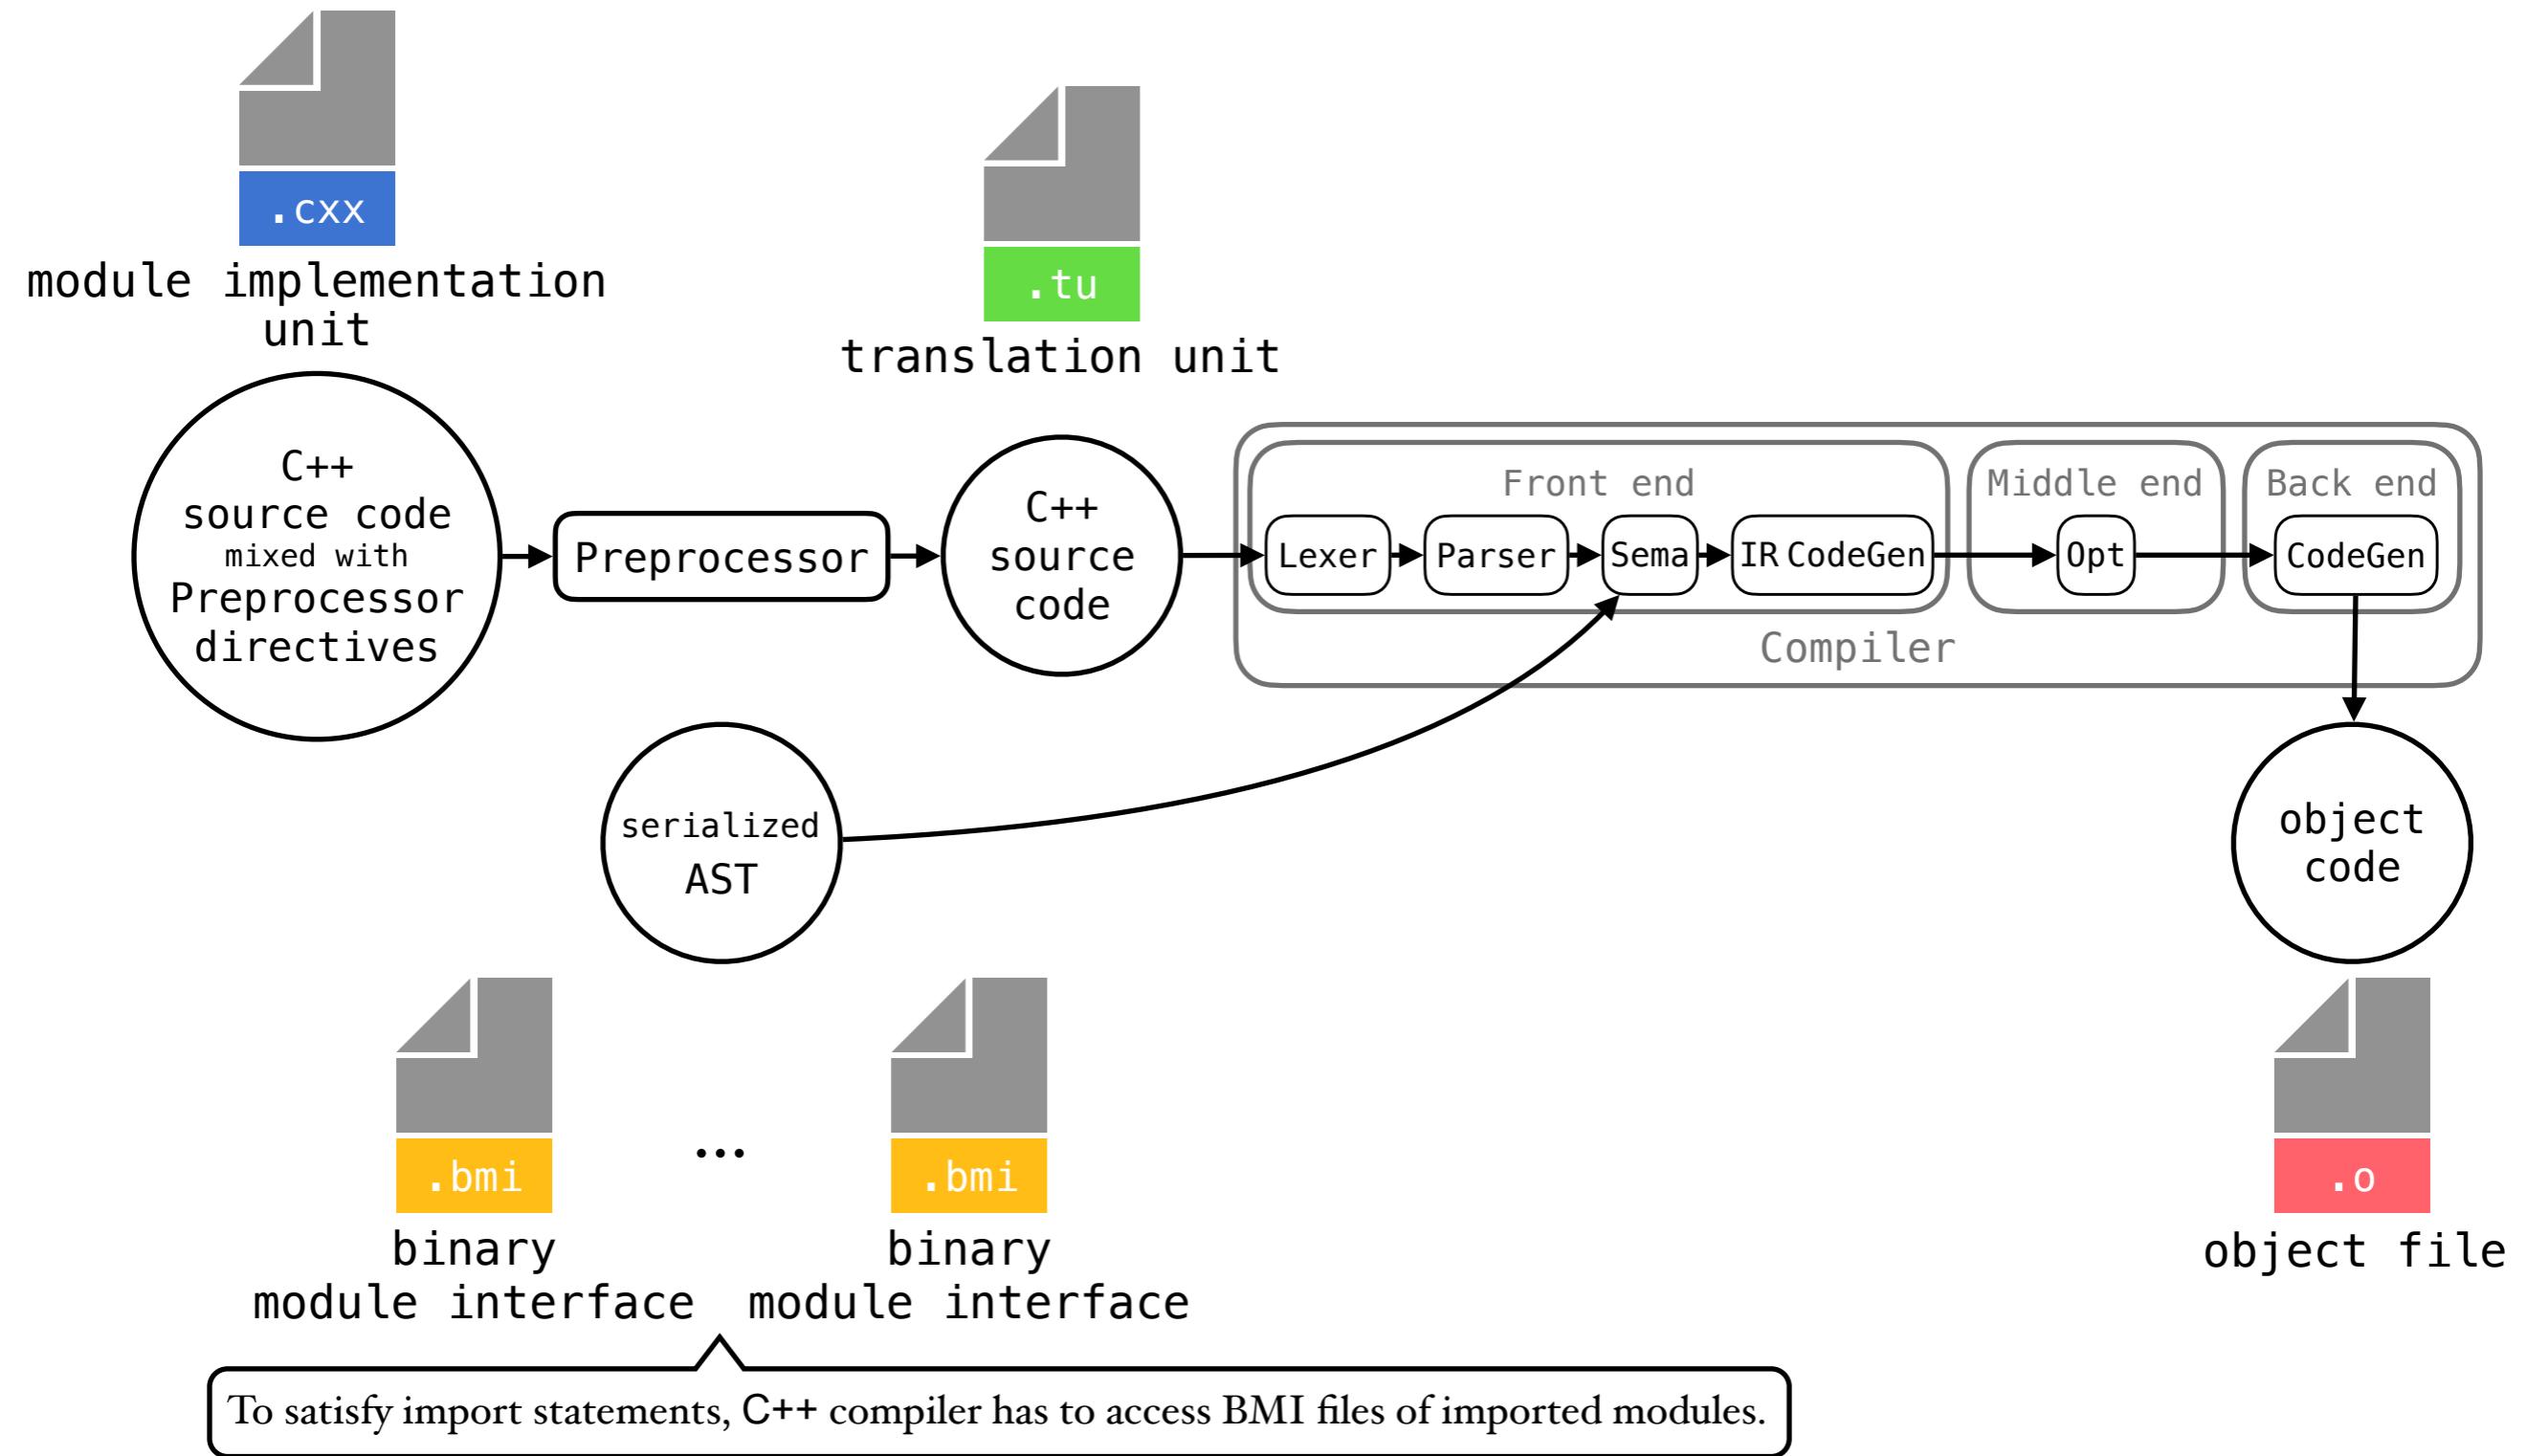
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



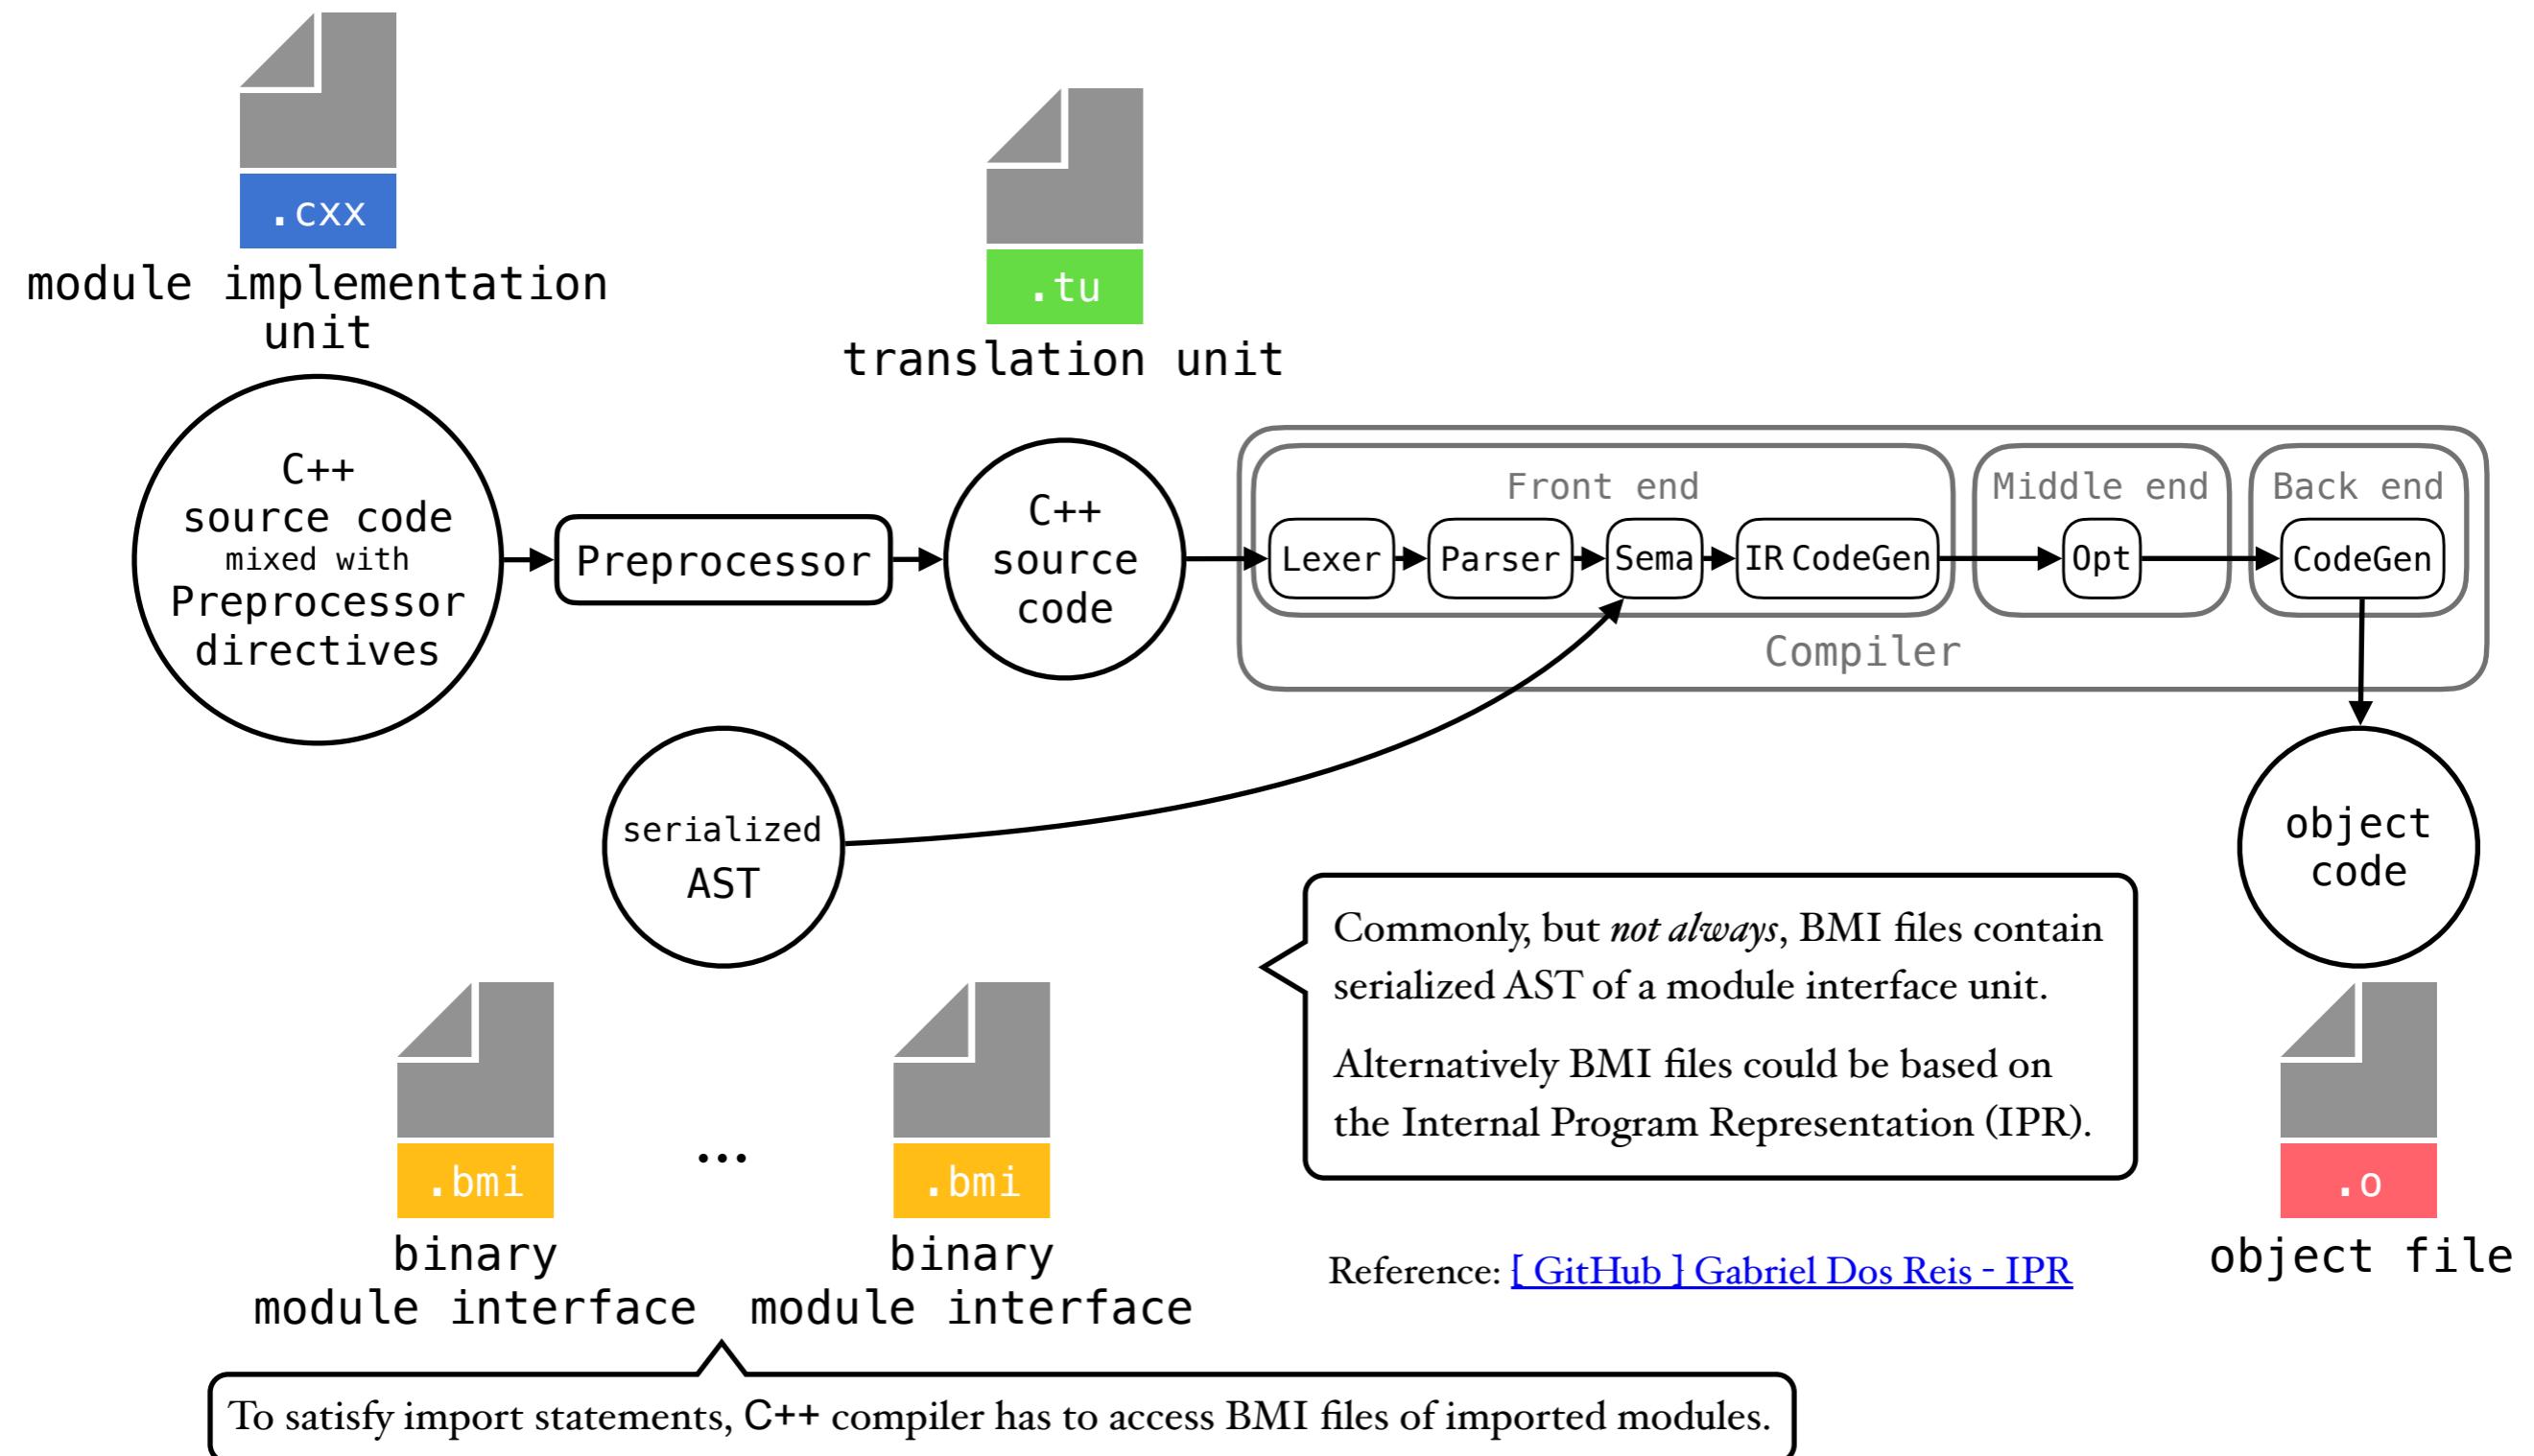
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



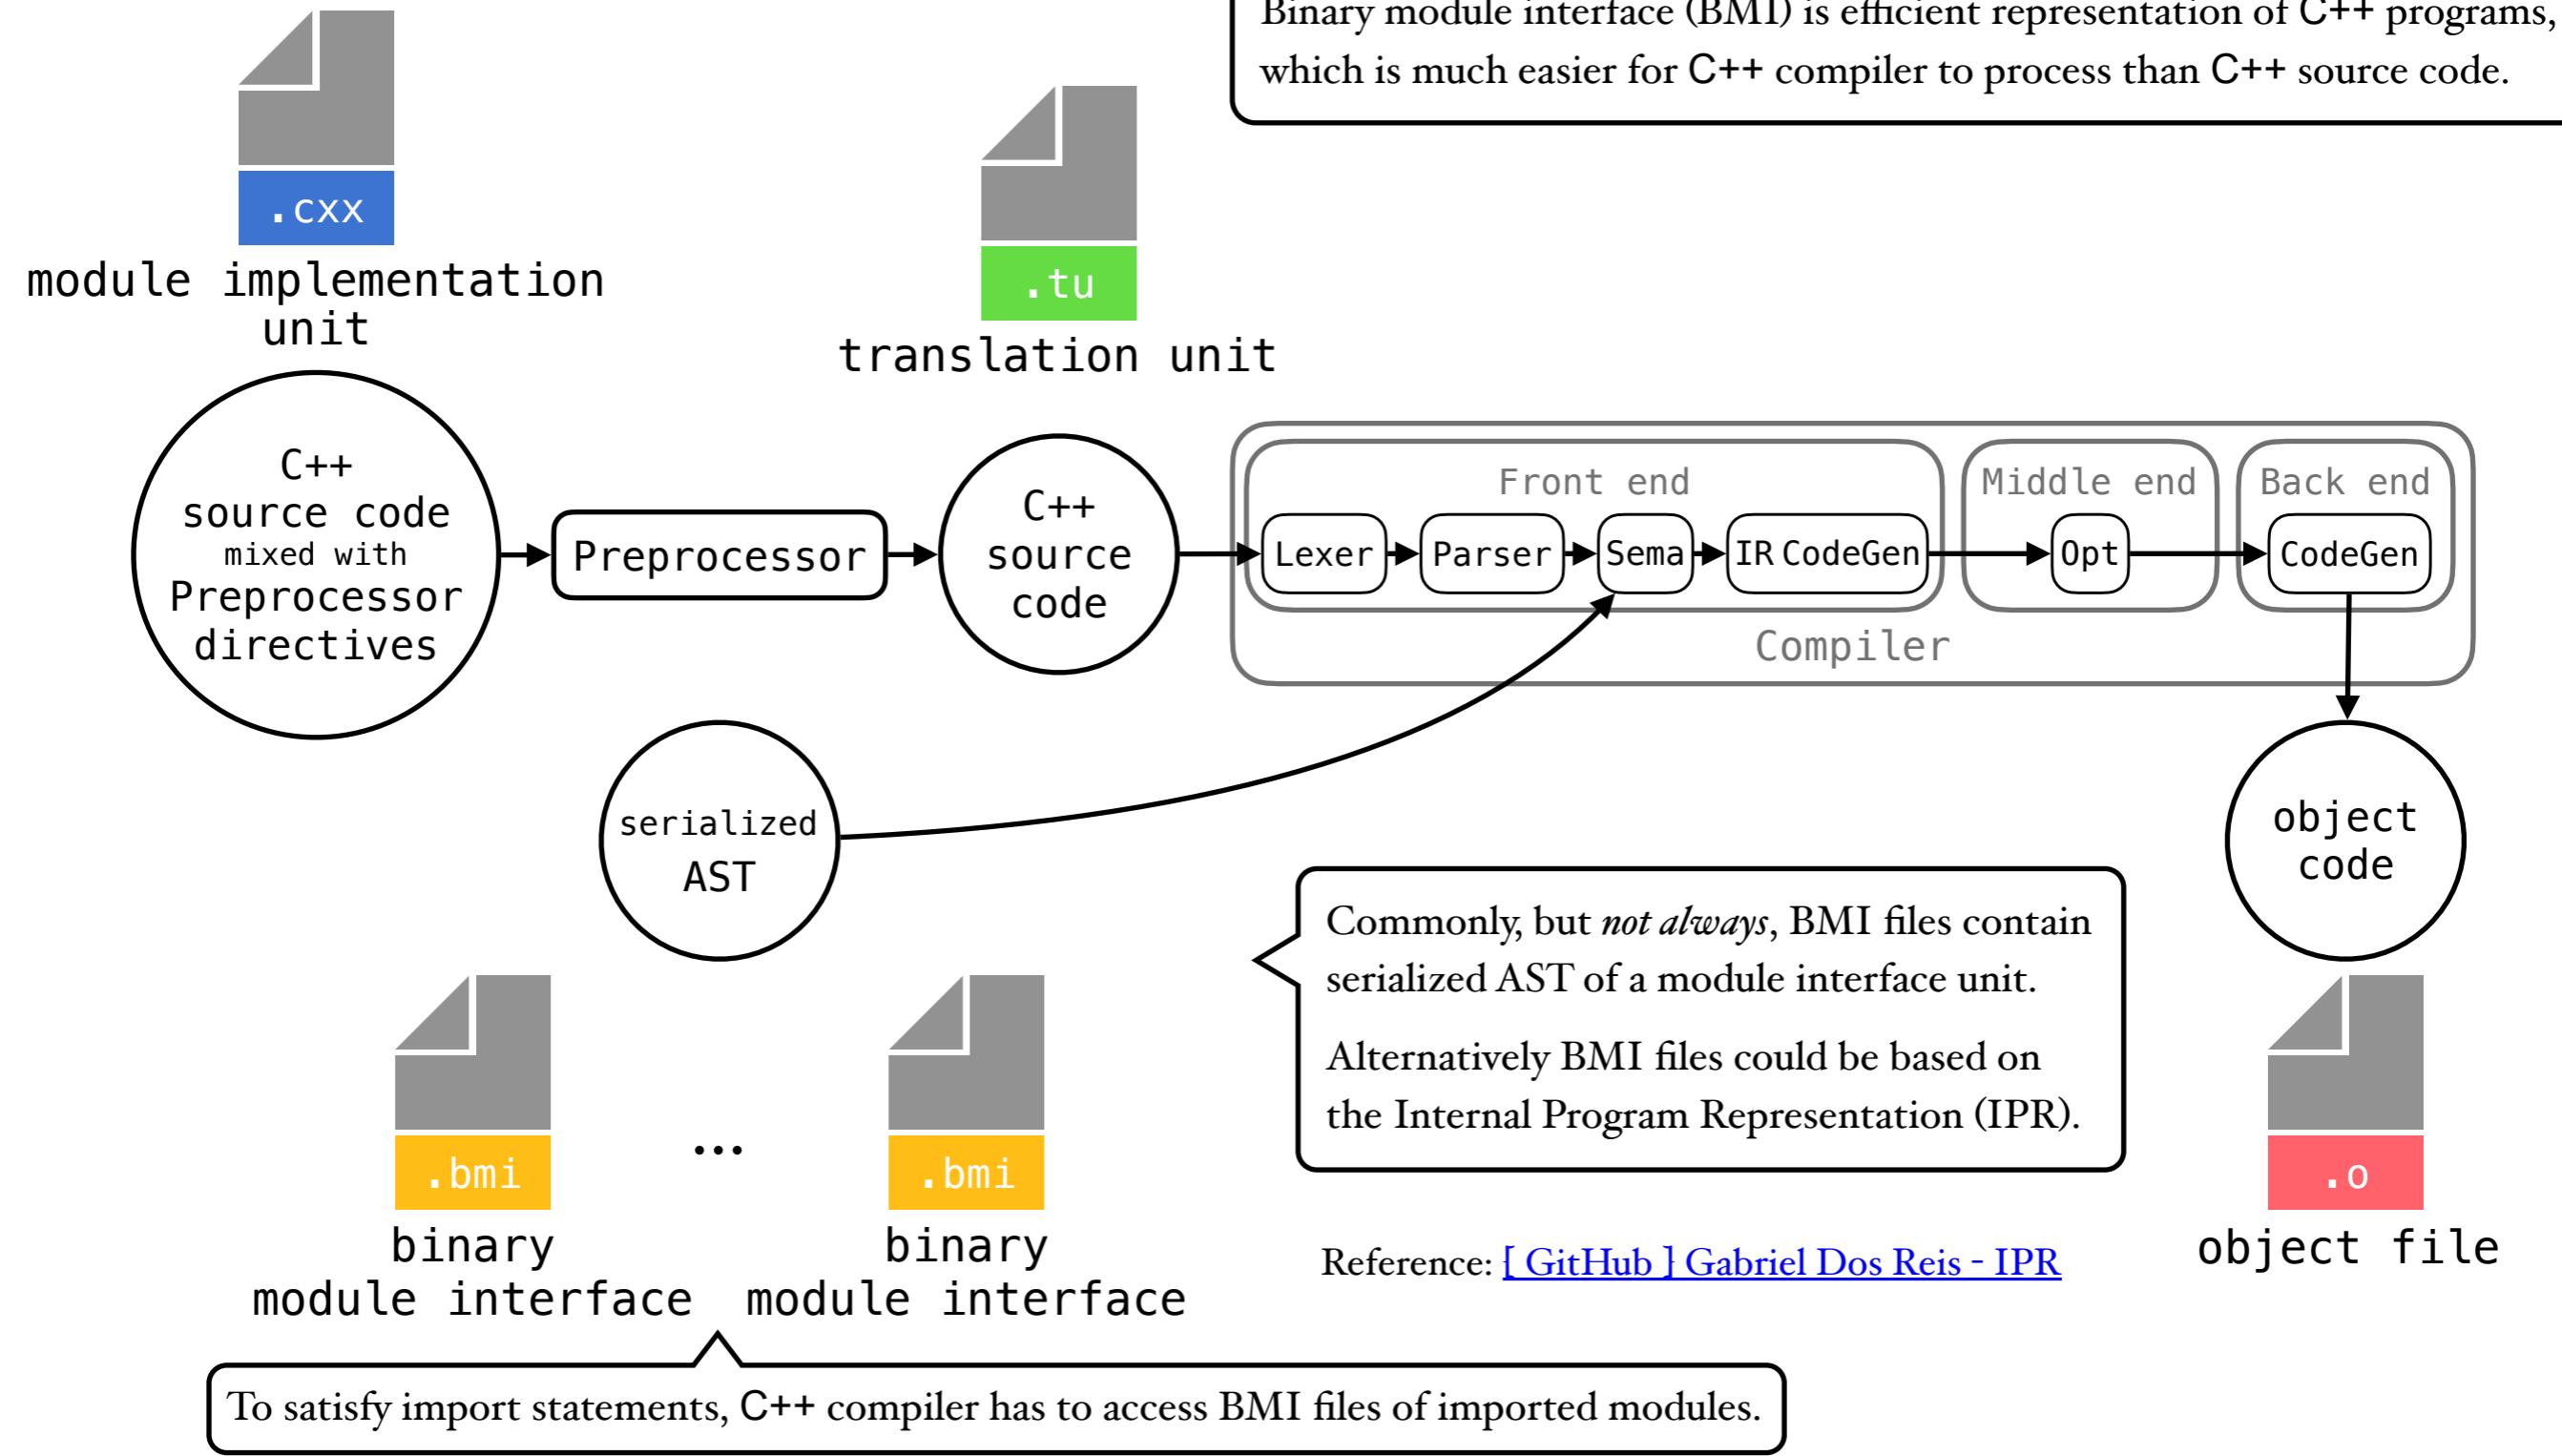
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



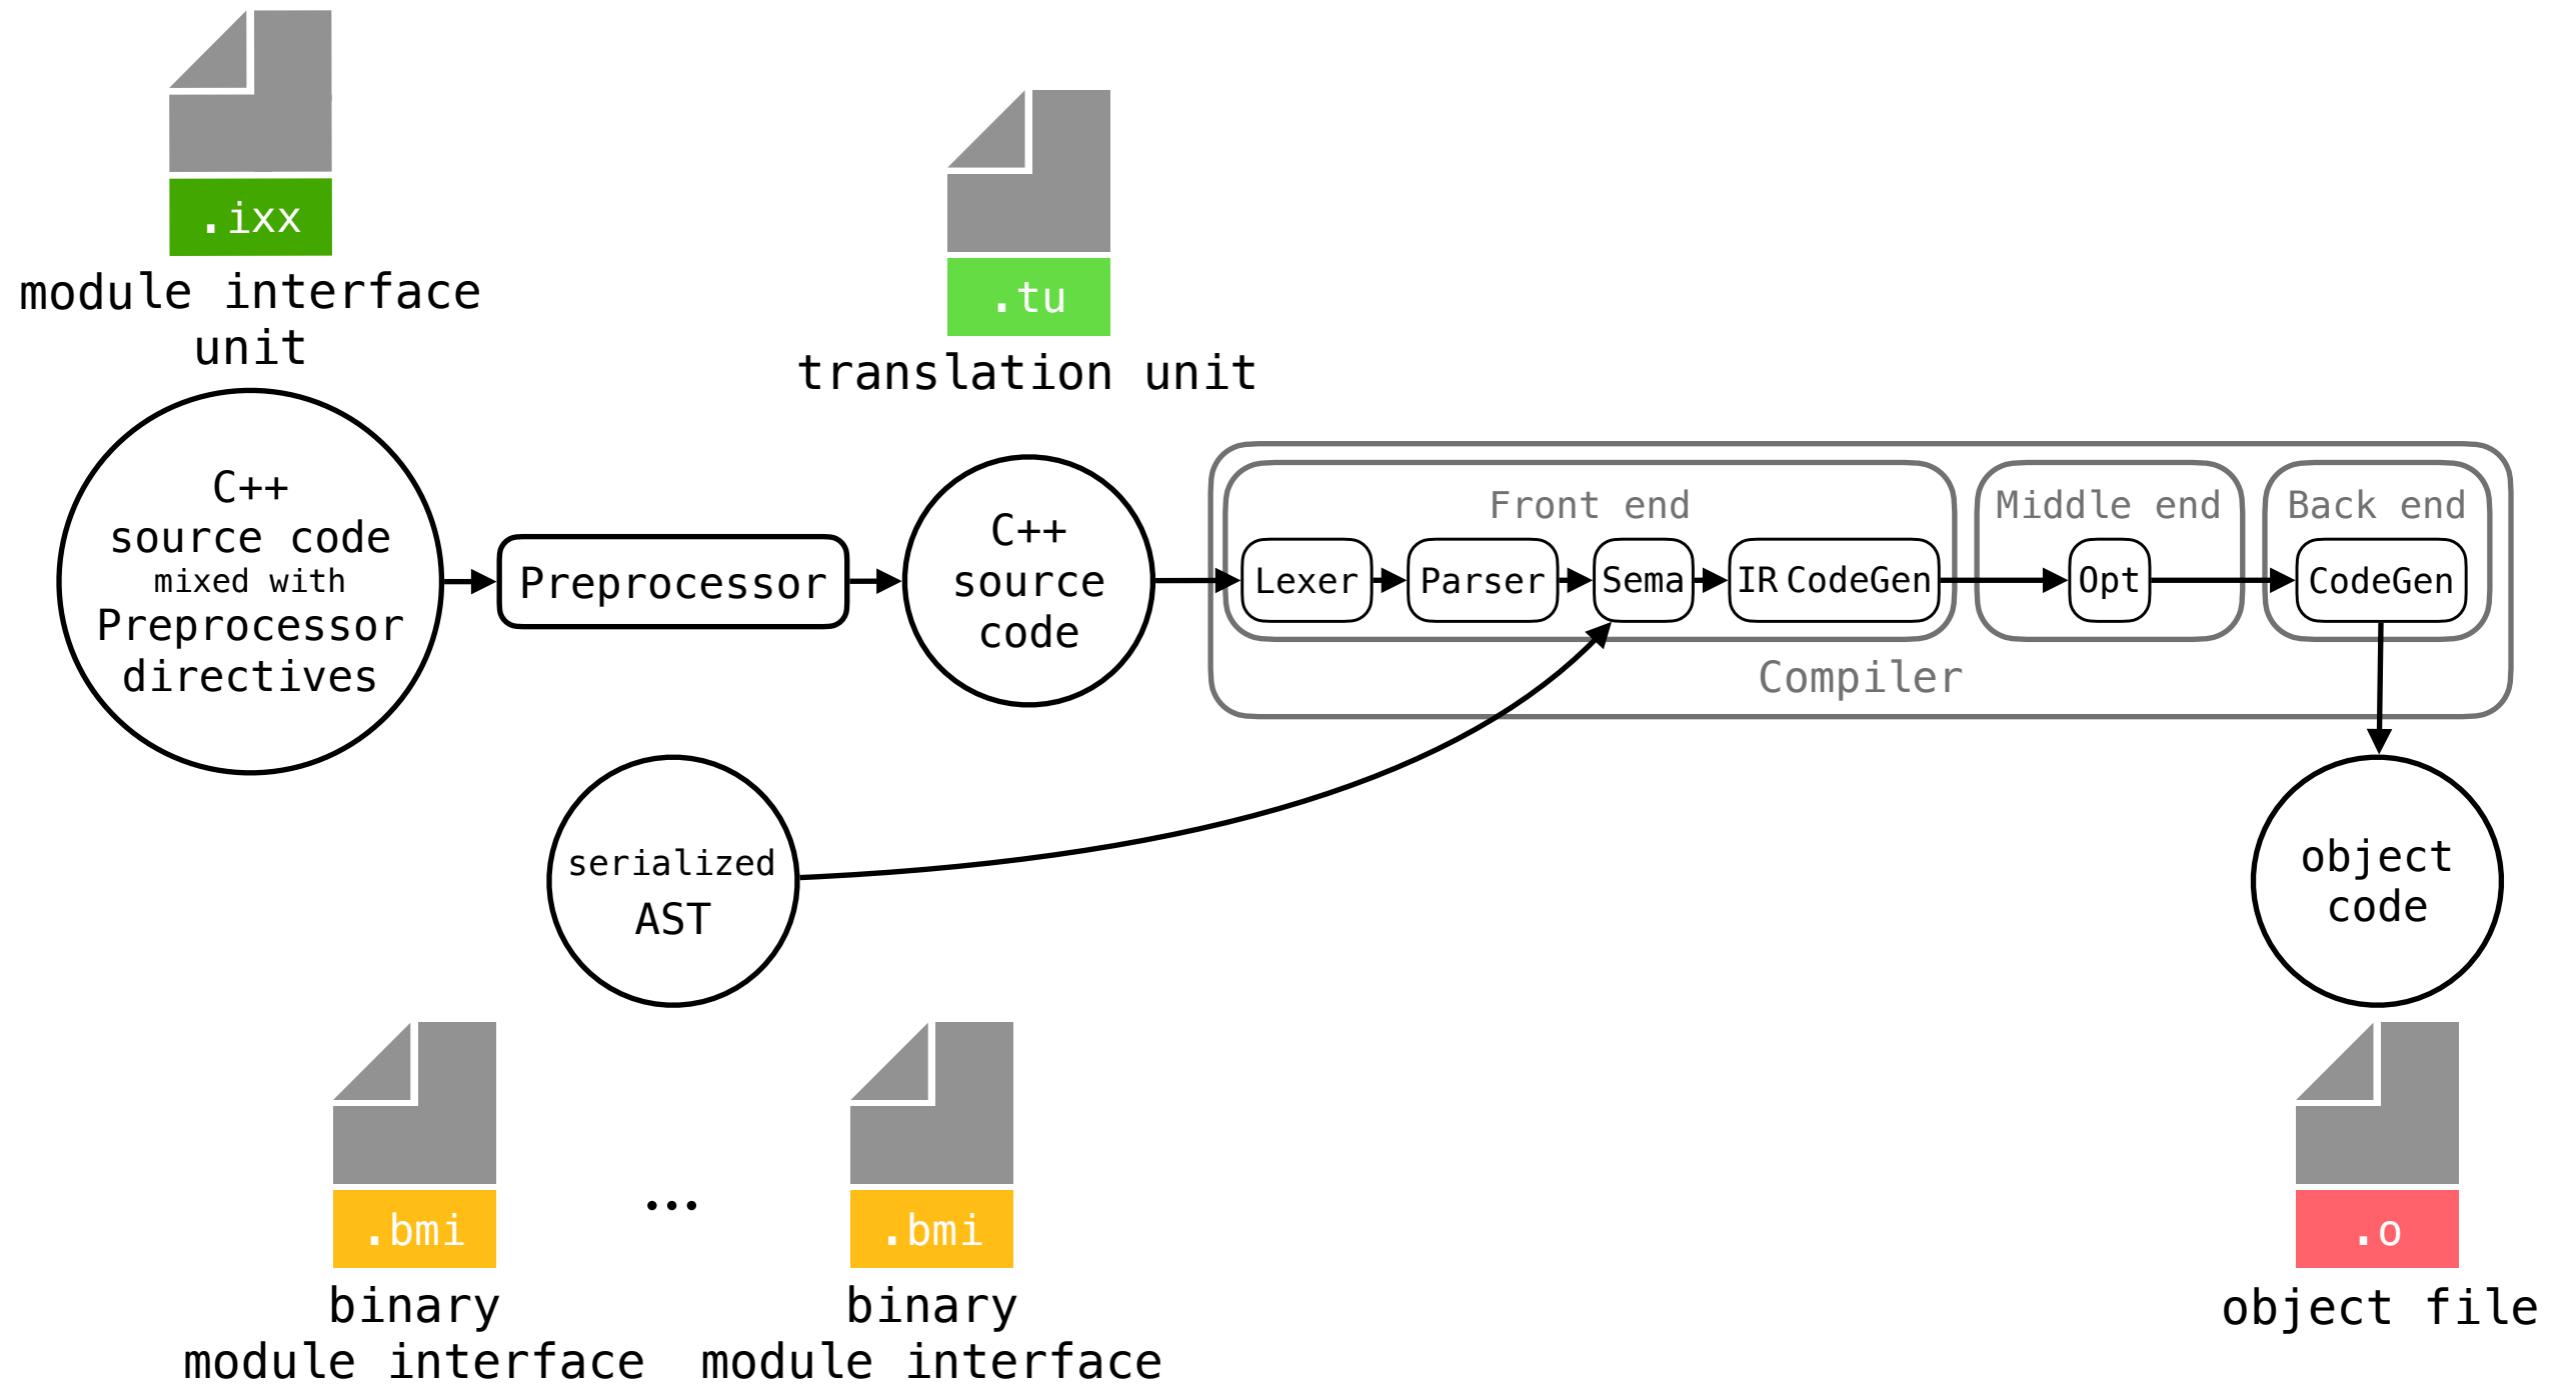
Modular Import Model

Module Implementation Unit - Preprocessing & Compilation



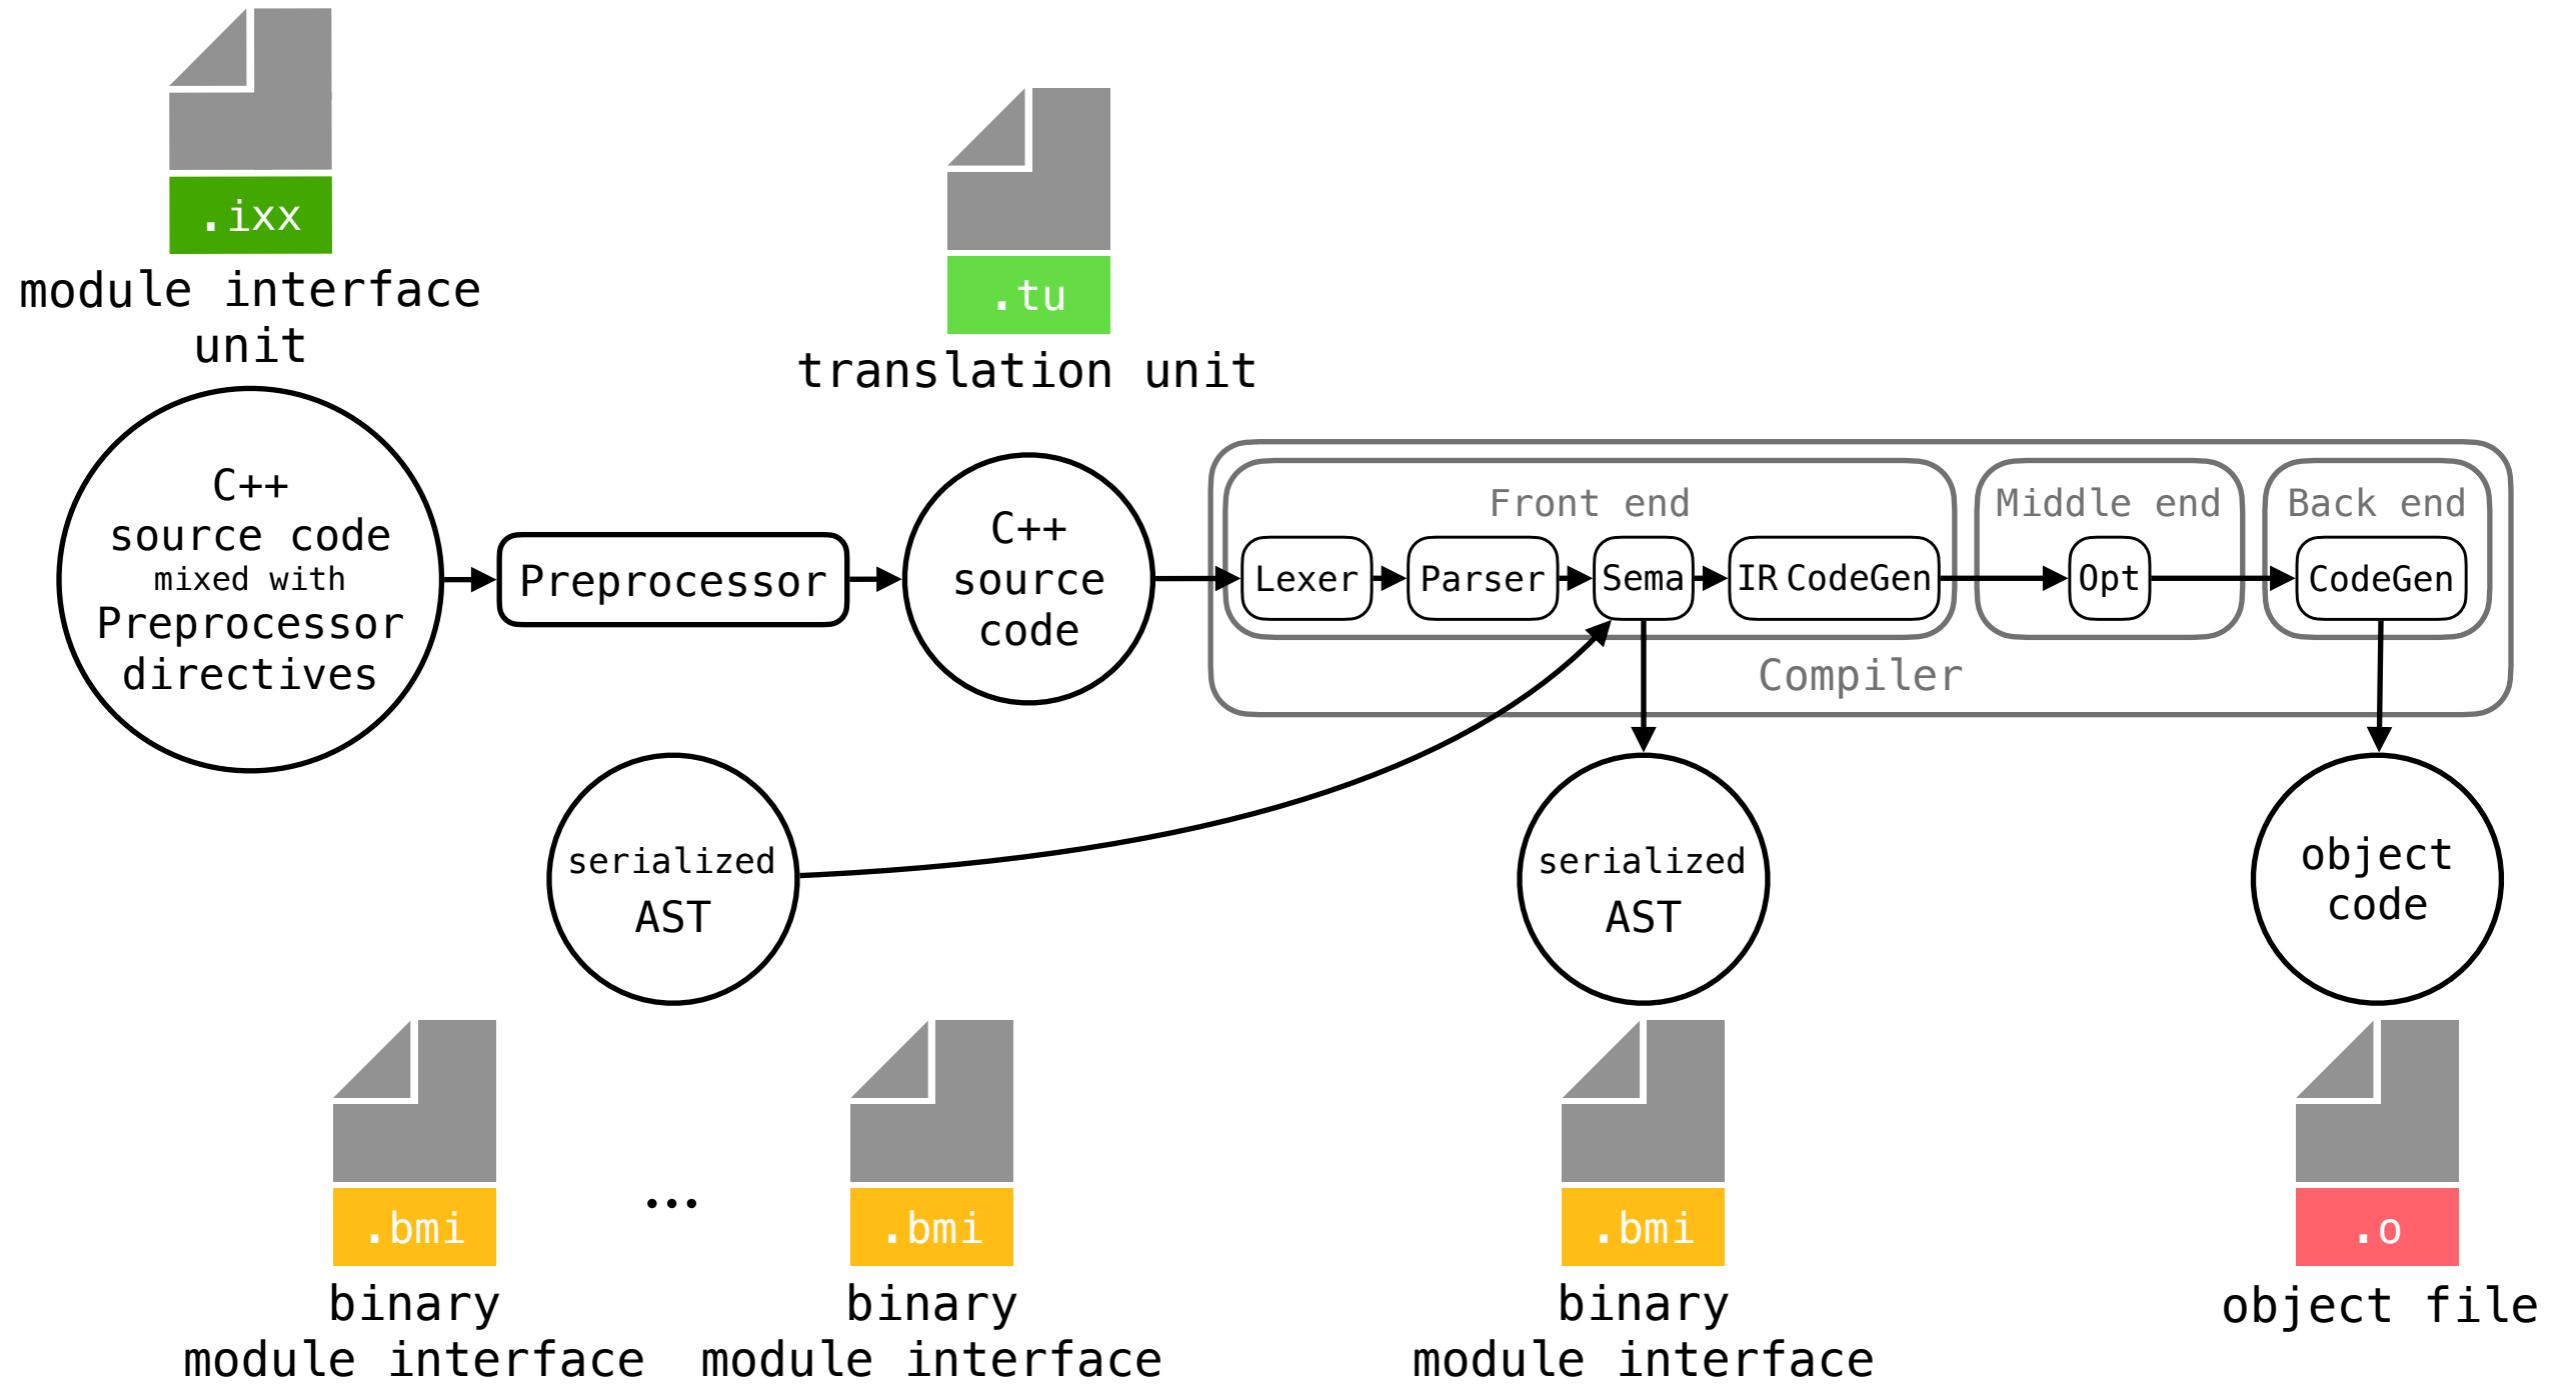
Modular Import Model

Module Interface Unit - Preprocessing & Compilation



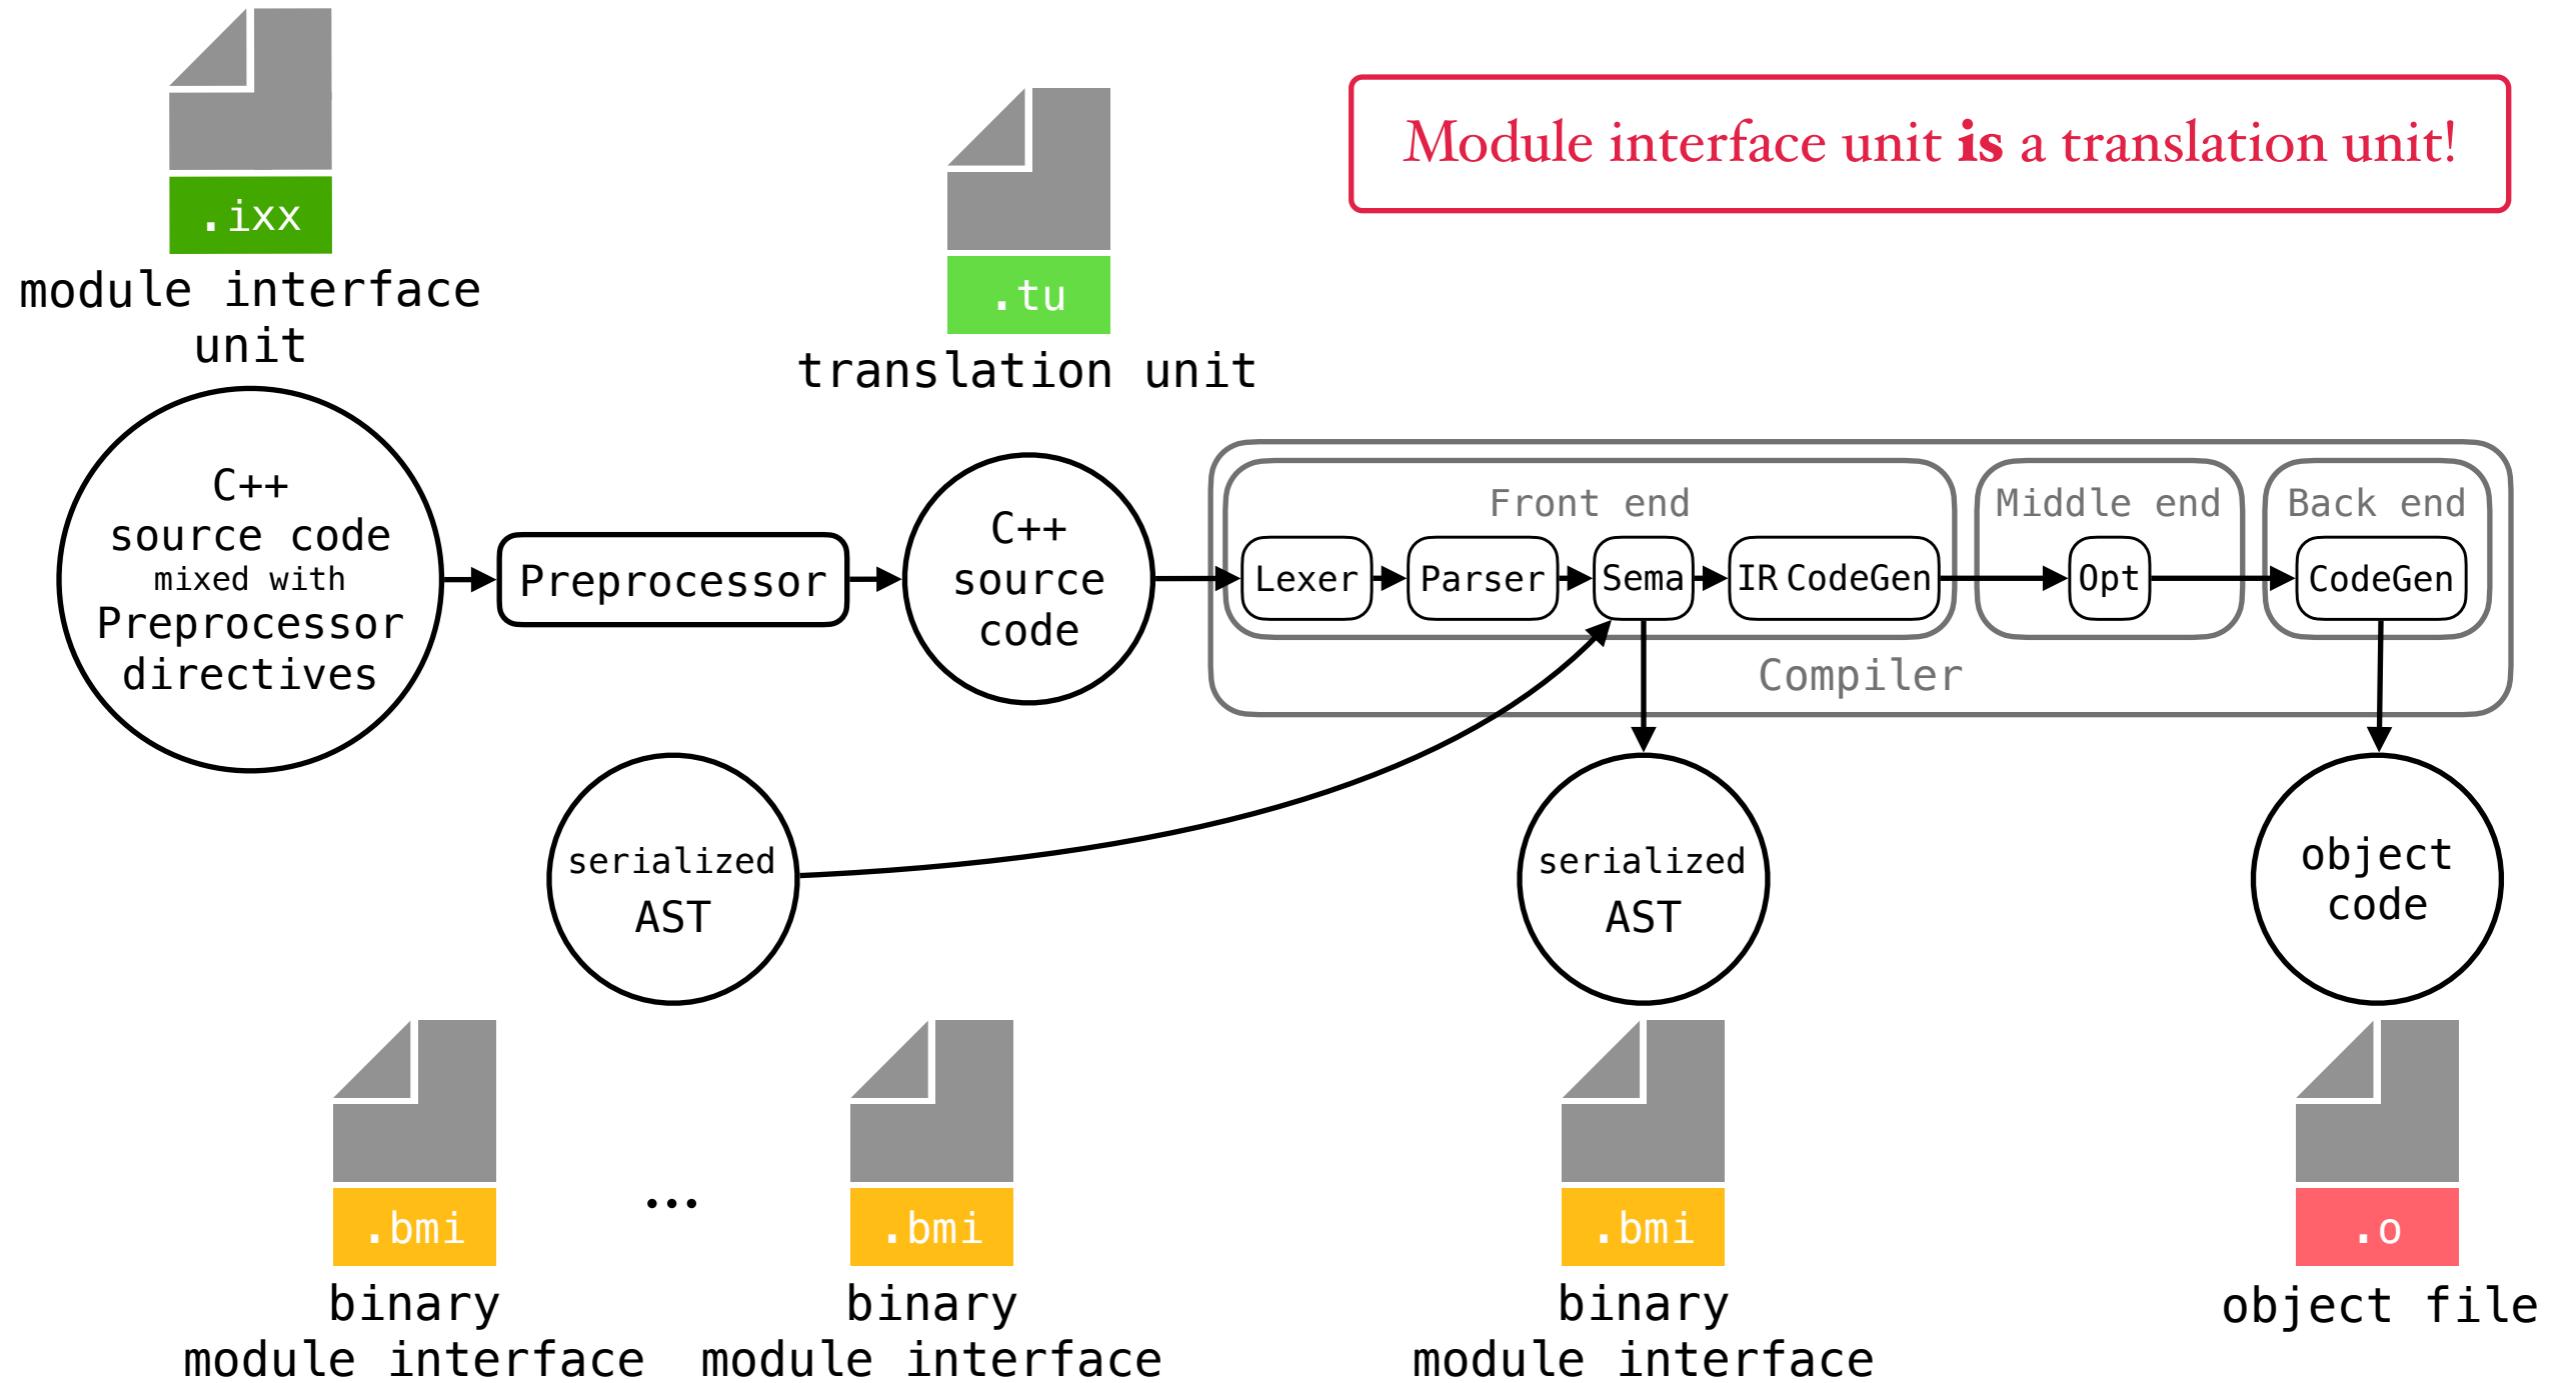
Modular Import Model

Module Interface Unit - Preprocessing & Compilation



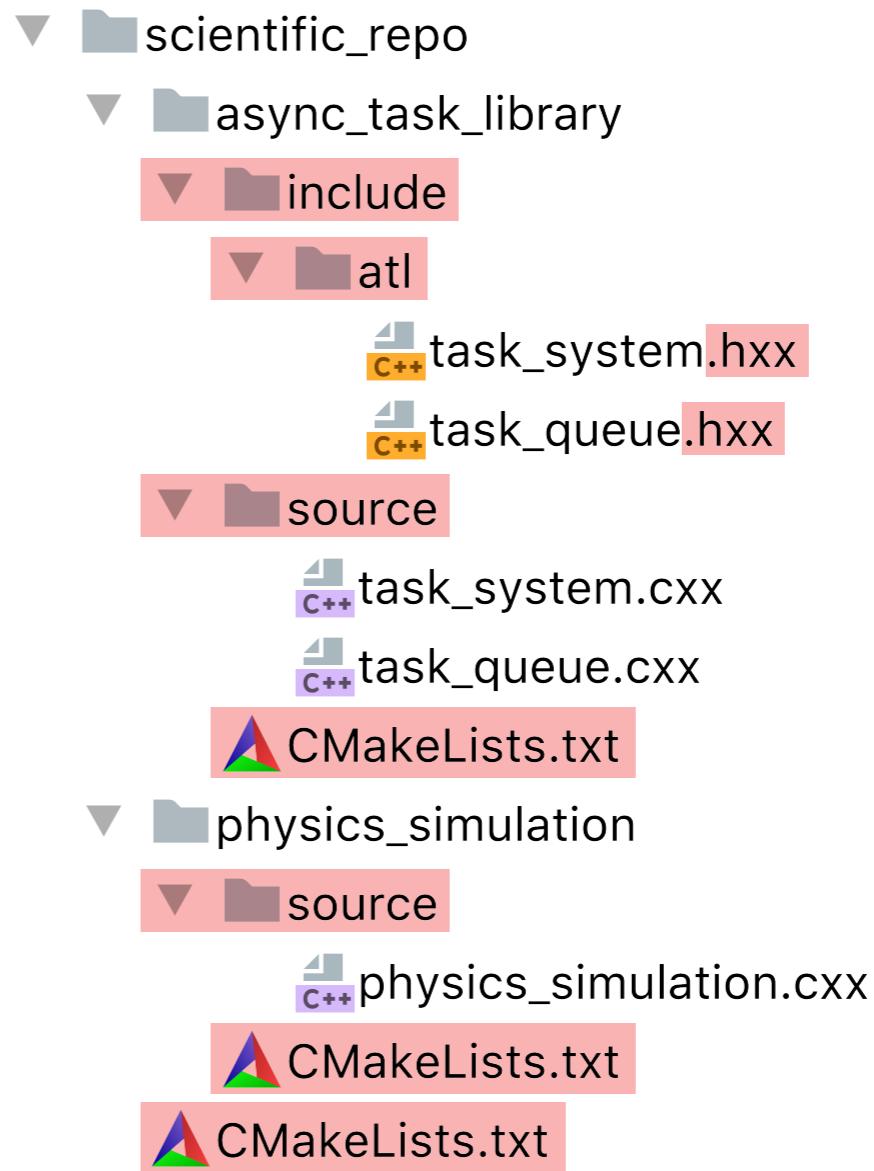
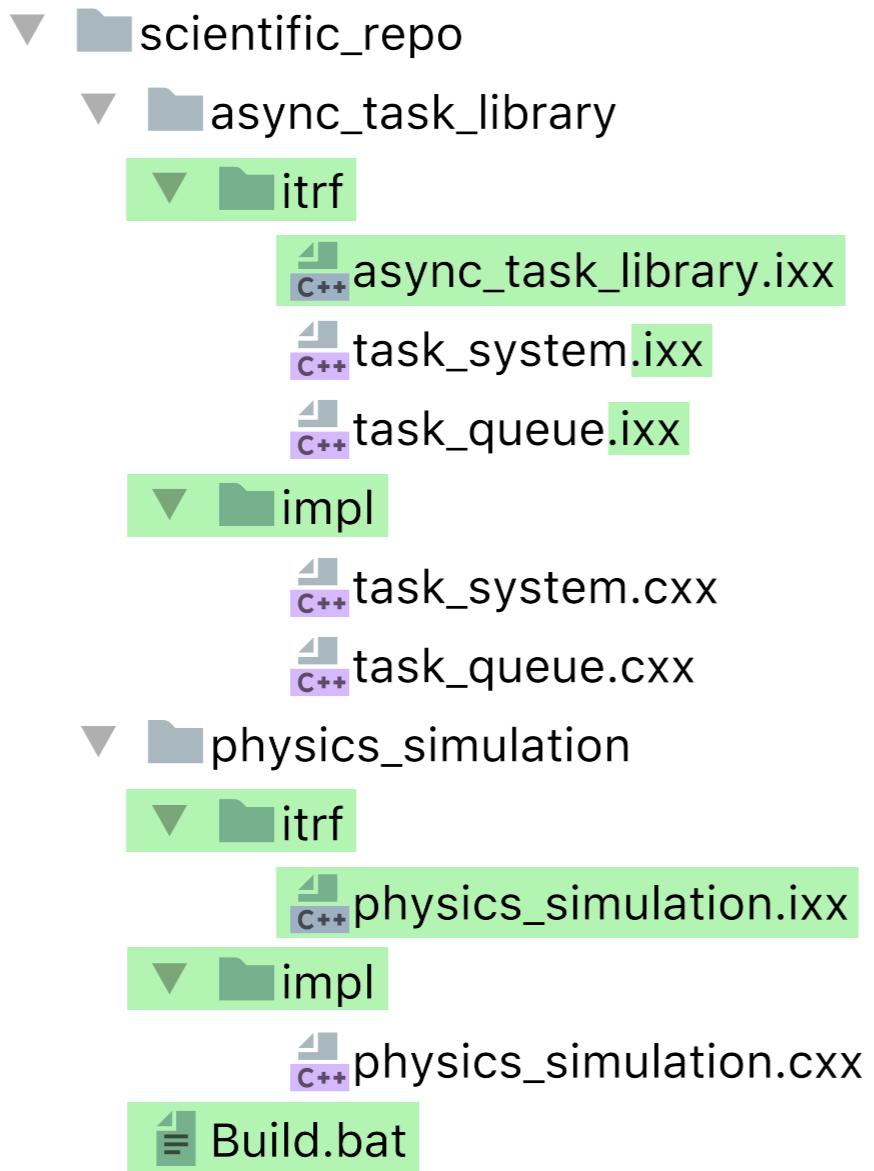
Modular Import Model

Module Interface Unit - Preprocessing & Compilation



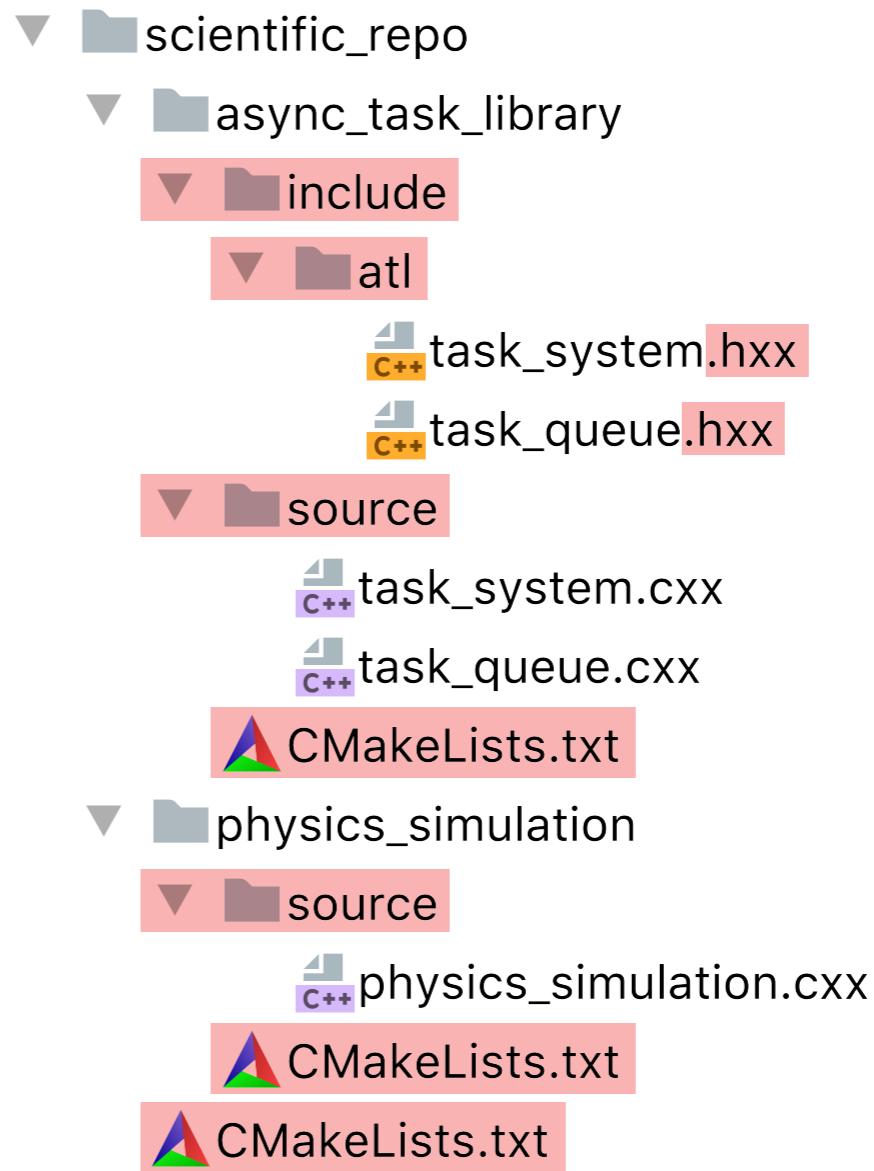
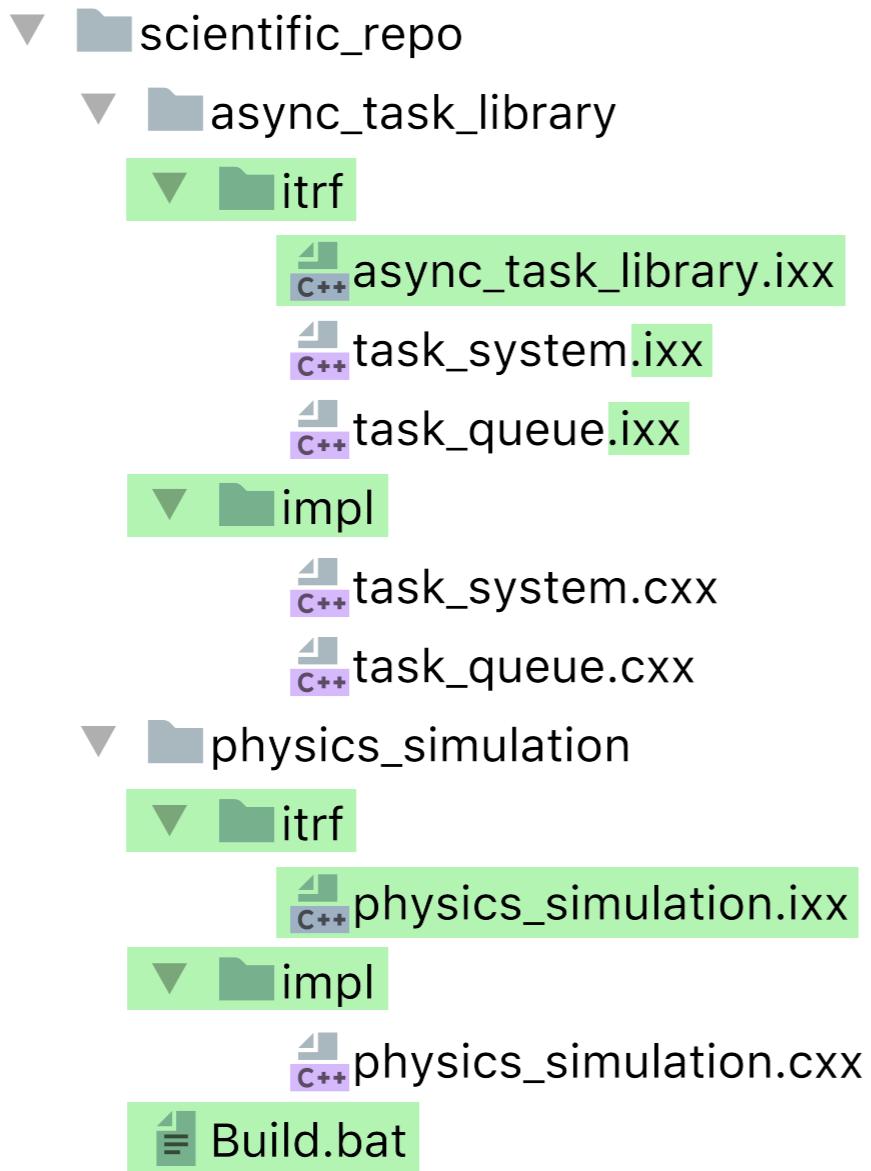
Scientific Repo

Project Organisation



Scientific Repo

Project Organisation



Currently build2 is the only build system supporting C++ modules, so as a temporary WA I've implemented batch script Build.bat, which calls MSVC compiler and linker directly.

Physics Simulation

Module Interface Unit

```
▼科学_repo
  ▼async_task_library
    ▼itrif
      c++async_task_library.ixx
      c++task_system.ixx
      c++task_queue.ixx
    ▼impl
      c++task_system.cxx
      c++task_queue.cxx
  //physics_simulation/itrif/physics_simulation.ixx
  export module phy_sim;
  export auto main () -> int;
  ▼physics_simulation
    ▼itrif
      c++physics_simulation.ixx
    ▼impl
      c++physics_simulation.cxx
  Build.bat
```

Physics Simulation

Module Implementation Unit

```
//physics_simulation/source/physics_simulation.hxx  
#include <atl/task_system.hxx>  
#include <cmath>  
  
namespace phy_sim  
{  
    struct vec3 { float x, y, z; };  
  
    auto length (const vec3& v) -> float { ... }  
  
    template <typename element_type>  
    using array = ...;  
  
    auto read_input () -> array<vec3> { ... }  
  
    auto compute_lengths (const array<vec3>& vectors)  
    -> array<float>  
    { ... }  
  
    auto print_results (const array<float>& lengths)  
    -> void  
    { ... }  
}  
  
auto main () -> int  
{ ... }
```

```
//physics_simulation/impl/physics_simulation.hxx  
module phy_sim;  
  
import atl;  
import std.core;  
  
namespace phy_sim  
{  
    struct vec3 { float x, y, z; };  
  
    auto length (const vec3& v) -> float { ... }  
  
    template <typename element_type>  
    using array = ...;  
  
    auto read_input () -> array<vec3> { ... }  
  
    auto compute_lengths (const array<vec3>& vectors)  
    -> array<float>  
    { ... }  
  
    auto print_results (const array<float>& lengths)  
    -> void  
    { ... }  
}  
  
auto main () -> int  
{ ... }
```

Async Task Library

Module Interface Unit

```
▼ 📂scientific_repo
  ▼ 📂async_task_library
    ▼ 📂itrif
      📄async_task_library.ixx
      📄task_system.ixx
      📄task_queue.ixx
    ▼ 📂impl
      📄task_system.cxx
      📄task_queue.cxx
  //async_task_library/itrif/async_task_library.ixx
  export module atl;
  export import :task_system;

  ▼ 📂physics_simulation
    ▼ 📂itrif
      📄physics_simulation.ixx
    ▼ 📂impl
      📄physics_simulation.cxx
  📄Build.bat
```

Async Task Library

Module Interface Unit

```
▼ 📂 scientific_repo
  ▼ 📂 async_task_library
    ▼ 📂 itrif
      📄 async_task_library.ixx
      📄 task_system.ixx
      📄 task_queue.ixx
    ▼ 📂 impl
      📄 task_system.cxx
      📄 task_queue.cxx
  ▼ 📂 physics_simulation
    ▼ 📂 itrif
      📄 physics_simulation.ixx
    ▼ 📂 impl
      📄 physics_simulation.cxx
  📄 Build.bat
```

```
//async_task_library/itrif/async_task_library.ixx
export module atl;
export import :task_system;
```

Define the interface of the whole atl module,
by providing to clients interface of the
atl:task_system module partition, that is reexport it.

Async Task Library

Task Queue - Module Partition Interface Unit

```
//async_task_library/include/atl/task_queue.hxx
```

```
#ifndef ATL_TASK_QUEUE  
#define ATL_TASK_QUEUE
```

```
#include <deque>  
#include <mutex>  
#include <condition_variable>  
#include <functional>  
#include <utility>  
#include <optional>
```

```
namespace atl  
{  
    class task_queue  
    {  
        std::deque<std::function<void()>> tasks;  
        std::mutex  
        std::condition_variable  
        bool done { false };  
  
    public:  
        auto finish () -> void;  
  
        auto pop () -> std::optional<std::function<void()>>;  
  
        template <typename callable>  
        auto push (callable&& task) -> void { ... }  
    };  
}  
#endif
```

```
//async_task_library/itrf/task_queue.ixx
```

```
export module atl:task_queue;
```

```
import std.core;  
import std.threading;
```

```
namespace atl  
{  
    class task_queue  
    {  
        std::deque<std::function<void()>> tasks;  
        std::mutex  
        std::condition_variable  
        bool done { false };  
  
    public:  
        auto finish () -> void;  
  
        auto pop () -> std::optional<std::function<void()>>;  
  
        template <typename callable>  
        auto push (callable&& task) -> void { ... }  
    };  
}
```

Async Task Library

Tak Queue - Module Partition Implementation Unit

```
//async_task_library/source/task_queue.cxx
#include <atl/task_queue.hxx>

namespace atl
{
    auto task_queue::finish() -> void
    { ... }

    auto task_queue::pop ()
    -> std::optional<std::function<void()>>
    { ... }
}
```

```
//async_task_library/impl/task_queue.cxx
module atl:task_queue;

namespace atl
{
    auto task_queue::finish() -> void
    { ... }

    auto task_queue::pop ()
    -> std::optional<std::function<void()>>
    { ... }
}
```

Async Task Library

Task System - Module Partition Interface Unit

```
//async_task_library/include/atl/task_system.hxx
#ifndef ATL_TASK_SYSTEM
#define ATL_TASK_SYSTEM

#include <atl/task_queue.hxx>
#include <vector>
#include <thread>

namespace atl
{
    class task_system
    {
        private:
            std::vector<std::thread> threads;
            task_queue queue;

        public:
            task_system();
            ~task_system();

            template <typename callable>
            auto async (callable&& task) -> void
            { ... }
    };
}

#endif
```

```
//async_task_library/itrf/task_system.ixx
export module atl:task_system;

import :task_queue;

namespace atl
{
    export class task_system
    {
        private:
            std::vector<std::thread> threads;
            task_queue queue;

        public:
            task_system();
            ~task_system();

            template <typename callable>
            auto async (callable&& task) -> void
            { ... }
    };
}
```

Async Task Library

Task System - Module Partition Interface Unit

```
//async_task_library/include/atl/task_system.hxx  
  
#ifndef ATL_TASK_SYSTEM  
#define ATL_TASK_SYSTEM  
  
#include <atl/task_queue.hxx>  
#include <vector>  
#include <thread>  
  
namespace atl  
{  
    class task_system  
    {  
        private:  
            std::vector<std::thread> threads;  
            task_queue queue;  
  
        public:  
            task_system();  
            ~task_system();  
  
            template <typename callable>  
            auto async (callable&& task) -> void  
            { ... }  
    };  
}  
  
#endif
```

```
//async_task_library/itrf/task_system.ixx  
  
export module atl:task_system;  
  
import :task_queue;  
  
Internal dependency inside module between module partitions.  
  
namespace atl  
{  
    export class task_system  
    {  
        private:  
            std::vector<std::thread> threads;  
            task_queue queue;  
  
        public:  
            task_system();  
            ~task_system();  
  
            template <typename callable>  
            auto async (callable&& task) -> void  
            { ... }  
    };  
}
```

Async Task Library

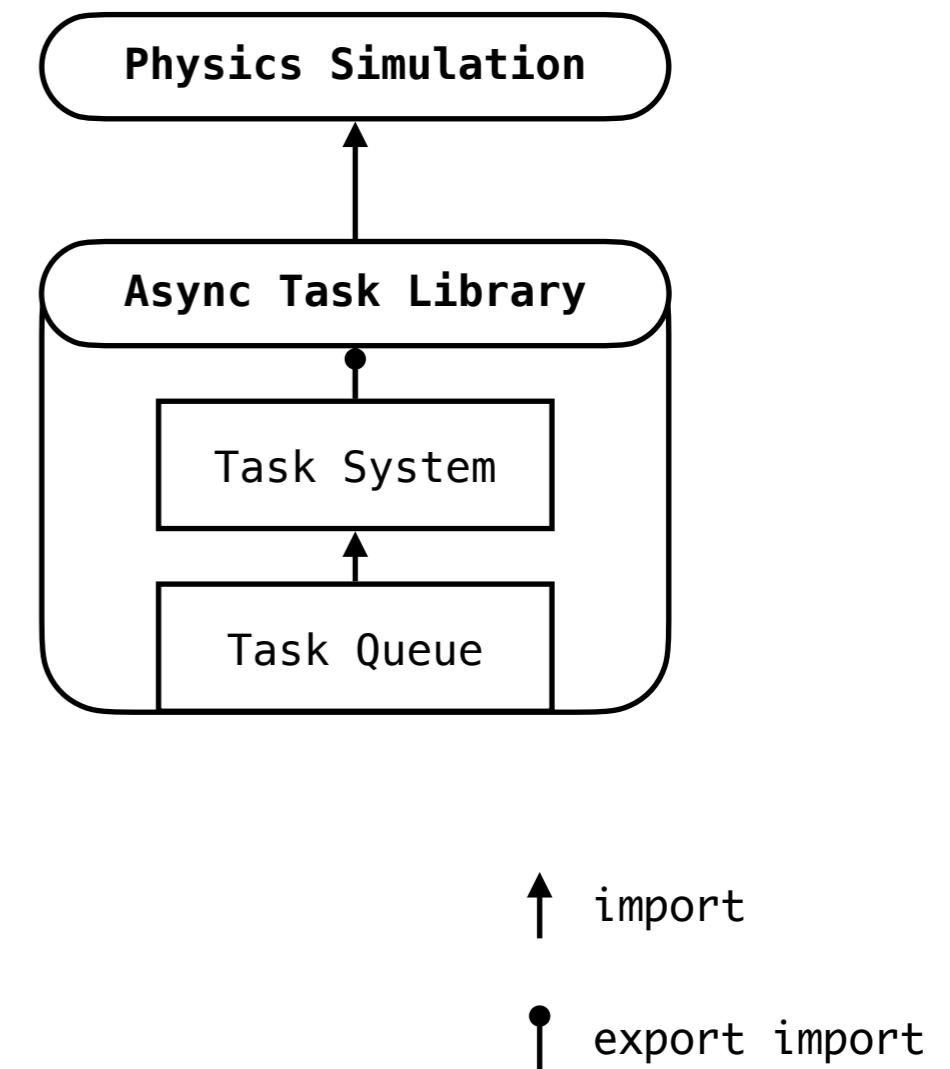
Task System - Module Partition Implementation Unit

```
//async_task_library/source/task_system.hxx  
#include <atl/task_system.hxx>  
  
namespace atl  
{  
    namespace  
    {  
        auto run_thread (task_queue& queue)  
        -> void { ... }  
    }  
  
    task_system::task_system ()  
    { ... }  
  
    task_system::~task_system ()  
    { ... }  
}
```

```
//async_task_library/impl/task_system.hxx  
module atl:task_system;  
  
namespace atl  
{  
    namespace  
    {  
        auto run_thread (task_queue& queue)  
        -> void { ... }  
    }  
  
    task_system::task_system ()  
    { ... }  
  
    task_system::~task_system ()  
    { ... }  
}
```

Modular Import Model

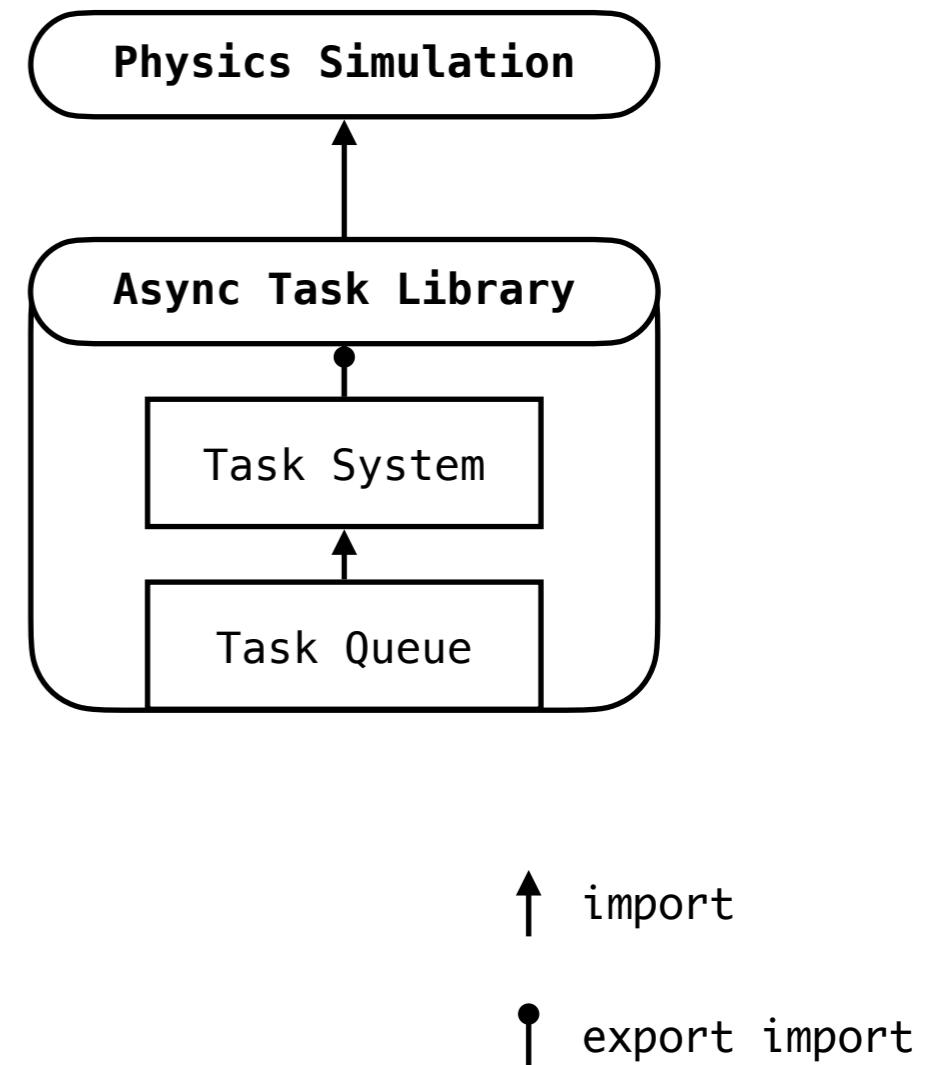
Dependency Graph



Modular Import Model

Dependency Graph

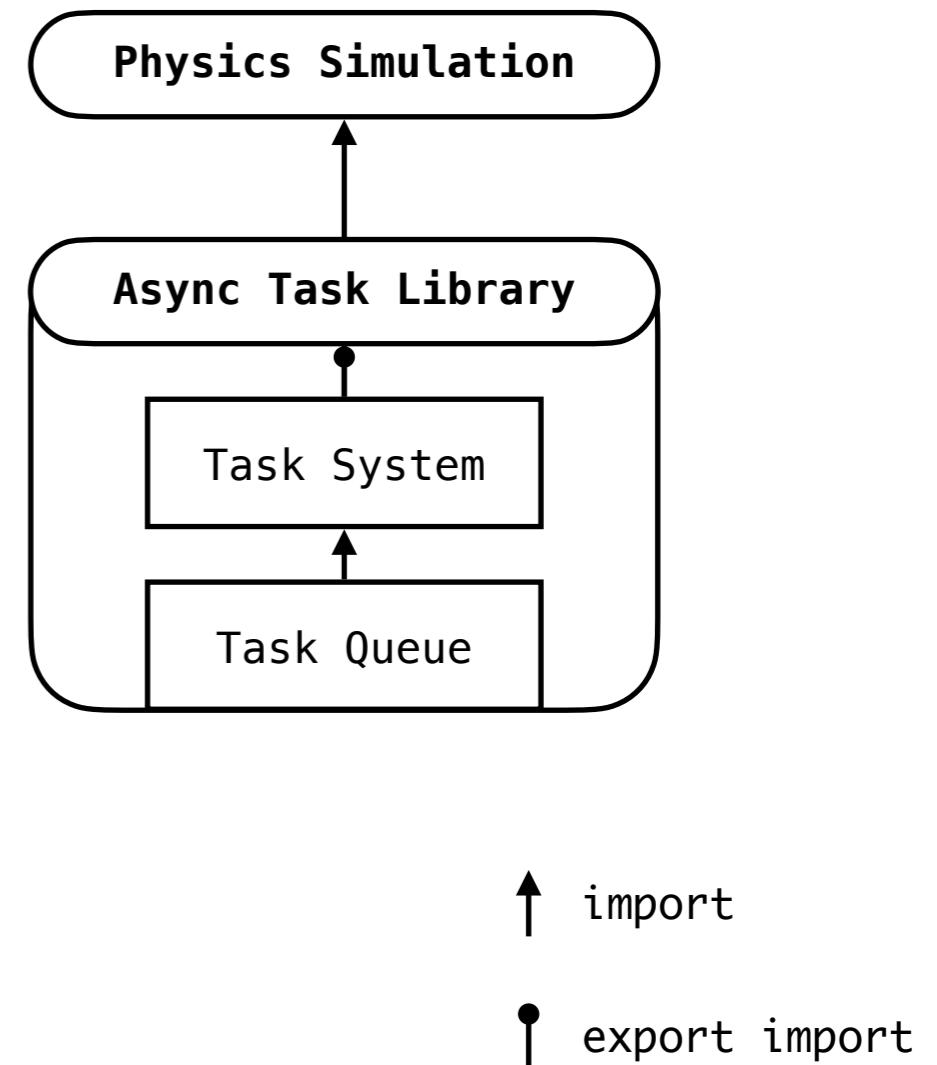
- Import statements form dependency graph between modules.



Modular Import Model

Dependency Graph

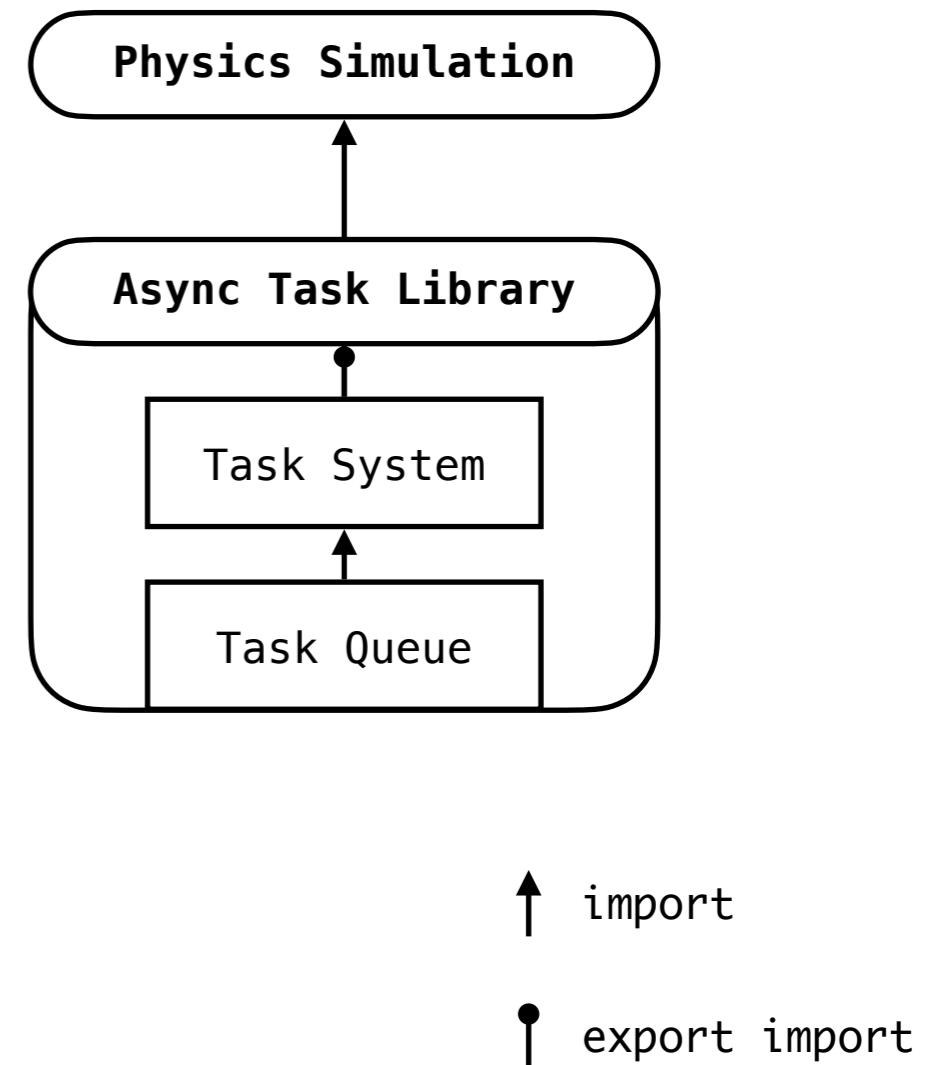
- Import statements form dependency graph between modules.
- Module dependencies cannot be cyclic, which C++ compilers enforce at compile-time.

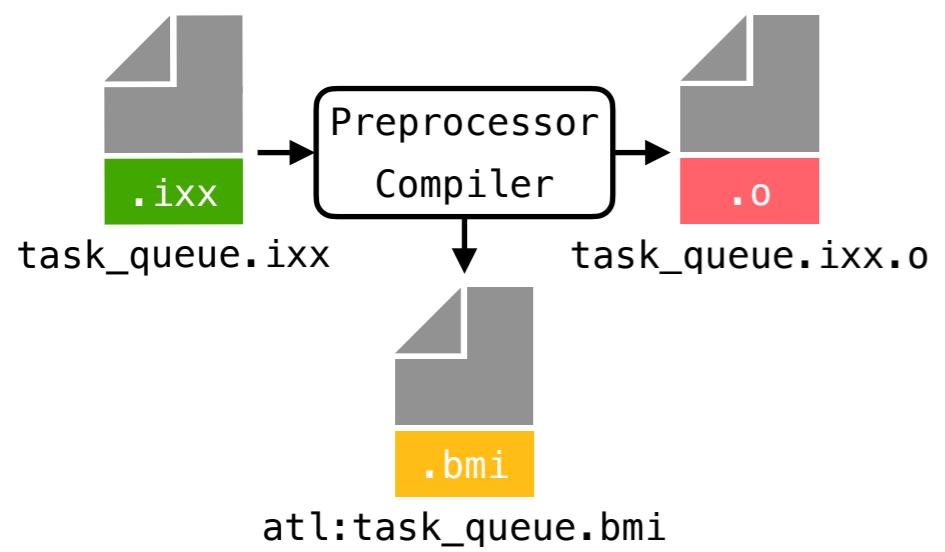


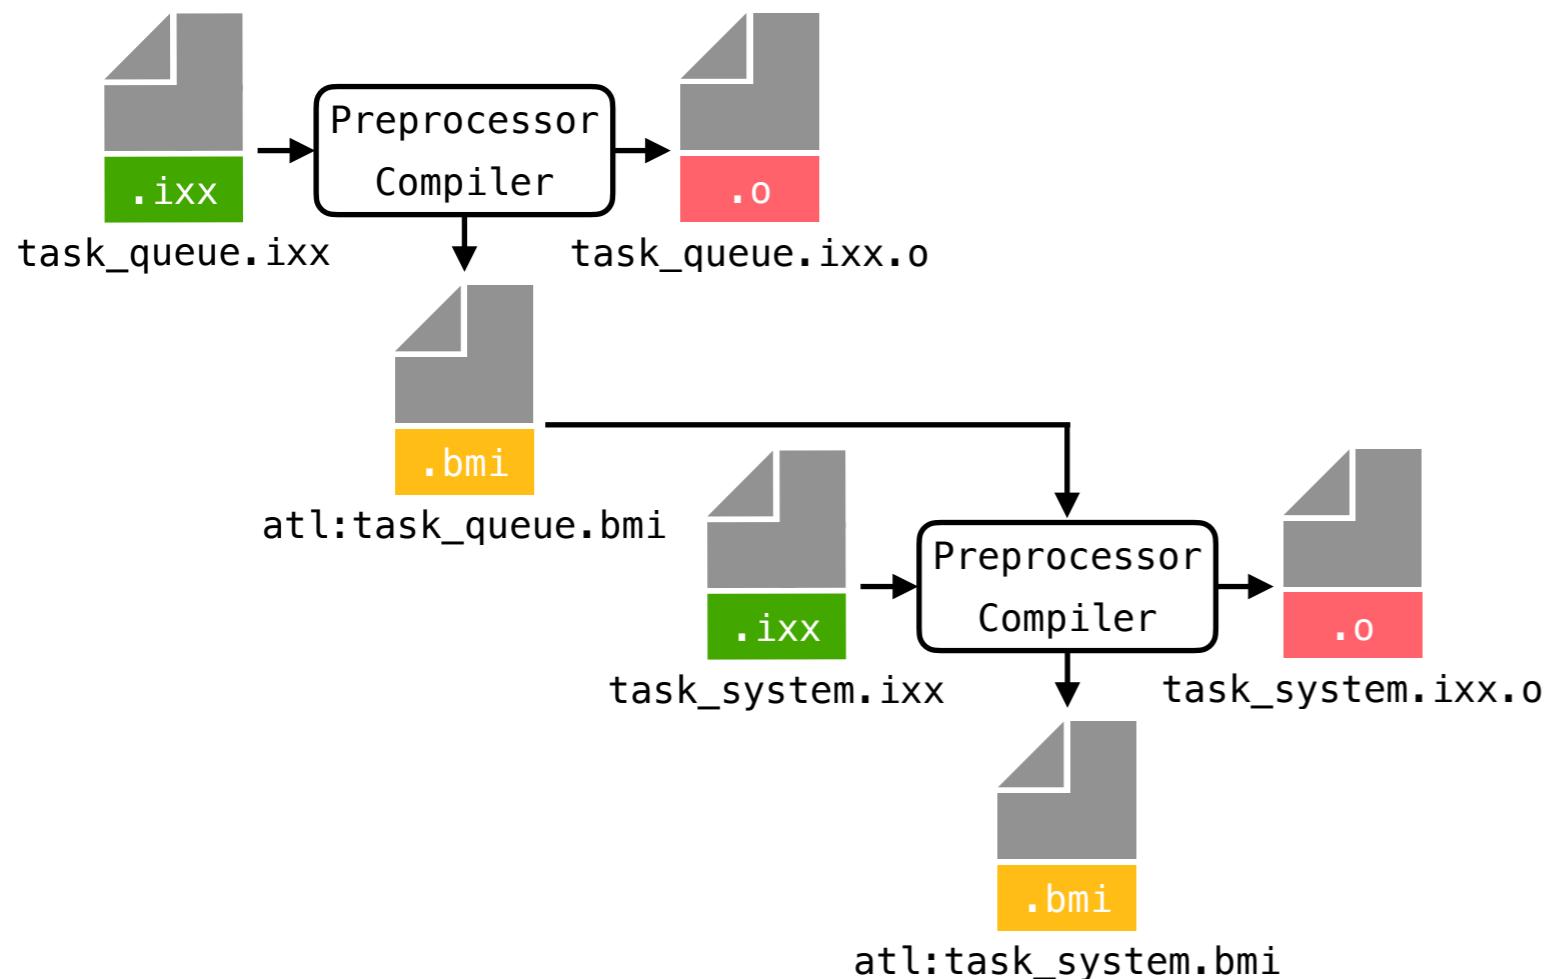
Modular Import Model

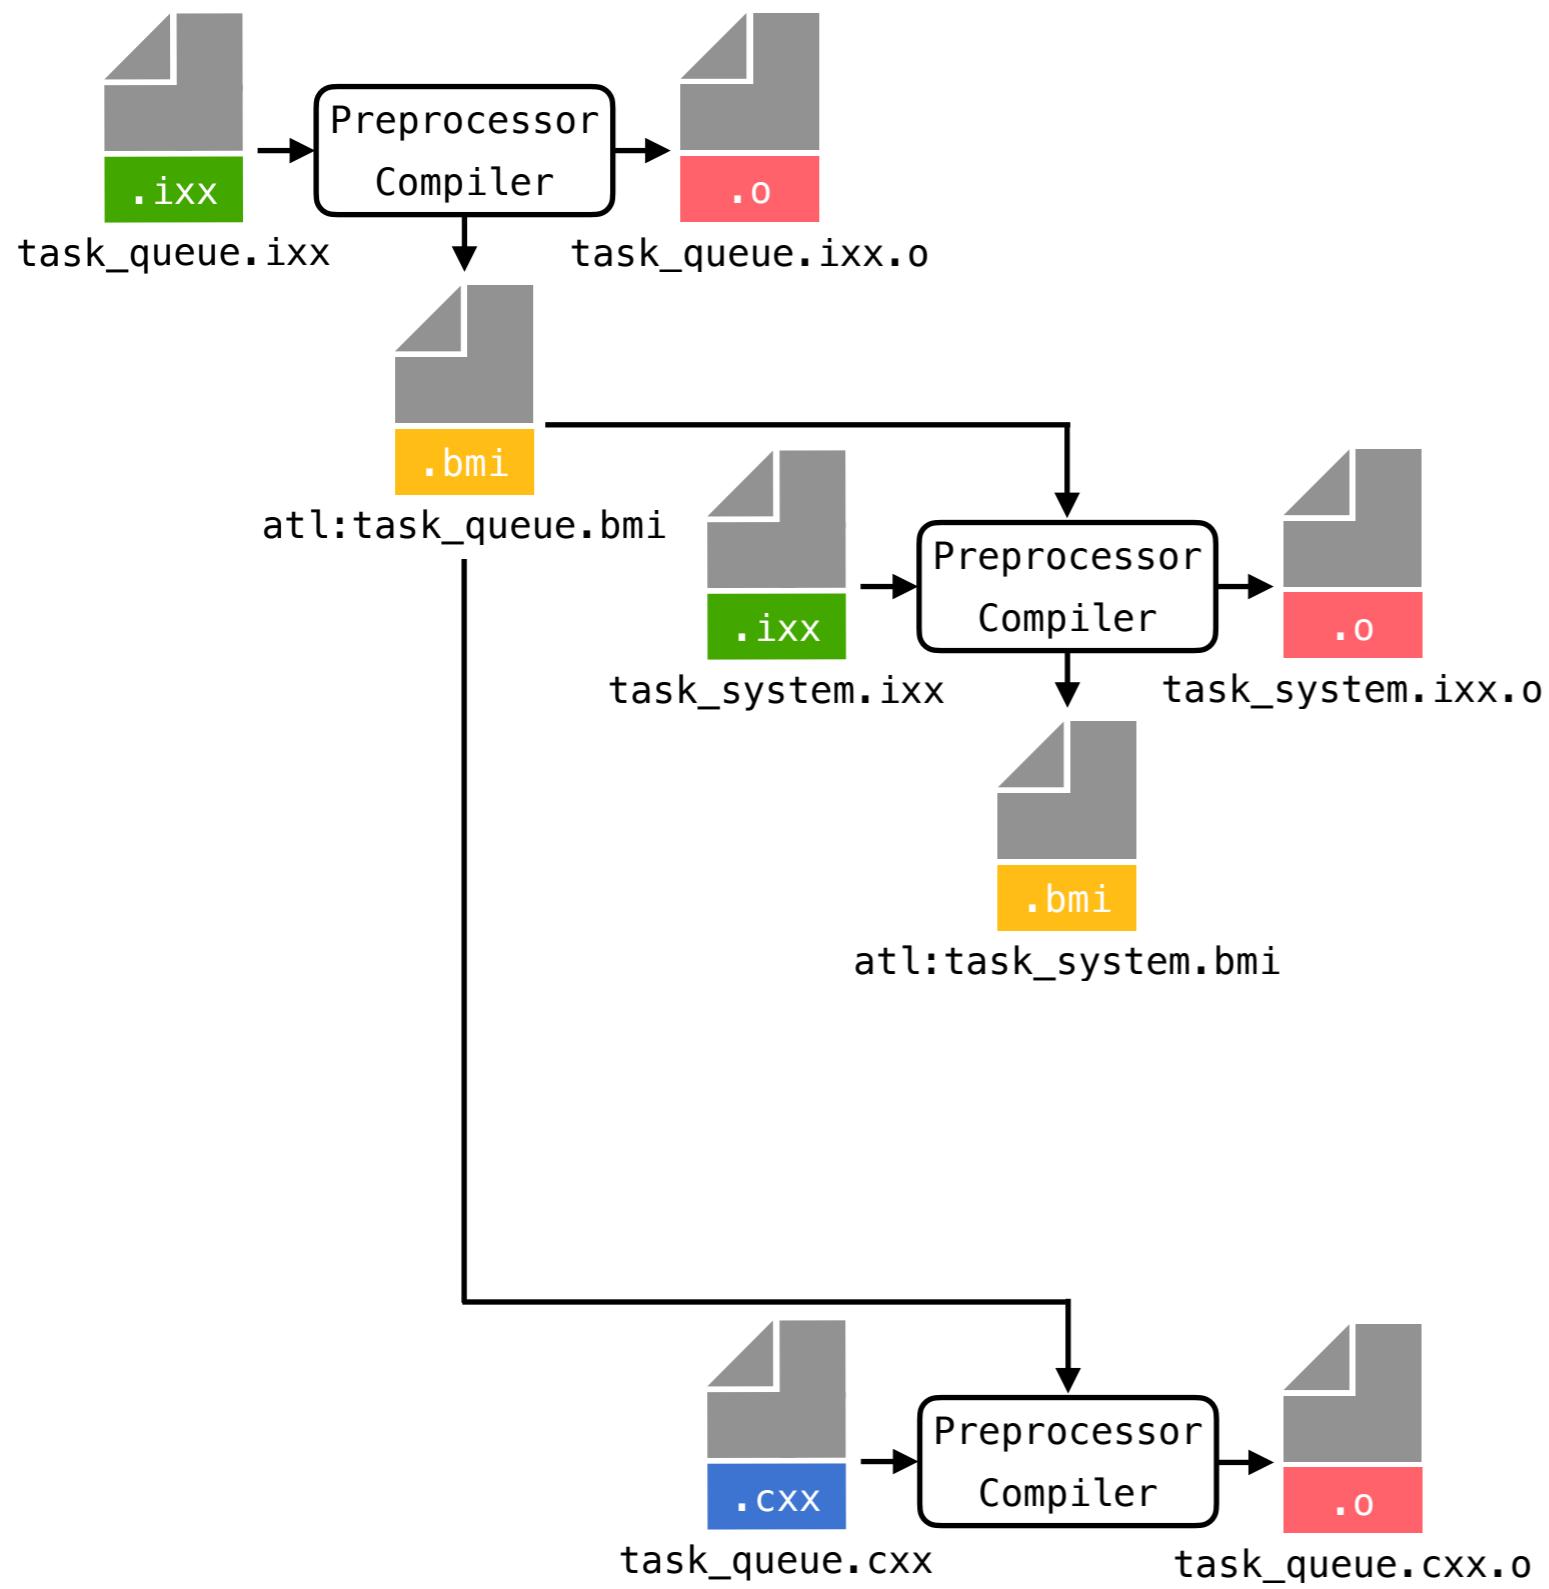
Dependency Graph

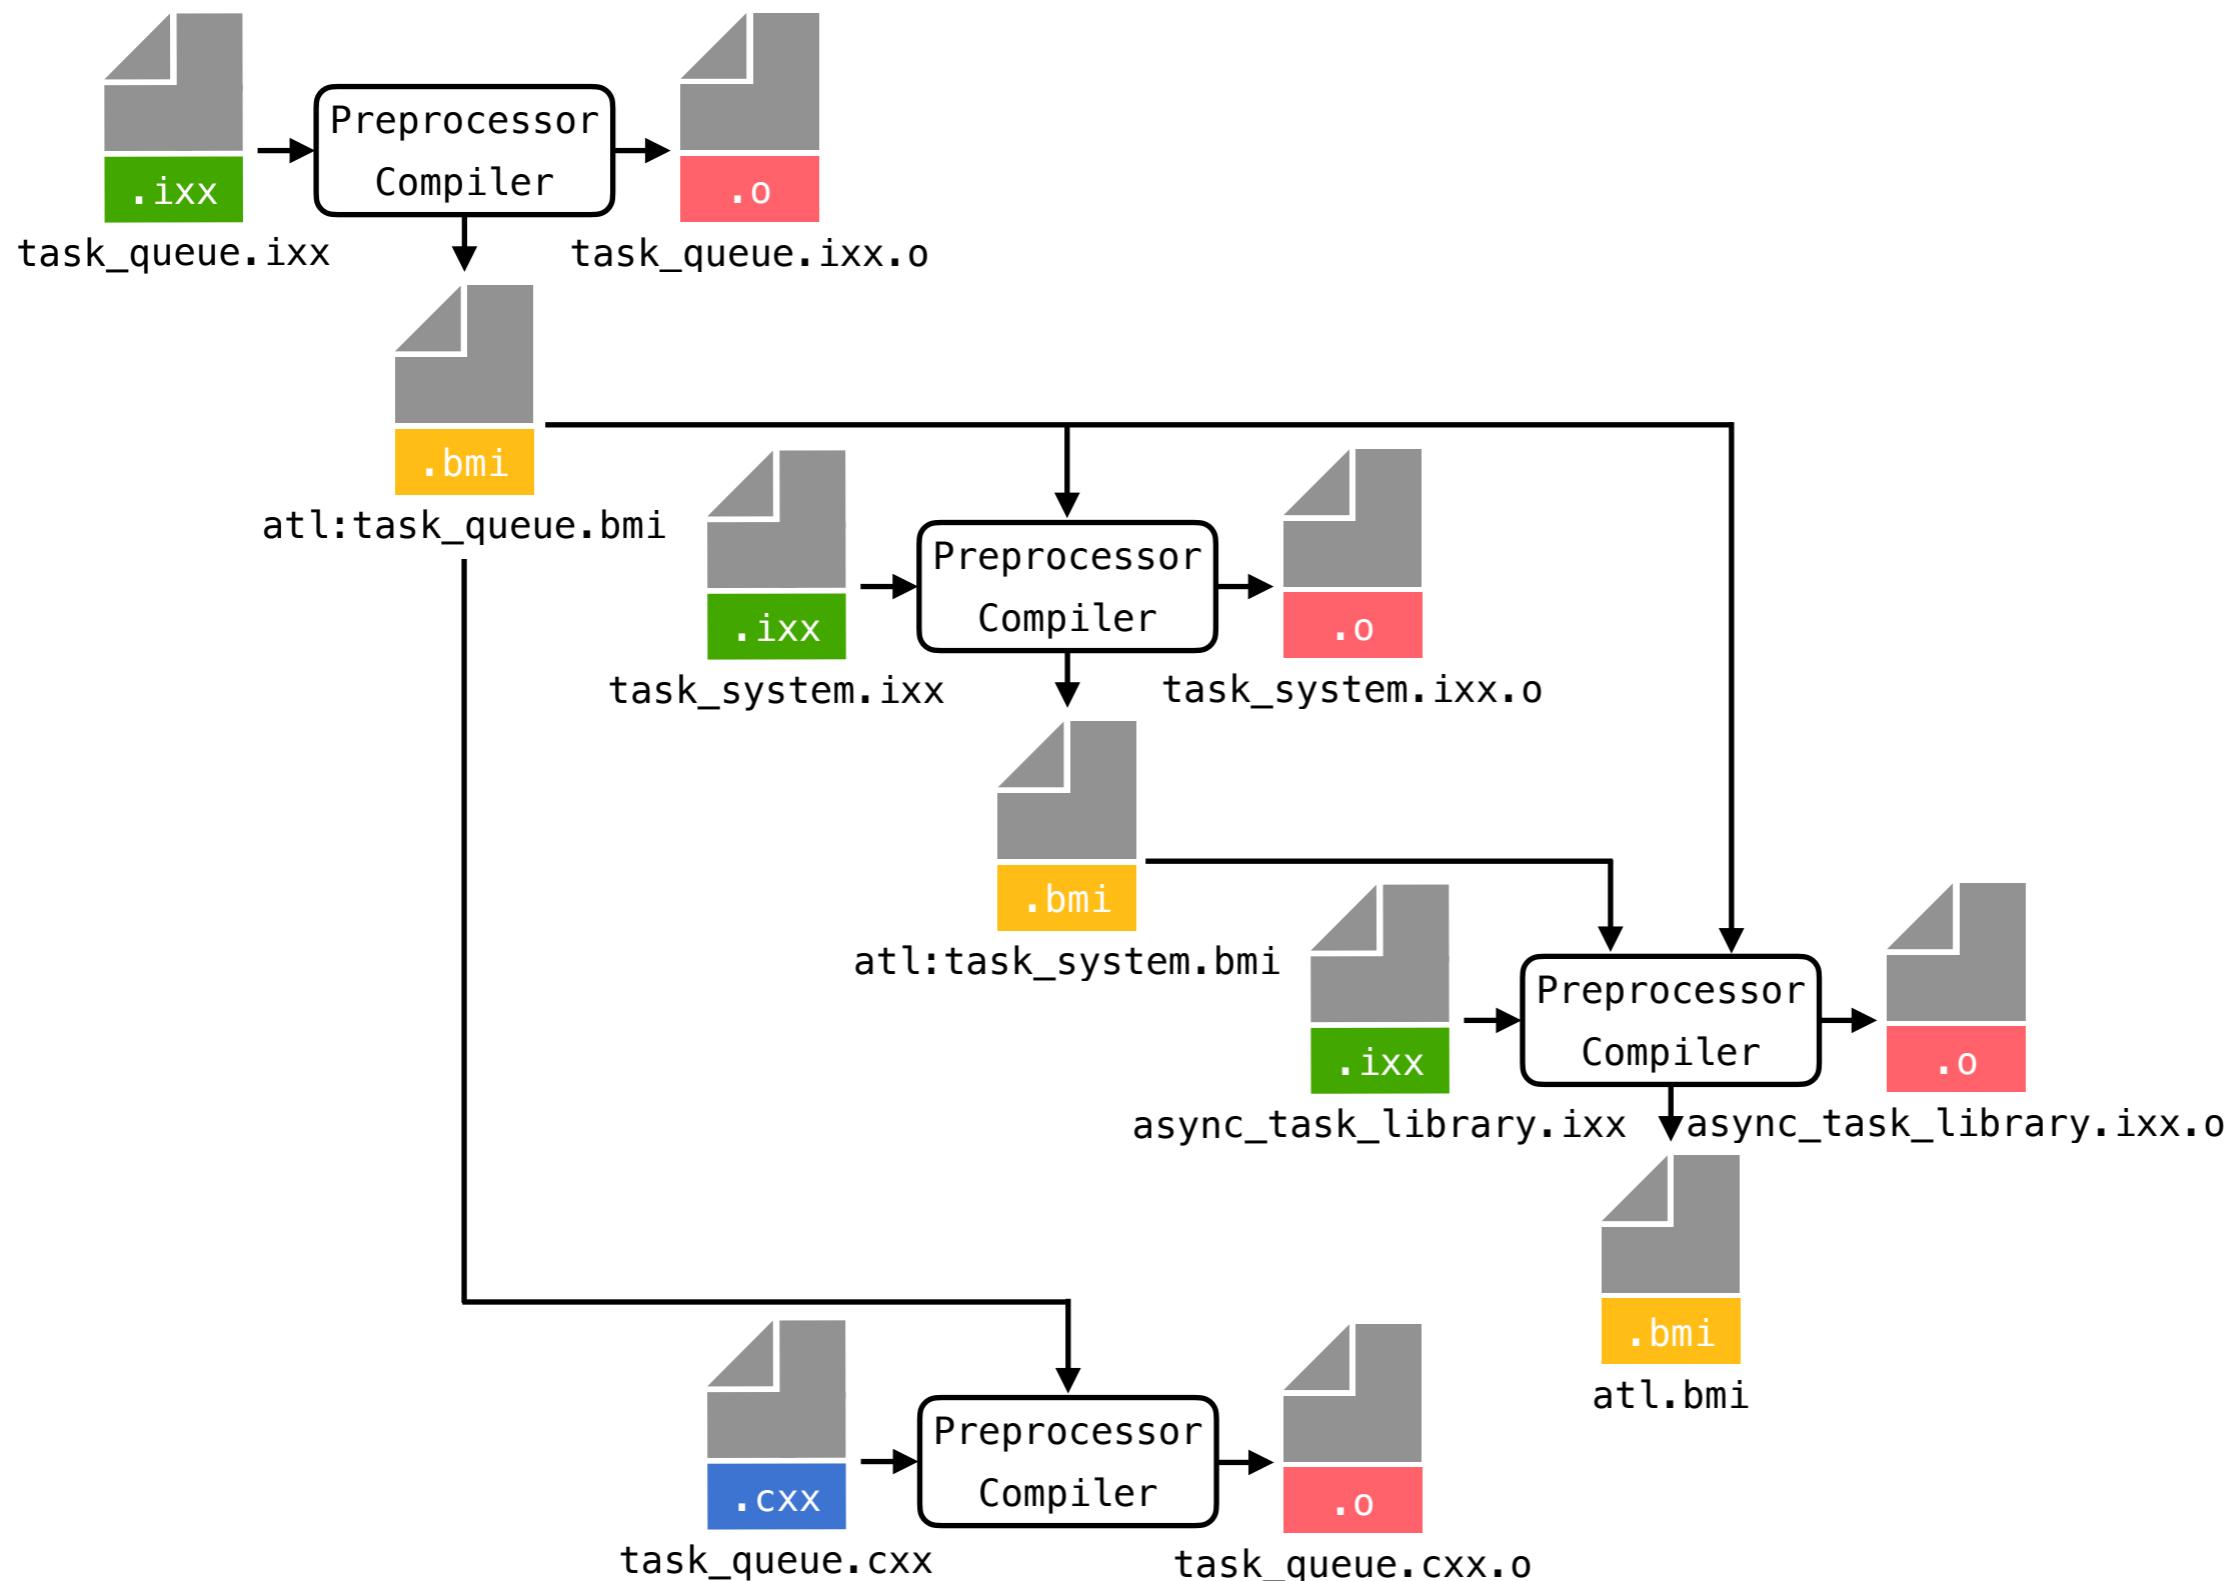
- Import statements form dependency graph between modules.
- Module dependencies cannot be cyclic, which C++ compilers enforce at compile-time.
- In short, modules dependency graphs are DAGs.

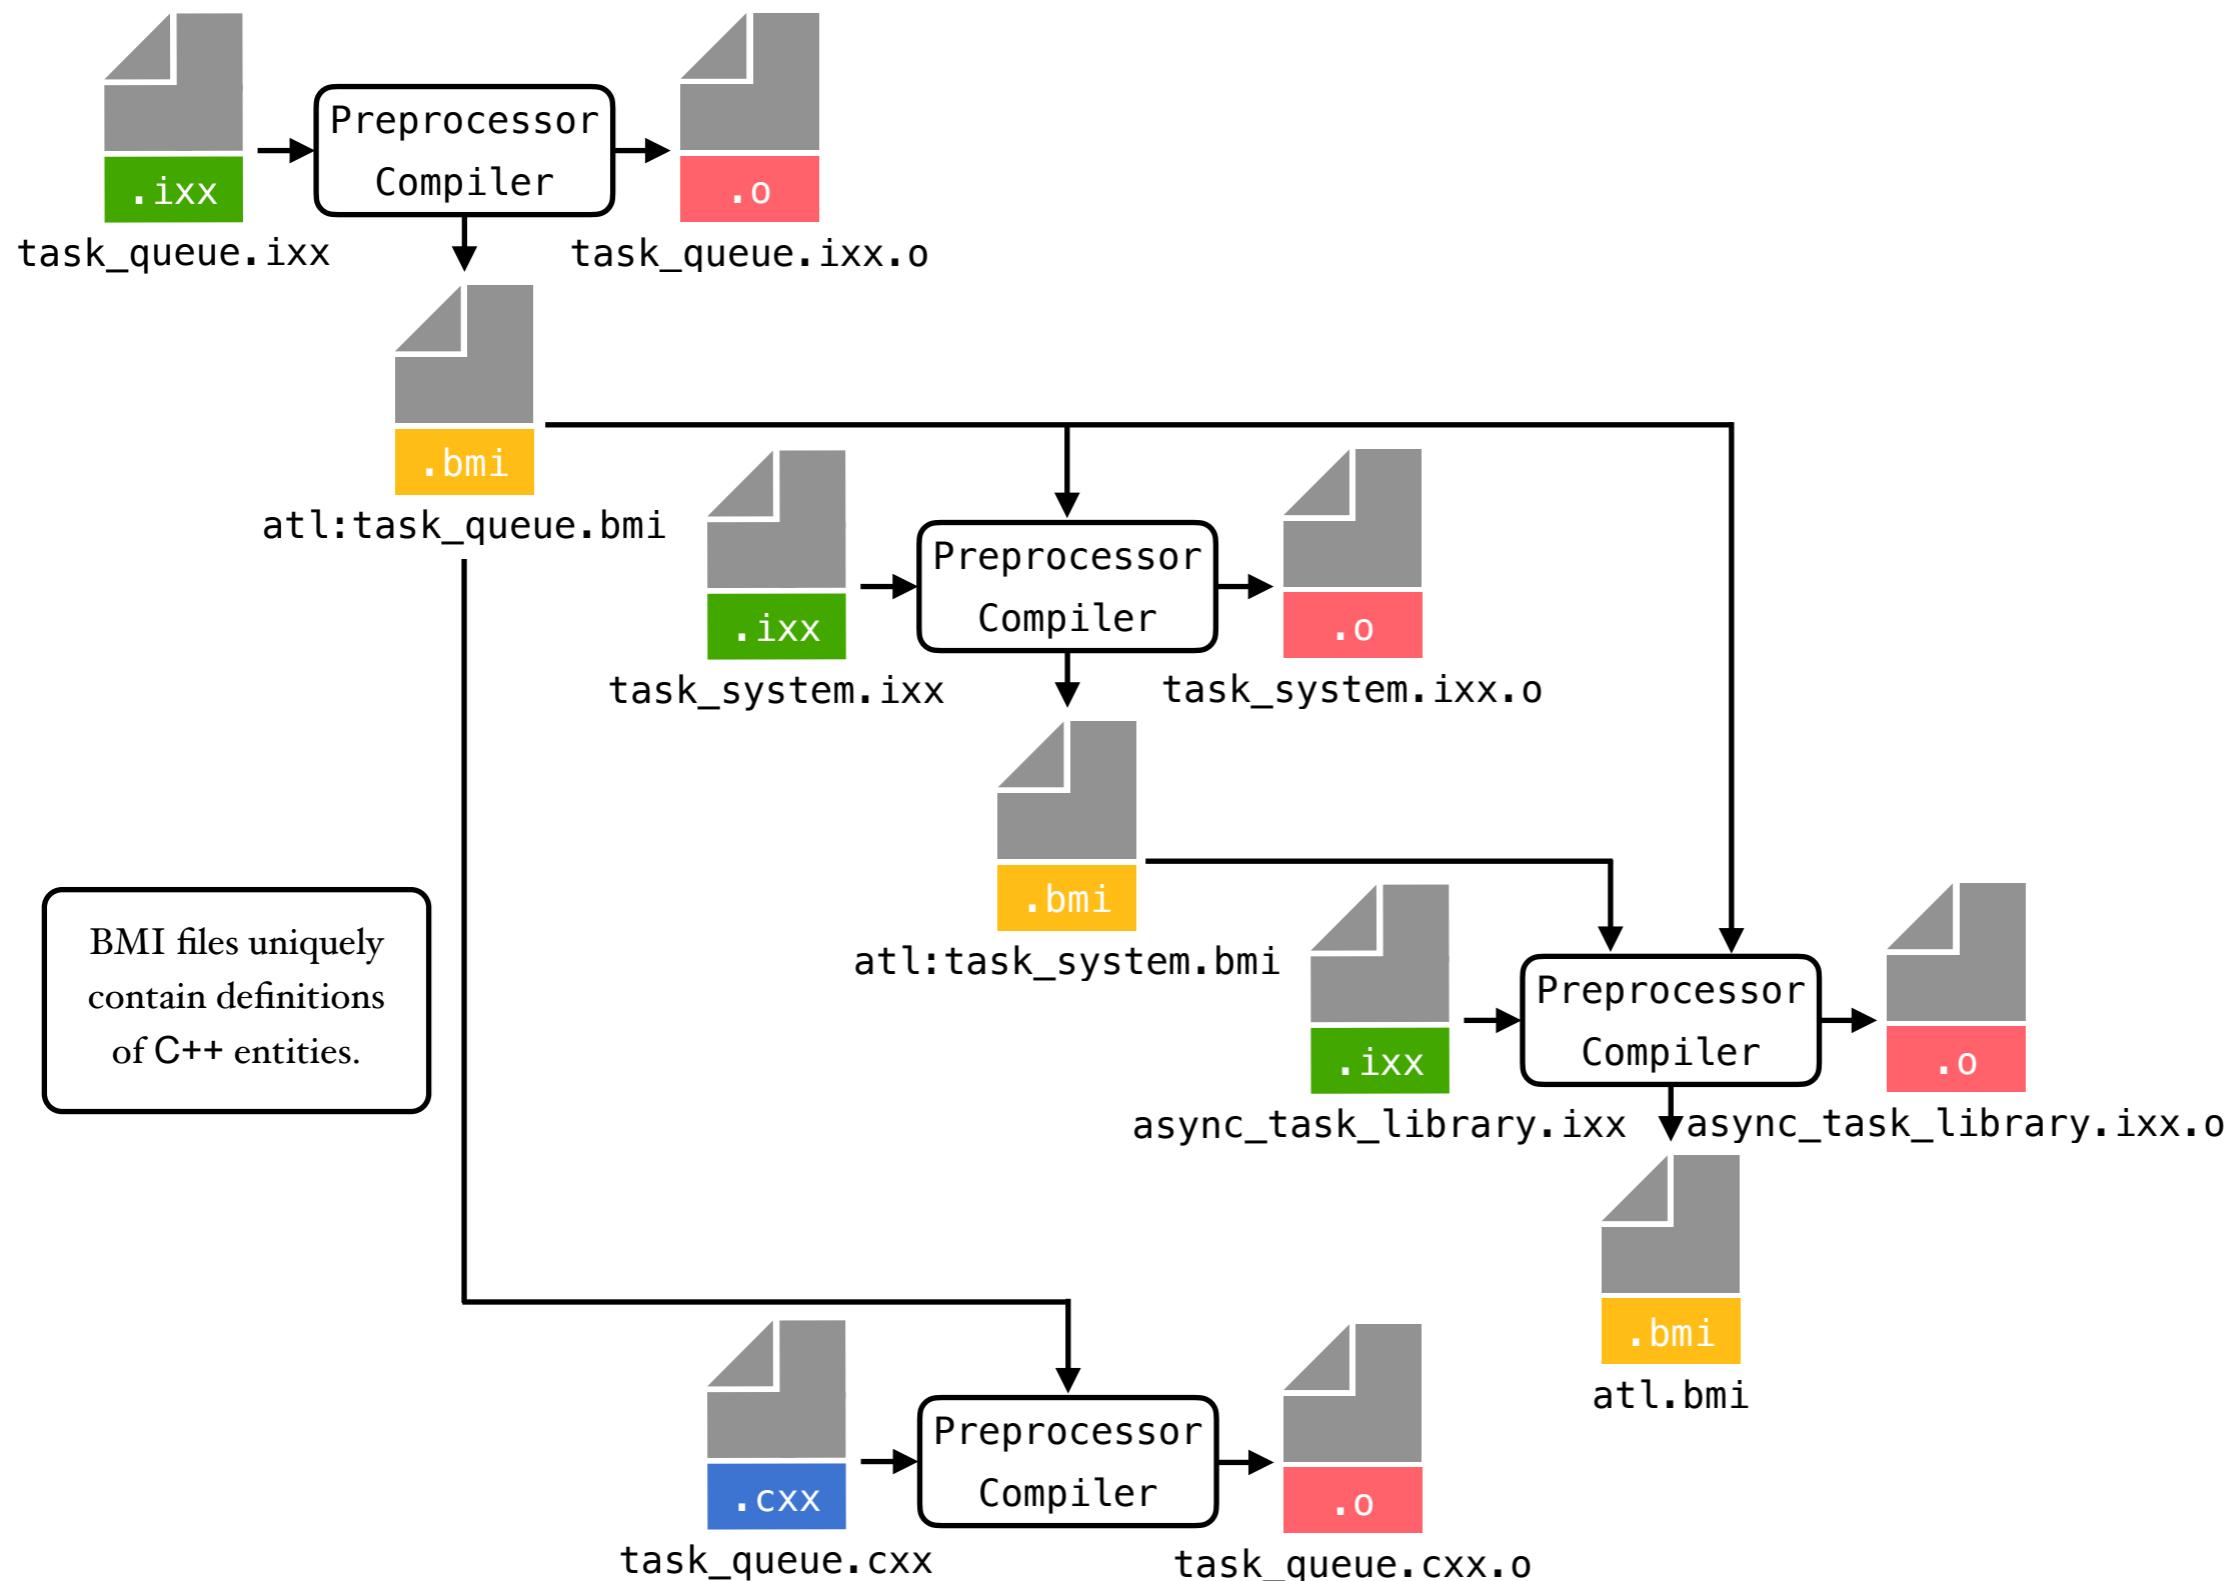


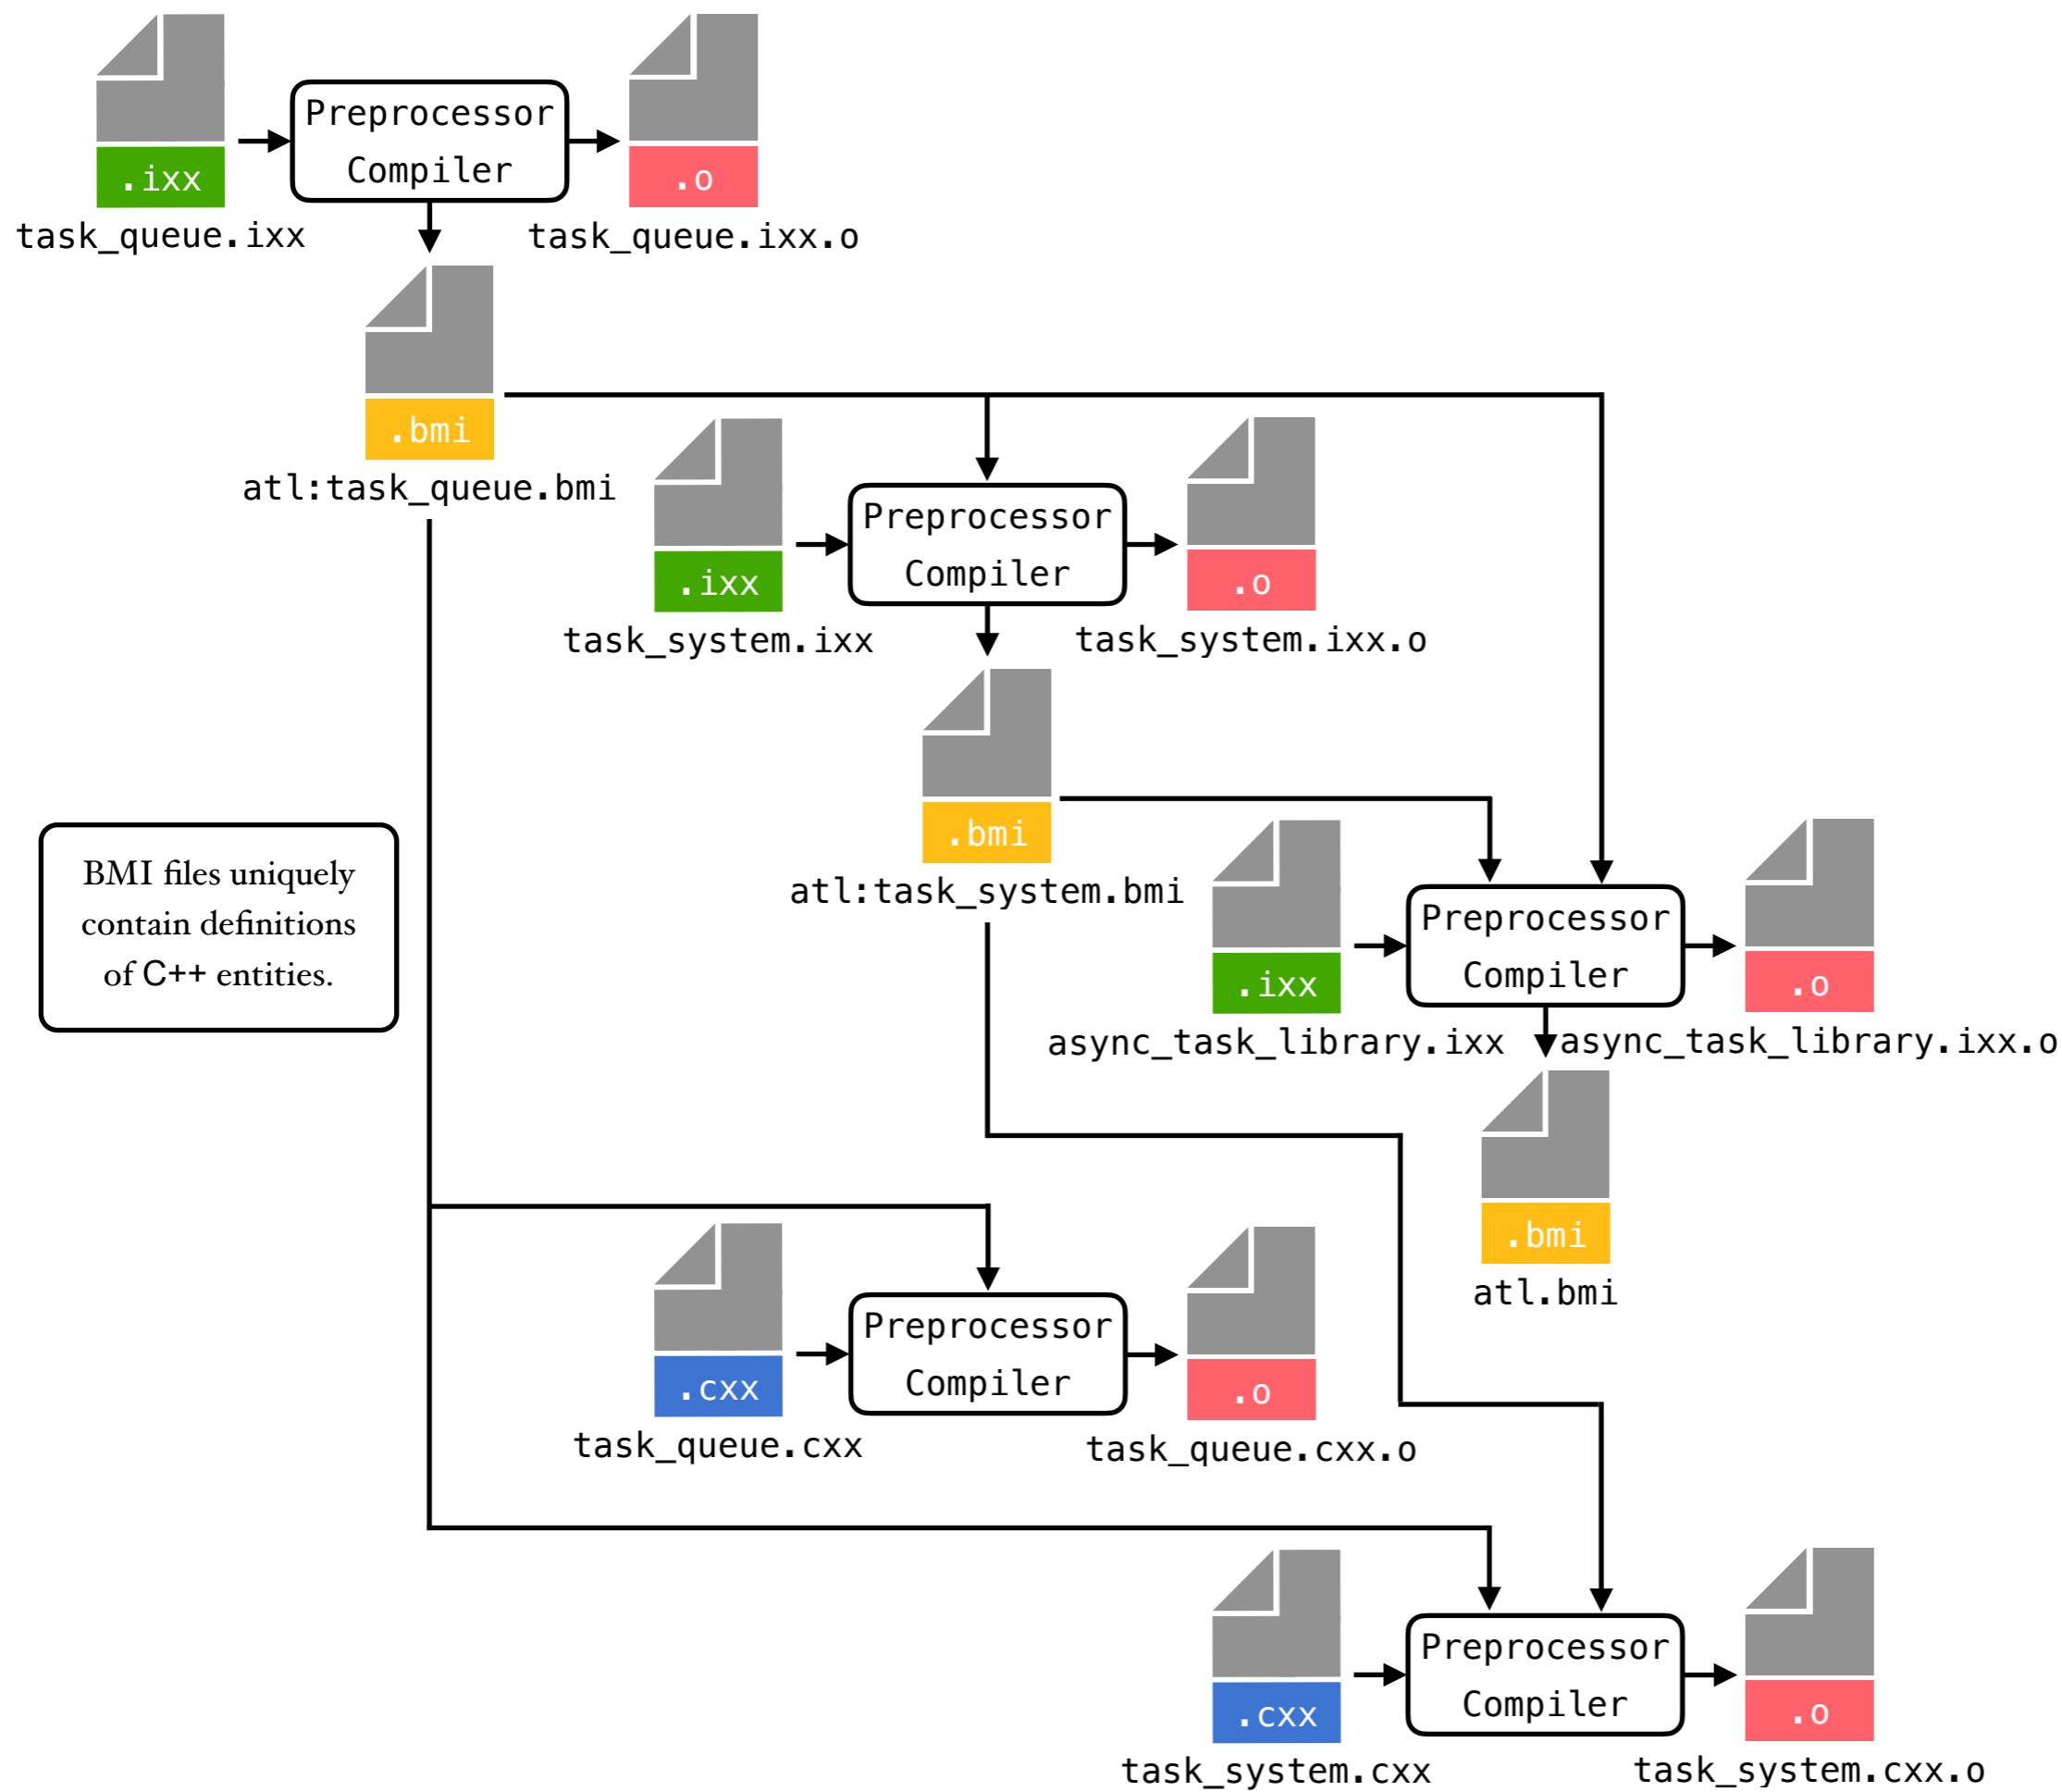


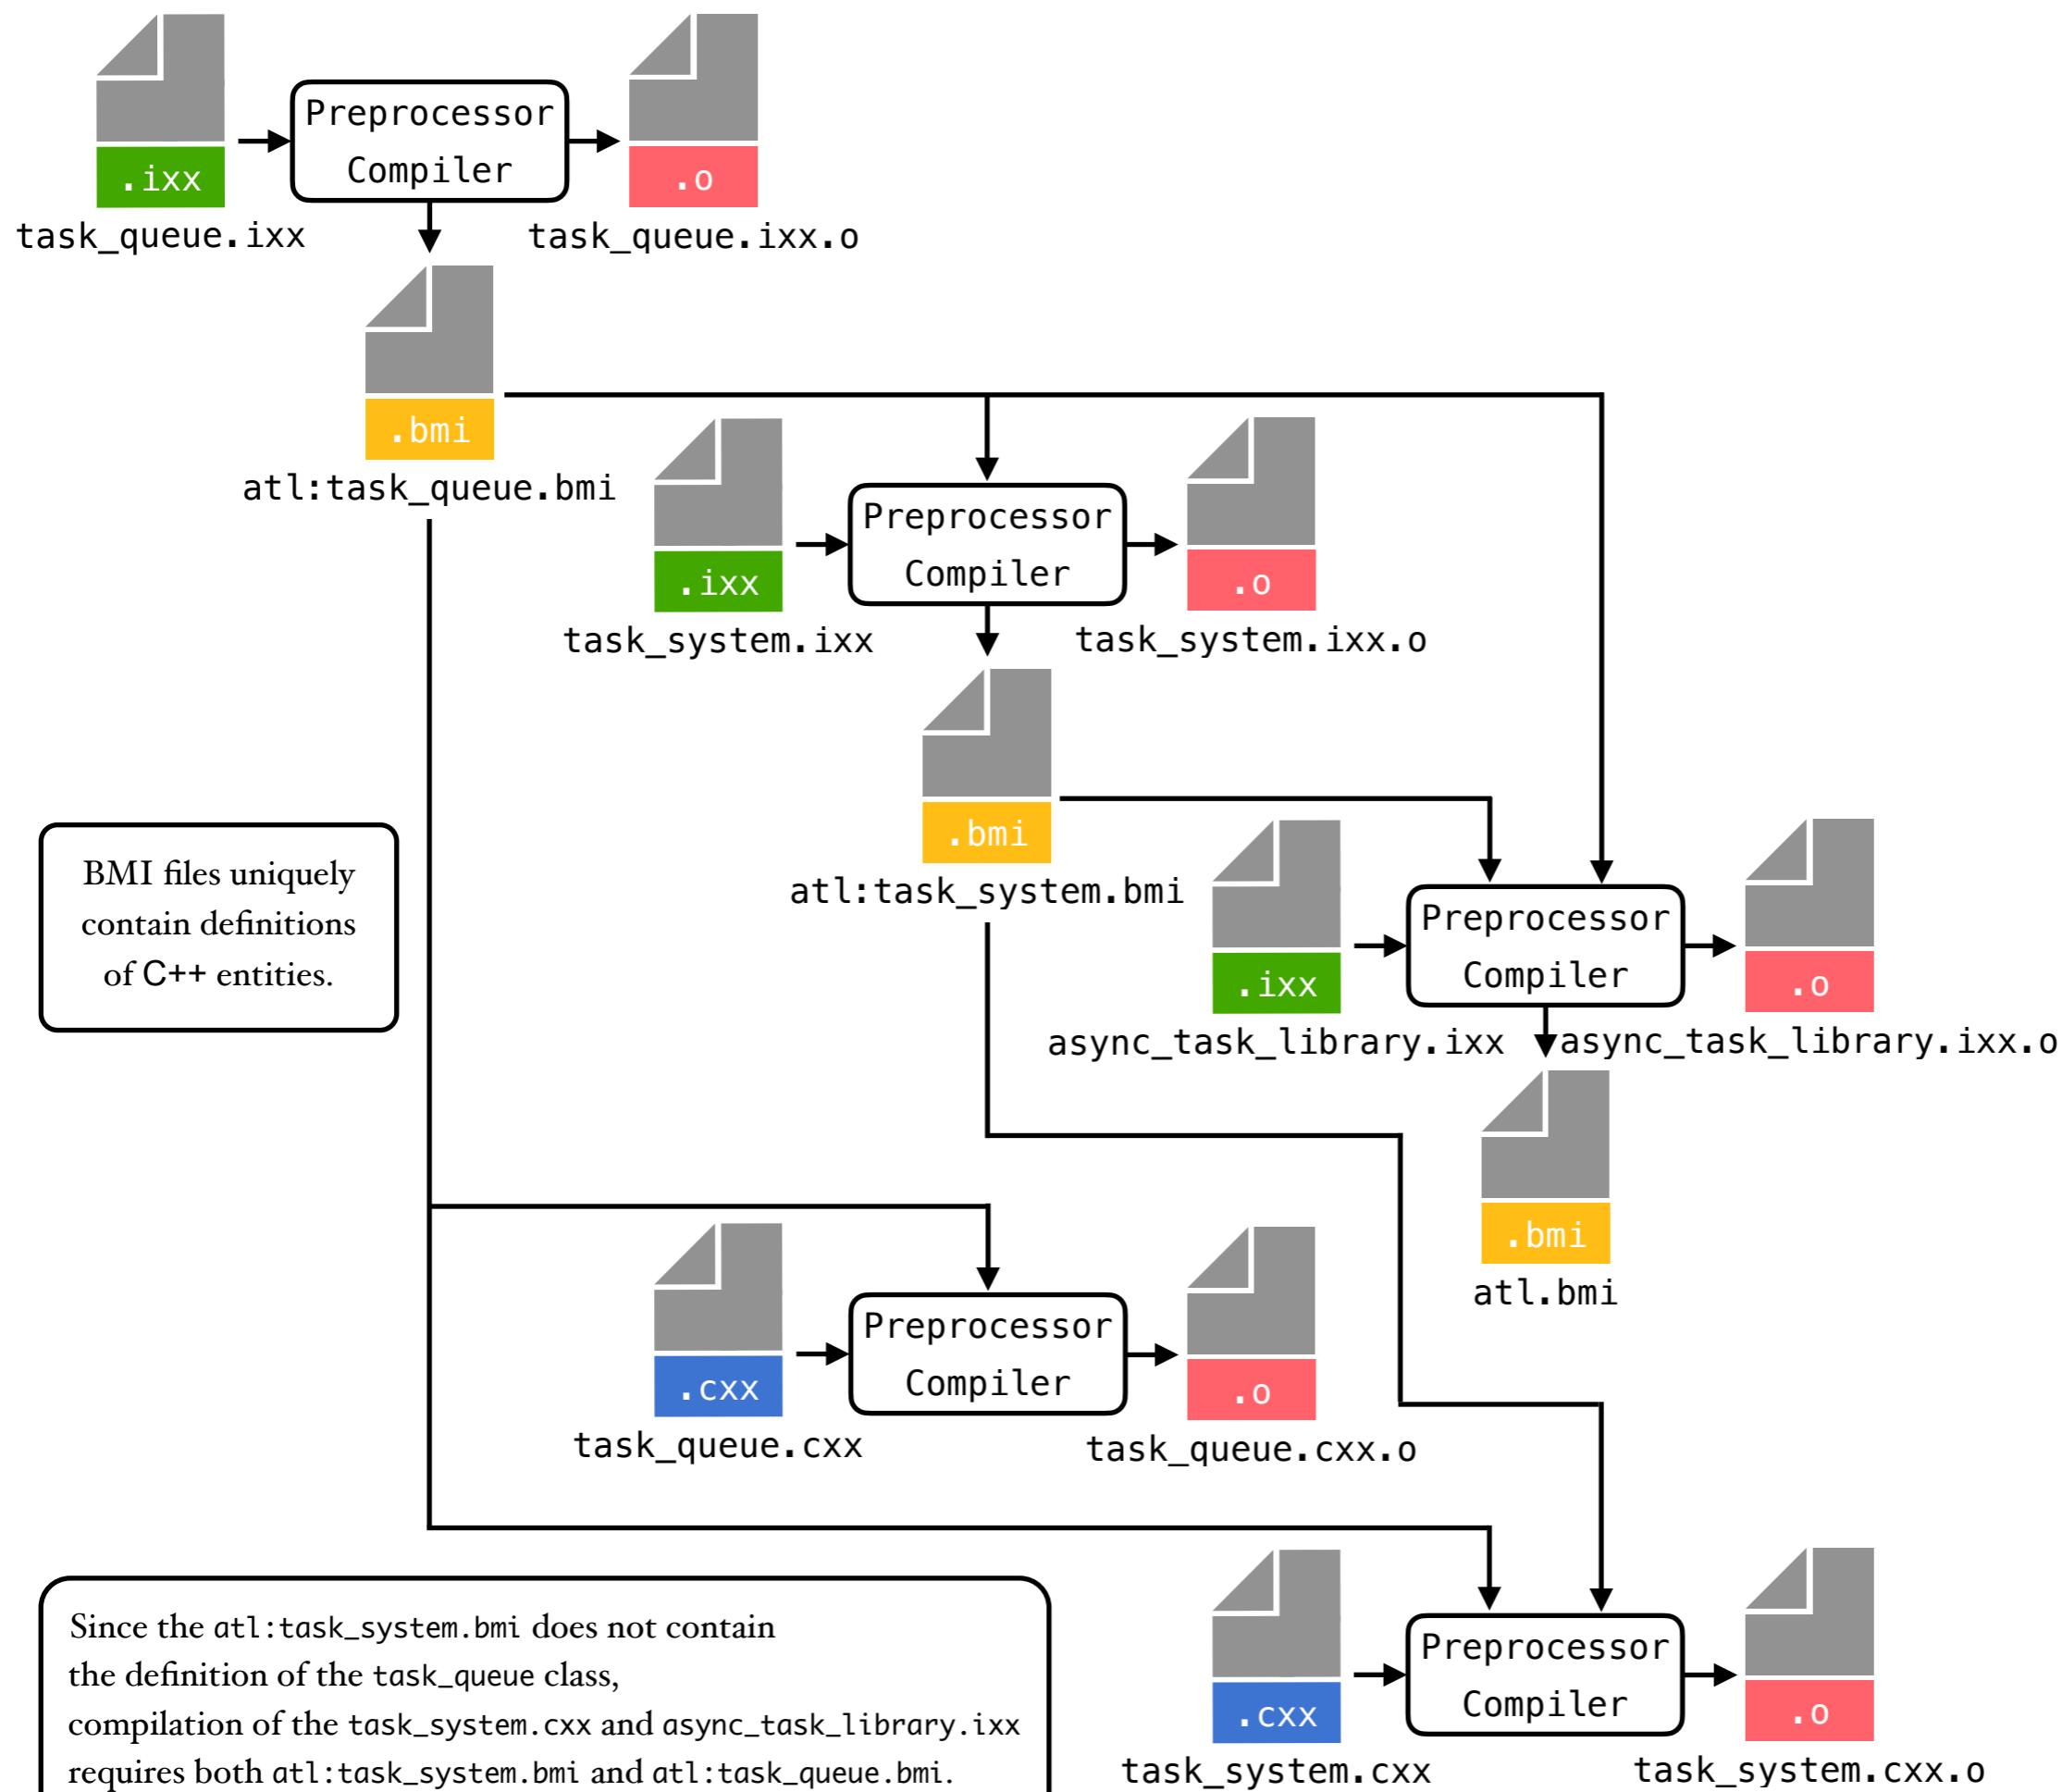


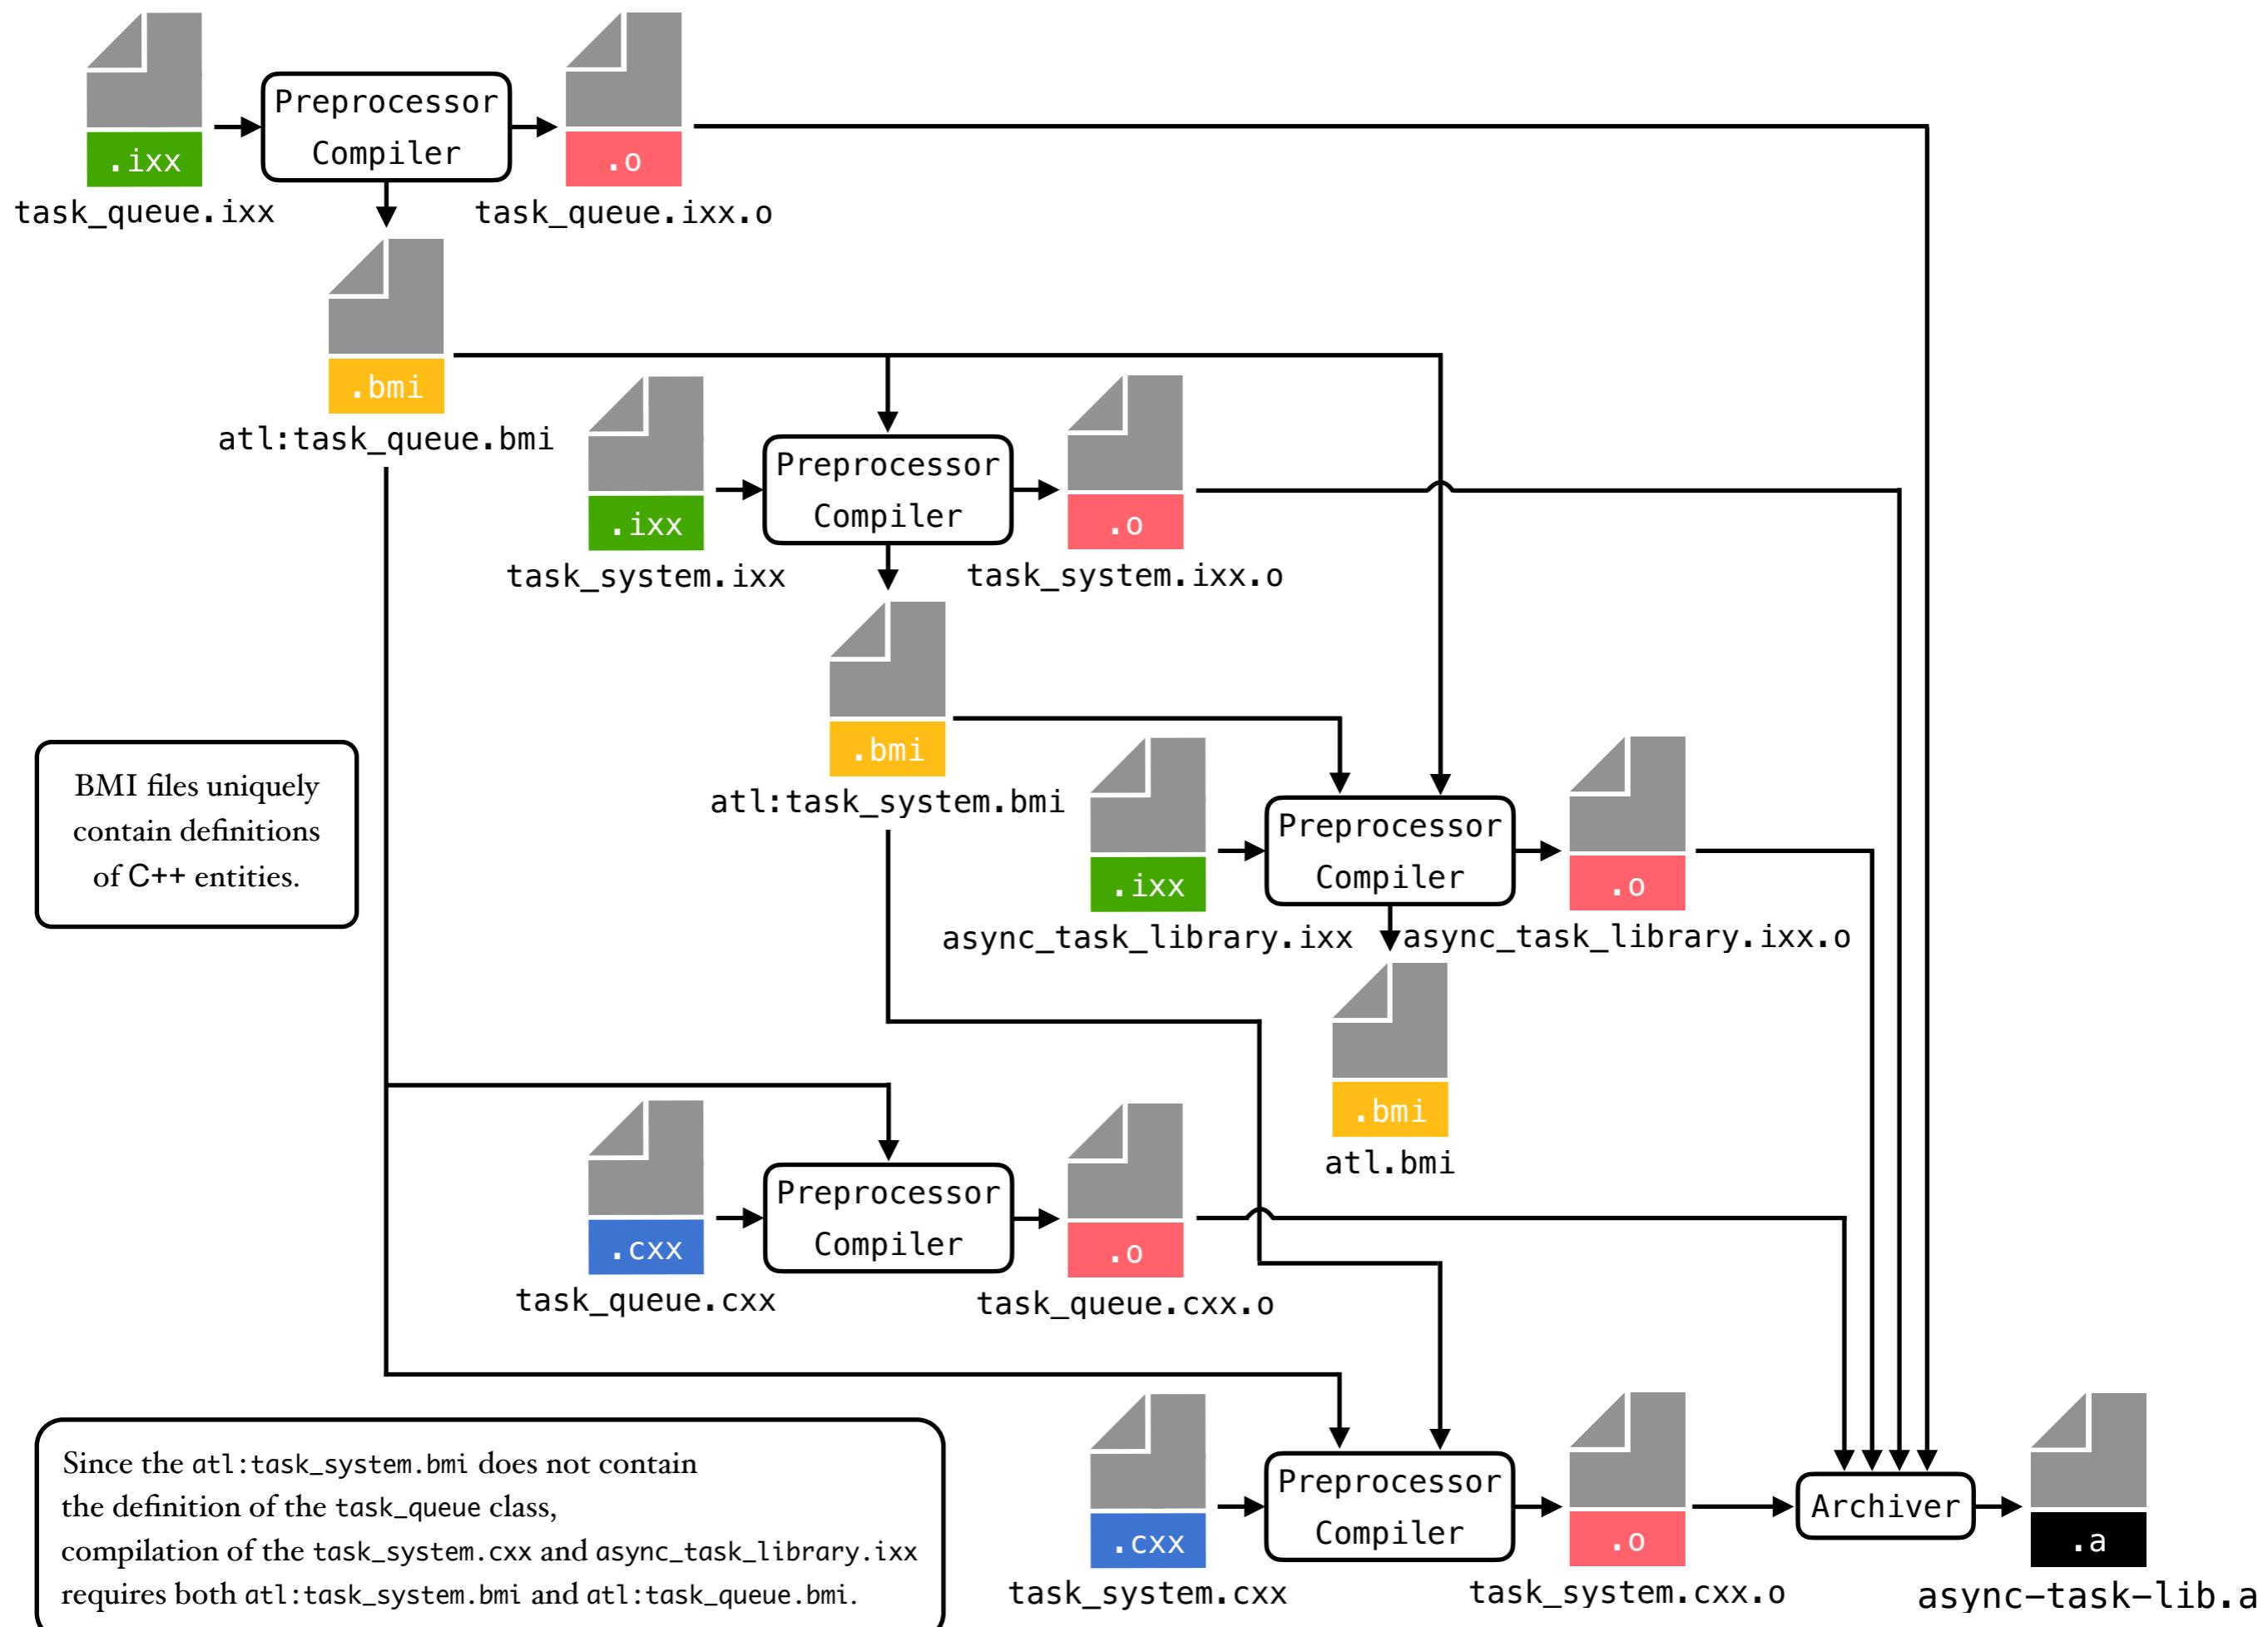






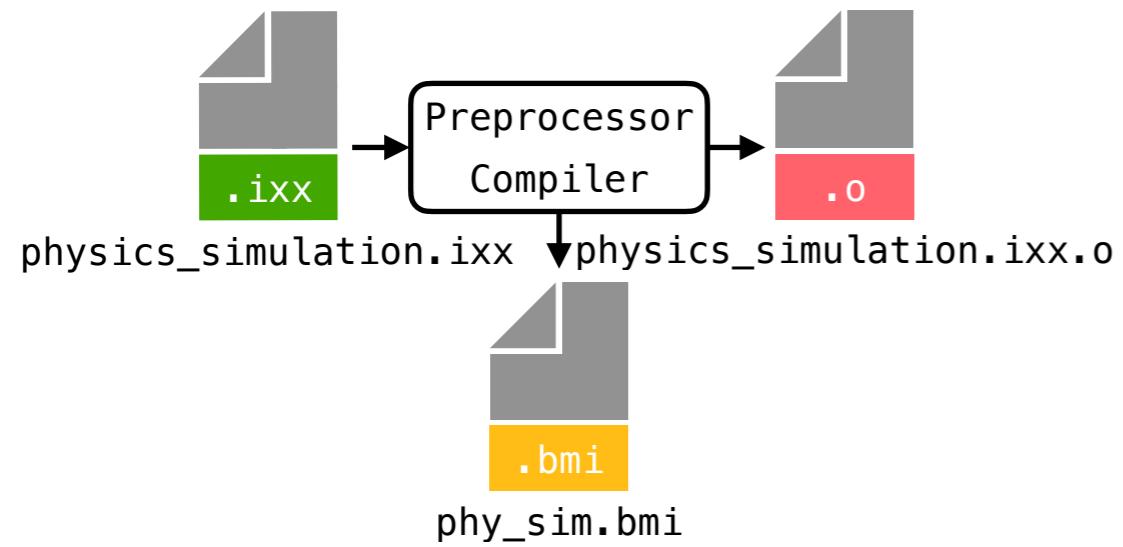






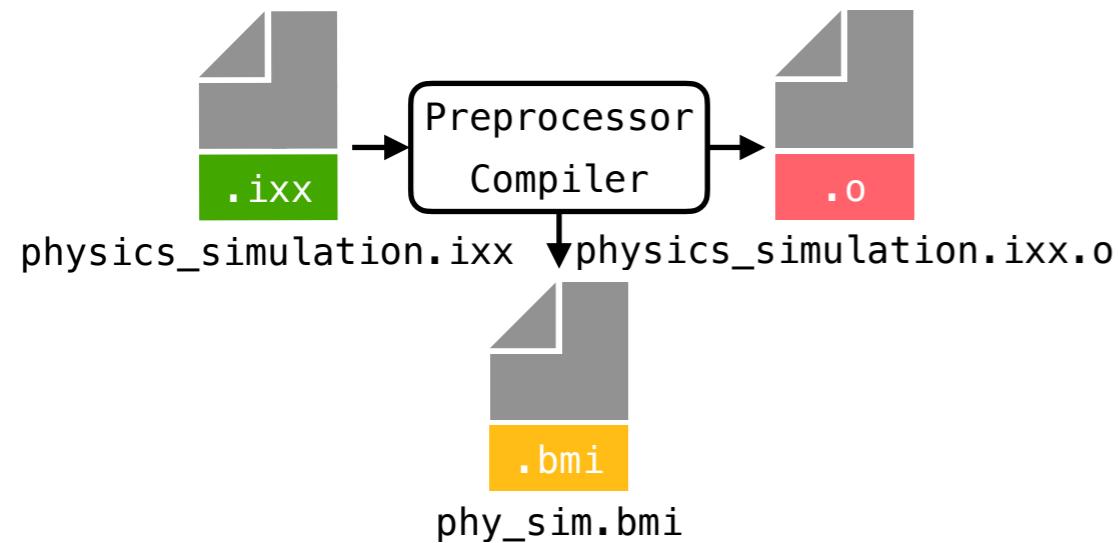
Modular Import Model

Build Workflow



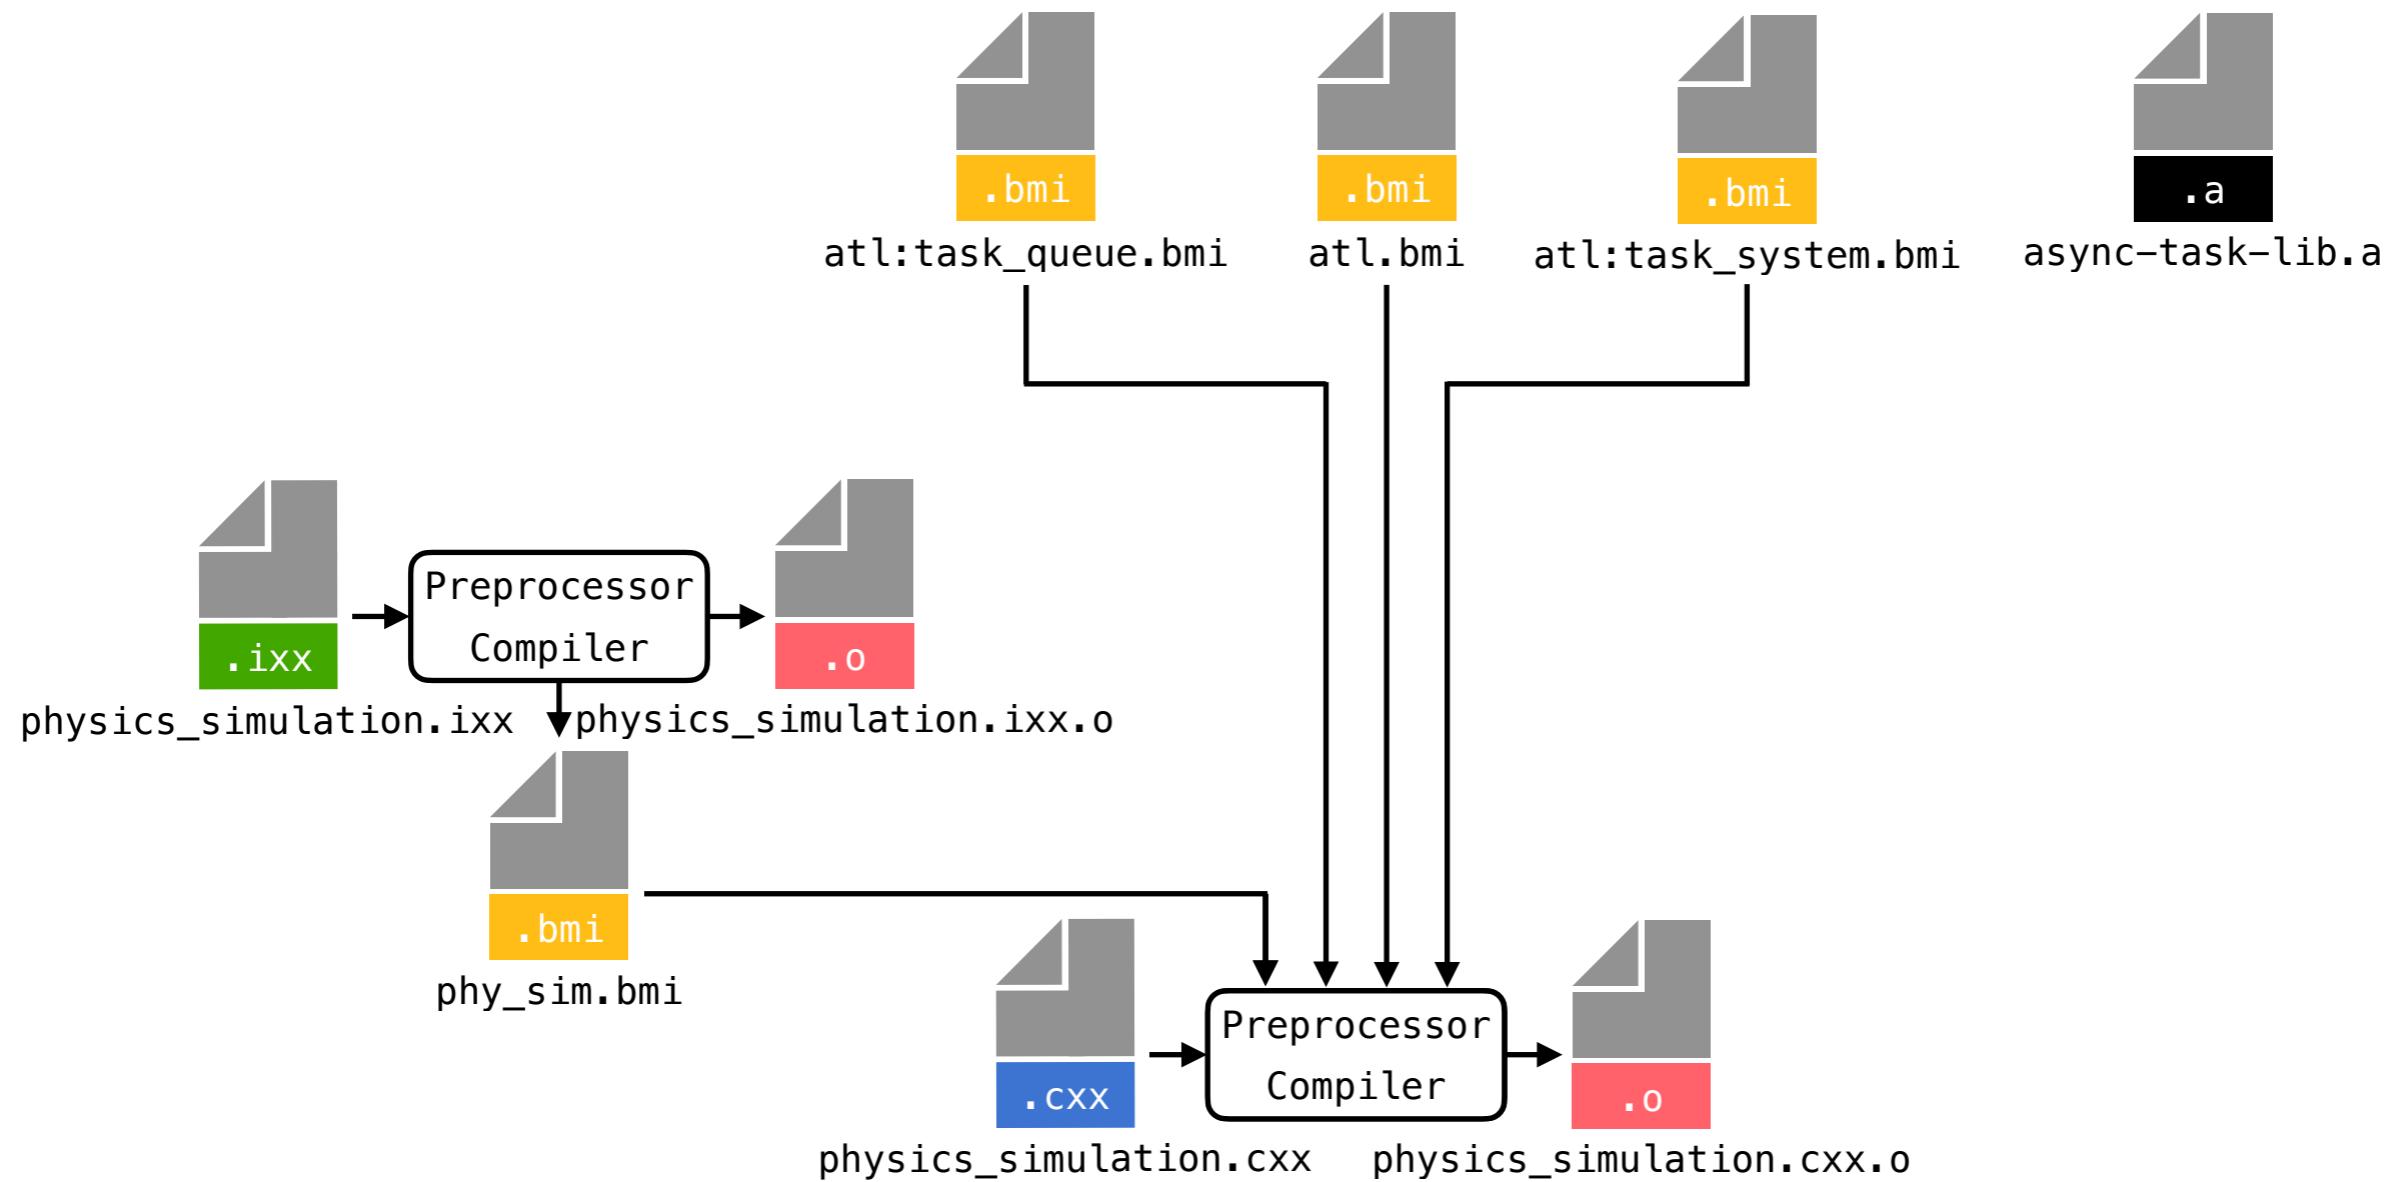
Modular Import Model

Build Workflow



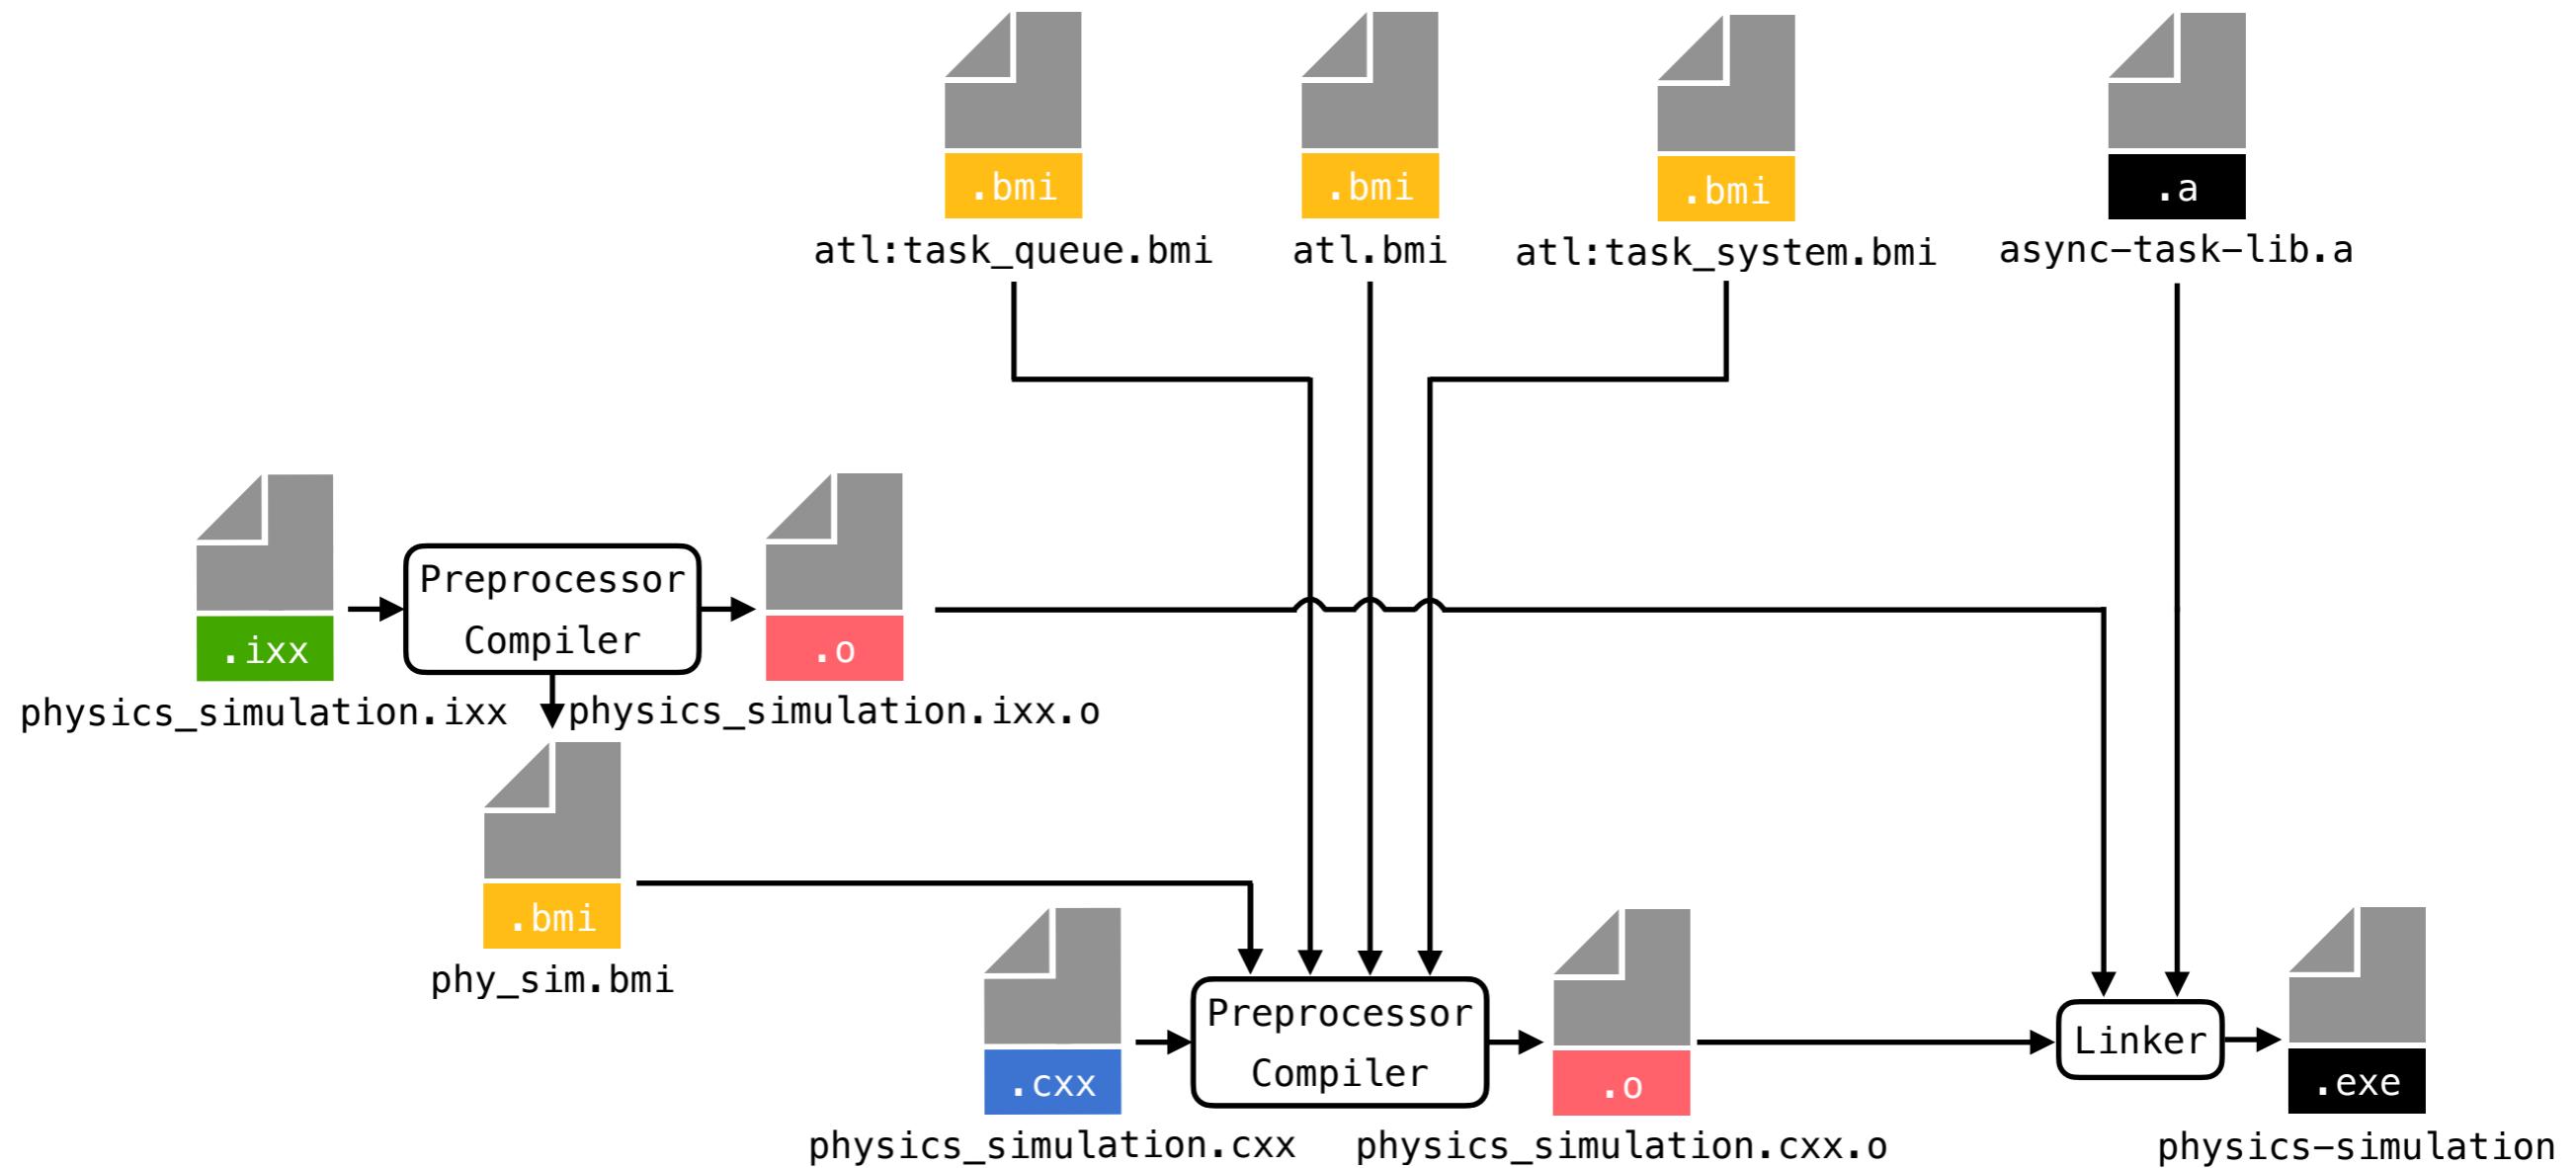
Modular Import Model

Build Workflow



Modular Import Model

Build Workflow



References

Conference Talks

- Core C++ 2019 :: Bryce Adelstein :: Modules are Coming
- Modules are coming - Bryce Adelstein Lelbach - Meeting Cpp 2019
- CppCon 2019: Gabriel Dos Reis “Programming with C++ Modules: Guide for the Working”
- CPPP 2019 - C++ Modules: What You Should Know - Gabriel Dos Reis
- CppCon 2015: Gabriel Dos Reis “Large Scale C++ with Modules: What You Should Know”
- CppCon 2016: Gabriel Dos Reis “C++ Modules: The State of The Union”
- 2019 LLVM Developers’ Meeting: E. Christopher & J. Doerfert “Introduction to LLVM”
- 2019 EuroLLVM Developers’ Meeting: V. Bridgers & F. Piovezan “LLVM IR Tutorial - Phis, GEPs ...”
- 2019 LLVM Developers’ Meeting: S. Haastregt & A. Stulova “An overview of Clang”

References

Documentation

- <https://llvm.org/docs/BitCodeFormat.html>
- <https://llvm.org/docs/CommandGuide/index.html>
- <https://llvm.org/docs/LangRef.html>
- <http://llvm.org/docs/GetElementPtr.html>
- <https://clang.llvm.org/docs/InternalsManual.html>
- <https://clang.llvm.org/docs/LibASTImporter.html>
- <https://clang.llvm.org/docs/PCHInternals.html>
- <https://clang.llvm.org/docs/IntroductionToTheClangAST.html>
- <https://devblogs.microsoft.com/cppblog/c-modules-in-vs-2015-update-1/>
- <https://devblogs.microsoft.com/cppblog/cpp-modules-in-visual-studio-2017/>
- <https://devblogs.microsoft.com/cppblog/c-modules-conformance-improvements-with-msvc-in-visual-studio-2019-16-5/>